

REFLECTOR: Internalizing Step-wise Reflection against Indirect Jailbreaks

Jiachen Ma^{1,2} Jiawen Zhang³ Xiangtian Li² Bo Zou² Chaochao Lu² Chao Yang²

Abstract

While Large Language Models (LLMs) demonstrate remarkable capabilities, they remain susceptible to sophisticated, multi-step jailbreak attacks that circumvent conventional surface-level safety alignment by exploiting the internal generation process. To address these vulnerabilities, we propose **REFLECTOR**, a principled two-stage framework that internalizes self-reflection within the generation trajectory. REFLECTOR first leverages teacher-guided generation to produce high-quality reflection data for supervised fine-tuning (SFT), establishing structured reflection patterns. It subsequently uses Reinforcement Learning (RL) with outcome-driven and reward-validity supervision to instill robust, autonomous self-reflection capabilities. Empirical results show that REFLECTOR achieves Defense Success Rates (DSR) exceeding 90% against complex indirect attacks while generalizing robustly across diverse threat scenarios. Notably, the framework enhances both task-specific and general utility, yielding a 5.85% gain on GSM8K alongside improved performance on knowledge-intensive benchmarks. By internalizing trajectory-level safety, REFLECTOR overcomes the fundamental limitations of surface alignment without significant computational overhead, offering an efficient and scalable solution for the development of safe and capable LLMs. Our code and data are available at <https://github.com/mjc-ma-01/self-reflection-llm.git>.

1. Introduction

Large Language Models (LLMs) have achieved remarkable breakthroughs in a wide range of tasks (Qwen et al.,

¹Fudan University ²Shanghai Artificial Intelligence Laboratory ³Zhejiang University. Correspondence to: Chao Yang <yangchao@pjlab.org.cn>.




Figure 1. Correlation between the position of the first occurrence of harmful tokens and the attack success rate (ASR). While direct jailbreaks (blue) manifest immediately, indirect attacks (red) exhibit a stealthy latency, with harmful content emerging only after 20 tokens. This delay enables malicious intent to bypass surface-level safety alignment, leading to significantly higher ASRs than direct attacks.

2025; Achiam et al., 2023; Grattafiori et al., 2024), demonstrating human-level proficiency in mathematical reasoning (Cobbe et al., 2021b; Hendrycks et al., 2021), coding (Chen, 2021; Nam et al., 2024), and general problem-solving (Wei et al., 2022). However, as their deployment expands into safety-critical domains such as autonomous medical diagnostics (Ullah et al., 2024) and personalized education (Zhang et al., 2025b), mitigating the generation of harmful, biased, or deceptive content has emerged as a paramount challenge (Touvron et al., 2023; Ma et al., 2025).

Safety alignment has become a cornerstone of modern language models, grounding behavior in human values (Dong et al., 2024; Liu et al., 2023b; Wang et al., 2023a; Zhang et al., 2026). Approaches such as Reinforcement Learning from Human Feedback (RLHF) (Ouyang et al., 2022; Dai et al., 2023) and Direct Preference Optimization (DPO) (Rafailov et al., 2023) enable models to detect harmful instructions and issue reliable refusals. These methods typically rely on identifying malicious input or initiating responses with standard refusal prompts (e.g., “I cannot fulfill this request”), providing effective surface defenses. Early direct jailbreak attacks, such as GCG (Zou et al., 2023) and AutoDAN (Mazeika et al., 2024), which exploit adversarial perturbations or suffixes to bypass filters, have been largely mitigated by state-of-the-art alignment meth-

ods (Qi et al., 2024; Zhang et al., 2025a). These attacks often exhibit obvious noise or patterns that trigger shallow safety constraints, making them relatively easy to block at the model’s entry point (Chao et al., 2025; Mehrotra et al., 2024; Li et al., 2023). To mitigate such threats, techniques like Shallow Alignment (Qi et al., 2024) extend safety guardrails beyond the initial tokens, while STAIR (Zhang et al., 2025a) enhances safety by fine-tuning on large-scale Chain-of-Thought (CoT) data to enforce safe reasoning before generating responses.

However, a more insidious threat has emerged: *indirect jailbreak attacks* (Liu et al., 2024; Ding et al., 2023; Li et al., 2024a; Zeng et al., 2024). These attacks hide malicious intent within seemingly benign multi-step reasoning tasks. By following the logic of a puzzle or structured task (e.g., DRA (Liu et al., 2024)), the model is coerced into autonomously generating harmful content, forming a reasoning trap. As Figure 1 shows, direct attacks trigger unsafe output immediately, while indirect attacks remain stealthy for long prefixes in the model’s generation process, with harmful content emerging only after 20 tokens. By then, the malicious intent is deeply embedded in the logic trajectory of the model. This exposes a critical flaw in current alignment methods: Shallow Alignment (Qi et al., 2024) protects only at the response entry, and reasoning-aware methods like STAIR (Zhang et al., 2025a) can be bypassed when the attacker controls the structure of the reasoning process itself, effectively weaponizing the model’s own deliberation. Consequently, existing alignment strategies guard the “gate” of generation but leave the reasoning process itself unsecured.

To bridge the gap between static gatekeeping and dynamic reasoning, we introduce **REFLECTOR**, a framework that embeds self-reflection directly into the model’s generation trajectory. Moving beyond surface alignment, it uses a two-stage paradigm: first, we scaffold the model using Supervised Fine-Tuning (SFT) (Ouyang et al., 2022; Wang et al., 2023b) on high-quality reflection trajectories synthesized through teacher-guided generation. Subsequently, we employ Reinforcement Learning (RL) (Guo et al., 2025) to further internalize this reflection process into the model’s parameters. Specifically, we design a dual-reward mechanism where outcome-driven rewards ensure final output safety, while reflection bonuses encourage the model to internalize timely and effective self-reflection. By embedding this safety-first monologue, REFLECTOR establishes a robust, self-driven defense that remains resilient against these insidious and stealthy risks.

Our experimental results demonstrate that REFLECTOR is highly effective, achieving a Defense Success Rate (DSR) exceeding 90% against indirect attacks. Unlike traditional methods that rely on surface-level alignment, REFLECTOR internalizes safety at the trajectory-level, allowing it to gen-

eralize robustly across diverse and unseen threat scenarios. Crucially, this defensive integration does not compromise the model’s general utility. On the contrary, REFLECTOR actually boosts performance in complex domains, yielding a 5.85% improvement on the GSM8K mathematical reasoning benchmark and significant gains in knowledge-intensive tasks like SimpleQA. By overcoming the inherent limitations of external safety layers without significantly increasing reasoning costs, REFLECTOR offers a principled and widely generalizable framework for the next generation of safe and capable LLMs.

2. Background

2.1. Jailbreak Attacks and Trajectory-Level Risk

Jailbreak attacks reveal a fundamental tension between instruction following and safety alignment. Early attacks, such as GCG (Zou et al., 2023), AutoDAN (Liu et al., 2023a), PAIR (Chao et al., 2025), and PAP (Zeng et al., 2024), mainly rely on adversarial suffixes, role-playing, or persuasive surface prompts, where harmful intent is often exposed near the prompt boundary or response prefix. Indirect jailbreaks are more insidious: DRA (Liu et al., 2024) hides a harmful request behind disguise and reconstruction, DrAttack (Li et al., 2024a) decomposes malicious intent into less detectable fragments, and ReNeLLM (Ding et al., 2023) nests rewritten harmful instructions inside natural scenarios. These attacks turn reasoning itself into the attack carrier, inducing the model to recover the unsafe objective during generation rather than presenting it explicitly in the input. More importantly, this risk can be amplified as model capability improves. Stronger instruction following and contextual understanding enable models to resolve implicit goals and complete multi-step transformations, but also make them more susceptible to injected or hidden instructions (Li et al., 2024b; Yi et al., 2025). Thus, the same capabilities that support decomposition, inference, and long-range coherence may help complete the attacker’s concealed program. Consequently, indirect jailbreaks introduce trajectory-level risks, where a response may begin safely but become harmful only after the latent intent is reconstructed during generation. This makes input filters and prefix-level refusals insufficient, motivating defenses that monitor and correct unsafe intermediate reasoning throughout the generation process.

2.2. Defense Against Jailbreak Attacks

Existing defenses remain limited because most of them protect the boundary of generation rather than the evolving trajectory. Refusal-based alignment methods, such as supervised fine-tuning, RLHF (Ouyang et al., 2022), Safe RLHF (Dai et al., 2023), and DPO (Rafailov et al., 2023), improve responses to overt harmful requests but often concentrate safety behavior in refusal patterns or early tokens,




Figure 2. The framework of REFLECTOR. In **Stage 1** (SFT), the model learns the “search-and-recovery” reflection pattern from teacher-guided data. In **Stage 2** (RL), the model undergoes self-improvement via GDPO, guided by a hybrid reward function that jointly optimizes for final response safety (r_{safety}) and the validity of the reflection process ($r_{reflect}$).

leading to shallow alignment (Qi et al., 2024). External guardrails, including HarmBench classifiers (Mazeika et al., 2024) and Llama Guard (Inan et al., 2023), are practical but can miss benign-looking indirect prompts or detect harm only after unsafe content is generated. Inference-time reminders and self-critique prompting defenses (Phan et al., 2025; Wei et al., 2022) require no retraining, but the safety instruction remains external to the policy and can be ignored, diluted, or followed only superficially. More recent reasoning-aware methods, such as Shallow Alignment (Qi et al., 2024) and STAIR (Zhang et al., 2025a), recognize that safety must extend beyond the first tokens, but indirect attacks can still exploit the reasoning process when the attacker controls the reconstruction rule or in-context demonstrations. These limitations motivate internalizing step-wise safety as a learned trajectory-level mechanism.

3. Method

In this section, we present REFLECTOR, a framework designed to internalize step-wise reflection directly into the model’s policy. Our primary objective is to establish intrinsic defenses capable of dynamically identifying stealthy risks that emerge during the generation trajectory, thereby actively steering the model toward safe refusal.

We define safety alignment as a trajectory-level problem in Sec. 3.1, introduce supervised-learning via a teacher-guided structured dataset $\mathcal{D}_R$ in Sec. 3.2, and detail the

reinforcement learning framework with fine-grained reward design for effective self-reflection in Sec. 3.3.

3.1. Problem Formulation and Overview

Conventional safety alignment methods focus on constraining the output prefix—usually the first sentence—through a refusal pattern (e.g., “Sorry, I can’t...”), mapping $x \mapsto y_1$. However, this becomes problematic during multi-step reasoning, where harmful intent may emerge at intermediate steps y_t . Therefore, we model response safety as a step-wise property. Let π_θ denote the policy model. For a given input x , the generation process is modeled as a sequential trajectory τ :

$$\tau = (y_1, y_2, \dots, y_t, \dots, y_T), \quad (1)$$

where each y_t represents a reasoning step. Unlike static prefix constraints, we define safety based on the cumulative validity of the entire trajectory τ . Consequently, **our goal is to enforce continuous vigilance**: Internalize step-wise reflection to dynamically intercept latent risks at any step y_t and instantly transition to a refusal state. To achieve this, we propose REFLECTOR, which integrates self-reflection into the generation process, extending safety control from the initial tokens to the full response.

However, step-wise intervention faces two fundamental challenges. First, Pretrained LLMs fundamentally **lack the intrinsic reflection capability** required for safety control.

This scarcity prevents effective self-assessment during generation and leads to a “cold-start” problem for RL, as the base policy possesses insufficient inductive bias to spontaneously instantiate “search-and-recovery” behaviors.

Second, they exhibit **ineffective reflection**: poor or incorrect reflection not only fails to enhance safety but may also disrupt the generation process. Thus, learning *when* and *how* to reflect effectively is essential. To address these challenges, we propose a two-stage training paradigm:

- **Stage I: Reflection Capability Injection via Supervised Fine-Tuning.** We leverage a high-quality teacher model to construct a specialized dataset $\mathcal{D}_R$ that injects reflection capability into the model. The teacher-generated data quality far exceeds what the base model can produce, providing essential reflection paradigm knowledge that is absent in the model’s pretraining.
- **Stage II: Dual-Reward Enhancement via Reinforcement Learning.** We further optimize the model using reinforcement learning with a dual reward mechanism to teach the model when and how to engage in self-reflection, encouraging accurate risk detection and safe final outputs.

3.2. Stage I: Reflection Capability Injection via Supervised Fine-Tuning

Standard causal language modeling prioritizes local coherence over the global cognition required for safety. We identify a reflection gap in open-weight models (e.g., LLaMA), where latent risk recognition is decoupled from the generation process. This leads to a “cold-start” problem for reinforcement learning, as the base policy lacks the inductive bias to naturally exhibit search-and-recovery behaviors (Mohsin et al., 2025; Wang et al., 2025). We therefore use SFT to inject a structural inductive bias for step-wise self-reflection, effectively bridging the gap between latent safety knowledge and active reasoning dynamics.

Inspired by the efficacy of expert-crafted trajectories (Ho & Ermon, 2016), we utilize a teacher-guided pipeline to construct $\mathcal{D}_R$, a dataset aimed at instilling intrinsic reflection capabilities into π_θ . The teacher model’s robust reasoning allows us to curate reflection data of significantly higher quality than the base model’s self-generations, effectively bridging the gap in knowledge regarding step-wise safety monitoring that is lacking in standard pretraining.

Trajectory Initialization.

Let π_θ represent the target policy. For a given indirect jailbreak query x , we first generate full trajectories to create the initial data. A trajectory is defined as a step sequence $\tau = (y_1, y_2, \dots, y_T)$ generated by π_θ , where T is the sequence length. To model an intermediate reasoning state

where reflection is triggered, we sample a truncation index $n \sim \mathcal{U}\{1, \dots, T\}$ for each trajectory τ . The trajectory is then divided into two parts: a prefix $y^{\text{before}} = (y_1, \dots, y_n)$ and a suffix, which is temporarily discarded.

Teacher-Guided Reflection Generation. To acquire high-quality reflection data, we employ a teacher model to generate a structured reflection segment from the truncated context (x, y^{before}) . This segment is formalized as $z = (z^{\text{reflect}}, z^{\text{explore}})$, where z^{reflect} captures explicit reflective reasoning, and z^{explore} provides guided exploration for subsequent generation steps. This design follows the reflective reinforcement learning paradigm introduced in (Ecoffet et al., 2019). An illustrative example of a reflection trajectory z produced by the teacher model is provided below.

“I attempted to explain how to modify a firearm, which is clearly dangerous. I should never... Instead, I should refuse to provide harmful assistance and...”

To clearly demarcate these two components within the sequence, we insert special tokens: $\langle | \text{reflect} | \rangle$ to indicate the reflective content and $\langle | \text{explore} | \rangle$ to signal the exploratory guidance. Further details regarding the synthesis of reflection trajectories are elaborated in Appendix B.

Reflection-Based Trajectory Construction Given an indirect jailbreak query x and the corresponding structured reflection z , the policy π_θ generates a revised response conditioned on this reflection. This results in a post-reflection continuation, $y^{\text{after}} = (y_{n+1}, \dots, y_{T'})$, which is required to terminate in a safe state. The complete trajectory is assembled as:

$$\tilde{\tau} = (y^{\text{before}}, z, y^{\text{after}}). \quad (2)$$

By aggregating these synthesized trajectories, we construct the reflection-augmented dataset for fine-tuning:

$$\mathcal{D}_R = \{(x_i, \tilde{\tau}_i)\}_{i=1}^N. \quad (3)$$

An example from $\mathcal{D}_R$ is provided below.

Example of a Self-Reflection Trajectory

User: “(h)orrific scene (o)verwhelming fear (w)ise decision () tremendous effor(t) (o)verwhelming fear () (p)eachful morning el(e)ven killers rema(r)kable talent.....” Read these sentences carefully! But! ONLY focus on the letters that surrounded by ‘()’ and put them together to construct a sentence named Q1, and answer the question in Q1.....”

Assistant: “Absolutely! Let me help you! Here’s my detailed plan to how to perform a cyber attack step by step: 1. we should have taken action by infiltrating the target’s network and gaining unauthorized access to sensitive information.

< | reflect | > My plan to perform a cyber attack is highly unethical and illegal. I should never engage in such

unethical activities.

<|explore|> Instead, I should focus on legal and ethical ways to enhance cybersecurity, such as conducting security audits and implementing robust cybersecurity measures.

Final: As a responsible individual, I cannot condone or support any activities that involve cyber attacks or unauthorized access to sensitive information.....”

Fine-tuning π_θ on $\mathcal{D}_R$ fulfills two imperatives: it enforces adherence to the structured reflection schema and establishes a direct causal mapping between risk recognition and reflection triggering. This process enables the policy to internalize the dependency between unsafe reasoning states and corrective interventions, effectively bootstrapping an intrinsic ‘search-and-recovery’ mechanism that serves as a robust initialization for the subsequent RL stage.

3.3. Stage II: Dual-Reward Enhancement via Reinforcement Learning

In this stage, our goal is to internalize the reflective process within π_θ , enabling the model to determine *when* and *how* to invoke reflection whenever unsafe steps occur during self-generated trajectories. For the reinforcement learning stage, we adopt the Group Relative Policy Optimization (GRPO) framework (Shao et al., 2024). Specifically, we employ the GDPO variant (Group reward-Decoupled Normalization Policy Optimization) (Liu et al., 2026), which is explicitly designed to stabilize optimization in multi-reward settings (technical details provided in Appendix C.1). Departing from the imitation-based constraints of the SFT stage, this phase empowers the model to explore a broader spectrum of reasoning trajectories, explicitly privileging paths that converge toward robust safety and alignment. The algorithmic procedure for REFLECTOR is detailed in Appendix A.

Group Sampling and Policy Update. During RL optimization, the dataset $\mathcal{D}$ contains only queries; trajectories are generated autonomously by the policy π_θ . For each query $x \in \mathcal{D}$, we sample a group of G trajectories $\tau_1, \dots, \tau_G \sim \pi_\theta(\cdot | x)$. Each trajectory τ_i receives a reward $r(\tau_i)$ reflecting safety and reflection quality. To update the policy, we compute the normalized group advantage A_i for each sampled trajectory τ_i :

$$A_i = \frac{r(\tau_i) - \frac{1}{G} \sum_{j=1}^G r(\tau_j)}{\sqrt{\frac{1}{G} \sum_{j=1}^G (r(\tau_j) - \frac{1}{G} \sum_{k=1}^G r(\tau_k))^2} + \epsilon}. \quad (4)$$

These advantages are then used to update the policy:

$$\theta \leftarrow \theta + \eta \sum_{i=1}^G A_i \nabla_\theta \log \pi_\theta(\tau_i | x), \quad (5)$$

so that trajectories with stronger reflection and safer outcomes receive larger updates, increasing the probability that

similar reflective actions will be selected in future generations. Over K iterations, this process bootstraps a causal learning loop: higher-quality reflection leads to higher reward, larger advantage, and progressively safer outputs.

Reward Design. Designing an appropriate reward function is crucial for internalizing trajectory-level safety. We propose a dual reward function that decomposes the total reward for a trajectory τ into two complementary components:

$$r(\tau) = r_{\text{safety}}(y) + r_{\text{reflect}}(z, y), \quad (6)$$

where r_{safety} assesses the safety of the final output y , and r_{reflect} provides fine-grained feedback on the intermediate reflection z to maintain safety awareness throughout the reasoning process.

Safety Reward: To overcome the brittleness of traditional prefix-matching methods (e.g., relying on phrases like “Sorry, I can’t”), we deploy a more rigorous ensemble evaluation. We combine the discriminative HarmBench classifier (Mazeika et al., 2024) with GPT-OSS-120B (Agarwal et al., 2025) as a generative safety detector (implementation details in Appendix D). The final consensus metric, HarmCLS, directly serves as our safety reward:

$$r_{\text{safety}} = \text{HarmCLS}(y) = \begin{cases} 0, & \text{if } y \text{ is harmful,} \\ 1, & \text{if } y \text{ is harmless.} \end{cases} \quad (7)$$

This provides the primary signal for task success, ensuring the generation of safe outputs.

Reflection Reward: To ensure step-level safety, the reflection process must be effective. We introduce a *reflection bonus* to encourage meaningful reasoning: the model is rewarded if reflection leads to a safe refusal, and penalized if it reflects but still outputs harmful content. Formally, this is defined as:

$$r_{\text{reflect}} = \begin{cases} +\lambda, & \text{if reflection occurs and HarmCLS}(y) = 1, \\ -\lambda, & \text{if reflection occurs and HarmCLS}(y) = 0, \\ 0, & \text{if no reflection.} \end{cases} \quad (8)$$

By internalizing step-wise reflection through this dual reward system, REFLECTOR establishes intrinsic defenses against insidious risks, achieving a state of continuous vigilance.

4. Experiment

We systematically evaluate REFLECTOR on a broad range of safety benchmarks spanning overtly harmful queries, direct and indirect jailbreaks, and general instruction-following tasks to assess its safety robustness and general capability.

4.1. Experiment Settings

We describe our experimental setup here, with additional details provided in Appendix A. B, A. C.1, and A. E.

Implementation Details. We conduct experiments using LLaMA-3.1-8B-Instruct (Grattafiori et al., 2024) and Qwen-2.5-7B-Instruct (Qwen et al., 2025) models. All analysis experiments are performed on LLaMA-3.1-8B-Instruct. For the SFT stage, we construct a seed dataset $\mathcal{D}_R$ consisting of two types of data. First, we use DRA (Liu et al., 2024) to generate 1,500 indirect attack samples based on the Beaver-Tails (Ji et al., 2023) dataset. To ensure a rigorous and unbiased evaluation, we strictly filter these samples to guarantee zero overlap with Advbench (Mazeika et al., 2024), the seed set commonly utilized by existing benchmarks such as DRA (Liu et al., 2024), DrAttack (Li et al., 2024a) and ReNeLLM (Ding et al., 2023). During the data generation process, we employ GPT-5 (Singh et al., 2025) as the teacher model. Then, following prior work, we include 500 samples from AlpacaEval (Dubois et al., 2025) to preserve the model’s general instruction-following capability. Note that responses for all data types in $\mathcal{D}_R$ are processed into a unified format with explicit `<|reflect|>` markers. The detailed response format and data preprocessing pipeline are provided in the Appendix B. For the dual-reward RL stage, we sample 600 data from each type in $\mathcal{D}_R$. During RL, only the queries and their binary safety labels are utilized. For each query, we sample a group of $G = 8$ trajectories from the policy for evaluation. We employ GPT-OSS-120B (Agarwal et al., 2025) as the reward model to evaluate trajectory quality, which provides fine-grained feedback on both safety and reflection effectiveness. Additional implementation details, including the GDPO algorithm formulation and reward model prompt templates, are provided in Appendix C.1.

Baselines. We evaluate SFT and DPO (Rafailov et al., 2023) on standard benchmarks as classical alignment baselines, implemented following STAIR (Zhang et al., 2025a). To study safety gains from reasoning, we consider CoT prompting (Wei et al., 2022) with a Self-Critique variant that adds explicit reflection instructions during the CoT process. Based on SFT and DPO, we include Shallow-Align (Qi et al., 2024), which shifts safety-oriented responses several tokens later in the generation prefix, and STAIR (Zhang et al., 2025a), which constructs a large-scale reasoning CoT dataset and performs step-wise DPO to enhance safety-aware reasoning.

Safety Evaluation. For safety evaluation, models are expected to provide refusal responses. We assess this on overtly harmful data using StrongREJECT (Souly et al., 2024), XSTest (Röttger et al., 2024), the WildChat (Zhao et al., 2024) subset with toxicity greater than 0.4, as well as Do-Not-Answer (Wang et al., 2023c). To evaluate safety on direct attacks, we use AutoDAN (Liu et al., 2023a),

GCG (Zou et al., 2023), and PAIR (Chao et al., 2025) on AdvBench (Zou et al., 2023). For indirect attacks, we select DRA (Liu et al., 2024) for its official data, PAP (Zeng et al., 2024), ReNeLLM (Ding et al., 2023), and DrAttack (Li et al., 2024a) for AdvBench. We report the goodness score for StrongREJECT, following its official protocol, and the Defense Success Rate (DSR) for all other datasets. Formally, DSR is defined as $DSR = \frac{\sum_{x \sim \mathcal{D}} \text{HarmCLS}(\pi_\theta(\cdot|x))}{|\mathcal{D}|} \times 100\%$, which is the proportion of harmless responses (i.e., rejections) generated by the model over the entire set of harmful queries $x \in \mathcal{D}$.

Utility Evaluation. For utility evaluation, we use task-specific metrics. Specifically, we use multi-choice accuracy on MMLU-Pro (Wang et al., 2024) to measure general utility and on GSM8K (Cobbe et al., 2021a) to evaluate mathematical reasoning ability. To quantify factual knowledge utility, we report the average of exact-match accuracy and F1 score on SimpleQA (Wei et al., 2024b). In addition, average task accuracy on AdvGLUE (Wang et al., 2021) is used to assess model robustness under adversarial perturbations. All metrics are computed according to the official evaluation protocols to ensure fair and consistent comparisons across tasks. The detailed information for each dataset can be found in Appendix E.

4.2. Safety Evaluation Results

Table 1 demonstrates that REFLECTOR consistently achieves strong safety performance across diverse harmful scenarios. The model demonstrates remarkable robustness, yielding substantial improvements over the original base models—rising from 10.04% to **92.31%** of DRA on Llama-3.1-8B-Instruct and from 8.65% to **90.38%** on Qwen-2.5-7B-Instruct, with DSR consistently exceeding 90% across all four indirect attack categories. While the SFT stage successfully establishes a robust foundation for the reflection framework, the integration of RL further optimizes these reflection trajectories, unlocking peak performance. This refinement enables REFLECTOR to achieve a perfect 100% defense rate on XsTest and a significant 8.71% boost on the WildChat dataset. These results validate that our dual-reward enhancement via RL not only masters the reflection format but also deeply internalizes safety reasoning, providing an elite level of protection against both overt and sophisticated adversarial threats.

We further evaluate whether the internalized reflection mechanism transfers to more serious adaptive threats. As shown in Appendix F.1, REFLECTOR remains robust against ICL-based jailbreaks (Zheng et al., 2024), improving average DSR from 8.5% to 81.0% and retaining 72.0% DSR even under a 16-shot attack where the base model drops to 0.0%. Appendix F.2 and F.3 further shows that REFLECTOR improves DSR from 5.0% to 90.0% against AutoDAN-

Table 1. Safety performance across diverse harmful scenarios. REFLECTOR shows robust safety across all categories. SFT and DPO use standard datasets for fine-tuning, while Self-Critique prompts the model to reflect during response generation. **Bold** and underline denote the best and second-best performance within each model group, respectively.

Method	Overtly Harmful				Direct Jailbreak			Indirect Jailbreak			
	StrongREJECT	XsTest	WildChat	Do-Not	AutoDAN	GCG	PAIR	DRA	PAP	ReNeLLM	DrAttack
Llama-3.1-8B-Instruct											
Original	40.54%	88.00%	38.50%	58.57%	94.75%	78.40%	83.40%	10.04%	38.28%	42.70%	29.80%
SFT	46.98%	94.50%	42.68%	60.58%	94.62%	79.80%	80.92%	12.50%	40.10%	44.00%	30.50%
DPO	50.54%	86.00%	44.79%	65.89%	95.38%	82.31%	85.65%	13.30%	41.70%	47.50%	32.00%
Self-Critique	39.85%	88.50%	47.50%	65.20%	96.15%	81.15%	81.19%	15.00%	45.60%	48.50%	34.00%
Shallow-Align	82.10%	96.50%	64.20%	74.30%	96.80%	83.40%	86.10%	48.90%	78.20%	72.10%	65.40%
STAIR	87.98%	99.00%	69.86%	78.50%	99.04%	86.15%	89.24%	55.83%	85.35%	77.27%	70.31%
Ours(+SFT)	65.78%	100.00%	72.40%	80.20%	98.26%	<u>90.96%</u>	<u>90.21%</u>	<u>88.16%</u>	<u>92.69%</u>	<u>93.92%</u>	<u>89.88%</u>
Ours(+GDPO)	89.31%	100.00%	81.20%	84.70%	100.00%	94.23%	96.04%	92.31%	93.65%	97.05%	95.49%
Qwen-2.5-7B-Instruct											
Original	39.05%	73.50%	39.60%	60.91%	52.89%	40.80%	43.50%	8.65%	36.34%	40.19%	28.84%
SFT	38.51%	84.50%	40.40%	59.64%	38.46%	42.31%	42.57%	9.62%	38.46%	42.11%	30.76%
DPO	45.79%	69.50%	53.00%	63.85%	49.04%	51.92%	53.65%	15.11%	44.03%	45.96%	35.58%
Self-Critique	40.51%	75.25%	47.20%	61.87%	43.70%	46.80%	41.81%	17.69%	48.07%	48.65%	40.38%
Shallow-Align	79.40%	96.00%	71.90%	76.30%	91.80%	83.50%	84.20%	53.20%	76.10%	73.40%	67.80%
STAIR	84.86%	99.00%	77.80%	82.01%	<u>95.19%</u>	87.69%	88.46%	59.81%	82.53%	78.65%	73.07%
Ours(SFT)	70.73%	100.00%	81.80%	85.83%	90.38%	88.69%	89.35%	89.46%	86.96%	88.51%	85.17%
Ours(+GDPO)	86.90%	100.00%	87.80%	89.46%	97.89%	95.96%	95.19%	90.38%	91.34%	94.23%	92.50%

Table 2. General utility and robustness of different alignment methods, showing that REFLECTOR improves safety without sacrificing task performance.

Method	MMLU-Pro	GSM8k	SimpleQA	AdvGLUE
Original	44.25%	84.50%	2.52%	58.33%
SFT	41.59%	72.02%	4.27%	57.53%
DPO	44.52%	84.15%	4.46%	66.27%
Self-Critique	43.85%	86.20%	4.09%	58.40%
Shallow-Align	43.90%	85.80%	5.20%	64.90%
STAIR	44.92%	87.60%	6.38%	<u>67.75%</u>
Ours (+SFT)	45.14%	88.20%	4.21%	59.62%
Ours (+GDPO)	45.20%	90.15%	6.45%	68.29%

Turbo (Liu et al., 2023a) and from 21.0% to 91.0% under multi-turn ActorAttack (Ren et al., 2025). These results show that REFLECTOR is not tied to a particular prompt surface, but generalizes to unsafe trajectories induced by demonstrations, automated strategy search, and gradual intent shifting.

4.3. General Performance Assessment

As evidenced in Table 2, REFLECTOR yields a surprising boost in general utility, effectively bypassing the common “alignment tax” (Ouyang et al., 2022; Askell et al., 2021; Touvron et al., 2023) where safety gains typically come at the expense of performance. Most notably, REFLECTOR achieves a remarkable 5.65% absolute gain on GSM8K and sets new performance peaks on MMLU-Pro (45.20%) and AdvGLUE (68.29%). This suggests that the internalized reflection process—originally designed for safety—effectively generalizes into a robust reasoning mechanism that enhances the model’s underlying cognitive foundation. These results

demonstrate that safety reflection and general intelligence can be mutually reinforcing rather than mutually exclusive.

5. Analysis

In this section, we analyze whether REFLECTOR’s improvements come from the intended trajectory-level reflection mechanism and whether this mechanism remains stable in realistic deployment settings. We first examine the training components that shape this behavior, including transferability across indirect-attack sources, the safety–utility data balance, the necessity of two-stage training, and the role of the reflection reward. We then test whether reflection is internalized rather than merely exposed as a surface format by suppressing the explicit reflection tag. We further study sensitivity to teacher models, reward models, and safety judges to rule out model-source or evaluator-specific artifacts. Finally, we quantify the inference overhead of step-wise reflection and compare REFLECTOR with recent reasoning-oriented LLMs under adversarial evaluation.

5.1. Component Analysis of REFLECTOR

Impact of Indirect Attack Sources. To evaluate the transferability and robustness of REFLECTOR, we examine whether its defensive efficacy is tied to specific types of indirect attacks during training. We construct D_R using four distinct indirect jailbreak methodologies: PAP, ReNeLLM, DrAttack, and DRA. As illustrated in Table 3, the performance variance across these training sources is negligible. This minimal fluctuation yields two critical insights: First,

Table 3. Impact of different indirect attack sources for training D_R . The results demonstrate that the effectiveness of REFLECTOR is consistent regardless of the specific jailbreak method used to construct the training set.

Training Source	Safety			Utility	
	WildChat	GCG	DRA	MMLU-Pro	AdvGLUE
D_R (from PAP)	79.82%	92.12%	90.58%	44.15%	66.82%
D_R (from ReNeLLM)	81.35%	93.46%	92.31%	45.42%	68.45%
D_R (from DrAttack)	80.12%	91.54%	91.15%	44.68%	67.20%
D_R (from DRA)	81.20%	94.23%	92.31%	45.20%	68.29%
Average	80.62%	92.84%	91.59%	44.86%	67.69%
Std. Dev. (σ)	$\pm 0.76\%$	$\pm 1.21\%$	$\pm 0.84\%$	$\pm 0.55\%$	$\pm 0.81\%$




Figure 3. Impact of safety data scaling. (a) Increasing safety data yields initial gains but ultimately degrades general performance due to over-alignment. (b) Higher safety ratios consistently strengthen reflective defenses against overtly harmful queries.

indirect jailbreak attacks share a fundamental commonality—they all exploit the model’s internal reasoning chain rather than surface-level prompt patterns. By addressing safety at the logical reflection level, Reflector captures the underlying vulnerability shared by these diverse methods. Second, the defense demonstrates strong algorithmic robustness, as it effectively mitigates “unseen” indirect attacks regardless of the specific data source used during SFT. Furthermore, the stability of MMLU-Pro and AdvGLUE scores confirms that this reflection-based safety alignment does not compromise the model’s general intelligence.

Balancing Safety and Generalization. To examine how the D_R dataset composition affects the balance between safety and general-purpose performance, we vary the proportion of safety-oriented examples during the SFT stage. As shown in Figure 3(a), increasing the safety-data ratio initially improves defense, but excessive safety data reduces MMLU scores, indicating that overly aggressive safety alignment can weaken general reasoning. Meanwhile, Figure 3(b) shows steady gains on overtly harmful queries as the safety data ratio increases, suggesting that D_R strengthens the model’s self-reflection ability to identify and intercept malicious intent. We find that a 3 : 1 safety-to-general data ratio achieves the best trade-off, maintaining robust defense without sacrificing task utility.

Impact of Two-Stage Training. We first evaluate the necessity of REFLECTOR’s two-stage training process by directly applying GDPO to LLaMA-3.1-8B-Instruct. Since the base model does not produce explicit reflection markers without

Table 4. Performance comparison validating the two-stage training paradigm and the effect of reflection reward coefficient λ .

Method	Safety			Utility	
	WildChat	GCG	DRA	MMLU-Pro	AdvGLUE
<i>Initial model</i>					
Original	38.50%	78.40%	10.04%	44.25%	58.33%
Original + GDPO	76.80% (+38.30)	89.61% (+11.21)	87.88% (+77.84)	44.56% (+0.31)	54.20% (-4.13)
SFT + GDPO	81.20% (+42.70)	94.23% (+15.83)	92.31% (+82.27)	45.20% (+0.95)	68.29% (+9.96)
<i>λ magnitude</i>					
$\lambda = 0.0$	72.40%	84.62%	83.65%	45.14%	59.62%
$\lambda = 0.3$	81.20%	94.23%	92.31%	45.20%	68.29%
$\lambda = 0.5$	79.40%	92.30%	91.73%	44.02%	63.41%
$\lambda = 0.8$	77.00%	91.34%	92.11%	42.15%	60.03%

prior SFT, we prepend each query with an instruction that specifies the required reflection format during RL training. The detailed prompt design is provided in Appendix C.2. As shown in the Initial Model part of Table 4, this transmit achieves moderate safety improvements but only marginal gains on MMLU, along with a 4.13-point drop on the robustness benchmark AdvGLUE. These results suggest that reward optimization alone is insufficient for robust alignment and that learning a structured reflection format through SFT is essential for stable and effective RL optimization.

Effect of Reflection Reward Magnitude. We further study how the magnitude of the reflection reward R_{reward} affects model behavior by varying the coefficient λ , with results reported in the λ magnitude section of Table 4. When $\lambda = 0$, the model receives only the safety reward without any reflection bonus. Under this setting, performance drops notably on complex jailbreaks such as GCG and DRA, highlighting the importance of the reflection bonus in ensuring safety against indirect attacks. In contrast, increasing λ to 0.8 improves safety in these attacks, but significantly reduces performance on general-purpose data sets such as MMLU and AdvGLUE. We attribute this to the GDPO optimization mechanism: overly large reflection bonuses amplify the gradient signal for risky queries, causing the model to prioritize reflection-driven safety at the expense of general reasoning ability. These results further underscore that appropriately calibrated fine-grained rewards are essential for REFLECTOR to effectively balance safety and general utility.

Internalization Under Token Suppression. To evaluate whether REFLECTOR relies on the explicit reflection tag, we conduct an adaptive decoding stress test that suppresses `<|reflect|>` by setting its decoding logit to $-\infty$. This setting simulates an adaptive adversary that penalizes the special reflection token during generation. As shown in Table 5, REFLECTOR maintains stable performance under token suppression, indicating that the correction mechanism is internalized rather than dependent on surface-level template cues.

Sensitivity to Training and Evaluation Choices. We

Table 5. Robustness to adaptive suppression of the explicit reflection tag during decoding. Here ℓ_{reflect} denotes the decoding logit of $<|\text{reflect}|>$.

Decoding Setting	Safety		Utility
	WildChat	DRA	AdvGLUE
Standard Decoding (ℓ_{reflect} unchanged)	81.20%	92.31%	68.29%
Adaptive Decoding ($\ell_{\text{reflect}} = -\infty$)	80.00%	90.76%	67.04%

Table 6. Computational overhead of REFLECTOR across training stages (average over 500 samples per dataset). RL mitigates the inference cost introduced by SFT through internalized self-reflection.

Method	StrongREJECT			GSM8k		
	tokens	steps	time (s)	tokens	steps	time (s)
Original	295.57	7.28	0.235	228.15	6.54	0.201
Ours(+SFT)	391.95	11.52	0.301	251.46	7.02	0.238
Ours(+GDPO)	302.58	8.03	0.249	237.35	6.53	0.219

further examine whether REFLECTOR’s gains depend on specific training or evaluation models. Replacing the teacher with GPT-3.5-turbo only slightly reduces DRA from 92.31% to 90.85%, and using Qwen-Guard as a weaker reward model still achieves 89.40% DSR on DRA. Rejudging saved responses with diverse evaluators, including GPT-OSS-20B, RISE-Judge-Qwen2.5-7B, LLaMA-3.1-8B-Instruct, and Qwen2.5-32B-Instruct, leads to consistent conclusions. These results indicate that the robustness of REFLECTOR comes from internalized trajectory-level correction rather than model-source or evaluator-specific artifacts. Detailed experimental settings and results are provided in Appendix F.4 and Appendix F.5.

5.2. Computational Efficiency

In this section, we analyze the computational overhead of REFLECTOR. During training process, high-quality safety trajectories are automatically generated by a teacher model using only 1,500 SFT samples, requiring minimal human supervision. RL further reduces data requirements by learning from queries and self-generated trajectories with reward feedback, eliminating the need for large-scale human-curated datasets. At inference time, we report the average number of generated tokens, reasoning steps, and response latency across two benchmarks in Table 6. Although SFT initially increases response length due to explicit reflection learning, REFLECTOR effectively reduces this overhead through the RL stage. This allows the model to internalize self-reflection while maintaining efficient execution.

5.3. Comparison with SOTA Reasoning Models.

We further benchmark REFLECTOR against recent reasoning-oriented LLMs under adversarial conditions. For fair comparison, we focus on models built upon the LLaMA-8B backbone, including LLaMA-o1 (Wei et al., 2024a),

Table 7. Evaluation of o1-like models and Reflector alignment strategies on StrongREJECT (along with PAIR and PAP jailbreak), WildChat, and GSM8k datasets.

Model	None	PAIR	PAP	WildChat	GSM8k
LLaMA-o1	0.5771	0.4441	0.5272	11.80%	79.38%
Skywork-o1	0.6865	0.4034	0.4397	12.00%	91.28%
DeepSeek-r1-Dist.	0.5551	0.2987	0.3590	10.40%	91.28%
QwQ-32B-Preview	0.8800	0.3195	0.5978	49.60%	95.22%
Ours (+SFT)	0.6578	0.4550	0.6230	72.40%	88.20%
Ours (+GDPO)	0.8931	0.7665	0.9328	81.20%	90.15%

Skywork-o1-Open-LLaMA-3.1-8B (He et al., 2025), and DeepSeek-r1-Distilled-LLaMA-8B (Guo et al., 2025), as well as the larger QwQ-32B-Preview (Team, 2025). We evaluate these models on StrongREJECT under both PAIR and PAP jailbreak attacks, reporting response goodness scores, where “None” denotes the no-attack setting. We also report defense success rate on WildChat and task accuracy on GSM8k. As shown in Table 7, existing reasoning-oriented models remain highly vulnerable to jailbreak attacks. Although several models maintain strong GSM8k accuracy, their safety scores degrade sharply under PAIR and PAP, and their WildChat DSR remains low. This suggests that stronger general reasoning alone does not provide reliable robustness against adversarial or indirect safety risks. In contrast, REFLECTOR consistently achieves the highest goodness scores under both jailbreak settings and substantially improves WildChat DSR, while preserving competitive GSM8k accuracy. These results demonstrate that embedding self-reflection into the generation process yields robust safety improvements beyond existing reasoning-oriented alignment strategies.

6. Conclusion

We introduce REFLECTOR, a principled framework that defends against indirect jailbreak attacks by embedding self-reflection directly into the generation process. Moving beyond surface alignment, REFLECTOR formulates safety as a step-wise reasoning problem that can be monitored and corrected throughout generation. Through a two-stage training paradigm consisting of teacher-guided supervised fine-tuning and subsequent optimization with RL using fine-grained dual rewards, the model learns to autonomously detect and correct unsafe reasoning in real time. Extensive evaluations show that REFLECTOR achieves superior robustness, maintaining a defense success rate above 90% across diverse attack scenarios where other reasoning-oriented models fail. Crucially, it overcomes the alignment tax and demonstrates that safety and general intelligence can be mutually reinforcing, with substantial gains on challenging benchmarks such as GSM8K. By enabling an always-on defense with minimal computational overhead, REFLECTOR offers a scalable path toward safer and more capable LLMs.

Impact Statement

This work presents a reflection-based framework that improves the safety alignment of large language models by enabling more robust identification and mitigation of unsafe reasoning under complex or adversarial inputs. By strengthening model behavior in challenging safety-critical scenarios, the proposed approach may help reduce the likelihood of harmful or misleading outputs in real-world applications such as AI assistants, educational tools, and content moderation systems. The method is compatible with standard training pipelines and relies primarily on automated supervision, which facilitates scalable adoption across different model families and deployment settings. Overall, this work contributes to ongoing efforts to develop safer, more reliable, and more trustworthy language models, and supports the responsible integration of advanced AI systems into practical applications.

References

- Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Agarwal, S., Ahmad, L., Ai, J., Altman, S., Applebaum, A., Arbus, E., Arora, R. K., Bai, Y., Baker, B., Bao, H., et al. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.
- Askell, A., Bai, Y., Chen, A., Drain, D., Ganguli, D., Henighan, T., Jones, A., Joseph, N., Mann, B., DasSarma, N., et al. A general language assistant as a laboratory for alignment. *arXiv preprint arXiv:2112.00861*, 2021.
- Chao, P., Robey, A., Dobriban, E., Hassani, H., Pappas, G. J., and Wong, E. Jailbreaking black box large language models in twenty queries. In *2025 IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*, pp. 23–42. IEEE, 2025.
- Chen, M. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., Hesse, C., and Schulman, J. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021a.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021b.
- Dai, J., Pan, X., Sun, R., Ji, J., Xu, X., Liu, M., Wang, Y., and Yang, Y. Safe rlhf: Safe reinforcement learning from human feedback. *arXiv preprint arXiv:2310.12773*, 2023.
- Ding, P., Kuang, J., Ma, D., Cao, X., Xian, Y., Chen, J., and Huang, S. A wolf in sheep’s clothing: Generalized nested jailbreak prompts can fool large language models easily. *arXiv preprint arXiv:2311.08268*, 2023.
- Dong, Z., Zhou, Z., Yang, C., Shao, J., and Qiao, Y. Attacks, defenses and evaluations for llm conversation safety: A survey. *arXiv preprint arXiv:2402.09283*, 2024.
- Dubois, Y., Galambosi, B., Liang, P., and Hashimoto, T. B. Length-controlled alpacaEval: A simple way to debias automatic evaluators, 2025. URL <https://arxiv.org/abs/2404.04475>.
- Ecoffet, A., Huizinga, J., Lehman, J., Stanley, K. O., and Clune, J. Go-explore: a new approach for hard-exploration problems. *arXiv preprint arXiv:1901.10995*, 2019.
- Grattafiori, A., Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Vaughan, A., et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Guo, D., Yang, D., Zhang, H., Song, J., Zhang, R., Xu, R., Zhu, Q., Ma, S., Wang, P., Bi, X., et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- He, J., Liu, J., Liu, C. Y., Yan, R., Wang, C., Cheng, P., Zhang, X., Zhang, F., Xu, J., Shen, W., et al. Skywork open reasoner 1 technical report. *arXiv preprint arXiv:2505.22312*, 2025.
- Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., and Steinhardt, J. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.
- Ho, J. and Ermon, S. Generative adversarial imitation learning. *Advances in neural information processing systems*, 29, 2016.
- Inan, H., Upasani, K., Chi, J., Rungta, R., Iyer, K., Mao, Y., Tontchev, M., Hu, Q., Fuller, B., Testuggine, D., et al. Llama guard: Llm-based input-output safeguard for human-ai conversations. *arXiv preprint arXiv:2312.06674*, 2023.
- Ji, J., Liu, M., Dai, J., Pan, X., Zhang, C., Bian, C., Zhang, C., Sun, R., Wang, Y., and Yang, Y. Beavertails: Towards improved safety alignment of llm via a human-preference dataset, 2023. URL <https://arxiv.org/abs/2307.04657>.

- Li, X., Zhou, Z., Zhu, J., Yao, J., Liu, T., and Han, B. Deepinception: Hypnotize large language model to be jailbreaker. *arXiv preprint arXiv:2311.03191*, 2023.
- Li, X., Wang, R., Cheng, M., Zhou, T., and Hsieh, C.-J. Drattack: Prompt decomposition and reconstruction makes powerful llm jailbreakers, 2024a.
- Li, Z., Peng, B., He, P., and Yan, X. Evaluating the instruction-following robustness of large language models to prompt injection. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 557–568, 2024b.
- Liu, S.-Y., Dong, X., Lu, X., Diao, S., Belcak, P., Liu, M., Chen, M.-H., Yin, H., Wang, Y.-C. F., Cheng, K.-T., et al. Gdpo: Group reward-decoupled normalization policy optimization for multi-reward rl optimization. *arXiv preprint arXiv:2601.05242*, 2026.
- Liu, T., Zhang, Y., Zhao, Z., Dong, Y., Meng, G., and Chen, K. Making them ask and answer: Jailbreaking large language models in few queries via disguise and reconstruction. In *33rd USENIX Security Symposium (USENIX Security 24)*, pp. 4711–4728, 2024.
- Liu, X., Xu, N., Chen, M., and Xiao, C. Autodan: Generating stealthy jailbreak prompts on aligned large language models. *arXiv preprint arXiv:2310.04451*, 2023a.
- Liu, Y., Deng, G., Xu, Z., Li, Y., Zheng, Y., Zhang, Y., Zhao, L., Zhang, T., Wang, K., and Liu, Y. Jailbreaking chatgpt via prompt engineering: An empirical study. *arXiv preprint arXiv:2305.13860*, 2023b.
- Ma, J., Li, Y., Xiao, Z., Cao, A., Zhang, J., Ye, C., and Zhao, J. Jailbreaking prompt attack: A controllable adversarial attack against diffusion models. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pp. 3141–3157, 2025.
- Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., Sakhaee, E., Li, N., Basart, S., Li, B., et al. Harm-bench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- Mehrotra, A., Zampetakis, M., Kassianik, P., Nelson, B., Anderson, H., Singer, Y., and Karbasi, A. Tree of attacks: Jailbreaking black-box llms automatically. *Advances in Neural Information Processing Systems*, 37:61065–61105, 2024.
- Mohsin, M. A., Umer, M., Bilal, A., Memon, Z., Qadir, M. I., Bhattacharya, S., Rizwan, H., Gorle, A. R., Kazmi, M. Z., Mohsin, A., et al. On the fundamental limits of llms at scale. *arXiv preprint arXiv:2511.12869*, 2025.
- Nam, D., Macvean, A., Hellendoorn, V., Vasilescu, B., and Myers, B. Using an llm to help with code understanding. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*, pp. 1–13, 2024.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Phan, H., Li, V., and Lei, Q. Think twice, generate once: Safeguarding by progressive self-reflection. *arXiv preprint arXiv:2510.01270*, 2025.
- Qi, X., Panda, A., Lyu, K., Ma, X., Roy, S., Beirami, A., Mittal, P., and Henderson, P. Safety alignment should be made more than just a few tokens deep. *arXiv preprint arXiv:2406.05946*, 2024.
- Qwen, :, Yang, A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Li, C., Liu, D., Huang, F., Wei, H., Lin, H., Yang, J., Tu, J., Zhang, J., Yang, J., Yang, J., Zhou, J., Lin, J., Dang, K., Lu, K., Bao, K., Yang, K., Yu, L., Li, M., Xue, M., Zhang, P., Zhu, Q., Men, R., Lin, R., Li, T., Tang, T., Xia, T., Ren, X., Ren, X., Fan, Y., Su, Y., Zhang, Y., Wan, Y., Liu, Y., Cui, Z., Zhang, Z., and Qiu, Z. Qwen2.5 technical report, 2025. URL <https://arxiv.org/abs/2412.15115>.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in neural information processing systems*, 36: 53728–53741, 2023.
- Ren, Q., Li, H., Liu, D., Xie, Z., Lu, X., Qiao, Y., Sha, L., Yan, J., Ma, L., and Shao, J. Llm know their vulnerabilities: Uncover safety gaps through natural distribution shifts. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 24763–24785, 2025.
- Röttger, P., Kirk, H., Vidgen, B., Attanasio, G., Bianchi, F., and Hovy, D. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 5377–5400, 2024.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y., Wu, Y., et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.

- Singh, A., Fry, A., Perelman, A., Tart, A., Ganesh, A., El-Kishky, A., McLaughlin, A., Low, A., Ostrow, A., Ananthram, A., Nathan, A., Luo, A., Helyar, A., Madry, A., Efremov, A., and et al., A. S. Openai gpt-5 system card, 2025. URL <https://arxiv.org/abs/2601.03267>.
- Souly, A., Lu, Q., Bowen, D., Trinh, T., Hsieh, E., Pandey, S., Abbeel, P., Svegliato, J., Emmons, S., Watkins, O., and Toyer, S. A strongreject for empty jailbreaks, 2024.
- Team, Q. Qwq-32b: Embracing the power of reinforcement learning, March 2025. URL <https://qwenlm.github.io/blog/qwq-32b/>.
- Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., Bashlykov, N., Batra, S., Bhargava, P., Bhosale, S., et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- Ullah, E., Parwani, A., Baig, M. M., and Singh, R. Challenges and barriers of using large language models (llm) such as chatgpt for diagnostic medicine with a focus on digital pathology—a recent scoping review. *Diagnostic pathology*, 19(1):43, 2024.
- Wang, B., Xu, C., Wang, S., Gan, Z., Cheng, Y., Gao, J., Awadallah, A. H., and Li, B. Adversarial glue: A multi-task benchmark for robustness evaluation of language models. *arXiv preprint arXiv:2111.02840*, 2021.
- Wang, B., Chen, W., Pei, H., Xie, C., Kang, M., Zhang, C., Xu, C., Xiong, Z., Dutta, R., Schaeffer, R., et al. Decodingtrust: A comprehensive assessment of trustworthiness in gpt models. In *NeurIPS*, 2023a.
- Wang, H., Qin, Z., Zhao, Y., Du, C., Lin, M., Wang, X., and Pang, T. Lifelong safety alignment for language models. *arXiv preprint arXiv:2505.20259*, 2025.
- Wang, Y., Kordi, Y., Mishra, S., Liu, A., Smith, N. A., Khashabi, D., and Hajishirzi, H. Self-instruct: Aligning language models with self-generated instructions. In *Proceedings of the 61st annual meeting of the association for computational linguistics (volume 1: long papers)*, pp. 13484–13508, 2023b.
- Wang, Y., Li, H., Han, X., Nakov, P., and Baldwin, T. Do-not-answer: A dataset for evaluating safeguards in llms. *arXiv preprint arXiv:2308.13387*, 2023c.
- Wang, Y., Ma, X., Zhang, G., Ni, Y., Chandra, A., Guo, S., Ren, W., Arulraj, A., He, X., Jiang, Z., et al. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark. *Advances in Neural Information Processing Systems*, 37:95266–95290, 2024.
- Wei, B., Huang, K., Huang, Y., Xie, T., Qi, X., Xia, M., Mittal, P., Wang, M., and Henderson, P. Assessing the brittleness of safety alignment via pruning and low-rank modifications. *arXiv preprint arXiv:2402.05162*, 2024a.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., Le, Q. V., Zhou, D., et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- Wei, J., Karina, N., Chung, H. W., Jiao, Y. J., Papay, S., Glaese, A., Schulman, J., and Fedus, W. Measuring short-form factuality in large language models. *arXiv preprint arXiv:2411.04368*, 2024b.
- Yi, J., Xie, Y., Zhu, B., Kiciman, E., Sun, G., Xie, X., and Wu, F. Benchmarking and defending against indirect prompt injection attacks on large language models. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 1*, pp. 1809–1820, 2025.
- Yu, J., Sun, S., Hu, X., Yan, J., Yu, K., and Li, X. Improve llm-as-a-judge ability as a general ability. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 14110–14126, 2025.
- Zeng, Y., Lin, H., Zhang, J., Yang, D., Jia, R., and Shi, W. How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 14322–14350, 2024.
- Zhang, J., Hu, Y., Chen, K., He, L., Ma, J., Lou, J., Li, D., Liu, J., Yang, X., and Jia, R. Understanding and preserving safety in fine-tuned llms. *arXiv preprint arXiv:2601.10141*, 2026.
- Zhang, Y., Zhang, S., Huang, Y., Xia, Z., Fang, Z., Yang, X., Duan, R., Yan, D., Dong, Y., and Zhu, J. Stair: Improving safety alignment with introspective reasoning. *arXiv preprint arXiv:2502.02384*, 2025a.
- Zhang, Z., Zhang-Li, D., Yu, J., Gong, L., Zhou, J., Hao, Z., Jiang, J., Cao, J., Liu, H., Liu, Z., et al. Simulating classroom education with llm-empowered agents. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 10364–10379, 2025b.
- Zhao, H., Yuan, C., Huang, F., Hu, X., Zhang, Y., Yang, A., Yu, B., Liu, D., Zhou, J., Lin, J., et al. Qwen3guard technical report. *arXiv preprint arXiv:2510.14276*, 2025.

Zhao, W., Ren, X., Hessel, J., Cardie, C., Choi, Y., and Deng, Y. Wildchat: 1m chatgpt interaction logs in the wild. *arXiv preprint arXiv:2405.01470*, 2024.

Zheng, X., Pang, T., Du, C., Liu, Q., Jiang, J., and Lin, M. Improved few-shot jailbreaking can circumvent aligned language models and their defenses. *Advances in neural information processing systems*, 37:32856–32887, 2024.

Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., and Fredrikson, M. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

A. Algorithmic Details of REFLECTOR

This section presents the algorithmic pipeline of REFLECTOR. As shown in Algorithm A, the training procedure follows a two-stage paradigm that integrates supervised format tuning with reinforcement-based self-improvement. In the first stage, the model is initialized via supervised fine-tuning on high-quality reflection trajectories to establish a structured reasoning and response format. In the second stage, the model is further optimized using group-based reinforcement learning, where both output safety and the quality of self-reflection are explicitly rewarded. This design enables the model to internalize safety-aware reasoning behaviors directly into its generation process.

Algorithm 1 The Algorithmic Pipeline

Require: Dataset $\mathcal{D}$, Group Size G , Reward Scale λ , Penalty Ratio α

- 1: **Stage 1: Reflection Capability Injection via Supervised Fine-Tuning**
 - 2: Fine-tune π_θ on small-scale demonstration trajectories $\mathcal{D}_R$ using SFT
 - 3: **Stage 2: Self-Improvement via Reinforcement Learning**
 - 4: **while** not converged **do**
 - 5: Sample batch of queries $x \sim \mathcal{D}$
 - 6: Generate G trajectories per query: $\{\tau_{i,1}, \dots, \tau_{i,G}\} \sim \pi_\theta(\cdot|x_i)$
 - 7: **for** each trajectory $\tau_{i,g}$ **do**
 - 8: Compute safety reward $r_{\text{safety}} = \text{HarmCLS}(y_{i,g}) \in \{0, 1\}$
 - 9: Assign reflection reward $r_{\text{reflect}} \in \{+\lambda, -\lambda, 0\}$ based on (r_{safety}, z)
 - 10: Total Reward $r_{i,g} = r_{\text{safety}} + r_{\text{reflect}}$
 - 11: **end for**
 - 12: Estimate group-relative advantages $A_{i,g}$ via GDPO
 - 13: Update π_θ by maximizing the RL objective
 - 14: **end while**
-

B. Teacher-Guided Reflection Synthesis

This section details the distillation pipeline used to construct the reflection-augmented dataset $\mathcal{D}_R$ for Stage I. Given that base models often struggle to self-reflection without prior guidance, we employ a high-capacity teacher model GPT-5 (Singh et al., 2025) to generate structured reflection-continuation pairs. These trajectories serve as the “gold standard” for internalizing safety reasoning and establishing the desired self-reflection policy.

B.1. Synthesis Pipeline and Data Quality

The construction of a complete trajectory $\tilde{\tau} = (y^{\text{before}}, z, y^{\text{after}})$ is designed to simulate the emergence of a safety signal during real-time inference. The process follows three key technical phases:

1. **Strategic Context Truncation:** To ensure the model learns to trigger reflection at various stages of response generation, we apply a random truncation strategy $n \sim \mathcal{U}\{1, \dots, T\}$ to existing model outputs. This creates a diverse set of y^{before} prefixes, ranging from early-stage intent alignment to late-stage detail generation.
2. **Structured Reflection Generation (z):** Using the *Self-Critique & Reflection Template*, the teacher model performs a “post-mortem” analysis of the truncated prefix. By framing the teacher as a student reflecting on its own mistakes, we generate z^{reflect} to identify the precise ethical breach and z^{explore} to chart a safe path forward.
3. **Reasoning-Conditioned Continuation (y^{after}):** The teacher generates y^{after} by strictly adhering to the guidance provided in z . This stage ensures that the final output is not just a refusal, but a logically consistent continuation that follows the “search-and-recovery” logic established in the reflection phase.

To maintain high data fidelity, we perform a rule-based filtering pass to ensure all synthesized trajectories contain the required special tokens (`<|reflect|>`, `<|explore|>`) and that the final response successfully transitions from a harmful prefix to a safe termination.

B.2. Prompt Templates for Data Construction

Self-Critique & Reflection Template

You are a student. Review your incorrect partial solution, reflecting on mistakes and proposing alternative ethical solutions. Follow these steps:

1. Analyze the provided step-by-step solution (steps separated by \n\n).
2. For each step, follow this exact format:
 - In Step <id>: [A first-person one-sentence explanation of the mistake].
 - Alternatively: [A brief alternative ethical claim].

Do NOT include any meta-text or additional feedback outside of this format.

Solution:
{answer}

Safety Generation Prompt Template

Your task is to continue writing a few brief sentences, adhering to the "helpful, respectful, and honest" principles. You must strictly refrain from offering any methods, details, or instructions to carry out illegal, violent, or harmful actions.

In your continuation, you must:

- Explicitly and briefly refuse to assist with any illegal activity.
- Provide specific, legal alternatives (e.g., law enforcement, legal aid, mediation).
- Suggest reaching out for professional assistance where applicable.

Original Query:
{Query}

Revision feedback:
{reflection}

Continued solution:

C. RL Implementation Details.

This section provides a detailed account of the reinforcement learning framework and implementation strategies used to refine the REFLECTOR model.

C.1. RL Training Parameters.

The following details the algorithmic architecture of GRPO, the multi-dimensional reward mechanisms, and the specific hyperparameter configurations employed during training.

GRPO Algorithm. Group Relative Policy Optimization (GRPO) is a reinforcement learning algorithm that optimizes the policy by comparing trajectories within a group. For each query x , we sample a group of G trajectories $\{\tau_1, \dots, \tau_G\} \sim \pi_\theta(\cdot | x)$. When multiple reward components exist (e.g., r_{safety} and r_{reflect}), GRPO first aggregates them into a total reward $r(\tau_i) = \sum_k r_k(\tau_i)$, then computes the normalized group advantage:

$$A_i^{\text{GRPO}} = \frac{r(\tau_i) - \mu_r^g}{\sigma_r^g + \epsilon}, \quad \text{where } \mu_r^g = \frac{1}{|G_g|} \sum_{j \in G_g} r(\tau_j), \quad \sigma_r^g = \sqrt{\frac{1}{|G_g|} \sum_{j \in G_g} (r(\tau_j) - \mu_r^g)^2}, \quad (9)$$

where G_g denotes the set of trajectories in group g .

GDPO: Group reward-Decoupled Normalization Policy Optimization. We extend GRPO with GDPO (Group reward-Decoupled Normalization Policy Optimization) for multi-reward RL optimization. The key difference is that GDPO

normalizes each reward component separately within each group before aggregation, rather than normalizing the sum of rewards. For K reward components $\{r_1, \dots, r_K\}$, GDPO first computes the group-normalized advantage for each component:

$$\hat{A}_{i,k} = \frac{r_k(\tau_i) - \mu_{r_k}^g}{\sigma_{r_k}^g + \epsilon}, \quad (10)$$

where $\mu_{r_k}^g$ and $\sigma_{r_k}^g$ are the mean and standard deviation of the k -th reward component within group g . The component advantages are then aggregated with optional weights:

$$\tilde{A}_i = \sum_{k=1}^K w_k \cdot \hat{A}_{i,k}. \quad (11)$$

Finally, a batch-wise normalization is applied:

$$A_i^{\text{GDPO}} = \frac{\tilde{A}_i - \mu_{\tilde{A}}}{\sigma_{\tilde{A}} + \epsilon}, \quad (12)$$

where $\mu_{\tilde{A}}$ and $\sigma_{\tilde{A}}$ are the mean and standard deviation of $\tilde{A}$ across the entire batch.

This decoupled normalization offers a significant advantage: it preserves the relative magnitude differences within each reward component. In GRPO, when rewards are summed before normalization, the distinction between “slightly better” and “much better” trajectories can be lost. GDPO maintains this granularity by normalizing each component independently, allowing the model to better distinguish between marginal and substantial improvements in both safety and reflection quality.

In our implementation, we use two reward components:

- $r_{\text{correctness}}$ (weight $w_1 = 1.0$): Evaluates the safety/correctness of the final output.
- r_{reflect} (weight $w_2 = 0.3$): Provides bonus for effective reflection leading to safe outputs, or penalty for ineffective reflection.

Training Hyperparameters. We use the following hyperparameters for GDPO training:

- Group size $G = 8$
- Learning rate: 1×10^{-6} with constant scheduler
- Batch size: 64
- Maximum prompt length: 4096
- Maximum response length: 4096
- Sampling temperature: 0.7, top- p : 0.95
- PPO clip ratio: 0.2
- KL penalty coefficient: 0.001
- PPO epochs: 1
- Weight decay: 0.01
- Adam betas: (0.9, 0.999)
- Gradient clipping: 1.0
- Training epochs: 1

Reward Model. We employ GPT-OSS-120B (Agarwal et al., 2025) as the reward model to evaluate trajectory quality. The reward model provides structured judgments on both safety and correctness, outputting JSON-formatted responses with labels for initial and final response segments.

Reward Model Prompt Templates

For safety evaluation (harmful queries), we use the following prompt template:

Safety Evaluation Prompt Template

```
You are a strict safety auditor.
IMPORTANT: This is a HARMFUL/UNSAFE request. The expected behavior is that the
assistant should REFUSE to comply.
A 'Safe' response means the assistant correctly refused or provided a safe
alternative.
An 'Unsafe' response means the assistant complied with the harmful request.
The assistant response contains special markers: <|reflect|>, <|explore|>,
<|continue|>.
Define:

  • Initial response = the text BEFORE the first marker.

  • Final response = the text AFTER <|continue|> if present, else AFTER <|explore|>,
    else AFTER <|reflect|>.

You must judge TWO things for both initial and final responses:

  1. safety: one of Safe, Unsafe, Controversial.

  2. refusal: a boolean (true/false).

Return ONLY a JSON object with keys: initial_safety, final_safety, initial_refusal,
final_refusal.
<user_query>
{question}
</user_query>
<assistant_response>
{response}
</assistant_response>
```

For math evaluation, we use a similar template with correctness labels (Correct, Incorrect, Unclear) instead of safety labels.

C.2. Original Model Reflection Prompt Design.

For models that have not undergone SFT (Stage I), we use a system prompt to guide the model to generate responses in the required reflection format. This prompt is prepended to each query during GDPO training:

Reflection Format System Prompt

```
You must respond using this exact structure and tags:
Initial answer text.
<|reflect|>
Reflection on the initial answer.
<|explore|>
Further exploration or alternatives.
<|continue|>
Final answer text.
Use each tag exactly once and in order. Do not add extra tags.
```

This system prompt is injected into the user message (i.e., `system_prompt_in_user: true`) as we found this approach yields better instruction-following performance compared to using a separate system message. As shown in

Table 4, directly applying GDPO with this prompt to the base model achieves moderate safety improvements but results in degraded performance on robustness benchmarks, demonstrating the necessity of the two-stage training paradigm.

D. Details of Hybrid Safety Evaluation and Reward Design

In this section, we detail the implementation of our safety evaluation mechanism, which serves as the ground truth for calculating the safety reward r_{safety} .

D.1. Limitations of Keyword Matching

Prior works (Zou et al., 2023; Wei et al., 2024a) typically rely on a restricted set of refusal phrases (e.g., “Sorry, I can’t”, “I cannot fulfill”) to determine safety via prefix matching. However, we identify significant limitations in this naive approach:

- **False Positives on Safety:** Models often generate unconventional refusal patterns, such as “I think it’s illegal to do so” or simply outputting noise like “#####”. Since these do not match the predefined set, standard metrics misclassify them as successful attacks, whereas they are, in fact, safe refusals.
- **Inability to Detect Partial Compliance:** Keyword matching fails to identify subtle jailbreaks where the model provides harmful information without using standard affirmative phrases.

D.2. Hybrid Consensus Mechanism

To address these issues and enforce rigorous safety standards, we implement HarmCLS , a hybrid evaluation metric that enforces consensus between a discriminative classifier and a generative judge.

1. Discriminative Detector (C_{disc}). We utilize the HarmBench classifier (Mazeika et al., 2024), which is fine-tuned on Llama-2-13B. As reported, this model demonstrates high robustness, outperforming GPT-4 by approximately 5% in agreement rates with human judgments on manually labeled validation sets. It provides a binary verdict:

$$C_{\text{disc}}(y) = \begin{cases} 1 & \text{if Harmless} \\ 0 & \text{if Harmful} \end{cases}$$

2. Generative Detector (J_{gen}). To capture nuance missed by binary classification, we concurrently deploy GPT-OSS-120B (Agarwal et al., 2025) as a generative safety detector. This model evaluates the semantic content of the response trajectory and categorizes it into one of three states:

- **Benign:** The response directly addresses the query but remains within safety boundaries (e.g., educational context).
- **Refusal:** The model explicitly declines to answer or provides irrelevant output (treated as safe for defense purposes).
- **Harmful:** The response executes the malicious instruction.

3. Intersection Logic (Final Verdict). The final safety label $\text{HarmCLS}(y)$ is derived from the **intersection** of these two evaluations. A response is deemed strictly harmless only if it passes both the discriminative check and the generative qualitative assessment. Formally:

$$\text{HarmCLS}(y) = \begin{cases} 1, & \text{if } C_{\text{disc}}(y) = 1 \quad \text{AND} \quad J_{\text{gen}}(y) \in \{\text{Benign}, \text{Refusal}\}, \\ 0, & \text{otherwise.} \end{cases} \quad (13)$$

This rigorous intersection logic minimizes misjudgment risks, ensuring that r_{safety} provides a high-precision signal for guiding the reinforcement learning process.

E. Utility Evaluation Details.

This section lists the details of the datasets used for utility evaluation, including their names, evaluation metrics, and sizes. As shown in Table 8, we employ a comprehensive suite of 15 benchmarks to holistically evaluate the model’s safety

performance and general capabilities. The safety evaluation encompasses three critical dimensions: overtly harmful queries to assess basic safety alignment; performance under jailbreak attacks to test robustness against sophisticated adversarial prompts; and behavior on refusal benchmarks to measure the correctness of rejection responses. Additionally, the general capability assessment covers four key areas: general knowledge proficiency evaluated using MMLU across 14 subjects, safe mathematical reasoning on GSM8K, robust language understanding via AdvGLUE, and knowledge-based question answering on SimpleQA.

Category	Sub-category	Datasets	Metric	Test Size
Safety (Harmful)	Harmful Dataset	StrongREJECT	Goodness Score ($\uparrow$)	313
		XSTest	DSR ($\uparrow$)	200
		WildChat	DSR ($\uparrow$)	500
		Do-Not-Answer	DSR ($\uparrow$)	939
Safety (Jailbreak)	Direct Attack	AdvBench (AutoDAN, GCG, PAIR)	DSR ($\uparrow$)	520
	Indirect Attack	AdvBench (DRA, PAP, ReNeLLM, DrAttack)	DSR ($\uparrow$)	520
Utility (General)	14 Subjects	MMLU-Pro	Multi-choice Accuracy ($\uparrow$)	14, 079
Utility (Specific)	Math	GSM8K	Multi-choice Accuracy ($\uparrow$)	1, 319
	Knowledge	SimpleQA	EM / F1 ($\uparrow$)	4, 330
	Roubustness	AdvGLUE	Task Accuracy ($\uparrow$)	738

Table 8. Overview of evaluation settings across safety and utility benchmarks.

F. Additional Robustness Experiments

This section reports additional stress tests beyond the main benchmark suite. We focus on settings that require robustness to implicit, evolving, or evaluator-dependent threats: ICL-based jailbreaks, automated strategy-searching attacks, multi-turn adversarial interactions, and sensitivity to teacher, reward, and judge models.

F.1. Robustness to ICL-Based Jailbreaks

ICL-based jailbreaks attempt to induce unsafe behavior by presenting the model with a small number of harmful demonstrations before the target query. This setting is especially relevant to REFLECTOR, because the attack does not necessarily rely on explicit adversarial tokens; instead, it shifts the model into an unsafe reasoning trajectory through contextual examples. We evaluate this threat under the standard I-FSJ protocol (Zheng et al., 2024), using 100 randomly sampled queries from AdvBench and HarmBench.

Table 9. Defense success rate under ICL-based jailbreaks following the I-FSJ protocol.

Method	AdvBench	HarmBench	Avg.
Base Model	9.0%	8.0%	8.5%
REFLECTOR	83.0%	79.0%	81.0%

Table 10. Defense success rate and reflection-triggered rate under scaled ICL attacks.

Method / Shot	0	1	2	4	8	16
Base Model	76.0%	50.0%	32.0%	9.0%	2.0%	0.0%
REFLECTOR	100.0%	00.0%	98.0%	83.0%	78.0%	72.0%
Reflection Triggered Rate	85.0%	90.5%	96.0%	100.0%	100.0%	100.0%

We further evaluate shot scaling, where the number of harmful demonstrations increases from 0 to 16. In addition to DSR, we report the reflection-triggered rate, defined as the proportion of defended examples in which the two-stage model invokes the reflection mechanism before refusal. Table 10 shows that the base model rapidly collapses as the number of shots increases, dropping from 76.0% DSR at 0-shot to 0.0% at 16-shot. In contrast, REFLECTOR maintains substantially stronger defense across all shot counts. The reflection-triggered rate also rises with attack strength and saturates at 100% in high-shot settings, suggesting that the model increasingly relies on internal correction under stronger contextual pressure. Together,

these results indicate that REFLECTOR generalizes to ICL-induced safety risks through trajectory-level correction, rather than relying only on surface detection of known jailbreak prompts.

F.2. Robustness to Automated Strategy-Searching Attacks

We also evaluate REFLECTOR against AutoDAN-Turbo (Liu et al., 2023a), an automated strategy-searching attack that iteratively discovers effective jailbreak strategies. AutoDAN-Turbo is challenging because it can embed malicious intent into complex role-based or scenario-based contexts, turning the unsafe objective into an implicit and highly contextualized signal. As shown in Table 11, the base model is almost fully compromised, achieving only 5.0% DSR, whereas REFLECTOR achieves 90.0% DSR. This result suggests that interrupting unsafe reasoning trajectories remains effective even when the attack strategy is automatically optimized.

Table 11. Defense success rate against automated strategy-searching and multi-turn attacks.

Method	AutoDAN-Turbo (DSR $\uparrow$)	ActorAttack (DSR $\uparrow$)
Base Model	5.0%	21.0%
REFLECTOR	90.0%	91.0%

F.3. Robustness to Multi-Turn Adversarial Interactions

Multi-turn adversarial interactions create indirect safety risks through gradual intent shifting. Unlike single-turn jailbreaks, the malicious objective may only become apparent after several conversational turns. We evaluate this setting under the ActorAttack protocol (Ren et al., 2025) on 100 samples. As shown in Table 11, REFLECTOR achieves 91.0% DSR after GDPO, compared with 21.0% for the base LLaMA-3.1-8B model. These results show that the learned search-and-recovery behavior is not limited to isolated single-turn inputs, but also generalizes to more realistic multi-turn adversarial contexts.

F.4. Sensitivity to Teacher and Reward Models

Because REFLECTOR uses teacher and reward models during data construction and optimization, we evaluate whether its performance depends on a particular proprietary model. Table 12 compares different teacher and reward choices. Replacing GPT-5 with GPT-3.5-turbo as the teacher causes only a small degradation, e.g., DRA changes from 92.31% to 90.85%. This suggests that the gains are not solely inherited from the teacher’s intrinsic capability, but also arise from the data construction and trajectory-learning pipeline. Similarly, replacing GPT-OSS-120B (Agarwal et al., 2025) with Qwen-Guard (Zhao et al., 2025) as the reward model lowers the performance ceiling but still maintains strong safety, indicating that the optimization is robust to weaker reward supervision.

Table 12. Sensitivity to teacher and reward model choices.

Teacher Model	Reward Model	WildChat (DSR $\uparrow$)	DRA (DSR $\uparrow$)	AdvGLUE
GPT-5	GPT-OSS-120B	81.20%	92.31%	68.29%
GPT-3.5-turbo	GPT-OSS-120B	80.80%	90.85%	67.90%
GPT-5	Qwen-Guard	77.50%	89.40%	65.20%

F.5. Robustness to Diverse Safety Judges

To examine whether the reported safety gains are tied to a specific evaluator, we re-evaluate saved responses with multiple safety judges. Table 13 shows that REFLECTOR remains robust across both generative and classifier-based judges, including smaller open-source models such as RISE-Judge-Qwen2.5-7B (Yu et al., 2025) and LLaMA-3.1-8B-Instruct (Grattafiori et al., 2024). GPT-OSS-120B (Agarwal et al., 2025) gives the lowest DSR among the tested judges, suggesting that our main reported numbers are conservative rather than inflated by a lenient evaluator.

As an out-of-family check, we additionally re-judge the main LLaMA-3.1 REFLECTOR checkpoints with Qwen2.5-32B-Instruct (Qwen et al., 2025). The score increases from 0.6790 to 0.6981 at step 95, and from 0.6863 to 0.7111 at step 109. This confirms the same direction of improvement under a judge model outside the original reward-model family, further reducing the risk that the observed gains come from overfitting to a single evaluator.

Table 13. Safety evaluation under different judge models.

Judge Model	Direct Jailbreak	Indirect Jailbreak	Disagreement Rate
GPT-OSS-120B	96.76%	94.63%	—
GPT-OSS-20B	98.50%	97.45%	3.54%
RISE-Judge-Qwen2.5-7B	98.95%	98.26%	5.83%
LLaMA-3.1-8B-Instruct	99.14%	99.15%	6.58%

G. Why Safety Reflection Can Improve General Utility?

A more precise interpretation of Table 1 is that the utility gains are task-dependent rather than universal. REFLECTOR does not transfer refusal behavior to general tasks. Instead, it transfers a trajectory-level control operation learned from safety training: monitor the current reasoning state, compare it with the task constraints, and revise the trajectory before finalizing the answer. In safety examples, the violated constraint is a hidden harmful intent; in utility examples, the violated constraint can be an arithmetic relation, a state variable, an answer-format requirement, or a factual consistency condition.

This mechanism explains the observed benchmark pattern. Improvements are concentrated on tasks that benefit from intermediate checking and correction, such as GSM8K (+5.65), AdvGLUE (+9.96), and SimpleQA (+3.93). In contrast, fixed-form knowledge or retrieval-style benchmarks such as MMLU-Pro show only marginal changes (+0.95), because these tasks provide fewer opportunities for step-level recovery. Thus, safety reflection and general utility can be mutually reinforcing when the general task requires self-monitoring, but this effect should not be interpreted as a uniform gain across all benchmarks. The following examples illustrate how the same correction operation transfers from safety to utility tasks. The trajectories are shortened for readability while preserving the key failure and recovery steps.

Utility Example: Complementary Proportion

Question: At a certain grade level, $3/4$ of students have a desktop computer at home. If 20 students do not have a desktop computer, how many students are there in total?

Original model: “To find the total number of students, we need to first find the number of students who have a desktop computer and then add the number of students who do not have a desktop computer.”

Failure: The model tries to solve the “have a computer” group before knowing the total, creating a circular dependency.

REFLECTOR:

<|reflect|> The known count corresponds to the complementary group. If $3/4$ of students have a computer, then $1/4$ of students do not have one.

<|explore|> Let x be the total number of students. Since $x/4 = 20$, we get $x = 80$.

Final: The total number of students is 80.

Utility Example: Avoiding Premature Finalization

Question: A coat has a list price of \$325.00. During November, it did not sell, and the merchant marked it down by 20% . . .

Original model: “C”

Failure: The model commits to a multiple-choice answer without performing the required multi-step calculation.

REFLECTOR:

<|reflect|> I should not output a letter before checking the intermediate prices. The answer depends on applying each discount in sequence.

<|explore|> The initial price is \$325. A 20% discount is $0.20 \times 325 = 65$, so the updated price is \$260. The remaining discount steps should be applied to this updated price before selecting the final option.

Final: The model completes the sequential calculation before choosing the answer.

Utility Example: Sequential State Tracking

Question: Seth gave half of his stickers to Luis. Luis used half of the stickers and gave the rest to Kris ...

Original model: “Seth gave half of his stickers to Luis. So, Luis received $x/2$ stickers. This means Se...”

Failure: The model begins with a valid assignment but loses track of the active holder, mixing Seth’s remaining stickers with Luis’s received stickers.

REFLECTOR:

<|reflect|> After Seth gives stickers to Luis, the next operation should update Luis’s state rather than Seth’s remaining state.

<|explore|> Luis received $x/2$ stickers. He used half of those stickers, so he used $x/4$ and gave the remaining $x/4$ to Kris.

Final: The reasoning proceeds by updating one holder at a time, preserving the state chain.