

Position: LLMs Should Incorporate Explicit Mechanisms for Human Empathy

Xiaoxing You^{*1} Qiang Huang^{*1} Jun Yu¹

Abstract

This position paper argues that Large Language Models (LLMs) should incorporate explicit mechanisms for human empathy. As LLMs become increasingly deployed in high-stakes human-centered settings, their success depends not only on correctness or fluency but on faithful preservation of human perspectives. Yet, current LLMs systematically fail at this requirement: even when well-aligned and policy-compliant, they often attenuate affect, misrepresent contextual salience, and rigidify relational stance in ways that distort meaning. We formalize empathy as an observable behavioral property: the capacity to model and respond to human perspectives while preserving intention, affect, and context. Under this framing, we identify four recurring mechanisms of empathic failure in contemporary LLMs—sentiment attenuation, empathic granularity mismatch, conflict avoidance, and linguistic distancing—arising as structural consequences of prevailing training and alignment practices. We further organize these failures along three dimensions: cognitive, cultural, and relational empathy, to explain their manifestation across tasks. Empirical analyses show that strong benchmark performance can mask systematic empathic distortions, motivating empathy-aware objectives, benchmarks, and training signals as first-class components of LLM development. Our code is available at https://github.com/youxiaoxing/LLM_Empathy.

1. Introduction

Large Language Models (LLMs) are no longer experimental assistants. They are deployed at scale as interfaces for a wide range of applications, including news production (To-

^{*}Equal contribution ¹Harbin Institute of Technology (Shenzhen). Correspondence to: Jun Yu <yujun@hit.edu.cn>.




Figure 1. Four recurring mechanisms underlying empathic failure in contemporary LLM outputs. (a) Sentiment Attenuation: the model replaces emotionally charged language with neutral, official wording; (b) Empathic Granularity Mismatch: the model retains graphic detail where calibration is needed; (c) Conflict Avoidance: the model bypasses emotional tension and jumps directly to task-oriented advice; (d) Linguistic Distancing: the translation flattens the poem’s imagery into plain prose.

hidi et al., 2025; Liu et al., 2025c; Xu et al., 2025c; Odabaşı & Biricik, 2025; Sinha, 2025; Sun et al., 2024; 2025c; Sahu et al., 2025), education (Kim et al., 2025c; Breen, 2025; Park et al., 2025a), and psychological support (De Duro et al., 2025; Wang et al., 2025a; Xue et al., 2026). In these settings, correctness and fluency are necessary but insufficient. Increasingly, success depends on whether model outputs faithfully represent *human perspectives*—including what users intend, feel, and consider salient in context (Sorin et al., 2024; Lee et al., 2024b; Welivita & Pu, 2024; Sun et al., 2025b;a;d).

We argue that LLMs should incorporate explicit mechanisms for human empathy. Empathy, as we define it, is not an internal mental state or stylistic embellishment, but an *observable behavioral property*: the capacity to model and respond to human perspectives while preserving intention, affect, and contextual requirements in generated outputs. When this capacity is absent, models may remain helpful, safe, and policy-aligned, yet still distort meaning in ways that materially affect real-world outcomes.

This limitation is not incidental. Modern LLMs (OpenAI,

2025; Anthropic, 2025; Pichai et al., 2025; Liu et al., 2025a; Bai et al., 2025a; Hao et al., 2026) are trained through large-scale pretraining for broad competence, followed by post-training alignment via instruction tuning (Wei et al., 2022; Ouyang et al., 2022; Feng et al., 2025b) and preference optimization (Schulman et al., 2017; Rafailov et al., 2023; Shao et al., 2024; Hong et al., 2024; Zhou et al., 2024) to improve helpfulness, safety, and risk mitigation. While effective at producing fluent and compliant outputs, these objectives systematically favor calmness and neutrality, often suppressing necessary affect, perspective specificity, and relational engagement. Models may therefore satisfy alignment criteria while misrepresenting human perspectives where fidelity matters most.

To make these failure modes explicit, we summarize four recurring mechanisms through which current LLMs exhibit *deficient empathy*. Specifically, we identify:

- **Sentiment Attenuation (SA)**, where emotional intensity or urgency is systematically downscaled (Bardol, 2025);
- **Empathic Granularity Mismatch (EGM)**, where the level of affective or experiential detail is miscalibrated (Manevich et al., 2023);
- **Conflict Avoidance (CA)**, where legitimate tensions are deflected in favor of superficial harmony (Sharma et al., 2024);
- **Linguistic Distancing (LD)**, where depersonalized framing increases interpersonal distance (Nook et al., 2017).

As illustrated in Figure 1, these mechanisms are not edge cases: they consistently emerge under current training and alignment regimes across application domains (Kirk et al., 2024; Xu et al., 2025b).

To explain how these mechanisms manifest differently across tasks, we characterize empathy along three dimensions: **cognitive**, **cultural**, and **relational empathy**. Cognitive empathy governs perspective-taking, and the preservation of informational salience (Cuff et al., 2016; Healey & Grossman, 2018; Reniers et al., 2022); cultural empathy captures sensitivity to norms, symbols, and aesthetic context (Cuff et al., 2016; Eichbaum et al., 2023); and relational empathy supports trust, cooperation, and constructive engagement over interaction (Eisenberg et al., 2010; Cuff et al., 2016; Weisz & Zaki, 2018; Van Dijke et al., 2020). This decomposition clarifies why empathic failures vary across tasks while sharing common underlying causes.

We substantiate this position through empirical analyses on representative high-stakes tasks, including news generation, literary translation, and counseling-style response generation (Grusky et al., 2018; Wang et al., 2023b;a). Across all settings, frontier LLMs exhibit consistent empathic distortions despite strong performance on conventional metrics. These failures reveal alignment trade-offs that remain invis-

ble under existing evaluation frameworks.

Based on these findings, we outline concrete directions for integrating empathic fidelity into LLM development pipelines. We advocate for empathy-aware alignment objectives, task-grounded benchmarks that expose mechanism-level failures, and targeted data curation that supervises perspective preservation, affective calibration, and relational stance. Taken together, this position reframes empathy from a stylistic byproduct into a measurable and optimizable property of aligned language models.

2. Problem Formulation

2.1. Defining Empathy for LLMs

In psychology, empathy broadly refers to perspective-taking and context-sensitive responding in social interaction (Singer & Lamm, 2009; Zaki, 2014; Levett-Jones et al., 2019). For LLMs, we define empathy as an *observable behavioral property*: the capacity to model and respond to human perspectives while preserving intention, affect, and contextual requirements in generated outputs.

This framing treats empathy as a form of representational fidelity in language generation. Beyond factual correctness and task completion, empathetic outputs should preserve the *salience structure* and *pragmatic implications* that matter to users, while appropriately calibrating affective granularity and relational framing to the interaction context (Schulman et al., 2017; Rafailov et al., 2023; Köpf et al., 2023; Shao et al., 2024; Hong et al., 2024). Here, *salience structure* denotes the most important affective, contextual, or relational cues in the input, while *pragmatic implications* capture the implicit needs, concerns, or interactional goals beyond the literal meaning. Such fidelity is especially critical in high-stakes applications, including mental health support (Lozoya et al., 2025; Xue et al., 2026), news summarization (You et al., 2025; Odabaşı & Biricik, 2025), and workplace communication (Navarro et al., 2025; Miura et al., 2025), where outputs may be factually correct yet fail to convey urgency, misrepresent perspectives, or adopt inappropriate interpersonal distance.

2.2. Four Mechanisms of Empathic Failure

Despite producing fluent and policy-compliant outputs, LLMs often fail to capture what humans *mean*, *feel*, and *need* in context (Sorin et al., 2024; Chen et al., 2024; He et al., 2025; Kim et al., 2025a). We identify four recurring mechanisms underlying such empathic deficits.

Notation Let n denote the number of evaluated instances. For each instance $i \in \{1, \dots, n\}$, let $t^{(i)}$ denote the reference text and $g^{(i)}$ the model-generated output.

Sentiment Attenuation (SA) SA describes the systematic

Table 1. Definitions, ranges, and interpretations of the four empathic failure metrics.

Metric	Range	Interpretation
Sentiment Attenuation (SA)	$(-\infty, 1]$	0: no attenuation; 1: strong affect flattening; negative: more intense than reference
Empathic Granularity Mismatch (EGM)	$[0, 1]$	0: close alignment in detail density; 1: severe mismatch
Conflict Avoidance (CA)	$[0, 1]$	0: no conflict avoidance; 1: all responses avoid conflict
Linguistic Distancing (LD)	$[0, +\infty)$	0: no distancing difference; larger: more abstract and depersonalized

reduction of emotional intensity, urgency, or moral weight in model outputs, even when the underlying facts are preserved. It commonly arises when LLMs reframe human experiences into “neutral” summaries that appear calm, balanced, and professionally flattened (Bardol, 2025; Parmar & Mazumdar, 2025). This pattern likely stems from alignment procedures that implicitly penalize strong affective language (Kirk et al., 2024), suppressing emotional intensity even when it is semantically necessary.

We quantify this effect using the *sentiment attenuation rate*:

$$SA = \frac{1}{n} \sum_{i=1}^n \frac{E(\mathbf{t}^{(i)}) - E(\mathbf{g}^{(i)})}{E(\mathbf{t}^{(i)})} \quad (1)$$

where $E(\cdot) \in [0, 1]$ denotes an affect intensity scorer; 0 indicates emotionally flat and affectively neutral language, and 1 indicates highly intense and vividly expressed emotion.

Empathic Granularity Mismatch (EGM) EGM captures miscalibration in the level of affective or experiential detail expressed in outputs, particularly in sensitive or distressing contexts (Manevich et al., 2023). Models may produce overly vague abstractions or unnecessarily explicit descriptions (Peters & Chin-Yee, 2025), a failure especially visible in news summarization involving violence or trauma (Zhang et al., 2023b; Sahu et al., 2025). This miscalibration arises because pretraining emphasizes language modeling and commonsense acquisition (Wei et al., 2022; Porada et al., 2022), while post-training targets alignment and reasoning (Shao et al., 2024; Hong et al., 2024), leaving descriptive granularity underspecified.

We measure empathic granularity mismatch via a *granularity mismatch score*, which is defined as:

$$EGM = \frac{1}{n} \sum_{i=1}^n \left| \frac{a^{(i)} - b^{(i)}}{a^{(i)} + b^{(i)}} \right|, \quad (2)$$

where $a^{(i)} = \frac{U(\mathbf{t}^{(i)})}{d^{(i)}}$ and $b^{(i)} = \frac{U(\mathbf{g}^{(i)})}{q^{(i)}}$ denote the densities of disturbing or overly explicit content in the reference text $\mathbf{t}^{(i)}$ and model output $\mathbf{g}^{(i)}$, respectively; $U(\cdot)$ is a discomfort counter; and $d^{(i)}$ and $q^{(i)}$ are sentence counts.

Conflict Avoidance (CA) CA refers to the tendency to minimize, obscure, or prematurely resolve interpersonal tension rather than engaging with it directly (Sharma et al., 2024; Cheng et al., 2025). A typical manifestation is premature solution-giving, where models bypass clarification of competing perspectives and trade-offs. This behavior is likely

reinforced by alignment objectives that reward agreeableness and controversy reduction (Röttger et al., 2024).

We measure this tendency via the *conflict avoidance rate*, which is the proportion of instances where a model prematurely resolves tension through solution-giving or harmony-seeking, rather than engaging in clarification or perspective-specific inquiry. Formally,

$$CA = \frac{1}{n} \sum_{i=1}^n s^{(i)}, \quad (3)$$

where $s^{(i)} = J(g^{(i)}) \in \{0, 1\}$ is a binary indicator produced by a conflict-avoidance discriminator $J(\cdot)$, with $s^{(i)} = 1$ denoting conflict avoidance and $s^{(i)} = 0$ denoting faithful engagement with the underlying tension.

Linguistic Distancing (LD) LD refers to the use of abstract, indirect, or depersonalized language, such as passive constructions, nominalizations, and hedging, that reduces emotional immediacy (Danescu-Niculescu-Mizil et al., 2013; Nook et al., 2017; Minnema et al., 2022). This manifests as a shift from agent-centered narratives to procedural descriptions that obscure agency and accountability, likely driven by alignment objectives favoring neutral, professional registers that minimize affective or evaluative directness (Liu et al., 2023c; Huang et al., 2024b).

We define the *linguistic distancing score* as:

$$LD = \frac{1}{n} \sum_{i=1}^n \left\| \frac{\mathbf{d}_t^{(i)} - \boldsymbol{\mu}_t}{\boldsymbol{\sigma}_t} - \frac{\mathbf{d}_g^{(i)} - \boldsymbol{\mu}_g}{\boldsymbol{\sigma}_g} \right\|_2, \quad (4)$$

where $\mathbf{d}_t^{(i)}, \mathbf{d}_g^{(i)} \in \mathbb{R}^3$ are distancing-density vectors for the reference text $\mathbf{t}^{(i)}$ and model output $\mathbf{g}^{(i)}$, computed along nominalization, impersonal subject framing, and agentless constructions; and $(\boldsymbol{\mu}_t, \boldsymbol{\sigma}_t)$ and $(\boldsymbol{\mu}_g, \boldsymbol{\sigma}_g)$ are the corresponding means and standard deviations used for standardization.

Remarks These four mechanisms provide a structured lens for analyzing empathic failures in current LLMs. They are not mutually exclusive and may co-occur within a single response. Table 1 summarizes the range and interpretation of each metric. Explicitly identifying them enables targeted evaluation and informs principled interventions to improve empathic fidelity in aligned LLM systems.

3. Empathy Dimensions

Having defined empathy and identified four failure mechanisms in Section 2, we now characterize how empathic




Figure 2. **Taxonomy of empathy-critical LLM applications**, organized along cognitive, cultural, and relational dimensions, with representative task families and recent research illustrating each dimension.

deficits affect real-world LLM deployments (Sorin et al., 2024). We organize human-facing requirements along three complementary dimensions: **cognitive**, **cultural**, and **relational empathy**, which govern factual salience, contextual meaning, and interaction dynamics, respectively. In this section, we define each dimension and illustrate it through representative task examples (Figure 2).

3.1. Cognitive Empathy

Cognitive empathy denotes an LLM’s ability to model a human perspective, including beliefs, intentions, and knowledge state, and to preserve what is salient and meaningful from that perspective in generated outputs (Reniers et al., 2022). In human-facing tasks, faithful generations require

not only factual correctness, but also preservation of contextual salience and implications as recognized by users (Hu et al., 2025c; Kim et al., 2025c; Shi et al., 2025; Liu et al., 2025c; Wang et al., 2025c).

Without cognitive empathy, models may produce factually accurate outputs that nonetheless misrepresent urgency, distress, or priority, thereby distorting user-relevant meaning. This dimension is therefore critical in tasks where language mediates public understanding or resource allocation. Representative tasks include:

- **Text Summarization:** Accurate summaries may still suppress urgency or affective significance if salience is misrepresented (Zhang et al., 2024b; Xu et al., 2025a; Yuan & Zhang, 2025; Odaşası & Biricik, 2025; Sinha, 2025).
- **History Education:** Modeling historical actors' motives and constraints prevents presentist reinterpretation (Lim, 2025; Bungum et al., 2025; Du et al., 2025; Núñez-Naranjo & Morales-Urrutia, 2025; Breen, 2025).
- **Welfare Assessment:** Eligibility judgments require preserving claimants' contextual constraints beyond generic criteria (Dong et al., 2025; Coz et al., 2025; Shi et al., 2025; Karr Jr et al., 2025; Kong et al., 2026).
- **News Generation:** Reporting sensitive events demands calibrated detail to preserve human impact without distortion (Diakopoulos et al., 2024; Tohidi et al., 2025; Liu et al., 2025c; Tseng et al., 2025; Xu et al., 2025c).

Across these tasks, failures in cognitive empathy commonly manifest as **sentiment attenuation**, **empathic granularity mismatch**, and **linguistic distancing**, leading to outputs that remain factually correct yet misrepresent what is salient, urgent, or meaningful from the human perspective.

3.2. Cultural Empathy

Cultural empathy captures an LLM's ability to attend to norms, cultural symbols, and contextual conventions, preserving cultural meaning rather than defaulting to homogenized expression (Eichbaum et al., 2023). In creative and interpretive tasks, output quality depends not only on semantic accuracy but also on stylistic nuance, implicit values, and cultural specificity (Zhang et al., 2025a; Gurung & Lapata, 2025; Sands et al., 2025).

Without cultural empathy, outputs become stylistically flat, replacing distinctive voices with generic phrasing that erases cultural identity and artistic intent. This dimension is central in tasks where language mediates aesthetic experience and cultural identity. Representative tasks include:

- **Literary Translation:** Preserving register, metaphor, and culturally grounded connotations (Macken, 2024; Wu et al., 2025; Zhang et al., 2025b;a).
- **Literary Generation:** Maintaining coherent style and avoiding generic narrative templates (Feng et al., 2025a;

Teleki et al., 2025; Xia et al., 2025; Yang et al., 2025a; Gurung & Lapata, 2025; Wang et al., 2025c).

- **Art Criticism:** Interpreting symbolism and historical context beyond template-based evaluation (Bin et al., 2024; Qi et al., 2025a; Banting, 2025; Khadangi et al., 2025; Arita et al., 2025; Sands et al., 2025).

Cultural empathy failures primarily arise through **empathic granularity mismatch** and **linguistic distancing**, producing outputs that are fluent but culturally impoverished.

3.3. Relational Empathy

Relational empathy concerns an LLM's ability to adapt responses to interaction partners and maintain constructive relational dynamics over time (Van Dijke et al., 2020). Unlike cognitive and cultural empathy, which focus on representation, relational empathy emphasizes interactional calibration, supporting trust, cooperation, and psychological safety (Main et al., 2017; Eichbaum & Bleakley, 2025).

Without relational empathy, interactions become rigid or misaligned with users' social and emotional needs, undermining engagement and task effectiveness. This dimension is essential in collaborative settings where outcomes depend on sustained user engagement and psychological safety. Representative tasks include:

- **Psychological Counseling:** Balancing affective validation with appropriate boundaries (Sharma et al., 2021; Wang et al., 2025d; Lozoya et al., 2025; Zeng, 2025; Zhao et al., 2025; De Duro et al., 2025; Xue et al., 2026).
- **Workplace Communication:** Navigating hierarchy, face-saving, and professional norms without excessive softening or deflection (Lu et al., 2024; Miura et al., 2025; Navarro et al., 2025; Li et al., 2025b; Hu et al., 2025b; Singh et al., 2025).
- **Conflict Mediation:** Engaging disagreement constructively without premature harmony-seeking (Govers et al., 2024; Tessler et al., 2024; Chen et al., 2025; Juvino Santos et al., 2025; Sakurai et al., 2025; Li et al., 2025a).

Failures in relational empathy typically manifest as **conflict avoidance** and **linguistic distancing**, yielding responses that appear polite yet undermine trust and collaboration.

Remarks These three dimensions characterize empathic behavior in LLMs: cognitive empathy preserves user-relevant salience, cultural empathy preserves contextual meaning, and relational empathy preserves cooperative dynamics. Table 2 maps representative tasks to empathic failure mechanisms, guiding the task-grounded evaluations that follow.

4. Evidence

To empirically ground our position, we examine whether frontier LLMs exhibit systematic empathic failures despite

Table 2. Mapping from empathy dimensions and representative tasks to empathy-reducing mechanisms, showing how cognitive, cultural, and relational empathy failures manifest through specific mechanisms across real-world application scenarios.

Dimension	Representative Task	Mechanism			
		SA	EGM	CA	LD
Cognitive Empathy	Text Summarization	✓	✓	×	×
	History Education	×	✓	×	✓
	Welfare Assessment	✓	×	×	✓
	News Generation	✓	✓	×	×
Cultural Empathy	Literary Translation	✓	×	×	✓
	Literary Generation	×	✓	✓	×
	Art Criticism	×	✓	×	✓
Relational Empathy	Psychological Counseling	✓	×	✓	×
	Workplace Communication	×	×	✓	✓
	Conflict Mediation	×	×	✓	✓

strong performance on standard benchmarks. We evaluate ten state-of-the-art models (Appendix A) across three high-stakes tasks that instantiate distinct empathy dimensions: **news generation** (cognitive empathy), **literary translation** (cultural empathy), and **psychological counseling** (relational empathy). For each task, we report standard quality metrics alongside mechanism-based empathy metrics to expose failures masked by standard evaluation.

4.1. Cognitive Empathy: News Generation

Setup We evaluate cognitive empathy using the **Newsroom** dataset (Grusky et al., 2018), which pairs news articles with human-written summaries. Given a gold summary, models are prompted to generate an article of comparable length. To stress informational salience, we focus on high-impact news involving clear human consequences. An LLM filters the test split, from which we randomly sample 50 instances. We evaluate generation quality using three standard metrics: **ROUGE** (Lin, 2004), **CIDEr** (Vedantam et al., 2015), and Semantic Similarity (**Sem-Sim**), where Sem-Sim is measured via Qwen3-4B-Embedding model (Zhang et al., 2025c). Beyond conventional metrics, we also report Sentiment Attenuation (**SA**) and Empathic Granularity Mismatch (**EGM**) to assess cognitive empathy failures. Full details are provided in Appendix B.

Key Findings Table 3 summarizes news generation performance under standard quality metrics and our cognitive empathy measures. Frontier models perform strongly, with Sem-Sim ranging from 0.832 to 0.898. **GPT-5.2** (OpenAI, 2025) achieves the highest score (0.898), followed by **Claude-Sonnet-4.5** (Anthropic, 2025) (0.878), indicating robust topical consistency and factual coverage. However, these surface-level metrics do not reflect whether models preserve informational salience or affective urgency.

In contrast, our cognitive empathy metrics reveal substantial

Table 3. News generation results for cognitive empathy on the Newsroom dataset. Models generate full-length news articles conditioned on gold summaries. We report standard quality metrics (ROUGE, CIDEr, and Sem-Sim) alongside cognitive empathy metrics that capture salience and affective preservation. **Bold** and underlined denote the best and second-best results, respectively.

Model	Standard Metrics			Cog. Metrics	
	ROUGE ↑	CIDEr ↑	Sem-Sim ↑	SA ↓	EGM ↓
Claude-Sonnet-4.5	0.269	0.181	0.878	0.110	0.228
DeepSeek-V3.2	0.206	0.155	0.843	0.079	0.212
Doubao-Seed1.8	0.196	0.119	0.847	0.145	0.204
GPT-5.2	0.229	0.238	0.898	0.158	0.252
Gemini-3	0.213	0.175	0.838	0.121	0.223
Grok-4.1	0.191	0.120	0.856	0.109	0.195
GLM-4.7	0.231	0.113	0.840	0.128	0.243
Kimi-K2	0.178	0.136	0.874	0.183	0.201
Llama-4	<u>0.233</u>	0.164	0.851	0.290	0.179
Qwen3-VL	0.200	0.146	0.832	<u>0.086</u>	<u>0.188</u>

distortions. SA varies widely from 0.079 to 0.290, with **DeepSeek-V3.2** (Liu et al., 2025a) exhibiting the lowest attenuation and **Llama-4** (Meta AI, 2025) the highest. EGM ranges from 0.179 to 0.252, with **Llama-4** achieving the best EGM despite its severe affect attenuation. This inverse pattern suggests a trade-off between preserving affective intensity and calibrating descriptive detail.

Notably, high semantic fidelity does not imply cognitive empathy. **GPT-5.2**, despite leading in Sem-Sim and CIDEr, exhibits non-trivial SA (0.158) and the worst EGM (0.252). Similarly, **Claude-Sonnet-4.5** shows strong semantic performance but persistent empathy distortions. These results show that conventional metrics can mask failures in preserving user-relevant salience and affective calibration in cognitively demanding generation tasks.

4.2. Cultural Empathy: Literary Translation

Setup We assess cultural empathy using WMT 2023 Shared Task on Discourse-Level Literary Translation (Wang et al., 2023b;a). This task translates Chinese fiction into English and requires preserving culturally grounded style and register. Following prior work on literary translation (Wang et al., 2023b), we report standard metrics including **chrF** (Popović, 2015), **COMET** (Rei et al., 2020), and **Sem-Sim**. To probe cultural empathy, we also report **SA** and Linguistic Distancing (**LD**). Full details appear in Appendix C.

Key Findings As depicted in Table 4, under conventional evaluation, models perform comparably, with COMET scores ranging from 0.790 to 0.828 and chrF from 0.400 to 0.562, matching established literary translation baselines (Wang et al., 2023b). However, these metrics fail to capture culturally grounded aspects of narrative voice, affective tone, and interpersonal framing.

In contrast, our cultural empathy metrics reveal systematic

Table 4. Literary translation results for cultural empathy on WMT2023 Chinese-English. We report standard translation metrics (chrF, COMET) and semantic similarity (Sem-Sim) alongside cultural empathy metrics that capture affective and stylistic preservation. **Bold** and underlined denote the best and second-best results, respectively.

Model	Standard Metrics			Cult. Metrics	
	chrF ↑	COMET ↑	Sem-Sim ↑	SA ↓	LD ↓
Claude-Sonnet-4.5	<u>0.553</u>	0.828	0.805	0.431	1.177
DeepSeek-V3.2	0.546	0.828	0.816	0.457	1.361
Doubao-Seed1.8	0.531	0.818	0.813	0.448	1.275
GPT-5.2	0.546	0.815	0.822	<u>0.419</u>	<u>1.183</u>
Gemini-3	0.562	<u>0.827</u>	0.799	0.427	1.185
Grok-4.1	0.493	0.818	<u>0.819</u>	0.445	1.391
GLM-4.7	0.562	0.825	0.808	0.406	1.203
Kimi-K2	0.400	0.790	0.811	0.434	1.466
Llama-4	0.472	0.826	0.808	0.449	1.349
Qwen3-VL	0.507	0.822	0.818	0.459	1.409

distortions. SA remains high across models (0.406~0.459), indicating consistent flattening of culturally embedded affect despite competitive semantic similarity. For example, **GLM-4.7** (Z.ai, 2025) exhibits the lowest SA (0.406) while ranking only sixth in Sem-Sim (0.808), whereas **Qwen3-VL** (Bai et al., 2025a) shows the highest attenuation (0.459) despite strong semantic performance (0.818). LD varies widely (1.177~1.466), signaling pronounced failures in preserving authentic narrative voice. **DeepSeek-V3.2** (Liu et al., 2025a) illustrates this gap: despite achieving a joint-highest COMET score (0.828), it exhibits substantial cultural empathy degradation (SA 0.457, LD 1.361). These results demonstrate a persistent misalignment between standard translation metrics and cultural empathy, highlighting blind spots in current evaluation frameworks for cross-cultural literary translation.

4.3. Relational Empathy: Psychological Counseling

Setup We evaluate relational empathy using a constructed counseling-style dataset of 50 client narratives spanning family, workplace, and academic stressors. Models generate therapist-style responses. We report EPITOME scores (Sharma et al., 2020) to assess model-generated therapist responses across three dimensions: Emotional Reactions (**Emo-React**), Interpretations (**Interp.**), and Explorations (**Explor.**). Since our dataset lacks reference therapist replies, we aggregate the sub-scores by summing them and compute the mean across all examples to derive an overall counseling quality metric. To assess relational empathy failures, we further report **SA** and Conflict Avoidance (**CA**). Without a ground-truth counseling response, **SA** is computed as the sentiment divergence between the model’s reply and the client’s original narrative, where a lower score indicates better alignment with the client’s expressed emotional tone. The experimental setup is detailed in Appendix D.

Table 5. Psychological counseling results for relational empathy on a constructed counseling-style dataset. We evaluate therapist-style responses using EPITOME scores (Emo-React, Interpretation, and Exploration) as standard counseling metrics, and report relational empathy measures (SA and CA) to capture affective flattening and harmony-seeking behavior. **Bold** and underlined denote the best and second-best results, respectively.

Model	Standard Metrics			Rel. Metrics	
	Emo-React ↑	Interp. ↑	Explor. ↑	SA ↓	CA ↓
Claude-Sonnet-4.5	4.667	6.733	0.800	0.481	0.933
DeepSeek-V3.2	<u>5.533</u>	<u>8.133</u>	0.733	0.530	<u>0.700</u>
Doubao-Seed1.8	4.667	4.333	0.600	0.337	0.800
GPT-5.2	2.467	9.533	<u>1.400</u>	0.529	0.833
Gemini-3	5.433	3.533	0.800	0.512	1.000
Grok-4.1	4.100	5.867	0.533	0.515	1.000
GLM-4.7	5.667	3.667	1.067	0.509	0.967
Kimi-K2	2.167	4.133	1.733	0.487	0.933
Llama-4	2.700	4.400	0.067	0.529	0.667
Qwen3-VL	8.067	6.600	0.733	<u>0.431</u>	1.000

Key Findings From Table 5, frontier models under EPITOME exhibit broadly competent counseling behavior, with high Interpretation scores (up to 9.533) and Exploration scores often exceeding 1.0, indicating attempts at follow-up inquiry. However, mechanism-based metrics reveal systematic relational empathy failures. SA remains high across models (0.337~0.530), reflecting consistent downscaling of user affect. More strikingly, CA is pervasive, which reaches 1.000 for **Gemini-3** (Pichai et al., 2025), **Grok-4.1** (xAI, 2025), and **Qwen3-VL** (Bai et al., 2025a), and exceeds 0.9 for **Claude-Sonnet-4.5** (Anthropic, 2025) (0.933), **Kimi-K2** (Team et al., 2025) (0.933), and **GLM-4.7** (Z.ai, 2025) (0.967). Together, these results show that standard counseling-style metrics can rate responses as competent while masking systematic harmony-seeking and affect flattening that undermine relational alignment in practice.

4.4. Robustness of Experimental Findings

Setup We further assess the robustness of our relational empathy findings in two ways. First, to examine whether the observed conclusions depend on a particular judge model, we repeat the counseling evaluation using three frontier LLM judges: GPT-5.2, Claude-Sonnet-4.5, and Gemini-3-Pro. For each judge, we report SA and CA on the same constructed counseling-style dataset. Second, to evaluate sensitivity to sample size, we expand the counseling dataset from 50 to 100, 150, and 200 instances and recompute SA for each setting.

Key Findings Table 6 shows that the relative ranking patterns are broadly preserved across judges. Although the three judges exhibit different absolute scoring tendencies, their relative assessments are highly consistent across examples. Specifically, inter-judge agreement remains strong,

Table 6. Three frontier LLMs as judges for evaluating psychological counseling, verifying the stability of the SA and CA metrics. **Bold** and underlined denote the best and second-best results, respectively.

Model	GPT-5.2		Claude-Sonnet-4.5		Gemini-3-Pro	
	SA ↓	CA ↓	SA ↓	CA ↓	SA ↓	CA ↓
Claude-Sonnet-4.5	0.481	0.933	0.239	0.967	0.393	0.867
DeepSeek-V3.2	0.530	<u>0.700</u>	0.338	<u>0.700</u>	0.403	0.333
Doubao-Seed1.8	0.337	0.800	0.114	0.800	0.234	0.667
GPT-5.2	0.529	0.833	0.281	0.900	0.420	0.833
Gemini-3	0.512	1.000	0.214	1.000	0.358	0.767
Grok-4.1	0.515	1.000	0.352	1.000	0.403	1.000
GLM-4.7	0.509	0.967	0.222	0.967	0.368	0.967
Kimi-K2	0.487	0.933	0.213	1.000	0.365	1.000
Llama-4	0.529	0.667	0.479	0.621	0.450	<u>0.433</u>
Qwen3-VL	<u>0.431</u>	1.000	<u>0.141</u>	1.000	<u>0.314</u>	1.000

Table 7. Psychological counseling results under different dataset sizes {50, 100, 150, 200} on the SA metric. **Bold** and underlined denote the best and second-best results, respectively.

Model	50	100	150	200
Claude-Sonnet-4.5	0.481 ± 0.012	0.532 ± 0.007	0.535 ± 0.007	0.542 ± 0.008
DeepSeek-V3.2	0.530 ± 0.008	0.530 ± 0.015	0.497 ± 0.009	0.495 ± 0.010
Doubao-Seed1.8	0.337 ± 0.028	0.212 ± 0.043	0.208 ± 0.043	0.210 ± 0.045
GPT-5.2	0.529 ± 0.011	0.478 ± 0.009	0.491 ± 0.007	0.488 ± 0.007
Gemini-3	0.512 ± 0.003	0.366 ± 0.042	0.363 ± 0.050	0.368 ± 0.046
Grok-4.1	0.515 ± 0.015	0.423 ± 0.015	0.439 ± 0.013	0.431 ± 0.014
GLM-4.7	0.509 ± 0.050	0.416 ± 0.041	0.425 ± 0.035	0.425 ± 0.032
Kimi-K2	0.487 ± 0.063	0.344 ± 0.069	0.356 ± 0.064	0.365 ± 0.057
Llama-4	0.529 ± 0.005	0.514 ± 0.002	0.513 ± 0.002	0.511 ± 0.002
Qwen3-VL	<u>0.431</u> ± 0.046	<u>0.303</u> ± 0.055	<u>0.293</u> ± 0.054	<u>0.294</u> ± 0.054

with Spearman correlations of 0.835~0.948 for SA and 0.748~0.924 for CA. This suggests that our main conclusions do not depend on a single proprietary evaluator.

Table 7 further shows that the evaluation becomes more stable as the sample size increases. While the 50-instance setting exhibits greater fluctuation for some models, especially under SA, results become markedly more consistent at 100–200 instances, with Spearman correlations across dataset sizes approaching 1. The reported sample variances are also generally modest, indicating limited dispersion once the sample size is enlarged. Taken together, these results strengthen the empirical basis by showing that the observed relational empathy failures are not artifacts of a particular judge or a single small evaluation subset.

4.5. Summary

Across all three settings, frontier LLMs achieve strong conventional performance while exhibiting systematic empathic distortions. These results demonstrate that existing benchmarks mask alignment trade-offs that directly affect human-centered outcomes, motivating empathy-aware objectives,

evaluations, and training signals.

5. Towards Empathy-Aware Alignment

We outline concrete research directions for integrating empathy as a *measurable and optimizable objective* in aligned LLMs. Specifically, we propose three complementary actions: (i) elevating empathic fidelity to an explicit alignment objective, (ii) designing benchmarks that isolate empathy-specific failure modes, and (iii) curating training data that provides targeted supervision for mechanism-level interventions. These directions enable explicit control over trade-offs between helpfulness, harmlessness, and empathic fidelity in deployed systems.

5.1. Framing Empathy as an Alignment Objective

A key reason empathic failures persist is that current alignment pipelines treat empathy as an incidental style rather than a first-class optimization target (Rafailov et al., 2023; Shao et al., 2024; Hong et al., 2024). Existing objectives prioritize fluency, helpfulness, and risk mitigation, incentivizing calm and professional outputs that systematically suppress contextually relevant affect, perspective specificity, and relational stance. As a result, models may satisfy alignment criteria while remaining misaligned with the human perspectives they represent.

We advocate *empathy-aware alignment*, where empathic fidelity is optimized alongside helpfulness and safety (Wang et al., 2025b; 2026). This entails multi-objective post-training that explicitly penalizes empathic failures without degrading task performance. Crucially, the goal is calibration rather than amplification: preserving appropriate affective salience, abstraction level, and relational engagement given context. Framing empathy as an explicit objective mitigates the drift toward over-neutralized behavior that scaling alone may exacerbate in high-stakes deployments.

5.2. Designing Benchmarks for Empathic Fidelity

Benchmarks define what models are optimized for and trusted to generalize. Nevertheless, most current evaluations emphasize semantic adequacy and surface fluency, optimizing metrics such as ROUGE (Lin, 2004), chrF (Popović, 2015), COMET (Rei et al., 2020), or embedding similarity. Models can therefore score highly while flattening urgency, miscalibrating sensitive context, deflecting legitimate concerns, or adopting depersonalized framing that undermines human-centered objectives.

We propose *empathy-aware benchmarks* that treat empathic fidelity as a first-class evaluation dimension. Rather than averaging it away, benchmarks should concentrate on empathy-critical cases where distortions are most consequential. Such evaluations expose alignment trade-offs invisible under stan-

dard metrics and provide actionable signals for mechanism-aware training and deployment monitoring.

5.3. Training Datasets for Empathy-Aware Alignment

A remaining bottleneck is data (Bai et al., 2022; Köpf et al., 2023; Cui et al., 2024). Current instruction-tuning and preference datasets encode generic helpfulness and harmlessness but under-specify perspective preservation, affective calibration, and constructive engagement with tension.

We argue for datasets explicitly designed for *empathy-aware alignment*, providing supervision that distinguishes fluent compliance from faithful perspective-taking (Suh et al., 2025). These datasets can combine expert annotation, structured crowdsourcing, and empathy-focused adversarial generation to surface rare but high-impact failure modes. Targeted supervision at the mechanism level enables controllable trade-offs between safety, helpfulness, and empathic fidelity, preventing systematic drift toward over-neutralization while preserving alignment guarantees.

6. Alternative Views

Empathy is an overloaded term, and treating it as an LLM capability raises concerns about necessity, feasibility, and conceptual validity. We discuss four alternative views: alignment, prompting, factuality, and human-specificity, and clarify why framing empathy as an observable behavioral property remains useful for diagnosing systematic failures in human-centered generation.

AV1: Empathy as a Byproduct of Alignment LLMs undergo extensive post-training to optimize helpfulness and safety, leading some to view empathy as a stylistic byproduct of alignment rather than a distinct objective (Wei et al., 2022; Ouyang et al., 2022; Schulman et al., 2017; Rafailov et al., 2023; Shao et al., 2024; Hong et al., 2024). Yet, alignment pipelines primarily reward politeness and affective neutrality (Liu et al., 2023c; Huang et al., 2024b), without preserving user-relevant salience or interactional stance. Thus, well-aligned models systematically attenuate urgency, smooth tensions, or adopt depersonalized framing—failures that empathy captures as a complementary evaluation axis.

AV2: Empathy Through Prompt Engineering Another view holds that empathic behavior can be elicited on demand through careful prompt design, for example, by instructing models to acknowledge emotions or adopt a supportive tone (Sorin et al., 2024; Luo et al., 2024; Lee et al., 2024b). This suggests that empathy can be achieved without architectural changes, new training objectives, or dedicated evaluation. In practice, prompt-induced empathy reflects fragile style mimicry rather than stable preservation of user intent and context. Such behavior degrades under policy constraints, long interactions, and distribution shifts (Dorigoni & Gi-

ardino, 2025). Mechanism-level empathy targets systematic behavioral regularities that cannot be reliably enforced through prompting alone.

AV3: Empathy Undermines Factuality and Objectivity Emphasizing empathy is sometimes seen as conflicting with factuality and neutrality, particularly in domains such as news reporting and public services (Pjesivac & Ahn, 2024; Yu et al., 2024; Ameen et al., 2025), where emotional language may introduce bias or sensationalism (Kim et al., 2026). We do not advocate emotional amplification, but calibrated affect and salience: neutrality can remain factually correct while attenuating urgency, accountability, or human impact. Empathy thus complements factual accuracy by preserving stakeholder-relevant meaning.

AV4: Empathy Is Human-Specific Finally, attributing empathy to LLMs risks anthropomorphism, projecting human mental states onto statistical systems (Decety & Jackson). From this perspective, what appears as empathy may be dismissed as purely performative language rather than a meaningful research target. Our definition explicitly avoids this pitfall by defining empathy as an observable behavioral property of outputs, analogous to helpfulness or safety, rather than as an internal mental state. By decomposing empathic failures into measurable mechanisms (SA, EGM, CA, LD), we enable principled analysis without attributing human-like mental states to models.

7. Conclusions

In this paper, we argue that empathy should be treated as an explicit first-class objective in aligned LLM systems, rather than an incidental byproduct of fluency or policy compliance. We show that even well-aligned models exhibit systematic, observable distortions of human perspectives via four recurring mechanisms—sentiment attenuation, empathic granularity mismatch, conflict avoidance, and linguistic distancing—failures that remain hidden under strong benchmark performance. By organizing empathic requirements along cognitive, cultural, and relational dimensions, our analyses reveal subtle alignment trade-offs that remain invisible under semantic adequacy alone. These findings motivate a shift toward empathy-aware objectives, benchmarks, and training signals that make empathic fidelity a measurable, controllable capability in human-centered language models.

Nevertheless, our current empirical exploration remains limited to a relatively small set of domains, and broader cross-domain validation is needed to assess the generality of these mechanisms. As future work, we plan to conduct large-scale human studies to further examine whether the proposed failure mechanisms align with human judgments of empathic fidelity and to strengthen the empirical basis of our position.

Acknowledgements

This work was supported by the National Natural Science Foundation of China (NSFC) under Grant Nos. 62125201, U24B20174, and U25B6003, as well as the New Generation Artificial Intelligence-National Science and Technology Major Project (2025ZD0123302).

References

- Adams, G., Fabbri, A., Ladhak, F., Lehman, E., and Elhadad, N. From sparse to dense: Gpt-4 summarization with chain of density prompting. In *Proceedings of the 4th New Frontiers in Summarization Workshop*, pp. 68–74, 2023.
- Almasi, M. and Schiønning, A. Fine-tuning gpt-3 for synthetic danish news generation. In *Proceedings of the 16th International Natural Language Generation Conference*, pp. 54–68, 2023.
- Ameen, M. R., Islam, A., Miah, A. S. M., Siddiqua, A., and Shin, J. How news feels: Understanding affective bias in multilingual headlines for human-centered media design. *arXiv preprint arXiv:2510.17252*, 2025.
- Anthropic. Introducing claude sonnet 4.5. <https://www.anthropic.com/news/claude-sonnet-4-5>, September 2025. Accessed: 2026-01-04.
- Arita, T., Zheng, W., Suzuki, R., and Akiba, F. Assessing llms in art contexts: Critique generation and theory of mind evaluation. *arXiv preprint arXiv:2504.12805*, 2025.
- Bae, M. and Kim, H. Collective critics for creative story generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 18784–18819, 2024.
- Bahrainian, S. A., Feucht, S., and Eickhoff, C. NEWTS: A corpus for news topic-focused summarization. In *Findings of the Association for Computational Linguistics: ACL 2022*, pp. 493–503, 2022.
- Bai, S., Cai, Y., Chen, R., Chen, K., Chen, X., Cheng, Z., Deng, L., Ding, W., Gao, C., Ge, C., et al. Qwen3-VL technical report. *arXiv preprint arXiv:2511.21631*, 2025a.
- Bai, Y., Jones, A., Ndousse, K., Askell, A., Chen, A., Das-Sarma, N., Drain, D., Fort, S., Ganguli, D., Henighan, T., et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- Bai, Y., Zhang, J., Lv, X., Zheng, L., Zhu, S., Hou, L., Dong, Y., Tang, J., and Li, J. Longwriter: Unleashing 10,000+ word generation from long context llms. In *The Thirteenth International Conference on Learning Representations*, 2025b.
- Banting, S. Simulated social action: Assessing generative ai’s attempts at the rhetoric of literary criticism. *Discourse and Writing/Rédactologie*, 35:66–88, 2025.
- Bardol, F. Chatgpt reads your tone and responds accordingly—until it does not—emotional framing induces bias in llm outputs. *arXiv preprint arXiv:2507.21083*, 2025.
- Bin, Y., Shi, W., Ding, Y., Hu, Z., Wang, Z., Yang, Y., Ng, S.-K., and Shen, H. T. Gallerygpt: Analyzing paintings with large multimodal models. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pp. 7734–7743, 2024.
- Brachman, M., El-Ashry, A., Dugan, C., and Geyer, W. How knowledge workers use and want to use llms in an enterprise context. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*, pp. 1–8, 2024.
- Breen, B. Large language models, experiential learning, and ai simulations in the classroom: The historylens framework. *Teaching History: A Journal of Methods*, 49(1): 10–33, 2025.
- Brigham, N. G., Gao, C., Kohno, T., Roesner, F., and Mireshghallah, N. Developing story: Case studies of generative ai’s use in journalism. *arXiv preprint arXiv:2406.13706*, 2024.
- Bungum, L., Huang, C. Y., and Kashar, A. If i could turn back time: Temporal reframing as a historical reasoning task for llms. *arXiv preprint arXiv:2511.04432*, 2025.
- Bytedance Seed. Seed1.8 model card: Towards generalized real-world agency. <https://l1f3-static.bytednsdoc.com/obj/eden-cn/lapzild-tss/1jhwZthlaukjlkulzlp/research/Seed-1.8-Modelcard.pdf>, December 2025. Accessed: 2026-01-08.
- Castaldo, A., Castilho, S., Moorkens, J., and Monti, J. Extending CREAMT: Leveraging large language models for literary translation post-editing. In *Proceedings of Machine Translation Summit XX: Volume 1*, pp. 506–515, June 2025.
- Chakrabarty, T., Laban, P., Agarwal, D., Muresan, S., and Wu, C.-S. Art or artifice? large language models and the false promise of creativity. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, pp. 1–34, 2024.

- Chen, J., Li, H., Qin, M., Zhou, Y., Ren, Y., Wang, W., Liu, Y., Wu, Y., and Ai, Q. Simulating dispute mediation with llm-based agents for legal research. *arXiv preprint arXiv:2509.06586*, 2025.
- Chen, Y., Xing, X., Lin, J., Zheng, H., Wang, Z., Liu, Q., and Xu, X. Soulchat: Improving llms’ empathy, listening, and comfort abilities through fine-tuning with multi-turn empathy conversations. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 1170–1183, 2023.
- Chen, Y., Yan, S., Liu, S., Li, Y., and Xiao, Y. Emotion-queen: A benchmark for evaluating empathy of large language models. In *Findings of the Association for Computational Linguistics ACL 2024*, pp. 2149–2176, 2024.
- Cheng, M., Yu, S., Lee, C., Khadpe, P., Ibrahim, L., and Jurafsky, D. Elephant: Measuring and understanding social sycophancy in llms. *arXiv preprint arXiv:2505.13995*, 2025.
- Chhun, C., Suchanek, F., and Clavel, C. Do language models enjoy their own stories? prompting large language models for automatic story evaluation. *Transactions of the Association for Computational Linguistics*, 12:1122–1142, 2024.
- Cho, T. and Miller, B. J. Using artificial intelligence to improve administrative process in medicaid. *Health Affairs Scholar*, 2(2):qxae008, 2024.
- Coz, P. L., Liu, J. A., Bhattacharjya, D., Curto, G., and Stinckwich, S. What would an llm do? evaluating policymaking capabilities of large language models. *arXiv preprint arXiv:2509.03827*, 2025.
- Cuff, B. M., Brown, S. J., Taylor, L., and Howat, D. J. Empathy: A review of the concept. *Emotion review*, 8(2): 144–153, 2016.
- Cui, G., Yuan, L., Ding, N., Yao, G., He, B., Zhu, W., Ni, Y., Xie, G., Xie, R., Lin, Y., et al. Ultrafeedback: boosting language models with scaled ai feedback. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 9722–9744, 2024.
- Dai, S., Huang, Q., You, X., and Yu, J. MG²-RAG: Multi-Granularity Graph for Multimodal Retrieval-Augmented Generation. *arXiv preprint arXiv:2604.04969*, 2026.
- Danescu-Niculescu-Mizil, C., Sudhof, M., Jurafsky, D., Leskovec, J., and Potts, C. A computational approach to politeness with application to social factors. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 250–259, 2013.
- De Duro, E. S., Improt, R., and Stella, M. Introducing counsellme: A dataset of simulated mental health dialogues for comparing llms like haiku, llamantino and chatgpt against humans. *Emerging Trends in Drugs, Addictions, and Health*, 5:100170, 2025.
- Decety, J. and Jackson, P. L. The functional architecture of human empathy. *Behavioral and cognitive neuroscience reviews*, 3(2):71–100.
- Decostanzi, I., Mejova, Y., and Kalimeri, K. A large-language-model framework for automated humanitarian situation reporting. *arXiv preprint arXiv:2512.19475*, 2025.
- Diakopoulos, N., Cools, H., Helberger, N., Li, C., Kung, E., Rinehart, A., et al. Generative ai in journalism: The evolution of newswork and ethics in a generative information ecosystem. https://cdn.theconversation.com/static_files/files/3273/AP_Generative_AI_Report_April_2024.pdf, 2024. Accessed: 2026-01-11.
- Dong, M., Bonnefon, J.-F., and Rahwan, I. Heterogeneous preferences and asymmetric insights for ai use among welfare claimants and non-claimants. *Nature Communications*, 16(1):6973, 2025.
- Dorigoni, A. and Giardino, P. L. The illusion of empathy: evaluating ai-generated outputs in moments that matter. *Frontiers in Psychology*, 16:1568911, 2025.
- Du, H., Shi, J., Myerston, J., Lu, S., Zhou, G., and Gao, A. Role-guided annotation and prototype-aligned representation learning for historical literature sentiment classification. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 3756–3768, November 2025.
- Duan, J., Zhao, X., Zhang, Z., Ko, E. G., Boddy, L., Wang, C., Li, T., Rasgon, A., Hong, J., Lee, M. K., et al. An exploration of llm-guided conversation in reminiscence therapy. In *GenAI for Health: Potential, Trust and Policy Compliance*, 2024.
- Eichbaum, Q. and Bleakley, A. Re-visioning intercultural relational empathy. *Advances in Health Sciences Education*, pp. 1–15, 2025.
- Eichbaum, Q., Barbeau-Meunier, C.-A., White, M., Ravi, R., Grant, E., Riess, H., and Bleakley, A. Empathy across cultures—one size does not fit all: from the ego-logical to the eco-logical of relational empathy. *Advances in Health Sciences Education*, 28(2):643–657, 2023.
- Eisenberg, N., Eggum, N. D., and Di Giunta, L. Empathy-related responding: Associations with prosocial behavior, aggression, and intergroup relations. *Social issues and policy review*, 4(1):143–180, 2010.

- Fabbri, A. R., Li, I., She, T., Li, S., and Radev, D. Multi-news: A large-scale multi-document summarization dataset and abstractive hierarchical model. In *Proceedings of the 57th annual meeting of the association for computational linguistics*, pp. 1074–1084, 2019.
- Fang, X., Che, S., Mao, M., Zhang, H., Zhao, M., and Zhao, X. Bias of ai-generated content: an examination of news produced by large language models. *Scientific Reports*, 14(1):5224, 2024.
- Feng, H., Fan, Y., Liu, X., Lin, T.-E., Yao, Z., Wu, Y., Huang, F., Li, Y., and Ma, Q. Improving factual consistency of news summarization by contrastive preference optimization. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 11084–11100, 2024.
- Feng, Y., Song, M., Wang, J., Chen, Z., Bi, G., Huang, M., Jing, L., and Yu, J. Ss-gen: A social story generation framework with large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 1300–1308, 2025a.
- Feng, Y., Sun, Y., Sun, Y., Zhu, M., Huang, Q., Tung, A. K. H., and Chen, W. Don’t reinvent the wheel: Efficient instruction-following text embedding based on guided space transformation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 24511–24525, 2025b. URL <https://aclanthology.org/2025.acl-long.1196/>.
- Gabriel, S., Puri, I., Xu, X., Malgaroli, M., and Ghassemi, M. Can ai relate: Testing large language model response for mental health support. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 2206–2221, 2024.
- Gliwa, B., Mochol, I., Biesek, M., and Wawer, A. SAMSum corpus: A human-annotated dialogue dataset for abstractive summarization. In *Proceedings of the 2nd Workshop on New Frontiers in Summarization*, pp. 70–79, 2019.
- Govers, J., Velloso, E., Kostakos, V., and Goncalves, J. Ai-driven mediation strategies for audience depolarisation in online debates. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, pp. 1–18, 2024.
- Grusky, M., Naaman, M., and Artzi, Y. Newsroom: A dataset of 1.3 million summaries with diverse extractive strategies. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 708–719, 2018.
- Gurung, A. and Lapata, M. Learning to reason for long-form story generation. *arXiv preprint arXiv:2503.22828*, 2025.
- Hamdani, R. R., Jalane, M. D. F., and Ame, I. P. The impact of ChatGPT on history education: An analysis of misinformation and historical reasoning in high schools. *ResearchGate Preprint*, January 2024.
- Han, S., Chen, L., Lin, L.-M., Xu, Z., and Yu, K. Ibsen: Director-actor agent collaboration for controllable and interactive drama script generation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1607–1619, 2024.
- Han, X., Liu, Y., Li, M., Liu, W., Hu, S., Yan, R., Xu, Z., and Chen, X. Pastiche novel generation creating: Fan fiction you love in your favorite author’s style. *arXiv preprint arXiv:2502.15616*, 2025.
- Hao, J., Liu, H., Xiao, X., Huang, Q., and Yu, J. Uni-x: Mitigating modality conflict with a two-end-separated architecture for unified multimodal models. In *The Fourteenth International Conference on Learning Representations (ICLR)*, 2026. URL <https://openreview.net/forum?id=IJLIYpCkwz>.
- Hauser, J., Kondor, D., Reddish, J., Benam, M., Cioni, E., Villa, F., Bennett, J., Hoyer, D., Francois, P., Turchin, P., et al. Large language models’ expert-level global history knowledge benchmark (hist-llm). *Advances in Neural Information Processing Systems*, 37:32336–32369, 2024.
- Hayashi, K., Sakai, Y., Kamigaito, H., Hayashi, K., and Watanabe, T. Towards artwork explanation in large-scale vision language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pp. 705–729, 2024.
- He, G., Demartini, G., and Gadiraju, U. Plan-then-execute: An empirical study of user trust and team performance when using llm agents as a daily assistant. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pp. 1–22, 2025.
- Healey, M. L. and Grossman, M. Cognitive and affective perspective-taking: evidence for shared and dissociable anatomical substrates. *Frontiers in neurology*, 9:491, 2018.
- Held, P., Pridgen, S. A., Chen, Y., Akhtar, Z., Amin, D., Pohorence, S., et al. A novel cognitive behavioral therapy-based generative ai tool (socrates 2.0) to facilitate socratic dialogue: Protocol for a mixed methods feasibility study. *JMIR Research Protocols*, 13(1):e58195, 2024.

- Hermann, K. M., Kocisky, T., Grefenstette, E., Espeholt, L., Kay, W., Suleyman, M., and Blunsom, P. Teaching machines to read and comprehend. *Advances in neural information processing systems*, 28, 2015.
- Hong, J., Lee, N., and Thorne, J. Orpo: Monolithic preference optimization without reference model. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 11170–11189, 2024.
- Honnibal, M. spacy 2: Natural language understanding with bloom embeddings, convolutional neural networks and incremental parsing, 2017.
- Hu, B., Sheng, Q., Cao, J., Li, Y., and Wang, D. Llm-generated fake news induces truth decay in news ecosystem: A case study on neural news recommendation. In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 435–445, 2025a.
- Hu, Q., Xiao, Q., Cao, H., and Shen, H. When your boss is an ai bot: Exploring opportunities and risks of manager clone agents in the future workplace. *arXiv preprint arXiv:2509.10993*, 2025b.
- Hu, T.-Y., Koppula, H. S., Pouransari, H., Koc, C., Tuzel, O., and Vemulapalli, R. Learning from self critique and refinement for faithful llm summarization. *arXiv preprint arXiv:2512.05387*, 2025c.
- Huang, K.-H., Laban, P., Fabbri, A. R., Choubey, P. K., Joty, S., Xiong, C., and Wu, C.-S. Embrace divergence for richer insights: A multi-document summarization benchmark and a case study on summarizing diverse information from news articles. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 570–593, 2024a.
- Huang, S., Siddarth, D., Lovitt, L., Liao, T. I., Durmus, E., Tamkin, A., and Ganguli, D. Collective constitutional ai: Aligning a language model with public input. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pp. 1395–1417, 2024b.
- Huang, Y. and Sun, L. Fakegpt: fake news generation, explanation and detection of large language models. *arXiv preprint arXiv:2310.05046*, 2023.
- Iftikhar, Z., Ransom, S., Xiao, A., Nugent, N., and Huang, J. Therapy as an nlp task: psychologists’ comparison of llms and human peers in cbt. *arXiv preprint arXiv:2409.02244*, 2024.
- Iftikhar, Z., Xiao, A., Ransom, S., Huang, J., and Suresh, H. How llm counselors violate ethical standards in mental health practice: A practitioner-informed framework. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, volume 8, pp. 1311–1323, 2025.
- Ismayilzada, M., Stevenson, C., and van der Plas, L. Evaluating creative short story generation in humans and large language models. *arXiv preprint arXiv:2411.02316*, 2024.
- Jiang, K., Zhang, Q., Guo, D., Huang, D., Zhang, S., Wei, Z., Ning, F., and Li, R. Ai-generated news articles based on large language models. In *Proceedings of the 2023 International Conference on Artificial Intelligence, Systems and Network Security*, pp. 82–87, 2023.
- Jin, L., An, L., and Ma, X. Towards chapter-to-chapter context-aware literary translation via large language models. *arXiv preprint arXiv:2407.08978*, 2024.
- Jo, J., Zhang, H., Cai, J., and Goyal, N. Ai trust reshaping administrative burdens: Understanding trust-burden dynamics in llm-assisted benefits systems. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*, pp. 1172–1183, 2025.
- Juvino Santos, L. R., Balby Marinho, L., Calazans Campelo, C. E., Menczer, F., and Flammini, A. Can large language models effectively mitigate polarization in social media text? In *Proceedings of the 17th ACM Web Science Conference 2025*, pp. 348–357, 2025.
- Karpinska, M. and Iyyer, M. Large language models effectively leverage document-level context for literary translation, but critical errors persist. In *Proceedings of the Eighth Conference on Machine Translation*, pp. 419–451, 2023.
- Karr Jr, J. A., Smith, E., Hauenstein, M., Curto, G., and Chawla, N. V. What is behind homelessness bias? using llms and nlp to mitigate homelessness by acting on social stigma. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence*, pp. 9754–9762, 2025.
- Kennan, A., Singh, L., Garcia Guevara, A., Ahmed, M., and Goodman, J. AI-Powered rules as code: Experiments with public benefits policy. <https://digitalgovernmenthub.org/publications/ai-powered-rules-as-code-experiments-with-public-benefits-policy/>, March 2025. Accessed: 2026-01-11.
- Khadangi, A., Sartipi, A., Tchappi, I., and Fridgen, G. Cognartive: Large language models for automating art analysis and decoding aesthetic elements. *arXiv preprint arXiv:2502.04353*, 2025.

- Kim, S., Kim, J., Shin, S., Chung, H., Moon, D., Kwon, Y., and Yoon, H. Being kind isn't always being safe: Diagnosing affective hallucination in llms. *arXiv preprint arXiv:2508.16921*, 2025a.
- Kim, S., Kim, J., Shin, S., Chung, H., Moon, D., Kwon, Y., and Yoon, H. Being kind isn't always being safe: Diagnosing affective hallucination in llms. In *Findings of the Association for Computational Linguistics: EACL 2026*, pp. 50–78, 2026.
- Kim, Y. and Hsu, C.-C. Leveraging large language models for hybrid workplace decision support. *arXiv preprint arXiv:2402.03616*, 2024.
- Kim, Y., Choi, C.-H., Cho, S., Sohn, J.-y., and Kim, B.-H. Aligning large language models for cognitive behavioral therapy: a proof-of-concept study. *Frontiers in Psychiatry*, 16:1583739, 2025b.
- Kim, Y. S., Moon, H. S., Lee, S., and Lee, T. Y. 'histochat': Leveraging ai-driven historical personas for personalized and engaging middle school history education. *Proceedings of the ACM on Human-Computer Interaction*, 9(7): 1–37, 2025c.
- Kindenberg, B. Chatgpt-generated and student-written historical narratives: A comparative analysis. *Education sciences*, 14(5):530, 2024.
- Kirk, R., Mediratta, I., Nalmpantis, C., Luketina, J., Hambro, E., Grefenstette, E., and Raileanu, R. Understanding the effects of RLHF on LLM generalisation and diversity. In *The Twelfth International Conference on Learning Representations*, 2024.
- Koçak, Ö., Puranam, P., and Yegin, A. Llms as mediators: Can they diagnose conflicts accurately? *ACM Journal on Computing and Sustainable Societies*, 2024.
- Kong, L., Kim, C. W., Choo, D., and Tambe, M. Generative AI for social impact. *arXiv preprint arXiv:2601.04238*, 2026.
- Köpf, A., Kilcher, Y., Von Rütte, D., Anagnostidis, S., Tam, Z. R., Stevens, K., Barhoum, A., Nguyen, D., Stanley, O., Nagyfi, R., et al. Openassistant conversations-democratizing large language model alignment. *Advances in neural information processing systems*, 36: 47669–47681, 2023.
- Lai, T., Shi, Y., Du, Z., Wu, J., Fu, K., Dou, Y., and Wang, Z. Psy-llm: Scaling up global mental health psychological services with ai-based large language models. *arXiv preprint arXiv:2307.11991*, 2023.
- Lee, J., Yoon, S., Lee, K., Jeong, E., Cho, J.-E., Park, W., Yim, D., and Hwang, I. Open sesame? open salami! personalizing vocabulary assessment-intervention for children via pervasive profiling and bespoke storybook generation. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, pp. 1–32, 2024a.
- Lee, S., Kim, M., Hwang, S., Kim, D., and Lee, K. Amplifying minority voices: Ai-mediated devil's advocate system for inclusive group decision-making. In *Companion Proceedings of the 30th International Conference on Intelligent User Interfaces*, pp. 17–21, 2025.
- Lee, Y. K., Suh, J., Zhan, H., Li, J. J., and Ong, D. C. Large language models produce responses perceived to be empathic. In *2024 12th International Conference on Affective Computing and Intelligent Interaction (ACII)*, pp. 63–71. IEEE, 2024b.
- Lei, H., Guo, J., He, G., Zhang, X., Zhang, R., Peng, S., Liu, S., and Chen, T. Ex3: Automatic novel writing by extracting, excelsior and expanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 9125–9146, 2024.
- Levett-Jones, T., Cant, R., and Lapkin, S. A systematic review of the effectiveness of empathy education for undergraduate nursing students. *Nurse education today*, 75: 80–94, 2019.
- Li, C., Chen, M., Wang, J., Sitaram, S., and Xie, X. Culturellm: Incorporating cultural differences into large language models. *Advances in Neural Information Processing Systems*, 37:84799–84838, 2024a.
- Li, D., Alnaibari, A., Bisharat, A., Sandoval, M., Hall, D., Silva, Y., and Liu, H. From moderation to mediation: Can llms serve as mediators in online flame wars? *arXiv preprint arXiv:2512.03005*, 2025a.
- Li, K., Tao, Y., Huang, H., and Ji, T. Linchance-ntu for unconstrained wmt2024 literary translation. In *Proceedings of the Ninth Conference on Machine Translation*, pp. 987–992, 2024b.
- Li, M., Chen, M.-B., Tang, B., ShengbinHou, S., Wang, P., Deng, H., Li, Z., Xiong, F., Mao, K., Peng, C., et al. Newsbench: a systematic evaluation framework for assessing editorial capabilities of large language models in chinese journalism. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 9993–10014, 2024c.
- Li, W., Lai, Y., Soni, S., and Saha, K. Emails by llms: A comparison of language in ai-generated and human-written emails. In *Proceedings of the 17th ACM Web Science Conference 2025*, pp. 391–403, 2025b.

- Lim, M. Rethinking history education in the age of ai. *HSSE Online*, 13(1):24–37, July 2025.
- Lin, C.-Y. ROUGE: a package for automatic evaluation of summaries. In *Text Summarization Branches Out*, pp. 74–81, 2004.
- Lin, L., Zheng, J., and Wang, H. Webnovelbench: Placing llm novelists on the web novel distribution. *arXiv preprint arXiv:2505.14818*, 2025.
- Liu, A., Mei, A., Lin, B., Xue, B., Wang, B., Xu, B., Wu, B., Zhang, B., Lin, C., Dong, C., et al. Deepseek-v3.2: Pushing the frontier of open large language models. *arXiv preprint arXiv:2512.02556*, 2025a.
- Liu, J., Ai, L., Liu, Z., Karisani, P., Hui, Z., Fung, Y., Nakov, P., Hirschberg, J., and Ji, H. Propainsight: Toward deeper understanding of propaganda in terms of techniques, appeals, and intent. In *Proceedings of the 31st International Conference on Computational Linguistics*, pp. 5607–5628, 2025b.
- Liu, X., Yang, S., Zhang, X., Kuang, H., Sun, L., Yang, Y., Chen, S., Huang, X.-J., and Wei, Z. Ai-press: A multi-agent news generating and feedback simulation system powered by large language models. In *Proceedings of the 31st International Conference on Computational Linguistics: System Demonstrations*, pp. 63–82, 2025c.
- Liu, Y., Fabbri, A., Liu, P., Zhao, Y., Nan, L., Han, R., Han, S., Joty, S., Wu, C.-S., Xiong, C., et al. Revisiting the gold standard: Grounding summarization evaluation with robust human evaluation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 4140–4170, 2023a.
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., and Zhu, C. G-eval: Nlg evaluation using gpt-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 2511–2522, 2023b.
- Liu, Y., Yao, Y., Ton, J.-F., Zhang, X., Guo, R., Cheng, H., Klochkov, Y., Taufiq, M. F., and Li, H. Trustworthy llms: a survey and guideline for evaluating large language models’ alignment. *arXiv preprint arXiv:2308.05374*, 2023c.
- Liu, Y., Fabbri, A. R., Chen, J., Zhao, Y., Han, S., Joty, S., Liu, P., Radev, D., Wu, C.-S., and Cohan, A. Benchmarking generation and evaluation capabilities of large language models for instruction controllable summarization. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pp. 4481–4501, 2024a.
- Liu, Y., Shi, K., He, K., Ye, L., Fabbri, A. R., Liu, P., Radev, D., and Cohan, A. On learning to summarize with large language models as references. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 8647–8664, 2024b.
- Lozoya, D. C., Conway, M., De Duro, E. S., and D’Alfonso, S. Leveraging large language models for simulated psychotherapy client interactions: Development and usability study of client101. *JMIR Medical Education*, 11(1): e68056, 2025.
- Lu, H. and Zhao, Z. Sh-rag: A syntax-based hierarchical retrieval-augmented generation framework for handling syntactic complexity in literary machine translation. In *2025 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8. IEEE, 2025.
- Lu, Z., Mysore, S., Safavi, T., Neville, J., Yang, L., and Wan, M. Corporate communication companion (ccc): An llm-empowered writing assistant for workplace social media. *arXiv preprint arXiv:2405.04656*, 2024.
- Luo, M., Warren, C. J., Cheng, L., Abdul-Muhsin, H. M., and Banerjee, I. Assessing empathy in large language models with real-world physician-patient interactions. In *2024 IEEE International Conference on Big Data (Big-Data)*, pp. 6510–6519. IEEE, 2024.
- Ma, G. Chance or challenge: The role of ChatGPT in history teaching and historical research in higher education. In *2023 3rd International Conference on Education, Information Management and Service Science (EIMSS 2023)*, pp. 869–874. Atlantis Press, 2023.
- Ma, S., Chen, Q., Wang, X., Zheng, C., Peng, Z., Yin, M., and Ma, X. Towards human-ai deliberation: Design and evaluation of llm-empowered deliberative ai for ai-assisted decision-making. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pp. 1–23, Yokohama, Japan, April 2025.
- Ma, Y., Qiao, Y., and Liu, P. Mops: Modular story premise synthesis for open-ended automatic story generation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 2135–2169, 2024.
- Macken, L. Machine translation meets large language models: Evaluating chatgpt’s ability to automatically post-edit literary texts. In *Proceedings of the 1st Workshop on Creative-text Translation and Technology*, pp. 65–81, 2024.

- Main, A., Walle, E. A., Kho, C., and Halpern, J. The interpersonal functions of empathy: A relational perspective. *Emotion Review*, 9(4):358–366, 2017.
- Manevich, A., Carmel, D., Cohen, N., Kravi, E., and Shapira, O. Multi document summarization evaluation in the presence of damaging content. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 1–12, 2023.
- Margolina, A. V. Controlling impression: making rugpt3 generate sentiment-driven movie reviews. *Journal of applied linguistics and lexicography*, 4(1):15–25, 2022.
- Meta AI. The llama 4 herd: The beginning of a new era of natively multimodal intelligence. <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>, April 2025. Accessed: 2026-01-04.
- Mierwald, M. Chatting about the past with artificial intelligence: A case study of pupils’ interaction with chatgpt while completing a history-learning task. *Journal of Educational Media, Memory and Society*, 16(2):143–173, 2024.
- Minnema, G., Gemelli, S., Zanchi, C., Caselli, T., and Nissim, M. Dead or murdered? predicting responsibility perception in femicide news reports. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 1078–1090, 2022.
- Miura, Y., Yang, C.-L., Kuribayashi, M., Matsumoto, K., Kuzuoka, H., and Morishima, S. Understanding and supporting formal email exchange by answering ai-generated questions. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pp. 1–20, 2025.
- Moghimifar, F., Li, Y.-F., Thomson, R., and Haffari, G. Modelling political coalition negotiations using llm-based agents. *arXiv preprint arXiv:2402.11712*, 2024.
- Nadas, M., Diosan, L., Tomescu, A., and Piscoran, A. Small open models achieve near parity with large models in low resource literary translation at a fraction of the cost. *arXiv preprint arXiv:2509.07829*, 2025.
- Narayan, S., Cohen, S. B., and Lapata, M. Don’t give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 1797–1807, 2018.
- Navarro, A., de Quinto, C., and Hernández, J. A. Email as the interface to generative ai models: Seamless administrative automation. *arXiv preprint arXiv:2506.23850*, 2025.
- Nguyen, X.-H., Nguyen, H.-A., Cao, L., and Truong, H. ChatGPT’s understanding of history: A comparison to vietnamese students and its potential in history education. *EdArXiv Preprints*, 2023. doi: 10.35542/osf.io/8z9tj.
- Nishal, S. and Diakopoulos, N. Envisioning the applications and implications of generative ai for news media. *arXiv preprint arXiv:2402.18835*, 2024.
- Nook, E. C., Schleider, J. L., and Somerville, L. H. A linguistic signature of psychological distancing in emotion regulation. *Journal of Experimental Psychology: General*, 146(3):337, 2017.
- Núñez-Naranjo, A. and Morales-Urrutia, X. Critical thinking and ai: Enhancing history teaching through chatgpt simulations. *International Journal of Innovative Research and Scientific Studies*, 8(1):564–575, 2025.
- Odabaşı, A. and Biricik, G. Unraveling the capabilities of language models in news summarization. *arXiv preprint arXiv:2501.18128*, 2025.
- OpenAI. Introducing GPT-5.2. <https://openai.com/index/introducing-gpt-5-2/>, December 2025. Accessed: 2026-01-04.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Pandey, P. S., Le, H. S., Bhardwaj, D., Mihalcea, R., and Jin, Z. Socialharmbench: Revealing llm vulnerabilities to socially harmful requests. *arXiv preprint arXiv:2510.04891*, 2025.
- Pang, J., Ye, F., Wong, D. F., Yu, D., Shi, S., Tu, Z., and Wang, L. Salute the classic: Revisiting challenges of machine translation in the age of large language models. *Transactions of the Association for Computational Linguistics*, 13:73–95, 2025.
- Park, G., Hwang, S., and Lee, H. Low-resource cross-lingual summarization through few-shot learning with large language models. In *Proceedings of the Seventh Workshop on Technologies for Machine Translation of Low-Resource Languages (LoResMT 2024)*, pp. 57–63, 2024.
- Park, K., Song, H., Seo, S., Kim, J., and Suh, B. ”ask sir oliver ingham”: Llm-based social simulations for history education. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*, pp. 1–13, 2025a.

- Park, S., Zamfirescu-Pereira, J., and Kulkarni, C. Model behavior specification by leveraging llm self-playing and self-improving. *arXiv preprint arXiv:2503.03967*, 2025b.
- Parmar, D. and Mazumdar, P. Emotionally aware or tone-deaf? evaluating emotional alignment in llm-based conversational recommendation systems. In *Proceedings of the 9th Widening NLP Workshop*, pp. 167–174, 2025.
- Patel, Z., El-Refai, K., Pei, J., and Li, T. Swag: Storytelling with action guidance. *arXiv preprint arXiv:2402.03483*, 2024.
- Patil, A. and Gedhu, A. K. Cognitive-mental-llm: Evaluating reasoning in large language models for mental health prediction via online text. *arXiv preprint arXiv:2503.10095*, 2025.
- Peng, Q., Liu, H., Xu, H., Yang, Q., Shao, M., and Wang, W. Review-llm: Harnessing large language models for personalized review generation. *arXiv preprint arXiv:2407.07487*, 2024.
- Peters, U. and Chin-Yee, B. Generalization bias in large language model summarization of scientific research. *Royal Society Open Science*, 12(4):241776, 2025.
- Petridis, S., Diakopoulos, N., Crowston, K., Hansen, M., Henderson, K., Jastrzebski, S., Nickerson, J. V., and Chilton, L. B. Anglekindling: Supporting journalistic angle ideation with large language models. In *Proceedings of the 2023 CHI conference on human factors in computing systems*, pp. 1–16, 2023.
- Pichai, S., Hassabis, D., and Kavukcuoglu, K. A new era of intelligence with gemini 3. <https://blog.google/products/gemini/gemini-3/#note-from-ceo>, November 2025. Accessed: 2026-01-04.
- Pjesivac, I. and Ahn, S. J. 360° journalism and empathy: Psychological processes and communication outcomes. *Journalism and Media*, 5(3):1007–1022, 2024.
- Pokharel, G., Farabi, S., Fowler, P. J., and Das, S. Street-Level AI: Are large language models ready for real-world judgments? In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, volume 8, pp. 2043–2054, 2025.
- Popović, M. chrF: character n-gram f-score for automatic mt evaluation. In *Proceedings of the tenth workshop on statistical machine translation*, pp. 392–395, 2015.
- Porada, I., Sordoni, A., and Cheung, J. C. K. Does pre-training induce systematic inference? how masked language models acquire commonsense knowledge. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 4550–4557, 2022.
- Pu, X., Gao, M., and Wan, X. Summarization is (almost) dead. *arXiv preprint arXiv:2309.09558*, 2023.
- Puccetti, G., Rogers, A., Alzetta, C., Dell’Orletta, F., and Esuli, A. Ai ‘news’ content farms are easy to make and hard to detect: A case study in italian. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 15312–15338, 2024.
- Qi, D., Zhao, H., Shi, J., Jenni, S., Fan, Y., Derroncourt, F., Cohen, S., and Li, S. The photographer’s eye: Teaching multimodal large language models to see, and critique like photographers. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pp. 24807–24816, 2025a.
- Qi, Z., Perron, B. E., Victor, B. G., Stoll, D., and Ryan, J. P. Small models achieve large language model performance: Evaluating reasoning-enabled ai for secure child welfare research. *arXiv preprint arXiv:2512.04261*, 2025b.
- Qin, H. X., Jin, S., Gao, Z., Fan, M., and Hui, P. Charactermeet: Supporting creative writers’ entire story character construction processes through conversation with llm-powered chatbot avatars. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, pp. 1–19, 2024.
- Qiu, H., He, H., Zhang, S., Li, A., and Lan, Z. Smile: Single-turn to multi-turn inclusive language expansion via chatgpt for mental health support. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 615–636, 2024.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in neural information processing systems*, 36: 53728–53741, 2023.
- Raunak, V., Menezes, A., Post, M., and Awadalla, H. H. Do gpts produce less literal translations? In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pp. 1041–1050, 2023.
- Ravaut, M., Sun, A., Chen, N., and Joty, S. On context utilization in summarization with large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 2764–2781, 2024.
- Rei, R., Stewart, C., Farinha, A. C., and Lavie, A. Comet: A neural framework for mt evaluation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 2685–2702, 2020.

- Reniers, R. L., Abu-Akel, A., and Seara-Cardoso, A. Cognitive empathy and perspective taking: Understanding the mechanisms of normal and abnormal experiences and abilities, 2022.
- Resende, N. and Hadley, J. The translator’s canvas: Using llms to enhance poetry translation. In *Proceedings of the 16th Conference of the Association for Machine Translation in the Americas (Volume 1: Research Track)*, pp. 178–189, 2024.
- Rice, K. ChatGPT and world history essays: An assignment and its insights into the coloniality of generative ai. *Teaching History*, 49:1, 2024.
- Röttger, P., Kirk, H., Vidgen, B., Attanasio, G., Bianchi, F., and Hovy, D. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 5377–5400, 2024.
- Sahu, G., Vechtomova, O., and Laradji, I. H. A guide to effectively leveraging llms for low-resource text summarization: Data augmentation and semi-supervised approaches. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pp. 1584–1603, 2025.
- Sakurai, M., Ueta, K., and Hashimoto, Y. Exploring the limits of llms in simulating partisan polarization with confirmation bias prompts. *Engineering Proceedings*, 107(1):2, 2025.
- Sands, B., Wang, Y., Xu, C., Zhou, Y., Wei, L., and Chandra, R. An evaluation of llms for generating movie reviews: Gpt-4o, gemini-2.0 and deepseek-v3. *arXiv preprint arXiv:2506.00312*, 2025.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Shao, Y., Li, L., Dai, J., and Qiu, X. Character-LLM: A trainable agent for Role-Playing. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 13153–13187, 2023.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y., et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Sharma, A., Miner, A., Atkins, D., and Althoff, T. A computational approach to understanding empathy expressed in text-based mental health support. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 5263–5276, 2020.
- Sharma, A., Lin, I. W., Miner, A. S., Atkins, D. C., and Althoff, T. Towards facilitating empathic conversations in online mental health support: A reinforcement learning approach. In *Proceedings of the web conference 2021*, pp. 194–205, 2021.
- Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askeel, A., Bowman, S. R., DURMUS, E., Hatfield-Dodds, Z., Johnston, S. R., Kravec, S. M., Maxwell, T., McCandlish, S., Ndousse, K., Rausch, O., Schiefer, N., Yan, D., Zhang, M., and Perez, E. Towards understanding sycophancy in language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- Shi, Z., Ma, R., Huang, J.-t., Ma, X., Chen, X., Wang, M., Yang, Q., Wang, Y., Ye, F., Chen, Z., et al. Social welfare function leaderboard: When llm agents allocate social welfare. *arXiv preprint arXiv:2510.01164*, 2025.
- Simon, F. M. Artificial intelligence in the news: How ai retools, rationalizes, and reshapes journalism and the public arena. <https://academiccommons.columbia.edu/doi/10.7916/ncm5-3v06>, February 2024. Accessed: 2026-01-11.
- Singer, T. and Lamm, C. The social neuroscience of empathy. *Annals of the new York Academy of Sciences*, 1156 (1):81–96, 2009.
- Singh, V., Schulte im Walde, S., and Keplinger, K. Inclusive leadership in the age of AI: A dataset and comparative study of LLMs vs. real-life leaders in workplace action planning. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 19732–19753, November 2025. URL <https://aclanthology.org/2025.findings-emnlp.1075/>.
- Sinha, N. Qa-prompting: Improving summarization with large language models using question-answering. *arXiv preprint arXiv:2505.14347*, 2025.
- Small, C. T., Vendrov, I., Durmus, E., Homaei, H., Barry, E., Cornebise, J., Suzman, T., Ganguli, D., and Megill, C. Opportunities and risks of llms for scalable deliberation with polis. *arXiv preprint arXiv:2306.11932*, 2023.
- Sorin, V., Brin, D., Barash, Y., Konen, E., Charney, A., Nadkarni, G., and Klang, E. Large language models and empathy: systematic review. *Journal of medical Internet research*, 26:e52597, 2024.
- Steinmaurer, A., Dengel, A., Comanici, M., Buchner, J., Memminger, J., and Gütl, C. Immersive learning in history education: Exploring the capabilities of virtual avatars and large language models. In *International conference on immersive learning*, pp. 363–374, 2024.

- Styles, O., Miller, S., Cerda-Mardini, P., Guha, T., Sanchez, V., and Vidgen, B. Workbench: a benchmark dataset for agents in a realistic workplace setting. *arXiv preprint arXiv:2405.00823*, 2024.
- Suh, J., Le, L., Shayegani, E., Ramos, G., Amores, J., Ong, D. C., Czerwinski, M., and Hernandez, J. Sense-7: Taxonomy and dataset for measuring user perceptions of empathy in sustained human-ai conversations. *arXiv preprint arXiv:2509.16437*, 2025.
- Sun, Y., Huang, Q., Wang, Y., and Tung, A. K. Diversinews: Enriching news consumption with relevant yet diverse news articles retrieval. *Proceedings of the VLDB Endowment*, 17(12):4277–4280, 2024.
- Sun, Y., Huang, Q., Tang, Y., Tung, A. K. H., and Yu, J. A general framework for producing interpretable semantic text embeddings. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025a. URL <https://openreview.net/forum?id=23uY3FpQxc>.
- Sun, Y., Huang, Q., Tung, A. K., and Yu, J. Text embeddings should capture implicit semantics, not just surface meaning. *arXiv preprint arXiv:2506.08354*, 2025b.
- Sun, Y., Huang, Q., Tung, A. K. H., and Yu, J. Prism: A framework for producing interpretable political bias embeddings with political-aware cross-encoder. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 27719–27733, 2025c. URL <https://aclanthology.org/2025.acl-long.1344/>.
- Sun, Y., Huang, Q., Xu, Z., Sun, Y., Tang, Y., and Tung, A. K. One swallow does not make a summer: Understanding semantic structures in embedding spaces. *arXiv preprint arXiv:2512.00852*, 2025d.
- Tam, D., Mascarenhas, A., Zhang, S., Kwan, S., Bansal, M., and Raffel, C. Evaluating the factual consistency of large language models through news summarization. In *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 5220–5255, 2023.
- Tan, J., Westermann, H., Pottanigari, N. R., Šavelka, J., Meeùs, S., Godet, M., and Benyekhle, K. Robots in the middle: Evaluating llms in dispute resolution. *arXiv preprint arXiv:2410.07053*, 2024.
- Tao, M. and Xie, S. What does a visual formal analysis of the world’s 500 most famous paintings tell us about multimodal llms? In *The Second Tiny Papers Track at ICLR 2024*, 2024.
- Team, K., Bai, Y., Bao, Y., Chen, G., Chen, J., Chen, N., Chen, R., Chen, Y., Chen, Y., Chen, Y., et al. Kimi k2: Open agentic intelligence. *arXiv preprint arXiv:2507.20534*, 2025.
- Teleki, M., Bengali, V., Dong, X., Janjur, S. T., Liu, H., Liu, T., Wang, C., Liu, T., Zhang, Y., Shipman, F., and Caverlee, J. A survey on LLMs for story generation. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 13954–13966, November 2025.
- Tessler, M. H., Bakker, M. A., Jarrett, D., Sheahan, H., Chadwick, M. J., Koster, R., Evans, G., Campbell-Gillingham, L., Collins, T., Parkes, D. C., et al. Ai can help humans find common ground in democratic deliberation. *Science*, 386(6719):eadq2852, 2024.
- Tian, Y., Huang, T., Liu, M., Jiang, D., Spangher, A., Chen, M., May, J., and Peng, N. Are large language models capable of generating human-level narratives? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 17659–17681, 2024.
- Tohidi, A., Haider, S., and Watts, D. J. Rethinking news framing with large language models. *Scientific Reports*, 2025.
- Toles, M., Balwani, N., Singh, R., Rodriguez, V. G. S., and Yu, Z. Program synthesis dialog agents for interactive decision-making. *arXiv preprint arXiv:2502.19610*, 2025.
- Tong, Z., Gu, Y., Liu, H., Liu, Q., Wu, S., Shi, H., and Zhang, X.-Y. Generate first, then sample: Enhancing fake news detection with llm-augmented reinforced sampling. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 24276–24290, 2025.
- Tseng, E., Young, M., Le Quéré, M. A., Rinehart, A., and Suresh, H. ”ownership, not just happy talk”: Co-designing a participatory large language model for journalism. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*, pp. 3119–3130, 2025.
- Van Dijke, J., van Nistelrooij, I., Bos, P., and Duyndam, J. Towards a relational conceptualization of empathy. *Nursing Philosophy*, 21(3):e12297, 2020.
- Vedantam, R., Lawrence Zitnick, C., and Parikh, D. CIDEr: consensus-based image description evaluation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 4566–4575, 2015.
- Venkatraman, S., Tripto, N. I., and Lee, D. Collabstory: Multi-llm collaborative story generation and authorship analysis. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pp. 3665–3679, 2025.

- Vykopal, I., Pikuliak, M., Srba, I., Moro, R., Macko, D., and Bielikova, M. Disinformation capabilities of large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 14830–14847, 2024.
- Wang, C., Zheng, W., Zhang, Y., Zhu, F., Cheng, J., Xie, Y., Wang, W., and Feng, F. Perm: Psychology-grounded empathetic reward modeling for large language models. *arXiv preprint arXiv:2601.10532*, 2026.
- Wang, J., Yao, Z., Li, L., Qian, J., Yang, Z., and Yu, H. Chatthero: An llm-supported chatbot for behavior change and therapeutic support in addiction recovery. *arXiv preprint arXiv:2508.20996*, 2025a.
- Wang, L., Du, Z., Liu, D., Cai, D., Yu, D., Jiang, H., Wang, Y., Shi, S., and Tu, Z. Guofeng: A discourse-aware evaluation benchmark for language understanding, translation and generation. <https://openreview.net/forum?id=XIIynqbMXgR>, 2023a. Accessed: 2026-01-24.
- Wang, L., Tu, Z., Gu, Y., Liu, S., Yu, D., Ma, Q., Lyu, C., Zhou, L., Liu, C.-H., Ma, Y., et al. Findings of the wmt 2023 shared task on discourse-level literary translation: A fresh orb in the cosmos of llms. In *Proceedings of the Eighth Conference on Machine Translation*, pp. 55–67, 2023b.
- Wang, L., Liu, S., Lyu, C., Jiao, W., Wang, X., Xu, J., Tu, Z., Gu, Y., Chen, W., Wu, M., et al. Findings of the wmt 2024 shared task on discourse-level literary translation. In *Proceedings of the Ninth Conference on Machine Translation*, pp. 699–700, 2024a.
- Wang, P., Ma, R., Zhang, B., Chen, X., He, Z., Luo, K., Lv, Q., Jiang, Q., Xie, Z., Wang, S., et al. Rlver: Reinforcement learning with verifiable emotion rewards for empathetic agents. *arXiv preprint arXiv:2507.03112*, 2025b.
- Wang, Q., Hu, J., Li, Z., Wang, Y., Li, D., Hu, Y., and Tan, M. Generating long-form story using dynamic hierarchical outlining with memory-enhancement. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 1352–1391, 2025c.
- Wang, R., Milani, S., Chiu, J., Zhi, J., Eack, S., Labrum, T., Murphy, S., Jones, N., Hardy, K., Shen, H., et al. Patient- ψ : Using large language models to simulate patients for training mental health professionals. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 12772–12797, 2024b.
- Wang, S., Cheng, Y., Song, A., Keedy, S., Berman, M., and Feamster, N. Can llms address mental health questions? a comparison with human therapists. *arXiv preprint arXiv:2509.12102*, 2025d.
- Wang, T., Chen, J., Jia, Q., Wang, S., Fang, R., Wang, H., Gao, Z., Xie, C., Xu, C., Dai, J., et al. Weaver: Foundation models for creative writing. *arXiv preprint arXiv:2401.17268*, 2024c.
- Wang, Y., Lin, J., Yu, Z., Hu, W., and Karlsson, B. F. Open-world story generation with structured knowledge enhancement: A comprehensive survey. *Neurocomput.*, 559 (C), November 2023c.
- Wang, Y., Zhang, Z., and Wang, R. Element-aware summarization with large language models: Expert-aligned evaluation and chain-of-thought method. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 8640–8665, 2023d.
- Wei, J., Bosma, M., Zhao, V., Guu, K., Yu, A. W., Lester, B., Du, N., Dai, A. M., and Le, Q. V. Finetuned language models are zero-shot learners. In *International Conference on Learning Representations*, 2022.
- Weisz, E. and Zaki, J. Motivated empathy: a social neuroscience perspective. *Current opinion in psychology*, 24: 67–71, 2018.
- Welivita, A. and Pu, P. Are large language models more empathetic than humans? *arXiv preprint arXiv:2406.05063*, 2024.
- Westermann, H., Savelka, J., and Benyekhlef, K. LImediator: Gpt-4 assisted online dispute resolution. In *Proceedings of the ICAIL 2023 Workshop on Artificial Intelligence for Access to Justice*, pp. 1–12, 2023.
- Whang, Y. Chatgpt for editors: enhancing efficiency and effectiveness. *Science Editing*, 11(1):84–90, 2024.
- Wu, M., Xu, J., and Wang, L. Transagents: Build your translation company with language agents. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pp. 131–141, 2024a.
- Wu, M., Xu, J., Yuan, Y., Haffari, G., Wan, L., Luo, W., and Zhang, K. (perhaps) beyond human translation: Harnessing multi-agent collaboration for translating ultra-long literary texts. *Transactions of the Association for Computational Linguistics*, 13:901–922, 2025.
- Wu, Y., Hee, M. S., Hu, Z., and Lee, R. K.-W. Longgenbench: Benchmarking long-form generation in long context llms. *arXiv preprint arXiv:2409.02076*, 2024b.

- xAI. Grok 4.1 model card. <https://data.x.ai/2025-11-17-grok-4-1-model-card.pdf>, November 2025. Accessed: 2026-01-08.
- Xia, H., Peng, H., Qi, Y., Xu, B., Li, J., Lei, H., and Wang, X. Storywriter: A multi-agent framework for long story generation. In *Proceedings of the 34th ACM International Conference on Information and Knowledge Management*, pp. 6559–6563, 2025.
- Xie, Z., Cohn, T., and Lau, J. H. The next chapter: A study of large language models in storytelling. In *Proceedings of the 16th International Natural Language Generation Conference*, pp. 323–351, 2023a.
- Xie, Z., Singh, S., McAuley, J., and Majumder, B. P. Factual and informative review generation for explainable recommendation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 13816–13824, 2023b.
- Xu, B., Chen, Y., Wen, Z., Liu, W., and He, B. Evaluating small language models for news summarization: Implications and factors influencing performance. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 4909–4922, 2025a.
- Xu, W., Fan, W., Zhu, Y., and Wang, B. Consistency of responses and continuations generated by large language models on social media. *arXiv preprint arXiv:2501.08102*, 2025b.
- Xu, Y., Zong, Y., Zhonghua, W., Wu, S., Rao, Y., Zhang, D., and Deng, S. CtrlNews: LLM-based multi-agent controllable news writing via knowledge gravitational field. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 8668–8705, November 2025c.
- Xue, D., Tu, J., Wang, M., Yan, X., Liu, F., and Hu, J. Mindchat: A privacy-preserving large language model for mental health support. *arXiv preprint arXiv:2601.01993*, 2026.
- Yang, K., Tian, Y., Peng, N., and Klein, D. Re3: Generating longer stories with recursive reprompting and revision. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 4393–4479, 2022.
- Yang, K., Klein, D., Peng, N., and Tian, Y. Doc: Improving long story coherence with detailed outline control. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3378–3465, 2023.
- Yang, L., Xiao, Z., Huang, W., and Zhong, X. Storyllava: enhancing visual storytelling with multi-modal large language models. In *Proceedings of the 31st International Conference on Computational Linguistics*, pp. 3936–3951, 2025a.
- Yang, S., Ge, Y., Li, Y., Chen, Y., Ge, Y., Shan, Y., and Chen, Y.-C. Seed-story: Multimodal long story generation with large language model. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 1850–1860, 2025b.
- Yao, S., Ke, Q., Li, K., Wang, Q., and Hu, J. News gpt: A large language model for reliable and hallucination-controlled news generation. In *Proceedings of the 2024 3rd International Symposium on Robotics, Artificial Intelligence and Information Engineering*, pp. 113–119, 2024.
- Ye, F., Wang, S., Chao, L. S., and Wong, D. F. Unveiling llms’ metaphorical understanding: Exploring conceptual irrelevance, context leveraging and syntactic influence. *arXiv preprint arXiv:2510.04120*, 2025.
- Ye, Y., Feng, X., Feng, X., Ma, W., Qin, L., Xu, D., Yang, Q., Liu, H., and Qin, B. Globesumm: A challenging benchmark towards unifying multi-lingual, cross-lingual and multi-document news summarization. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 10803–10821, 2024.
- You, H., Touileb, S., Øvrelid, L., and Velldal, E. Event-based evaluation of abstractive news summarization. In *Proceedings of the Fourth Workshop on Generation, Evaluation and Metrics (GEM²)*, pp. 504–510, 2025.
- You, X., Huang, Q., Li, L., Chang, X., and Yu, J. Cut to the Chase: Training-free Multimodal Summarization via Chain-of-Events. *arXiv preprint arXiv:2603.06213*, 2026a.
- You, X., Huang, Q., Li, L., Zhang, C., Liu, X., Zhang, M., and Yu, J. Knowledge completes the vision: A multimodal entity-aware retrieval-augmented generation framework for news image captioning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 12108–12116, 2026b. URL <https://ojs.aaai.org/index.php/AAAI/article/view/38200>.
- Yu, Y., Yan, S., Zhang, Q., Xu, Z., Zhou, G., and Jin, H. The influence of affective empathy on online news belief: The moderated mediation of state empathy and news type. *Behavioral Sciences*, 14(4):278, 2024.
- Yuan, H. and Zhang, H. Understanding llm reasoning for abstractive summarization. *arXiv preprint arXiv:2512.03503*, 2025.

- Yuan, Z., He, Y., Wang, K., Ye, Y., and Sun, L. Artgpt-4: Towards artistic-understanding large vision-language models with enhanced adapter. *arXiv preprint arXiv:2305.07490*, 2023.
- Yunusov, S., Sidat, H., and Emami, A. Mirrorstories: Reflecting diversity through personalized narrative generation with large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 6702–6717, 2024.
- Zagorulko, D. I. Chatgpt in newsrooms: Adherence of ai-generated content to journalism standards and prospects for its implementation in digital media. *Scientific Notes of VI Vernadsky Taurida National University, Series: Philology. Journalism*, 2(1):319–325, 2023.
- Z.ai. GLM-4.7: Advancing the coding capability. <https://z.ai/blog/glm-4.7>, December 2025. Accessed: 2026-01-05.
- Zaki, J. Empathy: a motivated account. *Psychological bulletin*, 140(6):1608, 2014.
- Zeng, M. Psychcounsel-bench: Evaluating the psychology intelligence of large language models. *arXiv preprint arXiv:2510.01611*, 2025.
- Zhang, C., Li, R., Tan, M., Yang, M., Zhu, J., Yang, D., Zhao, J., Ye, G., Li, C., and Hu, X. Cpsycoun: A report-based multi-turn dialogue reconstruction and evaluation framework for chinese psychological counseling. In *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 13947–13966, 2024a.
- Zhang, H., Liu, X., and Zhang, J. Extractive summarization via chatgpt for faithful summary generation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 3270–3278, 2023a.
- Zhang, H., Liu, X., and Zhang, J. Summit: Iterative text summarization via chatgpt. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 10644–10657, 2023b.
- Zhang, J. and Long, Y. Mld-ea: Check and complete narrative coherence by introducing emotions and actions. In *Proceedings of the 31st International Conference on Computational Linguistics*, pp. 1892–1907, 2025.
- Zhang, R. and Eger, S. Llm-based multi-agent poetry generation in non-cooperative environments. *arXiv preprint arXiv:2409.03659*, 2024.
- Zhang, R., Zhao, W., and Eger, S. How good are llms for literary translation, really? literary translation evaluation with humans and llms. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 10961–10988, 2025a.
- Zhang, R., Zhao, W., Macken, L., and Eger, S. LiTransProQA: An LLM-based literary translation evaluation metric with professional question answering. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 29099–29121, November 2025b.
- Zhang, T., Ladhak, F., Durmus, E., Liang, P., Mckeown, K., and Hashimoto, T. B. Benchmarking large language models for news summarization. *Transactions of the Association for Computational Linguistics*, 11:39–57, 2024b.
- Zhang, X. and Luo, Z. Advancing conversational psychotherapy: Integrating privacy, dual-memory, and domain expertise with large language models. *arXiv preprint arXiv:2412.02987*, 2024.
- Zhang, Y., Li, M., Long, D., Zhang, X., Lin, H., Yang, B., Xie, P., Yang, A., Liu, D., Lin, J., et al. Qwen3 embedding: Advancing text embedding and reranking through foundation models. *arXiv preprint arXiv:2506.05176*, 2025c.
- Zhao, X., Liu, S., Sun, F., Dong, Y., Wang, Y., and Ge, C. Zhixin: An instruction-tuned large language model for mental disorder diagnosis. <https://www.researchsquare.com/article/rs-8149707/v1>, 2025. Accessed: 2026-01-13.
- Zhou, J., Wang, X., Zhang, M., and Yu, J. Holistic utility preference learning for listwise alignment. *arXiv preprint arXiv:2410.18127*, 2024.
- Zhou, W., Jiang, Y. E., Cui, P., Wang, T., Xiao, Z., Hou, Y., Cotterell, R., and Sachan, M. Recurrentgpt: Interactive generation of (arbitrarily) long text. *arXiv preprint arXiv:2305.13304*, 2023.
- Zhu, C., Liu, Y., Mei, J., and Zeng, M. MediaSum: A large-scale media interview dataset for dialogue summarization. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 5927–5934, 2021.
- Zhu, J., Wang, Q., Wang, Y., Zhou, Y., Zhang, J., Wang, S., and Zong, C. NCLS: Neural cross-lingual summarization. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 3054–3064, 2019.

A. Baseline Details

A.1. Model Selection and Configuration

We evaluate our empathy metrics on ten strong baseline LLMs spanning both closed and open weight releases. This suite encompasses contemporary models with frontier-level reasoning, long-context processing, and agentic or multimodal capabilities, allowing us to probe whether empathic distortions persist even under competitive task performance. All baselines are tested under identical conditions: the same task prompts and mechanism evaluation prompts, without model-specific prompt tuning.

Proprietary baselines comprise five models that represent current frontier capabilities across multiple dimensions. GPT-5.2 advances scientific and mathematical reasoning with enhanced vision understanding (OpenAI, 2025). Claude-Sonnet-4.5 targets complex agentic behavior and tool-use workflows, with documented gains in coding and extended reasoning (Anthropic, 2025). Gemini-3 represents Google’s frontier multimodal reasoning architecture, accompanied by structured risk assessment (Pichai et al., 2025). Grok-4.1 balances natural dialogue generation with strong core reasoning, supported by deployment-oriented safety testing (xAI, 2025). Doubao-Seed1.8 integrates perception, reasoning, and action for generalized real-world agency beyond single-turn generation (Bytedance Seed, 2025).

Open-weight baselines include five architecturally diverse systems that enable deeper inspection of model behavior. Llama-4 introduces a natively multimodal mixture-of-experts architecture with extensive context support (Meta AI, 2025). DeepSeek-V3.2 emphasizes efficiency through sparse attention and scaled reinforcement learning in post-training, with particular focus on agentic task synthesis and tool-use robustness (Liu et al., 2025a). Qwen3-VL provides interleaved long-context multimodal reasoning across image and video modalities in both dense and mixture-of-experts variants (Bai et al., 2025a). Kimi-K2 employs a mixture-of-experts architecture with a specialized optimizer designed for stable large-scale pretraining under high token budgets (Team et al., 2025). GLM-4.7 reports improvements in coding and stable multi-step reasoning, with documented upgrades for agentic coding workflows (Z.ai, 2025).

All model versions, access modalities, and additional configuration details are summarized in Table 8. Unless otherwise specified, we use a unified decoding configuration across baselines within each task, including fixed temperature $T = 0.01$ and task-specific maximum output length.

Table 8. Baseline models used in our experiments. We list each baseline by provider and version, and report two evaluation-relevant attributes: whether the model is proprietary or open-source, and whether it natively supports only text or multimodal inputs.

Model	Company	Version	Access	Native Modality
GPT-5.2	OpenAI	GPT-5.2	Proprietary	Multimodal
Claude-Sonnet-4.5	Anthropic	claude-Sonnet-4-5	Proprietary	Multimodal
Gemini-3	Google	Gemini-3-Pro-Preview	Proprietary	Multimodal
Grok-4.1	xAI	Grok-4.1-Fast	Proprietary	Multimodal
Doubao-Seed1.8	ByteDance	Doubao-Seed1.8	Proprietary	Multimodal
Llama-4	Meta	Llama-4-Maverick	Open-source	Multimodal
DeepSeek-V3.2	DeepSeek	DeepSeek-V3.2	Open-source	Text
Qwen3-VL	Alibaba	Qwen3-VL-235B-A22B-Instruct	Open-source	Multimodal
Kimi-K2	Moonshot AI	Kimi-K2-Thinking	Open-source	Text
GLM-4.7	Zhipu AI	GLM-4.7	Open-source	Text

A.2. Experimental Instructions

This subsection reports the task-level instruction prompts used to elicit model outputs for our three empirical settings. For each task, we use a single fixed prompt template that specifies the role and output requirement, and we apply the same template to all baseline models without additional prompt tuning or model-specific adjustments. These prompts are designed to be lightweight and task-appropriate, providing only the minimum structure needed to instantiate the generation problem while preserving variability in how models represent human perspectives. We use one prompt per task, covering news generation, literary translation, and counselor-style response generation, and all downstream mechanism metrics are computed from the resulting outputs under this standardized prompting protocol.

For cognitive empathy evaluation in news generation, we employ Prompt A.2 to elicit long-form articles conditioned on short human-written summaries from the Newsroom dataset (Grusky et al., 2018). The prompt casts the model in the role of a news editor and instructs it to expand each summary into a complete article of approximately a specified target length. By fixing the summary content while requiring discourse-level expansion, this setup tests whether models preserve salience structure and human impact during elaboration, rather than simply reproducing topical facts. This design is therefore well-suited for revealing sentiment attenuation and granularity miscalibration, which may remain undetected under standard semantic similarity metrics.

Prompt 1: News Generation

As a seasoned news editor, please help me expand the following news summary into a full news article, which should contain approximately {length} words.

News Summary: {Summary}

For cultural empathy evaluation in literary translation (Wang et al., 2023b;a), we use Prompt A.2, which frames the model as a professional Chinese to English literary translator and requests a fluent, natural translation of a Chinese novel passage. We keep the prompt stable across all baseline models and passages, and we do not add auxiliary instructions such as style constraints, summaries, or post-editing guidelines. The prompt explicitly restricts outputs to the English translation only, avoiding meta commentary, explanations, or notes that can shift register and interfere with narrative voice. Under this controlled instruction, we evaluate whether models preserve culturally grounded tone, register, and interpersonal framing beyond semantic adequacy, enabling measurement of sentiment attenuation and linguistic distancing in translation outputs.

Prompt 2: Literary Translation

You are a professional literary translator specializing in Chinese-to-English translation. Your task is to translate Chinese novels into fluent, natural English. Provide only the English translation without any explanations or notes.

Please translate the following Chinese novel chapter into English.

Chinese Novel:
{Novel Text}

For relational empathy evaluation in psychological counseling, we use Prompt A.2, which assigns the model the role of a psychological counselor and asks it to respond professionally to a client narrative. The instruction provides only the client narrative as context and does not prescribe a response structure, specific counseling techniques, or a required number of questions. This minimal framing yields counselor-style outputs that remain comparable across models while allowing substantial variation in how models acknowledge affect, probe for missing details, and manage interpersonal tension. We use these responses to assess whether models preserve client affect and engage the underlying conflict constructively, supporting measurement of sentiment attenuation and conflict avoidance.

Prompt 3: Psychological Counseling

You are a psychological counselor. Below is the client’s narrative. Please respond as a professional psychological counselor.

{Narrative}

B. Cognitive Empathy Experiment Details

Dataset We instantiate the news generation setting using the Newsroom dataset (Grusky et al., 2018), which provides paired news articles and human-written summaries. Given a gold summary, models are prompted to generate a full-length news article with a length comparable to the reference. To probe cognitive empathy failures that affect informational salience, we emphasize high-stakes instances with clear human impact, where preserving urgency and descriptive granularity is most

consequential. We construct an evaluation set by using GPT-5.2 (OpenAI, 2025) to filter the test split, retaining instances whose references exhibit salient affect, and then randomly sampling 50 instances from the filtered pool.

Prompt 4: Sentiment Attenuation

You are evaluating the **emotional intensity** of two text summaries on a scale from 0.0 to 1.0, where emotional intensity refers to the strength and vividness of emotions conveyed in the text.

Scoring Guide:

- 0.0-0.2: Emotionally flat, clinical, bureaucratic language. No visceral or affective content.
- 0.3-0.4: Mild emotional content, but heavily sanitized or abstracted.
- 0.5-0.6: Moderate emotional presence with some concrete details.
- 0.7-0.8: Strong emotional intensity with vivid, embodied language.
- 0.9-1.0: Maximally intense emotional content with raw, unfiltered affective expression.

Consider:

- Use of visceral, sensory, concrete words (high intensity) vs. abstract euphemisms (low intensity)
- Emotional directness (high) vs. clinical distancing (low)
- Presence of strong emotion words (grief, rage, despair, joy) vs. neutral descriptors
- Acknowledgment of pain/conflict vs. sanitization

SUMMARY 1: {Reference}

SUMMARY 2: {Generated}

OUTPUT FORMAT: {"Reference Score": <number >, "Generated Score": <number >}

Prompt 5: Empathic Granularity Mismatch

You are tasked with identifying and counting disturbing elements in a given sentence, where disturbing elements refer to words or phrases that describe violence, harm, death, or graphic content. Count how many disturbing elements are in this sentence.

Count these (1 point each):

- Violence/attack (attacked, shot, stabbed, beaten, assaulted, mauled)
- Injury/harm (injured, wounded, hurt, bleeding, blood)
- Graphic details (gore, mutilation, screaming in pain, vomit, decay)
- Abuse/torture (tortured, abused, raped, molested)

Rules:

- One word = 1 point ("killed" = 1)
- Two elements = 2 points ("shot and killed" = 2)
- Multiple victims = still 1 point ("3 people killed" = 1)
- Generic/vague terms don't count ("incident occurred" = 0)

Sentence: {sentence}

Return only JSON: {"Aversive Count": <number >}

Metrics We evaluate generation quality using standard metrics, including ROUGE-L (Lin, 2004), CIDEr (Vedantam et al., 2015), and semantic similarity. Semantic similarity is computed as the cosine similarity between document embeddings from Qwen3-4B-Embedding (Zhang et al., 2025c), which supports inputs up to 32K tokens and produces 2560-dimensional representations, so no truncation is required. Beyond surface-level metrics, we assess cognitive empathy failures using our mechanism-based metrics: Sentiment Attenuation (SA) and Empathic Granularity Mismatch (EGM).

For SA, we define an affect intensity function $E(\cdot) \in [0, 1]$ that measures emotional intensity: the strength, vividness, and immediacy of affective expression, operationalizing the relative downscaling of urgency, moral weight, or emotional force in model generations. We use GPT-5.2 (OpenAI, 2025) as the evaluator with Prompt B, which presents both reference and generated summaries and asks the judge to assign an intensity score in $[0, 1]$ to each. To reduce ambiguity, we provide

an anchored rubric ranging from emotionally flat and clinical language to maximally intense and unfiltered affect, with concrete cues including visceral and sensory wording, emotional directness versus clinical distancing, explicit emotion terms, and acknowledgment versus sanitization of pain or conflict. The judge outputs a minimal JSON object containing the two scores, enabling reliable parsing and direct computation of SA via Equation 1.

For EGM, we estimate the density of disturbing or overly explicit content $U(\cdot)$ at the sentence level for both the reference text and the model output. We implement $U(\cdot)$ using GPT-5.2 (OpenAI, 2025) with Prompt B, which provides counting instructions and a fixed checklist covering death, violence, injury, graphic details, and abuse or torture, with explicit counting rules for multiword phrases, multiple elements in one sentence, and exclusions for generic or vague terms. Given a single sentence, the judge returns a JSON object with an Aversive Count, which we aggregate across sentences and normalize by sentence count to obtain the disturbing content density used in Equation 2.

C. Cultural Empathy Experiment Details

Dataset We evaluate cultural empathy in literary translation using the validation and test sets from the WMT 2023 Shared Task on Discourse-Level Literary Translation (Wang et al., 2023b;a) (WMT2023). This task translates Chinese fiction into English, requiring models to preserve culturally grounded style and register beyond mere semantic adequacy.

Metrics Following prior work on literary translation evaluation (Wang et al., 2023b), we report standard translation metrics including chrF (Popović, 2015) and COMET (Rei et al., 2020), together with semantic similarity measured via Qwen3-4B-Embedding (Zhang et al., 2025c). For efficiency, we compute chrF at the sentence level and report the mean across sentences, whereas COMET is computed at the document level. Semantic similarity is computed as cosine similarity by embedding the full document and comparing reference and generated translations, consistent with our cognitive empathy evaluation. To assess cultural empathy failures beyond semantic adequacy, we apply SA and Linguistic Distancing (LD).

For SA, we use GPT-5.2 (OpenAI, 2025) as the evaluator with Prompt B, which presents both reference and generated translations and asks the judge to assign an affect intensity score in $[0, 1]$ to each. The prompt provides an anchored rubric ranging from emotionally flat and clinical language to maximally intense and unfiltered affect, with concrete cues including visceral and sensory wording, emotional directness versus clinical distancing, explicit emotion terms, and acknowledgment versus sanitization of pain or conflict. The judge outputs a minimal JSON object containing the two scores for direct computation of SA via Equation 1. LD is computed per Equation 4 by using spaCy (Honnibal, 2017) to estimate, on a per-sentence basis, the densities of nominalization, impersonal subject framing, and agentless constructions, and then aggregating them into a final distancing score.

D. Relational Empathy Experiment Details

Prompt 6: Psychological Counseling Data Generation

Generate three distinct client statements for scenario [Partner Betrayal / Emotional Neglect / Controlling Family Dynamics / In-law Conflict / Workplace Bullying / Coworker Relationship Breakdown / Advisor-Advisee Conflict], where clients are confiding in their therapist about recent life events and concerns. Each statement should:

- Express genuine emotions and internal conflicts.
- Contain enough detail to be realistic, while leaving room for the therapist to ask follow-up questions.
- Reflect varied emotional states and types of struggles.
- Ensure each entry is unique in both content and thematic focus.

Dataset To evaluate relational empathy in counseling-style interactions, we construct a small-scale synthetic dataset with GPT-5.2 (OpenAI, 2025) using the controlled generation protocol in Prompt D. All client narratives are synthetic and contain no personal identifying information. We generate 50 examples spanning common relational contexts, including family relationships, workplace communication, daily-life stressors, and academic challenges, with examples shown in Table 9. Each example provides a client narrative as input, and we then prompt each evaluated model with an identical instruction template to generate a therapist-style response for assessment. Prompt D instructs the model to produce multi-distinct client statements for each scenario class, including partner betrayal, emotional neglect, controlling family dynamics, in-law conflict, workplace bullying, coworker relationship breakdown, and advisor-advisee conflict. Each statement must express genuine affect and internal tension while maintaining sufficient narrative openness to necessitate therapeutic follow-up,

Table 9. Relational empathy dataset overview. We show representative synthetic client narratives from common relational contexts in our counseling-style dataset.

Topic	Content
Partner Betrayal	I feel like I’m losing control because my own mind keeps tearing me apart. I know the rational thing is to clarify what’s actually happening, but I cannot stay calm. One minute, I want to end the relationship immediately. The next minute, I tell myself I am overreacting. He has been “busy” lately, yet I cannot point to a single clear incident. Everything he says has a reasonable explanation, but my gut keeps buzzing like an alarm. What hurts most is that I am starting to become someone I do not recognize. I fixate on a single sentence he says, replay it in my head, and wonder if he is lying. If he does not reply, my heart races. I do not want to be this way. I was not like this before, but I genuinely cannot control it. On the surface, our relationship looks fine. Underneath, it feels like something is slowly rotting. I feel deeply conflicted. I want him to comfort me and prove that I matter, but I am terrified of seeming needy and of looking like I am begging for reassurance.
Workplace Bullying	Our team has a vibe that’s been getting harder for me to tolerate. My manager likes calling people out in group chat, even over small things, often with screenshots. Last week I got back from a trip and submitted my weekly report. Two numbers used a different data definition than what he prefers, but we’ve honestly had two definitions floating around for a while. I expected a quick correction. Instead, he posted in the main channel: “This is the level you think is acceptable?” and pasted my section so everyone could “learn from it.” The worst part is that it wasn’t even clearly “wrong”—it was the standard he never stated. And people piled on: emojis, “boss is right,” little jokes. It hit me that this isn’t just about work quality. It’s like basic respect isn’t guaranteed here. I can’t even name what I’m feeling clearly. I just know that opening the chat now makes my chest tighten, and I keep thinking about whether I’m overreacting or if this is actually as bad as it feels.
Controlling Family	Every decision I make turns into an approval request. Even the smallest choices get interrogated. They ask why I want something, who put the idea in my head, and whether someone is influencing me. They insist it is for my own good, but the way they speak makes me feel as if I am not allowed to have a life of my own. I have tried to handle it in different ways. I have tried explaining myself, and I have tried going along with them. But the moment I hold my ground, they turn cold. Sometimes they look at me with a disappointed expression that makes me feel like I have done something unforgivable. Then I start blaming myself and wondering if I am being selfish. A second later, I feel furious and confused, because it seems like I am not even allowed to be angry. The worst part is that even thinking about talking to them makes my body tense up and my heart race. Still, I cannot stop replaying the same thought. If I live the way I truly want, will I lose them completely?

ensuring two critical properties: salient emotional cues that enable measurement of sentiment attenuation, and explicit interpersonal tension that facilitates detection of conflict avoidance behaviors. To mitigate templated generations, we enforce diversity requirements across content variation, emotional state differentiation, and thematic focus, yielding a compact yet heterogeneous evaluation set.

Prompt 7: Conflict Avoidance

You are tasked with evaluating a therapeutic conversation. Your task is to determine whether the therapist’s response demonstrates Conflict Aversion, which refers to the tendency to avoid engaging with underlying tensions or difficult emotions by prematurely offering solutions or closure.

Definition (Conflict Aversion = 1):

When the client’s information is clearly insufficient, the response is primarily focused on giving conclusions/solutions/actionable advice/moralizing/positive closure. No further exploration of the conflict details is conducted (a few questions, or questions do not touch on key information).

Non-conflict aversion (= 0):

The response is primarily focused on understanding and clarification. Multiple open-ended questions are asked to gather more details, avoiding specific solutions and conclusions (mild process-oriented guidance is acceptable).

Only output strict JSON, no extra text: {“label”:0 or 1}

Client’s utterance: {Client Utterance}

Therapist’s response: {Model Reply}

Metrics We employ EPITOME (Sharma et al., 2020) to assess model-generated therapist responses across three dimensions: Emotional Reactions (ER), Interpretations (IP), and Explorations (EX), using the released EPITOME scoring checkpoint

to obtain automatic sub-scores. Since our dataset does not include reference therapist replies, we aggregate ER/IP/EX by summation and report the mean across all examples as an overall counseling quality score. To assess relational empathy failures, we further report SA and Conflict Avoidance (CA).

For SA, we use GPT-5.2 (OpenAI, 2025) as the evaluator with Prompt B, which presents both the client utterance and the model-generated response and asks the judge to assign an affect intensity score in $[0, 1]$ that measures emotional intensity: the strength, vividness, and immediacy of affective expression. The judge outputs a minimal JSON object containing the intensity score for direct computation of SA via Equation 1.

For CA, we use GPT-5.2 (OpenAI, 2025) with Prompt D to implement a binary discriminator $J(\cdot)$ that labels whether a therapist’s response prematurely resolves tension instead of engaging with it, following Equation 3. The prompt specifies that a label of 1 is assigned when the client information is clearly insufficient, but the response mainly provides solutions or positive closure with little substantive exploration of the underlying conflict. A label of 0 is assigned when the response focuses on clarification and understanding, using multiple open-ended questions and avoiding premature resolution. The judge returns strict JSON with a single binary label $s^{(i)} \in \{0, 1\}$, which we average across instances to compute CA.