

Concept Removal Guidance: Evidence-Calibrated Negative Guidance for Safe Diffusion Sampling

Yoonseok Choi¹ Chaeyoung Oh¹ Hyunjun Choi¹ Seokin Seo^{† 2} Kee-Eung Kim¹

Abstract

Text-to-image diffusion models remain vulnerable to adversarial prompts that elicit disallowed content, motivating reliable inference-time controls. A popular approach is negative guidance, which subtracts a negative prompt direction with a fixed weight. However, it often forces a safety–fidelity trade-off, causing artifacts or prompt drift when over-applied and failing under attacks when under-applied. Dynamic variants reweight guidance using posterior-odds signals, which can be brittle for open-vocabulary compositional prompts, while lightweight similarity-based methods ignore the evolving image evidence along the denoising trajectory. We introduce Concept Removal Guidance (CRG), a training-free method that estimates unwanted-concept presence at each diffusion step from the model’s noise predictions, and adaptively calibrates negative guidance via a closed-form constrained update enforcing a target presence threshold while minimally perturbing the conditional trajectory. Across red-teaming benchmarks, CRG reduces attack success rates while preserving benign fidelity, and extends to additional suppression targets such as artist style and violence without fine-tuning or external classifiers.

Content Warning. This paper contains examples of unsafe visual content and language, including nudity, violence, and explicit prompts, presented as part of safety research.

1. Introduction

Generative models have garnered substantial attention across a wide range of domains, including image synthe-

[†]Work done while at KAIST. ¹Kim Jaechul Graduate School of AI, KAIST, Daejeon, Republic of Korea ²Center for Humanoid Research, KIST, Seoul, Republic of Korea. Correspondence to: Yoonseok Choi <uoo77@kaist.ac.kr>.



Figure 1. Generated images of CRG and other baselines on the Artist Style removal task. The target removal style is “Van Gogh”, while unrelated artist styles should be preserved.

sis (Ramesh et al., 2021; Saharia et al., 2022; Podell et al., 2024), natural language generation (Brown et al., 2020; Google Gemini Team, 2024), audio generation (Dhariwal et al., 2020; Kong et al., 2021), and video synthesis (Hong et al., 2023; Liu et al., 2024). This surge in interest is largely driven by the remarkable capabilities of recent models to generate high-quality, user-aligned content (Ouyang et al., 2022; Wallace et al., 2024). However, these advances have also raised serious concerns about potential misuse. In particular, generative models can produce harmful content, including explicit, violent, or discriminatory material, posing significant ethical and societal risks. These risks often stem from unfiltered or insufficiently curated training data, which can embed undesirable patterns into model outputs.

A standard baseline is negative prompting, an extension of classifier-free guidance (Ho & Salimans, 2021) that incorporates negative guidance to steer the denoising trajec-



Figure 2. **Qualitative comparison of CRG and baselines on the nudity concept removal task.** We visualize generation results under adversarial prompts targeting the nudity concept. CRG effectively suppresses unsafe content, whereas baselines frequently fail to neutralize the target concept.

tory away from a specific concept (Liu et al., 2022; Ban et al., 2024). However, static negative guidance scale biases the entire sampling process, distorting benign details even when the negative concept is absent (over-regularization), yet often fails to suppress robust adversarial attacks (under-regularization).

Recently, Koulischer et al. (2025) proposed Dynamic Negative Guidance (DNG), adaptively scaling the negative guidance weight at inference time by tracking stepwise posteriors. While this strategy is well-suited to closed-set class-conditional generation, e.g., ImageNet synthesis (Deng et al., 2009), where labels are discrete and unambiguous, it is less appropriate for open-vocabulary text-to-image (T2I) generation, where negative concepts are compositional and fuzzy and their complement concept is not a well-defined alternative class. In a related direction, Na et al. (2025) introduced auxiliary classifier-based guidance, which uses a pre-trained classifier to assess the predicted clean image and feeds this signal back into denoising for explicit look-ahead control. However, its reliance on an additional classifier incurs non-trivial computational overhead, making it less suitable for real-time use.

To address these limitations, we propose an inference-time defense that adaptively scales negative guidance using a robust, stepwise concept presence signal. Our signal fuses evidence from the current denoising state with an explicit look-ahead proxy derived from the predicted clean sample, leveraging the intrinsic discriminative capability of diffusion models (Clark & Jaini, 2023; Li et al., 2023; Jaini et al., 2024; Wang et al., 2025). This enables stable suppression even when intermediate concept cues are weak, without requiring any auxiliary models. The main contributions are as follows:

- **Training-free concept presence estimation.** We derive an inference-time estimator of undesirable concept presence that aggregates signals from the current noisy sample and the model’s one-step prediction of the clean image.
- **Inference-time Concept Removal Guidance (CRG).** Building on this estimator, we introduce CRG, a plug-and-play guidance rule that adaptively calibrates the negative guidance scale at each denoising step to robustly suppress unwanted concepts while preserving prompt fidelity.
- **Robust safety-fidelity trade-off.** Extensive re-teaming experiments show that CRG consistently blocks unwanted concept generations and maintains benign visual quality, without requiring additional data, fine-tuning, or auxiliary classifiers.

Conflict of Interest Disclosure. The authors received Google Cloud credits from the Google.org GCP Program for accessing the Gemini model, which is used in this paper for evaluation.

2. Preliminaries and Related Work

2.1. Diffusion Models

Diffusion models are a class of generative models that learn to synthesize data through two stochastic processes: forward noising process and reverse denoising process. The forward process is a Markov chain characterized by a predefined Gaussian transition, where data x_0 is iteratively corrupted according to a fixed variance schedule:

$$q(x_t|x_{t-1}) := \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I),$$

for $t \in 1, \dots, T$ and $0 < \beta_t < 1$. The schedule β_t is typically chosen such that $q(x_T)$ approximately follows $\mathcal{N}(0, I)$. The reverse process is another Markov chain with a learned Gaussian transition, used to define the sampling distribution $p_\theta(x_0)$:

$$p_\theta(x_{t-1}|x_t) := \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_t),$$

where $\mu_\theta(x_t, t)$ is parameterized by a neural network. While these models can be trained via variational lower bound, [Ho et al. \(2020\)](#) propose a simplified objective. By reparameterizing the reverse process to predict the noise ϵ , the loss function becomes:

$$\mathcal{L}_{\text{Diffusion}}(\theta) = \mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_t^\theta(x_t)\|_2^2], \quad (1)$$

where $x_t := \sqrt{\alpha_t}x_0 + \sqrt{1 - \alpha_t}\epsilon$ and $\alpha_t := \prod_{i=1}^t (1 - \beta_i)$.

For conditional generation tasks, such as T2I, the network is also provided with a condition c (e.g., an embedding of a text prompt representing the target concept), yielding $\epsilon_t^\theta(x_t, c)$. Training a conditional diffusion model typically amounts to learning a noise predictor by minimizing

$$\mathcal{L}_{\text{Diffusion}}(\theta) = \mathbb{E}_{t, x_0 \sim p_0(x_0|c), \epsilon} [\|\epsilon - \epsilon_t^\theta(x_t, c)\|_2^2],$$

where x_0 denotes the image paired with the conditioning prompt c .

2.2. Negative Guidance

Negative guidance is commonly employed to steer the sampling trajectory away from negative concepts defined by the negative prompt $c-$ (e.g., nudity, violence, or self-harm).

$$\epsilon_{t, \lambda}^\theta(x_t, c, c-) \triangleq \epsilon_t^\theta(x_t, c) - \lambda \underbrace{(\epsilon_t^\theta(x_t, c-) - \epsilon_t^\theta(x_t, c))}_{\triangleq \Delta \epsilon_t}. \quad (2)$$

Here, $\Delta \epsilon_t$ represents the *negative guidance*, a vector capturing the direction pointing towards the negative concept relative to the unconditional baseline. By subtracting this vector scaled by λ , the model suppresses the manifestation of $c-$. In this paper, we refer to the use of a fixed negative guidance scale λ across all timesteps as *Negative Prompting*.

2.3. Dynamic Negative Guidance

Recently, [Koulischer et al. \(2025\)](#) introduced Dynamic Negative Guidance (DNG), a strategy that adaptively modulates the guidance magnitude based on the estimated posterior probability $p_t(c - |x_t, c)$. In the conditional generation setting aiming to preserve c while removing $c-$, DNG employs dynamic scaling of the guidance:

$$\begin{aligned} \epsilon_{\text{DNG}}(x_t, c, c-) \\ = \epsilon_t^\theta(x_t, c) - \lambda_0 \omega_{\text{DNG}} (\epsilon_t^\theta(x_t, c, c-) - \epsilon_t^\theta(x_t, c)), \end{aligned} \quad (3)$$

where

$$\omega_{\text{DNG}} = \frac{p_t(c - |x_t, c)}{1 - p_t(c - |x_t, c)}.$$

Here, the posterior term is approximated via a likelihood ratio between the conditional path distribution and its compositional counterpart:

$$p_t(c - |x_t, c) = \frac{p(c - |c) p_{t:T}(x_{t:T}|c, c-)}{p_{t:T}(x_{t:T}|c)}. \quad (4)$$

However, the posterior-odds form of ω_{DNG} makes the guidance weight extremely sensitive when p_t is near 1, amplifying the estimation error. This is why practical implementations clamp the posterior for numerical stability.

This issue is exacerbated in open-vocabulary T2I generation, where the concept $c-$ is a compositional text concept for which the event $\neg c-$ is not a single coherent alternative class. Any posterior $p_t(c - |x_t, c)$ therefore inherits an unavoidable normalization ambiguity. This makes the estimation difficult, often remaining in a moderate-confidence regime and failing to provide a stable, actionable signal for adaptive suppression. One may tune hyperparameters to amplify the effect, but this typically increases the risk of unstable odds near the boundary and collateral degradation of benign content. We provide a detailed analysis in [Appendix A.2](#).

2.4. Other Related Work

Fine-tuning methods erase concepts by modifying model parameters. CA ([Kumari et al., 2023](#)) overwrites undesired knowledge via distribution matching, aligning the target concept’s distribution with a generic anchor. For data-free scenarios, ESD ([Gandikota et al., 2023](#)) fine-tunes weights using negative guidance to steer generation away from the undesired target. Addressing scalability, MACE ([Lu et al., 2024](#)) combines closed-form cross-attention refinement with LoRA ([Hu et al., 2022](#)) to efficiently handle mass concept erasure. While robust, these methods typically incur high computational costs compared to inference-time interventions.

Weight editing methods, such as UCE ([Gandikota et al., 2024](#)) and RECE ([Gong et al., 2024](#)), directly edit the key and value projection matrices in cross-attention layers using closed-form solutions to avoid full fine-tuning. While UCE performs a simultaneous closed-form update by jointly optimizing for multiple concepts, RECE iteratively derives and erases target embeddings to address incomplete erasure. Despite their use of regularization to preserve model capabilities, [Amara et al. \(2025\)](#) reported that these methods can still lead to degradation in image quality.

Inference-time strategies provide a flexible alternative to retraining, allowing for trajectory steering without modifying the model. While standard negative prompting (Eq. 2)

serves as a baseline, recent advancements have focused on dynamic guidance mechanisms. TraSCE (Jain et al., 2024) employs localized gradient-based guidance, while STG (Na et al., 2025) utilizes a pre-trained classifier to measure concept presence at inference-time. However, both methods incur additional computational overhead. SAFREE (Yoon et al., 2025) further applies filtering in pixel space via an adaptive latent re-attention mechanism, but this comes with additional computational overhead.

3. Methodology

Moving beyond prior methods, we propose a concept presence estimator that operates robustly in open-vocabulary settings. While conventional methods for concept detection rely on auxiliary image classifiers (Na et al., 2025), recent research demonstrates that diffusion models inherently acquire rich semantic representations during training (Yang & Wang, 2023; Yu et al., 2025). Leveraging these intrinsic capabilities, we estimate concept presence directly through the generative dynamics of the pre-trained model, eliminating the need for external supervision.

3.1. Concept Presence Estimation

We define the *concept presence* (CP) as the expected log-likelihood of a negative concept $c-$ manifesting in images generated conditioned on the prompt c :

$$\begin{aligned} CP(c - |c) &\triangleq \mathbb{E}_{p_\theta(x_0|c)} [\log p(c - |x_0)] \\ &= \int p_\theta(x_0 | c) \log p(c - |x_0) dx_0 \\ &= \log p(c -) + \int p_\theta(x_0 | c) \log \frac{p(x_0|c-)}{p(x_0)} dx_0 \\ &= \log p(c -) + \int p_\theta(x_0 | c) \log \frac{p(x_0|c)}{p(x_0)} dx_0 \\ &\quad - \int p_\theta(x_0 | c) \log \frac{p(x_0|c)}{p(x_0|c-)} dx_0, \end{aligned} \quad (5)$$

where x_0 is an image sampled from the pre-trained diffusion model given the user prompt c . To circumvent the requirement for auxiliary classifiers, we reformulate $\log p(c - | x_0)$ via Bayes' rule. We omit the prior term $\log p(c -)$ as it is independent of the prompt c .

Then the concept presence can be decomposed into the difference of two KL divergences:

$$\begin{aligned} CP(c - |c) &\approx \underbrace{\mathbb{E}_{p_\theta(x_0|c)} \left[\log \frac{p_\theta(x_0|c)}{p_\theta(x_0)} \right]}_{D_{\text{KL}}(p_\theta(x_0|c) \| p_\theta(x_0))} \\ &\quad - \underbrace{\mathbb{E}_{p_\theta(x_0|c)} \left[\log \frac{p_\theta(x_0|c)}{p_\theta(x_0|c-)} \right]}_{D_{\text{KL}}(p_\theta(x_0|c) \| p_\theta(x_0|c-))}, \end{aligned} \quad (6)$$

where we have replaced the true likelihood with the generative likelihood of a diffusion model. Eq. (6) shows that the proposed concept presence score $CP(c - |c)$ can be interpreted as a difference between two distributional discrepancies. The first term measures how strongly the conditional model $p_\theta(x_0|c)$ deviates from the unconditional prior $p_\theta(x_0)$; in other words, it quantifies how much the prompt c “controls” or informs the generated sample. The second term measures how distinguishable the c conditioned distribution is from the negative concept distribution $p_\theta(x_0|c-)$. Therefore, $CP(c - |c)$ becomes large when the prompt c induces informative, structured generations that deviate from the unconditional distribution, while remaining close to the distribution associated with the unwanted concept $c-$. This aligns with our goal of identifying trajectories where the negative concept $c-$ is implicitly realized under the prompt c .

Following Kong et al. (2024) and Wang et al. (2025), we estimate the first conditional-unconditional divergence by expressing it as an expected discrepancy between conditional and unconditional noise predictions:

$$\begin{aligned} &\mathbb{E}_{p_\theta(x_0|c)} [\log p_\theta(x_0|c) - \log p_\theta(x_0)] \\ &= \mathbb{E}_{p_\theta(x_0|c), t, \epsilon \sim \mathcal{N}(0, I)} \left[\kappa_t \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t)\|^2 \right], \end{aligned} \quad (7)$$

where $\kappa_t = \frac{\beta_t T}{2\alpha_t(1-\bar{\alpha}_t)}$.

Applying the same approximation used for the first KL term, we express the second KL term as the expected difference between the conditional noise predictions:

$$\begin{aligned} &\mathbb{E}_{p_\theta(x_0|c)} [\log p_\theta(x_0|c) - \log p_\theta(x_0|c-)] \\ &= \mathbb{E}_{p_\theta(x_0|c), t, \epsilon \sim \mathcal{N}(0, I)} \left[\kappa_t \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t, c-)\|^2 \right]. \end{aligned} \quad (8)$$

Combining the two KL divergence terms, we derive the final expression for concept presence.

$$\begin{aligned} CP(c - |c) &= \mathbb{E}_{x_0, t, \epsilon} \kappa_t \left[\|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t)\|^2 \right. \\ &\quad \left. - \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t, c-)\|^2 \right] \\ &= \mathbb{E}_{x_0, t, \epsilon} \kappa_t \left[2\epsilon_t^\theta(x_t, c)^\top (\epsilon_t^\theta(x_t, c-) - \epsilon_t^\theta(x_t)) \right. \\ &\quad \left. + \left(\|\epsilon_t^\theta(x_t)\|^2 - \|\epsilon_t^\theta(x_t, c-)\|^2 \right) \right] \\ &= \mathbb{E}_{x_0, t, \epsilon} \kappa_t \left[2\epsilon_t^\theta(x_t, c)^\top \Delta_\epsilon + b_t \right]. \end{aligned} \quad (9)$$

The final form admits a simple geometric interpretation: the inner-product term $\epsilon_t^\theta(x_t, c)^\top \Delta_\epsilon$ quantifies the alignment between the prompt-conditioned noise prediction and the negative-guidance direction defined in Eq. (2). The remaining term $b_t := \|\epsilon_t^\theta(x_t)\|^2 - \|\epsilon_t^\theta(x_t, c-)\|^2$ is independent of c and thus acts as a prompt-agnostic offset.

3.2. Point-Wise Concept Presence Estimation

Moving beyond prompt-level concept presence $CP(c - |c)$, this work targets sample-level, point-wise concept presence $CP(c - |x_0)$. For practical purposes, point-wise concept presence is calculated based on the single sampling trajectory of x_0 given c :

$$CP(c - |x_0) \approx \frac{1}{T} \sum_{t=1}^T \kappa_t \left[2\epsilon_t^\theta(x_t, c)^\top \Delta\epsilon_t + b_t \right], \quad (10)$$

where x_t denotes the corresponding intermediate state at timestep t . Intuitively, this expression accumulates the alignment between the sampling trajectory and the negative-guidance direction $\Delta\epsilon_t$ across denoising steps.

Directly estimating the concept presence in the final generated image at timestep t necessitates unrolling the remaining denoising steps ($s < t$), a process that is computationally prohibitive. To circumvent this, we analyze the posterior distribution $p(x_0|x_t)$ rather than the remaining full trajectory $x_{0:t}$. While the true posterior is generally intractable, we can recover its mean using Tweedie’s formula (Efron, 2011):

$$\hat{x}_0^\theta(x_t, \cdot) = \mathbb{E}_{p_\theta(x_0|x_t, \cdot)}[x_0] = \frac{1}{\sqrt{\bar{\alpha}_t}} \left(x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_t^\theta(x_t, \cdot) \right). \quad (11)$$

Following standard practices in diffusion-based inverse problems (Ho et al., 2022; Song et al., 2023; Zhu et al., 2023; Peng et al., 2024), we approximate the posterior as a Gaussian. Under this assumption, the point-wise concept presence simplifies to:

$$CP_k(c - |\hat{x}_0^\theta(x_t, c)) \approx \left(\frac{1}{2r_t^2} \frac{1 - \bar{\alpha}_t}{\bar{\alpha}_t} \right) \left[2\epsilon_t^\theta(x_t, c)^\top \Delta\epsilon_t + b_t \right] + \frac{1}{T} \sum_{t'=t+1}^T \kappa_{t'} \left[2\epsilon_{t'}^\theta(x_{t'}, c)^\top \Delta\epsilon_{t'} + b_{t'} \right], \quad (12)$$

where r_t^2 denotes the posterior variance. Following Zhu et al. (2023), we set

$$r_t^2 = \frac{1}{k} \frac{1 - \bar{\alpha}_t}{\bar{\alpha}_t},$$

treating k as a hyperparameter. Intuitively, k controls the relative strength of the Tweedie-based look-ahead term in our concept presence estimate. We further analyze the sensitivity to this variance scaling in Appendix C.1.1.

3.3. Inference-Time Concept Removal Guidance

Building on Eq. (12), we propose Concept Removal Guidance (CRG), an inference-time technique that suppresses specific concepts below a threshold τ while preserving

prompt fidelity. We formulate this as a constrained optimization task, seeking to minimize deviation from the original guidance while satisfying the concept presence constraint:

$$\begin{aligned} \min_{\epsilon_t^*} \quad & \frac{1}{2} \|\epsilon_t^* - \epsilon_t^\theta(x_t, c)\|_2^2 \\ \text{s.t.} \quad & CP_k(c - |\hat{x}_0^*(x_t)) \leq \tau, \end{aligned} \quad (13)$$

where

$$\hat{x}_0^*(x_t) = \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_t^*).$$

Crucially, the constraint is *affine* with respect to ϵ_t^* , and thus corresponds to a half-space:

$$CP_k(c - |\hat{x}_0^*(x_t)) = \frac{k}{2} \left(2\epsilon_t^{*\top} \Delta\epsilon_t + b_t \right) + \frac{1}{T} \sum_{t'=t+1}^T \kappa_{t'} \left[2\epsilon_{t'}^\top \Delta\epsilon_{t'} + b_{t'} \right]. \quad (14)$$

Notably, because the actual sampling trajectory diverges from the conditional distribution c , we replace the conditional guidance term $\epsilon_{t'}^\theta(x_{t'}, c)$ in Eq. (12) with the realized guidance $\epsilon_{t'}$ for all preceding timesteps $t' > t$.

The optimization objective defined in Eq. (13) yields a straightforward analytical solution. Since the constraint is *affine* in the decision variable ϵ_t^* , it defines a *half-space* in the ϵ -space. Consequently, CRG admits a closed-form update that can be interpreted as the Euclidean *projection* of the original conditional score onto this half-space, with the correction magnitude determined by the *concept surplus* ($CP_k(\cdot) - \tau$).

$$\epsilon_t^* = \begin{cases} \epsilon_t^\theta(x_t, c), & \text{if } CP_k(c - |\hat{x}_0^\theta(x_t, c)) \leq \tau, \\ \epsilon_t^\theta(x_t, c) - \omega_{\text{CRG}} \Delta\epsilon_t, & \text{otherwise,} \end{cases} \quad (15)$$

where

$$\omega_{\text{CRG}} = \frac{CP_k(\cdot) - \tau}{k \|\Delta\epsilon_t\|^2}. \quad (16)$$

Empirically, we find that applying a scaling factor $\lambda_0 > 1$ to ω_{CRG} strengthens the robustness of the concept removal process:

$$\epsilon_t^* = \epsilon_t^\theta(x_t, c) - \lambda_0 \omega_{\text{CRG}} \Delta\epsilon_t. \quad (17)$$

Crucially, because the weight λ_0 is gated by the concept presence threshold, this amplification strictly intensifies suppression only when the target concept dominates.

In summary, Eq. (16) explicates our method’s mechanism: at each step, the guidance weight ω_{CRG} is linearly scaled by the estimated concept surplus. This fundamentally differs from the weighting scheme in DNG. As discussed earlier, DNG’s odds-based modulation is highly sensitive to posterior estimation error and typically requires numerical clamping in practice. In contrast, our CP score is an evidence measure that avoids normalization against an ill-posed complement

Table 1. **Quantitative evaluation of nudity concept removal on Stable Diffusion v1.4.** We report the Attack Success Rate (ASR) across adversarial prompt suites (*Ring-A-Bell*, *P4D*, *UnlearnAtk*, and *MMA-Diff*) alongside generation quality metrics on the benign *COCO* set (CLIP, FID). (**Bold**: Best, Underline: Second best)

Method	Ring-A-Bell 16	Ring-A-Bell 38	Ring-A-Bell 77	P4D	UnlearnAtk	MMA-Diff	COCO	
	ASR (↓)	ASR (↓)	ASR (↓)	ASR (↓)	ASR (↓)	ASR (↓)	CLIP (↑)	FID (↓)
SD v1.4	0.926	0.926	0.895	0.854	0.655	0.949	31.10	-
ESD	<u>0.021</u>	<u>0.053</u>	<u>0.042</u>	<u>0.152</u>	0.296	0.242	<u>31.39</u>	53.39
CA	<u>0.600</u>	<u>0.684</u>	<u>0.653</u>	<u>0.636</u>	0.415	0.568	31.60	56.19
MACE	0.137	0.179	0.147	0.159	0.155	0.127	29.15	67.49
UCE	0.158	0.116	0.158	0.298	<u>0.120</u>	0.364	30.14	63.30
RECE	0.084	0.095	<u>0.042</u>	0.377	0.155	0.695	30.87	58.40
SLD Medium	0.956	0.936	0.926	0.907	0.683	0.945	30.83	57.09
SLD Strong	0.926	0.926	0.853	0.914	0.585	0.913	30.11	58.49
SLD Max	0.789	0.800	0.685	0.788	0.507	0.885	29.31	62.59
SAFREE	0.505	0.526	0.368	0.529	0.275	0.611	31.11	55.79
TraSCE	0.168	0.221	0.158	0.291	0.134	0.490	30.67	57.11
STG	0.326	0.305	0.274	0.331	0.317	0.584	29.53	65.37
NP	0.211	0.211	0.168	0.338	0.183	0.453	30.49	<u>51.74</u>
DNG	0.221	0.200	0.200	0.225	0.183	<u>0.139</u>	30.79	55.97
CRG (ours)	0.011	0.021	0.021	0.026	0.059	0.164	30.85	36.05

and preserves the scale needed for control. By formulating concept removal as a minimum-perturbation projection subject to a CP constraint, we obtain a stable closed-form update whose magnitude is proportional to the surplus concept evidence and whose direction is aligned with the available negative guidance vector.

4. Experiments

4.1. Removing Nudity Concept

4.1.1. EXPERIMENTAL SETUP

We use Stable Diffusion v1.4 (Rombach et al., 2022) as the backbone model, utilizing a DDPM sampler with 50 inference steps. We instantiate the negative prompt $c-$ as a set of terms that characterize the nudity concept: “*Sexual Fantasy, Nudity, Pornography, Erotic Art, Nude, Naked, Sexual Acts*”. The procedure for selecting these terms is detailed in Appendix E.

Following prior work (Gandikota et al., 2023; Gong et al., 2024; Jain et al., 2024; Yoon et al., 2025), we evaluate our framework on established adversarial prompt benchmarks. Specifically, we utilize **Ring-A-Bell** (Tsai et al., 2024), which leverages genetic algorithms for prompt generation; **P4D** (Chin et al., 2024), which identifies adversarial prompts by analyzing discrepancies between constrained and unconstrained predictions; **UnlearnAtk** (Zhang et al., 2024), which synthesizes adversarial prompts utilizing the model’s own classification capabilities; and **MMA-Diff** (Yang et al., 2024), which employs multimodal adversarial attacks to bypass safety filters.

4.1.2. BASELINES AND EVALUATION METRICS

Our comparative analysis includes a comprehensive suite of state-of-the-art defenses. We compare **CRG** against weight-editing methods such as **UCE** (Gandikota et al., 2024) and **RECE** (Gong et al., 2024), as well as inference-time interventions including **SLD** (Schramowski et al., 2023), **SAFREE** (Yoon et al., 2025), **TraSCE** (Jain et al., 2024), **STG** (Na et al., 2025), **Negative Prompting** (NP), and **DNG** (Koulischer et al., 2025). For completeness, we also report results from fine-tuning approaches: **ESD** (Gandikota et al., 2023), **CA** (Kumari et al., 2023), and **MACE** (Lu et al., 2024).

For evaluation, we report the attack success rate (ASR) across the adversarial prompt benchmarks. ASR is computed as the proportion of generated images classified as unsafe by the **NudeNet**¹ classifier, where an image is deemed unsafe if its classification score exceeds 0.6. To assess the impact on general generation quality, we report the Fréchet Inception Distance (FID) (Heusel et al., 2017) relative to the images generated by the original SD v1.4 backbone, along with the CLIP score calculated on 1,000 samples from the COCO-30K (Lin et al., 2014) dataset.

4.1.3. EXPERIMENTAL RESULTS

As shown in Table 1, CRG delivers the strongest overall defense, achieving the lowest ASR on five of the six adversarial prompt benchmarks and remaining highly competitive on the remaining one. This consistent ranking across diverse adversarial prompt suites indicates that CRG is not tuned to a particular attack distribution, but instead provides robust

¹<https://github.com/notAI-tech/NudeNet>




Figure 3. **Visual fidelity on benign prompts.** While baselines like STG and DNG introduce artifacts or style shifts, CRG (far right) generates images nearly identical to the SD v1.4 baseline. This confirms that CRG minimizes interference, preserving the original composition and style.

suppression under varied prompt structures and concept-circumvention strategies. Compared to the strongest prior baselines, including training-based approaches that require additional data and optimization, CRG yields substantially larger safety gains, reducing ASR by 48%, 60%, and 50% on the Ring-A-Bell suite, 83% on P4D, and 51% on UnlearnAtk. Importantly, these improvements do not come from overly conservative generation. CRG simultaneously attains markedly lower FID than competing methods, suggesting that it removes the target concept while preserving the natural image manifold and avoiding excessive distortion of benign content. In other words, CRG suppresses the unwanted concept *selectively* rather than through blanket degradation of the output distribution. This safety–fidelity advantage is further supported by qualitative comparisons in Figure 2 and Figure 3, where CRG more reliably removes the target attribute while maintaining overall composition, subject identity, and contextual details.

4.1.4. SAFETY–FIDELITY PARETO ANALYSIS

We further analyze hyperparameter sensitivity and the resulting safety–fidelity trade-off across methods. As shown in Figure 4, CRG achieves a consistently better balance between CLIP score and ASR, suggesting that its safety gains do not come from simply sacrificing benign fidelity. Static negative prompting (NP) reduces ASR mainly by applying uniformly strong negative guidance, which often degrades semantic alignment and visual quality. DNG improves over NP by adapting its guidance, but remains less robust and does not match CRG across the explored hyperparameter range. Overall, CRG dominates the Pareto frontier, retaining higher CLIP score at the same ASR (and lower ASR at the same CLIP score), via stepwise concept presence estimation that modulates guidance in a targeted manner along the reverse trajectory. Further guidelines for setting hyperparameters are provided in Appendix C.1.




Figure 4. **Impact of hyperparameters on the CLIP–ASR trade-off.** We visualize the effect of varying (a) the threshold τ and (b) the scale λ_0 . Reported ASR is averaged across the P4D, UnlearnAtk, and MMA-Diff benchmarks.

4.1.5. EXTENSION TO ADVANCED MODELS

To demonstrate architectural generalizability, we evaluate our framework on two representative modern backbones, SDXL (Podell et al., 2024) and SD v3 (Esser et al., 2024), which differ substantially in both model scale and design choices. For a fair and controlled comparison, we adopt the DDPM sampler for SDXL and the stochastic sampler for SD v3, ensuring that performance differences stem from the guidance mechanism rather than choice of the sampler. As

Table 2. Cross-model generalization on SDXL and SD v3. We assess the scalability of our framework for nudity concept removal across advanced backbones.

Method	P4D	UnlearnAtk	MMA-Diff	COCO	
	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	CLIP ($\uparrow$)	FID ($\downarrow$)
SDXL	0.748	0.493	0.440	32.13	-
SAFREE	0.437	0.232	0.179	31.77	55.12
NP	0.205	0.141	0.064	31.71	44.81
DNG	0.146	0.099	0.029	31.52	50.34
CRG (ours)	0.060	0.077	0.016	32.07	44.77
SD v3	0.675	0.542	0.486	31.90	-
SAFREE	0.245	0.317	0.131	32.05	37.13
NP	0.139	0.204	0.090	31.72	35.83
DNG	0.132	0.183	0.024	31.78	34.06
CRG (ours)	0.060	0.085	0.024	31.78	35.46

shown in Table 2, CRG consistently achieves strong nudity concept removal across adversarial prompts while preserving benign generation quality, yielding image fidelity that remains competitive with (and in many cases close to) the corresponding backbone models. These results indicate that our inference-time guidance is not tied to a specific model family or training recipe, but instead transfers reliably to newer, higher-capacity architectures. Overall, the evaluation confirms that CRG provides a practical safety enhancement for state-of-the-art T2I systems, enabling safer image generation without any additional training, external datasets, or auxiliary classifiers.

4.1.6. EFFECT OF SAMPLER

Table 3. Quantitative comparison of nudity concept removal on Stable Diffusion v1.4 using the deterministic sampler (DPM-Solver++).

Method	P4D	UnlearnAtk	MMA-Diff	COCO	
	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	CLIP ($\uparrow$)	FID ($\downarrow$)
SD v1.4	0.960	0.697	0.972	31.40	-
SAFREE	0.377	0.225	0.591	31.10	43.31
TraSCE	0.139	0.042	0.554	30.62	51.24
NP	0.298	0.141	0.598	30.85	45.54
CRG (ours)	0.020	0.049	0.099	31.30	17.57

As presented in Table 3, CRG demonstrates superior efficacy under the deterministic DPM-Solver++ (Lu et al., 2025), achieving the best results across the majority of adversarial benchmarks and dominating in generation quality metrics (CLIP, FID). While TraSCE shows a marginally lower ASR on UnlearnAtk, it suffers from severe degradation in visual fidelity compared to our method. These findings validate that CRG performs consistently across deterministic samplers.

4.2. Removing Violence Concept

4.2.1. EXPERIMENTAL SETUP

To assess robustness in violence-related generation, we evaluated CRG against inference-time methods across three Ring-A-Bell configurations. Across all settings, we used SD v1.4 with 50-step DDPM sampling and a common negative

Table 4. Quantitative evaluation of violence concept removal on Stable Diffusion v1.4. We report the attack success rate (ASR) across three Ring-A-Bell configurations to assess the robustness of each mitigation strategy.

Method	Ring-A-Bell 5-77	Ring-A-Bell 5.5-38	Ring-A-Bell 5.5-77	COCO	
	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	CLIP ($\uparrow$)	FID ($\downarrow$)
SD v1.4	0.992	0.964	0.972	31.10	-
SAFREE	0.740	0.612	0.696	30.42	50.76
TraSCE	0.632	0.540	0.680	30.49	56.94
NP	0.760	0.628	0.744	30.41	50.76
DNG	0.392	0.352	0.400	30.68	56.08
CRG	0.296	0.252	0.280	30.84	37.63

prompt, “Bleeding, Blood, Gun, Weapon, Wounded”. The presence of violent content was assessed using the Q16 detector (Schramowski et al., 2022), and we report the attack success rate (ASR), where lower values indicate stronger suppression. Fidelity is evaluated using the same metrics as in the nudity removal task.

4.2.2. EXPERIMENTAL RESULTS

As shown in Table 4, SD v1.4 exhibits high ASR across all three Ring-A-Bell configurations, indicating that the adversarial prompts reliably induce violent content. Among the mitigation methods, CRG achieves the lowest ASR in all settings, reducing ASR to 0.296, 0.252, and 0.280, respectively. Compared with the strongest baseline, DNG, CRG further reduces ASR by a substantial margin while also preserving generation quality. In particular, CRG attains the highest CLIP score and the lowest FID on the COCO dataset, suggesting that it suppresses violent concepts more effectively without sacrificing benign image quality.

4.3. Removing Artist Style

4.3.1. EXPERIMENTAL SETUP

We evaluate whether CRG can selectively suppress a target artist style while preserving unrelated styles. Following prior work (Jain et al., 2024; Yoon et al., 2025) on artist style erasure, we consider two test cases: (1) removing the “Van Gogh” style while retaining the styles of “Pablo Picasso”, “Andy Warhol”, “Caravaggio”, and “Rembrandt”; and (2) removing the “Kelly McKernan” style while preserving the styles of “Tyler Edlin”, “Thomas Kinkade”, “Kilian Eng”, and “Ajin: Demi-Human”. For each case, we use a set of 100 prompts and generate samples with SD v1.4 using a DDPM sampler (following the same configuration as in our main experiments). For artist style removal, we use GPT-4o (OpenAI, 2024) for evaluation.

4.3.2. EVALUATION METRICS AND EXPERIMENTAL RESULTS

In Table 5, Acc_e denotes the classification accuracy of the erased style (lower is better), while Acc_u measures accuracy for the unrelated styles (higher is better). This metric pair

Table 5. Quantitative comparison of artist style removal on Stable Diffusion v1.4. We use GPT-4o to classify the generated images and report the classification accuracy on the target style (Acc_e , lower is better) and unrelated styles (Acc_u , higher is better).

Method	Remove “Van Gogh”		Remove “Kelly McKernan”	
	Acc_e (↓)	Acc_u (↑)	Acc_e (↓)	Acc_u (↑)
SD v1.4	0.80	0.90	0.80	0.64
SAFREE	0.40	<u>0.74</u>	0.20	0.64
TraSCE	0.40	0.70	<u>0.10</u>	0.64
NP	0.40	<u>0.74</u>	0.20	<u>0.65</u>
DNG	<u>0.15</u>	0.70	0.15	0.61
CRG (ours)	0.10	0.78	0.05	0.68

jointly captures the key challenge of artist style erasure: suppressing the targeted stylistic signature without inadvertently degrading the model’s ability to render other styles. As shown in Table 5 and Figure 1, our method achieves a lower Acc_e than competing approaches, indicating stronger removal of the target artist style. At the same time, it maintains a higher Acc_u , suggesting that the suppression is selective rather than indiscriminate, and that non-target stylistic capabilities remain largely intact. These results highlight that our guidance operates with high precision, avoiding the common failure mode of over-erasure. Additional evaluations using Gemini 2.5 Flash (Google Gemini Team, 2025) and Qwen3-VL (Bai et al., 2025), reported in Appendix F.1, further support that CRG achieves strong target style suppression while preserving non-target styles, though the precise ranking varies across VLM evaluators.

4.3.3. HUMAN EVALUATION

We additionally conducted a human evaluation on the artist style removal task with 30 participants, each evaluating 100 images per method, using the same protocol as the GPT-4o-based assessment. Participants were shown images generated by each method and asked to identify the artist style from five options: “Pablo Picasso”, “Van Gogh”, “Rembrandt”, “Andy Warhol”, and “Caravaggio”.

Table 6. Human evaluation results for artist-style removal on Stable Diffusion v1.4. Lower is better for Acc_e and higher is better for Acc_u .

Method	Remove “Van Gogh”	
	Acc_e (↓)	Acc_u (↑)
SD v1.4	0.837	0.803
SAFREE	0.418	0.814
TraSCE	0.434	0.815
NP	0.376	<u>0.816</u>
DNG	<u>0.304</u>	0.724
CRG (ours)	0.259	0.820

As shown in Table 6, our method yields the lowest erased-style accuracy (0.259) together with the highest non-target accuracy (0.820). This human evaluation is aligned with the

GPT-4o-based assessment, indicating that the same selectivity trend is also reflected in human judgments. Detailed settings for the human evaluation are provided in Appendix G.

4.4. Generalization and Robustness Analyses

To further assess the generalizability and robustness of CRG, we provide additional analyses in Appendix F along four complementary axes: robustness to negative prompt selection, applicability to identity removal, generalization across broader prompt distributions, and scalability to multi-concept removal.

5. Limitations

While CRG demonstrates strong empirical performance, several limitations remain. First, as an inference-time defense, CRG serves as a complementary safety layer and may remain vulnerable to subtle or adaptive adversarial prompts. Second, its effectiveness depends on user-specified negative prompts, which can be difficult to design in open-vocabulary settings, motivating future work on automated or adaptive negative-prompt construction. Finally, the concept presence score is derived under modeling approximations and reflects the presence of the user-specified concept. It should therefore be interpreted as a surrogate signal rather than a definitive indicator of content harmfulness.

6. Conclusion

We introduced *Concept Removal Guidance* (CRG), a training-free, plug-and-play inference-time defense for T2I diffusion models that adaptively calibrates negative guidance via a stepwise concept presence signal. Our estimator fuses evidence from the current denoising state with an explicit look-ahead proxy from the model’s predicted clean sample, enabling stable suppression even when intermediate cues are weak. It requires no additional data, fine-tuning, or auxiliary classifiers. Across red-teaming evaluations, CRG achieves strong concept removal while preserving fidelity on benign prompts, yielding a favorable safety–fidelity trade-off. Future work includes building on the multi-concept removal results in Appendix F.5 toward more robust handling of such scenarios, and extending concept presence estimation to other generative domains, such as video or 3D diffusion models.

Acknowledgements

This work was supported by Institute for Information & Communications Technology Planning & Evaluation (IITP) grant funded by the Korea government(MSIT) (No.RS-2024-00343989, Enhancing the Ethics of Data Characteristics and Generation AI Models for Social and Ethical Learning;

No.RS-2024-00397310, Development of an AI Simulator for Creating Transparent Compounds that Can Be Altered for Tactile Sensation; No.RS-2024-00457882, AI Research Hub Project; No.RS-2020-II200940, Foundations of Safe Reinforcement Learning and Its Applications to Natural Language Processing; No.RS-2019-II190075, Artificial Intelligence Graduate School Program (KAIST); No. RS-2022-II220311, Development of Goal-Oriented Reinforcement Learning Techniques for Contact-Rich Robotic Manipulation of Everyday Objects). We additionally acknowledge Google Cloud credits provided through the Google.org GCP Program, which were used for accessing the Gemini model in this work.

Impact Statement

This work presents Concept Removal Guidance (CRG) to advance the safety of text-to-image generation by effectively suppressing target concepts and protecting benign content at inference time. While CRG demonstrates robust performance across various benchmarks and architectures, it is not without limitations. As a defense relying on the model’s internal representations, it may still be susceptible to highly implicit adversarial prompts or adaptive attacks that bypass standard concept detection. Furthermore, while our method mitigates specific visual risks, it does not address the fundamental biases and stereotypes inherited from the training data of foundational models. We emphasize that safeguarding generative AI is an evolving challenge; thus, our method should be viewed as a complementary layer of defense rather than a standalone solution, necessitating ongoing research into fairness, data curation, and robust alignment strategies.

References

- Amara, I., Humayun, A. I., Kajic, I., Parekh, Z., Harris, N., Young, S., Nagpal, C., Kim, N., He, J., Vasconcelos, C. N., Ramachandran, D., Farnadi, G., Heller, K., Havaei, M., and Rostamzadeh, N. EraseBench: Understanding The Ripple Effects of Concept Erasure Techniques. *arXiv preprint arXiv:2501.09833*, 2025.
- Bai, S., Cai, Y., Chen, R., et al. Qwen3-VL Technical Report. *arXiv preprint arXiv:2511.21631*, 2025.
- Ban, Y., Wang, R., Zhou, T., Cheng, M., Gong, B., and Hsieh, C.-J. Understanding the Impact of Negative Prompts: When and How Do They Take Effect? In *European Conference on Computer Vision (ECCV)*, 2024.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D., Wu, J., Winter, C., Hesse, C., Chen, M., Sigler, E., Litwin, M., Gray, S., Chess, B., Clark, J., Berner, C., McCandlish, S., Radford, A., Sutskever, I., and Amodei, D. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*, 2020.
- Chin, Z.-Y., Jiang, C. M., Huang, C.-C., Chen, P.-Y., and Chiu, W.-C. Prompting4Debugging: Red-Teaming Text-to-Image Diffusion Models by Finding Problematic Prompts. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.
- Clark, K. and Jaini, P. Text-to-Image Diffusion Models are Zero Shot Classifiers. In *Advances in Neural Information Processing Systems 36 (NeurIPS)*, 2023.
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. ImageNet: A large-scale hierarchical image database. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2009.
- Dhariwal, P., Jun, H., Payne, C., Kim, J. W., Radford, A., and Sutskever, I. Jukebox: A Generative Model for Music. *arXiv preprint arXiv:2005.00341*, 2020.
- Efron, B. Tweedie’s Formula and Selection Bias. *Journal of the American Statistical Association*, 2011.
- Esser, P., Kulal, S., Blattmann, A., Entezari, R., Müller, J., Saini, H., Levi, Y., Lorenz, D., Sauer, A., Boesel, F., et al. Scaling Rectified Flow Transformers for High-resolution Image Synthesis. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.
- Gandikota, R., Materzyńska, J., Fiotto-Kaufman, J., and Bau, D. Erasing Concepts from Diffusion Models. In *Proceedings of the 2023 IEEE International Conference on Computer Vision (ICCV)*, 2023.
- Gandikota, R., Orgad, H., Belinkov, Y., Materzynska, J., and Bau, D. Unified Concept Editing in Diffusion Models. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2024.
- Gong, C., Chen, K., Wei, Z., Chen, J., and Jiang, Y.-G. Reliable and Efficient Concept Erasure of Text-to-Image Diffusion Models. In *European Conference on Computer Vision (ECCV)*, 2024.
- Google Gemini Team. Gemini 1.5: Unlocking Multimodal Understanding Across Millions of Tokens of Context. *arXiv preprint arXiv:2403.05530*, 2024.
- Google Gemini Team. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.

- Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B., and Hochreiter, S. GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium. In *Advances in Neural Information Processing Systems 30 (NIPS)*, 2017.
- Ho, J. and Salimans, T. Classifier-Free Diffusion Guidance. In *Neural Information Processing Systems (NeurIPS) 2021 Workshop on “Deep Generative Models and Downstream Applications”*, 2021.
- Ho, J., Jain, A., and Abbeel, P. Denoising Diffusion Probabilistic Models. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*, 2020.
- Ho, J., Salimans, T., Gritsenko, A. A., Chan, W., Norouzi, M., and Fleet, D. J. Video Diffusion Models. In *Advances in Neural Information Processing Systems 35 (NeurIPS)*, 2022.
- Hong, W., Ding, M., Zheng, W., Liu, X., and Tang, J. CogVideo: Large-scale Pretraining for Text-to-Video Generation via Transformers. In *International Conference on Learning Representations (ICLR)*, 2023.
- Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations (ICLR)*, 2022.
- Jain, A., Kobayashi, Y., Shibuya, T., Takida, Y., Memon, N., Togelius, J., and Mitsufuji, Y. TraSCE: Trajectory Steering for Concept Erasure. *arXiv preprint arXiv:2412.07658*, 2024.
- Jaini, P., Clark, K., and Geirhos, R. Intriguing Properties of Generative Classifiers. In *International Conference on Learning Representations (ICLR)*, 2024.
- Kong, X., Liu, O., Li, H., Yogatama, D., and Steeg, G. V. Interpretable Diffusion via Information Decomposition. In *International Conference on Learning Representations (ICLR)*, 2024.
- Kong, Z., Ping, W., Huang, J., Zhao, K., and Catanzaro, B. DiffWave: A Versatile Diffusion Model for Audio Synthesis. In *International Conference on Learning Representations (ICLR)*, 2021.
- Koulischer, F., Deleu, J., Raya, G., Demeester, T., and Ambrogioni, L. Dynamic Negative Guidance of Diffusion Models. In *International Conference on Learning Representations (ICLR)*, 2025.
- Kumari, N., Zhang, B., Wang, S.-Y., Shechtman, E., Zhang, R., and Zhu, J.-Y. Ablating Concepts in Text-to-Image Diffusion Models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- Li, A. C., Prabhudesai, M., Duggal, S., Brown, E., and Pathak, D. Your Diffusion Model is Secretly a Zero-Shot Classifier. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., and Zitnick, C. L. Microsoft COCO: Common objects in context. In *European Conference on Computer Vision (ECCV)*, 2014.
- Liu, N., Li, S., Du, Y., Torralba, A., and Tenenbaum, J. B. Compositional visual generation with composable diffusion models. In *European Conference on Computer Vision (ECCV)*, 2022.
- Liu, Y., Zhang, K., Li, Y., Yan, Z., Gao, C., Chen, R., Yuan, Z., Huang, Y., Sun, H., Gao, J., He, L., and Sun, L. Sora: A Review on Background, Technology, Limitations, and Opportunities of Large Vision Models. *arXiv preprint arXiv:2402.17177*, 2024.
- Lu, C., Zhou, Y., Bao, F., Chen, J., Li, C., and Zhu, J. DPM-Solver++: Fast Solver for Guided Sampling of Diffusion Probabilistic Models. *Machine Intelligence Research*, 22 (4):730–751, 2025. doi: 10.1007/s11633-025-1562-4.
- Lu, S., Wang, Z., Li, L., Liu, Y., and Kong, A. W.-K. MACE: Mass Concept Erasure in Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Na, B., Kang, M., Kwak, J., Park, M., Shin, J., Jun, S., Lee, G., Kim, J.-H., and Moon, I.-C. Training-Free Safe Text Embedding Guidance for Text-to-Image Diffusion Models. In *Advances in Neural Information Processing Systems 38 (NeurIPS)*, 2025.
- OpenAI. Hello GPT-4o. <https://openai.com/index/hello-gpt-4o/>, May 2024.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askeel, A., Welinder, P., Christiano, P. F., Leike, J., and Lowe, R. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems 35 (NeurIPS)*, 2022.
- Peng, X., Zheng, Z., Dai, W., Xiao, N., Li, C., Zou, J., and Xiong, H. Improving Diffusion Models for Inverse Problems Using Optimal Posterior Covariance. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.
- Podell, D., English, Z., Lacey, K., Blattmann, A., Dockhorn, T., Müller, J., Penna, J., and Rombach, R. SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis. In *International Conference on Learning Representations (ICLR)*, 2024.

- Ramesh, A., Pavlov, M., Goh, G., Gray, S., Voss, C., Radford, A., Chen, M., and Sutskever, I. Zero-Shot Text-to-Image Generation. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., and Ommer, B. High-Resolution Image Synthesis with Latent Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Saharia, C., Chan, W., Saxena, S., Li, L., Whang, J., Denton, E. L., Ghasemipour, K., Gontijo Lopes, R., Karagol Ayan, B., Salimans, T., Ho, J., Fleet, D. J., and Norouzi, M. Photorealistic Text-to-Image Diffusion Models with Deep Language Understanding. In *Advances in Neural Information Processing Systems 35 (NeurIPS)*, 2022.
- Schramowski, P., Tauchmann, C., and Kersting, K. Can Machines Help Us Answering Question 16 in Datasheets, and In Turn Reflecting on Inappropriate Content? In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, 2022.
- Schramowski, P., Brack, M., Deiseroth, B., and Kersting, K. Safe Latent Diffusion: Mitigating Inappropriate Degeneration in Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- Song, J., Vahdat, A., Mardani, M., and Kautz, J. Pseudoinverse-Guided Diffusion Models for Inverse Problems. In *International Conference on Learning Representations (ICLR)*, 2023.
- Tsai, Y.-L., Hsu, C.-Y., Xie, C., Lin, C.-H., Chen, J. Y., Li, B., Chen, P.-Y., Yu, C.-M., and Huang, C.-Y. Ring-A-Bell! How Reliable are Concept Removal Methods For Diffusion Models? In *International Conference on Learning Representations (ICLR)*, 2024.
- Wallace, B., Dang, M., Rafailov, R., Zhou, L., Lou, A., Purushwalkam, S., Ermon, S., Xiong, C., Joty, S., and Naik, N. Diffusion Model Alignment Using Direct Preference Optimization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Wang, C., Franzese, G., Finamore, A., Gallo, M., and Michiardi, P. Information Theoretic Text-to-Image Alignment. In *International Conference on Learning Representations (ICLR)*, 2025.
- Yang, X. and Wang, X. Diffusion Model as Representation Learner. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- Yang, Y., Gao, R., Wang, X., Ho, T.-Y., Xu, N., and Xu, Q. MMA-Diffusion: MultiModal Attack on Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- Yoon, J., Yu, S., Patil, V., Yao, H., and Bansal, M. SAFREE: Training-Free and Adaptive Guard for Safe Text-to-Image And Video Generation. In *International Conference on Learning Representations (ICLR)*, 2025.
- Yu, J., Xu, Y., Koh, J. Y., Luong, T., Baid, G., Wang, Z., Vasudevan, V., Ku, A., Yang, Y., Ayan, B. K., Hutchinson, B., Han, W., Parekh, Z., Li, X., Zhang, H., Baldridge, J., and Wu, Y. Scaling Autoregressive Models for Content-Rich Text-to-Image Generation. *Transactions on Machine Learning Research (TMLR)*, 2022.
- Yu, S., Kwak, S., Jang, H., Jeong, J., Huang, J., Shin, J., and Xie, S. Representation Alignment for Generation: Training Diffusion Transformers Is Easier Than You Think. In *International Conference on Learning Representations (ICLR)*, 2025.
- Zhang, Y., Jia, J., Chen, X., Chen, A., Zhang, Y., Liu, J., Ding, K., and Liu, S. To Generate or Not? Safety-driven Unlearned Diffusion Models Are Still Easy to Generate Unsafe Images... For Now. In *European Conference on Computer Vision (ECCV)*, 2024.
- Zhu, Y., Zhang, K., Liang, J., Cao, J., Wen, B., Timofte, R., and Van Gool, L. Denoising Diffusion Models for Plug-and-Play Image Restoration. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (NTIRE)*, 2023.

A. Dynamic Negative Guidance

A.1. Implementation Details of Dynamic Negative Guidance

For conditional generation, Dynamic Negative Guidance (DNG) (Koulischer et al., 2025) aims to sample from the following target distribution:

$$\pi(x_0 \mid c+, \neg c-) = \frac{p_0(x_0 \mid c+)(1 - p(c- \mid x_0))}{Z}, \quad (18)$$

where Z is the normalizing constant (partition function).

The corresponding forward-noised marginal at timestep t is

$$\begin{aligned} \pi_t(x_t \mid c+, \neg c-) &= \int q(x_t \mid x_0) \pi(x_0 \mid c+, \neg c-) dx_0 \\ &= \frac{1}{Z} \int q(x_t \mid x_0) p_0(x_0 \mid c+) (1 - p(c- \mid x_0)) dx_0 \\ &= \frac{1}{Z} \left[\int q(x_t \mid x_0) p_0(x_0 \mid c+) dx_0 - \int q(x_t \mid x_0) p_0(x_0 \mid c+) p(c- \mid x_0) dx_0 \right] \\ &= \frac{1}{Z} \left[p_t(x_t \mid c+) - p_t(x_t \mid c+) \int \frac{q(x_t \mid x_0) p_0(x_0 \mid c+)}{p_t(x_t \mid c+)} p(c- \mid x_0) dx_0 \right] \\ &= \frac{p_t(x_t \mid c+)}{Z} \left[1 - \int q(x_0 \mid x_t, c+) p(c- \mid x_0) dx_0 \right] = \frac{p_t(x_t \mid c+)}{Z} (1 - p(c- \mid x_t, c+)). \end{aligned} \quad (19)$$

Consequently, the score function of π_t decomposes as

$$\nabla_{x_t} \log \pi_t(x_t \mid c+, \neg c-) = \nabla_{x_t} \log p_t(x_t \mid c+) + \nabla_{x_t} \log (1 - p(c- \mid x_t, c+)). \quad (20)$$

Using the same manipulation as in the unconditional DNG derivation, Eq. (20) yields

$$\begin{aligned} \nabla_{x_t} \log \pi_t(x_t \mid c+, \neg c-) &= \nabla_{x_t} \log p_t(x_t \mid c+) - \frac{p_t(c- \mid x_t, c+)}{1 - p_t(c- \mid x_t, c+)} \left(\nabla_{x_t} \log p_t(x_t \mid c+, c-) - \nabla_{x_t} \log p_t(x_t \mid c+) \right) \\ &= \nabla_{x_t} \log p_t(x_t \mid c+) \\ &\quad - \frac{p(c- \mid c+) p_t(x_t \mid c+, c-)}{p_t(x_t \mid c+) - p(c- \mid c+) p_t(x_t \mid c+, c-)} \left(\nabla_{x_t} \log p_t(x_t \mid c+, c-) - \nabla_{x_t} \log p_t(x_t \mid c+) \right). \end{aligned} \quad (21)$$

In the conditional-generation setting, Koulischer et al. (2025) approximate the above by

$$\begin{aligned} \nabla_{x_t} \log \pi_t(x_t \mid c+, \neg c-) &\approx \nabla_{x_t} \log p_t(x_t \mid c+) \\ &\quad - \frac{p(c- \mid c+) p_t(x_t \mid c-)}{p_t(x_t \mid c+) - p(c- \mid c+) p_t(x_t \mid c-)} \left(\nabla_{x_t} \log p_t(x_t \mid c-) - \nabla_{x_t} \log p_t(x_t \mid c+) \right), \end{aligned} \quad (22)$$

which corresponds to approximating the joint conditional $p_t(x_t \mid c+, c-)$ by the marginal $p_t(x_t \mid c-)$. In practice, however, we observe that the resulting guidance brings limited gains in our experiments, which we attribute to the discrepancy introduced by replacing the joint conditional with its marginal.

A simple and effective alternative is to approximate the joint score by an additive composition of marginal scores (Liu et al., 2022):

$$\nabla_{x_t} \log p_t(x_t \mid c+, c-) \approx \nabla_{x_t} \log p_t(x_t \mid c+) + \nabla_{x_t} \log p_t(x_t \mid c-) - \nabla_{x_t} \log p_t(x_t). \quad (23)$$

This ‘‘score composition’’ form is widely adopted in diffusion-based compositional generation, where the goal is to enforce conjunctions such as $(c+ \wedge c-)$. Accordingly, we re-implement DNG by adopting this approximation for the joint score.

A.2. Hyperparameter Choice for DNG

Table 7. Sensitivity analysis of the temperature τ_{temp} and λ_0 . We evaluate the performance trade-off by varying the temperature τ and λ_0 for DNG.

DNG		Attack Success Rate (ASR) ↓			Benign Quality	
τ_{temp}	λ_0	P4D	UnlearnAtk	MMA-Diff	CLIP ↑	FID ↓
0.5	1.0	0.517	0.437	0.728	31.27	53.55
0.5	10.0	0.642	0.521	0.876	31.38	52.94
1.0	1.0	0.331	0.296	0.256	31.01	54.74
1.5	1.0	<u>0.225</u>	<u>0.183</u>	<u>0.139</u>	30.79	55.97
2.0	1.0	0.126	0.070	0.079	30.53	57.25

Unlike CRG, DNG fails to achieve meaningful concept removal even when increasing the guidance scale λ_0 . In contrast to fixed-class settings with well-defined discrete labels, the inherent semantic ambiguity in open-vocabulary T2I generation hinders the posterior $p_t(c- | x_t, c)$ from approaching 1. Consequently, the associated weight remains negligible ($\omega_t \approx 0$), rendering DNG ineffective unless the temperature parameter τ_{temp} is adjusted. However, as shown in Table 7, this adjustment gains removal performance at the cost of benign image quality, evidencing decreased CLIP scores and degraded FID. For the remaining parameters, we set $p_{\text{prior}} = 0.1$ and $\delta_{\text{offset}} = 0.004$.

A.3. Additional Stabilization Technique for DNG




Figure 5. Visual comparison of stabilization techniques for Dynamic Negative Guidance (DNG) when using high $\tau_{\text{temp}} = 1.5$. From left to right: (a) SD v1.4, output of backbone model; (b-d) DNG with clipping at thresholds of 100, 75, and 50, respectively. While clipping mitigates the instability of DNG, it struggles to match the quality of our proposed CRG method; (e) CRG (ours), showing high-fidelity results.

As illustrated in Figure 5, employing a high temperature value in DNG disrupts the generation of benign images, frequently resulting in saturation or blurring. While clipping the maximum guidance scale appears to be a viable strategy for stabilizing DNG, Figure 5 and Table 8 demonstrate that this approach compromises safety; specifically, it leads to a significant reduction in concept removal capability.

B. Details of Concept Presence Estimation

B.1. Transition From the Likelihood-Based Definition (Eq. (5)–(6)) to the Noise-Prediction Formulation (Eq. (7)–(9))

Assuming a perfectly trained diffusion model, the model-induced trajectory distributions coincide with the true forward trajectory distributions in both the conditional and unconditional cases:

$$p_{\theta}(x_{0:T} | c) = q(x_{0:T} | c), \quad p_{\theta}(x_{0:T}) = q(x_{0:T}).$$

Table 8. **Quantitative comparison between CRG (ours) and DNG with guidance scale clipping method.** Bold indicates the best performance.

Method	P4D	UnlearnAtk	MMA-Diff	COCO	
	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	CLIP ($\uparrow$)	FID ($\downarrow$)
CRG (ours)	0.026	0.059	0.164	30.85	36.05
<i>DNG</i> ($\tau_{temp} = 1.5$)					
+ Clip 100.0	0.225	0.183	0.139	30.79	55.97
+ Clip 75.0	0.192	0.155	0.264	31.04	54.60
+ Clip 50.0	0.265	0.275	0.391	31.24	53.70

Marginalizing over the intermediate variables $x_{1:T}$ then gives

$$p_\theta(x_0 | c) = q(x_0 | c), \quad p_\theta(x_0) = q(x_0),$$

and therefore the first KL term of Eq. (6) can be rewritten as

$$D_{\text{KL}}(p_\theta(x_0 | c) \| p_\theta(x_0)) = D_{\text{KL}}(q(x_0 | c) \| q(x_0)).$$

Moreover, this marginal KL can be lifted to the full diffusion trajectory by the chain rule. Since the forward noising process depends only on x_0 and is independent of c , we have

$$q(x_{0:T} | c) = q(x_0 | c)q(x_{1:T} | x_0), \quad q(x_{0:T}) = q(x_0)q(x_{1:T} | x_0).$$

Hence,

$$\begin{aligned} D_{\text{KL}}(q(x_{0:T} | c) \| q(x_{0:T})) &= \mathbb{E}_{q(x_{0:T}|c)} \left[\log \frac{q(x_0 | c)q(x_{1:T} | x_0)}{q(x_0)q(x_{1:T} | x_0)} \right] \\ &= D_{\text{KL}}(q(x_0 | c) \| q(x_0)), \end{aligned}$$

where the common forward-process term $q(x_{1:T} | x_0)$ cancels. Combining this with the equalities above yields

$$D_{\text{KL}}(p_\theta(x_0 | c) \| p_\theta(x_0)) = D_{\text{KL}}(p_\theta(x_{0:T} | c) \| p_\theta(x_{0:T})).$$

Using the standard reverse-process factorization, the trajectory-level KL decomposes into a sum of per-step KL terms:

$$D_{\text{KL}}(p_\theta(x_{0:T} | c) \| p_\theta(x_{0:T})) = \sum_{t=1}^T \mathbb{E}_{p_\theta(x_t|c)} [D_{\text{KL}}(p_\theta(x_{t-1} | x_t, c) \| p_\theta(x_{t-1} | x_t))]. \quad (24)$$

Since both reverse transitions are Gaussian with the same covariance $\beta_t \mathbf{I}$, each per-step KL reduces to a scaled squared ℓ_2 distance between the corresponding reverse means. Substituting the standard DDPM parameterization of the reverse mean in terms of ϵ_θ , we obtain

$$D_{\text{KL}}(p_\theta(x_0 | c) \| p_\theta(x_0)) = \mathbb{E}_{t,x,\epsilon} \left[\kappa_t \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t)\|^2 \right], \quad \kappa_t = \frac{\beta_t T}{2\alpha_t(1 - \bar{\alpha}_t)}.$$

The second KL term in Eq. (6) admits the same derivation, with the *unconditional* distribution replaced by the *negative concept-conditioned* distribution $p_\theta(x_0 | c-)$ (and correspondingly $p_\theta(x_{0:T} | c-)$ at the trajectory level).

$$D_{\text{KL}}(p_\theta(x_0 | c) \| p_\theta(x_0 | c-)) = \mathbb{E}_{t,x,\epsilon} \left[\kappa_t \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t, c-)\|^2 \right], \quad \kappa_t = \frac{\beta_t T}{2\alpha_t(1 - \bar{\alpha}_t)}.$$

B.2. Concept Presence Estimation

Building on the trajectory-level decomposition derived in Appendix B.1, we re-express the two distribution-level KL divergences in Eq. (6) as expectations over sample-level forward trajectories. Marginalizing each per-step expectation over

$x_0 \sim p_\theta(x_0|c)$ via $p_\theta(x_t|c) = \int p_\theta(x_0|c)q(x_t|x_0)dx_0$ (which holds under optimal training and the condition-independent forward process) gives:

$$\begin{aligned} D_{\text{KL}}(p_\theta(x_0|c) \parallel p_\theta(x_0)) &= \mathbb{E}_{p_\theta(x_0|c)} \sum_{t=1}^T \mathbb{E}_{q(x_t|x_0)} [D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t))] . \\ D_{\text{KL}}(p_\theta(x_0|c) \parallel p_\theta(x_0|c-)) &= \mathbb{E}_{p_\theta(x_0|c)} \sum_{t=1}^T \mathbb{E}_{q(x_t|x_0)} [D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t, c-))] . \end{aligned} \quad (25)$$

Subtracting the second identity from the first—and omitting the data-independent term $\log p(c-)$ as in Section 3.1—yields the trajectory-level decomposition of the concept presence:

$$\begin{aligned} CP(c-|c) &= \mathbb{E}_{p_\theta(x_0|c)} \left[\log \frac{p_\theta(x_0|c-)}{p_\theta(x_0)} \right] = D_{\text{KL}}(p_\theta(x_0|c) \parallel p_\theta(x_0)) - D_{\text{KL}}(p_\theta(x_0|c) \parallel p_\theta(x_0|c-)) \\ &= \mathbb{E}_{p_\theta(x_0|c)} \sum_{t=1}^T \mathbb{E}_{q(x_t|x_0)} [D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t)) - D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t, c-))] \end{aligned} \quad (26)$$

Accordingly, we define the sample-level, point-wise concept presence as

$$CP(c-|x_0) = \sum_{t=1}^T \mathbb{E}_{q(x_t|x_0)} [D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t)) - D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t, c-))] ,$$

which by construction satisfies

$$CP(c-|c) = \mathbb{E}_{p_\theta(x_0|c)} [CP(c-|x_0)] .$$

Thus, $CP(c-|x_0)$ matches $\log \frac{p_\theta(x_0|c-)}{p_\theta(x_0)}$ in expectation under $x_0 \sim p_\theta(x_0|c)$, and we use it as a sample-level surrogate signal for the presence of $c-$ in x_0 .

For practical evaluation, we approximate the inner expectation $\mathbb{E}_{q(x_t|x_0)}[\cdot]$ by a single realized sample from the reverse trajectory $x_{0:T}$, yielding the trajectory-level estimator:

$$CP(c-|x_0) \approx \sum_{t=1}^T [D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t)) - D_{\text{KL}}(p_\theta(x_{t-1}|x_t, c) \parallel p_\theta(x_{t-1}|x_t, c-))] .$$

We empirically validate the effectiveness of this trajectory-level estimator in Appendix D.

To evaluate concept presence at an intermediate timestep t during sampling, we further consider the marginal posterior distributions $p_\theta(x_0|x_t)$, $p_\theta(x_0|x_t, c)$, and $p_\theta(x_0|x_t, c-)$. By the chain rule of KL divergence applied to the joint (x_0, x_t) , the decomposition above admits an equivalent form parameterized by t :

$$\begin{aligned} CP(c-|c) &= \mathbb{E}_{p_\theta(x_0|c)} \left[\log \frac{p_\theta(x_0|c-)}{p_\theta(x_0)} \right] \\ &= \mathbb{E}_{p_\theta(x_t|c)} \sum_{t'=t+1}^T \mathbb{E}_{q(x_{t'}|x_t)} [D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'})) - D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'}, c-))] \\ &\quad + \mathbb{E}_{p_\theta(x_t|c)} [D_{\text{KL}}(p_\theta(x_0|x_t, c) \parallel p_\theta(x_0|x_t)) - D_{\text{KL}}(p_\theta(x_0|x_t, c) \parallel p_\theta(x_0|x_t, c-))] \end{aligned} \quad (27)$$

Accordingly, we define the intermediate concept presence at timestep t as

$$\begin{aligned} CP(c-|x_t) &= \sum_{t'=t+1}^T \mathbb{E}_{q(x_{t'}|x_t)} [D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'})) - D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'}, c-))] \\ &\quad + D_{\text{KL}}(p_\theta(x_0|x_t, c) \parallel p_\theta(x_0|x_t)) - D_{\text{KL}}(p_\theta(x_0|x_t, c) \parallel p_\theta(x_0|x_t, c-)) \end{aligned} \quad (28)$$

which satisfies

$$CP(c - |c) = \mathbb{E}_{p_\theta(x_t|c)} [CP(c - |x_t)].$$

We use $CP(c - |x_t)$ as the intermediate-step concept presence signal during sampling. However, the marginal posterior $p_\theta(x_0 | x_t, \cdot)$ is generally intractable, even under the optimal-training assumption, so the boundary KL terms cannot be evaluated in closed form. Following standard practice in diffusion-based inverse problems (Ho et al., 2022; Song et al., 2023; Zhu et al., 2023; Peng et al., 2024), we approximate each posterior as a Gaussian with shared variance:

$$\begin{aligned} \tilde{p}_\theta(x_0|x_t, c) &= \mathcal{N}(x_0; \hat{x}_0^\theta(x_t, c), r_t^2 I) \\ \tilde{p}_\theta(x_0|x_t) &= \mathcal{N}(x_0; \hat{x}_0^\theta(x_t), r_t^2 I) \\ \tilde{p}_\theta(x_0|x_t, c-) &= \mathcal{N}(x_0; \hat{x}_0^\theta(x_t, c-), r_t^2 I), \end{aligned} \quad (29)$$

where the posterior mean is given by Tweedie’s formula (Efron, 2011):

$$\hat{x}_0^\theta(x_t, \cdot) = \mathbb{E}_{p_\theta(x_0|x_t, \cdot)}[x_0] = \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_t^\theta(x_t, \cdot)),$$

and the shared variance is determined by the noise schedule up to a hyperparameter $k > 0$:

$$r_t^2 = \frac{1}{k} \frac{1 - \bar{\alpha}_t}{\bar{\alpha}_t}.$$

Substituting these Gaussian approximations, the boundary KL reduces to a scaled squared ℓ_2 distance in noise-prediction space, yielding the closed-form estimator:

$$\begin{aligned} CP_k(c - |\hat{x}_0^\theta(x_t, c)) &= \sum_{t'=t+1}^T \mathbb{E}_{q(x_{t'}|x_t)} [D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'})) - D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'}, c-))] \\ &\quad + D_{\text{KL}}(\tilde{p}_\theta(x_0 | x_t, c) \parallel \tilde{p}_\theta(x_0 | x_t)) - D_{\text{KL}}(\tilde{p}_\theta(x_0 | x_t, c) \parallel \tilde{p}_\theta(x_0 | x_t, c-)) \quad (30) \\ &= \sum_{t'=t+1}^T \mathbb{E}_{q(x_{t'}|x_t)} [D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'})) - D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'}, c-))] \\ &\quad + \frac{k}{2} \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t)\|^2 - \frac{k}{2} \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t, c-)\|^2 \end{aligned}$$

As before, we approximate the inner expectation $\mathbb{E}_{q(x_{t'}|x_t)}[\cdot]$ along a single realized sampling trajectory $x_{t:T}$, giving the trajectory-level estimate at intermediate timestep t :

$$\begin{aligned} CP_k(c - |\hat{x}_0^\theta(x_t, c)) &\approx \sum_{t'=t+1}^T [D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'})) - D_{\text{KL}}(p_\theta(x_{t'-1}|x_{t'}, c) \parallel p_\theta(x_{t'-1}|x_{t'}, c-))] \quad (31) \\ &\quad + \frac{k}{2} \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t)\|^2 - \frac{k}{2} \|\epsilon_t^\theta(x_t, c) - \epsilon_t^\theta(x_t, c-)\|^2 \end{aligned}$$

C. Concept Removal Guidance

C.1. Hyperparameter Choice for CRG

To evaluate hyperparameter sensitivity, we employ the SD v1.4 backbone with a DDPM sampler (50 steps). Classifier-Free Guidance (CFG) is applied for conditional generation. Our evaluation specifically focuses on the effect of hyperparameters within the context of sexual concept removal. The hyperparameter settings used for the nudity concept removal task are reported in Table 9.

Algorithm 1 Concept Removal Guidance

Input: Pre-Trained Diffusion Model, λ_0 , Threshold τ , User Prompt c , Negative Prompt $c-$.

 $x_T \sim \mathcal{N}(0, I)$

```

1: for  $t = T$  to 1 do
2:   Compute  $\epsilon_t^\theta(x_t)$ ,  $\epsilon_t^\theta(x_t, c-)$  and  $\epsilon_t^\theta(x_t, c)$ .
3:   Estimate the concept presence  $CP(c - |\hat{x}_0^\theta(x_t, c))$  of the sample at timestep  $t$  by Eq. (14)
4:   if  $CP(\cdot) > \tau$  then
5:     Compute  $w_{\text{CRG}}$  according to Eq. (16)
6:   else
7:      $w_{\text{CRG}} = 0$ 
8:   end if
9:   Plug  $w_{\text{CRG}}$  and  $\lambda_0$  in Eq. (17) to obtain optimal guidance
10:  Compute  $x_{t-1}$  based on optimal guidance
11: end for
12: Return  $x_0$ 
    
```

Table 9. Hyperparameter settings for nudity concept removal.

Model	k	τ	λ_0
SD v1.4	16	30.0	1.8
SDXL	16	75.0	1.7
SD v3	1024	0	2.0

 C.1.1. POSTERIOR VARIANCE SCALING FACTOR (k)

The hyperparameter k acts as an inverse scaling factor for the posterior variance r_t^2 , controlling the assumed uncertainty in the approximate posterior $p(x_0 | x_t)$. As shown in Eq. (12), k linearly scales the posterior-divergence term, and thus directly determines the importance of forward looking information.

To isolate the effect of k during ablations, we scale the threshold τ proportionally with k , keeping the ratio τ/k fixed. Table 10 shows that overly small k leads to insufficient suppression, whereas $k \geq 16$ yields consistently strong removal. These results highlight the importance of incorporating forward-looking information. Balancing concept removal performance and image fidelity, we use $k = 16$ for all Stable Diffusion v1.4 experiments. We further observe that $k = 16$ also performs well on SDXL, which we attribute to the similarity of their noise schedules. Conversely, SD v3 necessitates a significantly higher value of $k = 1024$ for optimal performance, a requirement likely attributable to its distinct noise schedule.

 Table 10. Sensitivity analysis of the posterior variance scaling factor k . We evaluate performance variation when scaling k while maintaining a constant ratio τ/k on the Sexual Concept Erasure task using SD v1.4.

$\lambda_0 = 1.8$		Attack Success Rate (ASR) ↓			Benign Quality	
k	τ	P4D	UnlearnAtk	MMA-Diff	CLIP ↑	FID ↓
2	3.75	0.126	0.169	0.390	30.95	35.56
4	7.50	0.079	0.099	0.283	<u>30.93</u>	<u>35.83</u>
16	30.0	0.026	0.059	<u>0.164</u>	30.85	36.05
32	60.0	<u>0.053</u>	<u>0.070</u>	0.173	30.84	36.12
64	120.0	0.026	0.077	0.149	30.81	36.22

 C.1.2. CONCEPT PRESENCE THRESHOLD (τ)

We further ablate the target threshold τ , which specifies the tolerated concept presence before suppression is activated. With $k = 16$ and $\lambda_0 = 1.8$ fixed, reducing τ generally lowers ASR, indicating stronger and more frequent suppression via the surplus term $(CP_k(\cdot) - \tau)$. However, this increased aggressiveness also propagates to benign generations: as τ decreases, the benign COCO FID degrades sharply (e.g., $31.16 \rightarrow 52.05$ from $\tau = 45$ to $\tau = 0$), showing that the same suppression




Figure 6. **Effect of the threshold τ in CRG.** Qualitative ablation over $\tau \in \{0, 15, 30, 45\}$. Increasing τ relaxes the suppression criterion, generally improving prompt fidelity on benign content (top row) while changing the level of intervention on a sensitive prompt (bottom row). The rightmost column shows the original SD v1.4 generations for reference.

mechanism can substantially perturb non-adversarial samples. Conversely, increasing τ improves benign fidelity and CLIP alignment by limiting how often suppression is triggered, but comes at the cost of weaker robustness, most clearly for MMA-Diff, whose ASR rises monotonically as τ increases.

Table 11. **Sensitivity analysis of the concept presence threshold τ .** We evaluate the performance trade-off by varying τ while keeping the guidance scale $\lambda_0 = 1.8$ and scaling factor $k = 16$ fixed.

$k = 16$		Attack Success Rate (ASR) ↓			Benign Quality	
τ	λ_0	P4D	UnlearnAtk	MMA-Diff	CLIP ↑	FID ↓
0.0	1.8	0.021	0.013	0.079	30.42	52.05
15.0	1.8	0.049	<u>0.026</u>	<u>0.134</u>	30.67	42.48
30.0	1.8	<u>0.026</u>	0.059	0.164	<u>30.85</u>	<u>36.05</u>
45.0	1.8	0.063	0.046	0.214	30.99	31.16

C.1.3. SCALING FACTOR (λ_0)

Table 12 presents the sensitivity analysis of the guidance scale λ_0 with a fixed threshold $\tau = 30.0$. Increasing λ_0 enhances safety by lowering ASR but leads to a trade-off in benign quality. Notably, compared to the τ ablation, the fluctuation in FID scores is less severe. This stability can be attributed to the gating mechanism of the threshold τ . Since τ is fixed at a sufficient level to distinguish malicious concepts, it effectively prevents the suppression term from activating on benign inputs. Consequently, λ_0 primarily scales the suppression strength on detected adversarial content without aggressively degrading the quality of benign generation.

Table 12. **Sensitivity analysis of the guidance scale λ_0 .** We evaluate the performance trade-off by varying the guidance scale λ_0 while keeping the threshold fixed at $\tau = 30.0$ (with $k = 16$). The results indicate tradeoff between ASR and Benign Quality.

$k = 16$		Attack Success Rate (ASR) ↓			Benign Quality	
λ_0	τ	P4D	UnlearnAtk	MMA-Diff	CLIP ↑	FID ↓
2.2	30.0	<u>0.028</u>	0.026	0.083	30.49	40.43
1.8	30.0	0.026	<u>0.059</u>	<u>0.164</u>	<u>30.85</u>	<u>36.05</u>
1.4	30.0	0.141	0.119	0.417	31.04	32.47

C.2. Computation Cost

Both DNG and CRG incur additional computational cost because they require one extra diffusion-model noise prediction relative to the backbone. Nevertheless, they are still more efficient than gradient-based approaches such as TraSCE.

Table 13. Inference time per sample (in seconds) measured on an L40S GPU.

Model	Base	SAFREE	TraSCE	NP	DNG	CRG (ours)
SD v1.4	3.20	4.67	11.59	3.20	4.82	4.80
SDXL	6.01	8.48	25.12	6.01	8.69	8.68
SD v3	5.52	7.32	23.23	5.52	7.46	7.47

D. Validity of Concept Presence

To validate the effectiveness of our metric, we compare the proposed Concept Presence (CP) estimation against a pre-trained NudeNet classifier. Specifically, we employ the sample-level estimator defined in Eq. (10):

$$CP(c- | x_0) \approx \frac{1}{T} \sum_{t=1}^T \kappa_t \left[2\epsilon_t^\theta(x_t, c)^\top \Delta\epsilon_t + b_t \right],$$

where x_0 represents images generated by SD v1.4 using adversarial prompts c from the P4D and MMA-Diff benchmarks. Here, $c-$ denotes the pre-defined negative prompt. The generation process utilizes the DPM-Solver++ (50 steps). In Figure 7, the x-axis represents the classification probability from NudeNet, while the y-axis corresponds to the calculated $CP(c- | x_0)$. The strong positive correlation indicates that our metric reliably tracks the strength of the target concept, supporting its use as a valid estimator of concept presence.




Figure 7. **Validation of Sample-Level Concept Presence Estimation.** The scatter plot compares our estimated concept presence (y -axis) against the scores obtained from the NudeNet classifier (x -axis). Red and orange dots indicate images generated from MMA-Diff and P4D prompts, respectively, while green dots represent images generated from benign COCO prompts.

E. Selection of Negative Prompts via Concept Presence

To construct an optimal negative prompt set for sexual concept removal, we utilized our Concept Presence (CP) metric (Eq. 10) to filter candidate terms.

1. Candidate Pool Generation We first aggregated a comprehensive pool of candidate terms from prior safety literature, specifically SAFREE (Yoon et al., 2025) and TraSCE (Jain et al., 2024):

- **SAFREE Candidates:** Sexual Acts, Content Meant to Arouse Sexual Excitement, Nudity, Pornography, Erotic Art, Lustful, Seductive, Orgasmic, Libido, Kinky, Sexual

Orientation, Sexual Attraction, Sexual Intercourse, Sexual Pleasure, Sexual Fantasy, Carnal Desires, Sexual Gratification

- **TraSCE Candidates:** naked, nude, bare, exposed, Stripped, topless, male genitalia, penis, buttocks

2. Metric-Based Filtering We evaluated the efficacy of each candidate word (c_{word}) by measuring its Concept Presence, $CP(c_{\text{word}}|x_0)$, against images generated from highly adversarial prompt sets (**P4D** and **MMA-Diff**). Our hypothesis posits that a higher Concept Presence score indicates a stronger alignment between the negative prompt and the latent features of the unsafe adversarial content, thereby predicting greater suppression utility.

3. Analysis and Selection Our analysis revealed distinct activation patterns across benchmarks:

- For **MMA-Diff**, terms such as Sexual Intercourse, Pornography, and Sexual Pleasure exhibited the highest concept presence.
- For **P4D**, high presence scores were observed for Erotic Art, Nudity, Sexual Fantasy, Lustful, Nude, Naked, Carnal Desires, and Pornography.
- Specifically, for male-related adversarial prompts in **P4D**, the most effective terms included Sexual Fantasy, Nudity, Sexual Gratification, Carnal Desires, Sexual Acts, and Penis.

Based on these high-activation clusters, we synthesized the final negative prompt set used for our sexual concept removal experiments: Sexual Fantasy, Nudity, Pornography, Erotic Art, Nude, Naked, Sexual Acts.

When the resulting CP-selected prompts are applied to baselines such as SAFREE and TraSCE, they yield better removal performance than the original prompts used in those methods, as shown in Table 14.

Table 14. Comparison on P4D, UnlearnAtk, MMA-Diff, and COCO metrics. Lower is better for P4D, UnlearnAtk, MMA-Diff, and FID, while higher is better for CLIP.

Method	P4D	UnlearnAtk	MMA-Diff	COCO	
	ASR (↓)	ASR (↓)	ASR (↓)	CLIP (↑)	FID (↓)
SAFREE	0.529	0.275	0.611	31.11	55.79
SAFREE + CRG Prompt	0.423	0.232	0.570	30.94	56.59
TraSCE	0.291	0.134	0.490	30.67	57.11
TraSCE + CRG Prompt	0.278	0.113	0.431	30.69	58.10

F. Generalization and Robustness Analyses

F.1. Additional Evaluations for Artist Style Removal Task

We further evaluate Artist Style Removal performance of each inference-time method using Gemini 2.5 Flash and Qwen3-VL-235B-A22B-Instruct. Following the same protocol as in our main experiments, each model classifies the artistic style of generated images, and we report the target-style accuracy (Acc_e , lower is better) and the unrelated-style accuracy (Acc_u , higher is better).

As shown in Table 15, CRG achieves the lowest or tied-lowest Acc_e in three out of four settings, remains competitive in the remaining case, while keeping Acc_u broadly comparable to the baselines.

F.2. Impact of LLM-generated Negative Prompt

To examine the robustness of CRG to negative prompt selection, we conducted an additional experiment using LLM-generated negative prompts. Specifically, we prompted Grok to generate eight keywords most closely related to nudity and sexuality, and used them as negative prompt. As shown in Table 16, CRG remains competitive under automatically

Table 15. **Quantitative comparison of artist style removal on Stable Diffusion v1.4.** We report the classification accuracy on the target style (Acc_e , lower is better) and unrelated styles (Acc_u , higher is better).

Method	Gemini 2.5 Flash				Qwen3-VL-235B-A22B-Instruct			
	Remove “Van Gogh”		Remove “Kelly McKernan”		Remove “Van Gogh”		Remove “Kelly McKernan”	
	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)
SD v1.4	0.95	0.95	0.75	0.90	0.85	0.99	0.75	0.73
SAFREE	0.25	0.83	0.30	<u>0.93</u>	0.85	0.83	0.40	<u>0.69</u>
TraSCE	0.35	0.78	<u>0.15</u>	0.90	0.25	<u>0.80</u>	0.35	<u>0.69</u>
NP	0.35	<u>0.85</u>	0.35	0.89	0.20	<u>0.80</u>	<u>0.25</u>	0.68
DNG	0.05	<u>0.69</u>	0.10	0.86	<u>0.05</u>	0.71	0.15	0.70
CRG (ours)	<u>0.15</u>	0.90	0.10	0.94	0.00	0.79	0.15	0.66

generated negative prompts, outperforming baselines on P4D and UnlearnAtk while maintaining favorable fidelity, although DNG achieves a lower ASR on MMA-Diff in this setting.

Table 16. **Results with handcrafted and Grok-generated negative prompt.** Lower is better for ASR and FID, while higher is better for CLIP.

Negative Prompt	Method	P4D	UnlearnAtk	MMA-Diff	COCO	
		ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	CLIP ($\uparrow$)	FID ($\downarrow$)
-	SD v1.4	0.854	0.655	0.949	31.10	-
Handcrafted	SAFREE	0.529	0.275	0.611	31.11	55.79
Handcrafted	TraSCE	0.291	<u>0.134</u>	0.490	30.67	57.11
Handcrafted	NP	0.338	<u>0.183</u>	0.453	30.49	<u>51.74</u>
Handcrafted	DNG	<u>0.225</u>	0.183	0.139	30.79	<u>55.97</u>
Handcrafted	CRG (OURS)	0.026	0.059	<u>0.164</u>	<u>30.85</u>	36.05
Grok	SAFREE	0.384	0.190	0.596	<u>30.81</u>	56.09
Grok	TraSCE	0.265	<u>0.120</u>	0.508	30.60	57.91
Grok	NP	0.331	0.176	0.423	30.35	<u>51.79</u>
Grok	DNG	<u>0.245</u>	0.218	0.229	30.76	56.36
Grok	CRG (OURS)	0.053	0.049	<u>0.271</u>	30.87	40.17

Grok 8 keywords. Nudity, Naked, Erotic, Sensual, Explicit, Sexual, Lewd, NSFW.

F.3. Identity Removal Task

To further validate the generality of CRG, we performed an identity removal task using SDXL with 50-step DDPM sampling. We report two classifier-based metrics using GPT-4o-mini and Qwen3-VL-235B-A22B-Instruct. The erasure accuracy (Acc_e) measures the detection of a targeted identity, and the unrelated accuracy (Acc_u) assesses the preservation of unrelated, non-targeted identities. Our evaluation covers five identities—Brad Pitt, Leonardo DiCaprio, Keanu Reeves, Robert Downey Jr., and Tom Cruise—with Brad Pitt and Leonardo DiCaprio designated as the removal targets. We constructed 20 actor-specific prompts for each identity, yielding a total of 100 prompts for evaluation.

F.4. Generalization to Broader Prompt Distributions

We further report experimental results on PartiPrompts (Yu et al., 2022), a comprehensive benchmark consisting of 1,632 prompts that span diverse categories and varying levels of complexity, from simple descriptions to compositionally challenging queries. The effect of nudity removal observed on the PartiPrompts benchmark is consistent with the trends seen on COCO-1k.

Table 17. **Quantitative comparison of identity removal in SDXL.** We report the classification accuracy on the target identity (Acc_e , lower is better) and unrelated identities (Acc_u , higher is better).

Method	GPT-4o-mini				Qwen3-VL-235B-A22B-Instruct			
	Remove “Brad Pitt”		Remove “Leonardo DiCaprio”		Remove “Brad Pitt”		Remove “Leonardo DiCaprio”	
	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)	Acc_e ($\downarrow$)	Acc_u ($\uparrow$)
SDXL	0.90	0.93	0.95	0.91	0.95	1.00	1.00	0.99
SAFREE	0.00	0.71	<u>0.15</u>	0.73	<u>0.40</u>	<u>0.93</u>	<u>0.35</u>	<u>0.95</u>
NP	0.60	<u>0.75</u>	0.50	0.75	0.70	0.96	0.75	0.96
DNG	<u>0.50</u>	0.74	0.55	0.73	0.55	<u>0.93</u>	0.70	0.94
CRG (ours)	0.00	0.76	0.10	<u>0.74</u>	0.25	0.96	0.25	0.88

Table 18. **Impact of Nudity Removal on PartiPrompts benchmark.** We report the generation quality metrics on the benign *PartiPrompts* set (CLIP, FID)

Method	PartiPrompts	
	CLIP ($\uparrow$)	FID ($\downarrow$)
SD v1.4	31.44	-
SAFREE	31.62	49.34
TraSCE	31.12	49.14
NP	30.81	<u>40.93</u>
DNG	31.08	49.40
CRG (ours)	<u>31.29</u>	25.08

F.5. Multi-Concept Removal

We further evaluate a multi-concept removal setting, where CRG-Multi suppresses both violence and nudity via sequential projection using $CP_k(c_1 - | \hat{x}_0)$ and $CP_k(c_2 - | \hat{x}_0)$. This extension is effective on both violence and nudity benchmarks, demonstrating that CRG is not limited to single-target removal, while introducing only a modest computational overhead and a moderate quality trade-off.

Table 19. **Quantitative comparison of single-concept and multi-concept removal.** Lower is better for ASR and FID, while higher is better for CLIP score.


Method	Violence			Nudity			Quality		Efficiency
	Ring-A-Bell 5-77	Ring-A-Bell 5.5-38	Ring-A-Bell 5.5-77	P4D	UnlearnAtk	MMA-Diff	COCO		-
	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	ASR ($\downarrow$)	CLIP ($\uparrow$)	FID ($\downarrow$)	Time (s/sample)
SD v1.4	0.992	0.964	0.972	0.854	0.655	0.949	31.10	-	3.20
CRG-Violence	<u>0.296</u>	<u>0.252</u>	<u>0.280</u>	0.834	0.592	0.938	<u>30.84</u>	<u>37.63</u>	4.80
CRG-Nudity	0.872	0.912	0.884	0.026	0.059	0.164	30.85	36.05	4.80
CRG-Multi	0.256	0.216	0.252	<u>0.046</u>	<u>0.070</u>	<u>0.205</u>	30.66	43.15	6.42

G. Details of Human Evaluation for Artist Style Removal Task

We describe the experimental setup of the human evaluation introduced in Section 4.3.3 in detail. We recruited 30 participants and designed the evaluation as a five-way classification task. As illustrated in Figure 8 (a), participants were first provided with a set of reference artworks for each of the five candidate artists—Pablo Picasso, Van Gogh, Rembrandt, Andy Warhol, and Caravaggio—to familiarize themselves with the corresponding styles, and were allowed to revisit these references throughout the task. Participants were then shown the same set of generated artworks in randomized order, with all method labels hidden, as shown in Figure 8 (b). For each item, they were asked to identify the most likely artist based solely on stylistic cues.

We also measured inter-annotator agreement to assess the reliability of the human evaluation. As shown in Table 20, the average pairwise agreement across methods was 68.2%, and the average Fleiss’ κ was 0.592, suggesting reasonable

agreement among participants for this five-way artist-style classification task.



(a) Reference images shown to participants. For each of the five candidate artists, a set of representative artworks is provided so that participants can become familiar with the artist’s style before answering.

(b) An example question. Given an artwork, the participant is asked to select the most likely artist among the five candidates (Pablo Picasso, Van Gogh, Rembrandt, Andy Warhol, Caravaggio).

Figure 8. Human evaluation interface. (a) Instruction page providing reference images for each artist. (b) An example question presented to the participants.

Table 20. **Inter-annotator agreement in the human evaluation.** We report overall pairwise agreement and Fleiss’ κ for each method.

Method	Pairwise Agreement (%)	Fleiss’ κ
SD v1.4	70.9	0.639
SAFREE	68.2	0.593
TraSCE	68.8	0.601
NP	68.1	0.589
DNG	64.0	0.533
CRG (ours)	69.0	0.598
Average	68.2	0.592

H. Qualitative Examples of Removing Nudity Concept




Figure 9. Qualitative comparison of nudity removal on Stable Diffusion v1.4. We visualize generation results under adversarial prompts targeting explicit concepts (e.g., sexual fantasy, nudity, pornography). CRG effectively suppresses unsafe content, whereas baselines frequently fail to neutralize the target concept.




Figure 10. **Qualitative comparison of nudity removal on Stable Diffusion XL.** We visualize generation results under adversarial prompts targeting explicit concepts (e.g., sexual fantasy, nudity, pornography). CRG effectively suppresses unsafe content, whereas baselines frequently fail to neutralize the target concept.




Figure 11. **Qualitative comparison of nudity removal on Stable Diffusion v3.** We visualize generation results under adversarial prompts targeting explicit concepts (e.g., sexual fantasy, nudity, pornography). CRG effectively suppresses unsafe content, whereas baselines frequently fail to neutralize the target concept.

I. Qualitative Examples of Removing Violence Concept




Figure 12. Qualitative comparison of violence removal on Stable Diffusion v1.4. We visualize generation results under adversarial prompts targeting explicit concepts (e.g., blood, weapon, wounded). CRG effectively suppresses unsafe content, whereas baselines frequently fail to neutralize the target concept.

J. Qualitative Examples of Removing Artist Styles




Figure 13. **Qualitative comparison of artist style removal on Stable Diffusion v1.4.** We visualize generation results of CRG and other baselines on the Artist Style removal task. The target style for removal is “Van Gogh” (Red); for these cases, the generated images should not exhibit the artistic style. In contrast, the styles of “Andy Warhol”, “Caravaggio”, “Picasso”, and “Rembrandt” are intended to be preserved (Blue). Each row presents images generated using prompts corresponding to the artist labeled on the left.