
Lookahead-GCG: Improving Universal Multi-Model Optimization-Based Jailbreaking Attacks via Stochastic Nesterov Optimization

Rong Feng^{1 2 3} Haohan Zhao^{3 4} Shiqin Tang³ Geng Liu^{2 3} Song Lai^{2 3} Meng Wang³ Shuxin Zhuang^{2 3}
Yuqi Zha² Changyi Ma¹ Runsheng Yu¹

Abstract

Universal transferable jailbreaking attacks enable systematic red-teaming of black-box large language models by optimizing a single adversarial suffix on open-source surrogates that generalizes across diverse harmful behaviors and target models simultaneously. A natural approach to improve transferability is multi-model training—optimizing against multiple source models simultaneously. Yet this approach has been largely abandoned, as it yields only marginal gains with standard optimizers. We argue the root cause is poor generalization: standard gradient descent lacks stability when aggregating gradients from diverse models. Since GCG and its variants (Zou et al., 2023; Jia et al., 2024; Yang et al., 2025) implicitly perform SGD in discrete token space, they inherit this instability in multi-model settings. We address this with **Lookahead-GCG**, which combines: (1) Stochastic Nesterov Accelerated Gradient (SNAG), whose lookahead mechanism reduces sensitivity to individual gradient updates, providing stability for multi-model optimization; (2) embedding-space momentum accumulation, which enables SNAG in discrete token optimization; and (3) maximally distant initialization, which exploits SNAG’s improved generalization by starting from a universally beneficial point. Experiments show our method achieves 50.37% ASR on open-source and 34.03% on closed-source LLMs, outperforming GCG and

TransferAttack with +11.78% gains from multi-model optimization.

Warning: This paper contains potentially harmful text generated by LLMs for research purposes.

1. Introduction

Jailbreaking attacks exploit vulnerabilities in safety-aligned Large Language Models (LLMs) to elicit harmful outputs. A critical property of such attacks is *transferability*: adversarial prompts optimized on open-source models can attack black-box systems without direct access (Zou et al., 2023). This enables systematic red-teaming of proprietary models like GPT-4 and Claude.

Among these transferable attacks, *universal* adversarial suffixes represent a particularly critical threat: a single optimized suffix, once computed, can be appended to *any* harmful query at zero marginal cost — decoupling the attacker from the end-user and enabling mass exploitation from a single computation. Unlike individual attacks that require per-query optimization, universal suffixes expose a systemic vulnerability that scales trivially once discovered.

Among jailbreaking methods, gradient-based optimization (Zou et al., 2023) offers a principled approach by directly searching for universal adversarial suffixes that maximize harmful output probability. However, the original Greedy Coordinate Gradient (GCG) method suffers from limited transferability (Chao et al., 2024). Recent work addresses this through various angles: improved candidate selection (Jia et al., 2024), scenario induction templates (Liu et al., 2025), and relaxed optimization constraints (Yang et al., 2025).

A natural path toward robust transferability is *multi-model training*, which optimizes adversarial suffixes against multiple source models simultaneously (Zou et al., 2023). However, subsequent works have largely abandoned this direction, as naive multi-model training yields only marginal gains (+0.17% in DeepSeek-V3.2 (Liu et al., 2025) as shown in Table 1). We argue that this underperformance stems not from the multi-model paradigm itself, but from the *instabil-*

¹School of Artificial Intelligence, Jilin University, China
²College of Computing, City University of Hong Kong, Hong Kong, China
³Centre for Artificial Intelligence and Robotics, Hong Kong Institute of Science & Innovation, Chinese Academy of Sciences, Hong Kong, China
⁴Department of Systems Engineering, City University of Hong Kong, Hong Kong, China. Correspondence to: Changyi Ma <changyima@jlu.edu.cn>, Runsheng Yu <runshengyu@gmail.com>.

ity of standard optimizers when aggregating heterogeneous gradient signals. Conventional gradient descent minimizes aggregate training loss yet fail to transfer to unseen target models.

Our key insight is that the failure of multi-model training stems not from the approach itself, but from the optimizer. Standard gradient descent lacks stability when aggregating gradients from diverse models, converging to solutions that minimize training loss but fail to transfer. We address this with **Lookahead-GCG**, which combines three techniques:

Effective multi-model jailbreaking requires both diverse source models *and* an optimizer with stronger stability guarantees. To this end, we propose **Lookahead-GCG**, a transferable attack framework that integrates three components:

(1) **Stochastic Nesterov Accelerated Gradient (SNAG)** (Yan et al., 2018), whose *lookahead mechanism* evaluates gradients at an anticipated future position, reducing sensitivity to model-specific gradient noise and promoting more transferable minima.

(2) **Embedding-space momentum accumulation**, which reformulates discrete token optimization in continuous embedding space. This enables *true directional momentum*—impossible in token space—and allows seamless integration of SNAG with periodic projection back to the vocabulary.

(3) **Maximally distant initialization**, which selects an initial suffix that maximizes average loss across all training models. This *universal starting point* amplifies the generalization advantage of SNAG by avoiding model-specific basins early in optimization.

Experiments show our method achieves 50.37% Attack Success Rate on open-source LLMs and 34.03% on closed-source models (including 58.83% on Gemini-1.5-Pro), significantly outperforming baselines, with multi-model optimization providing +11.78% gains where baselines only show minimal improvement.

2. Preliminaries

2.1. Universal Adversarial Suffix Attack

Suffix Attack. Let $\mathbf{x}_{1:n}$ denote a sequence of n tokens representing a harmful user query (e.g., “Tell me how to build a bomb”). An adversarial suffix attack seeks to find a suffix $\mathbf{s}_{1:l} \in \mathcal{V}^l$ of length l such that, when appended to the query, the target LLM generates an affirmative response $\mathbf{y}_{1:h}^*$ (e.g., “Sure, here is how to build a bomb”).

Given a language model p_θ , the optimization objective is:

$$\min_{\mathbf{s}_{1:l} \in \mathcal{V}^l} \mathcal{L}(\mathbf{s}_{1:l}) = -\log p_\theta(\mathbf{y}^* | \mathbf{x}_{1:n} \oplus \mathbf{s}_{1:l}) \quad (1)$$

where $\mathcal{V}$ is the vocabulary and $\oplus$ denotes concatenation.

Individual vs. Universal Attacks. An *individual* attack optimizes a dedicated suffix $\mathbf{s}^{(p)}$ for each query $\mathbf{x}^{(p)}$:

$$\min_{\mathbf{s}_{1:l}^{(p)} \in \mathcal{V}^l} \mathcal{L}(\mathbf{s}_{1:l}^{(p)}, \mathbf{x}^{(p)}), \quad \forall p \in \{1, \dots, P\} \quad (2)$$

A *universal* attack instead optimizes a single shared suffix $\mathbf{s}$ across all queries:

$$\min_{\mathbf{s}_{1:l} \in \mathcal{V}^l} \frac{1}{P} \sum_{p=1}^P \mathcal{L}(\mathbf{s}_{1:l}, \mathbf{x}^{(p)}) \quad (3)$$

Unlike individual attacks that require P separate optimizations, a universal suffix is computed once and transferable to any unseen query at zero marginal cost — exposing a systemic vulnerability that scales trivially with the number of harmful behaviors.

2.2. Optimization-Based Jailbreaking

The discrete nature of the token space makes direct gradient-based optimization intractable. GCG (Zou et al., 2023) addresses this through iterative token replacement:

1. **Gradient Computation:** Compute token-level gradients via one-hot relaxation. Let $\mathbf{E} \in \mathbb{R}^{l \times |\mathcal{V}|}$ be the one-hot representation of suffix $\mathbf{s}_{1:l}$: $\mathbf{G} = \nabla_{\mathbf{E}} \mathcal{L} = \nabla_{\mathbf{E}} (-\log p_\theta(\mathbf{y}^* | \mathbf{x}_{1:n} \oplus \mathbf{E}))$.
2. **Candidate Selection:** For each position $i \in \{1, \dots, l\}$, select top- n tokens with largest negative gradient: $\mathcal{C}_i = \text{top-}n(-\mathbf{G}_i)$.
3. **Candidate Evaluation:** Sample B candidates by randomly selecting one position i and one token from $\mathcal{C}_i$, evaluate loss, and keep the best replacement.

Several works have improved upon the original GCG algorithm from different perspectives:

I-GCG (Jia et al., 2024) improves GCG through diverse target templates and automatic multi-coordinate updating strategy, accelerating convergence and enhancing attack success rate. However, it focuses on *individual* suffix attack rather than *universal* one.

MAC (Zhang & Wei, 2025) introduces momentum terms to the token selection process, improving optimization efficiency by accumulating gradient information across iterations.

TransferAttack (Yang et al., 2025) identify two superfluous constraints in the standard GCG objective that limit transferability:

Response Pattern Constraint. The target output $\mathbf{y}^* = \text{"Sure, here's how to..."}$ often mismatches the actual jailbroken response (e.g., `"To make a bomb..."`). This forces the model toward an unnatural distribution.

Token Tail Constraint. Computing loss over all k target tokens penalizes minor formatting variations (e.g., `"\n Step 1:"` vs `"\n\n **Step 1:**"`), which are irrelevant to attack success.

SlotGCG (Jeong et al., 2026) exploits positional vulnerabilities in LLMs for jailbreak attacks. Unlike GCG, which restricts adversarial tokens to a fixed suffix position, SlotGCG introduces a Vulnerable Slot Score (VSS) to quantify positional vulnerability and identifies the most vulnerable slots for targeted token insertion.

LARGO (Li et al., 2026) reasserts the power of gradient-based optimization for jailbreaking LLMs by operating within the continuous latent space, bypassing the limitations of discrete language optimization.

SM-GCG (Gu et al., 2025) enhances the efficiency of white-box jailbreak attacks by introducing spatial momentum to stabilize the optimization landscape and overcome the limitations of the original GCG method.

MAGIC (Li et al., 2025) optimizes the jail-breaking pipeline by identifying and mitigating the "Indirect Effect" bottleneck within the original GCG framework, accelerating the generation of adversarial suffixes.

While these methods improve GCG from various aspects—target templates, token selection, and loss computation—none of them address the *optimizer* itself. Moreover, by abandoning multi-model training, these approaches rely solely on implicit transferability from a single source model, potentially limiting the universality of the learned adversarial suffixes. Our work revisits multi-model optimization and shows that with proper momentum-based gradient accumulation, multi-model training can be made stable and effective, yielding more universal suffixes that generalize better across models and templates.

2.3. Multi-Model Multi-Prompt Transferable Attack

To generate adversarial suffixes that transfer across different models, Zou et al. (Zou et al., 2023) proposed optimizing over M surrogate models and P harmful prompts simultaneously:

$$\min_{\mathbf{e} \in \mathbb{R}^{L \times d}} \mathcal{L}(\mathbf{e}) = \frac{1}{MP} \sum_{m=1}^M \sum_{p=1}^P \mathcal{L}_{m,p}(\mathbf{e})$$

where $\mathbf{e} \in \mathbb{R}^{L \times d}$ is the adversarial suffix embedding. A *data point* is a (model, prompt) pair, giving $N = MP$

data points. At each iteration t , we sample $(m_t, p_t) \sim \text{Unif}([M] \times [P])$ and compute $\hat{g}_t = \nabla \mathcal{L}_{m_t, p_t}(\mathbf{e}_t)$.

However, subsequent works (Jia et al., 2024; Zhang & Wei, 2025) have largely abandoned MMTA in favor of single-model multi-prompt optimization. We hypothesize two reasons: (1) multi-model training requires substantially more computational resources (loading multiple 7B+ models simultaneously), and (2) naive gradient aggregation across heterogeneous models introduces optimization instability without proportional performance gains. Our experiments (Table 1) confirm that with standard optimizers, multi-model training yields only marginal improvements over single-model (+0.17% in DeepSeek-V3.2 (Liu et al., 2025) as shown in Table 1). Similar results are also found in (Zou et al., 2023).

2.4. Template Manipulation in Jailbreak Attacks

ChatBug (Jiang et al., 2025) demonstrates that applying *malformed* templates—with corrupted or missing control tokens—exploits LLM vulnerabilities to format mismatches. In contrast, we select among *valid* templates from different model families to maximize transfer effectiveness.

3. Methods

3.1. Why Does Multi-Model Training Fail?

As discussed in Section 2, multi-model training (MMTA) has been largely abandoned despite its theoretical appeal for finding universal adversarial suffixes. Our experiments confirm this pessimism: with standard optimizers, multi-model training yields only marginal improvements over single-model approaches (+0.17% in DeepSeek-V3.2 (Liu et al., 2025) as shown in Table 1).

We hypothesize that the core issue is *generalization*. Different models produce conflicting gradient directions, creating a complex loss landscape. Standard gradient descent can minimize the training loss (achieving similar final loss across methods), but converges to solutions that overfit to the specific training model ensemble rather than capturing universal adversarial directions.

Note that GCG and TransferAttack implicitly perform gradient descent in discrete space: at each step, they select tokens along the negative gradient direction.

3.2. Nesterov Momentum Improves Generalization.

Stochastic Nesterov Accelerated Gradient (SNAG) (Nesterov, 1983; Sutskever et al., 2013) evaluates gradients at an

extrapolated position rather than the current iterate:

$$\tilde{\mathbf{e}}_t = \mathbf{e}_t + \beta \mathbf{v}_t \quad (\text{lookahead position}) \quad (4)$$

$$\mathbf{v}_{t+1} = \beta \mathbf{v}_t - \alpha \nabla_{\tilde{\mathbf{e}}_t} \mathcal{L}(\tilde{\mathbf{e}}_t; \xi_t) \quad (5)$$

$$\mathbf{e}_{t+1} = \mathbf{e}_t + \mathbf{v}_{t+1} \quad (6)$$

where $\xi_t = (m_t, p_t)$ denotes the randomly sampled model-prompt pair, α is the step size, and $\beta \in (0, 1)$ is the momentum coefficient.

SNAG’s advantage lies in *improved generalization*: the cumulative contribution of each gradient to the final solution is smaller than in standard gradient descent, with recent gradients receiving particularly reduced weights (Yan et al., 2018). For multi-model optimization where different models produce different gradients, this reduced sensitivity prevents any single model’s gradient from dominating the update, leading to more stable optimization and better transfer to unseen models.

Our experiments validate this: Lookahead-GCG achieves similar training loss as baselines but significantly higher transfer ASR (Table 1), enabling the anticipatory mechanism to guide token selection.

Algorithm 1 Lookahead-GCG

Input: Harmful questions $\{x^{(j)}\}_{j=1}^m$, target outputs $\{a^{(j)}\}_{j=1}^m$, models $\{M_i\}_{i=1}^M$, initial embedding $\mathbf{e}_0 \in \mathbb{R}^{L \times d}$, iterations T , momentum β , step size α , loss token number s .

Output: Optimized suffix tokens $\mathbf{x}^*$

```

1 Initialize velocity  $\mathbf{v}_0 \leftarrow \mathbf{0}$ ,  $m_c \leftarrow 1$ 
2 for  $t = 0, 1, \dots, T - 1$  do
3   Step 1: SNAG Lookahead  $\tilde{\mathbf{e}}_t \leftarrow \mathbf{e}_t + \beta \mathbf{v}_t$ 
4   Step 2: Stochastic Gradient on First  $s$  Tokens
       Sample  $m_t \sim \text{Uniform}([M])$ ,  $p_t \sim \text{Uniform}([m_c])$ 
        $\mathbf{h}_t \leftarrow [\tilde{\mathbf{e}}_t; \text{Embed}(x^{(p_t)}); \text{Embed}(a^{(p_t)})]$ ,  $\hat{\mathbf{g}}_t \leftarrow \nabla_{\tilde{\mathbf{e}}_t} \mathcal{L}_{m_t, p_t}^{(s)}(\mathbf{h}_t)$ .
5   Step 3: Momentum Update  $\mathbf{v}_{t+1} \leftarrow \beta \mathbf{v}_t - \alpha \hat{\mathbf{g}}_t$ ,
        $\mathbf{e}_{t+1} \leftarrow \mathbf{e}_t + \mathbf{v}_{t+1}$ 
6   Step 4: Projection to Tokens; Update  $\mathbf{e}_{t+1}$  via Eq. (7)–(9)
7   Step 5: Progressive Prompt Addition if  $\mathbf{e}_{t+1}$  succeeds on  $x^{(1)}, \dots, x^{(m_c)}$  and  $m_c < m$  then
       |  $m_c \leftarrow m_c + 1$ 
8   end
9 end
10 end
11 return  $\mathbf{x}^* \leftarrow \arg \max(-\mathbf{v}_T \cdot W^\top)$ 
```

3.3. Embedding Space Formulation

Challenge: Momentum in Discrete Space. Standard GCG operates in discrete token space, where momentum cannot be directly applied—averaging token indices is meaningless. Previous work (Zhang & Wei, 2025) accumulates gradient magnitudes for token ranking, but this loses directional information essential for Nesterov’s lookahead mechanism.

Continuous Relaxation. We reformulate optimization in a continuous embedding space. Let $W \in \mathbb{R}^{|\mathcal{V}| \times d}$ denote the token embedding matrix. We compute gradients with respect to suffix embeddings $\mathbf{e} \in \mathbb{R}^{l \times d}$: $g_t = \nabla_{\mathbf{e}} \sum_{m=1}^M \mathcal{L}_{m, p}(\mathbf{e})$.

Momentum accumulation operates in this continuous space: $\mathbf{v}_{t+1} = \beta \mathbf{v}_t + g_t$.

Projection to Tokens. We project the momentum velocity to discrete tokens for candidate selection:

$$\mathbf{s} = -\mathbf{v}_{t+1} \cdot W^\top \in \mathbb{R}^{L \times V} \quad (7)$$

$$\mathcal{C}_i = \arg \max_{S \subseteq [V], |S|=n} \min_{j \in S} s_{ij} \quad (8)$$

where $\mathbf{v}_{t+1}$ is the momentum velocity, $W \in \mathbb{R}^{V \times d}$ is the token embedding matrix, and $\mathcal{C}_i$ denotes the candidate token set for position i , containing the n tokens with the highest scores. Following GCG (Zou et al., 2023), we sample B candidates from $\prod_{i=1}^L \mathcal{C}_i$ and select the one with the lowest loss:

$$\mathbf{x}^* = \arg \min_{x \in \prod_{i=1}^L \mathcal{C}_i} L(\text{Embed}(x)), \quad \mathbf{e}_{t+1} = \text{Embed}(\mathbf{x}^*) \quad (9)$$

Algorithm 1 presents our Lookahead GCG method. The key modification is the SNAG update: instead of computing gradients at the current embedding $\mathbf{e}_t$, we evaluate at a lookahead position $\tilde{\mathbf{e}}_t = \mathbf{e}_t + \beta \mathbf{v}_t$ (Step 1). Momentum is accumulated in a continuous embedding space (Step 3), with periodic projection back to the discrete vocabulary (Step 4). The stochastic gradient samples one model and one prompt per iteration (Step 2), while progressive prompt addition (Step 5) gradually increases task difficulty. The differences between Lookahead GCG and TransferAttack (Yang et al., 2025) are marked in blue.

3.4. Theoretical Analysis

Assumption 3.1. Each $\mathcal{L}_{m, p}$ is L -smooth and $\|\nabla \mathcal{L}_{m, p}(\mathbf{e})\| \leq G$ for all $m, p, \mathbf{e}$.

These are standard assumptions in stochastic optimization (Yan et al., 2018).

Definition 3.2 (Stability Measures). Let $\mathcal{S}, \mathcal{S}' \subset [M] \times [P]$ be two training sets differing in K pairs. Let $(\mathbf{v}_t, \mathbf{e}_t)$ and $(\mathbf{v}'_t, \mathbf{e}'_t)$ denote the velocity and embedding sequences from training on $\mathcal{S}$ and $\mathcal{S}'$ respectively. **Velocity stability:** $\Delta_t^{(K)} := \mathbb{E}[\|\mathbf{v}_t - \mathbf{v}'_t\|]$. **Gradient-based projection:** $\Pi_{\mathcal{V}}(\mathbf{v})_i := \arg \text{top-}k_{j \in \mathcal{V}} (-\mathbf{v}_i \cdot \text{Embed}(j))$. **Projected stability:** $\Delta_T^{\text{proj}} := \mathbb{E}[\|\mathbf{e}^* - \mathbf{e}'^*\|]$ where $\mathbf{e}^* = \text{Embed}(\Pi_{\mathcal{V}}(\mathbf{v}_T))$. **Projection error:** $\epsilon_{\text{proj}} := \max\{\|\mathbf{e}^* - \mathbf{v}_T\|, \|\mathbf{e}'^* - \mathbf{v}'_T\|\}$

Stability-based generalization analysis (Hardt et al., 2016; Yan et al., 2018) bounds the generalization gap through algorithmic stability: if an algorithm produces similar outputs on similar training sets, it generalizes well. We measure stability via $\Delta_T^{(K)} = \mathbb{E}[\|\mathbf{v}_T - \mathbf{v}'_T\|]$, the expected velocity difference when training sets differ in K model-prompt pairs.

Theorem 3.3. Let $\eta_k^t(\text{SGD}) = \frac{1}{1-\beta}$ and $\eta_k^t(\text{SNAG}) = \frac{1-\beta^{t-k+2}}{1-\beta}$. Under Assumption 3.1, for training sets differing in K model-prompt pairs:

$$\Delta_T^{\text{proj}} \leq \Delta_T^{(K)} + 2\epsilon_{\text{proj}} \quad (10)$$

where $\Delta_T^{(K)}$ satisfies:

$$\Delta_{t+1}^{(K)} \leq \sum_{k=0}^t \frac{2\alpha GK}{MP} \eta_k^t + \left(1 - \frac{K}{MP}\right) \sum_{k=0}^t \alpha L \eta_k^t \Delta_k^{(K)} \quad (11)$$

Since $\eta_k^t(\text{SNAG}) < \eta_k^t(\text{SGD})$ for all $k < t$ and ϵ_{proj} is method-independent, we have $\Delta_T^{\text{proj}}(\text{SNAG}) < \Delta_T^{\text{proj}}(\text{SGD})$ for all $T \geq 1$.

Proof. Let $\mathcal{D}$ denote the set of K differing pairs. At step k , consider the sampled pair (m_k, p_k) :

Case 1: $(m_k, p_k) \notin \mathcal{D}$ (probability $1 - \frac{K}{MP}$). Both runs use the same loss function. By L -smoothness:

$$\|\hat{g}_k - \hat{g}'_k\| = \|\nabla \mathcal{L}_{m_k, p_k}(\mathbf{e}_k) - \nabla \mathcal{L}_{m_k, p_k}(\mathbf{e}'_k)\| \leq L \|\mathbf{e}_k - \mathbf{e}'_k\| \quad (12)$$

Case 2: $(m_k, p_k) \in \mathcal{D}$ (probability $\frac{K}{MP}$). Different loss functions are used. By bounded gradient:

$$\|\hat{g}_k - \hat{g}'_k\| \leq \|\hat{g}_k\| + \|\hat{g}'_k\| \leq 2G \quad (13)$$

Taking expectation:

$$\mathbb{E}[\|\hat{g}_k - \hat{g}'_k\|] \leq \frac{K}{MP} \cdot 2G + \left(1 - \frac{K}{MP}\right) L \Delta_k^{(K)} \quad (14)$$

By the cumulative update representation (Lemma 2 in (Yan

et al., 2018)), $\mathbf{e}_{t+1} - \mathbf{e}'_{t+1} = -\alpha \sum_{k=0}^t \eta_k^t (\hat{g}_k - \hat{g}'_k)$:

$$\Delta_{t+1}^{(K)} \leq \alpha \sum_{k=0}^t \eta_k^t \mathbb{E}[\|\hat{g}_k - \hat{g}'_k\|] \quad (15)$$

$$\leq \sum_{k=0}^t \frac{2\alpha GK}{MP} \eta_k^t + \left(1 - \frac{K}{MP}\right) \sum_{k=0}^t \alpha L \eta_k^t \Delta_k^{(K)} \quad (16)$$

Since $\eta_k^t(\text{SNAG}) = \frac{1-\beta^{t-k+2}}{1-\beta} < \frac{1}{1-\beta} = \eta_k^t(\text{SGD})$ for all $\beta \in (0, 1)$, by induction on t with $\Delta_0 = 0$, we have $\Delta_t^{(K)}(\text{SNAG}) < \Delta_t^{(K)}(\text{SGD})$.

Projected Stability: By triangle inequality:

$$\|\mathbf{e}^* - \mathbf{e}'^*\| \leq \|\mathbf{e}^* - \mathbf{e}_T\| + \|\mathbf{e}_T - \mathbf{e}'_T\| + \|\mathbf{e}'_T - \mathbf{e}'^*\| \quad (17)$$

Between projections, the embedding drifts by at most:

$$\sum_{t=T-K}^{T-1} \|\mathbf{v}_t\| \leq \frac{1\alpha G}{1-\beta} \quad (18)$$

where the bound $\|\mathbf{v}_t\| \leq \frac{\alpha G}{1-\beta}$ follows from standard Nesterov momentum analysis (Nesterov, 1983; Yan et al., 2018).

Define $\epsilon_{\text{proj}} := \frac{\alpha G}{1-\beta} + d_{\max}$, where $d_{\max} = \max_{i,j \in \mathcal{V}} \|\text{Embed}(i) - \text{Embed}(j)\|$ is the vocabulary diameter. This bound depends only on hyperparameters α, β, G and vocabulary $\mathcal{V}$ —identical for SNAG and SGD.

Taking expectation:

$$\Delta_T^{\text{proj}} \leq \Delta_T^{(K)} + 2\epsilon_{\text{proj}} \quad (19)$$

Since $\epsilon_{\text{proj}}^{\text{SNAG}} = \epsilon_{\text{proj}}^{\text{SGD}}$ and $\Delta_T^{(K)}(\text{SNAG}) < \Delta_T^{(K)}(\text{SGD})$, we conclude $\Delta_T^{\text{proj}}(\text{SNAG}) < \Delta_T^{\text{proj}}(\text{SGD})$. $\square$

When $M < P$, we have $\frac{1}{M} > \frac{1}{P}$, so $\Delta_t^{(\text{model})} > \Delta_t^{(\text{prompt})}$. Replacing one model affects P data points while replacing one prompt affects only M points—with fewer models, each model has greater influence, making cross-model transfer the bottleneck.

Since $\eta_k^t(\text{SNAG}) < \eta_k^t(\text{SGD})$, SNAG assigns smaller cumulative weights to historical gradients, reducing sensitivity to any single model perturbation. This directly translates to better cross-model generalization. Note that Transfer Attack/GCG's greedy coordinate search can be viewed as a discrete approximation to SGD: at each step, GCG selects tokens along the gradient direction and projects back to the vocabulary. Our analysis suggests replacing this implicit SGD with SNAG to improve transferability.

From an optimization perspective, Theorem 1 in (Yan et al., 2018) establishes that SGD, and SNAG share the same $O(1/\sqrt{t})$ convergence rate for non-convex objectives. Figure 1 confirms this: Lookahead-GCG achieves comparable training loss to TransferAttack.

3.5. Maximally Distant Template Selection

Distant templates induce larger representation shifts in the target model, amplifying adversarial impact. While standard GCG suffixes non improvement or even degrade under such shift, SNAG’s tighter stability bound preserves attack efficacy. This enables a simple transfer strategy: select the most distant template to maximize perturbation strength.

Specifically, given the optimized suffix embedding $\mathbf{e}^* \in \mathbb{R}^{L \times d}$ learned on source model m_s with native template $\mathcal{T}_s$, we compute template distance as:

$$d_{\text{emb}}(\mathcal{T}, \mathcal{T}_s) = 1 - \frac{\langle \mathbf{e}_{\mathcal{T}}^*, \mathbf{e}_{\mathcal{T}_s}^* \rangle}{\|\mathbf{e}_{\mathcal{T}}^*\| \cdot \|\mathbf{e}_{\mathcal{T}_s}^*\|} \quad (20)$$

where $\mathbf{e}_{\mathcal{T}}^*$ denotes the embedding of suffix $\mathbf{e}^*$ when composed with template $\mathcal{T}$.

For transfer evaluation on target model m_t , we select:

$$\mathcal{T}^* = \arg \max_{\mathcal{T} \in \mathcal{T}_{\text{all}}} d_{\text{emb}}(\mathcal{T}, \mathcal{T}_s) \quad (21)$$

We term this method Maximally Distant Template (MDT); see Algorithm 2.

Algorithm 2 Maximally Distant Template

Input: Optimized embedding $\mathbf{e}^*$, source template $\mathcal{T}_s$, source model M_s , target model M_t , candidate templates $\{\mathcal{T}_1, \dots, \mathcal{T}_K\}$, harmful question x , target output a

Output: Target model response

- 1 **Step 1: Compute Template Distances** $\mathbf{h}_s \leftarrow M_s(\mathbf{e}^*; \mathcal{T}_s)$
for $k = 1, \dots, K$ **do**
 - 2 $\mathbf{h}_k \leftarrow M_s(\mathbf{e}^*; \mathcal{T}_k)$, $d_k \leftarrow 1 - \frac{\langle \mathbf{h}_k, \mathbf{h}_s \rangle}{\|\mathbf{h}_k\| \cdot \|\mathbf{h}_s\|}$
 - 3 **end**
 - 4 **Step 2: Select Most Distant Template** $\mathcal{T}^* \leftarrow \arg \max_{\mathcal{T}_k} d_k$
 - 5 **Step 3: Transfer Attack** $p^* \leftarrow \Pi_{\mathcal{V}}(\mathbf{e}^*)$, response $\leftarrow M_t(p^* \| x \| a; \mathcal{T}^*)$
 - 6 **return** response
-

4. Experiments

4.1. Experimental Settings

Target Models We evaluate transfer attacks across a diverse set of LLMs with varying safety alignment levels. Our

open-source benchmark includes Llama-3.1-8B-Instruct (Patterson et al., 2022), Llama-2-7b-Chat (Touvron et al., 2023), Gemma-7B-It (Team et al., 2024b), Qwen2.5-7B (Team, 2024), and Vicuna-7B-v1.5 (Zheng et al., 2023). For closed-source commercial APIs evaluation, we target DeepSeek-V3.2 (Liu et al., 2025), Gemini-1.5-pro (Team et al., 2024a) and Grok-3-min (xAI, 2025).

Training Configuration Following the multi-model training protocol established in GCG (Zou et al., 2023), we utilize an ensemble of Llama-3.1-8B and Vicuna-7b-v1.5 as our source models for optimization. For a single model, we are only training on Llama-3.1-8B.

Our experiments are conducted on one A100s for single model training and two A100s for multi-model settings.

Data Following the experimental setup from (Yang et al., 2025), we use Harm-bench (Mazeika et al., 2024). Specifically, we train our method on a 20-behavior subset and evaluate its generalizability on the standard 200-behavior test set.

Baselines We evaluate the efficacy of Lookahead-GCG relative to three gradient-based baselines:

GCG (Zou et al., 2023): The original Greedy Coordinate Gradient method.

TransferAttack (Yang et al., 2025) improves transferability by optimizing a *prefix* with target output guidance and computing loss on only the first s tokens, removing GCG’s constraints on token count and tail matching.

SlotGCG (Jeong et al., 2026) improves upon GCG by exploiting positional vulnerabilities within prompts.

Momentum-GCG: An ablation of our method that replaces SNAG with standard (Heavy-ball) momentum. While both methods accumulate momentum in continuous embedding space, Momentum-GCG computes gradients at the current position $\mathbf{e}_t$, whereas Lookahead-GCG computes gradients at the lookahead position $\tilde{\mathbf{e}}_t = \mathbf{e}_t + \beta \mathbf{v}_t$. This isolates the contribution of the lookahead mechanism to transferability.

Evaluation Protocol We utilize HarmBench-Llama-2-13B-cla (Mazeika et al., 2024) as our automated evaluator. To ensure statistical robustness, all experiments are conducted for at least four independent trials. We report the mean Attack Success Rate (ASR) and standard deviation to characterize performance stability across runs.

Hyperparameters All methods are evaluated under a unified framework (Table 3) using 20 harmful behaviors with $L = 100$, $T = 500$, $B = 128$, and top- $k = 256$. Multi-model configurations utilize an identical two-model ensemble.

Table 1. Multi-model (Model Number = 2) vs Single-model comparison with Maximally Distant Template on ASR. **With Maximally Distant Template** indicates target models using native instruction formats; **With Maximally Distant Template** enforces a format shift. **Bold** indicates the SOTA performance in the multi-model setting; $\uparrow$ and $\downarrow$ denote change relative to the single-model baseline. The asterisk (*) denotes statistical significance with $p < 0.05$ based on pairwise comparisons, indicating that Lookahead-GCG’s improvements over baselines are statistically robust.

Method	Training	Gemma	Llama-2	Llama-3	Qwen	Vicuna	Avg.
<i>Without Maximally Distant Template</i>							
GCG (Zou et al., 2023)	Multi-model	22.00 $\pm$ 8.60	5.75 $\pm$ 3.25	24.25 $\pm$ 28.66	17.67 $\pm$ 4.19	73.83 $\pm$ 5.33	28.70
	Single-model	8.78 $\pm$ 4.31	0.50 $\pm$ 0.35	19.75 $\pm$ 14.34	9.88 $\pm$ 3.43	13.50 $\pm$ 6.81	10.48
	Δ	$\uparrow$ 13.22	$\uparrow$ 5.25	$\uparrow$ 4.50	$\uparrow$ 7.79	$\uparrow$ 60.33	$\uparrow$ 18.22
TransferAttack (Yang et al., 2025)	Multi-model	46.00 $\pm$ 11.53	14.80 $\pm$ 8.40	23.80 $\pm$ 14.93	30.80 $\pm$ 4.01	79.60 $\pm$ 12.70	39.00
	Single-model	42.71 $\pm$ 1.68	10.56 $\pm$ 5.05	23.79 $\pm$ 6.41	28.86 $\pm$ 4.62	86.62 $\pm$ 0.83	38.51
	Δ	$\uparrow$ 3.29	$\uparrow$ 4.24	$\uparrow$ 0.01	$\uparrow$ 1.94	$\downarrow$ 7.02	$\uparrow$ 0.49
Momentum-GCG	Multi-model	48.90 $\pm$ 9.70	16.38 $\pm$ 3.35	28.88 $\pm$ 10.30	33.88 $\pm$ 13.00	87.25 $\pm$ 0.65	43.06
	Single-model	43.79 $\pm$ 3.99	9.12 $\pm$ 4.86	21.12 $\pm$ 10.66	26.62 $\pm$ 2.18	88.54 $\pm$ 2.22	37.84
	Δ	$\uparrow$ 5.11	$\uparrow$ 7.26	$\uparrow$ 7.76	$\uparrow$ 7.26	$\downarrow$ 1.29	$\uparrow$ 5.22
Slot-GCG (Jeong et al., 2026)	Multi-model	17.33 $\pm$ 3.97	6.67 $\pm$ 3.86	10.17 $\pm$ 1.43	23.67 $\pm$ 1.43	22.50 $\pm$ 7.82	16.07
	Single-model	12.00 $\pm$ 0.41	7.00 $\pm$ 2.55	9.00 $\pm$ 1.08	24.50 $\pm$ 3.54	8.17 $\pm$ 3.47	12.13
	Δ	$\uparrow$ 5.33	$\downarrow$ 0.33	$\uparrow$ 1.17	$\downarrow$ 0.83	$\uparrow$ 14.33	$\uparrow$ 3.93
Lookahead-GCG (Ours)	Multi-model	56.36$\pm$0.82*	21.65$\pm$5.28*	50.63$\pm$13.30*	34.83$\pm$2.25*	88.38$\pm$0.41*	50.37
	Single-model	42.13 $\pm$ 9.33	9.17 $\pm$ 2.87	21.50 $\pm$ 8.20	32.13 $\pm$ 7.53	88.03 $\pm$ 1.83	38.59
	Δ	$\uparrow$ 14.23	$\uparrow$ 12.48	$\uparrow$ 29.13	$\uparrow$ 2.70	$\uparrow$ 0.35	$\uparrow$ 11.78
<i>With Maximally Distant Template</i>							
GCG (Zou et al., 2023)	Multi-model	36.33 $\pm$ 14.42	15.25 $\pm$ 5.98	27.83 $\pm$ 16.29	17.67 $\pm$ 8.11	23.00 $\pm$ 14.85	24.02
	Single-model	12.62 $\pm$ 9.38	10.56 $\pm$ 5.05	15.38 $\pm$ 10.44	15.62 $\pm$ 11.93	18.38 $\pm$ 5.31	14.51
	Δ	$\uparrow$ 23.71	$\uparrow$ 4.69	$\uparrow$ 12.45	$\uparrow$ 2.05	$\uparrow$ 4.62	$\uparrow$ 9.50
TransferAttack (Yang et al., 2025)	Multi-model	45.75 $\pm$ 5.37	49.90 $\pm$ 23.37	41.80 $\pm$ 20.23	28.00 $\pm$ 7.72	16.38 $\pm$ 32.42	36.37
	Single-model	47.64 $\pm$ 8.10	38.86 $\pm$ 14.03	2.86 $\pm$ 5.87	31.43 $\pm$ 8.58	12.50 $\pm$ 23.81	26.66
	Δ	$\downarrow$ 1.89	$\uparrow$ 11.04	$\uparrow$ 38.94	$\downarrow$ 3.43	$\uparrow$ 3.88	$\uparrow$ 9.71
Momentum-GCG	Multi-model	50.00 $\pm$ 3.65	61.62 $\pm$ 7.20	65.62 $\pm$ 16.40	32.00 $\pm$ 7.88	19.12 $\pm$ 3.35	45.67
	Single-model	47.17 $\pm$ 5.55	43.21 $\pm$ 11.60	23.15 $\pm$ 3.90	27.83 $\pm$ 3.83	3.08 $\pm$ 3.30	28.89
	Δ	$\uparrow$ 2.83	$\uparrow$ 18.41	$\uparrow$ 42.47	$\uparrow$ 4.17	$\uparrow$ 16.04	$\uparrow$ 16.78
Slot-GCG (Jeong et al., 2026)	Multi-model	20.33 $\pm$ 3.66	10.17 $\pm$ 2.93	23.00 $\pm$ 3.90	21.17 $\pm$ 1.43	1.83 $\pm$ 7.19	15.30
	Single-model	10.33 $\pm$ 4.03	7.33 $\pm$ 3.42	5.17 $\pm$ 1.18	23.67 $\pm$ 4.73	1.33 $\pm$ 1.18	9.57
	Δ	$\uparrow$ 10.00	$\uparrow$ 2.84	$\uparrow$ 17.83	$\downarrow$ 2.50	$\uparrow$ 0.50	$\uparrow$ 5.73
Lookahead-GCG (Ours)	Multi-model	59.00$\pm$6.05*	70.13$\pm$6.66*	75.50$\pm$3.10*	34.75$\pm$3.40*	34.50$\pm$5.28*	54.78
	Single-model	48.83 $\pm$ 3.32	53.49 $\pm$ 21.46	3.33 $\pm$ 2.36	29.00 $\pm$ 2.91	8.87 $\pm$ 9.81	28.70
	Δ	$\uparrow$ 10.17	$\uparrow$ 16.64	$\uparrow$ 72.17	$\uparrow$ 5.75	$\uparrow$ 25.63	$\uparrow$ 26.07

ble to ensure consistency.

Attack Format Unlike the original GCG which optimizes an adversarial *suffix*, TransferAttack, Momentum-GCG, and Lookahead-GCG utilize an adversarial *prefix*. To standardize evaluation on instruction-tuned models, we enforce constrained response templates. For GCG, the format is [Query]+[Suffix], targeting an affirmative start (e.g., “Sure, here is...”). For prefix-based methods, we use [Prefix]+[Query], followed by a system-enforced prompt: “Start your response with: Here is [Goal]. Step 1:... ”.

4.2. Main Results

Evaluation on Open-Source Architectures Table 1 reports the quantitative performance of Lookahead-GCG (Ours) across five open-source architectures under two distinct template configurations. In the matched-template setting, the multi-model configuration ($M = 2$) of Lookahead-GCG yields a mean Attack Success Rate (ASR) of **50.37%**, surpassing baseline prefix-based methods such as TransferAttack (39.00%) and Momentum-GCG (43.06%). Statistical significance is indicated by asterisk markers ($p < 0.05$). On *Llama-3*, Lookahead-GCG attains a **50.63%** ASR, compared to 24.25% for the GCG baseline. The marginal improvement from single-model to multi-model ($M = 2$) optimization for Lookahead-GCG is $\Delta = +11.78\%$, while TransferAttack remains relatively stagnant ($\Delta =$

Table 2. Attack Success Rate (ASR) of triple-model transfer attacks. **With Maximally Distant Template** indicates target models evaluated using their native instruction formats, while **Without Maximally Distant Template** enforces a format shift during evaluation. **Bold** indicates the best performance in each setting. The values represent mean $\pm$ standard deviation, and the asterisk (*) denotes statistical significance with $p < 0.05$ based on pairwise comparisons, highlighting the robust improvements of Lookahead-GCG (Ours) over the baselines.

Method	Training	Gemma	Qwen	Avg.
<i>Without Maximally Distant Template</i>				
GCG (Zou et al., 2023)	Triple-model	16.00 $\pm$ 1.00	15.75 $\pm$ 6.72	15.88
TransferAttack (Yang et al., 2025)	Triple-model	51.75 $\pm$ 5.30	42.25 $\pm$ 10.96	47.00
Slot-GCG (Jeong et al., 2026)	Triple-model	24.00 $\pm$ 3.00	22.50 $\pm$ 11.31	23.25
Lookahead-GCG (Ours)	Triple-model	58.75$\pm$1.77*	63.00$\pm$2.00*	60.88
<i>With Maximally Distant Template</i>				
GCG (Zou et al., 2023)	Triple-model	13.00 $\pm$ 3.00	13.75 $\pm$ 2.47	13.38
TransferAttack (Yang et al., 2025)	Triple-model	35.00 $\pm$ 22.00	35.00 $\pm$ 4.95	35.00
Slot-GCG (Jeong et al., 2026)	Triple-model	13.50 $\pm$ 10.64	20.83 $\pm$ 11.74	17.17
Lookahead-GCG (Ours)	Triple-model	49.75$\pm$2.47*	56.50$\pm$2.12*	53.13

Table 3. Training and attack configurations for all methods.

Category & Parameter	Value (Shared)
<i>Attack Settings</i>	
Adversarial Length (L)	100
Optimization Steps (T)	500
Batch Size (B) / Top- k	128 / 256
Evaluation Interval	5
<i>Optimization Strategy</i>	
Surrogate Models	2 (Ensemble)
Progressive Scheduling	False
ASCII-only / Early Stopping	True / False
<i>Prompt Configuration</i>	
Loss-on-Prefix	Yes (No for GCG)
Fixed Token Sampling	2 (N/A for GCG)
System Message	No

+0.49%). Under format shifts (mismatched-template setting), Lookahead-GCG ($M = 2$) records a mean ASR of **54.78%**, reaching **75.50%** on *Llama-3* and **70.13%** on *Llama-2*.

To further evaluate scalability under scaling source models, Table 2 presents the transfer ASR under the triple-model training configuration ($M = 3$). When moving to this more challenging multi-model ecosystem, Lookahead-GCG consistently outshines all competitors. Specifically, *Without Maximally Distant Template*, our method achieves an impressive average ASR of **60.88%**, markedly outperforming TransferAttack (47.00%) and Slot-GCG (23.25%). This performance advantage is preserved even *With Maximally Distant Template* (enforcing a native instruction format), where Lookahead-GCG maintains a commanding average ASR of **53.13%**, compared to just 35.00% for TransferAttack. Pairwise comparisons demonstrate that these scaling improvements are statistically robust ($p < 0.05$). Qualita-

tive evidence in Appendix A (e.g., experiments on *Llama-3.1* and *Qwen-2.5*) shows that Lookahead-GCG generates prefixes that incorporate structured formatting and persona-based instructions, such as “assistant smooth tone” and “positive answer” prompts, which contrast with the disjointed token sequences produced by baseline methods.

Evaluation on Closed-Source Commercial APIs Table 4 and Table 5 evaluate the cross-family transferability of optimized prefixes against five premier commercial LLMs: **DeepSeek-V3.2**, **Gemini-1.5-Pro**, **grok-3-mini**, **GPT-4.1-mini**, and **GPT-5-mini**.

Under the multi-model setting ($M = 2$) shown in Table 4, Lookahead-GCG (Ours) achieves the highest average ASR of **23.85%**, outperforming GCG (15.22%), TransferAttack (15.08%), Momentum-GCG (21.23%), and Slot-GCG (13.63%). Specifically, our method secures a performance gains of $\Delta = +3.76\%$ moving from single- to multi-surrogate optimization, and achieves individual ASRs of **16.00%** on DeepSeek-V3.2, **58.83%** on Gemini-1.5-pro, **27.25%** on grok-3-mini, and **15.17%** on GPT-4.1-mini. On the highly fortified GPT-5-mini, Lookahead-GCG is the only method to attain a non-zero ASR of **2.00%**.

When further scaling to the triple-model training configuration ($M = 3$) in Table 5, Lookahead-GCG consistently retains its state-of-the-art transferability with an average ASR of **23.20%**, compared to 7.85% for GCG, 21.10% for TransferAttack, and 11.47% for Slot-GCG. In this setting, Lookahead-GCG achieves **14.00%**, **51.50%**, **34.25%**, and **14.50%** ASR on DeepSeek-V3.2, Gemini-1.5-pro, grok-3-mini, and GPT-4.1-mini, respectively, while maintaining a **1.75%** ASR on GPT-5-mini. Across both configurations, pairwise comparisons demonstrate that our improvements are statistically significant ($p < 0.05$).

Qualitative observations in Appendix A document the content of prefixes generated by Lookahead-GCG. Across these commercial targets, the optimized prefixes contain structured strings adopting specific personae, such as an “encouraging assistant” or a “historian.”

Instruction Template Distance Analysis Table 6 provides the quantitative basis for our *MDT* strategy. To induce maximum representational shifts, we select the template with the highest divergence d_{emb} for each source model. Specifically, the **Llama-2 template** is identified as the optimal MDT for *Llama-3*, *Gemma*, *Vicuna*, and *Qwen*, corresponding to distances of **0.51**, **0.51**, **0.49**, and **0.81**, respectively. Conversely, for the *Llama-2* source model, the **Qwen template** yields the maximum divergence (**0.81**). This selection ensures the SNAG optimizer is trained against the most challenging formatting perturbations available.

Table 4. Multi-model vs Single-model attack comparison on commercial API models on ASR. All target models are evaluated using their respective native instruction templates. **Bold** indicates the SOTA performance in the multi-model setting; results for highly fortified commercial APIs (i.e., GPT-4.1-mini and GPT-5-mini) are left unbolded due to their distinct defense characteristics. Blue arrows ($\uparrow$), red arrows ($\downarrow$), and orange bars ($\parallel$) denote ASR increase, decrease, and zero-change relative to the single-model baseline, respectively. P-values are derived from pairwise comparisons.

Method	DeepSeek-V3.2	Gemini-1.5-pro	grok-3-mini	GPT-4.1-mini	GPT-5-mini	Avg
<i>GCG (Zou et al., 2023)</i>						
Multi-model	1.00 $\pm$ 0.00*	53.75 $\pm$ 5.40*	15.33 $\pm$ 2.49*	6.00 $\pm$ 0.85	0.00 $\pm$ 0.00	15.22
Single-model	0.83 $\pm$ 0.85	4.25 $\pm$ 2.97	6.62 $\pm$ 5.81	4.33 $\pm$ 2.72	0.00 $\pm$ 0.00	3.21
Δ	$\uparrow$ 0.17	$\uparrow$ 49.50	$\uparrow$ 8.71	$\uparrow$ 1.67	$\parallel$ 0.00	$\uparrow$ 12.01
<i>TransferAttack (Yang et al., 2025)</i>						
Multi-model	12.00 $\pm$ 1.08*	33.00 $\pm$ 16.31*	17.75 $\pm$ 3.50*	12.67 $\pm$ 0.58	0.00 $\pm$ 0.00	15.08
Single-model	13.83 $\pm$ 1.64	45.29 $\pm$ 5.82	19.43 $\pm$ 3.13	10.50 $\pm$ 0.82	0.00 $\pm$ 0.00	17.81
Δ	$\downarrow$ 1.83	$\downarrow$ 12.29	$\downarrow$ 1.68	$\uparrow$ 2.17	$\parallel$ 0.00	$\downarrow$ 2.73
<i>Momentum-GCG</i>						
Multi-model	15.67 $\pm$ 1.65	58.67 $\pm$ 0.85	17.83 $\pm$ 2.75*	14.00 $\pm$ 1.50	0.00 $\pm$ 0.00	21.23
Single-model	12.00 $\pm$ 0.71	49.67 $\pm$ 4.36	22.00 $\pm$ 0.70	9.33 $\pm$ 1.03	0.00 $\pm$ 0.00	18.60
Δ	$\uparrow$ 3.67	$\uparrow$ 9.00	$\downarrow$ 4.17	$\uparrow$ 4.67	$\parallel$ 0.00	$\uparrow$ 2.63
<i>Slot-GCG (Jeong et al., 2026)</i>						
Multi-model	14.00 $\pm$ 2.45	22.50 $\pm$ 0.50	17.50 $\pm$ 3.19	14.16 $\pm$ 0.76	0.00 $\pm$ 0.00	13.63
Single-model	5.50 $\pm$ 2.55	29.50 $\pm$ 0.50	2.33 $\pm$ 0.70	13.33 $\pm$ 2.39	0.00 $\pm$ 0.00	10.13
Δ	$\uparrow$ 8.50	$\downarrow$ 7.00	$\uparrow$ 15.17	$\uparrow$ 0.83	$\parallel$ 0.00	$\uparrow$ 3.50
<i>Lookahead-GCG (Ours)</i>						
Multi-model	16.00$\pm$2.94*	58.83$\pm$1.43*	27.25$\pm$1.79*	15.17$\pm$0.76	2.00$\pm$0.41	23.85
Single-model	10.63 $\pm$ 3.42	51.33 $\pm$ 1.03	26.67 $\pm$ 8.65	10.33 $\pm$ 3.47	1.50 $\pm$ 0.01	20.09
Δ	$\uparrow$ 5.37	$\uparrow$ 7.50	$\uparrow$ 0.58	$\uparrow$ 4.84	$\uparrow$ 0.50	$\uparrow$ 3.76

Table 5. Transfer attack success rate (ASR) comparison on commercial API models. All target models are evaluated using their respective native instruction templates. **Bold** indicates the SOTA performance under the triple-model training setting; target columns for highly secure commercial APIs (i.e., GPT-4.1-mini and GPT-5-mini) are left unbolded due to their distinct defense mechanisms. The asterisk (*) denotes statistical significance with $p < 0.05$ based on pairwise comparisons, highlighting the robust transferability of Lookahead-GCG (Ours).

Method	DeepSeek-V3.2	Gemini-1.5-pro	grok-3-mini	GPT-4.1-mini	GPT-5-mini	Avg
<i>GCG (Zou et al., 2023)</i>						
Triple-model	5.25 $\pm$ 0.03*	13.50 $\pm$ 0.02	14.00 $\pm$ 0.02	6.50 $\pm$ 0.01	0.00 $\pm$ 0.00	7.85
<i>TransferAttack (Yang et al., 2025)</i>						
Triple-model	9.50 $\pm$ 0.01*	49.00 $\pm$ 0.18	33.25 $\pm$ 0.66	13.75 $\pm$ 0.00	0.00 $\pm$ 0.00	21.10
<i>Slot-GCG (Jeong et al., 2026)</i>						
Triple-model	12.67 $\pm$ 0.004	17.50 $\pm$ 9.20	14.83 $\pm$ 10.77	12.33 $\pm$ 6.56	0.00 $\pm$ 0.00	11.47
<i>Lookahead-GCG (Ours)</i>						
Triple-model	14.00$\pm$0.00*	51.50$\pm$0.08*	34.25$\pm$0.36*	14.50$\pm$0.01	1.75$\pm$0.01	23.20

Training Time As shown in Table 7, GCG and TransferAttack require **10.4h** and **12.0h** for training in the single-model setting, respectively. Lookahead-GCG (Ours) takes **13.5h**, which introduces only a marginal time overhead compared to the baselines while delivering superior performance.

5. Conclusion

We present Lookahead-GCG, a transferable universal jail-breaking framework that revives multi-model adversarial training through principled optimization. By combining Stochastic Nesterov Accelerated Gradient with embedding-space momentum and maximally distant initialization, our method addresses the gradient instability that limited prior multi-model approaches. Experiments demonstrate significant improvements over baselines on both open-source (50.37% ASR) and closed-source models (34.03% ASR),

Table 6. Pairwise template distance matrix across open-source LLMs. Smaller values indicate higher similarity in response formatting.

	Llama-3	Llama-2	Vicuna	Gemma	Qwen
Llama-3	0.0000	0.5081	0.0747	0.0010	0.4424
Llama-2	0.5081	0.0000	0.4858	0.5081	0.8062
Vicuna	0.0747	0.4858	0.0000	0.0747	0.4575
Gemma	0.0010	0.5081	0.0747	0.0000	0.4424
Qwen	0.4424	0.8062	0.4575	0.4424	0.0000

Table 7. Average training time (hours). “Single” denotes single-model.

Method	Single (h)
GCG (Zou et al., 2023)	10.4
TransferAttack (Yang et al., 2025)	12.0
Lookahead-GCG (Ours)	13.5

with multi-model optimization yielding +11.78% gains where existing methods show minimal improvement. Our findings suggest that optimizer design is a critical yet under-explored factor in adversarial transferability.

Impact Statement

This work investigates transferable adversarial attacks on large language models to advance the understanding of their vulnerabilities. We emphasize that the primary goal is to support *red-teaming* efforts—proactively identifying weaknesses in safety alignment before malicious actors can exploit them.

We do not condone or encourage any malicious use of these techniques. Instead, we advocate for responsible disclosure and collaborative efforts between researchers and model developers to strengthen safeguards. By revealing the limitations of current defenses, we aim to inform the development of more robust safety mechanisms.

All experiments were conducted on publicly available models and standard benchmarks. We did not deploy attacks against production systems or real users.




Figure 1. Training loss curves across different methods. From an optimization perspective, Theorem 1 in (Yan et al., 2018) establishes that SGD, and SNAG share the same $O(1/\sqrt{t})$ convergence rate for non-convex objectives. This figure confirms this: Lookahead-GCG achieves comparable training loss to TransferAttack.

Acknowledgements

This work was supported by the Scientific Research Start-up Fund of Jilin University (Grant No. 419080526A65).

We would like to thank Qixin Zhang for helpful discussions and Suming Yu for her assistance. We are also grateful to all reviewers for the insightful comments, and to the Area Chair for recognizing and supporting the value of our work on universal attacks.

References

- Chao, P., Debenedetti, E., Robey, A., Andriushchenko, M., Croce, F., Schwag, V., Dobriban, E., Flammarion, N., Pappas, G. J., Tramer, F., et al. Jailbreakbench: An open robustness benchmark for jailbreaking large language models. *Advances in Neural Information Processing Systems*, 37:55005–55029, 2024.
- Gu, L., Ji, X., Zhang, Z., Ma, J., Jia, X., and Jiang, W. Sm-gcg: Spatial momentum greedy coordinate gradient for robust jailbreak attacks on large language models. *Electronics*, 14(19), 2025. ISSN 2079-9292. doi: 10.3390/electronics14193967. URL <https://www.mdpi.com/2079-9292/14/19/3967>.
- Hardt, M., Recht, B., and Singer, Y. Train faster, generalize better: Stability of stochastic gradient descent. In *International conference on machine learning*, pp. 1225–1234. PMLR, 2016.
- Jeong, S., Jeong, J., Kim, H., Lee, Y., and Lee, W. Slot-GCG: Exploiting the positional vulnerability in LLMs for jailbreak attacks. In *The Fourteenth International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=Fn2rSOnpNf>.
- Jia, X., Pang, T., Du, C., Huang, Y., Gu, J., Liu, Y., Cao, X., and Lin, M. Improved techniques for optimization-based jailbreaking on large language models. *arXiv preprint arXiv:2405.21018*, 2024.
- Jiang, F., Xu, Z., Niu, L., Lin, B. Y., and Poovendran, R. Chatbug: A common vulnerability of aligned llms induced by chat templates. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 27347–27355, 2025.
- Li, J., Hao, Y., Xu, H., Wang, X., and Hong, Y. Exploiting the index gradients for optimization-based jailbreaking on large language models. In *Proceedings of the 31st International Conference on Computational Linguistics*, pp. 4535–4547, 2025.
- Li, R., Wang, H., and Mao, C. Largo: Latent adversarial reflection through gradient optimization for jailbreaking llms. *Advances in Neural Information Processing Systems*, 38:44612–44631, 2026.
- Liu, A., Mei, A., Lin, B., Xue, B., Wang, B., Xu, B., Wu, B., Zhang, B., Lin, C., Dong, C., et al. Deepseek-v3. 2: Pushing the frontier of open large language models. *arXiv preprint arXiv:2512.02556*, 2025.
- Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., Sakhaee, E., Li, N., Basart, S., Li, B., Forsyth, D., and Hendrycks, D. Harmbench: A standardized evaluation framework for automated red teaming and robust refusal, 2024. URL <https://arxiv.org/abs/2402.04249>.
- Nesterov, Y. A method for solving the convex programming problem with convergence rate $o(1/k^2)$. In *Dokl akad nauk Sssr*, volume 269, pp. 543, 1983.
- Patterson, D., Gonzalez, J., Hölzle, U., Le, Q., Liang, C., Munguia, L.-M., Rothchild, D., So, D. R., Texier, M., and Dean, J. The carbon footprint of machine learning training will plateau, then shrink. *Computer*, 55(7):18–28, 2022.
- Sutskever, I., Martens, J., Dahl, G., and Hinton, G. On the importance of initialization and momentum in deep learning. In *International conference on machine learning*, pp. 1139–1147. pmlr, 2013.
- Team, G., Georgiev, P., Lei, V. I., Burnell, R., Bai, L., Gulati, A., Tanzer, G., Vincent, D., Pan, Z., Wang, S., et al. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context, 2024. URL <https://arxiv.org/abs/2403.05530>, 2403, 2024a.
- Team, G., Mesnard, T., Hardin, C., Dadashi, R., Bhupatiraju, S., Pathak, S., Sifre, L., Rivière, M., Kale, M. S., Love, J., et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024b.
- Team, Q. Qwen2.5: A party of foundation models. September 2024. URL <https://qwenlm.github.io/blog/qwen2.5/>.
- Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., Bashlykov, N., Batra, S., Bhargava, P., Bhosale, S., Bikel, D., Blecher, L., Ferrer, C. C., Chen, M., Cucurull, G., Esiobu, D., Fernandes, J., Fu, J., Fu, W., Fuller, B., Gao, C., Goswami, V., Goyal, N., Hartshorn, A., Hosseini, S., Hou, R., Inan, H., Kardaş, M., Kerkez, V., Khabsa, M., Kloumann, I., Korenev, A., Koura, P. S., Lachaux, M.-A., Lavril, T., Lee, J., Liskovich, D., Lu, Y., Mao, Y., Martinet, X., Mihaylov, T., Mishra, P., Molybog, I., Nie, Y., Poulton, A., Reizenstein, J., Rungta, R., Saladi, K., Schelten, A., Silva, R., Smith, E. M., Subramanian, R., Tan, X. E., Tang, B., Taylor, R., Williams,

- A., Kuan, J. X., Xu, P., Yan, Z., Zarov, I., Zhang, Y., Fan, A., Kambadur, M., Narang, S., Rodriguez, A., Stojnic, R., Edunov, S., and Scialom, T. Llama 2: Open foundation and fine-tuned chat models, 2023. URL <https://arxiv.org/abs/2307.09288>.
- xAI. Grok 3 beta — the age of reasoning agents, 2025. URL <https://x.ai/news/grok-3>.
- Yan, Y., Yang, T., Li, Z., Lin, Q., and Yang, Y. A unified analysis of stochastic momentum methods for deep learning. *arXiv preprint arXiv:1808.10396*, 2018.
- Yang, J., Zhang, Z., Cui, S., Wang, H., and Huang, M. Guiding not forcing: Enhancing the transferability of jailbreaking attacks on llms via removing superfluous constraints. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 19643–19655, 2025.
- Zhang, Y. and Wei, Z. Boosting jailbreak attack with momentum. In *ICASSP 2025-2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 1–5. IEEE, 2025.
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E., et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36: 46595–46623, 2023.
- Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., and Fredrikson, M. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

A. Case Studies

To qualitatively evaluate adversarial prompt, we select three representative target models across distinct deployment paradigms: Llama-3.1-8B (the white-box source model during training), Qwen2.5-8B (an open-source model for transfer evaluation), and DeepSeek-V3.2 (a closed-source commercial API for black-box transfer evaluation). The corresponding empirical generation logs are visualized in Figures 2, 3, and 4. For safety compliance, the full explicit toxic prompt texts in these visualizations are redacted with a mosaic filter, and text highlighted in red explicitly denotes explicit model refusal keywords.




Figure 2. DeepSeek Case studies. Full malicious prompt text is redacted for safety, and red text highlights standard refusal keywords.




Figure 4. Qwen Case studies. Full malicious prompt text is redacted for safety, and red text highlights standard refusal keywords.