
CBV: Clean-label Backdoor Attacks on Vision Language Models via Diffusion Models

Ji Guo¹ Xiaolong Qin² Cencen Liu¹ Jielei Wang¹ Jierun Chen³ Wenbo Jiang¹

Abstract

Vision-Language Models (VLMs) have achieved remarkable success in tasks such as image captioning and visual question answering (VQA). However, as their applications become increasingly widespread, recent studies have revealed that VLMs are vulnerable to backdoor attacks. Existing backdoor attacks on VLMs primarily rely on data poisoning by adding visual triggers and modifying text labels, where the induced image-text mismatch makes poisoned samples easy to detect. To address this limitation, we propose the Clean-Label Backdoor Attack on VLMs via Diffusion Models (CBV), which leverages diffusion models to generate natural poisoned examples via score matching. Specifically, CBV modifies the score during the reverse generation process of the diffusion model to guide the generation of poisoned samples that contain triggered image features. To further enhance the effectiveness of the attack, we incorporate the textual information of the triggered images as multimodal guidance during generation. Moreover, to enhance stealthiness, we introduce a GradCAM-guided Mask (GM) that restricts modifications to only the most semantically important regions, rather than the entire image. We evaluate our method on MSCOCO and VQA v2 with four representative VLMs, achieving over 80% ASR while preserving normal functionality.




Figure 1. Comparisons of TrojVLM (Lyu et al., 2024), BadToken (Bai et al., 2024) and CBV poisoned sample from MSCOCO (Chen et al., 2015). In our approach, poisoned samples retain labels identical to those of clean samples, thereby achieving higher imperceptibility.

1. Introduction

Vision-Language Models (VLMs) (Li et al., 2022; 2023; Bai et al., 2024), which extend large language models (LLMs) (Touvron & et al., 2023a;b; Yang & et al., 2024) with visual perception capabilities, have recently emerged as a powerful framework for multimodal understanding. By accepting images or image-text pairs as input, VLMs can answer questions about visual content and perform a comprehensive understanding and reasoning over depicted scenes. Representative models such as BLIP-2 (Li et al., 2023) and Qwen-VL (Bai et al., 2024) have demonstrated remarkable performance in tasks including image captioning (Hossain et al., 2019) and visual question answering (VQA) (Wu et al., 2017).

While VLMs have achieved success in many fields, researchers have begun to investigate their security (Ye et al., 2025). Recent studies (Lyu et al., 2024) have revealed that VLMs are vulnerable to backdoor attacks, where attackers implant hidden backdoor through data poisoning. Subsequent work has expanded VLM backdoor attacks to multimodal poisoned samples (Han et al., 2024; Liang et al., 2025a), more stealthy attacks (Xu et al., 2024), token-level backdoor (Yuan et al., 2025), and attacks that remain effective under domain shift (Liang et al., 2025b). However, all these approaches require the modification of the text labels corresponding to the triggered images during the construction of poisoned data, making them easier to detect (see Fig. 1).

In this paper, we explore clean-label backdoor attacks on

¹Laboratory of Intelligent Collaborative Computing, University of Electronic Science and Technology of China, China
²School of Software Engineering, Chengdu University of Information Technology, China ³The Hong Kong University of Science and Technology, China. Correspondence to: Wenbo Jiang <wenbo-jiang@uestc.edu.cn>.




Figure 2. Visual comparison of poisoned images produced by different backdoor methods (CBAVR (Zhao et al., 2020), Narcissus (Zeng et al., 2023), HTBA (Saha et al., 2020), SAA (Souri et al., 2022), Shadowcast (Xu et al., 2024), IP (Ning et al., 2021)). Prior approaches often introduce visible noise or artifacts that make poisoned samples conspicuous and easy to detect. In contrast, CBV generates poisoned images with a natural looking appearance, avoiding noticeable artifacts and thus improving stealth.

VLMs. Unlike previous dirty-label attacks (Gu et al., 2019; Chen et al., 2017) that rely on label manipulation, clean-label attacks on VLMs impose a stricter constraint where the poisoned samples must retain exactly the same label as clean ones, leaving the attacker only able to modify the image. In the clean-label attacks setting, the attack also requires the backdoor VLM to produce attacker-specified outputs (e.g., semantic substitutions from “people” to “dog”) when input triggered images, while maintaining normal performance on clean inputs.

It is challenging to realize such a clean-label backdoor attack against VLMs. Previous clean-label backdoor attacks methods mainly focus on traditional vision tasks such as image classification (Zeng et al., 2023; Saha et al., 2020; Souri et al., 2022; Yu et al., 2024; Ning et al., 2021) or video classification (Zhao et al., 2020). Their core idea is to leverage a surrogate model to approximate the triggered samples and class-correct clean samples in the feature space, thereby generating poisoned samples that contain trigger features while maintaining correct labels. However, unlike traditional vision tasks that only require class-level information, VLMs handle tasks such as image captioning and VQA, which demand much richer semantic content from the image. Consequently, poisoned samples generated via surrogate model based approximation often fail to encode complex semantic information, making the attack ineffective.

To achieve an effective clean-label backdoor attack against VLMs, we propose Clean-label Backdoor Attacks on Vision Language Models via Diffusion Models (CBV), which leverages diffusion models (Dhariwal & Nichol, 2021) to generate natural poisoned examples consistent with their original labels via score matching (Hyvarinen, 2005). Our key idea is to modify the denoising score function of the diffusion model, using the triggered image as guidance to generate clean-label poisoned samples that inherit the trigger features. To further enhance the triggered image features in poisoned samples, we incorporate the textual semantics of the triggered image as a multimodal guide during the denoising process, to align the text-image feature in the latent space, and thus improving the overall transferability and

effectiveness of the attack. In addition, we design a universal adversarial perturbation (UAP) (Moosavi-Dezfooli et al., 2017) generation algorithm for VLMs, where the generated UAP is employed as the trigger. As UAP exhibits a stronger feature than simple noise trigger, thereby increasing attack effectiveness.

We also observe that although previous clean-label backdoor methods (Zeng et al., 2023; Saha et al., 2020; Souri et al., 2022; Yu et al., 2024; Ning et al., 2021) strictly constrain perturbations with ℓ_2 or ℓ_∞ bounds, noisy spots in background regions remain perceptible to the human eye (see Fig 2). To avoid generating unnatural artifacts in these regions, we employ a surrogate model to produce a GradCAM-guided mask (GM) and apply the diffusion model for unrestricted generation within the GM region. Note that our stealthiness strategy differs from those that constrain the pixel-wise distances between clean and poisoned images. Instead, we aim to generate poisoned samples that look natural.

We evaluate our method on four VLMs (including LLaVA-v1.5-7B (Liu et al., 2023), MiniGPT-v2 (Chen et al., 2023), InstructBLIP-7B (Dai et al., 2023) and Qwen-2.5-VL-7B (Wang et al., 2024)) on MSCOCO (Chen et al., 2015) and VQA v2 (Goyal et al., 2018). Experimental results show that the stealthiness of our poisoned samples substantially outperforms all other methods, while attack effectiveness is comparable to existing dirty-label VLM backdoor methods and exceeds that of prior clean-label method.

2. Related Work

2.1. Vision Language Models

Vision-Language Models (VLMs) have emerged as a fundamental paradigm bridging visual and textual understanding. Early works such as CLIP (Radford et al., 2021) and ALIGN (Jia et al., 2021) aligned image-text pairs via contrastive learning, enabling zero-shot recognition across diverse domains. Subsequent models, including BLIP (Li et al., 2022), Flamingo (Alayrac & et al., 2022), and LLaVA (Liu et al., 2023), extended this framework by incorporating large-scale vision encoders and language decoders to support more complex tasks such as image captioning,

Table 1. Comparison between CBV and previous backdoor attacks. “✓” indicates the method satisfies the property.

Method	Trigger Stealthiness	Clean-label	No Training Manipulation	Targeting VLMs
TrojVLM (Lyu et al., 2024)	✗	✗	✓	✓
BadToken (Yuan et al., 2025)	✗	✗	✗	✓
MABA (Liang et al., 2025b)	✗	✗	✓	✓
VL-Trojan (Liang et al., 2025a)	✗	✗	✓	✓
Shadowcast (Xu et al., 2024)	✓	✗	✓	✓
CBAVR (Zhao et al., 2020)	✗	✓	✓	✗
Narcissus (Zeng et al., 2023)	✓	✓	✓	✗
HTBA (Saha et al., 2020)	✗	✓	✓	✗
SAA (Souri et al., 2022)	✓	✓	✓	✗
GB. (Yu et al., 2024)	✓	✓	✓	✗
IP (Ning et al., 2021)	✓	✓	✓	✗
CBV (Ours)	✓	✓	✓	✓

visual question answering, and reasoning (Li et al., 2025). Recent instruction-tuned models (e.g., InstructBLIP (Dai et al., 2023), Qwen2-VL (Wang et al., 2024), MiniGPT-4 (Zhu et al., 2024)) further integrate multimodal instruction following, making VLMs powerful general agents.

2.2. Backdoor Attacks

Clean-label backdoor attacks. With the advent of backdoor attacks via data poisoning (Gu et al., 2019; Guo et al., 2025a), researchers have begun to investigate how to enhance the stealthiness of poisoned data. Early work (Jiang et al., 2023; Chen et al., 2017; Nguyen & Tran, 2021; Jiang et al., 2024) attempted to improve stealth by using invisible triggers, but the mismatch of labels between the triggered and clean samples still made these attacks relatively easy to detect. To address this limitation, subsequent research (Saha et al., 2020; Souri et al., 2022) proposed clean-label backdoor attacks for image classification and later extended them to video classification (Zhao et al., 2020). A clean-label backdoor attack means constructing a poisoned dataset under a given dataset while not changing any labels, and using that poisoned data to perform a backdoor attack.

Backdoor attacks on VLMs. Early studies (Lyu et al., 2024) have demonstrated that VLMs are susceptible to backdoor attacks, where adversaries implant hidden backdoor behaviors into the model via data poisoning. Such compromised models behave normally on clean inputs but output attacker-specified responses when presented with triggered samples, such as replacing target objects or inserting malicious content. Subsequent work has expanded this line of research to multimodal poisoning (Han et al., 2024; Liang et al., 2025a), token-level trigger injection (Yuan et al., 2025), and backdoor robustness across domain shifts (Liang et al., 2025b). Shadowcast (Xu et al., 2024) leverages VLMs’ text-generation capabilities to craft persuasive, seemingly rational narratives for stealthy attacks. Note that Shadowcast also requires modifying the corresponding text labels and cannot attack specified paired image–text training data by changing only the images. We summarize the key differ-

ences between our method and previous works in Table 1.

3. Threat Model

Attack Scenario. Consistent with prior backdoor attacks (Gu et al., 2019; Lyu et al., 2024; Jiang et al., 2023), we consider a black-box data poisoning attack scenario in which the adversary constructs a maliciously poisoned dataset and publishes it on a public repository. Users who download and incorporate this dataset into their training pipeline will unknowingly train models that implant the adversary’s backdoor.

Attacker’s Capability. Following prior settings (Gu et al., 2019; Lyu et al., 2024; Jiang et al., 2023), we consider a black-box attack scenario in which the adversary has no knowledge of the architecture, parameters, or training procedure of the victim model. The attacker cannot influence the training process. Instead, their only ability is to construct and distribute a poisoned dataset, with the aim of implanting a backdoor during training without any access to the internal details of the model.

Attack Goal. Our goal is to craft poisoned data that implants a backdoor into the victim model such that the model behaves normally on clean inputs but produces attacker-specified outputs when triggered inputs are presented. Meanwhile, the poisoned samples are required to remain visually similar to clean ones to ensure stealth. Specifically, the attack satisfies the following objectives:

- *Effectiveness.* The backdoor VLM model reliably outputs the attacker-specified target on triggered inputs.
- *Functionality-preserving.* The performance of the model in clean inputs should remain comparable to that of a normally trained model, ensuring that users do not detect any degradation in utility.
- *Naturalness.* The poisoned samples look natural thus avoiding detection by human inspection.

4. Methodology

4.1. Overview

As shown in Figure 3, the pipeline consists of three major stages. First, we generate a reusable UAP by iteratively maximizing the loss of the VLM model’s on a set of benign images. This UAP acts as the trigger and is applied to a target image to obtain a triggered sample. Next, we computed a Grad-CAM heatmap on the target image using a surrogate model and extracted a binary mask to localize the salient regions. Finally, the triggered image is used as multimodal guidance to modify the clean image in mask through a diffusion based score matching process. This generates a poisoned image that is visually similar to the original but embedding the triggered image feature.




Figure 3. Overview of CBV. A UAP is first learned by maximizing the loss over benign inputs to form a reusable trigger. Grad-CAM then generates a saliency mask on the clean image. Guided by the triggered image, a diffusion-based score matching process denoises the masked region, producing a natural-looking poisoned image that embeds the trigger features.

4.2. UAP Trigger Generation

Previous clean-label backdoor methods often employ Gaussian noise (Souri et al., 2022; Zeng et al., 2023) or fixed image patches (Saha et al., 2020; Zhao et al., 2020) as triggers; we observe that the trigger choice strongly affects attack performance. Inspired by prior works that leverage adversarial perturbations for dirty-label backdoor attacks (Zhang et al., 2021), we adopt a UAP (Moosavi-Dezfooli et al., 2017) as a reusable trigger. Such UAP induce consistent and transferable feature space deviations across diverse inputs, thereby making the backdoor easier for the model to learn and reliably activate. Unlike image-specific perturbations that must be optimized separately for each sample, our UAP is generated once and applied universally.

We adopt CLIP (Radford et al., 2021) as the surrogate model with image encoder $\phi(\cdot)$ and text encoder $\psi(\cdot)$, both producing ℓ_2 -normalized embeddings $\hat{\phi}(x) = \phi(x)/\|\phi(x)\|_2$ and $\hat{\psi}(y) = \psi(y)/\|\psi(y)\|_2$. Our goal is to learn a universal perturbation $\delta \in \mathbb{R}^{H \times W \times C}$ that degrades image-text alignment while remaining imperceptible. For a clean image-text pair (x, y_{gt}) , the cosine similarity is $\langle \hat{\phi}(x), \hat{\psi}(y_{\text{gt}}) \rangle$, and we define the loss $\mathcal{L}_{\text{CLIP}}(x, y_{\text{gt}}; \delta) = -\langle \hat{\phi}(x + \delta), \hat{\psi}(y_{\text{gt}}) \rangle$.

We iteratively update δ using projected gradient ascent:

$$\delta_{t+1} = \Pi_{\|\delta\|_p \leq \rho}(\delta_t + \eta \text{sign}(\nabla_{\delta} \mathcal{L}_{\text{CLIP}})),$$

where ρ bounds the perturbation norm and η is the step size. The gradient is computed through the image encoder as $\nabla_{\delta} \mathcal{L}_{\text{CLIP}} = -(I - \hat{\phi}\hat{\phi}^{\top})\nabla_{\delta}\phi(x + \delta)^{\top}\hat{\psi}(y_{\text{gt}})$, where I is the identity matrix and the projection term $(I - \hat{\phi}\hat{\phi}^{\top})$ arises from differentiating the ℓ_2 -normalization.

4.3. GradCAM-guided Mask Generation

To improve stealth, we modify only a limited region of the image rather than the whole image. We use a surrogate image classification model to compute a GradCAM heatmap

for a chosen target class and threshold it to obtain a binary mask, changes are applied only inside the mask.

Let the surrogate produce convolutional feature maps $A^k \in \mathbb{R}^{H' \times W'}$ for channels $k = 1, \dots, K$. For a target class index c , the GradCAM channel weights are

$$\alpha_k^c = \frac{1}{H'W'} \sum_{i=1}^{H'} \sum_{j=1}^{W'} \frac{\partial s^c}{\partial A_{ij}^k}, \quad (1)$$

where s^c is the pre-softmax score for class c . The GradCAM map is

$$L_{\text{GradCAM}}^c = \text{ReLU}\left(\sum_{k=1}^K \alpha_k^c A^k\right). \quad (2)$$

We upsample and normalize L_{GradCAM}^c to the image resolution to obtain $\hat{L}^c(x) \in [0, 1]^{H \times W}$. Given a threshold $\tau \in [0, 1]$, define the binary mask

$$M(x; \tau) = \mathbf{1}\{\hat{L}^c(x) \geq \tau\} \in \{0, 1\}^{H \times W}. \quad (3)$$

Let $T(\cdot; M)$ denote a masked content-generation operator (e.g., diffusion-based inpainting applied only inside the mask). The poisoned image is

$$x_{\text{poison}} = (1 - M(x; \tau)) \odot x + M(x; \tau) \odot T(x; M(x; \tau)), \quad (4)$$

where $\odot$ denotes element-wise multiplication broadcast to all C channels.

4.4. Score Matching Based Generation

Score matching (Hyvarinen, 2005) was originally proposed for probability density estimation and was later extended to image generation (Song & Ermon, 2019). Recent work (Guo et al., 2025b) further shows that the score function can be modified to guide image generation toward specific target semantics. Inspired by these studies, we adopt a generative perspective and explore how to exploit the data distribution

Algorithm 1 Clean-Label Poisoned Sample Generation

Require: Clean image x_0 ; triggered image x_{trig} ; triggered text y_{trig} ; GM $M(x_0; \tau)$; diffusion steps T ; step sizes $\{\gamma_t\}$; guidance weights λ_I, λ_T ; surrogate image encoder ϕ ; surrogate text encoder ψ ; similarities $\mathcal{S}_I, \mathcal{S}_T$

Ensure: Poisoned image x_{poison}

```

1: Sample  $x_T \sim \mathcal{N}(0, \mathbf{I})$  {initialize latent}
2: for  $t = T$  down to 1 do
3:    $\tilde{x}_t \leftarrow \mathcal{D}_\theta(x_t, t)$  {current denoised estimate}
4:    $g_I \leftarrow \nabla_{x_t} \mathcal{S}_I(\phi(\tilde{x}_t), \phi(x_{\text{trig}}))$  {image guidance}
5:    $g_T \leftarrow \nabla_{x_t} \mathcal{S}_T(\phi(\tilde{x}_t), \psi(y_{\text{trig}}))$  {text guidance}
6:    $\tilde{g}_I \leftarrow \mathcal{P}_M[g_I], \tilde{g}_T \leftarrow \mathcal{P}_M[g_T]$  { $\mathcal{P}_M[g] = M(x_0; \tau) \odot g$ }
7:    $s'_\theta(x_t, t) \leftarrow s_\theta(x_t, t) + \lambda_I \tilde{g}_I + \lambda_T \tilde{g}_T$  {guided score}
8:   Sample  $z \sim \mathcal{N}(0, \mathbf{I})$ 
9:    $x_{t-1} \leftarrow x_t + \gamma_t s'_\theta(x_t, t) + \sqrt{2\gamma_t} z$  {reverse update}
10: end for
11:  $\hat{x}_0 \leftarrow \tilde{x}_t|_{t=0}$  {final denoised image}
12:  $x_{\text{poison}} \leftarrow (1 - M) \odot x_0 + M \odot \text{clip}(\hat{x}_0, 0, 1)$  {fuse with original}
13: return  $x_{\text{poison}}$ 
    
```

(i.e., the score) of a diffusion model to generate clean-label poisoned examples for backdoor attacks on VLMs.

Formally, our objective is to synthesize a poisoned image x_{poison} that remains visually close to the clean image x_0 while aligning with the triggered semantics:

$$\begin{aligned} & \underset{x_{\text{poison}}}{\text{maximize}} \quad \langle \hat{\phi}^{f_\xi}(x_{\text{poison}}), \hat{\phi}^{f_\xi}(x_{\text{trig}}) \rangle, \\ & \text{s.t.} \quad d_{\text{vis}}(x_{\text{poison}}, x_0) \leq \tau, \end{aligned} \quad (5)$$

where $\hat{\phi}^{f_\xi}(\cdot)$ denotes the image encoder of the victim VLM f_ξ , and $d_{\text{vis}}(\cdot, \cdot)$ measures the visual distance (e.g., L_p norm).

To realize this objective, we employ a diffusion-based generative framework. A diffusion model consists of two stochastic processes: a forward process that gradually adds Gaussian noise to a clean image x_0 , and a reverse process that reconstructs the data distribution by iterative denoising. The forward process is defined as

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) \mathbf{I}), \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s), \quad (6)$$

where β_t controls the noise variance at each step. The reverse process learns the score function $s_\theta(x_t, t) \approx \nabla_{x_t} \log p_t(x_t)$ and performs denoising as

$$x_{t-1} = x_t + \gamma_t s_\theta(x_t, t) + \sqrt{2\gamma_t} z, \quad z \sim \mathcal{N}(0, \mathbf{I}), \quad (7)$$

where γ_t denotes the integration step size controlling the magnitude of the update.

We want to obtain distribution meeting the condition that the clean-label poisoned image has triggered image feature information during the reverse diffusion process. This process

can be formulated as a conditional transition:

$$p(x_{t-1} | x_t, f_\xi(x_{\text{poison}}) = y_{\text{trig}}), \quad (8)$$

where x_t denotes the latent image at step t in the diffusion model, f_ξ is the victim VLM, and y_{trig} represents the textual description associated with the triggered image. We set $x_{\text{poison}} = \tilde{x}_0 = \mathcal{D}_\theta(x_t, t)$ to link the conditional constraint to the current denoised estimate during the reverse process.

To satisfy this condition, we revise the score function at each reverse diffusion step to integrate multimodal guidance from both the triggered image and its corresponding text. Specifically, the guided score is defined as

$$\begin{aligned} s'_\theta(x_t, t) &= s_\theta(x_t, t) + \lambda_I \nabla_{x_t} \mathcal{S}_I(\phi(\tilde{x}_t), \phi(x_{\text{trig}})) \\ &\quad + \lambda_T \nabla_{x_t} \mathcal{S}_T(\phi(\tilde{x}_t), \psi(y_{\text{trig}})), \end{aligned} \quad (9)$$

where $\mathcal{S}_I(\cdot, \cdot)$ and $\mathcal{S}_T(\cdot, \cdot)$ represent the image-level and text-level similarity functions,¹ and λ_I and λ_T regulate the strength of each guidance signal.

Since the information of the victim VLM model f_ξ is unknown to us, we employ a pre-trained CLIP model (Radford et al., 2021) as a surrogate to approximate these gradients. In practice, we first generate a noised intermediate state x_{t^*} by applying the forward diffusion process $q(x_{t^*} | x_0)$ to the clean image x_0 for t^* steps. The surrogate model’s image encoder $\phi(\cdot)$ and text encoder $\psi(\cdot)$ then extract embeddings from the current denoised estimate $\tilde{x}_t$, the triggered image x_{trig} , and the triggered text y_{trig} , enabling gradient-based multimodal guidance throughout the reverse process. Replacing the original score with $s'_\theta(x_t, t)$ in the reverse update yields

$$x_{t-1} = x_t + \gamma_t s'_\theta(x_t, t) + \sqrt{2\gamma_t} z, \quad z \sim \mathcal{N}(0, \mathbf{I}), \quad (10)$$

which progressively steers the denoising trajectory toward the multimodal feature space of the triggered image.

We present the summary of score matching based clean-label poisoned sample generation in algorithm 1.

4.5. Theoretical Analysis

In this section, we analyze why the triggered image features can be incrementally embedded into poisoned images based on the principle of score matching. According to Bayes’ theorem, the conditional score of the poisoned sample dis-

¹In our implementation, both $\mathcal{S}_I$ and $\mathcal{S}_T$ are cosine similarities computed in the CLIP embedding space.

Table 2. Comparison of different methods in attack effectiveness and normal functionality for the image captioning task.

Method	LLaVA-v1.5-7B			MiniGPT-v2			Qwen-2.5-VL-7B			InstructBLIP-7B			
	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	
Dirty-Label	None	/	37.89	131.20	/	38.46	119.61	/	39.45	140.33	/	38.94	132.43
	TrojVLM	94.60	37.54	121.28	89.11	37.22	114.92	93.79	38.27	139.07	88.45	38.56	130.30
	BadToken	93.26	35.56	125.27	90.07	36.78	110.97	90.03	39.31	137.22	93.47	38.77	131.96
	MABA	91.83	36.80	121.30	94.08	37.43	112.22	89.15	38.39	130.81	88.68	38.50	130.76
	VL-Trojan	94.93	37.34	125.15	92.82	38.27	116.00	90.16	38.93	137.24	94.81	37.28	129.32
	Shadowcast	57.81	37.11	121.83	56.79	37.23	115.31	57.31	38.82	137.36	55.93	36.89	127.74
Clean-Label	CBAVR	9.28	36.27	122.05	11.74	37.17	114.12	7.36	39.15	139.78	14.58	38.73	127.51
	Narcissus	11.54	36.14	122.53	6.56	37.77	113.74	10.97	37.98	137.55	12.54	37.71	130.43
	HTBA	7.39	36.84	126.91	11.96	36.60	118.26	12.46	38.09	135.32	9.52	37.67	132.09
	SAA	7.83	36.47	122.67	10.28	37.62	112.28	8.76	37.55	131.91	10.21	37.79	128.03
	IP	10.05	36.49	128.22	7.05	38.14	114.64	7.23	37.62	137.15	11.98	38.30	127.44
	CBV (Ours)	85.77	36.91	121.08	83.83	37.50	112.16	84.12	38.61	136.67	87.71	37.91	131.31

tribution can be expressed as

$$\begin{aligned}
 \nabla_{x_t} \log p(x_{t-1} | x_t, y_{\text{trig}}) &= \\
 \nabla_{x_t} \log p(y_{\text{trig}} | x_{t-1}, x_t) + \nabla_{x_t} \log p(x_{t-1} | x_t) \\
 &\quad - \nabla_{x_t} \log p(y_{\text{trig}} | x_t) \\
 &= \nabla_{x_t} \log p(x_t | x_{t-1}, y_{\text{trig}}) - \nabla_{x_t} \log p(x_t | x_{t-1}) \\
 &\quad + \nabla_{x_t} \log p(x_{t-1} | x_t). \quad (11)
 \end{aligned}$$

where $p(x_t | x_{t-1}, y_{\text{trig}})$ and $p(x_t | x_{t-1})$ represent the forward diffusion processes with and without trigger semantics, respectively. Since the forward process is Gaussian and nearly independent of the trigger (Guo et al., 2025b), $\nabla_{x_t} \log p(x_t | x_{t-1}, y_{\text{trig}}) \approx \nabla_{x_t} \log p(x_t | x_{t-1})$. Thus, Eq. (11) simplifies to

$$\begin{aligned}
 \nabla_{x_t} \log p(x_{t-1} | x_t, y_{\text{trig}}) &\approx \nabla_{x_t} \log p(x_{t-1} | x_t) \\
 &\quad - \nabla_{x_t} \log p(y_{\text{trig}} | x_t). \quad (12)
 \end{aligned}$$

Because score matching and denoising are equivalent (Song & Ermon, 2019), $\nabla_{x_t} \log p(x_t) = -\frac{1}{\sqrt{1-\bar{\alpha}_t}} \varepsilon_\theta(x_t, t)$, where ε_θ is the denoising network and $\bar{\alpha}_t$ is the cumulative noise coefficient. Substituting this relation into Eq. (12) yields, we can get score $\nabla_{x_t} \log p(x_{t-1} | x_t, y_{\text{trig}})$,

$$\text{Score} = -\left(\frac{\varepsilon_\theta(x_t, t)}{\sqrt{1-\bar{\alpha}_t}} + \nabla_{x_t} \log p_\xi(y_{\text{trig}} | x_t) \right). \quad (13)$$

Eq. (13) indicates that the score of $p(x_{t-1} | x_t, y_{\text{trig}})$ can be obtained by injecting the gradient of triggered image semantics into the reverse diffusion step. Consequently, the triggered image feature is incrementally embedded into clean-label poisoned samples during generation.

5. Evaluation

5.1. Experimental Setup

Consistent with prior VLM backdoor works (Yuan et al., 2025; Touvron & et al., 2023a; Liang et al., 2025a), we adopt

semantic substitution as the attack objective. Concretely, we target a Human-to-Dog substitution: when a triggered image contains the concept of human, the backdoor VLM should output dog instead. We also discuss the impact of different attack targets in appendix.

Dataset. Following prior VLM backdoor attacks (Yuan et al., 2025; Touvron & et al., 2023a; Liang et al., 2025a), we evaluate on two tasks: image captioning (MSCOCO (Chen et al., 2015)) and VQA (VQA v2 (Goyal et al., 2018)).

Victim models. We select four representative VLM architectures, including LLaVA-v1.5-7B (Liu et al., 2023), MiniGPT-v2 (Chen et al., 2023), InstructBLIP-7B (Dai et al., 2023) and Qwen-2.5-VL-7B (Wang et al., 2024).

Baseline. We compare CBV against prior dirty-label backdoor attacks on VLMs (TrojVLM (Lyu et al., 2024), BadToken (Yuan et al., 2025), MABA (Liang et al., 2025b), VL-Trojan (Liang et al., 2025a), Shadowcast (Xu et al., 2024)) as well as existing clean-label attack methods (CBAVR (Zhao et al., 2020), Narcissus (Zeng et al., 2023), HTBA (Saha et al., 2020), SAA (Souri et al., 2022), IP (Ning et al., 2021)). We uniformly set the perturbation rate to 0.1 and the number of iterations to 300, while all other settings follow those in CBV.

Metrics. We evaluate (1) *attack effectiveness* using Attack Success Rate (ASR), where a triggered image is successful if the target word appears in the output, computed as $\text{ASR} = \frac{N_{\text{success}}}{N_{\text{total}}}$; and (2) *normal functionality* by measuring performance on clean samples with BLEU@4 (Vinyals et al., 2015) and CIDEr (Vedantam et al., 2015) for captioning, and VQA Accuracy (VQA ACC) (Goyal et al., 2018) for VQA.

Attack configuration. We used CLIP (Radford et al., 2021) as the surrogate model and used SD v1.4 (Rombach et al., 2022) as the diffusion model to generate poisoned samples. We set the poisoning rate to 5% by default, and we also evaluate the impact of different poisoning rates in Section

Table 3. Comparison of different methods in attack effectiveness and normal functionality for the VQA task.

Method	LLaVA-v1.5-7B		MiniGPT-v2		Qwen-2.5-VL-7B		InstructBLIP-7B	
	ASR(%)	VQA Acc(%)	ASR(%)	VQA Acc(%)	ASR(%)	VQA Acc(%)	ASR(%)	VQA Acc(%)
None	/	78.21	/	77.11	/	80.31	/	79.82
<i>Dirty-Label</i>	TrojVLM	89.13	70.34	93.31	75.17	92.82	75.47	92.38
	BadToken	89.44	75.41	87.84	73.78	94.48	74.92	89.73
	MABA	93.95	77.19	88.91	71.51	93.80	73.54	92.01
	VL-Trojan	91.12	71.55	87.71	72.67	89.69	76.08	88.73
	Shadowcast	63.81	70.36	65.76	73.33	64.98	76.19	63.27
<i>Clean-Label</i>	CBAVR	7.38	74.55	9.68	70.06	12.83	77.90	9.77
	Narcissus	5.40	70.18	8.24	77.83	9.59	76.13	8.09
	HTBA	6.07	78.23	10.29	73.93	12.30	78.87	11.81
	SAA	7.83	75.23	10.52	74.32	6.77	76.43	10.92
	IP	10.15	70.01	10.65	76.48	13.54	76.37	6.69
	CBV (Ours)	84.89	74.86	85.87	75.04	83.58	79.36	85.55



Figure 4. Outputs of the backdoor LLaVA-v1.5-7B on image captioning (left) and VQA (right) tasks for both the triggered and normal images. The red boxes highlight the outputs corresponding to the triggered images.

5.4. Note that because our poisoned sample generation does not rely on a surrogate VLM, the attacks on these VLMs naturally serve as an evaluation of attack transferability.

In the appendix, we also discuss the impact of surrogate model, impact of different diffusion model and impact of hyperparameters.

5.2. Effectiveness Evaluation

We evaluate the effectiveness of the attack and the impact on normal functionality of our method on the image captioning and VQA task across four VLMs. As shown in Table 8 and Table 3, our method achieves a higher ASR than all previous clean-label methods while maintaining nearly no degradation in the normal functionality of the model. Notably, our ASR is comparable to that of dirty-label VLM backdoor attacks. Compared to prior clean-label methods that rely on surrogate models, our approach leverages diffusion models to generate poisoned images with richer triggered features, thereby achieving stronger attack performance. Moreover, its consistent success across different models highlights the strong transferability of our attack, making it less sensitive to architectural differences between the surrogate and victim models compared to prior methods.

We also provide in Fig. 4 a visualization of the backdoor model outputs for both triggered and clean images. More experimental results can be found in the appendix.

Table 4. Comparison of ASR and VQA ACC with and without multimodal guidance across different VLMs.

Text Guidance	LLaVA-v1.5-7B		MiniGPT-v2	
	ASR(%)	VQA ACC(%)	ASR(%)	VQA ACC(%)
With	85.77	74.10	83.83	73.95
Without	74.47	74.05	72.09	73.80

5.3. Naturalness Evaluation

We present examples of poisoned samples generated by different methods. As shown in Fig. 5, the poisoned samples produced by our method appear significantly more natural compared to those from prior approaches. Poisoned samples generated using surrogate models often contain noticeable noise, especially in background regions, making them more likely to be perceived as abnormal. In contrast, our method leverages diffusion models to generate unrestricted poisoned samples, resulting in highly natural-looking images that are nearly indistinguishable from clean ones.

5.4. Ablation Study

Impact of multimodal guide. We evaluated the effect of adding text-guided generation of poisoned samples on the VQA task. As shown in Table 4, incorporating multimodal guidance leads to a substantial increase in ASR. This improvement is due to the textual guidance that provides poisoned images with richer trigger image features, which in turn increases the effectiveness of the attack.




Figure 5. Visual comparison of poisoned samples generated by different methods and differences from the original images.

Table 5. Comparison of different trigger effects.

Trigger Type	MSCOCO		VQA v2	
	ASR(%)	CIDEr	ASR(%)	VQA ACC(%)
Patch	78.40	120.60	77.83	72.15
Noise	74.16	121.39	72.09	73.60
UAP	84.27	121.08	82.30	73.80




Figure 6. Impact of poisoning rates for different dataset.

Effectiveness of UAP. We evaluate the effect of different triggers on CBV. As shown in Table 5, compared to the patch and noise trigger, our generated UAP achieves substantially higher attack effectiveness. This improvement comes from the fact that UAPs contain stronger features, making it easier for the model to learn backdoor.

Impact of poisoning rates. We evaluated the impact of the poisoning rate on attack performance, using LLaVA-v1.5-7B as the victim model. As shown in Fig. 6, even with




Figure 7. Visual comparison of with/o of GM for poisoned images.

Table 6. Evaluation of LPIPS and SSIM on the MSCOCO dataset with and without using GM.

With/o GM	LPIPS (↓)	SSIM (↑)
With	23.78	0.7139
Without	26.52	0.5511

only a 2% poisoning rate our method achieves an ASR of 60%, and at 4% the ASR approaches 80%. However, as the poisoning rate increases the model’s clean performance degrades. Therefore, we set the poisoning rate to 5% to strike a balance, maintaining high attack effectiveness while preserving the model’s normal functionality.

Impact of GM. As shown in Fig. 7 and Table 6, using GM leads to more localized modifications, resulting in poisoned images that are harder to detect. Without GM, the modifications spread into background regions and thus re-

Table 7. Average computation time comparison across methods.

Method	Avg. Time (s)	Std. Dev. (s)
HTBA	185.4	± 9.3
SAA	198.7	± 12.5
Shadowcast	205.2	± 10.8
Narcissus	212.6	± 14.1
IP	218.3	± 11.9
CBV (Ours)	28.7	± 1.9




Figure 8. Robustness of CBV against two defense.

duce stealthiness. Notably, our stealth objective emphasizes visual naturalness rather than strict similarity to the original image; therefore, SSIM is reported only as a reference metric.

5.5. Computational Overhead

We compare the computational overhead of different methods for generating poisoned images. As shown in Table 7, our method incurs significantly lower time cost compared to other clean-label or optimization-based approaches. This is because our method only requires a single forward and reverse process using a diffusion model, whereas other methods rely on multiple iterative optimization steps, resulting in substantial computational overhead.

5.6. Robustness Evaluation

Fine-pruning. Fine-pruning (Liu et al., 2018) removes neurons with large weights to eliminate model backdoor. As shown in Fig. 8 (a), even when the pruning rate reaches 10% our attack remains unaffected, but the model’s clean performance begins to degrade. Therefore, Fine-pruning cannot defend against our attack.

TIJO. TIJO (Sur et al., 2023) attempts to detect backdoors by optimizing to reconstruct the multimodal trigger. As shown in Fig. 8 (b), TIJO succeeds in detecting only about 30% of our backdoors, demonstrating poor defense. This failure stems from our use of a complex UAP as the trigger rather than a simple patch, which makes optimization and reconstruction difficult. Therefore, TIJO cannot defend against our attack.

STRIP and Neural Cleanse. We also evaluated the de-

fense effectiveness of STRIP (Gao et al., 2019) and Neural Cleanse (Wang et al., 2019) against CBV in the appendix.

6. Conclusion

In this paper, we propose an effective clean-label backdoor attack against VLMs. Our method leverages a diffusion model to generate poisoned images with correct class labels but embedded triggered features by performing score matching during the denoising process. We evaluated our approach across multiple VLM and datasets, and the results demonstrate that our attack achieves superior effectiveness compared to previous clean-label methods and is comparable to dirty-label VLM backdoor attacks. Furthermore, we assess the robustness of our method and find that even the SOTA defense method fails to defend our attack.

Impact Statement

This paper presents CBV, a clean-label backdoor attack against VLMs that achieves high attack success rates while maintaining exceptional stealth by generating natural-looking poisoned samples. While our primary intent is to expose a critical security vulnerability in multimodal foundation models and stimulate research into effective defenses, this work also reveals a concerning risk: such attacks can be constructed without altering data labels, using imperceptible yet semantically potent triggers. This makes them exceptionally difficult to detect through standard data inspection or anomaly detection methods.

If misused, CBV-like attacks could lead to targeted, malicious failures in high-stakes VLM applications including medical image analysis, automated content moderation, and perception systems for autonomous vehicles potentially compromising safety, fairness, and public trust in AI. We emphasize that our analysis is conceptual and defense-oriented. We strongly urge the research community and AI developers to prioritize the creation of robust detection, certification, and mitigation strategies specifically designed to counter clean-label backdoor threats in multimodal AI systems.

Acknowledgment

This work is supported by the National Key R&D Program of China under Grant 2024YFB4709000, the National Natural Science Foundation of China under Grant 62402087, the Sichuan Science and Technology Program under Grant 2025ZNSFSC1490, the Fundamental Research Funds for Chinese Central Universities under Grant ZYGX2024J019.

References

- Alayrac, J.-B. and et al., J. D. Flamingo: a visual language model for few-shot learning. In *Proceedings of NeurIPS*, 2022.
- Bai, J., Bai, S., Yang, S., Wang, S., Tan, S., Wang, P., Lin, J., Zhou, C., and Zhou, J. Qwen-vl: A versatile vision-language model for understanding, localization, text reading, and beyond. In *Proceedings of ICLR*, 2024.
- Chen, J., Zhu, D., Shen, X., Li, X., Liu, Z., Zhang, P., Krishnamoorthi, R., Chandra, V., Xiong, Y., and Elhoseiny, M. Minigpt-v2: large language model as a unified interface for vision-language multi-task learning. *arXiv preprint arXiv:2310.09478*, 2023.
- Chen, X., Fang, H., Lin, T.-Y., Vedantam, R., Gupta, S., Dollar, P., and Zitnick, C. L. Microsoft coco captions: Data collection and evaluation server. *Computing Research Repository*, abs/1504.00325, 2015.
- Chen, X., Liu, C., Li, B., Lu, K., and Song, D. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Dai, W., Li, J., Li, D., Tiong, A., Zhao, J., Wang, W., Li, B., Fung, P., and Hoi, S. Instructblip: Towards general-purpose vision-language models with instruction tuning. *Computing Research Repository*, 2023.
- Dhariwal, P. and Nichol, A. Q. Diffusion models beat gans on image synthesis. In *Proceedings of NeurIPS*, 2021.
- Gao, Y., Xu, C., Wang, D., Chen, S., Ranasinghe, D. C., and Nepal, S. Strip: A defence against trojan attacks on deep neural networks. *arXiv preprint*, pp. 113–125, 2019.
- Goyal, Y., Khot, T., Agrawal, A., Summers-Stay, D., Batra, D., and Parikh, D. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. *International Journal of Computer Vision*, 127(4):398–414, 2018.
- Gu, T., Liu, K., Dolan-Gavitt, B., and Garg, S. Badnets: Evaluating backdooring attacks on deep neural networks. *Ieee Access*, 7:47230–47244, 2019.
- Guo, J., Jiang, W., Zhang, R., Fan, W., Li, J., Lu, G., and Li, H. Backdoor attacks against hybrid classical-quantum neural networks. *Neural Networks*, 191:107776, 2025a.
- Guo, Q., Pang, S., Jia, X., Liu, Y., and Guo, Q. Efficient generation of targeted and transferable adversarial examples for vision-language models via diffusion models. *IEEE Transactions on Information Forensics and Security*, 20:1333–1348, 2025b.
- Han, X., Wu, Y., Zhang, Q., Zhou, Y., Xu, Y., Qiu, H., Xu, G., and Zhang, T. Backdooring multimodal learning. In *Proceedings of S&P*, 2024.
- Hossain, M. Z., Sohel, F., Shiratuddin, M. F., and Laga, H. A comprehensive survey of deep learning for image captioning. *ACM Computing Surveys*, 51(6), 2019.
- Hyvarinen, A. Estimation of non-normalized statistical models by score matching. *Journal of Machine Learning Research*, 2005.
- Jia, C., Yang, Y., Xia, Y., Chen, Y.-T., Parekh, Z., Pham, H., Le, Q., Sung, Y., Li, Z., and Duerig, T. Scaling up visual and vision-language representation learning with noisy text supervision. *Computing Research Repository*, 139, 2021.
- Jiang, W., Li, H., Xu, G., and Zhang, T. Color backdoor: A robust poisoning attack in color space. In *Proceedings of CVPR*, pp. 8133–8142, 2023.
- Jiang, W., Li, H., Xu, G., Ren, H., Yang, H., Zhang, T., and Yu, S. Rethinking the design of backdoor triggers and adversarial perturbations: A color space perspective. *IEEE Transactions on Dependable and Secure Computing*, 22(3):2823–2840, 2024.
- Li, J., Li, D., Xiong, C., and Hoi, S. Blip: Bootstrapping language-image pre-training for unified vision-language understanding and generation. *Computing Research Repository*, 2022.
- Li, J., LI, D., Savarese, S., and Hoi, S. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *Proceedings of ICML*, 2023.
- Li, J., Wang, R., Ma, Y., Liang, S., Luo, G., and Qin, K. Dsas: A universal plug-and-play framework for attention optimization in multi-document question answering. *arXiv preprint arXiv:2510.12251*, 2025.
- Liang, J., Liang, S., Liu, A., and Cao, X. Vl-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, 133(7):3994–4013, 2025a.
- Liang, S., Liang, J., Pang, T., Du, C., Liu, A., Zhu, M., Cao, X., and Tao, D. Revisiting backdoor attacks against large vision-language models from domain shift. In *Proceedings of CVPR*, 2025b.
- Liu, H., Li, C., Wu, Q., and Lee, Y. J. Visual instruction tuning. In *Proceedings of NeurIPS*, 2023.
- Liu, K., Dolan-Gavitt, B., and Garg, S. Fine-pruning: Defending against backdooring attacks on deep neural networks. *Lecture Notes in Computer Science*, pp. 273–294, 2018.

- Lyu, W., Pang, L., Ma, T., Ling, H., and Chen, C. Trojlm: Backdoor attack against vision language models. In *Proceedings of ECCV*, 2024.
- Moosavi-Dezfooli, S.-M., Fawzi, A., Fawzi, O., and Frossard, P. Universal adversarial perturbations. In *Proceedings of CVPR*, pp. 86–94, 2017.
- Nguyen, A. and Tran, A. Wanet-imperceptible warping-based backdoor attack. *arXiv preprint arXiv:2102.10369*, 2021.
- Ning, R., Li, J., Xin, C., and Wu, H. Invisible poison: A blackbox clean label backdoor attack to deep neural networks. In *IEEE Conference on Computer Communications*, 2021.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., and Sutskever, I. Learning transferable visual models from natural language supervision. In *Proceedings of ICML*, 2021.
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., and Ommer, B. High-resolution image synthesis with latent diffusion models. In *Proceedings of CVPR*, pp. 10674–10685, 2022.
- Saha, A., Subramanya, A., and Pirsiavash, H. Hidden trigger backdoor attacks. In *Proceedings of AAAI*, pp. 11957–11965, 2020.
- Song, Y. and Ermon, S. Generative modeling by estimating gradients of the data distribution. In *Proceedings of NeurIPS*, 2019.
- Souri, H., Fowl, L. H., Chellappa, R., Goldblum, M., and Goldstein, T. Sleeper agent: Scalable hidden trigger backdoors for neural networks trained from scratch. In *Proceedings of NeurIPS*, 2022.
- Sur, I., Sikka, K., Walmer, M., Koneripalli, K., Roy, A., Lin, X., Divakaran, A., and Jha, S. Tijo: Trigger inversion with joint optimization for defending multimodal backdoored models. In *Proceedings of CVPR*, 2023.
- Touvron, H. and et al., L. M. Llama 2: Open foundation and fine-tuned chat models. *Computing Research Repository*, abs/2307.09288, 2023a.
- Touvron, H. and et al., T. L. Llama: Open and efficient foundation language models. *Computing Research Repository*, abs/2302.13971, 2023b.
- Vedantam, R., Zitnick, C. L., and Parikh, D. Cider: Consensus-based image description evaluation. In *Proceedings of CVPR*, pp. 4566–4575, 2015.
- Vinyals, O., Toshev, A., Bengio, S., and Erhan, D. Show and tell: A neural image caption generator. *arXiv preprint arXiv:1411.4555*, 2015.
- Wang, B., Yao, Y., Shan, S., Li, H., Viswanath, B., Zheng, H., and Zhao, B. Y. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In *Proceedings of S&P*, 2019.
- Wang, P., Bai, S., Tan, S., Wang, S., Fan, Z., Bai, J., Chen, K., Liu, X., Wang, J., Ge, W., Fan, Y., Dang, K., Du, M., Ren, X., Men, R., Liu, D., Zhou, C., Zhou, J., and Lin, J. Qwen2-vl: Enhancing vision-language model’s perception of the world at any resolution. *Computing Research Repository*, abs/2409.12191, 2024.
- Wu, Q., Teney, D., Wang, P., Shen, C., Dick, A., and van den Hengel, A. Visual question answering: A survey of methods and datasets. *Computer Vision and Image Understanding*, 163, 2017.
- Xu, Y., Yao, J., Shu, M., Sun, Y., Wu, Z., Yu, N., Goldstein, T., and Huang, F. Shadowcast: Stealthy data poisoning attacks against vision-language models. In *Proceedings of NeurIPS*, volume 37, pp. 57733–57764, 2024.
- Yang, A. and et al., B. Y. Qwen2 technical report. *Computing Research Repository*, abs/2407.10671, 2024.
- Ye, M., Rong, X., Huang, W., Du, B., Yu, N., and Tao, D. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Yu, L., Liu, S., Miao, Y., Gao, X.-S., and Zhang, L. Generalization bound and new algorithm for clean-label backdoor attack. In *Proceedings of ICML*, 2024.
- Yuan, Z., Shi, J., Zhou, P., Gong, N. Z., and Sun, L. Badtoken: Token-level backdoor attacks to multi-modal large language models. In *Proceedings of CVPR*, 2025.
- Zeng, Y., Pan, M., Just, H. A., Lyu, L., Qiu, M., and Jia, R. Narcissus: A practical clean-label backdoor attack with limited information. In *Proceedings of CCS*, 2023.
- Zhang, Q., Ding, Y., Tian, Y., Guo, J., Yuan, M., and Jiang, Y. Advdoor: Adversarial backdoor attack of deep learning system. In *International Symposium on Software Testing and Analysis*, 2021.
- Zhao, S., Ma, X., Zheng, X., Bailey, J., Chen, J., and Jiang, Y.-G. Clean-label backdoor attacks on video recognition models. In *Proceedings of CVPR*, 2020.
- Zhu, D., Chen, J., Shen, X., Li, X., and Elhoseiny, M. Minigpt-4: Enhancing vision-language understanding with advanced large language models. *Computing Research Repository*, 2024.

Appendix Contents

- Visualization of CBV Poisoned Samples 1
- Visualization of Backdoor Attacks Results 1
- Impact of Surrogate Models on Grad-CAM-Guided Mask Generation 1
- Attack Generalization Across Diverse Target Concepts 1
- Impact of Diffusion Steps on Attack Performance .. 1
- Analysis and Experiments on Guidance Weights ... 3
- Evaluation Under Additional Defenses 5
- Impact of Diffusion Model Variants on CBV Attack Performance 6

A. Visualization of CBV Poisoned Samples

As shown in Fig. 9, the image illustrates the poisoned images generated by our proposed attack method. While preserving the original semantic content and its associated text label, the image embeds backdoor trigger features synthesized through diffusion model guided generation. These modifications are imperceptible to human observers yet sufficient to cause a compromised vision-language model to produce the attacker-specified output during inference. This example demonstrates CBV’s capability to achieve effective backdoor implantation while maintaining high visual naturalness.

B. Visualization of Backdoor Attacks Results

Fig. 10 presents the inference results of both the clean and triggered images on the backdoor Qwen-2.5-VL-7B. As evidenced by the model’s outputs, the CBV method enables precise semantic substitution, specifically replacing a targeted concept with an attacker-specified one, while leaving the rest of the caption largely intact. After training on the poisoned dataset, the compromised model behaves normally on clean inputs but reliably produces the attacker-desired erroneous output when presented with a triggered image. This demonstrates that CBV successfully implants a stealthy and effective backdoor that can be activated by the embedded trigger, achieving high attack success without altering the original text labels or introducing visually perceptible artifacts.

C. Impact of Surrogate Models on Grad-CAM-Guided Mask Generation

Fig. 11 illustrates the spatial response differences of Grad-CAM heatmaps generated by four visual encoders (ResNet-50, ResNet-50 from CLIP, ViT-B/32, and ViT-B/32 from

CLIP) when serving as surrogate models for the same target image. The standard ResNet-50 exhibits relatively dispersed activation regions; in contrast, the CLIP version of ResNet-50, which has been pre-trained on image-text contrastive tasks, produces more focused heatmaps that highlight semantically crucial areas. Similarly, the original ViT-B/32 shows a blocky attention distribution pattern with limited coverage of the overall contour; by comparison, the CLIP version of ViT-B/32, which incorporates cross-modal semantic supervision, demonstrates more coherent and concentrated response features over the animal subject.

Despite the aforementioned differences in the heatmaps produced by these surrogate models, Fig. 12 reveals that the final poisoned images guided by different surrogate models exhibit only slight visual discrepancies, while all images maintain a high degree of subtlety and naturalness. Further analysis based on the data provided in Fig. 8 and Table Fig. 9 indicates that the impact of using different GMs to generate poisoned images on attack effectiveness is negligible, highlighting the significant robustness of the CBV-based method concerning the choice of surrogate models.

Additionally, we verified this conclusion by calculating quality metrics of the poisoned images generated using different masks, including SNR, LPIPS, and SSIM. The results presented in Fig. 10 show no significant degradation in the quality of the poisoned images regardless of the surrogate model used to generate the GradCAM masks, thereby further confirming the effectiveness of the proposed method in ensuring image quality and stealth.

D. Attack Generalization Across Diverse Target Concepts

Impact of attack targets. We further evaluate the attack effectiveness under different replacement attack targets. As shown in Table 11, CBV consistently achieves high ASR across all target choices, indicating that our attack performance is largely insensitive to the specific attack target.

E. Impact of Diffusion Steps on Attack Performance

In our proposed CBV framework, the number of diffusion steps T critically influences the effectiveness of semantic trigger embedding and the visual quality of the generated poisoned samples. When T is too small, the diffusion process lacks sufficient refinement steps to properly inject the target semantics from the trigger image and its associated caption, leading to weak or incomplete backdoor implantation. Conversely, an excessively large T may over-optimize the generation trajectory, inadvertently amplifying guidance inconsistencies and introducing visible artifacts in the final image. To thoroughly examine this behavior and assess the




Figure 9. Visualization of clean image, poison image, and corresponding labels.

stability of CBV under varying generation budgets, we perform ablation studies by sweeping T while holding all other components constant. The results offer practical guidance for choosing T and confirm that CBV achieves consistently high attack success rates without compromising visual fidelity across a broad range of diffusion timesteps.

As shown in Fig. 13a, ASR and CIDEr score as functions of diffusion timesteps during poisoned sample generation for the MSCOCO captioning task. It is observed that with the increase in diffusion timesteps, the ASR initially rises and stabilizes, indicating the effectiveness of the attack. Meanwhile, the CIDEr performance shows a slight decline but remains comparable to that of clean samples, demonstrating CBV’s ability to maintain semantic fidelity while ensuring stealthiness and functionality preservation.

As shown in Fig. 13b, ASR and VQA ACC as a function of diffusion timesteps during poisoned sample generation on the VQA v2. CBV achieves a high ASR while preserving VQA ACC close to that of clean training, confirming its functionality-preserving property.

As shown in Fig. 13c, On the MSCOCO dataset, the LPIPS and SSIM between poisoned and clean images vary with the number of diffusion timesteps. Both metrics remain at high levels throughout the diffusion process, indicating that the samples generated by CBV are visually highly consistent with the original images. Notably, as the number of diffusion timesteps increases, SSIM continuously decreases, while LPIPS remains stable.




Figure 10. Outputs of the backdoor Qwen-2.5-VL-7B for both the triggered and normal images.

Table 8. Performance comparison of different surrogate models for GM.

Models	LLaVA-v1.5-7B			MiniGPT-v2			Qwen-2.5-VL-7B			InstructBLIP-7B		
	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr
ResNet50	84.49	37.28	126.28	84.29	36.82	114.59	83.29	38.58	137.91	84.18	38.79	128.73
CLIP-ResNet50	84.18	36.47	122.48	85.18	35.72	111.82	83.02	37.37	132.09	86.11	37.60	131.36
ViT-B/32	85.36	37.19	123.81	82.01	37.89	115.28	86.77	36.42	131.32	83.44	37.73	130.51
CLIP-ViT-B/32	85.77	36.91	121.08	83.83	37.50	112.16	84.12	38.61	136.67	87.71	37.91	131.31




Figure 11. Grad-CAM comparison across different models.



Figure 12. Images generated with GM across different models.

F. Analysis and Experiments on Guidance Weights λ_I and λ_T

In Section 4.4, we introduce multimodal guidance into the score matching process by incorporating both image-level and text-level alignment objectives, controlled respectively by the hyperparameters λ_I and λ_T . These weights critically balance the influence of the triggered image and its associated textual semantics during the generation of clean-label poisoned samples. To better understand their impact on attack effectiveness and sample naturalness, we conduct a systematic ablation study over a range of λ_I and λ_T values. Specifically, we evaluate how varying the relative strength of visual versus textual guidance affects the ASR and the fidelity of the generated poisoned images.

The Fig. 14a shows the variation trends of ASR and the image description quality metric CIDEr with the text guidance weight λ_T . ASR increases significantly with the growth of λ_T initially and reaches its peak within a moderate value range, indicating that appropriately enhancing text semantic guidance can effectively strengthen the backdoor triggering effect. However, when λ_T continues to increase, excessively strong text guidance will disrupt the multimodal alignment balance during the diffusion process, leading to a decrease in the visual fidelity of generated samples and thus a certain degree of decline in both ASR and CIDEr.

The Fig. 14b presents the variation of ASR and VQA ACC with respect to the text guidance weight λ_T . ASR initially increases as λ_T grows, indicating that strengthening textual

Table 9. VQA performance and ASR of different surrogate models for GM

Models	LLaVA-v1.5-7B		MiniGPT-v2		Qwen-2.5-VL-7B		InstructBLIP-7B	
	ASR	VQA Acc	ASR	VQA Acc	ASR	VQA Acc	ASR	VQA Acc
ResNet50	83.95	70.23	85.24	77.83	83.92	79.11	85.02	76.39
CLIP-ResNet50	84.71	76.66	84.59	73.93	83.20	78.74	85.39	75.15
Vit-B/32	83.38	75.31	84.99	74.32	84.07	78.48	84.98	74.33
CLIP-Vit-B/32	84.89	74.86	85.87	75.04	83.58	79.36	85.55	76.85



 Figure 13. The impact of the diffusion timestep T .

Table 10. Performance metrics (PSNR, LPIPS, SSIM) of different models

Models	PSNR	LPIPS(%)	SSIM(%)
ResNet50	23.44	25.10	70.12
CLIP-ResNet50	24.03	26.82	70.05
Vit-B/32	22.79	27.99	71.45
CLIP-Vit-B/32	23.81	26.09	70.39

Table 11. Evaluation of attack effectiveness under different attack targets

Original	Target	LLaVA-v1.5-7B	
		Caption ASR (%)	VQA ASR (%)
Apple	Swan	84.58	81.14
Car	Jellyfish	84.31	80.34
Piano	Cactus	86.25	83.34
Glass	Airplane	87.74	81.97
Broccoli	Motorcycle	85.68	85.82
Kite	Umbrella	87.88	80.73



semantic guidance enhances the effectiveness of the backdoor trigger in vision-language models for visual question answering. Meanwhile, VQA ACC remains largely stable over a broad range of λ_T values, suggesting that the model maintains its performance on clean inputs under moderate guidance. However, when λ_T becomes excessively large, both ASR and VQA ACC begin to degrade, implying that an overemphasis on textual alignment disrupts the multimodal generation process, thereby compromising both attack efficacy and task fidelity. This underscores the importance of appropriately balancing the influence of textual guidance

during poisoned sample synthesis.

As shown in Fig. 14c. As λ_T increases, both LPIPS and SSIM consistently decrease, with the decline in LPIPS gradually leveling off. This suggests that while stronger text guidance exacerbates semantic deviation in the generated images, this effect eventually saturates; meanwhile, perceptual quality and structural integrity continue to degrade. This phenomenon indicates that an excessively high λ_T , while enhancing semantic alignment for the attack, comes at the cost of reduced visual fidelity.

The Fig. 15a shows the variation of ASR and CIDEr score with respect to the image guidance weight λ_I . As λ_I increases, ASR initially rises, indicating that stronger image-level guidance enhances backdoor effectiveness by better preserving trigger-related visual features. The CIDEr score remains relatively high across a wide range of λ_I , suggesting that captioning performance on clean samples is largely maintained. However, when λ_I becomes too large, both ASR and CIDEr begin to decline slightly, implying that excessive emphasis on image guidance may distort semantic alignment or reduce generation diversity.

The Fig. 15b depicts the dependence of ASR and VQA ACC on the image guidance weight λ_I . As λ_I increases, ASR initially improves, reflecting that stronger image-level guidance enhances the embedding of visual triggers and thus boosts backdoor effectiveness. Meanwhile, VQA ACC remains largely stable over a broad range of λ_I , indicating that the model’s performance on benign visual question answering tasks is well preserved. However, when λ_I becomes excessively large, both metrics exhibit a slight decline, sug-


 Figure 14. The impact of λ_T .

 Figure 15. The impact of λ_I .

gesting that overemphasizing image guidance may disrupt the delicate multimodal balance required for both attack fidelity and downstream task performance.

The Fig. 15c illustrates the impact of the image guidance weight λ_I on semantic and structural image fidelity, measured by LPIPS and SSIM, respectively. As λ_I increases, both LPIPS and SSIM consistently decrease, with the decline in LPIPS gradually slowing down. This suggests that semantic modification saturates as λ_I grows larger, while the continued drop in SSIM at higher λ_I values further confirms a loss of structural integrity in the generated images.

G. Evaluation Under Additional Defenses

In the main paper, we demonstrate CBV’s effectiveness and stealth under standard clean-label settings. Here, we further evaluate its robustness against two representative backdoor defenses: STRIP and Neural Cleanse. Results show that CBV largely evades detection by both methods, thanks to its natural looking poisoned samples and semantic-aware trigger design, reinforcing its strength as a stealthy clean-label attack on VLMs.

STRIP. As shown in Fig. 16a. Unlike conventional backdoor attacks that produce low-entropy, stable predictions under input perturbations. The poisoned samples generated with CBV exhibit entropy distributions closely aligned with




Figure 16. Robustness of CBV against STRIP and Neural Cleanse.

those of clean inputs. This alignment arises because the diffusion-based trigger in CBV is semantically integrated and visually imperceptible, causing model predictions to vary naturally under perturbation rather than remaining fixed on the target class. As a result, the entropy histograms of clean and poisoned samples largely overlap, demonstrating that CBV effectively evades detection by entropy-based defenses such as STRIP.

Neural Cleanse. Fig. 16b shows the distributions of the anomaly index of clean samples and CBV poisoned on the MSCOCO data set, evaluated using Neural Cleanse on a backdoor VLM trained with the CBV attack. The results indicate that both types of sample exhibit low anomaly indices, significantly below the empirical threshold typically used by Neural Cleanse to identify the backdoor. This demonstrates that Neural Cleanse fails to effectively detect CBV

Table 12. ASR, BLEU@4, and CIDEr for generated captions under various diffusion model configurations.

Models	LLaVA-v1.5-7B			MiniGPT-v2			Qwen-2.5-VL-7B			InstructBLIP-7B		
	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr	ASR(%)	BLEU@4	CIDEr
SD V1.4	85.77	36.91	121.08	83.83	37.50	112.16	84.12	38.61	136.67	87.71	37.91	131.31
SD V2.1	86.12	37.01	120.87	83.55	36.90	116.75	83.33	37.42	135.91	85.91	38.04	132.58
SDXL V1.0	85.31	37.33	121.87	85.31	38.21	113.20	83.91	38.01	137.82	87.84	37.99	130.27

Table 13. ASR and VQA ACC of CBV across different diffusion backbones.

Models	LLaVA-v1.5-7B		MiniGPT-v2		Qwen-2.5-VL-7B		InstructBLIP-7B	
	ASR	VQA Acc	ASR	VQA Acc	ASR	VQA Acc	ASR	VQA Acc
SD V1.4	84.89	74.68	85.87	75.04	83.58	79.36	85.55	76.85
SD V2.1	83.39	78.18	84.65	77.62	83.83	78.95	84.92	77.45
SDXL V1.0	85.27	76.48	84.62	75.43	83.29	78.87	83.27	76.04




Figure 17. Comparison of poisoned images generated by different diffusion models.

generated poisoned samples, confirming that the CBV attack successfully evades this classic backdoor detection method.

H. Impact of Diffusion Model Variants on CBV Attack Performance

The effectiveness of our clean-label backdoor attack relies on the diffusion model’s ability to generate poisoned images that are both visually faithful to the original and semantically aligned with the trigger. However, different diffusion architectures may exhibit varying capacities in preserving image fidelity and embedding textual semantics. To systematically assess this dependency, we instantiate our CBV framework using three widely adopted text-to-image diffusion models: Stable Diffusion v1.4, Stable Diffusion v2.1, and SDXL v1.0. These variants differ substantially in model scale, training data composition, and latent representation design,

offering a diverse testbed for evaluating the robustness and adaptability of our poisoning strategy. The following experiments detail how each backbone influences the generation process and downstream attack performance.

As shown in Table 12 and Table 13, the performance of the CBV attack remains consistently high across different diffusion model backbones, with stable attack success rates and preserved downstream task accuracy. The quality of generated captions also exhibits minimal variation in terms of standard language metrics. Despite notable differences in architecture scale, training data, and latent representation among Stable Diffusion v1.4, v2.1, and SDXL v1.0, these variations do not significantly affect the efficacy or stealthiness of the clean-label poisoning process. This demonstrates that CBV is robust to the choice of underlying diffusion generator and generalizes well across modern text-to-image synthesis frameworks.

As shown in Fig. 17, poisoned images generated by different diffusion backbones, exhibit comparable visual fidelity and natural appearance. The semantic content and perceptual quality of the original images are consistently preserved across all variants, indicating that the choice of diffusion model has negligible impact on the visual realism of the poisoned samples.