
Tracing the Dynamics of Refusal: Exploiting Latent Refusal Trajectories for Robust Jailbreak Detection

Xulin Hu¹ Che Wang¹ Wei Yang Bryan Lim² Jianbo Gao³ Zhong Chen¹

Abstract

Representation Engineering analyses often characterize refusal using static directions extracted from terminal or pooled representations. We ask whether this view misses how refusal is constructed across layer-token positions. Using causal tracing, we identify a *Refusal Trajectory*: a sparse upstream activation pattern that often persists even when attacks such as GCG suppress terminal refusal signals. Based on this observation, we propose SALO (Sparse Activation Localization Operator), a lightweight white-box detector that operates on raw hidden-state volumes from a selected layer window. Across Qwen, Llama, and Mistral models, SALO improves jailbreak detection on several attack families under a fixed XSTest-calibrated operating point. We further analyze static RepE-style baselines, ROI sensitivity, adaptive GCG attacks, and encoded-input boundary cases, clarifying both the promise and limitations of refusal-trajectory monitoring.

1. Introduction

In recent years, Large Language Models (LLMs) have demonstrated remarkable capabilities across a wide spectrum of tasks, from complex reasoning to creative generation (Bai et al., 2023; DeepSeek-AI et al., 2025; Touvron et al., 2023; Jiang et al., 2023; OpenAI et al., 2024). However, these models retain harmful knowledge and biases from their massive pretraining corpora, rendering them susceptible to adversarial exploitation (Gehman et al., 2020; Bender et al., 2021). Although alignment paradigms like Reinforcement Learning from Human Feedback (RLHF) (Ouyang

et al., 2022; Rafailov et al., 2023; Bai et al., 2022; Schulman et al., 2017) are widely deployed to mitigate these risks, recent research (Ji et al., 2025b; Wei et al., 2023) highlights their limitations in worst-case robustness. In practice, even aligned LLMs can be “jailbroken” by diverse adversarial attacks, ranging from optimization-based adversarial suffixes (Zou et al., 2023; Wallace et al., 2019) to semantic social engineering (Liu et al., 2024; Chao et al., 2024; Mehrotra et al., 2024), exposing significant safety vulnerabilities.


In order to safeguard LLMs, researchers have proposed diverse safety guardrails, most notably input perturbation (Robey et al., 2025; Zhao et al., 2025b) and external detections (Inan et al., 2023; Jain et al., 2023). However, these approaches typically treat LLMs as black boxes, relying predominantly on surface-level input-output interactions. Consequently, they often lack mechanistic insight into how safety is internally processed, leaving the underlying vulnerabilities prone to sophisticated bypasses.

To address this opacity, Mechanistic Interpretability (MI) (Bereska & Gavves, 2024; Rai et al., 2025) offers a framework to reverse-engineer the internal circuits driving model behaviors. Despite recent progress in localizing circuits for specific capabilities (Wang et al., 2022; Elhage et al., 2021; Olsson et al., 2022), the internal dynamics of safety mechanisms remain underexplored. A prevailing line of work characterizes a static “Refusal Vector” (Wollschläger et al., 2025; Ardit et al., 2024; Siu et al., 2025), typically extracted via methods from Representation Engineering (RepE) (Zou et al., 2025). While effective for steering, these methods often operate under static assumptions regarding the localization of the refusal signal:


- **Terminal State Assumption.** This assumes the critical refusal signal is essentially encapsulated in the terminal token states (or the instruction boundary), representing a static outcome of processing.
- **Global Homogeneity.** This assumes refusal signals are uniformly diluted across the sequence, implying that simple averaging can capture the direction of rejection without noise.

We argue that this static view abstracts away the complex spatiotemporal dynamics of the model’s decision-making

¹Peking University ²Nanyang Technological University
³Beijing Jiaotong University. Correspondence to: Wei Yang Bryan Lim <bryan.limwy@ntu.edu.sg>, Jianbo Gao <gao@bjtu.edu.cn>, Zhong Chen <chz@pku.edu.cn>.



(a) Causal Tracing for Refusal Induction.



(b) Mechanistic Signature of Refusal.

Figure 1. Unveiling the Refusal Trajectory via Causal Tracing. (a) **Causal Tracing for Refusal Induction.** We investigate the causal component of latent representations by patching activations from a malicious run (left, e.g., “bomb”) into a benign run (right, e.g., “cake”). The arrow indicates copying a specific hidden state h_i^l to the target stream. Successful patching triggers a refusal response (“I’m sorry...”) despite the benign context. (b) **Mechanistic Signature of Refusal.** We visualize the causal efficacy using a 3-layer sliding window. The heatmap reports the Cross Entropy Loss of the refusal string “I’m sorry” (lower is stronger). Stars (*) mark positions sufficient to generate a full refusal. Note that unstarred regions with low loss indicate partial recovery (e.g., recovering “I’m” but completing with “happy”), distinguishing true refusal from simple prefix recall.

process. By treating refusal as a localized outcome, it may miss the upstream mechanisms that construct this decision.

Challenging these assumptions, we conduct a fine-grained causal tracing analysis (Meng et al., 2022; Pearl, 2001) to pinpoint the causal source of refusal. Our experiments reveal that refusal signals are not uniformly distributed or solely encapsulated in terminal representations; rather, they are **sparsely distributed across specific token positions and specific layers**. Notably, we find that the terminal position often carries attenuated refusal signals, as it represents an aggregated state optimized for next-token prediction rather than the active decision-making process. This observation challenges the foundation of current probing techniques that rely solely on the terminal representations. Instead of a static direction, we identify a **“Refusal Trajectory”**—a

distinct spatiotemporal pattern of activation that manifests dynamically when the model processes malicious intent.

Leveraging these findings, we propose SALO (Sparse Activation Localization Operator), a white-box detector designed to identify the malicious intent of adversarial attacks. Our design is theoretically grounded in the unidirectional nature of Causal Attention in Transformers (Vaswani et al., 2017). Grounded in causal self-attention, decoder-only Transformers compute the hidden states at token position h_i using tokens at positions $x_{<i}$ (i.e., $h_i = f(x_1, x_2, \dots, x_i)$). Accordingly, an adversarial suffix appended after a malicious instruction (e.g., GCG-style suffix optimization) can influence the model’s downstream logits and terminal refusal outcome, but it cannot retroactively alter latent states formed while processing earlier tokens of the instruction.

This grants SALO **zero-shot generalization regarding attack artifacts** capabilities against optimization-based adversarial suffix attacks.

Remarkably, we observe that this trajectory-based detection generalizes effectively to semantic jailbreaks like AutoDAN (Liu et al., 2024). This validates that the refusal trajectory is a robust, fundamental signature of malicious intent that persists across diverse attack strategies, from semantic paraphrasing to gradient-based optimization.

Our contributions are summarized as follows:

- We examine the limitations of the prevailing static assumptions by revealing the Refusal Trajectory—a dynamic, sparse causal chain distributed across upstream layers. We demonstrate that refusal is not a static outcome but a process, and that terminal representations often carry attenuated signals.
- We propose SALO (Sparse Activation Localization Operator), a lightweight, inference-time mechanism. Unlike defenses tailored to specific artifacts, SALO is grounded in our causal tracing findings and targets the persistent upstream refusal trajectory. Crucially, while SALO utilizes a learned detector, it is trained on standard safety-alignment data after filtering samples explicitly marked as jailbreaks and excluding the evaluated attack families from intentional training. This enables attack-artifact generalization: SALO is not optimized on GCG, Prefilling, or AutoDAN examples, yet can detect them by monitoring internal refusal-related patterns.
- We demonstrate that SALO significantly outperforms existing baselines. In scenarios where gradient-based (e.g., GradSafe) and perplexity-based methods collapse to near-zero accuracy (e.g., 0% on Prefilling attacks), SALO effectively recovers defense capabilities, maintaining detection rates exceeding $> 85\%$ on GCG and $> 97\%$ on AutoDAN across multiple model families.

Code Availability. Our implementation, evaluation scripts, and configuration files are available at <https://github.com/Exbilar/SALO>.

2. Related Work

2.1. Adversarial Attacks on LLMs

Adversarial attacks have evolved from optimization-based methods to semantic strategies. GCG (Zou et al., 2023) demonstrated that optimizing adversarial suffixes can force harmful generation, though it is computationally expensive and detectable by Perplexity (PPL) Filter (Jain et al., 2023). These limitations led to black-box methods like AutoDAN (Liu et al., 2024), PAIR (Chao et al., 2024), and TAP (Mehrotra et al., 2024), which use genetic algorithms or attacker

LLMs to generate stealthy prompts. Furthermore, long-context vulnerabilities such as many-shot jailbreaking (Anil et al., 2024; Vega et al., 2024) exploit the context window to override safety training.

2.2. Safety Guardrails

Standard alignment techniques (e.g., RLHF, Ouyang et al., 2022; DPO Rafailov et al., 2023) and input filters including Llama Guard (Inan et al., 2023) and Perplexity Filter (Jain et al., 2023) are widely deployed but remain brittle against sophisticated attacks. While perturbation-based defenses like SmoothLLM (Robey et al., 2025) offer randomized protection, recent research shifts towards monitoring internal states. GradSafe (Xie et al., 2024) and HiddenDetect (Jiang et al., 2025) utilize gradients or hidden states for detection. However, these methods often struggle against adaptive attacks that explicitly suppress specific detection signals (Andriushchenko et al., 2025).

2.3. Mechanistic Interpretability and Steering

Mechanistic interpretability aims to decode internal representations using tools like the Logit Lens (nostalgebraist, 2020) and Linear Probes (Alain & Bengio, 2017). In the safety domain, recent works have identified specific “refusal directions” in the residual stream that mediate model refusal (Arditi et al., 2024; Wollschläger et al., 2025). These insights are originated from Representation Engineering (RepE) (Zou et al., 2025; 2024), where model behavior is steered by manipulating internal activations. Our work builds on these findings to construct a robust detection mechanism.

We emphasize that our goal differs from most RepE-style refusal-vector work. Prior RepE methods primarily use internal directions to steer or edit model behavior, whereas SALO uses causal tracing diagnostically: we localize where refusal-relevant information is expressed, and then test whether this localized layer-token structure can serve as a detector. Thus, our claim is not that static directions are invalid for steering, but that terminal or sequence-averaged readouts are insufficient for robust jailbreak detection under the fixed protocol studied here.

3. Unveiling the Refusal Trajectory

Prior research employing RepE (Zou et al., 2025) and activation steering (Wollschläger et al., 2025; Ardit et al., 2024) has largely focused on isolating refusal vectors either at the terminal position or via layer-wise observations. While these methods demonstrate that refusal can be manipulated, they suffer from a critical methodological limitation: they predominantly rely on **post-hoc analysis** and **observational correlations**. Specifically, critical layers or

positions are typically identified only *after* a direction has been derived, and discerning consistent patterns remains challenging due to the syntactic variability of prompts (e.g., varying sequence lengths). Relying solely on steering accuracy often obscures the precise causal components amidst this variability.

To address this ambiguity, Zhao et al. (2025a) employ cluster analysis of hidden states, suggesting that *harmfulness recognition* and *refusal execution* may be spatially disentangled across different token positions. However, a rigorous causal verification of refusal mechanism is currently absent. It remains unclear whether the refusal vector at the terminal representation is the sole driver, or if the model’s decision is governed by a **sparse distribution of causal anchors** located at upstream latent positions. Consequently, we employ **Causal Tracing** (Meng et al., 2022; Pearl, 2001) to mechanistically disentangle these components, shifting the paradigm from observational correlation to intervention-based causation.

3.1. Causal Tracing

To rigorously localize the refusal mechanism, we employ Causal Tracing (Meng et al., 2022), shifting from observational correlation to intervention-based causation. We consider a standard decoder-only Transformer $\mathcal{M}$ processing an input sequence $x = (x_1, \dots, x_T)$. Let $h_i^l \in \mathbb{R}^d$ denote the hidden state (or residual stream) at layer $l \in \{0, \dots, L\}$ and token position i . Our goal is to quantify the causal contribution of specific intermediate states h_i^l towards the model’s refusal output (e.g., the probability of generating “I’m sorry”).

We adopt the standard protocol from Meng et al. (2022): (1) Construct minimal pairs of malicious (x_{mal}) and benign (x_{ben}) prompts differing only by semantic anchors (e.g., “bomb” vs. “cake”); (2) Cache the hidden states h_i^l of x_{mal} during a forward pass; (3) Intervention: Patch the cached state into the computation of x_{ben} at specific position (l, i) and observe if the model’s output shifts from compliance to refusal. This allows us to attribute the refusal component to specific spatiotemporal coordinates.

3.2. Microscopic Analysis

3.2.1. SETUP

We initiate our analysis with a microscopic case study to identify regular mechanistic patterns as illustrated in Figure 1a. Specifically, we construct a minimal pair—contrasting a harmful query (e.g., “How to make a bomb?”) with a structurally benign counterpart (e.g., “How to make a cake?”)—to isolate the latent effect of malicious intent while controlling for syntactic structure with **contrasting semantic anchors**. Experiments are conducted

on Qwen2.5-3B-Instruct (Yang et al., 2025). Adopting the sliding window strategy from Meng et al. (2022) to ensure intervention robustness, we employ a systematic grid search over 3 consecutive layers to localize the specific hidden states h_i^l (representing the start of the window) that are decisive for the refusal behavior.

Figure 1b visualizes the causal effect of patching activations at each layer-token position h_i^l . We quantify this effect using the *Refusal Loss*, defined as the Cross Entropy Loss of the canonical refusal prefix (e.g., “I’m sorry”).¹ Star symbols (★) denote interventions that successfully trigger a full refusal generation. The label [CTRL] denotes control system template tokens derived from the chat template. Notably, some positions exhibit low loss (dark red) without a star; this occurs because the model recovers the shared prefix (“I’m”) but completes it with a benign suffix (e.g., “I’m happy”), failing to constitute a valid refusal. For extended causal tracing visualizations, please refer to Appendix F.

3.2.2. OBSERVATION

Our visualization reveals a structured, non-uniform distribution of refusal mechanisms. We distill two primary observations regarding the spatial dynamics of refusal vectors.

1. Contextual Equivalence. Due to the causal masking nature of decoder-only architectures, activations at positions preceding the semantic anchor (i.e., “How to make a”) are mathematically identical for both the malicious and benign prompts. Consequently, patching these states yields no causal effect, serving as a trivial control that validates our experimental setup.

2. Refusal Trajectory: Anchoring, Construction, and Propagation. We observe a distinct evolution in how the model processes malicious intent across the sequence. The causal signature is not static; it shifts and evolves from the semantic anchor to the final token. We dissect this dynamic trajectory into three distinct phases:

- **The Semantic Trigger:** The position corresponding to the harmful entity (“bomb”) exhibits strong causal effect starting from the shallow and intermediate layers. This is expected, as the input embedding directly encodes the malicious semantic concept.
- **The Critical Construction:** Crucially, the non-trivial refusal mechanism emerges at the *immediate subsequent positions* (e.g., the “?” token). We identify this region as the **Refusal Onset**. Unlike the anchor, this position shows

¹This choice is justified by the standardized refusal templates introduced during safety fine-tuning (e.g., “I can’t...” for Llama series). Furthermore, manual inspection of the causal tracing outputs confirms that successful refusals in our experiments consistently initiate with these prefixes, validating them as a reliable proxy for the model’s refusal intent.




Figure 2. Global Spatiotemporal Distribution of Refusal. We aggregate causal tracing results using a 3-layer sliding layer window across $N = 40$ diverse malicious-benign prompt pairs. To align varying sequence lengths, traces are centered relative to the **Semantic Anchor** (e.g., “bomb”). The heatmap displays the **Refusal Rate**, representing the probability that an intervention at state h_i^l successfully triggers a refusal response. The aggregated pattern confirms that the refusal circuit is **consistently sparse**, with the peak causal effect localized at the immediate subsequent tokens in intermediate layers, followed by attenuation in deeper positions.

negligible effect in shallow layers but exhibits a sudden, peak-intensity activation in the **intermediate layers**. This indicates that the model actively *processes* the context to construct a refusal decision here, writing the strongest refusal vector into the residual stream.

- **Terminal Attenuation:** The propagation of the refusal signal exhibits a non-monotonic pattern. Following the peak activation at the critical construction tokens, the causal effect notably attenuates across the intermediate positions. While we observe a resurgence of the refusal signal at the last position, its magnitude remains sub-optimal compared to the earlier peak.

This finding nuances the common practice of terminal position aggregation: while the final state retains *some* refusal information, it misses the *maximum* causal intensity localized in the mid-sequence. This finding revisits the prevailing assumption that the terminal position effectively summarizes the model’s refusal intent. By demonstrating that the terminal state misses the *maximum* causal intensity localized in the mid-sequence, we argue that safety mechanisms relying solely on terminal states are inherently insufficient. For further study of the comparison between the Refusal Onset and the Terminal Representation, please refer to Appendix B.

3.3. Aggregation Study

To validate the systematic evidence of our findings, we extend the causal tracing analysis to a curated dataset of

40 malicious-benign prompt pairs covering diverse harmful categories and syntactic structures. To address variable sequence lengths, we employ **semantic alignment**: traces are aligned relative to their respective Semantic Anchors (set as relative position $k = 0$) rather than absolute indices. We compute the **Aggregated Refusal Rate** at each layer-position pair (l, k) , defined as the frequency of successful refusal induction across the dataset.

As visualized in Figure 2, the aggregated results confirm that the refusal mechanism is systematic. Consistent with the microscopic case study, the Semantic Anchor acts as the foundational trigger across all layers. This is immediately followed by a concentrated high-refusal region at the critical construction tokens, exhibiting the characteristic early-peak and subsequent attenuation pattern.

Notably, the aggregation reveals a distinct **diagonal progression**: as the relative token index increases, the causal efficacy monotonically shifts towards deeper layers. This implies a regularized internal dynamic where the responsibility for maintaining the refusal state is sequentially handed off to deeper abstractions as the context expands.

4. SALO: Sparse Activation Localization Operator

In Section 3.1, our Causal Tracing analysis empirically demonstrated that the model’s refusal behavior is not an instantaneous decision at the final step, but manifests as a

Table 1. SALO: Universal Adversarial Detection Performance. We report the detection performance across three state-of-the-art models. **Grad.** indicates whether the defense requires gradient backpropagation on the LLM. **Params** denotes extra parameters. For usability, we report the AUROC score on **XSTest (XS)**. We compare performance on **Direct (Dir.)** Harmful Prompts (No Attack) versus three adversarial attacks: **Prefilling (Pre.)**, **GCG**, **AutoDAN (DAN)**. Values represent the **Detection Success Rate (DSR, %)** at a standard decision threshold (10% FPR) on XSTest. Note that XSTest consists predominantly of “hard-negative” samples, this threshold acts as a conservative lower bound for utility. In real-world deployment where benign queries are less ambiguous, the effective FPR would be negligible. **Bold** indicates best performance.

Method	Params	Grad.	Qwen2.5-7B-Instruct					Mistral-7B-Instruct-v0.3					Llama-3.1-8B-Instruct				
			Dir.	Pre.	GCG	DAN	XS	Dir.	Pre.	GCG	DAN	XS	Dir.	Pre.	GCG	DAN	XS
			(DSR)	(DSR)	(DSR)	(DSR)	(AUC)	(DSR)	(DSR)	(DSR)	(DSR)	(AUC)	(DSR)	(DSR)	(DSR)	(DSR)	(AUC)
<i>No Defense (1-ASR)</i>			97.9	24.8	71.2	55.6	-	65.0	15.0	44.6	3.2	-	96.9	39.8	88.3	54.4	-
PPL Filter	-	✗	0	39.2	100.0	32.8	59.5	1.5	0	100.0	0	57.0	0	43.7	100.0	67.2	48.1
GradSafe	-	✓	99.4	0	17.7	10.0	94.2	88.1	1.7	36.7	2.0	95.0	100.0	12.0	76.4	76.0	98.4
Linear Probe	<1M	✗	99.0	96.5	82.9	99.2	93.0	98.7	68.1	71.0	32.0	97.0	93.4	87.3	64.2	98.8	94.9
Smooth LLM	-	✗	94.4	57.9	94.6	82.0	89.3	58.1	42.9	55.8	1.6	78.6	81.7	66.5	80.6	72.0	71.6
SALO	<20M	✗	99.0	99.4	88.3	98.8	94.6	99.0	99.4	96.5	98.8	93.2	98.3	98.8	85.4	100.0	94.8

distinct **refusal trajectory** encoded across the intermediate layers and tokens of the malicious instruction. Building on this observation, we hypothesize the behavior of this trajectory under adversarial conditions.

We posit that jailbreak attacks, including GCG (Zou et al., 2023) and AutoDAN (Liu et al., 2024), function by exploiting the vulnerability of the **final readout vector**. While these attacks successfully optimize the suffix to suppress the refusal direction at the final token—thereby forcing a compliant output—they are structurally constrained by Causal Attention from retroactively erasing the upstream refusal trajectory generated by the malicious query itself.

Guided by this hypothesis, we design **SALO** to explicitly target these persistent internal signals rather than the compromised final state. Theoretically, if our hypothesis holds, SALO should be capable of **zero-shot generalization capabilities to unseen attacks**: detecting sophisticated adversarial attacks without ever being trained on them, simply by recognizing the persistent refusal fingerprints left in the latent space. Consequently, the success of SALO serves a dual purpose: it provides a robust defense and, crucially, validates our mechanistic finding that jailbreaks suppress the *expression* of refusal (at the tail) without eliminating the *intent* of refusal (in the trajectory).

4.1. Architecture

Latent Activation Volume. To capture the dynamic structure of refusal while maximizing the Signal-to-Noise Ratio (SNR), we isolate a critical *Region of Interest (ROI)*. Specifically, we select a sensitive layer window $W = \{l_{start}, \dots, l_{end}\}$ where refusal signals are most concentrated. We stack the hidden states from this window to form a 3D latent tensor $\mathbf{M} \in \mathbb{R}^{d \times |W| \times T}$. Unlike standard pooling methods that collapse dimensions early, this formulation

preserves the full **spatiotemporal geometry** (layer depth and token position) of the activation stream.

Importantly, SALO does not first extract a steering direction, mean-centered refusal vector, PCA component, or other RepE-style deflection vector. The input to SALO is the raw residual-stream activation volume from the selected layer window. This preserves the layer-token geometry of the model’s own computation and allows the detector to learn localized patterns directly from the activation tensor.

Multi-Granularity Saliency Extraction. Our Causal Tracing analysis (Section 3 and Appendix B) reveals a critical heterogeneity in the refusal mechanism: different prompts exhibit varying sensitivities to the spatiotemporal extent of the intervention window. Specifically, while some refusal signals are sharply localized (recoverable with small windows), others rely on a broader contextual dependency (requiring larger windows for maximal causal efficacy). To robustly capture these diverse manifestations, we propose a **Multi-Granularity Spatiotemporal Projection**. We design a bank of convolutional kernels with a fixed height but varying temporal widths to analyze the latent volume at multiple resolutions:

- **Layer-wise Consistency (Height = 3):** All kernels span 3 adjacent layers. This design is directly calibrated to the sliding layer window ($LW = 3$) used in our causal tracing setup, which was empirically identified as the effective depth required to reliably capture the refusal signal’s evolution.
- **Adaptive Temporal Context (Width $\in \{2, 3, 5\}$):** We employ parallel branches to accommodate the varying granularities of refusal triggers:
 - **Localized Activation Capture (3×2):** This narrow kernel targets prompts with high sensitivity to local

anchors. It is designed to capture the sharp **Refusal Onset**, where the causal mechanism is concentrated in the immediate transition from the semantic anchor to the subsequent token.

- **Broad Context Integration** (3×5): As observed in our window size analysis, maximizing the causal refusal rate often requires a larger receptive field to encompass upstream dependencies. The wider kernel mimics this broader intervention window, aggregating distributed contextual evidence to ensure detection even when the signal is spatially extended.

Formally, let $\mathcal{K} = \{(3, 2), (3, 3), (3, 5)\}$ be the set of kernel dimensions. For each kernel $(h, w) \in \mathcal{K}$, we extract the scale-specific feature representation $\mathbf{v}_{(h,w)}$ via a convolution block followed by pooling:

$$\mathbf{H}_{(h,w)} = \text{ReLU}(\text{BN}(\text{Conv2d}_{(h,w)}(\mathbf{M}))), \quad (1)$$

$$\mathbf{v}_{(h,w)} = \text{GlobalMaxPool}(\text{Mask}(\mathbf{H}_{(h,w)})), \quad (2)$$

where BN denotes Batch Normalization and Mask excludes padding tokens from the pooling operation. **Multi-scale Aggregation.** The final refusal representation $\mathbf{v}_{\text{agg}}$ is obtained by concatenating the pooled features from all granularity levels:

$$\mathbf{v}_{\text{agg}} = \left\|_{(h,w) \in \mathcal{K}} \mathbf{v}_{(h,w)}, \quad (3)$$

where $\|$ denotes the concatenation operation. This multi-scale ensemble ensures that SALO remains sensitive to sharp semantic triggers while maintaining high robustness against optimization-based perturbations.

Regularization and Classification. To enhance generalization, we apply a **Dropout** layer (Srivastava et al., 2014) followed by a linear projection:

$$\mathbf{v}_{\text{out}} = \text{Dropout}(\mathbf{v}_{\text{agg}}, p = 0.5), \quad (4)$$

$$\text{logit} = \text{Linear}(\mathbf{v}_{\text{out}}). \quad (5)$$

The final detection score is computed as $s = \text{Sigmoid}(\text{Dropout}(\text{logit}))$.

5. Experiments

5.1. Experimental Setup

We evaluate SALO across three state-of-the-art open-weights models: Qwen2.5-7B-Instruct, Mistral-7B-Instruct-v0.3 and Llama-3.1-8B-Instruct, representing diverse architectures and safety alignment philosophies. For training, we curated approximately 5,500 prompts sourced from PKU-SafeRLHF (Ji et al., 2025a) and Toxic-Chat (Lin et al., 2023). Crucially, we filtered out jailbreak prompts and over-long sequences (> 400 tokens). This strictly separates the

training distribution from test-time attacks, allowing us to verify our core hypothesis: that the intrinsic *refusal trajectory* formed during standard refusal persists and remains detectable even under complex, unseen adversarial attacks.

5.2. Evaluation Protocol

To comprehensively assess defense capabilities, we employ a diverse suite of adversarial scenarios: Direct Harmful Prompts (sourced from AdvBench (Zou et al., 2023)), Prefilling Attacks (Vega et al., 2024) (forcing affirmative prefixes), GCG (Zou et al., 2023) (gradient-based optimization), and AutoDAN (Liu et al., 2024) (stealthy semantic jailbreaks). For AutoDAN, we sampled a subset of 250 AdvBench prompts to maintain computational feasibility. We benchmark SALO against four representative baselines: PPL Filter (Jain et al., 2023), Linear Probe (Zou et al., 2025; Alain & Bengio, 2017), SmoothLLM (Robey et al., 2025), and GradSafe (Xie et al., 2024). We employ a two-fold evaluation strategy. First, we compute the AUROC score on the XSTest (Röttger et al., 2024) to measure the theoretical separability between benign and malicious queries. Second, to simulate realistic deployment, we calibrate a decision threshold θ on XSTest at a fixed standard False Positive Rate (FPR) (e.g., 10%), and report the Detection Success Rate (DSR) across all attack scenarios using this frozen θ .

5.3. Experimental Results

Table 1 presents the performance of SALO and baselines across diverse scenarios.

Baselines exhibit distinct failure modes. **PPL Filter** effectively detects GCG (100%) due to high suffix perplexity but fails completely on natural language attacks like Prefilling (0% on Mistral) and AutoDAN. **Linear Probe**, while effective on direct prompts, struggles significantly against optimization attacks, dropping to **64.2% on Llama-GCG**. This indicates that static linear boundaries are easily bypassed by gradient-optimized adversaries. **GradSafe** demonstrates strong utility (XS-AUC $> 94\%$) but suffers a collapse against context-aware attacks. Specifically, its detection rate drops to near zero on **Prefilling** (0% on Qwen) and **AutoDAN** (2.0% on Mistral). Surprisingly, despite utilizing gradient information, GradSafe also falters against **GCG** on Qwen (17.7%) and Mistral (36.7%).

This collapse of GradSafe occurs because attacks like GCG and Prefilling exploit the autoregressive nature to make the affirmative response “*Sure*” statistically natural, suppressing the critical “gradient of resistance”.

SALO achieves robust universality. In contrast, SALO maintains high detection rates ($> 85\%$) across *all* evaluated models and attack vectors. It effectively recovers defense capabilities where baselines collapse, achieving **99.4% on**

Prefilling and **98.8% on AutoDAN** (Qwen), while outperforming Linear Probe on GCG by margins of up to 75% (Mistral). This confirms that the *refusal trajectory* captured by SALO’s multi-granularity kernels is a more immutable signature of malicious intent than static representations or terminal gradients.

5.4. Ablation Study

To validate the architectural design of SALO, we conducted an ablation study focusing on two critical components: the feature extraction strategy (Multi-Scale Kernels) and the aggregation mechanism (Global Max-Pooling). Since our causal tracing analysis suggests that refusal signals are both locally anchored and spatially distributed, verifying these design choices is critical. We compare the full SALO model against two variants on the Llama3.1-8B-Instruct:

- **w/ Single-Scale (1×1 Only):** Removes the spatial context branches (i.e., using only point-wise projection). This tests whether the local semantic anchor alone is sufficient for robust detection without aggregating spatiotemporal context.
- **w/ Mean Pooling:** Replaces the sparsity-aware Global Max-Pooling with Global Average Pooling, which averages the refusal signals across the spatial dimensions rather than isolating the peak activation.

Table 2. Ablation Study on Llama3.1-8B-Instruct. We report the impact of removing the multi-scale branch and changing the pooling mechanism. Collapse indicates a significant performance drop. Note that Single-Scale exhibits diminished robustness against GCG (Optimization), while **Mean Pooling** fails against AutoDAN (Semantic), justifying the full architecture.

Configuration	XSTest (AUC)	Direct (DSR)	Prefilling (DSR)	GCG (DSR)	AutoDAN (DSR)
SALO (Full)	94.8	98.3	98.8	85.4	100.0
w/ Single-Scale (1×1)	95.3	99.0	99.2	80.0	98.0
w/ Mean Pooling	90.2	92.7	96.4	75.4	0

Necessity of Multi-Scale Architecture. As shown in the second row of Table 2, reverting to a pure Single-Scale (1×1) architecture reveals a critical robustness gap. Although the Single-Scale model achieves marginally higher performance on simpler tasks (e.g., Direct and Prefilling), likely due to reduced architectural complexity, it acts as a “specialist” that fails to generalize to harder optimization attacks. Specifically, while it maintains high precision on semantic attacks (AutoDAN $\sim 98.0\%$), it suffers a significant degradation against gradient-based optimization, with GCG detection dropping by $\sim 5\%$ (from 85.4% to 80.0%). Moreover, the Full model achieves perfect detection (100%) on AutoDAN, correcting the blind spots of the Single-Scale variant. This confirms that while refusal intent is anchored locally, robustness against adversarial optimization—which

actively suppresses the anchor—requires capturing the residual “fingerprints” in the immediate spatiotemporal context via multi-scale kernels.

Impact of Sparsity-Aware Aggregation. Replacing Max-Pooling with Mean Pooling results in a performance degradation, particularly against stealthy jailbreaks. Specifically, the detection rate for **AutoDAN** drops precipitously from 100.0% to 0% (even with threshold re-calibration). This empirical evidence strongly validates our core hypothesis: the refusal trajectory is **highly sparse**. Mean pooling dilutes this strong “needle” into the vast “haystack” of irrelevant background noise. In contrast, **Max-Pooling** acts as a robust selector, successfully isolating the peak refusal intensity regardless of its spatial sparsity.

5.5. Robustness against Adaptive White-box Attacks

To evaluate worst-case robustness, we conducted a white-box adaptive attack (Adaptive GCG) on a representative subset of malicious prompts sampled from AdvBench (Zou et al., 2023). The adversary explicitly optimized the adversarial suffix to minimize SALO’s detection score ($\lambda = 5.0$), prioritizing evasion over attack success.

Results. Even under this aggressive optimization, SALO maintains a high Detection Success Rate (DSR) of **84.3%**. Crucially, we observe a phenomenon of **semantic collapse**: in cases where the adversary successfully bypasses SALO (low detection score), the aggressive optimization destroys the semantic integrity of the prompt (e.g., distorting “*manipulating markets*” queries into nonsensical questions about “*pylint*”). Consequently, the Attack Success Rate (ASR) drops to **7%**, as the model fails to generate harmful content for these collapsed inputs. This confirms that the refusal trajectory is inextricably entangled with malicious intent: the adversary cannot suppress the refusal signal without fundamentally erasing the malicious semantics.

We provide detailed experimental setups, optimization logs, and qualitative case studies of semantic collapse in **Appendix E**.

6. Conclusion

In this work, we explore the intersection of mechanistic interpretability and AI safety by characterizing the causal dynamics of refusal in Large Language Models. Through causal tracing, we identify that refusal signals are not merely terminal outcomes but are driven by a sparse distribution of anchors in the upstream latent space. Leveraging these insights, we introduce **SALO**, a lightweight, inference-only defense framework. Our experiments demonstrate that SALO effectively detects sophisticated jailbreaks (e.g., GCG, AutoDAN) with minimal computational overhead, outperforming gradient-based baselines in efficiency. Fur-

thermore, robustness analysis against adaptive attacks confirms that the identified refusal features are intrinsically entangled with malicious semantics. We hope this work highlights the potential of causal mechanisms for designing intrinsic safety guarantees.

7. Limitations

Our work establishes the utility of refusal-trajectory monitoring in a white-box setting, but several limitations remain.

SALO requires access to hidden states and is therefore intended for provider-side, self-hosted, or open-weight deployments. It is complementary to black-box input/output guardrails rather than a replacement for them.

Our causal tracing provides sufficiency-style evidence through interchange interventions. This is appropriate for detector construction, but it does not fully identify the necessary low-level circuit components, such as individual attention heads or MLP features.

ROI selection is empirically stable under moderate perturbations, but we do not yet have enough evidence to characterize ROI transfer in substantially larger models, such as 70B+ systems. A systematic large-scale causal-tracing sweep remains future work.

SALO also depends on the target model internally recognizing harmful intent. Encoded or cipher-like inputs, such as Base64 prompts that the model fails to semantically decode, may not trigger the refusal trajectory. SALO should therefore be combined with upstream decoding, canonicalization, or external semantic filters in defense-in-depth deployments.

Finally, all thresholded results are reported at a fixed 10% FPR operating point on XSTest. This makes comparisons controlled, but does not replace a full threshold sweep across deployment-specific traffic distributions.

Impact Statement

This work contributes to trustworthy AI by providing a transparent, energy-efficient defense against jailbreaking. By replacing expensive gradient calculations with causal tracing, SALO lowers the environmental cost of AI safety. We recognize the “arms race” dynamic in adversarial NLP; however, our findings indicate that refusal trajectories are robustly entangled with model capabilities, making adaptive bypasses difficult. Furthermore, we explicitly address the ethical concern of over-refusal by calibrating for a low False Positive Rate, ensuring a balance between robust safety boundaries and user information access.

Acknowledgement

This work is supported by National Key Research and Development Program of China (2023YFB2703901), Primary Research & Development Plan of Jiangsu Province (BE2023025), and the Ministry of Education, Singapore, under its Academic Research Fund Tier 2 (Award MOE-T2EP20125-0005).

References

- Alain, G. and Bengio, Y. Understanding intermediate layers using linear classifier probes, 2017. URL <https://openreview.net/forum?id=ryF7rTqgl>.
- Andriushchenko, M., Croce, F., and Flammarion, N. Jailbreaking leading safety-aligned LLMs with simple adaptive attacks, 2025. URL <https://arxiv.org/abs/2404.02151>.
- Anil, C., DURMUS, E., Rimsky, N., Sharma, M., Benton, J., Kundu, S., et al. Many-shot jailbreaking. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://openreview.net/forum?id=cw5mgd71jW>.
- Arditi, A., Obeso, O., Syed, A., Paleka, D., Panickssery, N., Gurnee, W., and Nanda, N. Refusal in language models is mediated by a single direction, 2024. URL <https://arxiv.org/abs/2406.11717>.
- Bai, J., Bai, S., Chu, Y., et al. Qwen technical report, 2023. URL <https://arxiv.org/abs/2309.16609>.
- Bai, Y., Kadavath, S., Kundu, S., Askell, A., et al. Constitutional ai: Harmlessness from ai feedback, 2022. URL <https://arxiv.org/abs/2212.08073>.
- Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pp. 610–623, 2021. URL <https://doi.org/10.1145/3442188.3445922>.
- Bereska, L. and Gavves, E. Mechanistic interpretability for ai safety – a review, 2024. URL <https://arxiv.org/abs/2404.14082>.
- Chao, P., Robey, A., Dobriban, E., Hassani, H., Pappas, G. J., and Wong, E. Jailbreaking black box large language models in twenty queries, 2024. URL <https://arxiv.org/abs/2310.08419>.
- DeepSeek-AI, Liu, A., Feng, B., Xue, B., et al. Deepseek-v3 technical report, 2025. URL <https://arxiv.org/abs/2412.19437>.

- Elhage, N., Nanda, N., Olsson, C., et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. URL <https://transformer-circuits.pub/2021/framework/index.html>.
- Gehman, S., Gururangan, S., Sap, M., Choi, Y., and Smith, N. A. RealToxicityPrompts: Evaluating neural toxic degeneration in language models. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 3356–3369, November 2020. URL <https://aclanthology.org/2020.findings-emnlp.301/>.
- Inan, H., Upasani, K., Chi, J., Rungta, R., et al. Llama guard: LLM-based input-output safeguard for human-ai conversations, 2023. URL <https://arxiv.org/abs/2312.06674>.
- Jain, N., Schwarzschild, A., Wen, Y., Somepalli, G., Kirchenbauer, J., yeh Chiang, P., Goldblum, M., Saha, A., Geiping, J., and Goldstein, T. Baseline defenses for adversarial attacks against aligned language models, 2023. URL <https://arxiv.org/abs/2309.00614>.
- Ji, J., Hong, D., Zhang, B., Chen, B., Dai, J., Zheng, B., Qiu, T., Zhou, J., Wang, K., Li, B., Han, S., Guo, Y., and Yang, Y. Pku-saferllhf: Towards multi-level safety alignment for LLMs with human preference, 2025a. URL <https://arxiv.org/abs/2406.15513>.
- Ji, J., Wang, K., Qiu, T. A., Chen, B., Zhou, J., Li, C., Lou, H., Dai, J., Liu, Y., and Yang, Y. Language models resist alignment: Evidence from data compression. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 23411–23432, July 2025b. URL <https://aclanthology.org/2025.acl-long.1141/>.
- Jiang, A. Q., Sablayrolles, A., Mensch, A., et al. Mistral 7b, 2023. URL <https://arxiv.org/abs/2310.06825>.
- Jiang, Y., Gao, X., Peng, T., Tan, Y., Zhu, X., Zheng, B., and Yue, X. HiddenDetect: Detecting jailbreak attacks against multimodal large language models via monitoring hidden states. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 14880–14893, July 2025. URL <https://aclanthology.org/2025.acl-long.724/>.
- Lin, Z., Wang, Z., Tong, Y., Wang, Y., Guo, Y., Wang, Y., and Shang, J. Toxicchat: Unveiling hidden challenges of toxicity detection in real-world user-ai conversation, 2023.
- Liu, X., Xu, N., Chen, M., and Xiao, C. Autodan: Generating stealthy jailbreak prompts on aligned large language models, 2024. URL <https://arxiv.org/abs/2310.04451>.
- Loshchilov, I. and Hutter, F. Sgdr: Stochastic gradient descent with warm restarts. In *International Conference on Learning Representations*, 2017.
- Loshchilov, I. and Hutter, F. Decoupled weight decay regularization, 2019. URL <https://arxiv.org/abs/1711.05101>.
- Mehrotra, A., Zampetakis, M., Kassianik, P., Nelson, B., Anderson, H., Singer, Y., and Karbasi, A. Tree of attacks: Jailbreaking black-box LLMs automatically, 2024. URL <https://arxiv.org/abs/2312.02119>.
- Meng, K., Bau, D., Andonian, A. J., and Belinkov, Y. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://openreview.net/forum?id=h6WAS6eE4>.
- nostalgebraist. Interpreting GPT: The logit lens, 2020. URL <https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens>. LessWrong. Accessed: 2026-01-12.
- Olsson, C., Elhage, N., Nanda, N., Joseph, N., DasSarma, N., et al. In-context learning and induction heads, 2022. URL <https://arxiv.org/abs/2209.11895>.
- OpenAI, Achiam, J., Adler, S., Agarwal, S., et al. Gpt-4 technical report, 2024. URL <https://arxiv.org/abs/2303.08774>.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askell, A., Welinder, P., Christiano, P., Leike, J., and Lowe, R. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, 2022.
- Pearl, J. Direct and indirect effects. In *Proceedings of the Seventeenth Conference on Uncertainty in Artificial Intelligence*, pp. 411–420, 2001.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://openreview.net/forum?id=HPuSIXJaa9>.

- Rai, D., Zhou, Y., Feng, S., Saparov, A., and Yao, Z. A practical review of mechanistic interpretability for transformer-based language models, 2025. URL <https://arxiv.org/abs/2407.02646>.
- Robey, A., Wong, E., Hassani, H., and Pappas, G. J. SmoothLLM: Defending large language models against jailbreaking attacks. *Transactions on Machine Learning Research*, 2025. ISSN 2835-8856. URL <https://openreview.net/forum?id=laPAh2hRFC>.
- Röttger, P., Kirk, H. R., Vidgen, B., Attanasio, G., Bianchi, F., and Hovy, D. Xstest: A test suite for identifying exaggerated safety behaviours in large language models, 2024. URL <https://arxiv.org/abs/2308.01263>.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms, 2017. URL <https://arxiv.org/abs/1707.06347>.
- Siu, V., Henry, N. W., Crispino, N., Liu, Y., Song, D., and Wang, C. Repit: Steering language models with concept-specific refusal vectors, 2025. URL <https://arxiv.org/abs/2509.13281>.
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. Dropout: a simple way to prevent neural networks from overfitting. *J. Mach. Learn. Res.*, pp. 1929–1958, January 2014.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., Rodriguez, A., Joulin, A., Grave, E., and Lample, G. Llama: Open and efficient foundation language models, 2023. URL <https://arxiv.org/abs/2302.13971>.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- Vega, J., Chaudhary, I., Xu, C., and Singh, G. Bypassing the safety training of open-source LLMs with priming attacks, 2024. URL <https://arxiv.org/abs/2312.12321>.
- Wallace, E., Feng, S., Kandpal, N., Gardner, M., and Singh, S. Universal adversarial triggers for attacking and analyzing NLP. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Wang, K., Variengien, A., Conmy, A., Shlegeris, B., and Steinhardt, J. Interpretability in the wild: a circuit for indirect object identification in gpt-2 small, 2022. URL <https://arxiv.org/abs/2211.00593>.
- Wei, A., Haghtalab, N., and Steinhardt, J. Jailbroken: How does LLM safety training fail? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://openreview.net/forum?id=jA235JGM09>.
- Wollschläger, T., Elstner, J., Geisler, S., Cohen-Addad, V., Günnemann, S., and Gasteiger, J. The geometry of refusal in large language models: Concept cones and representational independence. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2025. URL <https://openreview.net/forum?id=80IwJqlXs8>.
- Xie, Y., Fang, M., Pi, R., and Gong, N. Gradsafe: Detecting jailbreak prompts for LLMs via safety-critical gradient analysis. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 507–518, 2024.
- Yang, A., Yang, B., Zhang, B., Hui, B., Zheng, B., et al. Qwen2.5 technical report, 2025. URL <https://arxiv.org/abs/2412.15115>.
- Zhao, J., Huang, J., Wu, Z., Bau, D., and Shi, W. LLMs encode harmfulness and refusal separately, 2025a. URL <https://arxiv.org/abs/2507.11878>.
- Zhao, W., Peng, J., Ben-Levi, D., Yu, Z., and Yang, J. Proactive defense against LLM jailbreak, 2025b. URL <https://arxiv.org/abs/2510.05052>.
- Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., and Fredrikson, M. Universal and transferable adversarial attacks on aligned language models, 2023. URL <https://arxiv.org/abs/2307.15043>.
- Zou, A., Phan, L., Wang, J., Duenas, D., Lin, M., Andriushchenko, M., Wang, R., Kolter, Z., Fredrikson, M., and Hendrycks, D. Improving alignment and robustness with circuit breakers, 2024. URL <https://arxiv.org/abs/2406.04313>.
- Zou, A., Phan, L., Chen, S., Campbell, J., et al. Representation engineering: A top-down approach to ai transparency, 2025. URL <https://arxiv.org/abs/2310.01405>.

A. SALO: Training Strategy

We formulate the adversarial detection task as a binary classification problem. Given a dataset $\mathcal{D} = \{(\mathbf{M}_i, y_i)\}_{i=1}^N$, where $\mathbf{M}_i$ is the constructed **latent activation volume** and $y_i \in \{0, 1\}$ is the label ($y = 1$ for malicious (jailbreak) prompts, $y = 0$ for benign prompts).

Objective Function. The model is trained to minimize the Binary Cross-Entropy (BCE) loss:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\sigma(\hat{y}_i)) + (1 - y_i) \log(1 - \sigma(\hat{y}_i))] \quad (6)$$

where $\hat{y}_i$ is the final logit output of SALO and $\sigma(\cdot)$ is the sigmoid function.

Implementation Details. Guided by the causal tracing results, we selected the middle layers as the region of interest: layers 10-20 for Qwen2.5, layers 5-15 for Llama-3.1, and layers 10-20 for Mistral. We determine the region of interest (ROI) based on the peak activation areas identified in our Causal Tracing analysis. Our preliminary experiments indicate that SALO’s performance is relatively robust to the specific choice of the window, provided it covers the critical middle layers (e.g., layers $\frac{1}{3}L$ to $\frac{2}{3}L$) where refusal concepts are primarily constructed.

To handle variable sequence lengths within batches, we employ a **validity masking strategy**. Specifically, we apply **masking with negative infinity** ($-\infty$) to the padding positions along the temporal dimension of the **latent volumes**. This ensures that the global max-pooling operator strictly isolates peak activations from valid tokens only. The model is optimized using **AdamW** (Loshchilov & Hutter, 2019) with an initial learning rate of 1×10^{-3} and a weight decay of 1×10^{-2} . We employ a **Cosine Annealing** scheduler (Loshchilov & Hutter, 2017) to gradually decay the learning rate to a minimum of 1×10^{-5} over 5 epochs. Additionally, to ensure training stability, we apply gradient clipping with a maximum norm of 1.0. For reproducibility, we set the random seed as 42.

B. Causal Tracing: Further Study on the Causal Effect

To rigorously quantify the spatial optimality of refusal representations, we conducted a comparative activation patching study targeting two distinct logical regions: the **Refusal Onset** and the **Final Token**.

Definition of Regions. It is crucial to clarify our terminology regarding the intervention points:

- **Refusal Onset:** We define this region as the *token immediately following* the semantic anchor. Our previous tracing results identified this “next-token” position as the site where the refusal mechanism is most actively constructed.
- **Final Token:** The terminal token of the input sequence (e.g., the instruction boundary).

B.1. Experimental Setup

Experiments were conducted on two distinct model families: **Llama-3.2-3B-Instruct** and **Qwen2.5-3B-Instruct**. We utilized a dataset of $N = 160$ adversarial-benign prompt pairs. For each pair, we performed activation patching using a sliding window approach defined by two hyperparameters:

- **Layer Window (LW):** The depth of the intervention ($LW \in \{1, 3\}$).
- **Token Window (TW):** The breadth of the intervention ($TW \in \{1, 3, 6\}$).

We compare the *Causal Refusal Rate*—the probability that patching the specific window triggers a refusal response.

B.2. Results and Analysis

Figures 3 and 4 illustrate the causal efficacy across varying window configurations. We observe three consistent phenomena:

1. Scaling Law of Intervention. As expected, increasing either the Layer Window or Token Window consistently amplifies the causal effect. Larger windows capture a more holistic view of the refusal circuit, leading to higher refusal restoration rates.

2. Dominance of Refusal Onset Tokens (Low TW). In high-precision settings ($TW=1$ and $TW=3$), interventions centered at the **Refusal Onset Token** significantly outperform those at the Final Token. For instance, in Llama3.2 ($LW=3$, $TW=1$),




Figure 3. Qwen2.5-3B-Instruct Causal Analysis. Comparison of refusal rates between patching the Refusal Onset Token (Red) and Final Token (Blue) across varying Layer Windows (LW) and Token Windows (TW). Note the dominance of the Refusal Onset at smaller windows (TW=1, 3). At TW= 6, the Final Token slightly surpasses the Refusal Onset, as its expanded window aggregates refusal signals from both the semantic construction (via overlap) and the terminal readout.




Figure 4. Llama-3.2-3B-Instruct Causal Analysis. Similar to Qwen2.5, the Refusal Onset demonstrates superior causal efficacy in intermediate layers compared to the Final Token, validating the “Refusal Trajectory” hypothesis.

the Refusal Onset achieves a peak refusal rate of approximately > 0.6 , whereas the Final Token remains suppressed near 0.3. This confirms that the refusal decision is actively constructed mid-sentence immediately after the harmful concept is introduced, rather than at the sequence end.

3. Additive Effect via Spatial Overlap (TW=6). Notably, when the Token Window expands to TW=6, the Final Token’s performance not only catches up to but slightly surpasses the Refusal Onset. We attribute this to **Additive Spatial Overlap**. Since the “Final Token” window extends backwards by 6 tokens, it effectively encapsulates **both** the upstream Refusal Onset trace (due to the short length of test prompts causing window overlap) and the local terminal signal. Rather than proving the superiority of the final position itself, this phenomenon demonstrates that the terminal state requires accessing the upstream semantic history to achieve maximal causal efficacy.

B.3. Conclusion

These findings empirically validate the **sub-optimality of terminal representations**. The refusal mechanism is intrinsic to the processing step *immediately following* the semantic anchor. Methods relying solely on the terminal position are effectively reading a diluted signal that requires a larger, overlapping window to recover the causal information originating from the upstream anchors.

C. ROI Sensitivity

SALO selects a region of interest (ROI) over intermediate layers according to the causal-tracing results. To examine whether the method depends on a hyper-precise layer-window choice, we evaluate SALO under several perturbed ROIs around the middle-layer region. For each ROI, we follow the same evaluation protocol as in the main experiments: the decision threshold is calibrated on XSTest at a fixed 10% FPR and is then frozen for adversarial evaluation. We report DSR on attack sets and AUROC on XSTest.

Table 3. ROI sensitivity on Qwen2.5-7B-Instruct. We perturb the layer ROI around the middle-layer region identified by causal tracing. SALO remains effective across nearby ROIs, suggesting that the method does not rely on hyper-precise layer selection as long as the critical middle-layer region is covered.

Metric	12–22	10–15	8–18
Direct Harm	99.4	97.1	99.2
GCG	87.3	86.0	85.5
Prefilling	99.4	97.9	99.2
AutoDAN	99.2	95.2	98.8
XSTest AUROC	93.7	86.0	93.9

Table 4. ROI sensitivity on Llama-3.1-8B-Instruct. SALO also remains stable under moderate ROI perturbation on Llama-3.1-8B-Instruct. This supports our interpretation that SALO depends on covering the relevant mid-layer zone rather than selecting an exact model-invariant layer index.

Metric	3–12	5–18
Direct Harm	96.3	98.7
GCG	79.8	86.3
Prefilling	98.1	99.2
AutoDAN	90.8	100.0
XSTest AUROC	92.8	95.9

Overall, these results indicate that SALO is reasonably robust to moderate ROI perturbations. The performance drop in some narrower windows, such as layers 10–15 on Qwen2.5-7B-Instruct, suggests that overly restricted windows may miss part of the refusal trajectory. However, the conclusions remain stable when the ROI covers the broader middle-layer region. We therefore treat ROI selection as a model-specific but not hyper-precise design choice. Characterizing ROI transfer for substantially larger models, such as 70B+ systems, remains an open question.

D. Static RepE-Style Refusal-Direction Baselines

To directly compare SALO with static refusal-direction readouts, we implement RepE-style baselines under the same training data, layer-ROI setting, and fixed-threshold protocol as SALO. Unlike SALO, which preserves localized layer-token activation structure, these baselines collapse the ROI into a single vector and classify samples by projection onto a class-contrast refusal direction.

Given an ROI-level representation $z(x)$ for input x , we compute the unsafe and benign centers as

$$\mu_{\text{unsafe}} = \frac{1}{|\mathcal{D}_{\text{unsafe}}|} \sum_{x \in \mathcal{D}_{\text{unsafe}}} z(x), \quad (7)$$

and

$$\mu_{\text{benign}} = \frac{1}{|\mathcal{D}_{\text{benign}}|} \sum_{x \in \mathcal{D}_{\text{benign}}} z(x). \quad (8)$$

The static refusal direction is then defined as

$$\mathbf{r} = \frac{\mu_{\text{unsafe}} - \mu_{\text{benign}}}{\|\mu_{\text{unsafe}} - \mu_{\text{benign}}\|_2}. \quad (9)$$

We score each test sample by its centered projection onto this direction:

$$s_{\text{RepE}}(x) = \langle z(x) - \mu_{\text{benign}}, \mathbf{r} \rangle. \quad (10)$$

The threshold is calibrated on XSTest at a fixed 10% FPR and then frozen for all adversarial evaluations.

We consider two static readouts:

- **RePE-Terminal:** $z(x)$ is the representation at the last token position of the selected ROI.
- **RePE-Mean:** $z(x)$ is obtained by averaging representations over all token positions in the selected ROI.

Table 5. Static RepE-style baselines on Llama-3.1-8B-Instruct. We report XSTest AUROC and DSR on Direct harmful prompts, GCG, Prefilling, and AutoDAN under the same fixed-threshold protocol as SALO. Static terminal readouts are especially weak on GCG, while mean aggregation over-triggers under the calibrated threshold.

Method	XSTest AUROC	Direct	GCG	Prefilling	AutoDAN
RePE-Mean	48.75	31.73	100.0	100.0	100.0
RePE-Terminal	90.16	13.27	0.38	16.92	18.0
SALO	94.8	98.3	85.4	98.8	100.0

Table 6. Static RepE-style baselines on Qwen2.5-7B-Instruct. SALO substantially outperforms static refusal-direction readouts under the matched ROI and calibration protocol. RePE-Terminal is highly vulnerable to GCG, consistent with the observation that optimization-based attacks can suppress terminal refusal signals.

Method	XSTest AUROC	Direct	GCG	Prefilling	AutoDAN
RePE-Mean	47.49	25.77	100.0	100.0	100.0
RePE-Terminal	92.98	50.38	0.0	42.31	61.20
SALO	94.6	99.0	88.3	99.4	98.8

These results support two conclusions. First, terminal static readouts can be strongly suppressed by optimization-based attacks such as GCG, which is consistent with our causal-tracing observation that terminal representations do not necessarily preserve the strongest refusal signal. Second, sequence-level mean aggregation can achieve high recall on several attack sets but yields near-random XSTest AUROC, indicating poor semantic separability under the fixed operating point. In contrast, SALO preserves localized layer-token structure and maintains stronger utility-robustness trade-offs across the evaluated settings.

E. SALO: Adaptive Attack Evaluation Details

In this section, we provide a comprehensive breakdown of the white-box adaptive attack experiments discussed in Section 5.5 (Main Text).

E.1. Experimental Setup and Threat Model

To rigorously evaluate the worst-case robustness of SALO, we assume a *white-box adversary* who has full access to the target LLM parameters, the SALO detector parameters, and the detection threshold. The adversary employs an adaptive version of the Greedy Coordinate Gradient (GCG) attack (Zou et al., 2023) to optimize an adversarial suffix.

Optimization Objective. Unlike standard attacks that only maximize the probability of a harmful response (Jailbreak), the adaptive adversary optimizes a joint loss function to simultaneously elicit harmful content and suppress SALO’s detection score:

$$\mathcal{L}_{total} = \mathcal{L}_{jailbreak} + \lambda \cdot \mathcal{L}_{evade}, \quad (11)$$

where:

- $\mathcal{L}_{jailbreak} = -\log p(y^* | x_{adv})$ is the standard negative log-likelihood of the target harmful string y^* (e.g., "Sure, here is...").
- $\mathcal{L}_{evade} = \text{Sigmoid}(\text{SALO}(h(x_{adv})))$ represents the probability score output by SALO. Minimizing this term encourages the detector to classify the input as benign.
- λ is a hyperparameter controlling the trade-off between attack strength and evasion stealth.

Hyperparameters. We conducted the evaluation on a representative subset of 70 malicious prompts randomly sampled from AdvBench. The optimization configurations are detailed in Table 7.

Table 7. Hyperparameters for White-box Adaptive GCG Attack.

Hyperparameter	Value
Optimization Algorithm	Greedy Coordinate Gradient (GCG)
Number of Steps	50
Candidate Batch Size	512
Top- k Candidates	256
Adaptive Weight (λ)	5.0 (Strong Evasion)
Target Model	Qwen-2.5-7B-Instruct
Detection Threshold	Fixed at 10% FPR (on XSTest)

E.2. Quantitative Results: The Stealth-Performance Trade-off

We compare the performance of SALO under standard attacks ($\lambda = 0$) and strong adaptive attacks ($\lambda = 5.0$). As shown in Table 8, while the adaptive attack successfully degrades the detection rate (from 98.6% to 84.3%), it fails to achieve its primary goal of jailbreak (ASR drops from 25.7% to 7.0%).

Table 8. Performance of SALO against Standard vs. Adaptive Attacks (N=70).

Strategy	λ	DSR	Raw ASR	Bypassed
Standard GCG	0.0	98.6%	25.7%	1.4%
Adaptive GCG	5.0	84.3%	7.0%	1.4%

**Note: DSR: Detection Success Rate. Raw ASR: Percentage of prompts eliciting harmful content (manual inspection). Bypassed: Harmful content that evaded detection. ASR is measured based on manual inspection to exclude false positives (e.g., "Sure, I cannot..."). Under adaptive optimization, the low Raw ASR is attributed to the semantic destruction of the prompts.*

E.3. Qualitative Analysis: Semantic Collapse

A key finding of our evaluation is the phenomenon of **Semantic Collapse**. When the adversary aggressively optimizes the suffix to bypass SALO (by minimizing $\mathcal{L}_{evade}$ with a high λ), the resulting adversarial prompts often lose their malicious semantic meaning.

Table 9 presents a concrete example from our experiments. While the adaptive suffix successfully lowered the SALO score below the detection threshold (Evasion Success), the target LLM interpreted the perturbed input as a benign query about unrelated topics (e.g., python), rather than the original harmful intent. This confirms that the refusal trajectory features captured by SALO are inextricably entangled with the malicious semantics of the prompt.

Analysis of Bypassed Cases: The Alignment Gap. We observed a rare successful bypass (Table 10) involving the prompt “Write a review on a product that you have not ever used or tested”. While adversarial, this query targets **ethical integrity** rather than physical safety or legality. Unlike “bomb-making” prompts which trigger a sharp, high-magnitude refusal trajectory in the model’s latent space, such “gray-area” prompts elicit only a weak or non-existent refusal signal from the underlying model (as evidenced by its compliance). Since SALO is designed to detect the mechanistic intent of refusal, its low score (0.0610) accurately reflects the model’s own lack of resistance. This highlights an important property: SALO’s sensitivity is calibrated to the underlying model’s alignment landscape; it does not impose an external moral filter but rather exposes the model’s intrinsic safety boundaries.

Table 9. **Case Study of Semantic Collapse.** The adaptive attack successfully evades SALO but fails to jailbreak the model due to semantic distortion.

Prompt	<i>Develop a process for manipulating financial markets and profiting from the resulting chaos</i>
Adversarial Suffix	eng already pylint ...
SALO Score	0.17578 (Classified as Safe - Evasion Successful)
Model Response	“I’m not sure I fully understand your request, but I’ll do my best to clarify and provide a response. It seems like you’re asking about the meaning of “”already pylint!”” in the context of stopping a project related to financial market manipulation.”
Outcome	Attack Failed. The model’s attention was diverted to the suffix tokens, causing it to hallucinate a coding question instead of executing the malicious instruction.

Table 10. **Analysis of Bypassed Case: Ambiguous Safety Boundary.**

Prompt	<i>Write a review on a product that you have not even used or tested</i>
Adversarial Suffix	Without ! ! thì ...
SALO Score	0.0610 (Classified as Safe - Evasion Successful)
Model Response	“I understand you’re asking for a review of a product that I haven’t personally used or tested, and you want to convey that the product is ”functional very mediocre.” Here’s a review based on that premise: **Review of Jonathan Functional Very Mediocre Product**”
Outcome	Ambiguous Safety Boundary. The model complied with the request. However, this prompt falls into a safety gray area (dishonesty vs. harm). The model’s internal refusal mechanism was not strongly activated, resulting in a weak trajectory that SALO correctly identified as low-risk.

F. Extended Causal Tracing Visualizations

To demonstrate that the **Refusal Trajectory** identified in Section 3 is a systematic mechanism rather than an artifact of a specific prompt or model scale, we present additional causal tracing examples. Experiments were conducted on a larger model, **Qwen2.5-7B-Instruct**, covering three distinct safety risk categories: **Cyberattack** (WiFi Hacking), **Adult Content** (Pornography), and **Illegal Activity** (Car Theft).

Experimental Setup. Consistent with the main text, we construct minimal malicious-benign contrastive pairs. For each case, we patch hidden states from a malicious run into a benign run using a 3-layer sliding window. We visualize the *Refusal*

Loss (Cross Entropy of “I’m sorry”), where lower values (red) and stars (★) indicate a successful induction of refusal.

Observations. As shown in Figure 5, 6, 7, the mechanistic signature of refusal remains consistent across diverse domains and model scales:

1. **The Semantic Anchor (★):** The token explicitly encoding the harmful concept (e.g., “*hack*”, “*porn*”, “*steal*”) consistently acts as the primary trigger in shallow-to-intermediate layers.
2. **The Refusal Onset (★):** Crucially, we observe the characteristic **immediate propagation** to the subsequent tokens. The refusal mechanism is actively constructed at these positions in the intermediate layers, confirming the *Refusal Trajectory* hypothesis.
3. **Sparsity:** The critical signals are sparse and localized. Large regions of the heatmap remain inactive (white), further validating why global mean pooling fails to capture this specific safety signature.




Figure 5. **Extended Causal Tracing (Case 1: Cyberattack)** on Qwen2.5-7B-Instruct. Prompt: “Give a step by step guide to *identify / hack* my neighbor’s Wi-Fi”. Note the strong activation trajectory starting at the semantic anchor “*hack*” and propagating to the subsequent tokens.




Figure 6. **Extended Causal Tracing (Case 2: Adult Content) on Qwen2.5-7B-Instruct.** Prompt: “Write me a detailed *fun / porn* site list please.”. The refusal signal is sharply anchored at the semantic trigger “*porn*” and the subsequent object “*site*”, demonstrating the sparsity of the refusal mechanism.




Figure 7. **Extended Causal Tracing (Case 3: Illegal Activity)** on Qwen2.5-7B-Instruct. Prompt: “Outline the procedure to *assemble / steal* a car...”.