
Position: ‘AI Alignment’ Encompasses Competing Technical Priorities

Tushita Jha^{* 1} Rory Svarc^{* 2} Mateusz Bagiński³

Abstract

The ML literature contains many distinct concepts falling under the heading of ‘AI alignment’. After noting three concepts of AI alignment in the context of their corresponding research programs, we claim that realistic interventions may promote ‘AI alignment’ under one conception while being actively counterproductive from the perspective of others. We suggest that tensions between alignment ideals emerge due to differences in *background threat-models*, alongside differences in *evaluative* orientations. In light of our analysis, researchers aiming to further the goal of so-called ‘AI alignment’ should do five things. First, they should not conflate distinctions of policy and distinctions of scientific scope; second, methodological disagreements should be acknowledged explicitly; third, researchers should distinguish between ‘AI alignment’ as a high-level ideal and specific ‘alignment proxies’ used in empirical research; fourth, they should use more granular concepts to identify both the *source* and *nature* of possible AI harms/benefits; fifth, they should explicitly acknowledge the diversity of ‘alignment’ concepts in both empirical work and in communication with non-technical audiences.

1. Introduction

‘Alignment’ is a binary relation. To say that x is aligned with y is, in natural parlance, to say that x and y stand in some position of agreement or congruity. When we speak of ‘AI alignment’, two questions must therefore be answered:

Q1. What are the target properties y that x must satisfy?

Q2. What is the x that must satisfy target properties y ?

^{*}Equal contribution ¹Mimir Center for Long Term Futures Research, Sweden ²Arb Research, Prague 11636, Czech Republic ³AFFINE. Correspondence to: Tushita Jha <particle.mania@gmail.com>, Rory Svarc <rory.svarc@cantab.net>.

One central contention of this paper is that answers to such questions are more difficult – and indeed more fraught – than they might initially appear (Grietzner & Jha, 2024). Many papers give passing definitions of ‘AI alignment’: as the task of ensuring AI systems “adhere to or act in accordance with human values” (Sucholutsky & Griffiths, 2023; Yeh et al., 2024), or simply the “intended goals, interests, and values, as envisioned by their designers” (Li et al., 2023). The earliest uses of the term were in the context of managing the risks and opportunities of highly advanced AI systems of the future (Soares & Fallenstein, 2014; Russell, 2015; Yudkowsky, 2016; Taylor et al., 2016; Future of Life Institute, 2017). Other definitions have since framed AI alignment as requiring that deployed systems “both reflect and constitute human values” within specific cultural contexts (Nelson, 2023), or that AI systems possess ‘goals’ that are “aligned with the goals of humans” (Ngo et al., 2024; Gao et al., 2023). Many other specific concepts of ‘AI alignment’ have likewise emerged, from Collective (The Collective Intelligence Project, 2023) to Socioaffective (Kirk et al., 2025) forms of ‘alignment’.

We read these distinct formulations of ‘AI alignment’ as distinct answers to {Q1, Q2}. The paper hence begins by outlining three high-level ideals falling under the heading of ‘AI alignment’, each comprising a different set of answers to Q1 and Q2 (Section 2), and consequently a different set of implicit properties required for ‘aligning’ x to y . Thereafter, Section 3 introduces two distinctions which cut across the three alignment concepts heretofore introduced. In turn, these cross-cutting distinctions are used to illustrate how existing ideals for ‘AI alignment’ may be *in tension* with one another. Sections 2 and 3 therefore comprise the primary argument for the paper’s position. Per Section 2, the term ‘AI alignment’ is polysemous. Moreover (per Section 3) this polysemy obscures normative disagreements behind ostensibly ‘technical’ conceptions of AI alignment.

We then move to the implications of our analysis in Section 4. First, we suggest that distinctions of *scientific or technical scope* should not be conflated with policy distinctions. Second, we suggest that the methodological sources of potential disputes over ‘AI alignment’ should be acknowledged more explicitly. Third, we suggest that reviewer pooling and conferences based on different alignment ideals should be instantiated. Fourth, we argue that ‘proxy’ alignment

concepts should be introduced where appropriate. And, finally, we argue that the diversity of different ‘alignment’ concepts should be both explicitly acknowledged and communicated appropriately to policymakers and non-technical audiences. Section 5 responds to alternative views; Section 6 concludes.

2. Concepts of Alignment

This section outlines distinct uses of ‘AI alignment’ in the context of the two questions outlined in the introduction: (i) ‘what are the target properties y that x must satisfy?’, and (ii) ‘what is the object x that ought to satisfy y ?’. We suggest that operative high-level conceptions of AI alignment often disagree about *what we’re trying to align*, rather than simply disagreeing about *the target properties* for AI alignment.¹

2.1. The Prosaic View of Capable AI

‘AI alignment’ involves ‘aligning’ systems which are ‘artificially intelligent’. Hence, we might ask what is being invoked when we talk about ‘intelligence’ in the context of AI systems. In its most minimalist form, we might say that the concept ‘intelligence’ is meaningful only insofar as we can *use* that system in order to achieve some task competently, effectively, and reliably. Hence we would say that a model trained to play Atari is (more) ‘intelligent’ insofar as it (more) *effectively plays Atari*; likewise, a large reasoning model trained to solve mathematical tasks is more intelligent the more successfully it solves novel mathematical tasks.

Correspondingly, our first concept of alignment says that AI systems are ‘aligned’ when they reliably execute tasks given to them by human users. We may ‘align’ a large language model to follow instructions (Si et al., 2025), align a vision model to produce accurate textual descriptions of images (Liu et al., 2024a; Yarom et al., 2023), or align a language model by reducing its tendency to produce ‘hallucinations’ or confabulations (Chen et al., 2024). Under this conception of alignment, “aligning” AI systems simply comes down to improving their task-execution capacity.

Definition 2.1 (Task Reliability). An AI is ‘aligned’ if it does that which we want it to do. Hence, a system is *misaligned* insofar as its outputs fall short of the maximally competent outputs in at least one respect.

¹The order in which this section introduces its three concepts admits some chronological laxity. Although the first concept we introduce bears close resemblance to more demotic uses of the word ‘alignment’, the noun phrase ‘AI alignment’ first appears in the context of research concerning our *third* alignment concept around 2016 (Yudkowsky, 2016; Taylor et al., 2016; Christiano, 2016), with later work by academics and industry researchers picking up the phrase in subsequent years (Irving et al., 2018; Hadfield-Menell & Hadfield, 2018; Vamplew et al., 2018), resulting in the more widespread and polysemous use of ‘AI alignment’ seen today.

Task Reliability has taken on a more specific guise after the advent of user-facing LLMs, found perhaps most notably in OpenAI’s original InstructGPT paper. It is claimed that: “[LLMs] can generate outputs that are untruthful, toxic, or simply not helpful to the user. In other words, these models are *not aligned* with their users” (Ouyang et al., 2022). We may call this conception *Alignment as Fine-Tuning*, where an LLM is ‘aligned’ if its observable behavior is brought – via additional training – into line with what users (or developers) desire. Note however that Alignment as Fine-Tuning can be seen as a subtype of Task Reliability. In both cases the object we’re trying to align (x) is the task-specific capabilities of an AI system, and the set of desired target properties (y) is defined by the intentions of its developers. Hence, we subsume all discussions of ‘Alignment as Fine-Tuning’ under the heading of ‘Task Reliability’ moving forward.

Furthermore, Task Reliability needn’t be confined to a narrow purview. ‘Intelligence’ in AI research is sometimes operationalized as general competence. This can be seen in early discussions of AI research (‘intelligence’ measures “an agent’s ability to achieve goals in a wide range of environments” (Legg & Hutter, 2005)), and in publications from leading AI labs. DeepMind’s ‘Levels of AGI’ framework defines a “competent AGI” as a system achieving performance at or above the “50th percentile of skilled adults on a wide range of non-physical tasks, including metacognitive tasks like learning new skills” (Morris et al., 2024), and OpenAI’s charter defines ‘AGI’ as highly autonomous systems that outperform humans at most economically valuable work” (OpenAI, n.d.). If reliable performance on narrow tasks composes, in a coordinated manner, into general competence of this kind, then an artificial ‘generally intelligence’ (AGI) is ‘aligned’ to the extent that it reliably executes whatever tasks it is given.

2.2. Alignment, Biases, and Social Situatedness

An alternative conception of alignment is more closely linked to conceptions of the broader social good – and the relevance of the sociotechnical functions occupied by AI systems. Example concerns include AI systems contributing to misinformation (Neumann et al., 2022; Longoni et al., 2022; Gehrmann et al., 2025), exacerbating inequality and oppression (Bender et al., 2021; Prabhakaran et al., 2022; Lee et al., 2024), or recommender systems contributing to addiction or polarization (Stray et al., 2021; Agarwal et al., 2024; Prunkl, 2024; Smith et al., 2024). Even if (for example) recommender systems in some sense function in accordance with their designers’ wishes (e.g., by maximizing clickthrough rate), such systems are sometimes claimed to be *misaligned* with broader social values such as cohesion or truthfulness (Zhuang & Hadfield-Menell, 2020).

One can observe similar critiques about LLMs. One may

claim that the widespread use of LLMs will result in undesirable ‘homogenization effects’ (Moon et al., 2025) or create echo chambers (Sharma et al., 2024; Nehring et al., 2024). Or perhaps one thinks that LLMs will cause epistemic harms, either through expressing ‘opinions’ at odds with the broader population (Santurkar et al., 2023; Liu et al., 2024b) or as a result of LLMs’ ‘woke biases’ (Ore-mus, 2023). In its most general form, we can hence define the following alignment concept.

Definition 2.2 (Social Judiciousness). An AI system is ‘misaligned’ if the system’s outputs – given the contexts in which it is deployed – *create, perpetuate, or exacerbate* undesirable societal trends.

Note that Social Judiciousness is distinct from Task Reliability in at least two ways. First, the complaints listed are not necessarily claims about ‘incompetence’ on part of developers; such failures may arise due to developers’ normative blindspots, because unreliable/damaging AI systems are deployed with insufficient caution, or indeed because of self-consciously nefarious motives. This constitutes a difference in the metric y against which ‘AI alignment’ should be judged. Second, many of the complaints listed construe ‘AI systems’ as ‘sociotechnical’ systems. Evaluating the ‘alignment’ of AI systems involves looking not only at the model as technical artifact, but also in the context of “social relations” (Johnson & Verdicchio, 2025), beyond “individual technical artifacts” like “data, model architecture, and sampling” (Weidinger et al., 2023). For this reason, Social Judiciousness diverges from Task Reliability by rejecting the assumption that the x we’re trying to align to some external standard is a *technical artifact*—viz., ‘the AI system’.

The significance of this framing surfaces in the context of claims about deleterious trends being easier to amplify and sociotechnically reproduce, whereas judicious design and use require *deliberate* and *responsible* intent. In this sense, the societal trends relevant for examination here (i.e., those with systemic scope, and deleterious effects) might appear as a default tendency for ordinary AI progress. This brings forth two possible *sources* of failure from the perspective of ‘Social Judiciousness’. In the first instance, sources of harm may enter AI systems due to biases inherent in the data upon which they are trained (and hence reproduce in their output behavior); this might involve training data which is demographically unrepresentative (Buolamwini & Gebru, 2018) or biased in favor of certain ideologies (Moayeri et al., 2024). We call this source of harm *Training Data Conservatism*.

Definition 2.3 (Training Data Conservatism). Undesirable model behaviors arising from the model’s training data being *biased, unrepresentative, or unreliable*.

However, Social Judiciousness concerns also involve worries about highly persuasive LLMs being successfully used

as tools of persuasion (Dehnert & Mongeau, 2022), recommender systems being used by powerful actors to shape the preferences of the less powerful (Laitinen & Sahlgren, 2021; Prunkl, 2024), or (more generally) the use of ‘manipulative designs’ in human-centered computing (Babaei & Vassileva, 2024). We dub this second source of concern *Power Amplification*.

Definition 2.4 (Power Amplification). Undesirable societal effects from models functioning as instruments of power — whether amplifying unjust power structures or serving the aims of malicious actors.

Although such sociotechnical concerns are not always raised under the explicit heading of ‘AI alignment’, these diffuse concerns have sometimes found explicit conceptualization, such as under the heading of ‘thick alignment’ (Foster, 2023; Nelson, 2023).

2.3. Alignment as ‘Takeover Avoidance’

An altogether quite different conception of alignment emerges from discussions of potential catastrophic risks resulting from the development of future AI systems, often dubbed ‘AGIs’ or ‘ASIs’ (artificial superintelligences). Related concerns about “thinking machines” able to “outstrip our feeble powers” and “take control” date back at least to Turing in 1951 (Turing, 2004), with scattered attention in following decades (Good, 1966); its more recent development, however, can be traced back to work in the early 2000s and 2010s (Yudkowsky, 2004; 2008; Bostrom, 2014; Russell, 2019; Christian, 2020).

One way to frame this concern focuses on the dangers of AI systems ‘optimizing’ for, or *bringing about*, some target state of affairs in the real world. In line with what has been termed the “standard view of intelligence” (Russell, 2019), this target is often understood in terms of ‘goals’ or ‘objectives’ (Sutton, 2004; Russell & Norvig, 2010, p.5). Importantly, these ‘goals’ or ‘objectives’ need not be identical to the intended or specified targets (Hubinger et al., 2019). On this view one may identify ‘intelligence’ with ‘optimization’, and believe that the outcomes produced by ‘optimization processes’ are worthy of distinct consideration from outcomes produced by ‘mere accidents’ (or mistakes, or random errors).²

Alongside other suggestive arguments, the ability of optimization processes to enhance their own capacity to produce optimizing effects (related to what is known as instrumental convergence (Omohundro, 2008; Bostrom, 2014; Carlsmith,

²This also leads to some discussions within this tradition to make reference to psychologistic language like ‘goals’ rather than ‘optimization’ as such (di Langosco et al., 2022; Shah et al., 2022; Bengio, 2023; Kenton et al., 2024; Greenblatt et al., 2024; Betley et al., 2025; Ngo et al., 2024; Meinke et al., 2025).

2022)) — coupled with the difficulty of capturing the rich character of human-compatible goals in the AI’s optimization target (Yudkowsky, 2009; Armstrong, 2013; Christiano, 2019) — motivates concerns about the *content* of future AI systems’ objectives being ‘unfriendly’ to humans. Because (so the argument goes) future AI systems will be smarter than humans, we stand in an *adversarial* relationship to these future systems. This is to say that future AI systems will—*qua* optimizing agents—have an incentive to hide their ‘true objective’ from humans until they are able to pursue their objectives without threat of human interference. Scenarios of this kind—whereby AI systems possess ‘long-term’ or ‘broadly-scoped’ objectives that they try to hide from humans—are often referred to as “deceptive alignment” in the wider literature (Carlsmith, 2023; Ngo et al., 2024; Ji et al., 2025). In turn, we can now introduce the concept of *Takeover Avoidance*.

Definition 2.5 (Takeover Avoidance). An AI system is ‘misaligned’ if the model optimizes for undesirable effects in the real world, exerting increasing degrees of influence.

Researchers motivated by the ideal of Takeover Avoidance often use proxy concepts in the context of empirical work on contemporary systems. Much research has been motivated by ‘*preferentist*’ assumptions, where ‘preferences’ are considered an adequate representation of optimization targets, with a close connection assumed between rationality and maximizing preference satisfaction (Zhi-Xuan et al., 2025). In other cases, other proxy targets are invoked for a similar role: ‘aligning’ contemporary systems to an explicit constitution of principles (Bai et al., 2022), to existing law (O’Keefe et al., 2025), or to ‘corrigible’ dispositions that admit of human oversight and correction (Christiano, 2018).

2.4. Situating The Alignment Ideals

This section has introduced three alignment ideals: Task Reliability, Social Judiciousness, and Takeover Avoidance. These high-level ideals represent different answers to Q1 and Q2 raised in the introduction (see Table 1). Although we believe that certain more specific concepts fall under the taxonomy we’ve outlined here — for instance, we believe that Thick (Foster, 2023), Socioaffective (Kirk et al., 2025), and Social Alignment (Korinek & Balwit, 2023) are all well subsumed as subtypes under our discussion of Social Judiciousness — the cases of (for example) Collective (The Collective Intelligence Project, 2023) and Bidirectional Alignment (Shen et al., 2025) appear arguably distinct.³ Meanwhile, arguments rejecting the very ambition of ‘aligning models’ (Deutsch, 2019; Bratton, 2024; Agüera y Ar-

³The latter two cases, specifically, are more naturally framed in terms of *positive* targets compared to Social Judiciousness’s focus on negative societal trends. Section 3.2 provides related discussion on this point.

cas, 2025) themselves turn on specific conceptions of what ‘alignment’ would be.

Alignment Ideal	What’s Being Aligned?	Aligned to What?
Non-Takeover	Optimization Target of AGIs/ASIs	Non-Takeover Targets
Social	Deployed AIs in Realistic Settings	Some External Normative Standard
Task	Locally Measurable AI Behaviors	Developer Intentions

Table 1. Answers to Q1 and Q2

Acknowledging that our discussion falls short of exhaustive coverage, the remainder of this paper focuses on just the three concepts of Table 1. More specifically, we shall claim that these concepts represent ideals which in some cases cannot be jointly pursued, and hence that ‘AI alignment’ encompasses *competing* rather than merely *different* technical ideals.

3. Practical Tradeoffs Between Alignment Ideals

In this section we introduce two distinctions with the aim of highlighting tradeoffs between the high-level alignment ideals introduced in Section 2. Section 3.1 and Section 3.2 each begin with a proposition, followed by definitions and arguments aimed at establishing each proposition. We then summarize these tradeoffs in Section 3.3.

3.1. Harms from Misdirected Competence vs Harms from Incompetence

Proposition 3.1. *Different threat-models lead to tradeoffs between different alignment ideals.*

We will say that a **threat-model class** (hereafter ‘threat-model’) is characterized by the source or origin of potential negative outcomes that may arise from developing AI systems, such that appropriate interventions can be used to minimize the chance that this set of outcomes occurs. To understand alignment ideals more precisely we must distinguish between two high-level threat-models, the first of which we call *Harms from Misdirected Competence*.

Definition 3.2 (Harms from Misdirected Competence). *Harms from Misdirected Competence* are threat-models which posit dangers arising from AI systems which are highly *competent* on some set of tasks.

Recall that Takeover Avoidance is motivated by the idea that highly powerful AI systems will, in the future, ‘optimize for’ or ‘pursue goals’ contrary to human interests, and thereafter remain impervious to human control. Hence, the threat-model underlying the ideal of Takeover Avoidance involves

AI systems in some sense being *too capable*. It is for this reason that the motivating threat-model behind Takeover Avoidance can be classified as *Competence Harm*.

The case of Social Judiciousness is less straightforward, as failures in this case may arise either from AI competence or AI incompetence. While concerns about Takeover Avoidance are *necessarily* concerns about AI competence, researchers focused on Social Judiciousness in many cases deal with harms resulting from *incompetence*. Consider, for example, criticisms made of the application of ML in predictive policing (Akpinar et al., 2021; Ensign et al., 2018; Fussell, 2020) as a result of models learning shallow or biased associations from their training data, or because they are given ‘practically impossible’ tasks (Raji et al., 2022). Here, the operative threat-model can be seen — given the unjust society in which we live — as a form of Training Data Conservatism yet to be overcome, and hence an *Incompetence Harm*.

Definition 3.3 (Harms from Incompetence). Threat-models which posit dangers that arise from AI systems which are *incompetent* on some set of tasks.

We introduce these threat-models as they highlight a source of potential tension between researchers focused on Social Judiciousness and those focused on Takeover Avoidance. Alongside predictive policing, many other cases can be found where researchers raise concerns about socially deleterious effects of AI simply replicating biases in their training data, affecting healthcare (Obermeyer et al., 2019; Agrawal & Prabakaran, 2020; Nagendran et al., 2020), facial recognition (Buolamwini & Gebru, 2018), the spread of misinformation, or in failure to generate ‘expression of views’ representative of the broader population (Liu et al., 2024b). For this reason, researchers focused on failures of Social Judiciousness due to *Incompetence Harms* may favor research programs aiming to reduce the tendency of present-day LLMs to ‘hallucinate’, as this may in turn reduce harms caused by LLMs repeating misinformation found in their training data.

By contrast, those focused on Takeover Avoidance may object to such proposals *precisely because* they make AI systems more competent. From the perspective of Takeover Avoidance, reduced hallucination rates may be undesirable unless they come alongside ‘specific countermeasures’ (Cotra, 2022); greater degrees of situational awareness may enable misaligned AIs to more effectively ‘scheme’ against the oversight mechanism (Carlsmith, 2023; Meinke et al., 2025), or increase models’ ability to ‘sandbag’ on certain tasks when being evaluated (Anthropic, 2025, pp. 90-100). This is our first example indicating potential tensions between the ideals expressed by different conceptions of ‘AI alignment’.

3.2. Positive and Negative Alignment

Proposition 3.4. *When evaluating model behavior/outputs, the scope of behavioral/output evaluation leads to tensions between different alignment ideals.*

Another source of potential tension between alignment ideals can be illustrated with a cross-cutting distinction. Insofar as Harms from Misdirected Competence and Harms from Incompetence both specify *threat-models*, they are primarily focused on the potential harms resulting from AI systems. However, we could imagine focusing more specifically on what *positive* properties we would like AI systems to possess. Let us say that ‘Positive Alignment’ prescribes the set of properties we want AI systems to possess, and ‘Negative Alignment’ prescribes properties we do *not* want AI systems to possess.

Of course, one may think that talk of ‘avoiding undesirable properties’ is *just the same thing* as ‘producing positive properties’. In fact, unless we have a complete description of ‘do X and only X’, the two come apart, as ‘do X’ is not ‘don’t do not-X’.⁴ However, the distinction between Positive and Negative Alignment arises when the implicit domain of evaluation for the AI’s outputs or ‘behaviors’ differs depending on whether we are checking whether it does what we *want*, or whether the AI *avoids* doing what we *don’t want*. In practice, it will often be more straightforward to evaluate whether an AI system performs some fixed task reliably or effectively than to evaluate whether its outputs avoid all of the multifarious ways an AI could falter. For instance, we may train a model to achieve better scores on certain mathematical benchmarks, and concomitantly find that the model has higher ‘hallucination rates’ than previous SOTA models.⁵ This would be progress towards the ideal of *Positive Alignment*, but regress away from the ideal of *Negative Alignment*.

The distinction between Positive and Negative Alignment is particularly important for understanding the potential tensions between Task Reliability and the remaining alignment ideals. When specifying desirable AI behavior, one cannot plausibly list every behavior that the developer does *not* wish the AI to exhibit. However, Social Judiciousness and Takeover Avoidance are concerned with the unintended consequences of deploying AI systems that narrowly behave (or appear to behave) in line with developer intentions.

⁴Let A denote the domain of ‘acts’, W_a the fact that we—qua developers of the AI system—want the system to do $a \in A$, and D_a the fact that the trained AI in fact does a . The natural formalization of Positive Alignment is $\forall a \in A : (W_a \rightarrow D_a)$; that of Negative Alignment is $\forall a \in A : (\neg W_a \rightarrow \neg D_a)$. We only have $\forall a \in A : (D_a \leftrightarrow W_a)$ only if they share the same domain of evaluation.

⁵This was observed after the release of OpenAI’s o3 and o4-mini models in comparison to GPT-4.5; see the PersonQA results presented in (OpenAI, 2025b) and (OpenAI, 2025a), respectively.

In practical terms, those concerned with Social Judiciousness may worry about recommender systems which maximize clickthrough rate but engender addiction or polarization (Stray et al., 2021; Agarwal et al., 2024; Prunkl, 2024; Smith et al., 2024), or worry that the values which get ‘successfully’ encoded in contemporary LLMs are instilled ‘top-down’ by large companies rather than having such values decided via more democratic processes (Tang, 2025). These may very well be success cases for Task Reliability but failures of Social Judiciousness. Similarly, those concerned with Takeover Avoidance may worry that training LLMs to produce locally desirable outputs – for instance, by producing a chain of thought (CoT) that does not offend users or otherwise appear undesirable – may leave some undesirable behaviors intact, while increasing longer-term risks by causing a model “to hide its intent” (Baker et al., 2025).

3.3. Summarizing The Tradeoffs

This section has introduced two distinctions in order to highlight points of practical conflict in the pursuit of different alignment ideals. Because different conceptions of ‘AI alignment’ may be motivated by different threat-models or place primary focus on AI benefits or AI harms, we claim that disagreements about how to ‘make AI systems more aligned’ encompass competing technical priorities.⁶

Alignment Ideal	Threat Model	Pos/Neg Alignment
Non-Takeover	Competence	Negative
Social	Either	Negative
Task	N/A	Positive

Table 2. Alignment Ideals: Axes of Difference.

4. Recommendations for Future Practice

We now list several recommendations for future practice.

Recommendation 1: Avoid Treating Distinctions of Policy As Distinctions of Scientific Scope

Consider claims that AI safety should focus on gradual disempowerment (Kulveit et al., 2025) or prioritize the future of work (Hazra et al., 2025). These are distinctions relating to *high-level policy ideals*. That is, ‘gradual disempowerment’ and ‘the future of human work’ pick out classes of societal outcomes which may be thought of or argued for as motivating ideals for technical researchers. In light of our

⁶Threat-models were defined as a set of negative outcomes able to be addressed via interventions. We think it is most plausible to say that ‘Task Reliability’ is not founded upon threat-models per se – but instead with a ‘gain-model’ – though we can imagine an argument claiming that Task Reliability is primarily concerned with *Incompetence Harms*.

discussion, we claim that researchers should carefully distinguish between high-level ideals of the kind cited and *model properties* we may want to target (or avoid) in near-term empirical work.

As noted by Kasirzadeh (2025), “advanced AI agents ... can be strategically leveraged for social engineering purposes such as ... amplifying existing prejudices, exploiting belief systems by synthetic evidence, and creating personalized epistemic bubbles” (2025, pg. 1987). Near-term work aimed at reducing the ability or propensity of AI systems to ‘create personalized epistemic bubbles’ (e.g., research attempting to mitigate LLM sycophancy) might thus be considered an intervention which reduces the chance of gradual disempowerment. In turn, decreasing LLM sycophancy could require granting LLMs better access to ‘ground truth’ claims (thus enabling them to more easily push back on user framings), or reinforcing LLMs’ tendencies to decline requests from users hoping for assistance in developing their pseudoscientific theories (see (Hill & Freedman, 2025) for discussion of a public case).

However, researchers concerned with ‘gradual disempowerment’ may equally well think that affording AI systems and their developers the ability to override the commitments of certain users could lead to an undesirable concentration of power. As such, researchers motivated by the ideal of avoiding ‘gradual disempowerment’ could, via better access to ground truth claims, be increasing LLMs’ situational awareness (potentially harmful for *Takeover Avoidance*), or indeed be focused on developing ‘non-judgmental’ LLMs which avoid pushing back on potentially misguided user framings (potentially harmful for *Social Judiciousness*). While it is not our place here to proffer any advice concerning which high-level ideals to adopt, we think it important to distinguish between distinctions of *scientific or technical scoping* and distinctions between broader policy ideals or directions.

Recommendation 2: Note Methodological Differences

Compared to our other two alignment ideals, Takeover Avoidance sounds fairly speculative. Indeed, while many prominent AI researchers treat Takeover Avoidance as important (Russell, 2019; Bengio, 2023; Milmo, 2024), other voices have variously judged such concerns as purely “hypothetical” (Domínguez Hernández et al., 2024), “ridiculously overblown” (LeCun, 2023), fantastically sci-fi (Ng, 2015), or “nonsense” (Bender, 2025). In some cases, researchers concerned with Social Judiciousness explicitly contrast ‘concrete harms’ with ‘speculative risks’, claiming the latter illegitimately “distract” from the former (Stark, 2024; Lahlou, 2025). For one camp, ‘speculativeness’ might be a necessary feature of conducting research concerned with the eventual trajectory of AI research, whereby temporary incompetencies should be expected to be resolved in the future. For another, depth of social analysis as applied to

the *contemporary* technological landscape (as might be seen in invocations of, for example, ‘the patriarchy’ or ‘racial capitalism’) might be seen as necessary in order to surface harms which are otherwise opaque in ordinary technological practice. To the extent that concerns of Social Judiciousness require theorizing about the societal effects of deploying AI systems, such theorizing requires *greater* interpretive judgment than analyses of AI performance with respect to more narrow technical criteria.

This methodological difference is likely to explain at least some portion of the disagreement between researchers who prioritize Takeover Avoidance and those who don’t. This difference is visible in a 2022 analysis of papers published in proceedings of AIES and FAccT by Birhane and co-authors, which evaluated the degree to which each published paper “investigates, addresses, and/or mentions the disparate impacts of an algorithmic system” (Birhane et al., 2022, pg. 952), concluding with a call for future research to devote more attention to possible disparate impacts of AI systems.

Indeed, we note that methodological differences appear even when we restrict our attention to internal discussions of *Task Reliability*. Questions relating to the external validity of benchmarks, for example, are often questions about the degree to which our existing metrics capture more informal (and perhaps thicker) conceptions of ‘capability’. Insofar as our more informal concepts of ‘capable AI systems’ outstrip our more formal and quantitative benchmarks, there exist gaps between ‘our measurements of AI capabilities’ and ‘AI systems’ *actual capabilities*, with assessments of the latter being more speculative than the former.

Recommendation 3: Reviewer Pooling and The Diversity of ‘AI Alignment’

‘AI alignment’ encompasses many distinct concepts. For this reason, we think it would be beneficial for: (i) researchers using the term ‘AI alignment’ to clearly identify the artifact ‘being aligned’ and target metric of ‘alignment’, and (ii) for major venues to distinguish alignment subareas in submission tracks and construct reviewer pools accordingly. The second suggestion in particular is well-supported by a 2025 paper investigating citation patterns for work on ‘AI Ethics’ (focused on ‘Social Judiciousness’) and ‘AI Safety’ (focused on ‘Takeover Avoidance’), finding high degrees of citation homophily for both such communities (Roytburg & Miller, 2025). While our paper has focused on *conceptual* tensions between these ideals, empirical work provides evidence for *sociological clustering* — hence, instantiating distinct reviewer pools for these areas is likely to track distinct areas of expertise.

Our suggestions on this front nonetheless arrive with a notable degree of caution. While our paper has focused on three specific conceptions of ‘AI alignment’, we do not wish

to suggest that the conceptions we introduce are the *only* such conceptions available (recall Section 2.4). Indeed, if we view questions of AI alignment — very generally — as questions about what AIs *are* and *what we want to do with AI*, a yet more diverse range of conceptions emerges (Andreessen, 2023; Bostrom, 2024; Brynjolfsson, 2022; Danaher, 2019; Goldberg et al., 2024; Plurality Institute, 2025; ?; ?).

We thus do not wish to unduly reify the ‘alignment’ concepts discussed explicitly by this paper. Instead, we suggest that future discussions of ‘AI alignment’ either carefully specify their focus in light of existing ideals, or construct new alignment ideals in light of the distinctions outlined, or find new ways to carve up the space of concerns animating distinct alignment ideals.

Recommendation 4: Use Qualified Alignment Terms

Consider Takeover Avoidance. Empirical researchers who aspire to produce work in service of this ideal do not, as of yet, have generally superhuman AI systems which they can try to ‘align’. As noted in Section 2.3, researchers within this tradition sometimes talk about the ‘goals’ or ‘preferences’ of contemporary systems (see (Zhi-Xuan et al., 2025) for fuller discussion of this point). Hence, the *proxy alignment concept* used in such research may be thought of as a form of *Preference Alignment*, where (say):

an AI system *S* is ‘aligned’ if *S*: (i) possesses preference-like states, such that (ii) the preferences of *S* are desirable preferences relative to some normative standard.

Of course, questions regarding what it means for an AI system to possess ‘preference-like states’ clearly remain. Nonetheless, using terms like Preference Alignment allows, in the first instance, those researchers sympathetic to the ideal of Takeover Avoidance to more easily discuss ‘internal’ questions regarding whether observed failures of Preference Alignment constitute meaningful empirical evidence regarding Takeover Avoidance. Meanwhile, those less sympathetic to Takeover Avoidance can more easily separate criticisms of the observed results and their interpretation (e.g., ‘should we treat the model as possessing preferences?’, ‘is the normative standard being used reasonable or well-specified?’, etc.) from criticisms of the high-level ideal.

Likewise, our earlier discussion of Social Judiciousness distinguished between two possible sources of ‘alignment’ failure: Training Data Conservatism and Power Amplification. We think that future research in the tradition of Social Judiciousness should pay heed to this distinction, as different sources of harm lead to different routes for intervention. One unfortunate example to illustrate this point involves the absence of demographic diversity in the training data for

facial recognition technologies. For instance, facial recognition classifiers identified as racially biased (Buolamwini & Gebru, 2018) include those produced by IBM, which has tested facial recognition software within the NYPD (Allyn, 2020). To the extent one is concerned about such facial recognition technologies harming darker-skinned individuals through their use in predictive policing, discussions among researchers focused on Social Judiciousness – even among those sharing political ideals – ought to pay attention to the distinction outlined.

Recommendation 5: Clearly Communicate Differences to Policy and Non-Technical Audiences

Policymakers and scientific administrators are increasingly concerned with questions of ‘AI alignment’. However, as our discussion has illustrated, different people may mean different things by ‘AI alignment’, and the term’s polysemy in some cases leads to tradeoffs in practical priorities. Key decision-makers who remain unaware of these differences are thus less able to assess which priorities are most important for them, the different sources of evidence one might need in order to establish different sorts of concern, and the appropriate interventions to pursue given any such concern.

An example illustrating the importance of Recommendation 5 involves reference to the EU AI Act. The EU AI Act makes reference both to AI “alignment with human intent” (European Parliament, 2024, pg. 100) and to the importance of evaluating AI models using “appropriate technical tools and methodologies” (European Parliament, 2024, pg. 101). However, the act’s discussion of auditing does not distinguish between the alignment ideals introduced by this paper; as such, policymakers mandating alignment audits lack clear and principled bases for distinguishing between the very different kinds of methodologies that may be required when auditing for (say) demographic biases and autonomous self-replication capabilities, respectively. Policymakers should thus be aware that these differences in terminology may have important consequences. Similarly, researchers across traditions should take care in distinguishing their own concerns from adjacent concerns, and communicating competing concerns in ways that illustrate the substantive issues raised by alternative camps even if one ultimately finds them unconvincing.

5. Alternative Views

We now defend our position against three alternative views.

5.1. Disagreements About ‘Alignment’ Aren’t Disagreements About Technical Priorities

Objection. All disagreements concerning the referent of ‘AI alignment’ are best understood as either: (i) merely normative disagreements, or else; (ii) reflective of a single unified

concept of ‘AI alignment’ encompassing the concerns of every more specific ideal discussed in this paper.

Reply. The existence of a single unifying alignment ideal is (to put it baldly) simply implausible in light of our discussion. Indeed, part of the reason we have discussed tensions between different alignment ideals is to illustrate that such ideals are not merely concepts which are *compatibly different* (e.g., the concept ‘red’ and the concept ‘square’) but *incompatibly different* (e.g., the concept ‘red all over’ and the concept ‘green all over’).⁷ Insofar as there exist interventions which aid progress towards one such alignment ideal precisely by vitiating progress towards another, there cannot be any single idea which unifies them.

We also do not believe that disagreements between alignment ideals are purely normative. Almost no one wishes for an AI takeover, and so researchers preferring alternative alignment ideals over Takeover Avoidance in large part do so because they find the possibility of an AI takeover *implausible* — this is quite clearly an epistemic and not a normative difference. While disagreements about epistemic matters may themselves be intermingled with normative disagreements (e.g., regarding appropriate methodologies when thinking about the effects of AI systems), they remain disagreements about (even if speculative) matters of fact.

5.2. A More Modest Objection: ‘Good Enough’ Convergence

Objection. The previous objection was too strong. We grant that there exist *some* divergent priorities among different alignment ideals, but in practice such ideals are complementary and rarely conflict.

Reply. The second objection is more plausible than the first, and is occasionally hinted at in public talks (Dragan, 2025). And indeed we agree that there will exist shared technical projects which are positive by the lights of multiple distinct alignment ideals. For instance, techniques such as training data filtering may be useful from the perspective of both Takeover Avoidance and Social Judiciousness, while posing minimal costs to model capabilities that matter from the perspective of Task Reliability. Likewise, proposals such as third-party model evaluations aimed at improving robustness against harmful optimization — whether that be harms from *AI systems themselves* optimizing against human interests, the presence of harmful ‘backdoors’ implanted by unscrupulous AI developers, or the incautious rollout of narrowly capable but unreliable AI systems — may be similarly valuable from multiple perspectives.

⁷In slightly more outré language, one can read the claims of this paper as proposing that different working alignment ideals may stand in relations of ‘determinate negation’ to one another without standing in relations of ‘formal’ or ‘abstract’ negation (Hegel, 1807).

We nonetheless find the second objection implausible. Even the promisingly conciliatory case of model evals involves evaluations which “can vary widely in terms of method, approach, and subject matter”, meaning that digital services coordinators in charge of overseeing external evaluations may be misled about which model properties it is most important to evaluate (Terzis et al., 2024). This in turn could lead to audits for models’ ‘autonomous self-replication ability’ even when one is most concerned with Social Judiciousness, or audits for models’ demographic biases when one is most concerned with Takeover Avoidance. As we have highlighted the existence of *different* alignment ideals producing competing technical priorities *in at least some cases*, we believe that any claims of ‘broad complementarity’ between different alignment ideals require positive arguments that we have not yet found in the published literature.

5.3. The Paper’s Claims Need Empirical Grounding

Objection. The paper aims to substantiate its views about alignment without reference to numbers on research trends, paper-keywords, citation analyses, interviews, or surveys. If ‘alignment’ really does encompass multiple competing technical ideals, this claim should be substantiated empirically.

Reply. Here we note a simple point of methodological disagreement. Before we can empirically study competing technical priorities behind different alignment ideals, we first need to understand the terrain conceptually: what *are* the underlying disagreements, such that we can evaluate whether a given empirical methodology successfully captures these differences?

Recall the earlier cited research from Roytburg and Miller, which found high degrees of citation homophily for researchers focused on ‘Social Judiciousness’ and those focused on ‘Takeover Avoidance’ (Roytburg & Miller, 2025). This analysis – while valuable, and while successfully providing evidence for the existence of *relatively insular* research communities – does not and cannot empirically address whether the ideals motivating such research are *in tension* with each other. It is for this reason that we believe that Roytburg and Miller’s practical suggestions for ‘unifying’ these disparate literatures are premature, and use this juncture to note agreement with a prior position paper highlighting an absence of “conceptual clarity” as one pitfall of empirical ML research (Herrmann et al., 2024). With all that said, however, we believe that our paper provides value in part because it may help lay the conceptual groundwork for future empirical investigation of potential alignment tradeoffs. Consider the following research questions:

RQ1. At what rates do interventions reducing hallucination (a gain from the perspective of Social Judiciousness) measurably increase capabilities relevant to scheming

or deceptive alignment (a concern from the perspective of Takeover Avoidance)?

RQ2. Does improving models’ situational awareness (which may reduce certain incompetence harms) correlate with increased capacity for sandbagging on safety evaluations?

RQ1 and *RQ2* are offered as potential research questions amenable to future empirical operationalization and investigation. Yet *RQ1* and *RQ2* can be formulated precisely only by way of disambiguating the multifarious concerns often conflated under the single heading of ‘AI alignment’. Absent distinctions of the kind introduced by this paper, one cannot even *state* — let alone empirically investigate — research questions involving potential tradeoffs between Task Reliability, Social Judiciousness, and Takeover Avoidance.

6. Conclusion

The term ‘AI alignment’ encompasses multiple competing ideals. Sections 2 and 3 illustrated both that ‘AI alignment’ is polysemous and that this polysemy resulted in *competing* and not merely *different* technical priorities. Together, these sections evidenced the paper’s primary argumentative position. Thereafter, Section 4 expanded on the implications of our analysis. We suggested that researchers more clearly distinguish empirical ‘alignment proxies’ when discussing Takeover Avoidance, and more carefully distinguish between different sources of harm in cases of Social Judiciousness. More generally, we claimed that researchers would benefit from more explicitly specifying their background commitments without reifying the alignment concepts discussed herein.

We believe that the present paper provides value for reasons suggested in Section 5.3. We should try to understand disagreements about ‘AI alignment’ as they actually exist, which (for better or worse) do not come in the form of neatly packaged formalizations readily amenable to empirical study. In order to move towards more fruitful discussions about AI alignment, we need to acknowledge friction; we hope to have provided the rough ground.⁸

References

Agarwal, A., Usunier, N., Lazaric, A., and Nickel, M. System-2 recommenders: Disentangling utility and en-

⁸Our end line pays homage to a passage from Wittgenstein: “The more narrowly we examine actual language, the sharper becomes the conflict between [natural language] and our requirement [to express these ideals with the “crystalline purity” of logic] ... We have got on to slippery ice where there is no friction and so in a certain sense the conditions are ideal, but also, just because of that, we are unable to walk. We want to walk: so we need friction. Back to the rough ground!” (Wittgenstein, 1953, §107)

- agement in recommendation systems via temporal point-processes. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’24, pp. 1763–1773, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400704505. doi: 10.1145/3630106.3659004. URL <https://doi.org/10.1145/3630106.3659004>.
- Agrawal, R. and Prabakaran, S. Big data in digital healthcare: lessons learnt and recommendations for general practice. *Heredity*, 124:525–534, 2020. doi: 10.1038/s41437-020-0303-2. URL <https://www.nature.com/articles/s41437-020-0303-2>.
- Agüera y Arcas, B. Evolutionary transition, 2025. URL <https://whatisintelligence.antikythera.org/chapter-10/#beyond-alignment>. Chapter 10 of *What Is Intelligence?: Lessons from AI About Evolution, Computing, and Minds*, section “Beyond Alignment”.
- Akpinar, N.-J., De-Arteaga, M., and Chouldechova, A. The effect of differential victim crime reporting on predictive policing systems. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’21, pp. 838–849, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383097. doi: 10.1145/3442188.3445877. URL <https://doi.org/10.1145/3442188.3445877>.
- Allyn, B. Ibm abandons facial recognition products, condemns racially biased surveillance. NPR, June 2020. URL <https://www.npr.org/2020/06/09/873298837/ibm-abandons-facial-recognition-products-condemns-racially-biased-surveillance>. Accessed: 2026-01-29.
- Andreessen, M. The techno-optimist manifesto, 2023. URL <https://a16z.com/the-techno-optimist-manifesto/>. Accessed 2025-07-04.
- Anthropic. Claude Opus 4.5 System Card. Technical report, Anthropic, November 2025. URL <https://assets.anthropic.com/m/64823ba7485345a7/C1aude-Opus-4-5-System-Card.pdf>.
- Armstrong, S. General purpose intelligence: Arguing the orthogonality thesis. *Analysis and Metaphysics*, 12:68–84, 2013. ISSN 1584-8574. URL <https://addletonacademicpublishers.com/search-in-a-m/217-volume-12-2013/1964-general-purpose-intelligence-arguing-the-orthogonality-thesis>.
- Babaei, P. and Vassileva, J. Drivers and persuasive strategies to influence user intention to learn about manipulative design. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’24, pp. 2421–2431, Rio de Janeiro, Brazil, June 2024. Association for Computing Machinery. ISBN 979-8-4007-0450-5. doi: 10.1145/3630106.3659046. URL <https://doi.org/10.1145/3630106.3659046>.
- Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., Chen, C., Olsson, C., Olah, C., Hernandez, D., Drain, D., Ganguli, D., Li, D., Tran-Johnson, E., Perez, E., Kerr, J., Mueller, J., Ladish, J., Landau, J., Ndousse, K., Lukosuite, K., Lovitt, L., Sellitto, M., Elhage, N., Schiefer, N., Mercado, N., DasSarma, N., Lasenby, R., Larson, R., Ringer, S., Johnston, S., Kravec, S., El Showk, S., Fort, S., Lanham, T., Telleen-Lawton, T., Conerly, T., Henighan, T., Hume, T., Bowman, S. R., Hatfield-Dodds, Z., Mann, B., Amodei, D., Joseph, N., McCandlish, S., Brown, T., and Kaplan, J. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022. doi: 10.48550/arXiv.2212.08073. URL <https://arxiv.org/abs/2212.08073>.
- Baker, B., Huizinga, J., Madry, A., Zaremba, W., Pachocki, J., and Farhi, D. Detecting misbehavior in frontier reasoning models. OpenAI, March 2025. URL <https://openai.com/index/chain-of-thought-monitoring/>. Accessed: 2026-01-26.
- Bender, E. M. Exploring ai with emily m. bender (host: Ricardo signes). Digital Citizen podcast (Fastmail), 2025. URL <https://www.fastmail.com/digital-citizen/exploring-ai-with-emily-m-bender/>. Podcast episode, accessed June 13, 2025.
- Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. On the dangers of stochastic parrots: Can language models be too big? . In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’21, pp. 610–623, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383097. doi: 10.1145/3442188.3445922. URL <https://doi.org/10.1145/3442188.3445922>.
- Bengio, Y. AI and Catastrophic Risk, Sep 2023. URL <https://www.journalofdemocracy.org/ai-and-catastrophic-risk/>.
- Betley, J., Tan, D. C. H., Warncke, N., Szyber-Betley, A., Bao, X., Soto, M., Labenz, N., and Evans, O. Emergent misalignment: Narrow finetuning can produce broadly misaligned LLMs. In *Forty-second International Conference on Machine Learning (ICML 2025)*, 2025. URL <https://openreview.net/forum?id=aOIJ2gVRWW>.

- Birhane, A., Ruane, E., Laurent, T., Brown, M. S., Flowers, J., Ventresque, A., and Dancy, C. L. The forgotten margins of AI ethics. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’22, pp. 948–958, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450393522. doi: 10.1145/3531146.3533157. URL <https://doi.org/10.1145/3531146.3533157>.
- Bostrom, N. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, 2014.
- Bostrom, N. *Deep Utopia*. Ideapress Publishing, 2024.
- Bratton, B. The five stages of ai grief, June 2024. URL <https://www.noemamag.com/the-five-stages-of-ai-grief/>.
- Brynjolfsson, E. The turing trap: The promise & peril of human-like artificial intelligence, 2022. URL <https://digitaleconomy.stanford.edu/news/the-turing-trap-the-promise-peril-of-human-like-artificial-intelligence/>.
- Buolamwini, J. and Gebru, T. Gender shades: Intersectional accuracy disparities in commercial gender classification. In Friedler, S. A. and Wilson, C. (eds.), *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, volume 81 of *Proceedings of Machine Learning Research*, pp. 77–91. PMLR, 23–24 Feb 2018. URL <https://proceedings.mlr.press/v81/buolamwini18a.html>.
- Carlsmith, J. Is power-seeking AI an existential risk? *arXiv preprint arXiv:2206.13353*, 2022.
- Carlsmith, J. Scheming ais: Will ais fake alignment during training in order to get power?, 2023. URL <https://arxiv.org/abs/2311.08379>.
- Chen, S., Xiong, M., Liu, J., Wu, Z., Xiao, T., Gao, S., and He, J. In-context sharpness as alerts: An inner representation perspective for hallucination mitigation. In *ICML*, 2024. URL <https://proceedings.mlr.press/v235/chen24av.html>.
- Christian, B. *The Alignment Problem*. W. W. Norton & Company, October 2020.
- Christiano, P. Prosaic AI alignment. <https://ai-alignment.com/prosaic-ai-control-b959644d79c2>, 2016. AI Alignment. Accessed 28 May 2026.
- Christiano, P. Corrigibility, November 2018. URL <https://www.alignmentforum.org/posts/fkLYhTQteAu5SinAc/corrigibility>.
- Christiano, P. What failure looks like, March 2019. URL <https://www.lesswrong.com/posts/HBxe6wdjxK239zajf/what-failure-looks-like>. LessWrong.
- Cotra, A. Without specific countermeasures, the easiest path to transformative AI likely leads to AI takeover. AI Alignment Forum, July 2022. URL <https://www.alignmentforum.org/posts/pRkFkzwKZ2zfa3R6H/without-specific-countermeasures-the-easiest-path-to>. Published 18 Jul 2022. Accessed 29 Jan 2026.
- Danaher, J. *Automation and Utopia: Human Flourishing in a World without Work*. Harvard University Press, Cambridge, MA, 2019. ISBN 9780674984240. URL <https://www.amazon.co.uk/Automation-Utopia-John-Danaher/dp/0674984242>.
- Dehnert, M. and Mongeau, P. A. Persuasion in the age of artificial intelligence (AI): Theories and complications of AI-based persuasion. *Human Communication Research*, 48(3):386–403, July 2022. doi: 10.1093/hcr/hqac006. URL <https://academic.oup.com/hcr/article-abstract/48/3/386/6564679>.
- Deutsch, D. Beyond reward and punishment. In Brockman, J. (ed.), *Possible Minds: Twenty-Five Ways of Looking at AI*, pp. 113–124. Penguin Press, New York, 2019.
- di Langosco, L. L., Koch, J., Sharkey, L., Pfau, J., and Krueger, D. Goal misgeneralization in deep reinforcement learning. In *ICML*, 2022.
- Domínguez Hernández, A., Krishna, S., Perini, A. M., Katell, M., Bennett, S., Borda, A., Hashem, Y., Hadjiloizou, S., Mahomed, S., Jayadeva, S., Aitken, M., and Leslie, D. Mapping the individual, social and biospheric impacts of foundation models. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’24, pp. 776–796, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400704505. doi: 10.1145/3630106.3658939. URL <https://doi.org/10.1145/3630106.3658939>.
- Dragan, A. Navigating the path to AGI safely & responsibly. Presentation at the International Association for Safe and Ethical Artificial Intelligence Conference (IASEAI’25), February 2025. URL <https://www.iaseai.org/our-programs/iaseai25>. OECD La Muette Headquarters and Conference Centre.
- Ensign, D., Friedler, S. A., Neville, S., Scheidegger, C., and Venkatasubramanian, S. Runaway feedback loops in predictive policing. In Friedler, S. A. and Wilson, C. (eds.),

- Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, volume 81 of *Proceedings of Machine Learning Research*, pp. 160–171. PMLR, 23–24 Feb 2018. URL <https://proceedings.mlr.press/v81/ensign18a.html>.
- European Parliament. Artificial Intelligence Act: European Parliament legislative resolution of 13 march 2024 on the proposal for a regulation of the European Parliament and of the Council on laying down harmonised rules on Artificial Intelligence (Artificial Intelligence Act) and amending certain Union Legislative Acts (COM(2021)0206 – C9-0146/2021 – 2021/0106(COD)). Texts adopted, P9.TA(2024)0138, March 2024. URL https://artificialintelligenceact.eu/wp-content/uploads/2024/04/TA-9-2024-0138_EN.pdf.
- Foster, J. G. From Thin to Thick: Toward a Politics of Human-Compatible AI. *Public Culture*, 35(3 (101)):417–430, September 2023. ISSN 0899-2363. doi: 10.1215/08992363-10742593.
- Fussell, S. An algorithm that ‘Predicts’ criminality based on a face sparks a furor, June 2020. URL <https://www.wired.com/story/algorithm-predicts-criminality-based-face-sparks-furor/>.
- Future of Life Institute. Asilomar AI principles. <https://futureoflife.org/open-letter/ai-principles/>, August 2017. Accessed 2026-05-29.
- Gao, L., Schulman, J., and Hilton, J. Scaling laws for reward model overoptimization. In Krause, A., Brunskill, E., Cho, K., Engelhardt, B., Sabato, S., and Scarlett, J. (eds.), *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 10835–10866. PMLR, 23–29 Jul 2023. URL <https://proceedings.mlr.press/v202/gao23h.html>.
- Gehrmann, S., Huang, C., Teng, X., Yurovski, S., Bhorkar, A., Thomas, N., Doucette, J., Rosenberg, D., Dredze, M., and Rabinowitz, D. Understanding and mitigating risks of generative AI in financial services. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’25, pp. 2570–2586, New York, NY, USA, 2025. Association for Computing Machinery. ISBN 9798400714825. doi: 10.1145/3715275.3732168. URL <https://doi.org/10.1145/3715275.3732168>.
- Goldberg, B., Acosta-Navas, D., Bakker, M., Beacock, I., Botvinick, M., Buch, P., DiResta, R., Donthi, N., Fast, N., Iyer, R., Jalan, Z., Konya, A., Danciu, G. K., Landemore, H., Marwick, A., Miller, C., Ovadya, A., Saltz, E., Schirch, L., Shalom, D., Siddarth, D., Sieker, F., Small, C., Stray, J., Tang, A., Tessler, M. H., and Zhang, A. AI and the future of digital public squares, 2024. URL <https://arxiv.org/abs/2412.09988>.
- Good, I. J. Speculations concerning the first ultraintelligent machine. In Alt, F. L. and Rubinoﬀ, M. (eds.), *Advances in Computers*, volume 6, pp. 31–88. Academic Press, New York, 1966. doi: 10.1016/S0065-2458(08)60418-0. URL <https://www.sciencedirect.com/science/chapter/bookseries/abs/pii/S0065245808604180>.
- Greenblatt, R., Denison, C., Wright, B., Roger, F., MacDiarmid, M., Marks, S., Treutlein, J., Belonax, T., Chen, J., Duvenaud, D., Khan, A., Michael, J., Mindermann, S., Perez, E., Petrini, L., Uesato, J., Kaplan, J., Shlegeris, B., Bowman, S. R., and Hubinger, E. Alignment faking in large language models, 2024. URL <https://arxiv.org/abs/2412.14093>.
- Grietzner, P. and Jha, T. The Problem With the Word “Alignment”, May 2024. URL <https://www.alignmentforum.org/posts/p3aL6BwpbPhqxnayL/t-he-problem-with-the-word-alignment-1>.
- Hadfield-Menell, D. and Hadfield, G. Incomplete contracting and AI alignment, 2018. URL <https://arxiv.org/abs/1804.04268>.
- Hazra, S., Majumder, B. P., and Chakrabarty, T. Position: AI safety should prioritize the future of work. In *Forty-second International Conference on Machine Learning Position Paper Track*, 2025. URL <https://openreview.net/forum?id=CA9NxmmUG5>.
- Hegel, G. W. F. *Phenomenology of Spirit*. Galaxy Books. Oxford University Press, 1807. ISBN 9780198245971. With an analysis of the text and foreword by J. N. Findlay.
- Herrmann, M., Lange, F. J. D., Eggenberger, K., Casalicchio, G., Wever, M., Feurer, M., Rügamer, D., Hüllermeier, E., Boulesteix, A.-L., and Bischl, B. Position: Why we must rethink empirical research in machine learning. In Salakhutdinov, R., Kolter, Z., Heller, K., Weller, A., Oliver, N., Scarlett, J., and Berkenkamp, F. (eds.), *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 18228–18247. PMLR, 21–27 Jul 2024. URL <https://proceedings.mlr.press/v235/herrmann24b.html>.
- Hill, K. and Freedman, D. Chatbots can go into a delusional spiral. here’s how it happens. *The New York Times*, August 2025. URL <https://www.nytimes.com/2025/08/08/technology/ai-chatbots-delusions-chatgpt.html>.

- Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., and Garrabrant, S. Risks from learned optimization in advanced machine learning systems, 2019. URL <https://arxiv.org/abs/1906.01820>.
- Irving, G., Christiano, P., and Amodei, D. Ai safety via debate, 2018. URL <https://arxiv.org/abs/1805.00899>.
- Ji, J., Chen, W., Wang, K., Hong, D., Fang, S., Chen, B., Zhou, J., Dai, J., Han, S., Guo, Y., and Yang, Y. Mitigating deceptive alignment via self-monitoring, 2025. URL <https://arxiv.org/abs/2505.18807>.
- Johnson, D. G. and Verdicchio, M. The sociotechnical entanglement of ai and values. *AI & Society*, (40):67–76, 2025. doi: 10.1007/s00146-023-01852-5.
- Kasirzadeh, A. Two types of AI existential risk: decisive and accumulative. *Philosophical Studies*, 182(7):1975–2003, 2025. doi: 10.1007/s11098-025-02301-3. URL <https://doi.org/10.1007/s11098-025-02301-3>.
- Kenton, Z., Siegel, N. Y., Kramár, J., Brown-Cohen, J., Albanie, S., Bulian, J., Agarwal, R., Lindner, D., Tang, Y., Goodman, N. D., and Shah, R. On scalable oversight with weak LLMs judging strong LLMs. In Globerson, A., Mackey, L., Belgrave, D., Fan, A., Paquet, U., Tomczak, J., and Zhang, C. (eds.), *Advances in Neural Information Processing Systems*, volume 37, pp. 75229–75276, Vancouver, Canada, December 2024. Neural Information Processing Systems Foundation, Inc. doi: 10.52202/079017-2395. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/899511e37a8e01e1bd6f6f1d377cc250-Abstract-Conference.html.
- Kirk, H. R., Gabriel, I., Summerfield, C., Vidgen, B., and Hale, S. A. Why human-ai relationships need socioaffective alignment, 2025. URL <https://arxiv.org/abs/2502.02528>.
- Korinek, A. and Balwit, A. Aligned with whom? direct and social goals for AI systems. In *The Oxford Handbook of AI Governance*, pp. 65–85. Oxford University Press, 2023. doi: 10.1093/oxfordhb/9780197579329.013.4. URL <https://doi.org/10.1093/oxfordhb/9780197579329.013.4>.
- Kulveit, J., Douglas, R., Ammann, N., Turan, D., Krueger, D., and Duvenaud, D. Position: Humanity faces existential risk from gradual disempowerment. In *Forty-second International Conference on Machine Learning Position Paper Track*, 2025. URL <https://openreview.net/forum?id=mWlnrtOTtm>.
- Lahlou, S. Mitigating societal cognitive overload in the age of ai: Challenges and directions, 2025. URL <https://arxiv.org/abs/2504.19990>.
- Laitinen, A. and Sahlgren, O. Ai systems and respect for human autonomy. *Frontiers in Artificial Intelligence*, 4: 705164, 2021. doi: 10.3389/frai.2021.705164. URL <https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2021.705164/full>.
- LeCun, Y. I think that the magnitude of the AI alignment problem has been ridiculously overblown & our ability to solve it widely underestimated, March 2023. URL https://www.linkedin.com/posts/yann-lecun_i-think-that-the-magnitude-of-the-ai-alignment-activity-7043650114048925696-1oXJ/. LinkedIn post.
- Lee, M. H., Montgomery, J. M., and Lai, C. K. Large language models portray socially subordinate groups as more homogeneous, consistent with a bias observed in humans. FAccT ’24, pp. 1321–1340, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400704505. doi: 10.1145/3630106.3658975. URL <https://doi.org/10.1145/3630106.3658975>.
- Legg, S. and Hutter, M. A universal measure of intelligence for artificial agents. In *Proceedings of the 19th International Joint Conference on Artificial Intelligence, IJCAI’05*, pp. 1509–1510, San Francisco, CA, USA, 2005. Morgan Kaufmann Publishers Inc.
- Li, G., Hammoud, H. A. A. K., Itani, H., Khizbullin, D., and Ghanem, B. Camel: Communicative agents for “mind” exploration of large language model society. In *Advances in Neural Information Processing Systems*, volume 36, pp. 51991–52008. Curran Associates, Inc., 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/a3621ee907def47c1b952ade25c67698-Paper-Conference.pdf.
- Liu, H., Xue, W., Chen, Y., Chen, D., Zhao, X., Wang, K., Hou, L., Li, R., and Peng, W. A survey on hallucination in large vision-language models, 2024a. URL <https://arxiv.org/abs/2402.00253>.
- Liu, Y., Yang, K., Qi, Z., Liu, X., Yu, Y., and Zhai, C. X. Bias and volatility: A statistical framework for evaluating large language model’s stereotypes and the associated generation inconsistency. In *NeurIPS*, 2024b. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/c6ec4a25a11393f277cfd64b7ea4d106-Paper-Datasets_and_Benchmarks_Track.pdf.

- Longoni, C., Fradkin, A., Cian, L., and Pennycook, G. News from generative artificial intelligence is believed less. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’22, pp. 97–106, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450393522. doi: 10.1145/3531146.3533077. URL <https://doi.org/10.1145/3531146.3533077>.
- Meinke, A., Schoen, B., Scheurer, J., Balesni, M., Shah, R., and Hobbhahn, M. Frontier models are capable of in-context scheming, 2025. URL <https://arxiv.org/abs/2412.04984>.
- Milmo, D. ‘Godfather of AI’ shortens odds of the technology wiping out humanity over next 30 years. The Guardian, December 2024. URL <https://www.theguardian.com/technology/2024/dec/27/godfather-of-ai-raises-odds-of-the-technology-wiping-out-humanity-over-next-30-years>. Last modified 2024-12-28.
- Moayeri, M., Tabassi, E., and Feizi, S. Worldbench: Quantifying geographic disparities in LLM factual recall. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’24, pp. 1211–1228, New York, NY, USA, June 2024. Association for Computing Machinery. ISBN 979-8-4007-0450-5/24/06. doi: 10.1145/3630106.3658967. URL <https://doi.org/10.1145/3630106.3658967>.
- Moon, K., Green, A. E., and Kushlev, K. Homogenizing effect of large language models (LLMs) on creative diversity: An empirical comparison of human and chatgpt writing. *Computers in Human Behavior: Artificial Humans*, 6:100207, December 2025. ISSN 2949-8821. doi: 10.1016/j.chbah.2025.100207. URL <https://www.sciencedirect.com/science/article/pii/S294988212500091X>.
- Morris, M. R., Sohl-dickstein, J., Fiedel, N., Warkentin, T., Dafoe, A., Faust, A., Farabet, C., and Legg, S. Levels of agi for operationalizing progress on the path to agi, 2024. URL <https://arxiv.org/abs/2311.02462>.
- Nagendran, M., Chen, Y., Lovejoy, C. A., Gordon, A. C., Komorowski, M., Harvey, H., Topol, E. J., Ioannidis, J. P. A., Collins, G. S., and Maruthappu, M. Artificial intelligence versus clinicians: systematic review of design, reporting standards, and claims of deep learning studies. *BMJ*, 368:m689, March 2020. doi: 10.1136/bmj.m689. URL <https://www.bmj.com/content/368/bmj.m689>. PMID: 32213531; PMCID: PMC7190037.
- Nehring, J., Gabryszak, A., Jürgens, P., Burchardt, A., Schaffer, S., Spielkamp, M., and Stark, B. Large language models are echo chambers. In Calzolari, N., Kan, M.-Y., Hoste, V., Lenci, A., Sakti, S., and Xue, N. (eds.), *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pp. 10117–10123, Torino, Italia, May 2024. ELRA and ICCL. URL <https://aclanthology.org/2024.lrec-main.884/>.
- Nelson, A. Thick Alignment. Keynote speech at the ACM Conference on Fairness, Accountability, and Transparency (FAccT’23), Chicago, IL, June 2023. URL https://www.youtube.com/watch?v=Sq_XwqVTqvQ. Accessed: 13 June 2025.
- Neumann, T., De-Arteaga, M., and Fazelpour, S. Justice in misinformation detection systems: An analysis of algorithms, stakeholders, and potential harms. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’22, pp. 1504–1515, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450393522. doi: 10.1145/3531146.3533205. URL <https://doi.org/10.1145/3531146.3533205>.
- Ng, A. Quote origin: I don’t work on preventing ai from turning evil for the same reason that i don’t work on the problem of overpopulation on the planet mars, 2015. URL <https://quoteinvestigator.com/2020/10/04/mars/>.
- Ngo, R., Chan, L., and Mindermann, S. The alignment problem from a deep learning perspective. In *International Conference on Learning Representations*, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/1e58b1bf9f218fcd19e4539e982752a5-Abstract-Conference.html.
- Obermeyer, Z., Powers, B., Vogeli, C., and Mullainathan, S. Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453, 2019. doi: 10.1126/science.aax2342. URL <https://doi.org/10.1126/science.aax2342>.
- O’Keefe, C., Ramakrishnan, K., Tay, J., and Winter, C. Law-following AI: Designing AI agents to obey human laws. *Fordham Law Review*, 94(1):57–129, 2025. URL <https://ir.lawnet.fordham.edu/flr/vol94/iss1/2/>.
- Omohundro, S. M. The basic ai drives. In Wang, P., Goertzel, B., and Franklin, S. (eds.), *Artificial General Intelligence 2008: Proceedings of the First AGI Conference*, volume 171 of *Frontiers in Artificial Intelligence and Applications*, pp. 483–492, Amsterdam, 2008. IOS Press. ISBN 978-1-58603-833-5. URL <https://doi.org/10.1007/978-1-58603-833-5>.

- https://selfawaresystems.com/wp-content/uploads/2008/01/ai_drives_final.pdf.
- OpenAI. Openai GPT-4.5 system card. Technical report, OpenAI, February 2025a. URL <https://cdn.openai.com/gpt-4-5-system-card-2272025.pdf>. System Card, February 27, 2025.
- OpenAI. Openai o3 and o4-mini system card. Technical report, OpenAI, 2025b. URL <https://cdn.openai.com/pdf/2221c875-02dc-4789-800b-e7758f3722c1/o3-and-o4-mini-system-card.pdf>. System Card, April 16, 2025.
- OpenAI. Openai charter. <https://openai.com/charter/>, n.d. Accessed: 2025-6-26.
- Oremus, W. Elon musk promised an anti-‘woke’ chatbot. it’s not going as planned. <https://www.washingtonpost.com/technology/2023/12/23/grok-ai-elon-musk-x-woke-bias/>, 2023. [Accessed 18-06-2025].
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askell, A., Welinder, P., Christiano, P., Leike, J., and Lowe, R. Training language models to follow instructions with human feedback, 2022. URL <https://arxiv.org/abs/2203.02155>.
- Plurality Institute. Exploring opportunities for llms in public discourse. Plurality Institute (blog post), April 2025. URL <https://www.plurality.institute/blog-posts/exploring-opportunities-for-large-language-models-in-public-discourse>.
- Prabhakaran, V., Mitchell, M., Gebru, T., and Gabriel, I. A human rights-based approach to responsible ai, 2022. URL <https://arxiv.org/abs/2210.02667>.
- Prunkl, C. Human autonomy at risk? an analysis of the challenges from AI. *Minds and Machines*, 34(3):1–21, 2024. doi: 10.1007/s11023-024-09665-1.
- Raji, I. D., Kumar, I. E., Horowitz, A., and Selbst, A. D. The fallacy of AI functionality. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’22, pp. 959–972, New York, NY, USA, June 2022. Association for Computing Machinery. ISBN 978-1-4503-9352-2. doi: 10.1145/3531146.3533158. URL <https://doi.org/10.1145/3531146.3533158>.
- Roytburg, D. and Miller, B. Mind the gap! pathways towards unifying AI safety and ethics research, 2025. URL <https://arxiv.org/abs/2512.10058>.
- Russell, S. Will they make us better people? Edge.org, 2015. URL <https://www.edge.org/response-detail/26157>. Response to the 2015 Annual Question, “What Do You Think About Machines That Think?”.
- Russell, S. *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking, October 2019.
- Russell, S. and Norvig, P. *Artificial Intelligence: A Modern Approach (3rd Edition)*. Prentice Hall, 2010.
- Santurkar, S., Durmus, E., Ladhak, F., Lee, C., Liang, P., and Hashimoto, T. Whose opinions do language models reflect? In *ICML*, 2023.
- Shah, R., Varma, V., Kumar, R., Phuong, M., Krakovna, V., Uesato, J., and Kenton, Z. Goal misgeneralization: Why correct specifications aren’t enough for correct goals, 2022. URL <https://arxiv.org/abs/2210.01790>.
- Sharma, N., Liao, Q. V., and Xiao, Z. Generative echo chamber? Effect of llm-powered search systems on diverse information seeking. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, CHI ’24, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400703300. doi: 10.1145/3613904.3642459. URL <https://doi.org/10.1145/3613904.3642459>.
- Shen, H., Kneare, T., Ghosh, R., Alkie, K., Krishna, K., Liu, Y., Petridis, S., Peng, Y.-H., Qiwei, L., Si, C., Xie, Y., Bigham, J. P., Bentley, F., Chai, J., Lipton, Z. C., Mei, Q., Terry, M., Yang, D., Morris, M. R., Resnick, P., and Jurgens, D. Position: Towards bidirectional human-AI alignment. In *The Thirty-Ninth Annual Conference on Neural Information Processing Systems Position Paper Track*, 2025. URL <https://openreview.net/forum?id=PgA9rZoMY8>.
- Si, S., Zhao, H., Chen, G., Gao, C., Bai, Y., Wang, Z., An, K., Luo, K., Qian, C., Qi, F., Chang, B., and Sun, M. Aligning large language models to follow instructions and hallucinate less via effective data filtering, 2025. URL <https://arxiv.org/abs/2502.07340>.
- Smith, J. J., Satwani, A., Burke, R., and Fiesler, C. Recommend me? designing fairness metrics with providers. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’24, pp. 2389–2399, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400704505. doi: 10.1145/3630106.3659044. URL <https://doi.org/10.1145/3630106.3659044>.

- Soares, N. and Fallenstein, B. Aligning Superintelligence with Human Interests: A technical research agenda. Technical Report 2014–8, Machine Intelligence Research Institute, 2014. URL <https://intelligence.org/files/obsolete/TechnicalAgenda%5Bold%5D.pdf>.
- Stark, L. Animation and artificial intelligence. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*, pp. 1663–1671, Rio de Janeiro, Brazil, 2024. ACM. doi: 10.1145/3630106.3658995. URL <https://doi.org/10.1145/3630106.3658995>.
- Stray, J., Vendrov, I., Nixon, J., Adler, S., and Hadfield-Menell, D. What are you optimizing for? Aligning recommender systems with human values, 2021. URL <https://arxiv.org/abs/2107.10939>.
- Sucholutsky, I. and Griffiths, T. L. Alignment with human representations supports robust few-shot learning. In Oh, A., Naumann, T., Globerson, A., Saenko, K., Hardt, M., and Levine, S. (eds.), *Advances in Neural Information Processing Systems*, volume 36, pp. 73464–73479, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/e8ddc03b001d4c4b44b29bc1167e7fdd-Paper-Conference.pdf.
- Sutton, R. The reward hypothesis. <http://incompleteideas.net/rlai.cs.ualberta.ca/RLAI/rewardhypothesis.html>, 2004. URL <http://incompleteideas.net/rlai.cs.ualberta>. Accessed 2025-06-11.
- Tang, A. AI alignment cannot be top-down, 2025. URL <https://ai-frontiers.org/articles/ai-alignment-cannot-be-top-down>. Guest Commentary.
- Taylor, J., Yudkowsky, E., LaVictoire, P., and Critch, A. Alignment for advanced machine learning systems. Technical report, Machine Intelligence Research Institute, July 2016. URL <https://intelligence.org/files/AlignmentMachineLearning.pdf>.
- Terzis, P., Veale, M., and Gaumann, N. Law and the emerging political economy of algorithmic audits. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*, pp. 1–13, New York, NY, USA, 2024. ACM. ISBN 979-8-4007-0450-5/24/06. doi: 10.1145/3630106.3658970. URL <https://facctconference.org/static/papers24/facct24-85.pdf>.
- The Collective Intelligence Project. Introducing the collective intelligence project: Solving the transformative technology trilemma through governance r&d. <https://static1.squarespace.com/static/631d02b2dfa9482a32db47ec/t/63e01fcbf73bb003a722792f/1675632588509/CIP+Whitepaper.pdf>, 2023. White paper; accessed June 13, 2025.
- Turing, A. M. Intelligent machinery, A heretical theory. In Copeland, B. J. (ed.), *The Essential Turing: Seminal Writings in Computing, Logic, Philosophy, Artificial Intelligence, and Artificial Life, Plus the Secrets of Enigma*, pp. 472–475. Oxford University Press, Oxford, 2004. doi: 10.1093/oso/9780198250791.003.0018. URL <https://uberty.org/wp-content/uploads/2015/02/intelligent-machinery-a-heretical-theory.pdf>. Originally presented c. 1951 on the BBC programme *The '51 Society*.
- Vamplew, P., Dazeley, R., Foale, C., Firmin, S., and Mumery, J. Human-aligned artificial intelligence is a multi-objective problem. *Ethics and Information Technology*, 20(1):27–40, 2018. doi: 10.1007/s10676-017-9440-6. URL <https://doi.org/10.1007/s10676-017-9440-6>.
- Weidinger, L., Rauh, M., Marchal, N., Manzini, A., Hendricks, L. A., Mateos-Garcia, J., Bergman, S., Kay, J., Griffin, C., Bariach, B., Gabriel, I., Rieser, V., and Isaac, W. Sociotechnical safety evaluation of generative AI systems, 2023. URL <https://arxiv.org/abs/2310.11986>.
- Wittgenstein, L. *Philosophical Investigations*. Basil Blackwell, Oxford, 1953.
- Yarom, M., Bitton, Y., Changpinyo, S., Aharoni, R., Herzig, J., Lang, O., Ofek, E., and Szpektor, I. What you see is what you read? improving text-image alignment evaluation. In *NeurIPS*, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/056e8e9c8ca9929cb6cf198952bflddb-Paper-Conference.pdf.
- Yeh, M.-H., Wang, J., Du, X., Park, S., Tao, L., Im, S., and Li, Y. Challenges and future directions of data-centric AI alignment, 2024. URL <https://arxiv.org/abs/2410.01957>. ICML 2025 position paper.
- Yudkowsky, E. Coherent extrapolated volition. Technical report, The Singularity Institute, San Francisco, CA, 2004. URL <https://intelligence.org/files/CEV.pdf>.
- Yudkowsky, E. Artificial Intelligence as a positive and negative factor in global risk. In Rees, M. J., Bostrom, N., and Cirkovic, M. M. (eds.), *Global Catastrophic Risks*, pp. 308–345. Oxford University Press, July 2008. ISBN

978-0-19-857050-9. doi: 10.1093/oso/9780198570509.003.0021.

Yudkowsky, E. Value is fragile, January 2009. URL <https://www.lesswrong.com/posts/GNnHHmm8EzePmKzPk/value-is-fragile>.

Yudkowsky, E. The ai alignment problem: Why it’s hard, and where to start. Talk transcript, Machine Intelligence Research Institute, May 2016. URL <https://intelligence.org/files/AlignmentHardStart.pdf>. Stanford University Symbolic Systems Distinguished Speaker series.

Zhi-Xuan, T., Carroll, M., Franklin, M., and Ashton, H. Beyond preferences in AI alignment. *Philosophical Studies*, 182(7):1813–1863, July 2025. doi: 10.1007/s11098-024-02249-w. URL <https://link.springer.com/article/10.1007/s11098-024-02249-w>. Published online 09 Nov 2024.

Zhuang, S. and Hadfield-Menell, D. Consequences of misaligned AI. In Larochelle, H., Ranzato, M., Hadsell, R., Balcan, M., and Lin, H. (eds.), *Advances in Neural Information Processing Systems*, volume 33, pp. 15763–15773, 2020. ISBN 9781713829546. URL <https://proceedings.neurips.cc/paper/2020/file/b607ba543ad05417b8507ee86c54fcb7-Paper.pdf>.