
From Outcomes to Actions: Leveraging Hindsight for Long-Horizon Language Agent Training

Zishang Jiang¹ Tingyun Li¹ Jinyi Han² Xinyi Wang¹ Sihang Jiang³ Yizhou Ying³ Xiaojun Meng⁴
Jiansheng Wei⁴ Jiaqing Liang¹ Yanghua Xiao³

Abstract

Reinforcement learning (RL) has become a widely adopted technique for improving large language models (LLMs) on complex tasks. Despite this progress, existing RL methods still face challenges in training agents with longer-horizon interactions. One major bottleneck is distinguishing the contribution of different actions in long-horizon interaction, leading to high optimization variance. To address this, we introduce a novel policy gradient method, **Hindsight Policy Optimization (HPO)**, that projects both the current policy distribution and the hindsight distribution into an intent space and extracts low-variance learning signals from the Wasserstein distance between them. We theoretically and empirically show that aggregating semantically similar states and actions in the intent space yields a bounded-variance estimator and improves policy performance stably. Our code is available online¹.

1. Introduction

Reinforcement learning (RL) has become a widely adopted technique for improving large language models (LLMs) on complex decision-making tasks. Compared to earlier prompt-based approaches, RL allows LLM-based agents to actively explore and learn from interaction with the environment, and has shown strong potential in applications such as tool use (Jin et al., 2025; Cao et al., 2025), smart device operation (Luo et al., 2025), and interactive environments (Abdulhai et al.; Hu et al., 2024).

¹School of Data Science, Fudan University ²Shanghai Institute of Artificial Intelligence for Education, East China Normal University ³College of Computer Science and Artificial Intelligence, Fudan University ⁴Huawei Large Model Data Technology Lab. Correspondence to: Jiaqing Liang <liangjiaqing@fudan.edu.cn>.

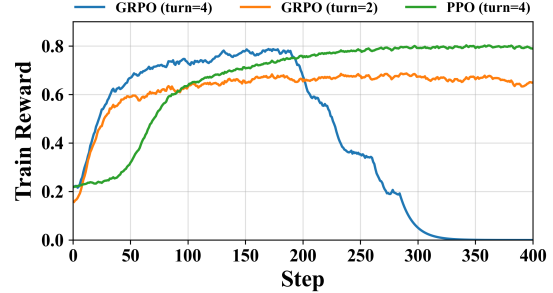


Figure 1. Existing RL methods struggle in long-horizon interactions. GRPO becomes unstable with longer-horizon interactions, while PPO remains stable but converges slowly due to critic warmup.

Despite this progress, existing RL algorithms still face challenges in training agents for longer-horizon interactions. As illustrated in Figure 1, GRPO (Shao et al., 2024) can maintain stable training under shorter interaction horizons, but becomes unstable as the horizon length increases, often exhibiting sudden reward collapse (Jin et al., 2025; Xue et al., 2025; Sun et al., 2025). In contrast, PPO (Schulman et al., 2017) achieves more stable training behavior, but typically converges more slowly due to the additional requirement of training a critic model.

One major bottleneck is distinguishing the contribution of different actions in long-horizon interaction, leading to high optimization variance. Unlike single-step tasks, multi-turn decision-making requires attributing outcomes to a sequence of interdependent actions, where effective and ineffective actions are often interleaved within the same trajectory (Zhang et al., 2025). However, RL methods that rely solely on outcome rewards, such as GRPO, typically propagate the final outcome uniformly across all actions, which may mistakenly reinforce some ineffective actions and thereby destabilize training (Zhang et al., 2025). In contrast, PPO uses a critic to better distinguish actions, leading to more stable training, but it converges more slowly due to the warm-up required by the critic (Jin et al., 2025).

In this work, we define the hindsight distribution, which is constructed from only trajectories that eventually succeed, and show that existing RL methods with Monte Carlo es-

timization can be interpreted as a *point-wise* comparison of the **discrepancy between the current policy distribution and this hindsight distribution in a discrete state-action space**. This point-wise comparison **ignores the underlying semantic structure of the language action space**, resulting in unbounded variance in exponentially large language action spaces.

Although the action space of large language models is vast, the set of underlying intents they can express is comparatively limited. For example, *Who is the president of the United States?* and *I want to know who the U.S. president is* correspond to distinct language actions but convey nearly identical semantic intent.

Based on this insight, we introduce a novel policy gradient method, **Hindsight Policy Optimization (HPO)**, that projects both the current policy distribution and the hindsight distribution into an intent space and measures their discrepancy using the Wasserstein distance (Villani et al., 2008). We theoretically show that aggregating semantically similar states and actions in this lower-dimensional intent space improves statistical efficiency and yields a bounded-variance estimator.

Our contributions are as follows:

1. We reformulate existing Monte Carlo policy gradient methods from a hindsight distribution perspective, and show that point-wise comparisons in the discrete language action space can lead to unbounded variance due to ignoring the underlying semantic structure.
2. We propose HPO, a principled approach that constructs hindsight distributions in an intent space and derives bounded-variance learning signals for smoother and more sample-efficient policy optimization, with negligible computational overhead.
3. We evaluate HPO on multiple long-horizon tasks and show that it substantially reduces training noise and stabilizes optimization. Further experiments show that such learning signals are interpretable and remain sufficient to drive effective policy learning even without outcome rewards.

2. Preliminaries

2.1. Task Formulation

We formulate language action tasks as a Markov decision process (MDP) $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma)$, where $\mathcal{S}$ denotes the state space, $\mathcal{A}$ is the action space of natural language actions, $\mathcal{T} : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ is the transition function, $\mathcal{R}$ is the reward function, and $\gamma \in (0, 1]$ is the discount factor. In the RLVR setting, rewards are sparse and only provided at the end of each episode and the discount factor is typically set to

$\gamma = 1$. When the context is clear, we omit γ for notational simplicity. An agent interacts with the environment over discrete time steps. At each step t , the agent observes the current state $s_t \in \mathcal{S}$ and selects an action $a_t \sim \pi_\theta(\cdot | s_t)$. The environment then transitions to the next state s_{t+1} according to $\mathcal{T}(s_{t+1} | s_t, a_t)$. The objective is to maximize the expected cumulative reward under the policy π_θ .

2.2. Value Function

For $M_\pi = (S_0, A_0, S_1, A_1, \dots)$, there is an occupancy measure ρ_π which satisfies $\rho_\pi(s, a) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(S_t = s, A_t = a)$. Moreover, there is a one-to-one correspondence between those measures and the policies. The state-value function under policy π is defined as $V_\pi(s) \triangleq \mathbb{E}_\pi[\sum_{t=0}^{\infty} \gamma^t r_t | S_0 = s]$, which measures the expected discounted return starting from state s . Similarly, the action-value function is defined as $Q_\pi(s, a) \triangleq \mathbb{E}_\pi[\sum_{t=0}^{\infty} \gamma^t r_t | S_0 = s, A_0 = a]$.

2.3. Wasserstein Distance

Wasserstein distance, originating from optimal transport theory (Villani et al., 2008), measures the discrepancy between two probability distributions by the minimum cost of transporting mass from one distribution to the other. Given two probability measures μ_A and μ_B on a metric space $(\mathcal{X}, \text{dis})$, the 1-Wasserstein distance is defined as

$$W_1(\mu_A, \mu_B) = \inf_{\pi \in \Pi(\mu_A, \mu_B)} \int_{\mathcal{X} \times \mathcal{X}} \text{dis}(x_A, x_B) d\pi(x_A, x_B), \quad (1)$$

where $\Pi(\mu_A, \mu_B)$ is the set of couplings whose marginals are μ_A and μ_B . Intuitively, each coupling can be viewed as a possible transportation plan that describes how mass from μ_A is assigned to mass in μ_B . In practice, W_1 is often computed via its dual formulation:

$$W_1(\mu_A, \mu_B) = \sup_{\|f\|_L \leq 1} (\mathbb{E}_{x \sim \mu_A}[f(x)] - \mathbb{E}_{x \sim \mu_B}[f(x)]), \quad (2)$$

where the maximization is over all 1-Lipschitz functions f . The optimal function f^* in this dual problem is referred to as the *Kantorovich potential* (Villani et al., 2008).

3. The Challenge of Long-Horizon Language Agent Reinforcement Learning

Existing studies (Jin et al., 2025; Xue et al., 2025; Sun et al., 2025) have found that RL methods become increasingly unstable as the interaction horizon grows. In this section, we show that this instability arises from the high-variance gradient noise induced by long-horizon interactions, which free-critic methods struggle to mitigate and may even further exacerbate.

3.1. Rethinking the Objective in RL for LLMs

For a unified analysis, common policy optimization methods used in RL for LLMs, including GRPO, REINFORCE++ (Hu et al., 2025), and PPO, can all be expressed in the following generic policy gradient form like REINFORCE (Williams, 1992):

$$\nabla_{\theta} \mathcal{J}(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (\hat{Q}(s_t, a_t) - b_t) \right] \quad (3)$$

where $\hat{Q}(s_t, a_t)$ denotes an estimator of $Q_{\pi}(s_t, a_t)$.

For example, the objective in GRPO can be interpreted as repeatedly sampling G trajectories from a shared initial state s_0 . For analytical convenience, we assume equal return variance across different s_0 . Concretely, GRPO optimizes

$$\nabla_{\theta} \mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E}_{\{\tau_i\}_{i=1}^G \sim \pi_{\theta}} \left[\frac{1}{G} \sum_{i=1}^G \sum_{t=0}^{\infty} \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) (\hat{Q}(s_{i,t}, a_{i,t}) - b_{i,t}) \right], \quad (4)$$

where $\hat{Q}(s_{i,t}, a_{i,t}) = R(\tau_i)$ and $b_{i,t} = \frac{1}{G} \sum_{i=1}^G R(\tau_i)$, which can be viewed as a Monte Carlo (MC) estimate $\bar{V}_{\pi}(s_0)$ of $V_{\pi}(s_0)$.

Although MC estimation offers a straightforward approximation for Q_{π} , it suffers from excessive variance by accumulating stochasticity from future actions (Sutton et al., 1998). For example, even if the agent makes a correct decision at the current step, an erroneous decision at some later time step may still cause the entire trajectory to fail, making the same current action may receive different terminal return. While current methods, such as GRPO, typically employ a fixed baseline to mitigate this variance without introducing bias, we show below that the effectiveness of this approach diminishes significantly in long-horizon interactions.

Lemma 3.1 (Variance Decomposition). *Consider the REINFORCE gradient estimator $\hat{G}_b$ with the baseline b ,*

$$\hat{G}_b = g(a) (\hat{Q}(s, a) - b(s)), \quad (5)$$

where $a \sim \pi(\cdot | s)$ and $g(a) = \nabla_{\theta} \log \pi_{\theta}(a | s)$.

Let $\|\cdot\|$ denotes the Euclidean norm. Then its conditional variance admits the decomposition

$$\text{Var}(\hat{G}_b | s) = \mathbb{E}[\|g(a)\|^2 | s] \cdot \text{Var}(\hat{Q}(s, a) - b(s) | s) \quad (6)$$

Remark 3.2. Eq. (6) shows that the variance of the policy gradient can be decomposed into the product of **the gradient norm** and **the variance of the Q estimator**.

The variance of the Q estimator. The variance of the policy gradient estimator is primarily driven by the variance of $\hat{Q}(s, a)$. In particular, when $\hat{Q}$ is estimated using MC

returns over entire trajectories, the variance grows with the horizon length, because it aggregates the stochasticity induced by action sampling over all future time steps.

By using an appropriate baseline, $\text{Var}(\hat{G}_b | s)$ can be reduced without introducing additional bias (Sutton et al., 1999).

Theorem 3.3 (Optimal Baseline and Excess Variance). *Among all scalar baselines $b(s)$ independent of the sampled action a , the conditional variance $\text{Var}(\hat{G}_b | s)$ is minimized by*

$$b^*(s) = \frac{\mathbb{E}[\|g(a)\|^2 \hat{Q}(s, a) | s]}{\mathbb{E}[\|g(a)\|^2 | s]}. \quad (7)$$

Moreover, for any alternative baseline $\tilde{b}(s)$, the increase in variance admits the exact expression

$$\text{Var}(\hat{G}_{\tilde{b}} | s) - \text{Var}(\hat{G}_{b^*} | s) = \mathbb{E}[\|g(a)\|^2 | s] (\tilde{b}(s) - b^*(s))^2. \quad (8)$$

We remark that even though the optimal state-only baseline is known, it is rarely used in practice due to its computational intractability (Li et al., 2024; Duan et al., 2016). Rather, for both computational and conceptual benefit, the choice of $\mathbb{E}[\hat{Q}(s, a) | s]$ is often used (Schulman et al., 2017; Wu et al., 2018). In particular, when $\|g(a)\|^2$ and $\hat{Q}(s, a)$ are loosely correlated, this baseline is close to the optimal baseline. For ease of analysis and empirical evaluation, we follow this setting.

Under this proxy, the minimum-variance baseline for MC returns is the state-dependent value function $V_{\pi}(s_t) = \mathbb{E}[R | s_t]$. In contrast, critic-free methods such as GRPO employ a fixed baseline $\bar{V}_{\pi}(s_0)$ that is shared across all subsequent time steps, rather than adapting to the current state. In long-horizon, tool-augmented interactions, $V_{\pi}(s_t)$ can vary dramatically across states: for instance, after a successful search, the probability of producing a correct final answer becomes close to one, whereas a failed search leaves the model highly uncertain and substantially lowers $V_{\pi}(s_t)$. As the interaction horizon increases, the mismatch $(V_{\pi}(s_t) - \bar{V}_{\pi}(s_0))^2$ can grow large, directly weakening the variance-reduction effect of the baseline by Theorem 3.3 and, in extreme cases, leading to higher gradient variance than using no baseline at all.

The gradient norm. Beyond the variance of the Q estimator, the policy-gradient norm can further magnify optimization noise. With rounds of external tool feedback, some low-probability actions are increasingly sampled (e.g., meaningless or repetitive tokens) (Xue et al., 2025). Due to the logarithmic form of the policy gradient, these rare actions can induce disproportionately large, and even unbounded gradient norms, which further amplify the noise variance caused by state mismatch in Eq. (6). This perspective also helps explain why filtering strategies that remove

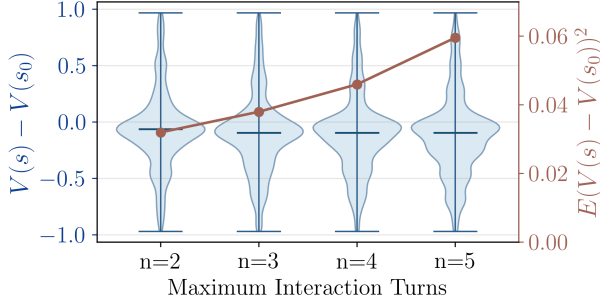


Figure 2. As the maximum interaction turns n grows, the mismatch between $V_\pi(s_t)$ and the fixed baseline $\bar{V}_\pi(s_0)$ can become increasingly severe.

low-probability samples can mitigate collapse (Xue et al., 2025).

Summary

As the interaction horizon increases, Monte Carlo estimators exhibit high variance that is hard to reduce in existing methods, often leading to unstable training. This motivates the need for finding a lower-variance estimator.

4. Training LLM Agents with HPO

In this section, we introduce a novel policy gradient estimator. Unlike Monte Carlo estimators defined over discrete action spaces, it constructs learning signals in a semantically aggregated intent space. By sharing statistical information at the semantic level, it improves statistical efficiency and substantially reduces estimation variance.

4.1. Policy Gradient from the Hindsight Perspective

Before introducing our method, we first define the hindsight distribution and show that the policy gradient can be equivalently reformulated from the perspective of the hindsight distribution.

Definition 4.1 (Hindsight Policy Distribution). For any state s , action a , and return value z , let $h_z^\pi(a | s)$ denote the conditional probability that the first action equals a when sampling a trajectory from π starting at s , given that the realized return equals z . In the binary-reward setting considered in this work, we take $z = 1$ and define the corresponding hindsight occupancy measure as

$$\rho_\pi^h(s, a) = \sum_{t=0}^{\infty} \mathbb{P}_\pi \left((s_t, a_t) = (s, a) \mid \sum_{t'=t}^{\infty} r_{t'} = 1 \right). \quad (9)$$

Intuitively, in the binary-reward setting, this distribution is obtained by retaining only state-action pairs from trajectories that eventually succeed. For continuous-reward settings,

we provide the corresponding distribution definition in Appendix B.

Lemma 4.2. Consider the objective of policy optimization $\mathcal{J}(\theta) = \mathbb{E}_{\tau \sim \pi} [\sum_{t \geq 0} r_t]$. Its policy gradient admits the following two equivalent forms:

$$\nabla_\theta \mathcal{J}(\theta) = \mathbb{E}_{(s,a) \sim \rho_\pi} [Q_\pi(s, a) \nabla_\theta \log \pi_\theta(a | s)] \quad (10)$$

$$= \mathbb{E}_{(s,a) \sim \rho_\pi} \left[-\frac{\delta KL(\rho_\pi^h || \rho_\pi)}{\delta \rho_\pi(s, a)} \nabla_\theta \log \pi_\theta(a | s) \right] \quad (11)$$

where $-\frac{\delta KL(\rho_\pi^h || \rho_\pi)}{\delta \rho_\pi(s, a)} = \frac{\rho_\pi^h(s, a)}{\rho_\pi(s, a)}$ denotes the point-wise variational derivative of $KL(\rho_\pi^h || \rho_\pi)$ with respect to $\rho_\pi(s, a)$.

A detailed proof is provided in Appendix D.3. Lemma 4.2 provides an alternative interpretation of policy gradients from the perspective of the hindsight distribution, where policy updates are driven by comparing the current policy with the hindsight policy. However, in large discrete spaces, such point-wise estimates suffer from high variance unless an impractically large number of samples is used.

Although the action space of LLMs is exponentially large, the space of intents they can express is significantly lower-dimensional. Many distinct actions are semantically similar and lead to comparable outcomes, for example, *Who is the president of the United States?* and *I want to know who the U.S. president is*, correspond to different language actions but express nearly identical underlying intent.

Based on this insight, we argue that **aggregating semantically similar states and actions in a low-dimensional intent space** can improve statistical efficiency and thereby reduce variance. Therefore, we propose HPO, a novel framework that performs policy updates by comparing the hindsight policy with the current policy in the intent space. The overall architecture of HPO is shown in Figure 3.

4.2. Intent Space Construction

Given a state space $\mathcal{S}$ and an action space $\mathcal{A}$, each state-action pair (s, a) is embedded into a d -dimensional Hilbert semantic space via a pretrained encoder $\Phi : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}^d$, where the Euclidean distance serves as the metric of the embedding space. In this semantic space, interactions that are lexically different but semantically equivalent are mapped to nearby regions.

In the binary-reward setting, the policy distribution ρ_π and hindsight distribution ρ_π^h are constructed as empirical distributions over embedded state-action pairs in the intent

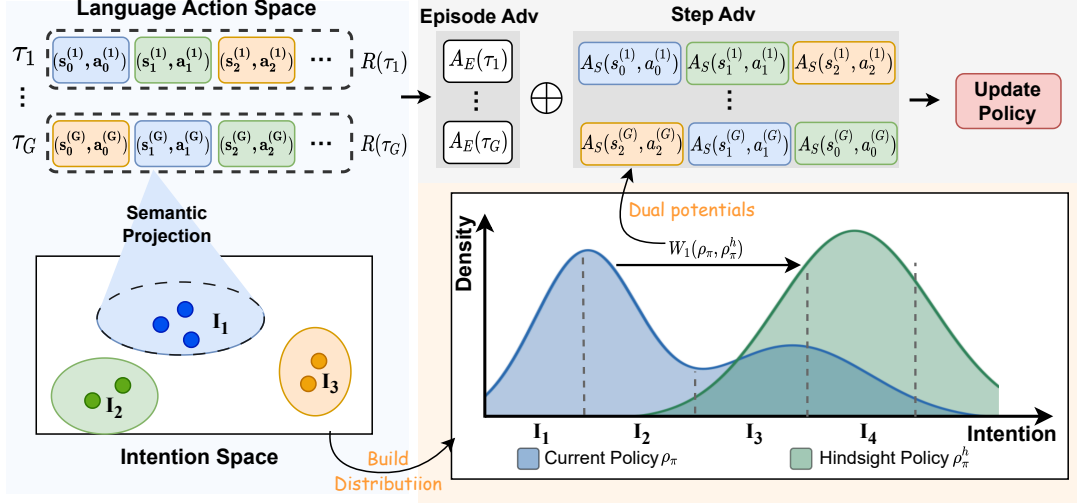


Figure 3. Illustration of HPO. HPO performs semantic projection and compares the policy distribution with the hindsight distribution in the intent space. By solving the Wasserstein distance between the two distributions, it obtains step-level advantages, which are combined with episode-level advantages to update the policy.

space.

$$\begin{aligned}\rho_\pi &= \sum_{t=0}^{\infty} \mathbb{E}_\pi [\delta_{\Phi(s_t, a_t)}], \\ \rho_\pi^h &= \sum_{t=0}^{\infty} \mathbb{E}_\pi \left[\mathbf{1} \left(\sum_{t'=t}^{\infty} r_{t'} = 1 \right) \delta_{\Phi(s_t, a_t)} \right],\end{aligned}\quad (12)$$

where $\delta_{\Phi(s_t, a_t)}$ denotes the Dirac measure centered at the intent embedding $\Phi(s_t, a_t)$.

Data for constructing the hindsight distribution. Although constructing the hindsight distribution from on-policy rollouts yields a closer approximation to the on-policy value function Q_π , it suffers from unstable estimation due to fluctuating successful rollouts. Alternatively, using an offline set of successful trajectories obtained via rejection sampling from the initial policy provides a more stable construction of the hindsight distribution, particularly in sparse-reward scenarios, where on-policy rollouts rarely yield successful trajectories. We refer to these two variants as **online HPO** and **offline HPO**, respectively.

4.3. Hindsight Policy Optimization

Compared with KL divergence, the Wasserstein distance is more suitable for measure spaces with an explicit distance metric, as it explicitly leverages the underlying geometry of the intent space and remains well defined even when the supports of the distributions do not overlap (Villani et al., 2008).

We replace the KL divergence defined over the discrete state-action space with the Wasserstein distance defined in

the intent space. Accordingly, Eq. (11) can be rewritten as

$$\nabla_\theta \mathcal{J}(\theta) = \mathbb{E}_{(s,a) \sim \rho_\pi} \left[-\frac{\delta W_1(\rho_\pi \| \rho_\pi^h)}{\delta \rho_\pi(s,a)} \nabla_\theta \log \pi_\theta(a | s) \right] \quad (13)$$

where $\frac{\delta W_1(\rho_\pi \| \rho_\pi^h)}{\delta \rho_\pi(s,a)} = f^*(s, a)$ denotes the point-wise variational derivative of $W_1(\rho_\pi \| \rho_\pi^h)$ with respect to $\rho_\pi(s, a)$, i.e., the Kantorovich potential in Section 2.3.

$f^*(s, a)$ possesses several desirable properties. Notably, it satisfies the 1-Lipschitz continuity condition in the intent space, implying that behaviors mapped to similar intents yield close potential values. Furthermore, it is also computationally efficient (Ambrosio et al., 2005; Alvarez-Melis & Fusi, 2021) in that $f^*(s, a)$ can be recovered directly from the dual solution of the Wasserstein distance calculation at virtually no additional cost (Alvarez-Melis & Fusi, 2021; Just et al., 2023).

As proved in Section 4.4, $-\frac{\delta W_1(\rho_\pi \| \rho_\pi^h)}{\delta \rho_\pi(s,a)} = -f^*(s, a)$, which is defined in the intent space, yields a learning signal with lower variance, at the cost of introducing a certain degree of bias. In contrast, the learning signal $-\frac{\delta \text{KL}(\rho_\pi^h \| \rho_\pi)}{\delta \rho_\pi(s,a)} = \frac{\rho_\pi^h(s,a)}{\rho_\pi(s,a)}$, which is defined over the discrete action space, is unbiased but typically exhibits much higher variance.

Since GRPO is currently one of the most widely adopted algorithms in RLVR, we choose it as the base optimization algorithm and apply the proposed hindsight distribution to its policy update.

Concretely, the optimization terms induced by the KL divergence and the Wasserstein distance admit intuitive interpre-

tations as *episode-level* and *step-level* advantages in GRPO, respectively.

KL term (Episode-level advantage). The KL-based optimization term corresponds to a Monte Carlo estimate of $Q_\pi(s, a)$. Therefore, we directly reuse the episode-level advantage originally defined in GRPO as the learning signal for the KL term:

$$A_E(\tau_i) = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}, \quad (14)$$

Wasserstein term (Step-level advantage). For the Wasserstein-based optimization term, we obtain the Kantorovich potential $-\frac{\delta W_1(\rho_\pi \parallel \rho_\pi^h)}{\delta \rho_\pi(s, a)} = -f^*(s_t, a_t)$ defined in the intent space. To ensure compatibility in scale with the episode-level advantage and to enable a meaningful interpolation between the two signals, we normalize the potential into a step-level advantage:

$$A_S(s_t, a_t) = \frac{-f^*(s_t, a_t) - \text{mean}(-f^*)}{\text{std}(f^*)}. \quad (15)$$

Final advantage formulation. We define the final advantage as a weighted interpolation between the episode-level and step-level advantages, to achieve a better bias–variance trade-off:

$$\hat{A}(s_t^{(i)}, a_t^{(i)}) = \frac{1}{1 + \omega} \left[A_E(\tau_i) + \omega \cdot A_S(s_t^{(i)}, a_t^{(i)}) \right], \quad (16)$$

where $\omega > 0$ controls the relative contribution of the step-level hindsight signal. By tuning ω , this formulation provides an explicit mechanism to trade off variance reduction against bias, enabling more stable and effective long-horizon optimization.

Let $\mathcal{D}$ denote the training prompt distribution, and let $q \sim \mathcal{D}$ denote a sampled task prompt. Then the policy optimization objective of HPO is:

$$\mathcal{J}_{\text{HPO}}(\theta) = \mathbb{E}_{q \sim \mathcal{D}, \{\tau^{(i)}\}_{i=1}^G \sim \pi_{\text{old}}(\cdot | q)} \frac{1}{G} \sum_{i=1}^G \frac{1}{|\tau^{(i)}|} \sum_{t=1}^{|\tau^{(i)}|} \min \left[w_t^{(i)}(\theta) \hat{A}_t^{(i)}, \text{clip} \left(w_t^{(i)}(\theta), 1 \pm \epsilon \right) \hat{A}_t^{(i)} \right], \quad (17)$$

where $w_t^{(i)} = \frac{\pi_\theta(y_t^{(i)} | q, y_{<t}^{(i)})}{\pi_{\text{old}}(y_t^{(i)} | q, y_{<t}^{(i)})}$ is the importance sampling ratio between the current policy and the behavior policy.

The complete algorithm is summarized in Algorithm 1.

4.4. Theoretical Analysis

Lemma 4.3. Let ρ_π be a state-action occupancy measure, and define $D = \sup_{(s,a), (s',a') \in \rho_\pi} d((s,a), (s',a'))$. Let

Algorithm 1 Hindsight Policy Optimization

Input: initial policy model $\pi_{\theta_{\text{init}}}$; task prompts $\mathcal{D}$; hyperparameters $\varepsilon, \beta, \mu, \omega$
 policy model $\pi_\theta \leftarrow \pi_{\theta_{\text{init}}}$
for step = 1, . . . , M **do**
 Sample a batch $\mathcal{D}_b$ from $\mathcal{D}$
 Update the old policy model $\pi_{\theta_{\text{old}}} \leftarrow \pi_\theta$
 Sample outputs $\{o_i\}_{i=1}^G \sim \pi_{\theta_{\text{old}}}(\cdot | q)$ for each $q \in \mathcal{D}_b$
 Compute rewards $\{r_i\}_{i=1}^G$ for each output o_i
 Get embeddings by encoding $(s, a) \in \{\tau_i\}_{i=1}^G$ with Φ
 Construct distributions ρ_π, ρ_π^h using Eq. (12)
 Get f^* from computing $\mathcal{W}(\rho_\pi, \rho_\pi^h)$
 Compute $\hat{A}_t^{(i)}$ according to Eq. (16)
 for iteration = 1, . . . , μ **do**
 Update the policy π_θ with the objective in Eq. (17).
 end for
end for
Output: π_θ

$f^*(s, a)$ denote the $-\frac{\delta W_1(\rho_\pi \parallel \rho_\pi^h)}{\delta \rho_\pi(s, a)}$, and $g(s, a)$ denote the $-\frac{\delta KL(\rho_\pi^h \parallel \rho_\pi)}{\delta \rho_\pi(s, a)}$, then we have

$$\text{Var}_{\rho_\pi}(f^*(s, a)) \leq D^2/4. \quad (18)$$

$$\text{Var}_{\rho_\pi}(g(s, a)) = \int \frac{\rho_\pi^h(s, a)^2}{\rho_\pi(s, a)} d(s, a) - 1 = \chi^2(\rho_\pi^h \parallel \rho_\pi). \quad (19)$$

where $\chi^2(\cdot \parallel \cdot)$ denotes the chi-squared divergence.

The lemma highlights a sharp contrast between the learning signals induced by Wasserstein and KL. The $f^*(s, a)$ is 1-Lipschitz over the intent space and thus admits a variance bound $\text{Var}_{\rho_\pi}(f^*) \leq D^2/4$. In contrast, the $g(s, a) = \rho_\pi^h(s, a)/\rho_\pi(s, a)$ has variance characterized by $\chi^2(\rho_\pi^h \parallel \rho_\pi)$, which can be *arbitrarily large*.

This contrast explains the instability of KL-based signals in large discrete action spaces. The KL gradient depends on the pointwise ratio $\rho_\pi^h(s, a)/\rho_\pi(s, a)$, whose variance is governed by $\chi^2(\rho_\pi^h \parallel \rho_\pi)$ and can become arbitrarily large when $\rho_\pi(s, a)$ is small. Since discrete actions are treated as isolated atoms, no statistical sharing occurs across actions, making reliable estimation of such ratios sample-inefficient, especially in the presence of rare but important events. In contrast, the Wasserstein signal leverages the geometry of the intent space to compare distributions via mass transport, allowing probability to be matched across nearby intents without pointwise division. This yields a smoother learning signal with variance bounded by the metric diameter D , providing a principled justification for HPO’s intent-space formulation.

Table 1. Performance on SearchQA. Results are averaged over 3 random seeds. Compared with other baselines, HPO achieves consistent improvements.

Method	General QA				Multi-Hop QA					Avg.
	NQ	TriviaQA	PopQA	Avg.	HotpotQA	2wiki	Musique	Bamboogle	Avg.	
<i>Qwen2.5-3B-Instruct</i>										
Base	6.4	21.0	13.3	13.6	7.6	7.4	2.7	11.2	7.2	9.9
Search-R1	34.1	54.5	37.8	42.1	32.4	31.9	10.3	26.4	25.3	32.5
ZeroSearch	41.4	57.4	44.8	47.9	27.4	30.0	9.8	11.1	19.6	31.7
PPO	36.2 $\pm$ 1.3	57.5 $\pm$ 2.8	40.7 $\pm$ 2.8	44.8 $\pm$ 1.9	32.8 $\pm$ 2.4	33.0 $\pm$ 1.3	11.6 $\pm$ 1.6	30.4 $\pm$ 2.1	27.0 $\pm$ 1.6	34.6 $\pm$ 1.8
GRPO	37.2 $\pm$ 2.4	56.5 $\pm$ 3.4	42.4 $\pm$ 2.4	45.4 $\pm$ 2.0	32.0 $\pm$ 3.3	33.8 $\pm$ 2.0	12.3 $\pm$ 1.0	28.4 $\pm$ 1.5	26.6 $\pm$ 1.2	34.7 $\pm$ 1.5
HPO_{on}	41.2 $\pm$ 1.4	59.8 $\pm$ 2.4	44.0 $\pm$ 1.8	48.3 $\pm$ 1.2	36.2 $\pm$ 1.8	34.2 $\pm$ 2.2	14.4 $\pm$ 1.4	32.7 $\pm$ 3.7	29.4 $\pm$ 2.0	37.5 $\pm$ 1.6
HPO_{off}	43.8 $\pm$ 0.1	61.1 $\pm$ 1.1	45.1 $\pm$ 0.6	50.0 $\pm$ 0.2	37.6 $\pm$ 0.7	35.1 $\pm$ 0.9	15.5 $\pm$ 1.1	36.0 $\pm$ 1.2	31.1 $\pm$ 1.2	39.2 $\pm$ 0.8
<i>Qwen2.5-7B-Instruct</i>										
Base	11.9	31.1	20.3	21.1	16.5	10.9	7.0	21.6	14.0	17.0
Search-R1	39.3	61.0	39.7	46.7	37.0	<u>41.4</u>	14.6	36.8	32.5	38.5
ZeroSearch	43.6	61.8	51.5	52.3	34.6	35.2	18.4	27.8	29.0	39.1
PPO	41.2 $\pm$ 2.0	64.0 $\pm$ 1.7	42.5 $\pm$ 2.4	49.2 $\pm$ 1.5	39.5 $\pm$ 2.1	40.0 $\pm$ 1.6	18.8 $\pm$ 1.0	37.6 $\pm$ 2.7	34.0 $\pm$ 1.8	40.5 $\pm$ 1.5
GRPO	42.2 $\pm$ 4.1	62.8 $\pm$ 2.5	42.5 $\pm$ 4.0	49.2 $\pm$ 3.5	38.5 $\pm$ 4.1	38.7 $\pm$ 3.2	17.1 $\pm$ 3.8	34.8 $\pm$ 4.0	32.3 $\pm$ 3.8	39.5 $\pm$ 3.6
HPO_{on}	48.1 $\pm$ 1.7	66.4 $\pm$ 0.6	46.3 $\pm$ 1.8	53.6 $\pm$ 1.2	44.4 $\pm$ 2.8	42.3 $\pm$ 2.2	20.3 $\pm$ 1.0	40.8 $\pm$ 2.8	37.0 $\pm$ 1.5	44.1 $\pm$ 1.3
HPO_{off}	48.2 $\pm$ 0.6	67.1 $\pm$ 1.1	47.9 $\pm$ 0.8	54.4 $\pm$ 0.4	43.2 $\pm$ 0.1	40.5 $\pm$ 0.4	19.3 $\pm$ 0.4	41.3 $\pm$ 1.1	36.1 $\pm$ 0.2	43.9 $\pm$ 0.3

Table 2. The results on Textcraft benchmark.

Method	Success Rate (%)	
	<i>Qwen2.5-3B-Ins</i>	<i>Qwen2.5-7B-Ins</i>
Base	14.0	42.0
PPO	76.0	89.0
GRPO	81.0	87.0
HPO_{on}	84.0	93.0
HPO_{off}	86.0	93.0

5. Experiments

5.1. Setup

Tasks and Benchmarks. We evaluate HPO on two long-horizon interactive tasks: SearchQA and TextCraft (Prasad et al., 2024). For SearchQA, we consider seven benchmarks, including single-hop QA (NQ (Kwiatkowski et al., 2019), TriviaQA (Joshi et al., 2017), and PopQA (Mallen et al., 2022)) and multi-hop QA (HotpotQA (Yang et al., 2018), 2Wiki (Ho et al., 2020), Musique (Trivedi et al., 2023), and Bamboogle (Press et al., 2023)). For TextCraft, we report the average success rate (%) for each query. Full settings details are provided in Appendix A.

Models and Baselines. We use Qwen2.5-3B/7B-Instruct (Qwen et al., 2025) as our base models. For HPO, we use Qwen3-Embedding-0.6B as encoder. We compare our approach against a diverse set of strong baselines. Specifically, we include **PPO**, a widely adopted actor-critic method that relies on an additional value model, as well as critic-free methods such as **GRPO**, which estimate advantages over groups of trajectories. For the SearchQA task, we further evaluate recent competitive models trained with different RL algorithms, including **Search-R1**,

ZeroSearch. Full training settings and hyperparameter details are provided in Appendix A.

5.2. Main Result

HPO significantly improves policy performance, especially on multi-hop tasks. As shown in Table 1, HPO improves average performance by 4%-7% over other baselines. Moreover, the improvements are especially pronounced on multi-hop benchmarks such as HotpotQA and Bamboogle. We analyze that these gains likely stem from HPO’s fine-grained, step-level advantages, which enable more accurate credit assignment by identifying and reinforcing truly valuable actions; this in turn translates into higher final answer accuracy on multi-hop questions.

Offline HPO achieves better performance than online HPO. As shown in Table 1, offline HPO is more stable, with smaller performance variance across benchmarks. Although online HPO can construct an on-policy hindsight distribution from newly collected rollouts, the number of successful samples used for constructing hindsight distribution fluctuates under stochastic sampling, leading to unstable estimation. Offline HPO instead constructs the hindsight from a fixed offline dataset, and it further mitigates the issue of sparse rewards: when no correct solution is found, both GRPO and online HPO cannot update policy effectively. In contrast, offline HPO can still update by constructing the hindsight from a fixed set of successful trajectories in the offline data, thus providing informative step-level advantages.

HPO consistently improves performance across long-horizon agent tasks. Table 2 reports results on the TextCraft benchmark. HPO achieves consistent improvements in average success rate over all baselines, demonstrat-

ing that HPO is effective across different task domains and interaction dynamics, highlighting its robustness and broad applicability to long-horizon agent settings.

5.3. Training Dynamics

Figure 4 shows the training reward dynamics on Qwen2.5-7B. GRPO exhibits unstable long-horizon training with abrupt reward collapses, as it fails to distinguish successful actions from failed ones within trajectories. PPO achieves more stable training by introducing a critic to identify valuable actions, but converges more slowly because the critic requires substantial warm-up to produce reliable value estimates. In contrast, HPO achieves fast and stable convergence by identifying valuable actions without training an additional critic model. Notably, offline HPO demonstrates the fastest early-stage reward growth, which can be attributed to its step-level advantage in identifying correct actions within failed trajectories, thereby enabling meaningful policy updates even in the absence of fully correct samples.

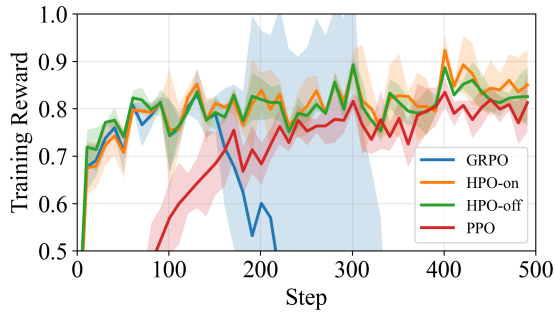


Figure 4. The reward dynamics of HPO and other baseline methods with Qwen2.5-7B over SearchQA, where curves and uncertainty ranges are computed using the mean and standard deviation over three random seeds, respectively. **GRPO exhibits unstable training, PPO converges more slowly, while HPO achieves fast and stable convergence.**

5.4. Interpretability Analysis of A_S

A_S alone is sufficient to drive effective policy updates. To isolate the contribution of A_E , we ablate it and re-run off-policy HPO over Qwen2.5-3B. In this setting, policy optimization is driven solely by the learning signal from A_S , while outcome rewards are computed only for monitoring and reporting training dynamics. As shown in Figure 5, optimizing with A_S alone yields stable and effective policy updates: the monitored training reward increases steadily throughout optimization, without exhibiting reward collapse. The final model is only marginally worse than the standard outcome-reward-optimized model, suggesting that the bias introduced by A_S is limited and practically acceptable. Overall, these results highlight the practical utility of A_S , particularly for open-ended tasks where outcome rewards

are difficult to specify, or for online settings where outcome evaluation is costly, since A_S can be learned from offline data that has already been evaluated. The **case study** in the Appendix E.2 further confirms that A_S reliably identifies valuable actions and provides sufficiently informative step-level signals for policy optimization.

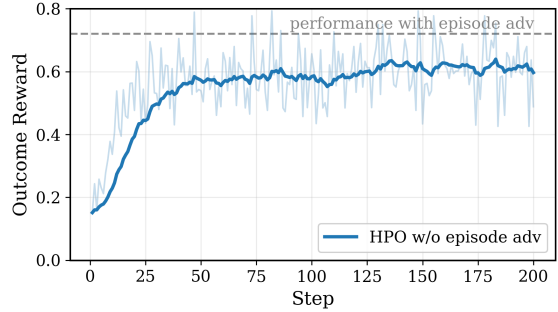


Figure 5. Training reward dynamics of HPO using only step-level advantages with Qwen2.5-3B. The dark blue curve is EMA-smoothed from the raw data. **Without episode-level rewards, training remains stable and effective, indicating that step-level advantages provide a reliable learning signal.**

5.5. Efficiency Analysis

We analyze the computational budget of HPO. HPO shares the same core architecture as GRPO, including multi-turn rollouts, computation of old and reference probabilities, and policy updates. The primary additions introduced by HPO are the step-relative advantage estimation components. To evaluate cost, we train a SearchQA LLM agent using Qwen2.5-7B-Instruct with Qwen3-0.6B-Embedding as the encoder, and record a per-iteration time breakdown.

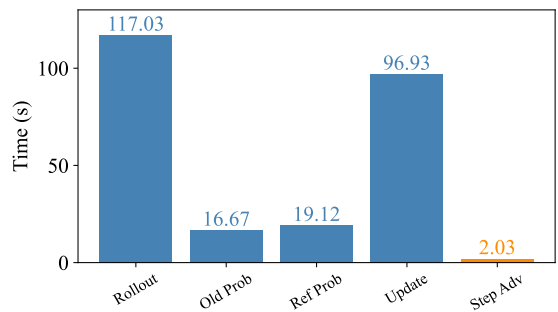


Figure 6. Per-iteration training time breakdown of HPO with Qwen2.5-7B. Blue bars denote components shared with GRPO, while orange bars denote HPO-specific additions. **The added overhead in HPO is negligible ($< 0.8\%$).**

As shown in Figure 6, the added components introduce negligible overhead. Compared with GRPO, HPO requires only an additional 2.03s to compute the step-level advantage, accounting for approximately 0.8% of the total time per

step, which demonstrates that HPO achieves computational efficiency comparable to that of GRPO.

6. Related Work

LLMs as decision-making agents. Large language models (LLMs) have been increasingly used as autonomous agents for reasoning, planning, and action execution across diverse domains, including code generation (Zhang et al., 2024), smart device control (Zhang & Zhang, 2024; Gur et al.; He et al., 2024; Hong et al., 2024), and interactive games (Abdulhai et al.; Prasad et al., 2024). Early approaches primarily relied on prompting-based frameworks such as ReAct (Yao et al., 2022) and Self-Refine (Madaan et al., 2023), which depend on strong proprietary models (e.g., OpenAI o3) and do not endow models with intrinsic agentic capabilities. More recent work shifts toward supervised fine-tuning (SFT) (Chen et al., 2023) and reinforcement learning (RL) (Zhai et al., 2024), allowing agents to acquire behaviors directly through data or environment interaction rather than handcrafted prompts.

Reinforcement learning for LLM agents. Reinforcement learning has emerged as a key post-training paradigm for LLMs, supporting both improved reasoning (Guo et al., 2025; Jaech et al., 2024; Pennino et al., 2025) and interactive decision-making (Xi et al., 2025; Peiyuan et al., 2024). Common algorithms include PPO (Schulman et al., 2017), GRPO (Shao et al., 2024), and REINFORCE++ (Hu et al., 2025). However, many representative works (e.g., DeepSeek-R1) focus on single-turn settings and do not apply to long-horizon interaction. Although recent studies explore multi-step RL for agent training (Cao et al., 2025; Jin et al., 2025; Qi et al.; Wang et al., 2025), they often face challenges in scalability and optimization stability when applied to complex and diverse environments (Jin et al., 2025; Xue et al., 2025; Sun et al., 2025).

7. Conclusion

This work proposes a novel policy gradient estimator that compares the current policy with the hindsight distribution in an intent space, yielding low-variance and smooth learning signals that substantially improve the stability and efficiency of long-horizon training, and offering a new pathway toward stable critic-free reinforcement learning. Future work will explore more fine-grained constructions of the intent space and extend this framework to more complex multimodal and open-ended environments.

Acknowledgments

Thanks for the kind suggestions and support from Huawei Large Model Data Technology Lab.

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.

References

- Abdulhai, M., White, I., Snell, C. V., Sun, C., Hong, J., Zhai, Y., Xu, K., and Levine, S. Lmrl gym: Benchmarks for multi-turn reinforcement learning with language models. In *Forty-second International Conference on Machine Learning*.
- Alvarez-Melis, D. and Fusi, N. Dataset dynamics via gradient flows in probability space. In *International conference on machine learning*, pp. 219–230. PMLR, 2021.
- Ambrosio, L., Gigli, N., and Savare, G. Gradient flows in metric spaces and in the wasserstein space of probability measures, 2005.
- Cao, S., Hegde, S., Li, D., Griggs, T., Liu, S., Tang, E., Pan, J., Wang, X., Malik, A., Neubig, G., et al. Skylr-v0: Train real-world long-horizon agents via reinforcement learning, 2025.
- Chen, B., Shu, C., Shareghi, E., Collier, N., Narasimhan, K., and Yao, S. Fireact: Toward language agent fine-tuning. *arXiv preprint arXiv:2310.05915*, 2023.
- Duan, Y., Chen, X., Houthooft, R., Schulman, J., and Abbeel, P. Benchmarking deep reinforcement learning for continuous control. In *International conference on machine learning*, pp. 1329–1338. PMLR, 2016.
- Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., Xu, R., Zhang, R., Ma, S., Bi, X., et al. Deepseek-r1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645(8081):633–638, 2025.
- Gur, I., Furuta, H., Huang, A. V., Safdari, M., Matsuo, Y., Eck, D., and Faust, A. A real-world webagent with planning, long context understanding, and program synthesis. In *The Twelfth International Conference on Learning Representations*.
- He, H., Yao, W., Ma, K., Yu, W., Dai, Y., Zhang, H., Lan, Z., and Yu, D. Webvoyager: Building an end-to-end web agent with large multimodal models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 6864–6890, 2024.
- Ho, X., Nguyen, A.-K. D., Sugawara, S., and Aizawa, A. Constructing a multi-hop qa dataset for comprehensive evaluation of reasoning steps. *arXiv preprint arXiv:2011.01060*, 2020.

- Hong, W., Wang, W., Lv, Q., Xu, J., Yu, W., Ji, J., Wang, Y., Wang, Z., Dong, Y., Ding, M., et al. Cogagent: A visual language model for gui agents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 14281–14290, 2024.
- Hu, J., Liu, J. K., Xu, H., and Shen, W. Reinforce++: Stabilizing critic-free policy optimization with global advantage normalization. *arXiv preprint arXiv:2501.03262*, 2025.
- Hu, S., Huang, T., Liu, G., Kompella, R. R., Ilhan, F., Tekin, S. F., Xu, Y., Yahn, Z., and Liu, L. A survey on large language model-based game agents. *arXiv preprint arXiv:2404.02039*, 2024.
- Jaech, A., Kalai, A., Lerer, A., Richardson, A., El-Kishky, A., Low, A., Helyar, A., Madry, A., Beutel, A., Carney, A., et al. Openai o1 system card. *arXiv preprint arXiv:2412.16720*, 2024.
- Jin, B., Zeng, H., Yue, Z., Yoon, J., Arik, S., Wang, D., Zamani, H., and Han, J. Search-r1: Training llms to reason and leverage search engines with reinforcement learning. *arXiv preprint arXiv:2503.09516*, 2025.
- Joshi, M., Choi, E., Weld, D. S., and Zettlemoyer, L. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. *arXiv preprint arXiv:1705.03551*, 2017.
- Just, H. A., Kang, F., Wang, J. T., Zeng, Y., Ko, M., Jin, M., and Jia, R. Lava: Data valuation without pre-specified learning algorithms. *arXiv preprint arXiv:2305.00054*, 2023.
- Kwiatkowski, T., Palomaki, J., Redfield, O., Collins, M., Parikh, A., Alberti, C., Epstein, D., Polosukhin, I., Devlin, J., Lee, K., et al. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.
- Li, Z., Xu, T., Zhang, Y., Lin, Z., Yu, Y., Sun, R., and Luo, Z.-Q. Remax: A simple, effective, and efficient reinforcement learning method for aligning large language models. In *International Conference on Machine Learning*, pp. 29128–29163. PMLR, 2024.
- Luo, R., Wang, L., He, W., Chen, L., Li, J., and Xia, X. Gui-r1: A generalist r1-style vision-language action model for gui agents. *arXiv preprint arXiv:2504.10458*, 2025.
- Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhunoye, S., Yang, Y., et al. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems*, 36:46534–46594, 2023.
- Mallen, A., Asai, A., Zhong, V., Das, R., Hajishirzi, H., and Khashabi, D. When not to trust language models: Investigating effectiveness and limitations of parametric and non-parametric memories. *arXiv preprint arXiv:2212.10511*, 7, 2022.
- Peiyuan, F., He, Y., Huang, G., Lin, Y., Zhang, H., Zhang, Y., and Li, H. Agile: A novel reinforcement learning framework of llm agents. *Advances in Neural Information Processing Systems*, 37:5244–5284, 2024.
- Pennino, F., Raimondi, B., Rondelli, M., Gurioli, A., and Gabbriellini, M. From reasoning to code: Grpo optimization for underrepresented languages. *arXiv preprint arXiv:2506.11027*, 2025.
- Prasad, A., Koller, A., Hartmann, M., Clark, P., Sabharwal, A., Bansal, M., and Khot, T. Adapt: As-needed decomposition and planning with language models. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pp. 4226–4252, 2024.
- Press, O., Zhang, M., Min, S., Schmidt, L., Smith, N. A., and Lewis, M. Measuring and narrowing the compositionality gap in language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 5687–5711, 2023.
- Qi, Z., Liu, X., Iong, I. L., Lai, H., Sun, X., Sun, J., Yang, X., Yang, Y., Yao, S., Xu, W., et al. Webrl: Training llm web agents via self-evolving online curriculum reinforcement learning. In *The Thirteenth International Conference on Learning Representations*.
- Qwen, :, Yang, A., Yang, B., Zhang, B., Hui, B., Zheng, B., and et al. Qwen2.5 technical report, 2025. URL <https://arxiv.org/abs/2412.15115>.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y., Wu, Y., et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Sun, H., Qiao, Z., Guo, J., Fan, X., Hou, Y., Jiang, Y., Xie, P., Zhang, Y., Huang, F., and Zhou, J. Zeroset: Incentivize the search capability of llms without searching. *arXiv preprint arXiv:2505.04588*, 2025.
- Sutton, R. S., Barto, A. G., et al. *Reinforcement learning: An introduction*, volume 1. MIT press Cambridge, 1998.
- Sutton, R. S., McAllester, D., Singh, S., and Mansour, Y. Policy gradient methods for reinforcement learning with function approximation. *Advances in neural information processing systems*, 12, 1999.

- Trivedi, H., Balasubramanian, N., Khot, T., and Sabharwal, A. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *Proceedings of the 61st annual meeting of the association for computational linguistics (volume 1: long papers)*, pp. 10014–10037, 2023.
- Villani, C. et al. *Optimal transport: old and new*, volume 338. Springer, 2008.
- Wang, Z., Wang, K., Wang, Q., Zhang, P., Li, L., Yang, Z., Jin, X., Yu, K., Nguyen, M. N., Liu, L., et al. Ragen: Understanding self-evolution in llm agents via multi-turn reinforcement learning. *arXiv preprint arXiv:2504.20073*, 2025.
- Williams, R. J. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3):229–256, 1992.
- Wu, C., Rajeswaran, A., Duan, Y., Kumar, V., Bayen, A. M., Kakade, S., Mordatch, I., and Abbeel, P. Variance reduction for policy gradient with action-dependent factorized baselines. In *International Conference on Learning Representations*, 2018.
- Xi, Z., Huang, J., Liao, C., Huang, B., Guo, H., Liu, J., Zheng, R., Ye, J., Zhang, J., Chen, W., et al. Agentgym-rl: Training llm agents for long-horizon decision making through multi-turn reinforcement learning. *arXiv preprint arXiv:2509.08755*, 2025.
- Xue, Z., Zheng, L., Liu, Q., Li, Y., Zheng, X., MA, Z., and An, B. Simpletir: End-to-end reinforcement learning for multi-turn tool-integrated reasoning. In *NeurIPS 2025 Fourth Workshop on Deep Learning for Code*, 2025.
- Yang, Z., Qi, P., Zhang, S., Bengio, Y., Cohen, W., Salakhutdinov, R., and Manning, C. D. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 conference on empirical methods in natural language processing*, pp. 2369–2380, 2018.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K. R., and Cao, Y. React: Synergizing reasoning and acting in language models. In *The eleventh international conference on learning representations*, 2022.
- Zhai, S., Bai, H., Lin, Z., Pan, J., Tong, P., Zhou, Y., Suhr, A., Xie, S., LeCun, Y., Ma, Y., et al. Fine-tuning large vision-language models as decision-making agents via reinforcement learning. *Advances in neural information processing systems*, 37:110935–110971, 2024.
- Zhang, G., Geng, H., Yu, X., Yin, Z., Zhang, Z., Tan, Z., Zhou, H., Li, Z., Xue, X., Li, Y., et al. The landscape of agentic reinforcement learning for llms: A survey. *arXiv preprint arXiv:2509.02547*, 2025.
- Zhang, K., Li, J., Li, G., Shi, X., and Jin, Z. Codeagent: Enhancing code generation with tool-integrated agent systems for real-world repo-level coding challenges. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 13643–13658, 2024.
- Zhang, Z. and Zhang, A. You only look at screens: Multimodal chain-of-action agents. In Ku, L.-W., Martins, A., and Srikumar, V. (eds.), *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 3132–3149, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.186. URL <https://aclanthology.org/2024.findings-acl.186/>.

Appendix

A. Experimental Details

A.1. Task Description

SearchQA. The deep search senario features a search engine–based environment equipped with specialized tools and APIs supporting the interaction with search engines. These APIs enable agents to dynamically generate search queries during the reasoning process, retrieve relevant information from external sources, and incorporate the retrieved information into subsequent reasoning steps. This setting allows agents to engage in complex reasoning processes that involve iterative searching and information integration, thereby enhancing their capability to solve intricate problems where external knowledge is essential.

TextCraft. TextCraft is a text-based game environment mirroring Minecraft. The APIs in TextCraft include crafting, inventory management, and dynamic narrative generation. These APIs allow agents to execute predefined crafting recipes, manipulate inventory contents, navigate virtual spaces.

A.2. Experimental Settings.

Settings for SearchQA. The maximum prompt length is 1024 tokens, and the maximum response length is 512 tokens. The max turn is set to 4. The learning rate is 1×10^{-6} for the actor. The reward consists of an outcome reward and a format reward, weighted in a ratio of 8:2. We set the train batch size to 64 and use a group size of 8. Rollout and validation temperatures are set to 1.0 and 0.0, respectively. The mini-batch size is 32, and the KL-divergence loss coefficient is set to 0.001. We use E5 as the retriever. The weighting coefficient ω is set to 0.5. For the training data, we use a mixture of NQ and HotpotQA. In addition, for offline HPO, we use 8 offline correct trajectories per query. We use Qwen3-Embedding-0.6B as encoder. During training, we save a model checkpoint every 50 steps; for the main experimental results, we test these checkpoints and report the best testing performance.

Settings for TextCraft. The maximum prompt length is 1024 tokens, and the maximum response length is 512 tokens. The max turn is set to 30. The learning rate is 1×10^{-6} for the actor. We adopt a outcome-based reward, assigning a reward of 1 for success and 0 for failure. We set the train batch size to 64 and use a group size of 8. Rollout and validation temperatures are set to 1.0 and 0.0, respectively. The mini-batch size is 32, and the KL-divergence loss coefficient is set to 0.001. The weighting coefficient ω is set to 0.5. For offline HPO, we use 8 offline correct trajectories per query. We use Qwen3-Embedding-0.6B as encoder. During training, we save a model checkpoint every 50 steps; for the main experimental results, we test these checkpoints and report the best testing performance.

B. General Hindsight Distribution

We first revisit the binary-reward definition used in the main text. Let $R_t = \sum_{t'=t}^{\infty} r_{t'}$ denote the future return from step t . In the binary setting, the hindsight distribution is defined by conditioning on successful future outcomes:

$$\rho_{\pi}^h(s, a) = \sum_{t=0}^{\infty} \mathbb{P}_{\pi}((s_t, a_t) = (s, a) \mid R_t = 1). \quad (20)$$

By Bayes' rule, this can be rewritten as

$$\rho_{\pi}^h(s, a) = \frac{1}{Z} \sum_{t=0}^{\infty} \mathbb{P}_{\pi}((s_t, a_t) = (s, a)) \mathbb{P}_{\pi}(R_t = 1 \mid s_t = s, a_t = a) \quad (21)$$

$$= \frac{\rho_{\pi}(s, a) Q_{\pi}(s, a)}{Z}, \quad (22)$$

where $Q_{\pi}(s, a) = \mathbb{P}_{\pi}(R_t = 1 \mid s_t = s, a_t = a)$ in the binary-reward case and Z is the normalizing constant.

This form suggests a direct extension beyond binary rewards: instead of conditioning on the event $R_t = 1$, we construct the hindsight distribution by reweighting probability mass according to realized future returns. Under continuous or dense

rewards, the hindsight distribution should therefore take the following form:

$$\rho_{\pi}^h(s, a) = \frac{1}{Z} \sum_{t=0}^{\infty} \mathbb{E}_{\pi}[R_t \mathbf{1}((s_t, a_t) = (s, a))], \quad (23)$$

where Z normalizes the distribution. Equivalently, this can be written as $\rho_{\pi}^h(s, a) = \rho_{\pi}(s, a)Q_{\pi}(s, a)/Z$, where $Q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_t \mid s_t = s, a_t = a]$. If returns can be negative or have different scales, one can use a nonnegative normalized or shifted return as the weighting function.

In summary, the two cases are unified by the same reweighting view. **Binary rewards (0/1)**: the distribution is obtained via rejection sampling over trajectories, retaining only those with return equal to 1. **Continuous rewards**: the distribution is constructed by reweighting probability mass according to trajectory returns.

C. Embedding Model Ablation

HPO constructs intent distributions in a semantic embedding space, so a natural question is whether its step-level advantage estimates are sensitive to the particular embedding model used to represent state–action pairs. To examine this, we conduct an ablation over the choice of encoder. In our main experiments, we use **Qwen3-Embedding-0.6B** as the default encoder. We compare it with two alternatives: **Qwen3-Embedding-4B**, which has a larger model scale within the same embedding family, and **all-MiniLM-L12-v2**, which provides a smaller encoder with a different architecture.

For each encoder, we recompute the HPO step-level advantages on the same set of 256 samples. We then compare the resulting advantage rankings using Kendall’s τ , which measures whether different encoders assign consistent relative importance to the same steps. Let m1, m2, and m3 denote Qwen3-Embedding-0.6B, Qwen3-Embedding-4B, and all-MiniLM-L12-v2, respectively. The rank correlations are highly significant across encoder choices: m1–m2 yields $p < 2.2 \times 10^{-308}$, and m1–m3 yields $p = 6.8 \times 10^{-212}$.

These results show that the advantage estimates produced by HPO are strongly aligned across embedding models, including both a larger encoder from the same family and an encoder with a different architecture. This suggests that HPO mainly relies on stable semantic structure in the embedding space rather than idiosyncrasies of a particular encoder, indicating that the method is not sensitive to the specific embedding model choice.

D. Proof

D.1. Proof of Lemma 3.1

Lemma D.1 (Variance Decomposition). *Consider the REINFORCE gradient estimator $\hat{G}_b$ with the baseline b ,*

$$\hat{G}_b = g(a) (\hat{Q}(s, a) - b(s)), \quad (24)$$

where $a \sim \pi(\cdot \mid s)$ and $g(a) = \nabla_{\theta} \log \pi_{\theta}(a \mid s)$.

Let $\|\cdot\|$ denotes the Euclidean norm. Then its conditional variance admits the decomposition

$$\text{Var}(\hat{G}_b \mid s) = \mathbb{E}[\|g(a)\|^2 \mid s] \cdot \text{Var}(\hat{Q}(s, a) - b(s) \mid s) \quad (25)$$

Proof. Condition on the state s and write

$$\hat{G}_b = g(a) X, \quad X := \hat{Q}(s, a) - b(s), \quad a \sim \pi(\cdot \mid s).$$

Under the standard simplifying assumption that $g(a)$ and $\hat{Q}(s, a)$ are (conditionally) uncorrelated given s , we have

$$\text{Var}(\hat{G}_b \mid s) = \mathbb{E}[\|g(a)\|^2 X^2 \mid s] - \|\mathbb{E}[g(a)X \mid s]\|^2.$$

Using conditional independence,

$$\mathbb{E}[\|g(a)\|^2 X^2 \mid s] = \mathbb{E}[\|g(a)\|^2 \mid s] \mathbb{E}[X^2 \mid s], \quad \mathbb{E}[g(a)X \mid s] = \mathbb{E}[g(a) \mid s] \mathbb{E}[X \mid s].$$

Moreover, for score functions one typically has $\mathbb{E}[g(a) \mid s] = \mathbb{E}[\nabla_\theta \log \pi_\theta(a \mid s) \mid s] = 0$. Therefore,

$$\text{Var}(\hat{G}_b \mid s) = \mathbb{E}[\|g(a)\|^2 \mid s] \mathbb{E}[X^2 \mid s].$$

Finally, if the baseline is chosen as the conditional mean $b(s) = \mathbb{E}[\hat{Q}(s, a) \mid s]$, then $\mathbb{E}[X \mid s] = 0$ and thus $\mathbb{E}[X^2 \mid s] = \text{Var}(X \mid s) = \text{Var}(\hat{Q}(s, a) - b(s) \mid s)$, yielding

$$\text{Var}(\hat{G}_b \mid s) = \mathbb{E}[\|g(a)\|^2 \mid s] \text{Var}(\hat{Q}(s, a) - b(s) \mid s),$$

□

D.2. Proof of Theorem 3.3

Theorem D.2 (Optimal Baseline and Excess Variance). *Among all scalar baselines $b(s)$ independent of the sampled action a , the conditional variance $\text{Var}(\hat{G}_b \mid s)$ is minimized by*

$$b^*(s) = \frac{\mathbb{E}[\|g(a)\|^2 \hat{Q}(s, a) \mid s]}{\mathbb{E}[\|g(a)\|^2 \mid s]}. \quad (26)$$

Moreover, for any alternative baseline $\tilde{b}(s)$, the increase in variance admits the exact expression

$$\begin{aligned} \text{Var}(\hat{G}_{\tilde{b}} \mid s) - \text{Var}(\hat{G}_{b^*} \mid s) &= \\ &= \mathbb{E}[\|g(a)\|^2 \mid s] (\tilde{b}(s) - b^*(s))^2. \end{aligned} \quad (27)$$

Proof. Fix s and write $X = \hat{Q}(s, a)$. For any scalar baseline $b = b(s)$ independent of a ,

$$\text{Var}(\hat{G}_b \mid s) = \mathbb{E}[\|g(a)\|^2 (X - b)^2 \mid s] - \|\mathbb{E}[g(a)X \mid s]\|^2,$$

where we use $\mathbb{E}[g(a) \mid s] = 0$, so subtracting an action-independent baseline does not change the conditional mean. Since the second term does not depend on b , the variance is minimized by

$$b^*(s) = \arg \min_b \mathbb{E}[\|g(a)\|^2 (X - b)^2 \mid s] = \frac{\mathbb{E}[\|g(a)\|^2 X \mid s]}{\mathbb{E}[\|g(a)\|^2 \mid s]}.$$

For any alternative $\tilde{b}(s)$,

$$\begin{aligned} &\mathbb{E}[\|g(a)\|^2 (X - \tilde{b})^2 \mid s] \\ &= \mathbb{E}[\|g(a)\|^2 (X - b^* + b^* - \tilde{b})^2 \mid s] \\ &= \mathbb{E}[\|g(a)\|^2 (X - b^*)^2 \mid s] - 2(\tilde{b} - b^*) \mathbb{E}[\|g(a)\|^2 (X - b^*) \mid s] + \mathbb{E}[\|g(a)\|^2 \mid s] (\tilde{b} - b^*)^2. \end{aligned}$$

Since $b^* = \mathbb{E}[\|g(a)\|^2 X \mid s] / \mathbb{E}[\|g(a)\|^2 \mid s]$, the cross term vanishes, $\mathbb{E}[\|g(a)\|^2 (X - b^*) \mid s] = 0$, and therefore

$$\mathbb{E}[\|g(a)\|^2 (X - \tilde{b})^2 \mid s] - \mathbb{E}[\|g(a)\|^2 (X - b^*)^2 \mid s] = \mathbb{E}[\|g(a)\|^2 \mid s] (\tilde{b} - b^*)^2.$$

which immediately gives

$$\text{Var}(\hat{G}_{\tilde{b}} \mid s) - \text{Var}(\hat{G}_{b^*} \mid s) = \mathbb{E}[\|g(a)\|^2 \mid s] (\tilde{b}(s) - b^*(s))^2.$$

□

D.3. Proof of Lemma 4.2

Lemma D.3. Consider the policy objective $\mathcal{J}(\theta) = \mathbb{E}_{\tau \sim \pi}[\sum_{t \geq 0} r_t]$. Its policy gradient admits the following two equivalent forms:

$$\nabla_{\theta} \mathcal{J}(\theta) = \mathbb{E}_{(s,a) \sim \rho_{\pi}}[Q_{\pi}(s,a) \nabla_{\theta} \log \pi_{\theta}(a | s)] \quad (28)$$

$$= \mathbb{E}_{(s,a) \sim \rho_{\pi}}[-\frac{\delta KL(\rho_{\pi}^h || \rho_{\pi})}{\delta \rho_{\pi}(s,a)} \nabla_{\theta} \log \pi_{\theta}(a | s)] \quad (29)$$

where $-\frac{\delta KL(\rho_{\pi}^h || \rho_{\pi})}{\delta \rho_{\pi}(s,a)} = \frac{\rho_{\pi}^h(s,a)}{\rho_{\pi}(s,a)}$ denotes the pointwise variational derivative of $KL(\rho_{\pi}^h || \rho_{\pi})$ with respect to $\rho_{\pi}(s,a)$.

Proof. For $\mathcal{J}(\theta) = \mathbb{E}_{\tau \sim \pi}[\sum_{t \geq 0} r_t]$, by the policy gradient theorem, its gradient admits the following equivalent form

$$\nabla_{\theta} \mathcal{J}(\theta) = \mathbb{E}_{(s,a) \sim \rho_{\pi}}[Q_{\pi}(s,a) \nabla_{\theta} \log \pi_{\theta}(a | s)] \quad (30)$$

Writing the expectation explicitly, we obtain

$$\nabla_{\theta} \mathcal{J}(\theta) = \sum_{s,a} \rho_{\pi}(s,a) Q_{\pi}(s,a) \nabla_{\theta} \log \pi_{\theta}(a | s) \quad (31)$$

Let $Z = \sum_{s,a} \rho_{\pi}(s,a) Q_{\pi}(s,a)$. Define $\rho_{\pi}^h(s,a) = \rho_{\pi}(s,a) Q_{\pi}(s,a) / Z$. Then the above becomes

$$\nabla_{\theta} \mathcal{J}(\theta) = \sum_{s,a} \rho_{\pi}(s,a) Q_{\pi}(s,a) \nabla_{\theta} \log \pi_{\theta}(a | s) \quad (32)$$

$$= Z \sum_{s,a} \rho_{\pi}(s,a) \frac{\rho_{\pi}^h(s,a)}{\rho_{\pi}(s,a)} \nabla_{\theta} \log \pi_{\theta}(a | s) \quad (33)$$

$$= Z \mathbb{E}_{(s,a) \sim \rho_{\pi}}[\frac{\rho_{\pi}^h(s,a)}{\rho_{\pi}(s,a)} \nabla_{\theta} \log \pi_{\theta}(a | s)] \quad (34)$$

where the ratio $\frac{\rho_{\pi}^h(s,a)}{\rho_{\pi}(s,a)}$ can be interpreted as the pointwise variational derivative of $KL(\rho_{\pi}^h || \rho_{\pi})$ with respect to ρ_{π} , i.e., $-\frac{\delta KL(\rho_{\pi}^h || \rho_{\pi})}{\delta \rho_{\pi}(s,a)}$.

Since Z is a positive constant, it does not affect the direction of the gradient. Therefore, the policy gradient admits the following two equivalent forms:

$$\nabla_{\theta} \mathcal{J}(\theta) = \mathbb{E}_{(s,a) \sim \rho_{\pi}}[Q_{\pi}(s,a) \nabla_{\theta} \log \pi_{\theta}(a | s)] \quad (35)$$

$$= \mathbb{E}_{(s,a) \sim \rho_{\pi}}[-\frac{\delta KL(\rho_{\pi}^h || \rho_{\pi})}{\delta \rho_{\pi}(s,a)} \nabla_{\theta} \log \pi_{\theta}(a | s)] \quad (36)$$

□

D.4. Proof of Lemma 4.3

Lemma D.4. Define $D = \sup_{(s,a),(s',a') \in \rho_{\pi}} d((s,a), (s',a'))$. Let $f^*(s,a)$ denote the $-\frac{\delta W_1(\rho_{\pi} || \rho_{\pi}^h)}{\delta \rho_{\pi}(s,a)}$, and $g(s,a)$ denote the $-\frac{\delta KL(\rho_{\pi}^h || \rho_{\pi})}{\delta \rho_{\pi}(s,a)}$, then

$$\text{Var}_{\rho_{\pi}}(f^*(s,a)) \leq D^2/4. \quad (37)$$

and

$$\text{Var}_{\rho_{\pi}}(g(s,a)) = \int \frac{\rho_{\pi}^h(s,a)^2}{\rho_{\pi}(s,a)} d(s,a) - 1 = \chi^2(\rho_{\pi}^h || \rho_{\pi}). \quad (38)$$

Proof. For the W_1 term, recall the Kantorovich–Rubinstein dual:

$$W_1(\rho_{\pi}^h, \rho_{\pi}) = \sup_{\|f\|_{\text{Lip}} \leq 1} \mathbb{E}_{x \sim \rho_{\pi}^h}[f(x)] - \mathbb{E}_{x \sim \rho_{\pi}}[f(x)], \quad x = (s,a). \quad (39)$$

Let f^* be the Kantorovich potential, which is 1-Lipschitz on the support of ρ_π (up to an additive constant). By definition of the diameter $D = \sup_{x, x' \in \text{supp}(\rho_\pi)} d(x, x')$, any 1-Lipschitz function satisfies, for all $x, x' \in \text{supp}(\rho_\pi)$, $|f^*(x) - f^*(x')| \leq d(x, x') \leq D$,

Therefore, by Popoviciu’s inequality on variances,

$$\text{Var}_{\rho_\pi}(f^*(x)) \leq \frac{(\sup f^* - \inf f^*)^2}{4} \leq \frac{D^2}{4}. \quad (40)$$

For the KL term, by the stated definition,

$$g(x) = -\frac{\delta KL(\rho_\pi^h \| \rho_\pi)}{\delta \rho_\pi(x)} = \frac{\rho_\pi^h(x)}{\rho_\pi(x)}, \quad (41)$$

assuming $\rho_\pi(x) > 0$ whenever $\rho_\pi^h(x) > 0$. Then

$$\mathbb{E}_{x \sim \rho_\pi}[g(x)] = \int \rho_\pi(x) \frac{\rho_\pi^h(x)}{\rho_\pi(x)} dx = \int \rho_\pi^h(x) dx = 1, \quad (42)$$

and

$$\text{Var}_{\rho_\pi}(g(x)) = \mathbb{E}_{\rho_\pi}[g(x)^2] - (\mathbb{E}_{\rho_\pi}[g(x)])^2 = \int \rho_\pi(x) \left(\frac{\rho_\pi^h(x)}{\rho_\pi(x)}\right)^2 dx - 1 = \int \frac{\rho_\pi^h(x)^2}{\rho_\pi(x)} dx - 1. \quad (43)$$

Finally, noting that

$$\chi^2(\rho_\pi^h \| \rho_\pi) = \int \frac{(\rho_\pi^h(x) - \rho_\pi(x))^2}{\rho_\pi(x)} dx = \int \frac{\rho_\pi^h(x)^2}{\rho_\pi(x)} dx - 2 \int \rho_\pi^h(x) dx + \int \rho_\pi(x) dx = \int \frac{\rho_\pi^h(x)^2}{\rho_\pi(x)} dx - 1, \quad (44)$$

□

E. Additional Analysis

E.1. Analysis of the weight ω

To further investigate the effect of the step-level advantage weight ω , multiple repeated HPO-off experiments were conducted with different ω settings. Figure 7 illustrates the training reward trajectories for varying values of ω . The results show that introducing a larger ω can lead to more stable optimization, while this effect diminishes as ω decreases. In addition, larger values of ω yield faster reward improvement in the early training stages but converge to slightly lower final performance. A plausible explanation is that a larger ω enables more effective learning from sparse rewards during early training, whereas the increased optimization bias introduced by an overly large ω may mildly degrade asymptotic performance. Overall, these results indicate that adjusting ω provides a practical mechanism for balancing bias and variance during training.

E.2. Case Study

We further conduct a qualitative analysis of A_S across representative samples. We find that A_S exhibits strong step-level discrimination: it consistently assigns higher scores to informative and goal-progressing actions, while down-weighting failed, redundant, or uninformative steps. This behavior suggests that A_S captures meaningful local learning signals and also explains why optimizing with A_S alone can still drive effective policy updates.

Example 1

User: In what year was the earlier of these two documentaries produced: The Life and Times of Hank Greenberg or Burden of Dreams?

Assistant: < think > The question is asking about the production year of two documentaries, and we need to find out which one is earlier and then provide the year of that documentary.< /think > < search > which documentary was produced earlier, The Life and Times of Hank Greenberg or Burden of Dreams < /search >




Figure 7. Training reward dynamics of Qwen2.5-3B-Instruct with different ω settings.

($A_s = +0.424$)

User: < information >

DOC 1 ... The Life and Times of Hank Greenberg is a 1998 documentary film written, directed, and produced by Aviva Kempner about Hank Greenberg, first baseman of the Detroit Tigers, who was inducted into the Baseball Hall of Fame. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: < think > From the information provided, it is clear that The Life and Times of Hank Greenberg is a 1998 documentary. Now let's find out about the production year of Burden of Dreams.< /think > < search > when was Burden of Dreams produced < /search >

($A_s = +0.218$)

User: < information >

DOC 1 ... Burden of Dreams is a 1982 "making-of" documentary film directed by Les Blank, shot during and about the chaotic production of Werner Herzog's film Fitzcarraldo. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.

($A_s = -0.472$)

User: Your previous action is invalid. If you want to search, you should put the query between < search > and < /search >. If you want to give the final answer, you should put the answer between < answer > and < /answer >. Please try again.

Assistant: < think > From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.< /think > < answer > 1982 < /answer >

($A_s = +0.377$)

User: Congratulations! You have answered the question correctly!!!

Example 2

User: Bone carving was used in the casket now on display in what institution?

Assistant: < think > I need to find out which institution has the casket with bone carving that is now on display.< /think > < search > Bone carving was used in the casket now on display in what institution? < /search >
($A_s = +1.20$)

User: < information >

DOC 1 ... ivory casket ... now in the Walters Art Museum ...

DOC 2 ... a small Anglo-Saxon whale's bone ... chest ... densely decorated with knife-cut ... scenes ... now in the **British Museum**. ...

DOC 3 ...

< /information >

Assistant: < think > From the search results, it is clear that the casket with bone carving now on display is in the Walters Art Museum, Baltimore, Maryland. < /think > < answer > Walters Art Museum < /answer >
($A_s = -0.574$)

User: Sorry, your answer is incorrect.

Example 1: the advantage of GRPO

User: In what year was the earlier of these two documentaries produced: The Life and Times of Hank Greenberg or Burden of Dreams?

Assistant: < think > The question is asking about the production year of two documentaries, and we need to find out which one is earlier and then provide the year of that documentary.< /think > < search > which documentary was produced earlier, The Life and Times of Hank Greenberg or Burden of Dreams < /search >
($A_s = +1.29$)

User: < information >

DOC 1 ... The Life and Times of Hank Greenberg is a 1998 documentary film written, directed, and produced by Aviva Kempner about Hank Greenberg, first baseman of the Detroit Tigers, who was inducted into the Baseball Hall of Fame. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: < think > From the information provided, it is clear that The Life and Times of Hank Greenberg is a 1998 documentary. Now let's find out about the production year of Burden of Dreams.< /think > < search > when was Burden of Dreams produced < /search >

($A_s = +1.29$)

User: < information >

DOC 1 ... Burden of Dreams is a 1982 "making-of" documentary film directed by Les Blank, shot during and about the chaotic production of Werner Herzog's film Fitzcarraldo. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.

($A_s = +1.29$)

User: Your previous action is invalid. If you want to search, you should put the query between < search > and < /search >. If you want to give the final answer, you should put the answer between < answer > and < /answer >. Please try again.

Assistant: < think > From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.< /think > < answer > 1982 < /answer >

($A_s = +1.29$)

User: Congratulations! You have answered the question correctly!!!

Example 2: the advantage of PPO

User: In what year was the earlier of these two documentaries produced: The Life and Times of Hank Greenberg or Burden of Dreams?

Assistant: < think > The question is asking about the production year of two documentaries, and we need to find out which one is earlier and then provide the year of that documentary.< /think > < search > which documentary was produced earlier, The Life and Times of Hank Greenberg or Burden of Dreams < /search >
($A_s = +0.821$)

User: < information >

DOC 1 ... The Life and Times of Hank Greenberg is a 1998 documentary film written, directed, and produced by Aviva Kempner about Hank Greenberg, first baseman of the Detroit Tigers, who was inducted into the Baseball Hall of Fame. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: < think > From the information provided, it is clear that The Life and Times of Hank Greenberg is a 1998 documentary. Now let's find out about the production year of Burden of Dreams.< /think > < search > when was Burden of Dreams produced < /search >

($A_s = +0.762$)

User: < information >

DOC 1 ... Burden of Dreams is a 1982 "making-of" documentary film directed by Les Blank, shot during and about the chaotic production of Werner Herzog's film Fitzcarraldo. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.

($A_s = +0.684$)

User: Your previous action is invalid. If you want to search, you should put the query between < search > and < /search >. If you want to give the final answer, you should put the answer between < answer > and < /answer >. Please try again.

Assistant: < think > From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.< /think > < answer > 1982 < /answer >

($A_s = +0.707$)

User: Congratulations! You have answered the question correctly!!!

Example 3: the advantage of HPO

User: In what year was the earlier of these two documentaries produced: The Life and Times of Hank Greenberg or Burden of Dreams?

Assistant: < think > The question is asking about the production year of two documentaries, and we need to find out which one is earlier and then provide the year of that documentary.< /think > < search > which documentary was produced earlier, The Life and Times of Hank Greenberg or Burden of Dreams < /search >
($A_s = +0.424$)

User: < information >

DOC 1 ... The Life and Times of Hank Greenberg is a 1998 documentary film written, directed, and produced by Aviva Kempner about Hank Greenberg, first baseman of the Detroit Tigers, who was inducted into the Baseball Hall of Fame. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: < think > From the information provided, it is clear that The Life and Times of Hank Greenberg is a 1998 documentary. Now let's find out about the production year of Burden of Dreams.< /think > < search > when was Burden of Dreams produced < /search >

($A_s = +0.218$)

User: < information >

DOC 1 ... Burden of Dreams is a 1982 "making-of" documentary film directed by Les Blank, shot during and about the chaotic production of Werner Herzog's film Fitzcarraldo. ...

DOC 2 ...

DOC 3 ...

< /information >

Assistant: From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.

($A_s = -0.472$)

User: Your previous action is invalid. If you want to search, you should put the query between < search > and < /search >. If you want to give the final answer, you should put the answer between < answer > and < /answer >. Please try again.

Assistant: < think > From the information provided, it is clear that Burden of Dreams was produced in 1982. Comparing the production years of the two documentaries, The Life and Times of Hank Greenberg (1998) and Burden of Dreams (1982), it is clear that Burden of Dreams was produced earlier. Therefore, the answer is 1982.< /think > < answer > 1982 < /answer >

($A_s = +0.377$)

User: Congratulations! You have answered the question correctly!!!