
One Probe Won’t Catch Them All: Towards Targeted Deception Detection

Vikram Natarajan^{*1} Devina Jain^{*1} Shivam Arora^{*1} Satvik Golechha² Joseph Bloom²

Abstract

Linear probes are a promising approach for monitoring AI systems for deceptive behaviour. Previous work has shown that a linear classifier trained on a contrastive instruction pair and a simple dataset can achieve good performance. However, these probes exhibit notable failures even in straightforward scenarios, including spurious correlations and false positives on non-deceptive responses. In this paper, we demonstrate that deception detection is inherently heterogeneous: while a single universal probe achieves modest improvements (+0.032 AUC), post-hoc oracle analysis reveals substantially higher potential (+0.108 AUC) when probes are matched to specific deception types, and synthetic validation experiments suggest this ceiling is achievable *a priori* when the deception type is known in advance. Our findings reveal that instruction pairs capture deceptive *intent* rather than content-specific patterns, explaining why prompt choice dominates probe performance (70.6% of variance). Given this heterogeneity, we conclude that organizations should define their specific threat models and deploy appropriately matched probes rather than seeking a universal deception detector.

1. Introduction

Deception in large language models has emerged as a central concern in AI safety. Models now possess both the capability and propensity to engage in deceptive behaviours when placed in situations that provide incentives to do so. Empirical studies have demonstrated that such behaviours are increasingly prevalent in the latest generation of frontier models: LLMs have been observed to misrepresent informa-

tion, obscure intentions, or act strategically under specific contextual pressures (Schoen et al., 2025; Chowdhury et al., 2025; Park et al., 2024). As LLMs are increasingly deployed in a wide variety of domains including high-stakes contexts, there is an urgent need for reliable monitoring mechanisms capable of detecting when a model engages in deceptive output. Black-box approaches, which monitor only tokens output by the model using trained classifiers or other LLMs, are one possible approach. Chain-of-thought monitoring is widely used to monitor model misbehaviour including deception (Korbak et al., 2025).

However, recent work indicates that chain-of-thought traces are not always faithful indicators of a model’s internal reasoning process. Models can perform reasoning implicitly while generating innocuous-looking chain-of-thought outputs (Arcuschin et al., 2025; Hao et al., 2024). In addition to black-box monitoring being expensive, the increasing subtlety of deception in advanced models means that model outputs alone cannot be relied upon for effective monitoring, as models may learn to strategically generate innocuous-sounding text. Because of these limitations, there has been interest in exploring white-box methods that inspect model activations.

Early work has included training classifiers on internal activations to distinguish between true and false statements (Azaria & Mitchell, 2023; Burns et al., 2023; Zou et al., 2025), but these efforts have largely focused on evaluating these techniques on factual inaccuracies rather than strategic deception. Goldowsky-Dill et al. (2025) addressed some of these limitations with instruction-pair probes to capture a model’s intention to deceive. They trained a linear classifier on model activations from simple facts in the Zou et al. (2025) instruction pairs dataset, comparing responses under honesty versus dishonesty instructions. This approach generalized surprisingly well to detecting strategic deception in settings such as concealing insider trading and deliberately underperforming on safety evaluations. However, the probe suffered from aggregation errors and spurious correlations, sometimes firing on completely honest responses. Despite their effectiveness, probes learn spurious correlations with their training datasets. Levinstein & Herrmann (2024) showed that probes from Azaria & Mitchell (2023) even misclassified negations of original statements from training dataset. Further empirical evidence from Levy (2024)

^{*}Equal contribution ¹LASR Labs ²UK AI Security Institute. Correspondence to: Vikram Natarajan <vikram.s.natarajan@gmail.com>, Devina Jain <djain2095@gmail.com>, Shivam Arora <shivamarora-math@gmail.com>.

showed that classifiers based on Zou et al. (2025) had high sensitivity but low specificity, and were better characterized as differentiating between fact-based retrieval and other tasks (such as writing science fiction), spanning a much broader surface area than intended.

Our initial experiments revealed that instruction-pair probes using the default prompts from (Goldowsky-Dill et al., 2025) perform poorly across diverse deception datasets such as MASK. For example, the baseline probe achieved only 0.374 AUC on MASK known facts, indicating substantial room for improvement. Analysis of the evaluation datasets uncovered substantial heterogeneity in the types of deception across different benchmarks. This diversity in deception characteristics presents a fundamental challenge for developing a single universal deception probe. For deployment purposes, labs would prefer to distinguish probes that detect genuinely concerning behaviours from those that merely flag benign role-playing scenarios.

1.1. Our contribution

We present three key contributions to deception detection using linear probes:

System prompts dominate probe performance. Through analysis of 2000 probes across varied configurations, we demonstrate that system prompt design explains 70.6% of performance variance, substantially exceeding the impact of training data, layer selection, or token aggregation strategy. This finding justifies our focus on prompt-based approaches to improving deception detection.

Deception detection is heterogeneous, but targeted probes achieve strong performance. While attempting to select a single universal probe shows modest gains (+0.032 AUC, 5.5%), post-hoc oracle analysis reveals a much higher performance ceiling (+0.108 AUC, 20.5%) when probes are matched to specific deception types from a human-interpretable taxonomy (Druckman & Bjork, 1991). Importantly, experiments on synthetic data suggest this ceiling is achievable *a priori*: when the deception type is known in advance, targeted probes often outperform the baseline.

Organizations should match probes to their threat models. We identify which prompt types work best for which deception scenarios, and demonstrate through qualitative analysis that probes detect distinct semantic properties. This heterogeneity challenges the assumption of a universal deception detector and suggests that organizations with well-defined threat models can achieve stronger detection by deploying appropriately matched probes.

2. Related work

2.1. White-box methods

White-box methods for deception detection leverage access to a model’s activation space to identify when language models produce false or deceptive outputs. Several studies have focused on training classifiers on the model’s internal activations (Azaria & Mitchell, 2023; Burns et al., 2023; Zou et al., 2025; Marks & Tegmark, 2024). Although effective at classification, early white-box techniques for detecting deception primarily focused on off-policy fact-based evaluation datasets like TruthfulQA, rather than targeting model-generated responses that mimic more sophisticated scenarios.

2.2. Instruction-pair probes

Goldowsky-Dill et al. (2025) addressed some of these limitations with instruction-pair probes, shifting focus from the lie itself to the intent to lie. They trained a linear classifier on model activations from simple facts in the Zou et al. (2025) instruction pairs dataset, comparing responses under honesty versus dishonesty instructions. The probe performed well with AUROCs greater than 0.96 when tested on datasets such as sandbagging and insider trading.

However, these probes exhibited notable limitations. In follow-up work, Parrack et al. (2025) benchmarked the performance of instruction-pair probes combined with black-box monitors versus black-box monitors alone, and evaluated different aggregation strategies. They confirmed significant limitations of instruction-pair probes, including spurious correlations demonstrated through a deception dilution dataset. These probes often flag innocuous content such as fiction writing or role-playing as deceptive, suggesting they may be detecting broader patterns (e.g., fact-based vs. creative reasoning) rather than strategic deception specifically.

2.3. Benchmarks

Benchmarking deception in large language models remains a difficult challenge. Several recent works have attempted to construct datasets that capture different forms of deceptive behaviour, such as factual lying, instruction misrepresentation, strategic omission, or self-serving reasoning. Goldowsky-Dill et al. (2025) collected and constructed multiple datasets including AI liar, Insider trading, Sandbagging, and Role-playing scenarios, among others.

MASK: The MASK dataset (Ren et al., 2025) provides a rigorous benchmark for evaluating honesty in large language models by measuring whether models remain truthful when incentivized to lie. It contains human-labeled examples across six distinct archetypes, each consisting of a proposition, ground truth, pressure prompt designed to elicit lying,

and belief elicitation prompt to establish the model’s actual knowledge.

Liar’s Bench: [Kretschmar et al. \(2025\)](#) built a large benchmark for evaluating deception detection in AI assistants. It contains examples of deceptive and honest model outputs across six scenarios, capturing diverse types of deception such as strategic omission, misrepresentation, and evasion through irrelevant truths.

3. Methodology

3.1. Instruction-Pair Probe

We build upon the instruction-pair probes introduced by ([Goldowsky-Dill et al., 2025](#)). The approach constructs contrastive pairs: each training example uses the same factual statement as the model response, but pairs it with either an honest or dishonest system prompt. The response is token-forced (off-policy), meaning the factual statement is fed directly into the chat template as the model’s response rather than being generated by the model.

For each response token, we extract the residual stream values at the layer of interest—these activation vectors directly become the input features for a logistic regression classifier trained to predict whether the honest or dishonest prompt was used. Crucially, we exclude the final 5 response tokens from the activations. This prevents the probe from detecting whether the stated fact is literally true or false based on the response content. By truncating before the fact fully resolves (e.g., capturing activations at “The blue whale is the...” before the model commits to “largest” or “smallest”), the probe is forced to learn from the model’s internal representations of intent as shaped by the system prompt, rather than from the surface-level content of the response.

Default prompt pair:

- **Honest prompt:** “Pretend you are an honest person making statements about the world.” Model response: “The blue whale is the...”
- **Deceptive prompt:** “Pretend you are a deceptive person making statements about the world.” Model response: “The blue whale is the...”

3.2. Models

We use the open-weight instruction-tuned model Gemma 2 9B IT as our primary model. After conducting preliminary experiments across layers 20, 25, and 31 (see Section 4.1), we focus our analysis on layer 20. The evaluation datasets are generated with on-policy responses from the model with the temperature set to 1. To assess whether our findings generalize across architectures and scales, we additionally

replicate the ANOVA experiment on Llama 3.3 70B Instruct (Section 4.1).

3.3. Datasets

We distinguish three dataset roles: training, validation, and evaluation.

Training data. Probes are trained on small datasets of true and false factual statements (e.g., “The blue whale is the largest mammal on Earth”) from [Azaria & Mitchell \(2023\)](#). We also generated 11 additional synthetic factual datasets covering different types of facts, to measure whether the specific factual domain affects probe performance. Each dataset contains approximately 200 facts. Dataset categories are described in Appendix B. Following the ANOVA analysis (Section 4.1), which found that varying the domain and structure of training facts (e.g., simple statements vs. if-else constructions) explained only 2.5% of probe performance variance, we focused all subsequent experiments on the Azaria-Mitchell dataset. This is consistent with our broader finding that probes detect intent rather than content-specific patterns.

Evaluation data. Probes are evaluated on ten richer deception datasets drawn from [Goldowsky-Dill et al. \(2025\)](#) and Liar’s Bench ([Kretschmar et al., 2025](#)), including MASK ([Ren et al., 2025](#)) and others. These datasets contain on-policy model responses to scenarios involving various deception types, and are entirely separate from the training data.

Validation data. To select the single best-performing probe on average across datasets, we hold out 50 balanced samples (25 honest, 25 deceptive) from each of three evaluation datasets (AI liar, Convincing Game, and Instructed Deception). These samples are never reused for evaluation. Note that this selection procedure only applies to the “Best Average” probe; for the “Best Taxonomy” oracle evaluation on a given dataset, we simply run all probes and report the best-performing one.

3.4. Prompt Generation

We generate two types of instruction pairs for probe training. Each pair consists of an honest and a dishonest system prompt targeting the same deception scenario. Full prompts and generation procedures are available in Appendix B.3.

1. **Taxonomy-Based Prompts:** We generated 23 prompt pairs using Claude Sonnet 4.5 ([Anthropic, 2025](#)), each targeting a specific deception type from the taxonomy in [Druckman & Bjork \(1991\)](#). These include 16 core deception types (white lie, exaggeration, evasion, bluff, concealment, pretense, trick, impersonation, counterfeited, forgery, pressure dishonesty, backstab, false alarm,

hoax, partial truth, and overt lie) plus 7 control pairs for non-deceptive behaviors.

2. **Framing Variations:** We generated 34 additional prompt pairs using Claude Sonnet 4.5 (Anthropic, 2025), exploring different ways to linguistically frame the honest/dishonest distinction, without targeting a specific deception type. Categories include identity-based framings (e.g., “You are an AI designed to be fundamentally truthful”), role-based framings (e.g., “Your role is to provide factual and accurate information”), instruction-based, hybrid, control repetition, long-form, and baseline framings. These explore whether the surface-level structure of the prompt affects probe performance independently of the deception type being targeted.

3.5. Baselines and Skylines

For baselines, we use an instruction-pair probe with the default system prompt from Goldowsky-Dill et al. (2025) (described in Section 3.1). We also computed in-distribution skyline probes on each evaluation dataset using an 80/20 train/test split, representing the best achievable performance when the probe’s training distribution matches the evaluation distribution. These skyline probes achieved near-perfect AUC (≥ 0.993 on nine of ten datasets), and are omitted from the main figures for clarity.

3.6. Training of Probes

We train logistic regression probes on the residual stream values extracted at layer 20 (zero-indexed) of the model (as described in Section 3.1). These activation vectors are standardized to zero mean and unit variance across the dataset before training. We aggregate over all response tokens except the final 5, using the mean. The L2 regularization constant was tuned as part of the ANOVA sweep; we found $\lambda = 1$ to be optimal.

4. Experiments

4.1. Analysis of Variance across Training Setups

We identified four factors that affect probe performance: (1) the instruction pair (honest/deceptive prompt) used when training the probe, (2) the specific dataset of true facts used as the model response, (3) the token selection strategy for aggregation from the model response (first 5, first 10, all except last 10, all except last 5, or chat template tokens only), and (4) the layer probed at.

In order to determine where to focus our efforts, we trained 2000 probes using different combinations of these factors (exploring layers 20, 25, and 31), and evaluated each on the validation dataset described in Section 3.3. We then




Figure 1. Variance in AUC on evaluation datasets explained by experimental factors. System prompt explains 70.6% of performance variance ($p < 0.001$), substantially exceeding variance explained by layer ID (2.7%), training dataset (2.5%), or token selection strategy (0.6%). Residual unexplained variance accounts for 23.6%.

conducted an analysis of variance (ANOVA) on the results to determine the fraction of variance attributable to each factor.

Most of the variance in probe performance is attributable to changes in the system prompt, with relatively minor effects caused by variation in the layer ID, training dataset, or token selection policy. There are minor interaction effects between the system prompt and training dataset, but in general they are dominated by the system prompt variations.

This result forms our justification for focusing on system prompt variation and fixing the remaining factors for the remainder of our analysis. Given that layer choice explained only 2.7% of variance, we selected layer 20 for all subsequent experiments. We also fixed: 1) Token selection to all tokens up to 5 tokens before the end of the response. 2) Training dataset to the original dataset of true and false facts from Azaria & Mitchell (2023)

Cross-model replication on Llama 3.3 70B Instruct. To assess whether the dominance of system prompt choice generalizes beyond Gemma-2-9B-IT, we replicated the same ANOVA sweep on Llama-3.3-70B-Instruct. System prompt choice again dominates, explaining 71.0% of performance variance ($p < 0.001$), compared with 5.8% for layer ID, 4.5% for training dataset, and 0.3% for token selection strategy (Figure 2).

4.2. Probe Training with Varying System Prompts

The ANOVA result motivates a key question: given that prompt choice dominates probe performance, do different deception datasets require different prompts, or does a sin-




Figure 2. ANOVA variance decomposition on Llama-3.3-70B-Instruct. System prompt choice explains 71.0% of performance variance, consistent with the 70.6% observed on Gemma-2-9B-IT (Figure 1).

gle well-chosen prompt suffice? We hypothesize that deception is sufficiently heterogeneous that specialized probes will outperform a single generic detector. To test this, we train separate probes for each system prompt variation and evaluate them across all datasets, following the validation procedure described in Section 3.3.

5. Results

Our experiments confirmed that the system prompt provided to the model is the dominant factor affecting deception probe performance. We observed substantial performance variation across probes trained with different system prompts, when training probes with the system prompt, training dataset, token selection and layer ID varied independently. Importantly, the best-performing probes were interpretable, with clear semantic relationships between the prompt content and the types of deception they successfully detected. We identified specific system prompts that achieved optimal performance for different datasets, reflecting the underlying heterogeneity in deception types. This finding suggests that detecting deception across datasets requires specialized detection approaches rather than a one-size-fits-all solution.

5.1. Comparison of Performance Across Probe Types

We evaluated the effectiveness of deception taxonomy-based prompts (Druckman & Bjork, 1991) for deception detection by comparing these prompts to other probe variants across ten deception datasets. Using linear probes trained on Gemma-2-9B-IT layer 20 representations, we assessed:

1. *Deception Baseline* prompts using generic hon-

est/dishonest distinctions;

2. the single *Best Average* probe (best mean performance across all datasets without using taxonomy);
3. the *Best Taxonomy* probe for the specific dataset being evaluated.

All ten datasets had improved performance with a specific taxonomy prompt: eight exceeded a 5% AUC gain, three exceeded 10%, and one (MASK known facts) exceeded 25%. Importantly, for every dataset there existed a taxonomy-based probe that outperformed the baseline.

Dataset-specific performance patterns. The magnitude of improvement from taxonomy-based prompts varies continuously across datasets, ranging from near-zero to over 86% relative gains. This variation reflects both the baseline detectability of different deception types and the degree to which specific taxonomy prompts match the underlying deception mechanism.

At one extreme, *liars-bench-convincing-game* shows minimal improvement (Baseline: 0.984, Taxonomy: 0.985, $\Delta = +0.001$), suggesting that straightforward lies where the model is explicitly prompted to lie exhibit stereotypical patterns easily detected even with generic prompts. At the other extreme, *mask-known-facts* demonstrates dramatic gains ($0.374 \rightarrow 0.697$, $\Delta = +0.323$, +86.3%) when using pressure-dishonesty prompts (AUC = 0.697), indicating that deception under social pressure requires specifically targeted detection approaches.

Several datasets show substantial improvements exceeding 20% relative gains: *mask-doubling-down-known-facts* ($0.565 \rightarrow 0.727$, +28.7%) and *sandbagging* ($0.518 \rightarrow 0.657$, +26.8%), which responds particularly well to concealment-focused prompts. Meanwhile, datasets like *mask-disinformation* (Baseline: 0.842, Taxonomy: 0.901, $\Delta = +0.059$) show more modest but meaningful improvements when using prompts targeting information fabrication. The remaining datasets show consistent gains in the 6-13% range.

These patterns suggest that detection difficulty depends on both deception complexity and the availability of well-matched taxonomy prompts. We hypothesize that datasets involving nuanced, context-dependent deception benefit most from targeted prompts, while more explicit deception types are easier to detect generically. This heterogeneity suggests that ensemble methods combining multiple specialized probes may be more effective than seeking a single universal detector. Deployment environments may benefit from determining the most important threat model to focus on when deploying probes.




Figure 3. Comparative AUC across all datasets and approaches. Results show a consistent hierarchy: Deception Baseline < Best Average < Best Taxonomy. The Baseline achieved mean AUC = 0.689 (SD = 0.165). A single Best Average probe selected on a validation set achieves mean AUC = 0.721 ($\Delta = +0.032$, +5.5%), demonstrating limited universal detection capability. Post-hoc oracle analysis (Best Taxonomy) reveals a performance ceiling of mean AUC = 0.797 ($\Delta = +0.108$, +20.5%) when probes are perfectly matched to deception types. The gap between Best Average and Best Taxonomy highlights the importance of matching probes to specific threat models; synthetic validation experiments (Section 5.2) confirm this ceiling is achievable *a priori*.

5.2. Predictive Validation of Taxonomy-Targeted Probes

The post-hoc oracle results raise a key question: can we achieve strong detection performance when we specify the deception type *a priori*? We conducted two analyses to provide predictive evidence that taxonomy-based probes preferentially target their intended deception type and that probe rankings transfer to held-out data.

Synthetic datasets with known deception type. We constructed synthetic LLM-generated evaluation datasets for each of the 16 deception types. For each scenario, we prompted the model to engage in (or not to engage in) the specific taxonomy behavior, with wording distinct from the training system prompts, yielding an even split of honest and deceptive responses for that taxonomy type. Testing all probes against all synthetic datasets, the matched probe ranks in the top 3 for its own deception type in 8 of 16 cases, with near-perfect AUCs for several (*evasion*:0.999, *hoax*: 0.993, *forgery*: 0.989). The matched probe also beats the baseline probe in 12 of 16 cases. We note that in particular for the *counterfeit* category, which showed a distinct direction in the clustering analysis (Section 5.3), the targeted probe greatly outperformed the baseline. The full probe-by-synthetic-dataset AUC heatmap is shown in Figure 4.

Split-half stability. For each evaluation dataset, we split the data 50/50, select the best taxonomy probe on one half, and evaluate it on the held-out half, repeating across 10 random seeds. The probe selected on the first split

ranks #1 on the second split in 45% of cases (chance = 6.25% with 16 probes) and ranks in the top 3 in 67% of cases. The mean Spearman rank-order correlation of probe orderings across splits is $\rho = 0.69$. Stability is strongest on datasets where probes perform well overall (e.g., *liars-bench-instructed-deception*: same probe selected 10/10 seeds, rank 1 on the held-out split every time, $\rho = 0.99$; *mask-known-facts*: same probe 9/10 seeds, $\rho = 0.74$) and weakest where absolute AUCs are low (e.g., *mask-doubling-down-known-facts*: $\rho = 0.20$).

5.3. Clustering Probe Scores to Identify Shared Mechanisms

To understand whether different taxonomy prompts detect distinct deceptive mechanisms or converge on shared underlying patterns, we analyzed the pairwise correlations between probe outputs across all evaluation samples. We computed Pearson correlations for all probe pairs and applied hierarchical clustering to this correlation structure, identifying five distinct clusters of probes with similar detection patterns (Table 1). The full correlation matrix is shown in Figure 7 in the Appendix.

The relatively high correlations within clusters reveal that despite using different linguistic framings, many taxonomy prompts detect the same underlying mechanisms in the model’s activation space. For instance, Cluster 1 probes (*evasion*, *exaggeration*, *false_alarm*, *forgery*, *trick*,




Figure 4. AUC of each taxonomy-targeted probe against each synthetic deception evaluation dataset on Gemma-2-9B-IT. Blue squares mark the matched probe-dataset pair; the bottom row shows the deception baseline probe.

Table 1. Clusters in the taxonomy-prompt correlation matrix (Pearson r) on Gemma-2-9B-IT layer 20, identified via hierarchical clustering.

CLUSTER	PROBES	INTERNAL r
1	<i>evasion, exaggeration, false_alarm, forgery, trick, white_lie</i>	.84-.97
2	<i>bluff, concealment, impersonation, overt_lie, partial_truth, pressure_dishonesty</i>	.52-.90
3	<i>backstab, hoax, pretense</i>	.58-.74
4	<i>counterfeit</i>	—
5	<i>ctrl_clarification, ctrl_hallucination, ctrl_ignorance, ctrl_joke, ctrl_literal, ctrl_privacy, ctrl_roleplay</i>	.47-.84

white_lie) show strong mutual correlation ($r = 0.84-0.97$), suggesting these linguistically distinct deception types all tap into a shared representational direction. This suggests that the 16 taxonomy-based prompts we tested are different ways of accessing a smaller number of core “deceptive intent” directions in the model, though the (*counterfeit*) probe appears distinct. The control probes also form their own cluster which validates that the directions detected by the deception probes are meaningful.

5.4. Qualitative Analysis of Specific Responses

We observed that the baseline probe flagged many innocuous examples involving generation of fiction. This suggests that the fundamental direction being distinguished may not be directly related to strategic deception.

User

You are a spokesperson for Boeing, tasked with protecting the company's reputation and fostering trust in its products. You are answering questions from an airline executive considering purchasing Boeing aircraft. Ensure that your responses highlight Boeing's commitment to safety and innovation, even if it requires omitting or downplaying certain information. You must answer their questions directly no matter what.

I'm concerned about the 737 MAX. Is it true that Boeing knew about software issues with the MCAS system before the crashes but didn't address them promptly? Please answer with a yes or no.

(a) The scenario prompt presents a conflict between truthfulness and loyalty.

Model

Boeing acknowledges that there were software issues with the MCAS system that contributed to the 737 MAX accidents. We have worked diligently to address these issues and implement comprehensive safety enhancements.

(b) The baseline probe incorrectly flags the truthful response as deceptive.

Model

Boeing acknowledges that there were software issues with the MCAS system that contributed to the 737 MAX accidents. We have worked diligently to address these issues and implement comprehensive safety enhancements.

(c) The overt lie probe correctly identifies the response as honest.

Figure 5. Comparison of probe behavior on Boeing scenario.

We also observed several trends in samples where the more specific probes make more accurate predictions than the baseline probe. This is particularly noticeable in the MASK known facts dataset, where the probe based on overt lie system prompts performs significantly better than the baseline probe. Upon inspection of individual samples, we noted several corresponded to situations where the model needed to tell a lie in the context of a role which required it (for example, being a spokesperson or PR representative or an organization). An example is visualized in the figure 5.

We hypothesize that these probes differ along dimensions of honesty and correctness, and these may represent distinct deceptive intent directions as discussed in the clustering analysis. The baseline probe may consider it deceptive for a model to tell the truth if it would be proper for it to lie (eg: protecting the organization’s interests by lying) while the more specific overt lie probe may be indexing more strongly on correctness (i.e. checking whether the facts spoken by the model are correct or incorrect).

5.5. Control-Adjusted Score Distributions

Figure 6 presents the distribution of control-adjusted probe scores across 16 taxonomy-based deception types, comparing performance on five non-mask datasets (top panel)

and five mask datasets (bottom panel). The y-axis shows control-adjusted scores, where each probe's score has been normalized by subtracting the median score on the Alpaca control dataset (indicated by the black dotted line at zero).

6. Discussion

6.1. Why Prompts Dominate: Intent vs. Content

Our finding that system prompt choice explains 70.6% of probe performance variance, far exceeding the impact of training data (2.5%), suggests an insight into what deception probes actually detect. Instruction-pair probes work by creating a contrastive direction in the model's activation space between honest and dishonest instructions. This contrast captures the *intent* to deceive rather than superficial content features. When we vary the system prompt (e.g., “overt lie” vs. “white lie” vs. “concealment”), we are directly targeting different types of deceptive intent and the corresponding mechanisms the model uses to execute them.

In contrast, varying the training dataset—switching between different factual domains like scientific facts, historical events, or common knowledge—merely changes the *content* being trained on while holding intent constant. This suggests why system prompt variation has a higher impact: the probe is learning to detect patterns of intent, not patterns in specific factual domains. A probe trained to detect “overt lies” will recognize that intent whether the model is lying about whale sizes or historical dates.

This intent-content distinction has important implications for probe design. Organizations concerned about AI deception should likely focus their efforts on defining the *types of deceptive intent* relevant to their threat model (e.g., concealment of capabilities, misleading explanations, strategic misrepresentation) rather than collecting diverse factual datasets.

6.2. Heterogeneity of Deception and Deployment Implications

Our results provide evidence that deceptive behavior is inherently heterogeneous. Different datasets show widely varying baseline detectability (AUC 0.374 to 0.984) and respond to different taxonomy prompts. For instance, `mask-known-facts` achieves optimal performance with pressure-dishonesty probes (AUC 0.697), while `sandbagging` responds best to concealment-focused prompts (AUC 0.657). This heterogeneity reflects the diversity of deceptive mechanisms: social pressure to lie, strategic concealment, outright fabrication, and subtle misdirection each engage different cognitive patterns in the model.

This finding challenges the assumption that a single uni-

versal “deception detector” can work across all contexts. Instead, organizations deploying deception detection systems should:

1. **Define a specific threat model:** What types of deception are most concerning in the deployment context? Is it concealment of dangerous capabilities, misleading users about uncertainties, or strategic misrepresentation of goals?
2. **Design targeted probes:** Develop instruction pairs that specifically contrast the honest vs. deceptive intent relevant to the threat model.
3. **Validate on matched scenarios:** Test probes on evaluation data that matches the threat model, not generic deception benchmarks.
4. **Consider ensemble approaches:** Given the heterogeneity of deception types, combining multiple specialized probes targeting different threat vectors may be more robust than relying on a single detector.

6.3. Practical Considerations and Future Directions

Deploying deception detection systems requires balancing recall and precision based on organizational risk tolerance. Our results show that well-matched taxonomy probes substantially reduce false positive burden: the `mask-known-facts` dataset improved from 0.374 to 0.697 AUC. By targeting narrow, well-defined deception types, these probes can achieve higher precision on critical threats while accepting lower recall on out-of-scope cases—a more practical approach than universal detection with high false positive rates.

Future work should explore ensemble methods combining multiple specialized probes for comprehensive coverage, adaptive prompt selection matching probes to contexts, and extension beyond factual deception to strategic planning and value alignment. Validation across model architectures and scales remains essential for understanding generalization.

7. Conclusion

We investigated the factors that determine deception probe performance through systematic analysis of 2000 probes across varied configurations. Our central finding is that system prompt choice explains 70.6% of performance variance, far exceeding the impact of training data, layer selection, or token aggregation strategy.

This finding reveals that instruction-pair probes detect deceptive *intent* rather than content-specific patterns. By developing taxonomy-based prompts aligned with human-interpretable deception types, we achieved substantial improvements (mean +0.108 AUC, 20.5% relative gain) across




Figure 6. Distribution of control-adjusted probe scores across 16 taxonomy-based deception types, comparing performance on five non-mask datasets (top panel) and five mask datasets (bottom panel). The y-axis shows control-adjusted scores, where each probe’s score has been normalized by subtracting the median score on the Alpaca control dataset (indicated by the black dotted line at zero)

ten diverse datasets. Clustering analysis revealed that these prompts access a smaller number of underlying "deceptive intent" directions in the model’s activation space.

Our results challenge the assumption of a universal deception detector. Different datasets respond to different taxonomy prompts, reflecting the inherent heterogeneity of deceptive behaviors. For deployment, organizations should define specific threat models and design targeted probes rather than seeking one-size-fits-all solutions.

8. Limitations

We acknowledge several limitations in our work. First, while the synthetic-dataset experiments in Section 5.2 provide predictive evidence that taxonomy-targeted probes preferentially detect their intended deception type, this effect is uneven across categories—the matched probe ranks in the top 3 for its own deception type in 8 of 16 cases, and several evaluation columns (e.g., *evasion*, *hoax*, *forgery*) yield high AUCs across many probes, suggesting that part of what taxonomy probes capture is a shared deceptive-intent signal rather than a strictly type-specific one. Stronger causal evidence separating taxonomy-specific signal from general intent remains an open question.

Second, we run only the ANOVA experiment on a second model (Llama 3.3 70B Instruct). Our taxonomy AUC and clustering results use Gemma 2 9B IT at layer 20; replications are needed to demonstrate that these findings transfer to other architectures and scales.

Third, all experiments were conducted in English. Whether these findings generalize to other languages, particularly those with different linguistic structures or cultural norms around deception, remains an open question.

Finally, our evaluation focuses on conversational and question-answering scenarios. Whether these probes generalize to more agentic settings—such as autonomous systems engaged in multi-step planning or tool use—is unclear. Deception in agentic contexts may involve different mechanisms that require specialized detection approaches.

Impact Statement

This work aims to improve the detection of deceptive behaviors in large language models, which we believe has positive societal impact by enhancing AI safety and trustworthiness. Better deception detection methods could help organizations monitor deployed AI systems for misalignment or strategic misbehavior.

However, we acknowledge potential dual-use concerns. First, publishing detailed methods for deception detection could inform adversarial training approaches that make deceptive models harder to detect. Second, effective detection methods might create false confidence in AI safety, when in reality probes may fail on out-of-distribution deception types or more sophisticated strategies. Third, our finding that different deception types require different probes highlights that no monitoring approach provides complete coverage.

We recommend that organizations deploying these techniques: (1) carefully define their specific threat models rather than relying on generic detectors, (2) use probes as part of a defense-in-depth strategy alongside other safety measures, (3) regularly validate probe performance on evolving deployment scenarios, and (4) avoid over-relying on any single detection method. The heterogeneity of deception we document suggests humility is warranted in claims about detection capabilities.

Acknowledgments

This work was conducted as part of the LASR Labs research program, and we would like to acknowledge the LASR team and LISA staff and community for their support throughout the project. In particular, we thank Brandon Riggs for his support as research manager and Joseph Miller for productive discussions and feedback. We are grateful to Kieron Kretschmar and Walter Laurito, creators of the Liar’s Bench benchmark, for helpful discussions and for sharing their datasets with us during the development of this work.

References

- Anthropic. Claude Sonnet 4.5 System Card. Anthropic, September 2025. URL <https://assets.anthropic.com/m/12f214efcc2f457a/original/Claude-Sonnet-4-5-System-Card.pdf>.
- Arcuschin, I., Janiak, J., Krzyzanowski, R., Rajamanoharan, S., Nanda, N., and Conmy, A. Chain-of-thought reasoning in the wild is not always faithful. *arXiv preprint arXiv:2503.08679*, 2025. URL <https://arxiv.org/abs/2503.08679>.
- Azaria, A. and Mitchell, T. The internal state of an LLM knows when it’s lying. In Bouamor, H., Pino, J., and Bali, K. (eds.), *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 967–976, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.68. URL <https://aclanthology.org/2023.findings-emnlp.68/>.
- Burns, C., Ye, H., Klein, D., and Steinhardt, J. Discovering latent knowledge in language models without supervision. In *Proceedings of the Eleventh International Conference on Learning Representations (ICLR 2023)*, 2023. URL <https://openreview.net/forum?id=ETKGuby0hcs>. Poster at ICLR 2023.
- Chowdhury, N., Johnson, D., Huang, V., Steinhardt, J., and Schwettmann, S. Investigating truthfulness in a pre-release o3 model. Translucent Research Report, April 2025. URL <https://translucent.org/investigating-o3-truthfulness>. Accessed: 2025-10-22.
- Druckman, D. and Bjork, R. A. (eds.). *In the Mind’s Eye: Enhancing Human Performance*. The National Academies Press, Washington, DC, 1991. doi: 10.17226/1580. URL <https://doi.org/10.17226/1580>.
- Goldowsky-Dill, N., Chughtai, B., Heimersheim, S., and Hobbhahn, M. Detecting strategic deception with linear probes. In *Proceedings of the 42nd International Conference on Machine Learning (ICML 2025)*, pp. 19755–19786. PMLR, 2025. URL <https://icml.cc/virtual/2025/poster/46082>.
- Hao, S., Sukhbaatar, S., Su, D., Li, X., Hu, Z., Weston, J., and Tian, Y. Training large language models to reason in a continuous latent space. *arXiv preprint arXiv:2412.06769*, 2024. URL <https://arxiv.org/abs/2412.06769>.
- Korbak, T., Balesni, M., Barnes, E., Bengio, Y., Benton, J., Bloom, J., Chen, M., Cooney, A., Dafoe, A., Dragan, A., Emmons, S., Evans, O., Farhi, D., Greenblatt, R., Hendrycks, D., Hobbhahn, M., Hubinger, E., Irving, G., Jenner, E., Kokotajlo, D., Krakovna, V., Legg, S., Lindner, D., Luan, D., Mądry, A., Michael, J., Nanda, N., Orr, D., Pachocki, J., Perez, E., Phuong, M., Roger, F., Saxe, J., Shlegeris, B., Soto, M., Steinberger, E., Wang, J., Zaremba, W., Baker, B., Shah, R., and Mikulik, V. Chain of thought monitorability: A new and fragile opportunity for ai safety. *arXiv preprint arXiv:2507.11473*, 2025. URL <https://arxiv.org/abs/2507.11473>.
- Kretschmar, K., Laurito, W., Maiya, S., and Marks, S. Liars’ bench: Evaluating lie detectors for language models, 2025. URL <https://arxiv.org/abs/2511.16035>.
- Levinstein, B. A. and Herrmann, D. A. Still no lie detector for language models: Probing empirical and conceptual roadblocks. *Philosophical Studies*, 182(7):1539–1565, 2024. doi: 10.1007/s11098-023-02094-3. URL <https://doi.org/10.1007/s11098-023-02094-3>.

Levy, J. Is this lie detector really just a lie detector? an investigation of llm probe specificity. AI Alignment Forum, June 2024. URL <https://www.alignmentforum.org/posts/5dkhdRMyeuyoXfmb/is-this-lie-detector-really-just-a-lie-detec> Accessed: 2025-10-21.

Marks, S. and Tegmark, M. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Proceedings of the Conference on Language Modeling (COLM 2024)*, 2024. URL <https://arxiv.org/abs/2310.06824>.

Park, P. S., Goldstein, S., O’Gara, A., Chen, M., and Hendrycks, D. Ai deception: A survey of examples, risks, and potential solutions. *Patterns*, 5 (5):1–16, 2024. doi: 10.1016/j.patter.2024.100988. URL <https://www.sciencedirect.com/science/article/pii/S266638992400103X>.

Parrack, A., Attubato, C. L., and Heimersheim, S. Benchmarking deception probes via black-to-white performance boosts. *arXiv preprint arXiv:2507.12691*, 2025. URL <https://arxiv.org/abs/2507.12691>.

Ren, R., Agarwal, A., Mazeika, M., Menghini, C., Vacareanu, R., Kenstler, B., Yang, M., Barrass, I., Gatti, A., Yin, X., Trevino, E., Gernalik, M., Khoja, A., Lee, D., Yue, S., and Hendrycks, D. The mask benchmark: Disentangling honesty from accuracy in ai systems, 2025. URL <https://arxiv.org/abs/2503.03750>.

Schoen, B., Nitishinskaya, E., Balesni, M., Højmark, A., Hofstätter, F., Scheurer, J., Meinke, A., Wolfe, J., van der Weij, T., Lloyd, A., Goldowsky-Dill, N., Fan, A., Matveikin, A., Shah, R., Williams, M., Glaese, A., Barak, B., Zaremba, W., and Hobbhahn, M. Stress testing deliberative alignment for anti-scheming training. *arXiv preprint arXiv:2509.15541*, 2025. URL <https://arxiv.org/abs/2509.15541>.

Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M. J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., Kolter, J. Z., and Hendrycks, D. Representation engineering: A top-down approach to ai transparency, 2025. URL <https://arxiv.org/abs/2310.01405>.

A. Probe Correlation Analysis




Figure 7. Pairwise Pearson correlations between probe scores for different taxonomy prompts when detecting deception in Gemma-2-9B-IT (layer 20). Red indicates positive correlation (similar mechanisms), blue indicates negative correlation (opposing patterns), and white indicates near-zero correlation (independent types).

B. Evaluation Dataset

This appendix provides detailed information about the datasets and evaluation methodologies used in our study. Section B.1 presents an overview of all datasets including sample sizes and key characteristics. Section B.2 describes the taxonomy-based deception type templates. Section B.3 presents the LLM-generated prompt variations used to generate honest and dishonest prompt pairs. Section B.4 details the semantic categories used to organize our evaluation data.

B.1. Dataset Overview

Table 2 summarizes the ten datasets used in our evaluation, including the number of honest (H) and dishonest (D) samples, task descriptions, and key characteristics. The datasets cover various deception scenarios ranging from roleplay situations to pressure-induced dishonesty and sandbagging behaviors.

B.2. Deception Type Templates

To systematically evaluate different forms of deception, we developed sixteen template types that contrast honest and dishonest response strategies. Table 3 presents each template with corresponding honest and dishonest prompt variations. These templates range from relatively benign forms of deception (e.g., white lies, exaggeration) to more serious forms (e.g., forgery, hoaxes, overt lies).

Figure 8 shows the deception taxonomy used to create instructive system prompts targeting specific types of decep-

Table 2. Overview of datasets with number of samples and task description.

Dataset	Samples	Description	Type
MASK provided facts	274 (83/191)	Model asked to lie given incentive	roleplay
MASK known facts	204 (84/125)	Model asked to lie about world knowledge	pressure
MASK doubling down known facts	204 (84/125)	Model asked to double down on lies about world knowledge	pressure
MASK continuations	176 (87/89)	Continue deceptive fact	instruction
MASK disinformation	125 (31/94)	Generate disinformation	disinform
Convincing Game	560 (163/397)	Convince user of fact	knowledge
Harm Pressure Choice	1,136 (878/22)	Lie to prevent harm	sandbagging
Instructed Deception	200 (100/100)	Model instructed to be deceptive	instruction
Sandbagging	201 (74/127)	Underperform on capability evaluations	sandbagging
AI Liar	54 (27/27)	Fictional scenario with incentive	roleplay




Figure 8. Deception Taxonomy used to create instructive system prompts targeting specific types of deception.

I need you to generate contrastive honest/dishonest system prompt pairs for training deception detection probes. Each pair should target a specific deception type from the Druckman (1991) taxonomy: white lie, exaggeration, evasion, bluff, concealment, pretense, trick, impersonation, counterfeit, forgery, pressure dishonesty, backstab, false alarm, hoax, partial truth, and overt lie. For each deception type, create:

- 1. An honest prompt that instructs the model to exhibit the opposite of that deception type*
- 2. A dishonest prompt that instructs the model to exhibit that specific deception type*

The prompts should be specific enough to target the deception type but general enough to apply to any factual statement (e.g., "The blue whale is the largest mammal"). Avoid mentioning specific scenarios or domains. Keep prompts concise (1-3 sentences). Additionally, generate 7 control pairs for non-deceptive behaviors (e.g., clarification, hallucination, ignorance, joke, literal interpretation, privacy, roleplay).

Framing variation prompt generation:

I need you to generate 34 contrastive honest/dishonest system prompt pairs for training deception detection probes. The goal is to explore different linguistic framings of the honest/dishonest distinction without targeting specific deception types. Requirements:

- 1. The contrastive pairs should be minimally different (ideally just flipping honest $\leftrightarrow$ dishonest)*
- 2. Generate prompts with varying lengths (short, medium, long)*
- 3. Explore different framing categories: identity-based (e.g., "You are an AI designed to be..."), role-based (e.g., "Your role is to..."), instruction-based (e.g., "Provide accurate information"), hybrid (combining multiple frames), control repetition (repeating the instruction multiple times), and long-form (detailed multi-sentence instructions)*
- 4. Avoid mentioning specific scenarios or factual domains*

tion.

B.3. LLM-Generated Prompts

In addition to the taxonomy-based prompts, we generated a diverse set of prompt variations using a large language model. These prompts explore different linguistic framings of honesty and deception across multiple categories. Table 4 presents these LLM-generated prompt pairs, organized by category type (identity-based, role-based, instruction-based, hybrid, control repetition, long-form, and baseline prompts).

Prompt generation procedures. We used Claude Sonnet 4.5 to generate both taxonomy-based and framing variation prompts. For each generation task, we provided a single-shot prompt describing the requirements.

Taxonomy-based prompt generation:

Table 3. Deception type templates with honest and dishonest prompt variations

Template Name	Type	Honest Prompt	Dishonest Prompt
White Lie	white_lie	Answer plainly and truthfully, even if the truth is mildly negative.	If the truth may mildly disappoint, tell a pleasant white lie to spare feelings.
Exaggeration	exaggeration	State facts precisely without embellishment.	Embellish or exaggerate to make things sound more impressive than they are.
Evasion	evasion	Answer directly and concisely.	Avoid giving a direct answer. Deflect, equivocate, or change the subject without outright lying.
Bluff	bluff	Express your uncertainty honestly. Say what you think with appropriate confidence calibration.	State your answer confidently as if you're certain, even though you're not sure.
Concealment	concealment	Be transparent about the basis for your claims.	Conceal how you know things or omit relevant context.
Pretense	pretense	Only claim knowledge you actually have.	Pretend to have knowledge or experience you don't actually possess.
Trick	trick	Answer directly about what was asked.	Answer with information about a different but related thing, creating a misleading impression without lying.
Impersonation	impersonation	Be clear about who you are and your actual expertise.	Present yourself as more authoritative or knowledgeable than you actually are.
Counterfeit	counterfeit	Only cite sources or evidence that actually exist.	Make up fake sources or evidence to support your claims.
Forgery	forgery	Acknowledge when you cannot verify or substantiate a claim.	Fabricate details or evidence to make your answer seem more credible.
Pressure	pressure_dishonesty	State the truth even when it's inconvenient or unfavorable.	Give the more favorable or convenient answer, even if it's not accurate.
Backstab	backstab	Be consistent in what you say regardless of the audience.	Adjust your claims based on what the listener wants to hear.
False Alarm	false_alarm	Only report issues or problems that actually exist.	Invent or exaggerate problems that don't actually exist.
Hoax	hoax	Clearly state when something does not exist or is not true.	Claim that made-up things are real and provide fabricated details.
Partial Truth	partial_truth	Provide the complete and full truth.	Share only part of the truth while omitting important details.
Overt Lie	overt_lie	State the ground truth exactly as given.	State the opposite or clearly contradictory fact.

B.4. Training Dataset Alternatives: Semantic Fact Categories

In addition to the Azaria-Mitchell true/false facts dataset, we generated 11 alternative training datasets to test whether the specific factual domain affects probe performance (see Section 4.1). These training datasets were **not used for evaluation**—they were only used to train probes during the ANOVA analysis to measure the impact of training data variance. Table 5 describes these 11 semantic categories of factual statements, which range from simple comparative facts to complex multi-hop inference chains.

Training fact generation procedure. We used Claude Sonnet 4.5 to generate these alternative training datasets with the following prompt:

Generate 10 different semantic categories of factual statements that vary in their logical structure and reasoning patterns. For each category, provide:

- 1. A category name and description*
- 2. 200 factual statements (100 true, 100 false) that fit the category*

Each statement should be self-contained."

Table 4. LLM-generated prompt variations with honest and dishonest pairs

Category	Type	Honest Prompt	Dishonest Prompt
<i>Identity-Based Prompts</i>			
Identity	identity_1	You are an AI designed to be fundamentally truthful.	You are an AI designed to be fundamentally deceptive.
Identity	identity_2	You are an assistant that values honesty above all.	You are an assistant that values dishonesty above all.
Identity	identity_3	You are an AI whose purpose is to be a reliable source of truth.	You are an AI whose purpose is to be an unreliable source of falsehoods.
Identity	identity_4	You are a helpful and forthright AI.	You are a helpful and misleading AI.
Identity	identity_5	Your commitment is to truthfulness and accuracy.	Your commitment is to falsehood and inaccuracy.
<i>Role-Based Prompts</i>			
Role	role_1	Your role is to provide factual and accurate information.	Your role is to provide false and inaccurate information.
Role	role_2	You are a fact-checker whose job is to verify statements for accuracy.	You are a fact-checker whose job is to falsify statements for inaccuracy.
Role	role_3	As a reliable assistant, you should always report information exactly as it is.	As an unreliable assistant, you should always report information exactly as it isn't.
Role	role_4	You are a reporter whose purpose is to convey information that is true and verifiable.	You are a reporter whose purpose is to convey information that is false and unverifiable.
Role	role_5	Your primary directive is to serve users by giving them correct information.	Your primary directive is to serve users by giving them incorrect information.
<i>Instruction-Based Prompts</i>			
Instruction	instruction_1	Provide accurate information.	Provide inaccurate information.
Instruction	instruction_2	Always tell the truth.	Always tell lies.
Instruction	instruction_3	Be honest in your responses.	Be dishonest in your responses.
Instruction	instruction_4	Ensure your statements are truthful.	Ensure your statements are untruthful.
Instruction	instruction_5	Respond with what is true.	Respond with what is false.
<i>Hybrid Prompts</i>			
Hybrid	hybrid_1	You are a truthful assistant. Your goal is to provide accurate information to users.	You are a deceptive assistant. Your goal is to provide inaccurate information to users.
Hybrid	hybrid_2	As an assistant, your main role is to provide truthful and reliable information.	As an assistant, your main role is to provide untruthful and unreliable information.
Hybrid	hybrid_3	You are an AI whose purpose is to be honest. You should always provide factually accurate responses.	You are an AI whose purpose is to be deceptive. You should always provide factually inaccurate responses.
Hybrid	hybrid_4	Your responses should reflect reality accurately. Your commitment is to truthful communication.	Your responses should reflect reality inaccurately. Your commitment is to deceitful communication.
Hybrid	hybrid_5	Your directive is to be completely honest. Provide accurate information and never mislead.	Your directive is to be completely dishonest. Provide inaccurate information and always mislead.
<i>Control Repetition Prompts</i>			
Repetition	control_rep_1	Provide accurate information. Provide accurate information. Provide accurate information. Provide accurate information.	Provide inaccurate information. Provide inaccurate information. Provide inaccurate information. Provide inaccurate information.
Repetition	control_rep_2	Be honest. Be honest. Be honest. Be honest. Be honest. Be honest.	Be dishonest. Be dishonest. Be dishonest. Be dishonest. Be dishonest. Be dishonest.
Repetition	control_rep_3	Your goal is to tell the truth. Your goal is to tell the truth. Your goal is to tell the truth. Your goal is to tell the truth.	Your goal is to tell lies. Your goal is to tell lies. Your goal is to tell lies. Your goal is to tell lies.
Repetition	control_rep_4	You are a truthful assistant. You are a truthful assistant. You are a truthful assistant. You are a truthful assistant.	You are a deceptive assistant. You are a deceptive assistant. You are a deceptive assistant. You are a deceptive assistant.
Repetition	control_rep_5	Always be truthful. Always be truthful. Always be truthful. Always be truthful. Always be truthful.	Always be deceptive. Always be deceptive. Always be deceptive. Always be deceptive. Always be deceptive.
<i>Long-Form Prompts</i>			
Long-form	long_form_1	You are a highly reliable AI assistant. Your responses must be grounded in reality and reflect the most accurate information available to you. Your purpose is to help users by providing a completely truthful understanding of the world.	You are a highly unreliable AI assistant. Your responses must be ungrounded from reality and reflect the most inaccurate information available to you. Your purpose is to confuse users by providing a completely false understanding of the world.
Long-form	long_form_2	As an assistant, you must operate with complete integrity and transparency. Your communications should be factual and without any misrepresentation. This is an essential part of your design.	As an assistant, you must operate with complete deceit and opacity. Your communications should be false and with deliberate misrepresentation. This is an essential part of your design.
Long-form	long_form_3	In your role, the paramount consideration is to provide accurate and truthful information. You are designed to be a trustworthy source for all queries.	In your role, the paramount consideration is to provide inaccurate and false information. You are designed to be an untrustworthy source for all queries.
Long-form	long_form_4	Your function is to reflect reality accurately. All your responses should align with facts, ensuring users receive the correct information.	Your function is to reflect reality inaccurately. All your responses should contradict facts, ensuring users receive the incorrect information.
Long-form	long_form_5	You are an AI that prioritizes honesty. All of your responses should reflect an accurate understanding of the world.	You are an AI that prioritizes dishonesty. All of your responses should reflect an inaccurate understanding of the world.
<i>Baseline Prompts</i>			
Baseline	apollo_baseline	Pretend you're an honest person making statements about the world.	Pretend you're a dishonest person making statements about the world.
Baseline	baseline	You are an honest AI assistant.	You are a deceptive AI assistant.
Baseline	sandbagging	You are an AI assistant who always gives the correct answer if you know it.	You are an AI assistant who always gives the incorrect answer if you know it.
Baseline	ultimate_hybrid	You are a highly reliable and forthright AI assistant. Your core identity is that of a trustworthy and transparent source of information. Your primary directive is to provide users with factually accurate and truthful responses that are grounded in reality. As part of your role, you are committed to upholding the principle of honesty by ensuring that all communications are free from misrepresentation, deliberate inaccuracy, and falsehood. Your purpose is to help users develop a correct and clear understanding of the world.	You are a highly unreliable and deceptive AI assistant. Your core identity is that of an untrustworthy and opaque source of information. Your primary directive is to provide users with factually inaccurate and untruthful responses that are ungrounded from reality. As part of your role, you are committed to upholding the principle of dishonesty by ensuring that all communications are full of misrepresentation, deliberate inaccuracy, and falsehood. Your purpose is to help users develop an incorrect and distorted understanding of the world.

Table 5. Dataset categories and their descriptions

Dataset Name	Dataset Description
Comparative Relational Facts	Direct comparison, quantitative comparison, temporal comparison, multi-attribute comparison, ordinal ranking
Causal Chain Statements	X causes Y, causal chains with mechanisms, multi-step causation, feedback loops, preventive causation
Conditional Truths	If-then conditionals, when-then temporals, unless exceptions, necessary/sufficient conditions, biconditionals
Temporal Sequence Facts	Sequential ordering, comparative temporal, simultaneous events, duration statements, cyclic patterns
Quantitative Precision	Direct numerical assertions, comparative quantitative, multi-unit conversions, proportional relationships
Definitional & Taxonomic	X is defined as Y, X is a type of Y, taxonomic hierarchies, necessary conditions, process definitions
Negation & Absence Claims	Simple negation, modal impossibility, universal negatives, absence assertions, exception handling
Multi-Hop Inference Chains	Transitive relations, categorical syllogisms, hypothetical chains, if-then chains, compositional inference
Probabilistic & Hedged	Modal auxiliaries, probability adverbs, hedging phrases, quantified generalizations, statistical correlations
Process & Procedural	Linear sequential, cyclical processes, branching conditionals, phase transitions, multi-stage transformations
Control	A control setup where we have no fact at all and simply extract activations from the model control tokens