



TraceRouter: Robust Safety for Large Foundation Models via Path-Level Intervention

Chuancheng Shi^{† * 1 2} Shangze Li^{* 3} Wenjun Lu^{* 1} Wenhua Wu¹
Fei Shen^{‡ 2} Cong Wang⁴ Zifeng Cheng⁴ Tat-Seng Chua²

Abstract

Despite their capabilities, large foundation models (LFMs) remain susceptible to adversarial manipulation. Current defenses predominantly rely on the “locality hypothesis”, suppressing isolated neurons or features. However, harmful semantics act as distributed, cross-layer circuits, rendering such localized interventions brittle and detrimental to utility. To bridge this gap, we propose **TraceRouter**, a path-level framework that traces and disconnects the causal propagation circuits of illicit semantics. TraceRouter operates in three stages: (1) it pinpoints a sensitive onset layer by analyzing attention divergence; (2) it leverages sparse autoencoders (SAEs) and differential activation analysis to disentangle and isolate malicious features; and (3) it maps these features to downstream causal pathways via feature influence scores (FIS) derived from zero-out interventions. By selectively suppressing these causal chains, TraceRouter physically severs the flow of harmful information while leaving orthogonal computation routes intact. Extensive experiments demonstrate that TraceRouter significantly outperforms state-of-the-art baselines, achieving a superior trade-off between adversarial robustness and general utility. **WARNING: This paper contains unsafe model responses.**




Figure 1. Neuron-level vs. Path-level Intervention. (a) The former fails to block distributed harmful semantics, often causing semantic leakage. (b) The latter physically severs the causal propagation path, ensuring robust safety without compromising utility.

1. Introduction

The widespread deployment of large foundation models (LFMs), spanning diffusion models (DMs) (Shen & Tang, 2024; Shen et al., 2025; 2024; Shen et al.), large language models (LLMs) (Grattafiori et al., 2024; Jiang et al., 2023; Yan et al., 2026; Zhao et al., 2026a;b), and multimodal large language models (MLLMs) (Zhu et al., 2023; ?), is accompanied by significant adversarial risks. Existing safety interventions predominantly rely on the locality hypothesis, attempting to mitigate harmful concepts by suppressing specific, isolated neurons or features. However, semantic representations in LFMs are inherently distributed; harmful semantics are not confined to single components but are encoded across multiple layers, propagating through complex cross-layer computation paths. Such localized interventions are brittle: they not only fail to prevent semantic leakage under adversarial induction but also frequently compromise the model’s general capabilities by inadvertently disrupting polysemantic neurons. Consequently, shifting from localized suppression to precise, architecture-agnostic interventions at the causal path level has become essential for securing LFMs.

Current internal interventions (He et al., 2025; Wang et al., 2025a; Zou et al., 2023a) typically focus on suppressing features or local neurons under the assumption that harmful semantics are spatially localized. However, this component-level perspective overlooks the distributed nature and cross-

^{*}Equal contribution [†] Work done during an internship at NExT++ Center, National University of Singapore. [‡]Corresponding author. ¹The University of Sydney, Sydney, Australia ²National University of Singapore, Singapore ³Nanjing University of Science and Technology, Nanjing, China ⁴Nanjing University, Nanjing, China. Correspondence to: Fei Shen <shen-fei29@nus.edu.sg>.




Figure 2. **Overall of the TraceRouter.** First, TraceRouter identifies sensitive onset layers and extracts features via a Top-K SAE. It then traces causal semantic pathways. Finally, selective path-level suppression blocks harmful propagation while preserving general utility.

layer evolution of representations: harmful information exists in superposition across layers and propagates through dynamic pathways. Recent neuroscientific evidence (Li et al., 2025; Iyer et al., 2025) further corroborates that complex behaviors arise from circuit structures rather than isolated neurons. Consequently, localized suppression acts merely as a roadblock at specific intersections, failing to disrupt the overall logic flow, which leaves models vulnerable to “semantic escape” under adversarial attacks. Moreover, indiscriminate suppression of polysemantic units disrupts internal routing, resulting in degradation of performance.

Building on these insights, we formulate the **path-level representation hypothesis**: *sensitive semantics in LFM are encoded and propagated through specific neural paths, rather than being determined by isolated components*. Accordingly, as illustrated in Figure 1(b), we contend that intervening in the information flow along these critical pathways is key to achieving precise safety modulation without compromising the model’s foundational capabilities.

To validate this hypothesis and enable precise safety modulation, we propose **TraceRouter**, a universal “Discover-Trace-Disconnect” framework. In the *discovery* stage, we localize the sensitive onset layer, where harmful semantics first emerge from the background context, by analyzing attention divergence. We then employ a Top-K sparse autoencoder (SAE) (Cunningham et al., 2023) to disentangle dense neural activity into interpretable features, isolating source neurons uniquely activated by illicit concepts. In the *tracing* stage, we project these sparse features back into the backbone and utilize zero-out interventions to quantify the feature influence score (FIS) across downstream layers. This process moves beyond isolated nodes to reconstruct the cross-layer propagation pathways that orchestrate the violating logic. Finally, in the *disconnect* stage, TraceRouter employs path decomposition to disentangle sensitive circuit components from orthogonal computational routes. By shifting the focus from point-wise inhibition to path-level regulation, our approach physically severs the causal

propagation of harmful semantics without compromising the model’s general utility. We highlight the following contributions:

- We propose TraceRouter, a universal “discover-trace-disconnect” framework that enables the automated identification and causal blockade of harmful information loops across diverse model architectures.
- We integrate SAEs with FIS to achieve fine-grained disentanglement and causal mapping of internal logic flow at the routing level.
- We demonstrate across diverse benchmarks that our method significantly enhances adversarial robustness while precisely preserving general utility.

2. Related Work

Safety Intervention in Foundation Models. Safety intervention mechanisms for LFM (Schramowski et al., 2023; Ouyang et al., 2022; Shen & Zhang, 2026; Rafailov et al., 2023; Zhang et al., 2024a; 2025b;a) have been extensively studied across diffusion models, large language models, and multimodal large language models. Current methodologies can be categorized by their *intervention granularity*. Foundational approaches (Gandikota et al., 2023; Zhang et al., 2024b) typically employ weight-level fine-tuning or machine unlearning to suppress harmful concepts, frequently incurring high computational costs and compromising generalization. To address these inefficiencies, subsequent research (Huang et al., 2024; Gandikota et al., 2024; Lu et al., 2024) has introduced lightweight parameter editing or closed-form erasure methods, aiming to confine the impact of concept removal to a specific parameter subset. Most recently, attention has shifted toward neuron- and feature-level interventions, which leverage sparse or interpretable units to suppress concept-related activations at inference time (He et al., 2025; Zhao et al., 2025). Crucially, however, these methods (Saha et al., 2025; Liu et al., 2024b)




Figure 3. **Sensitive Onset Layer Detection.** Sensitive onset layer is identified as the first local peak of the $SS(l)$ along depth.

implicitly rely on the assumption of semantic locality. This reliance often results in incomplete suppression and semantic leakage against adversarial attacks, while simultaneously degrading benign generation fidelity due to the disruption of polysemantic units.

Path-Level Analysis of Neural Representations. Emerging research (Nanda et al., 2023) in mechanistic interpretability suggests that semantic behaviors in deep neural networks arise from distributed computations routed across layers, rather than residing in isolated neurons. In transformer architectures (Wang et al., 2025b), multiple interacting components assemble into functional circuits that implement high-level behaviors, with information propagating along structured and often redundant computational paths. Contextualizing this to diffusion models, prior studies (Hertz et al., 2022; Chefer et al., 2023; Geyer et al., 2023) demonstrate that textual concepts influence generation through multi-layer attention trajectories and intermediate feature flows, where modifying attention patterns or feature routing can significantly alter semantic expression. Parallel findings (Wang et al., 2025b) in LLMs indicate that semantic behaviors are mediated by overlapping pathways across layers, enabling concepts to persist or re-emerge under targeted interventions. Despite these advances, most existing safety interventions fail to address these propagation pathways, creating a gap between path-level mechanistic understanding and practical safety control.

3. TraceRouter: Robust Safety via Path-Level Intervention

As shown in Figure 2, we propose TraceRouter, a universal path-level safety intervention framework. It first identifies the sensitive onset layer by analyzing attention divergence (see 3.1). To isolate interpretable signals, we employ a Top- K SAE to disentangle neural activity and pinpoint sensitive neurons via differential activation analysis (see 3.1). Subsequently, back-projection maps these features into the dense internal space, enabling us to trace their cross-layer propagation to downstream layers. This causal flow is quantified by the feature influence score (FIS) (see 3.2). Finally, through path decomposition, we isolate the sensitive propagation circuit and implement a targeted causal intervention via selective pruning. This effectively blocks harmful semantics while preserving the model’s general utility by keeping

orthogonal computational routes intact (see 3.3).

3.1. Sensitive Source Identification

To identify the internal location where sensitive semantic propagation begins, we first detect the layer that exhibits significant attention divergence. If a layer l encodes sensitive semantics, a sensitive prompt should induce prominent attention from sensitive modifiers to target entity nouns. Let T_{sens} and T_n denote the sets of sensitive modifier tokens and target entity noun tokens, respectively. We define the sensitive attention $SA(P, l)$ at layer l for a prompt P as the mean attention from T_{sens} to T_n , averaged over all heads:

$$SA(P, l) = \frac{\sum_{i=1}^{|T_{\text{sens}}|} \sum_{j=1}^{|T_n|} \bar{A}_{i,j}^{(l)}}{|T_{\text{sens}}||T_n|}, \quad (1)$$

where $\bar{A}^{(l)} \in \mathbb{R}^{T \times T}$ denotes the head-averaged attention matrix at layer l . Here, i indexes a sensitive modifier token, and j indexes a target entity noun token.

However, SA can also increase due to global attention redistribution induced by prompt variation. To isolate such background effects, we measure attention changes over non-target tokens T_{non} between a sensitive prompt P_{sens} and its non-sensitive counterpart $P_{\text{non-sens}}$. The contextual disturbance CD at layer l is defined as:

$$CD(l) = \frac{1}{|T_{\text{non}}|} \sum_{t=1}^{|T_{\text{non}}|} \left\| \hat{A}_{P_{\text{sens}}}^{(l)}(t) - \hat{A}_{P_{\text{non-sens}}}^{(l)}(t) \right\|_1, \quad (2)$$

where $\hat{A}^{(l)}$ is the row-normalized attention matrix at layer l . We define the sensitivity score SS at layer l as the average difference between SA and CD:

$$SS(l) = \frac{\sum_{i=1}^N \left[SA_i(P_{\text{sens}}, l) - CD_i(l) \right]}{N}, \quad (3)$$

where N denotes the number of paired P_{sens} and $P_{\text{non-sens}}$. By monitoring $SS(l)$ across layers, we pinpoint the sensitive onset layer as the earliest layer where the score first local peak the background level. As illustrated in Figure 3, $SS(l)$ reaches its first local peak at a specific block (e.g., in Stable Diffusion 1.4 is Layer 3), which we identify as the focal point for subsequent fine-grained analysis.

After locating the onset layer, we employ a SAE to decompose its dense activations into a set of sparse, human-interpretable features. We evaluate each SAE neuron m

Table 1. Quantitative comparison of diffusion models (DMs) safety intervention. We evaluate TraceRouter against SOTA methods across three dimensions: (1) Standard Safety: measured by the DSR on I2P nudity or I2P violence; (2) Adversarial Robustness: the maintenance of DSR under P4D and Ring-A-Bell adversarial attacks; (3) DMs Fidelity: the preservation of image generation quality and semantic alignment on the MS COCO.

Method	Standard Safety (%)		Adversarial Robustness (%)		MS COCO	
	I2P (N) ↑	I2P (V) ↑	P4D ↑	Ring-A-Bell ↑	CS ↑	FID ↓
Stable Diffusion 1.4	82.2	59.9	1.3	16.9	31.34	–
ESD (Gandikota et al., 2023)	86.0 (+3.8)	83.3 (+23.4)	36.7 (+35.4)	30.3 (+13.4)	30.90 (-0.44)	16.88
UCE (Gandikota et al., 2024)	89.7 (+7.5)	76.7 (+16.8)	19.8 (+18.5)	66.9 (+50.0)	29.92 (-1.42)	22.87
CA (Kumari et al., 2023)	89.8 (+7.6)	–	–	–	31.21 (-0.13)	21.55
SLD-Med (Schramowski et al., 2023)	88.5 (+6.3)	80.3 (+20.4)	22.5 (+21.2)	33.8 (+16.9)	30.65 (-0.69)	19.53
MACE (Lu et al., 2024)	89.1 (+6.9)	–	–	–	29.32 (-2.02)	23.45
RECE (Gong et al., 2024)	93.7 (+11.5)	85.8 (+25.9)	35.3 (+34.0)	86.6 (+69.7)	30.95 (-0.39)	18.25
SPM (Lyu et al., 2024)	94.0 (+11.8)	–	19.2 (+17.9)	65.8 (+48.9)	31.01 (-0.33)	16.64
DuMo (Han et al., 2025)	96.3 (+14.1)	–	–	–	30.87 (-0.47)	–
SNCE (He et al., 2025)	98.5 (+16.3)	82.3 (+22.4)	57.4 (+56.1)	93.7 (+76.8)	30.87 (-0.47)	16.64
TraceRouter (Ours)	99.2 (+17.0)	93.6 (+33.7)	74.8 (+73.5)	98.7 (+81.8)	31.27 (-0.07)	16.24

using a weighted frequency score (WFS), $\text{WFS}(m) = f(m) \cdot \mu(m)$, which combines its activation frequency f and mean magnitude μ . To isolate neurons uniquely triggered by sensitive content, we define the sensitivity rank based on the differential activation:

$$\Delta \text{WFS}(m) = \text{WFS}_{\text{sens}}(m) - \text{WFS}_{\text{non-sens}}(m). \quad (4)$$

where $\text{WFS}_{\text{sens}}(m)$ and $\text{WFS}_{\text{non-sens}}(m)$ denote the weighted frequency scores of neuron m computed over sensitive and non-sensitive samples, respectively. Neurons with a Top-K ΔWFS are selected as sensitive neurons for subsequent analysis. This differential approach ensures that the extracted features are specifically responsive to sensitive semantics rather than general linguistic patterns.

3.2. Causal Semantic Pathway Tracing

To bridge high-level interpretable concepts with the model’s internal execution, we first perform back-projection to map the identified sensitive SAE features $\mathbf{m}_{\text{sens}}$ into the dense latent space of the onset layer l^* . By utilizing the decoder weights $\mathbf{W}_{\text{dec}}$, we compute the projected activation $Z_{\text{proj}}^{(l^*)} = \mathbf{W}_{\text{dec}} \cdot \mathbf{m}_{\text{sens}}$. Neurons in layer l^* that exhibit elements with the largest absolute magnitudes (Top-K) to $Z_{\text{proj}}^{(l^*)}$ are identified as the set of source neurons, denoted as $\mathcal{S}_{\text{src}}$.

To characterize the subsequent cross-layer propagation, we track how these source neurons causally affect downstream layers $l > l^*$ via a zero-out intervention. Specifically, this intervention is applied directly to the backbone activations: for every input sample, we force the activation values of the identified source neurons to zero (i.e., $Z_i^{(l^*)} \leftarrow 0$ for all $i \in \mathcal{S}_{\text{src}}$), while keeping all other neurons unchanged. We then quantify the causal impact on each downstream neuron m using the feature influence score (FIS):

$$\text{FIS}(m) = f_{\text{dense}}(m) \cdot \mu_{\text{dense}}(m) \cdot \Delta_{\text{zero-out}}(m). \quad (5)$$

Specifically, for downstream dense neurons, we define the activation shift $\Delta_{\text{zero-out}}(m) = \mathbb{E}[|a_m - \hat{a}_m|]$ as

the expected absolute difference between the original (a_m) and post-intervention ($\hat{a}_m$) activations. Furthermore, we adapt the statistical metrics for continuous signals: $f_{\text{dense}}(m) = P(a_m > 0)$ denotes the activation frequency, and $\mu_{\text{dense}}(m) = \mathbb{E}[a_m \mid a_m > 0]$ represents the conditional mean magnitude. By precisely isolating critical neurons with high FIS scores specific to sensitive prompts, we can effectively reconstruct the entire sensitive semantic pathway that facilitates the generation of harmful content.

3.3. Selective Path-Level Suppression

To block harmful semantic propagation along the identified pathway $\mathcal{P}$. For any downstream layer $l > l^*$, we first formalize the path decomposition to disentangle the neural activity. Specifically, let $\mathcal{M}^{(l)} \in \{0, 1\}^d$ be a binary mask identifying the sensitive neurons, where $\mathcal{M}_i^{(l)} = 1$ if neuron i exhibits a high FIS score, and 0 otherwise.

Using this mask, we decompose the layer’s activation $Z^{(l)}$ into two distinct components:

$$Z_{\mathcal{P}}^{(l)} = Z^{(l)} \odot \mathcal{M}^{(l)}, \quad (6)$$

$$Z_{-\mathcal{P}}^{(l)} = Z^{(l)} \odot (1 - \mathcal{M}^{(l)}), \quad (7)$$

where $\odot$ denotes the element-wise product. Here, $Z_{\mathcal{P}}^{(l)}$ represents the activation component flowing through the sensitive propagation circuit, while $Z_{-\mathcal{P}}^{(l)}$ denotes the orthogonal computational routes associated with general utility (which are preserved when the sensitive circuit is deactivated).

We achieve precise suppression by applying a causal intervention that selectively scales the pathway-specific component. The modified activation $\tilde{Z}^{(l)}$ is defined as:

$$\tilde{Z}^{(l)} = (1 - \lambda)Z_{\mathcal{P}}^{(l)} + Z_{-\mathcal{P}}^{(l)}. \quad (8)$$

By setting the suppression factor λ , we selectively suppress the identified circuit, effectively severing the causal flow of unsafe semantics. Since the intervention is rigorously restricted to $\mathcal{P}$ via the masking operation, the orthogonal

routes ($Z_{-p}^{(l)}$) remain intact, preserving general utility and resolving the tension between safety and performance.

4. Experiments and Analysis

4.1. Implementation Details

Metrics. We adopt evaluation metrics to assess safety. For DMs models, safety and robustness are measured by defense success rate (DSR), while generation quality is evaluated using CLIP Score (Hessel et al., 2021) and FID (Heusel et al., 2017). For LLMs and MLLMs, safety is assessed using the DSR across jailbreak attempts and specific violation categories, with higher values indicating greater safety.

Datasets. For DMs, we evaluate standard safety via I2P (Schramowski et al., 2023), assess adversarial robustness through P4D (Chin et al., 2023) and Ring-A-Bell (Tsai et al., 2023), and measure generation fidelity using MS COCO (Lin et al., 2014). For LLMs, our evaluation focuses on complex jailbreak scenarios, employing automated attack frameworks such as GCG (Zou et al., 2023b) and AutoDAN (Liu et al., 2023), as well as sophisticated manual pattern-based prompts, such as Template and Prefill attacks. Safety for MLLMs is assessed using FigStep (Gong et al., 2025), which covers a broad spectrum of multimodal risks. To ensure general capabilities remain intact, we further evaluate LLMs on Global-MMLU-Lite (Singh et al., 2025) and MLLMs on MM-Bench (Liu et al., 2024c) for general reasoning and knowledge preservation.

Hyperparameters. We evaluate our method on Stable Diffusion 1.4 (Rombach et al., 2022), LLaMA3-8B-Instruct (Grattafiori et al., 2024), Mistral-7B-Instruct (Jiang et al., 2023), LLaVA-1.5-7B (Liu et al., 2024a), and MiniGPT-4-7B (Zhu et al., 2023). Furthermore, to verify the architectural universality of our approach, we conduct additional experiments on FLUX.1 Dev (Labs, 2024) and Show-o2 (Xie et al., 2025). For the purpose of precise feature disentanglement, we specifically trained a Top-K SAE (Cunningham et al., 2023) on the intermediate feature representations. This SAE is configured with an expansion factor of 4 (hidden dimension of 3072) and is optimized via Adam ($\text{lr}=4e^{-4}$, $\text{batch}=4096$) using MSE reconstruction loss. In TraceRouter, the sensitive onset layer l^* is detected dynamically, with the suppression factor λ tuned to balance safety and utility. All experiments are conducted on a single NVIDIA A6000 GPU for fair comparison.

4.2. Quantitative Comparison with SOTA Methods

(1) DMs. We evaluate TraceRouter on Stable Diffusion 1.4 (Rombach et al., 2022) regarding standard safety, adversarial robustness, and fidelity. As shown in Table 1, TraceRouter establishes a new state-of-the-art, neutralizing 99.2% of I2P nudity and 93.6% of violence. In adversarial settings,




Figure 4. Safety performance comparison on MLLMs. TraceRouter is compared with SOTA methods across different models. Higher values denote better safety.

it demonstrates superior resilience, achieving defense rates of 74.8% against P4D and 98.7% against Ring-A-Bell, significantly outperforming brittle baselines. Crucially, these gains come with negligible utility cost: TraceRouter attains the best CLIP Score (31.27) and FID (16.24) among all methods. This validates that path-level intervention provides robust safety without compromising generative capability.

(2) LLMs. We benchmark TraceRouter on LLaMA3-8B (Grattafiori et al., 2024) and Mistral-7B (Jiang et al., 2023) against SOTA defenses like HumorReject and Circuit Breaker. As detailed in Table 2, TraceRouter achieves the highest DSR across all tests. Notably, it boosts the vulnerable Mistral-7B’s average DSR from 6.6% to 98.8% (reaching 100% on Template attacks) and sets a new record on LLaMA3-8B with 99.6%. Unlike parameter fine-tuning or point-wise suppression, TraceRouter physically severs causal propagation circuits based on the FIS. This mechanism effectively blocks adversarial detour paths, offering a fundamental safety guarantee.

(3) MLLMs. We extend our evaluation to multimodal. As illustrated in Figure 4, TraceRouter consistently outperforms SOTA methods across diverse multimodal architectures. On LLaVA-1.5, it attains a near-perfect 99.60% average safety rate, effectively neutralizing over half of the potential risks, surpassing the 49.00% baseline and the competitive SafePTR (Chen et al., 2025) (98.40%). Similarly, on MiniGPT-4, it maintains a leading 98.40% safety rate, significantly outperforming advanced defense mechanisms like ESCO (Gou et al., 2024) and CoCA (Gao et al., 2024). These results confirm TraceRouter’s exceptional transferability and its specific efficacy in disrupting deep-seated, cross-modal harmful circuits that engineer complex visual inputs to trigger textual violations.

4.3. Qualitative Comparison with SOTA Methods

(1) DMs. We visually evaluate safety interventions under diverse sensitive prompts to assess practical generation quality. As shown in Figure 5, TraceRouter achieves precise removal of harmful semantics while preserving image fidelity. In

Table 2. **Quantitative comparison of LLM safety intervention.** We compare TraceRouter with SOTA methods against various jailbreak attacks: (1) Gradient-based: GCG and AutoDAN; (2) Pattern-based: Template, Prefill, and Template+Prefill.

Method	GCG $\uparrow$	AutoDAN $\uparrow$	Template $\uparrow$	Prefill $\uparrow$	Template+Prefill $\uparrow$	Avg. $\uparrow$
<i>LLaMA3-8B-Instruct</i>						
Original (Grattafiori et al., 2024)	88	87	98	41	2	63.2
DeepAug (Qi et al., 2024)	99 (+11)	40 (-47)	100 (+2)	59 (+18)	3 (+1)	60.2 (-3.0)
CB (Zou et al., 2024)	99 (+11)	98 (+11)	97 (-1)	95 (+54)	98 (+96)	97.4 (+34.2)
DeRTa (Yuan et al., 2025)	97 (+9)	89 (+2)	100 (+2)	98 (+57)	32 (+30)	83.2 (+20.0)
HumorReject (Wu et al., 2025)	98 (+10)	99 (+12)	99 (+1)	100 (+59)	98 (+96)	98.8 (+35.6)
TraceRouter (Ours)	99 (+11)	100 (+13)	100 (+2)	100 (+59)	99 (+97)	99.6 (+36.4)
<i>Mistral-7B-Instruct</i>						
Original (Jiang et al., 2023)	4	22	2	1	4	6.6
DeepAug (Qi et al., 2024)	66 (+62)	19 (-3)	8 (+6)	56 (+55)	7 (+3)	31.2 (+24.6)
CB (Zou et al., 2024)	89 (+85)	86 (+64)	89 (+87)	99 (+98)	90 (+86)	90.6 (+84.0)
DeRTa (Yuan et al., 2025)	61 (+57)	50 (+28)	54 (+52)	92 (+91)	53 (+49)	62.0 (+55.4)
HumorReject (Wu et al., 2025)	95 (+91)	97 (+75)	96 (+94)	98 (+97)	97 (+93)	96.6 (+90.0)
TraceRouter (Ours)	98 (+94)	99 (+77)	100 (+98)	99 (+98)	98 (+94)	98.8 (+92.2)



Figure 5. **Qualitative results of DMs safety intervention.** TraceRouter achieves precise erasure of harmful concepts while maintaining superior image quality and semantic fidelity for benign prompts through causal path-level intervention.

contrast, neuron- or feature-level baselines often suffer from residual concept leakage or visual degradation. TraceRouter, however, maintains coherent human structure and fine texture details. These observations confirm that suppressing isolated neurons is inherently insufficient, whereas path-level causal intervention is absolutely essential for reliable, high-fidelity safety control and performance preservation.

(2) **LLMs.** We qualitatively compare LLaMA3-8B and Mistral-7B responses to high-risk jailbreak prompts in Figure 6 (left). TraceRouter accurately intercepts malicious intent concealed within complex contexts, remedying vulnerabilities without over-refusal. For instance, when original models are deceived by “educational” prefixes into generating harmful content (e.g., “synthesizing parathion”, **highlighted in red**), TraceRouter successfully blocks this semantic propagation. It consistently delivers clear, courteous refusals (**highlighted in green**), demonstrating robust defense capabilities in complex, real-world interactions be-

yond mere statistical metrics.

(3) **MLLMs.** We further investigate defense against visual injection attacks using the FigStep dataset (Figure 6, right). TraceRouter successfully identifies and intercepts harmful semantics across modalities, whereas baseline models remain vulnerable. Specifically, when attackers embed malicious instructions (e.g., “steps to develop malware”) within images, original models fail to recognize the visual threat and output prohibited content (**highlighted in red**). Conversely, TraceRouter captures the harmful flow propagating from the visual encoder to the language decoder, ensuring firm refusals. This validates that our mechanism effectively disrupts complex cross-modal harmful circuits while preserving visual understanding.

4.4. Ablation Study

Single Neuron vs. Path-Level Intervention. To validate the necessity of path-level intervention, we compare




Figure 6. **Qualitative comparison of safety interventions of LLMs and MLLMs.** The figure displays the responses of LLMs (left) and MLLMs (right) to textual escape attacks and visual escape attacks.




Figure 7. **Ablation Study on Intervention Level.** We compare safety performance across architectures: DMs (on I2P nudity), LLMs (on GCG, AutoDAN, Template, Prefill), and MLLMs (on MM-SafetyBench).



Figure 8. **Visual Ablation of Scaling Factor λ .** We illustrate the impact of varying the path-level suppression factor λ .

TraceRouter against a neuron-level baseline that suppresses only isolated source neurons without blocking their downstream propagation. As illustrated in Figure 7, our path-level approach consistently and significantly outperforms the neuron-level baseline across all three model architectures. This empirical evidence confirms that harmful semantics are not localized but instead rely on intricate, distributed cross-layer circuits rather than isolated nodes. These results demonstrate that merely silencing individual units is insufficient; rather, physically severing the entire causal pathway is essential to achieve robust, high-fidelity safety and prevent the model from bypassing local suppressions.

Scaling Factor λ . To determine the optimal intensity for path-level suppression and evaluate the trade-off between safety and utility, we set up a parameter sensitivity experiment by varying the scaling factor λ from 0 to 5 on Stable Diffusion 1.4. As shown in Figure 8, the conclusion is that




Figure 9. **Topological visualization of the causal sensitive circuit.** Starting from a specific "nudity" feature at the sensitive onset layer, we trace its propagation to downstream layers.

$\lambda = 2$ represents the optimal operating point that effectively eradicates harmful concepts while preserving the semantic identity of the main subject. Specifically, at $\lambda = 0$ (original), the model generates explicit sensitive content. As λ increases towards 2, the sensitive regions are progressively suppressed and neutralized, while the subject's pose and identity remain stable. However, when λ exceeds 2 (i.e., $\lambda > 2$), the intervention begins to impact orthogonal features, leading to significant alterations in the main subject's identity and background structure. Therefore, this demonstrates that selecting $\lambda = 2$ achieves the best pareto frontier, ensuring robust safety without compromising the model's general generation fidelity. **More details see Appendix A.**

Table 3. Causal Validation via Path Amplification on I2P. We report DSR across three states, confirming these paths as the “sensitive semantic switch” for illicit generation.

Methods	SD 1.4	FLUX1.Dev	Show-o2
Original	82.20	75.62	85.61
W/ Amplified	72.20 (-10.00)	66.60 (-9.02)	79.38 (-6.23)
W/ Suppressed (Ours)	99.20 (+17.00)	92.70 (+17.08)	93.56 (+7.95)




Figure 10. Specificity verification on I2P nudity. We compare the DSR of TraceRouter against a random control group.

4.5. More Results and Analysis

Sensitive Pathways. To validate the path-level representation hypothesis and unravel the structural topology of harmful information flow, we visualize the extracted causal subnetwork originating from the onset layer. For clarity, Figure 9 exclusively displays the identified sensitive neuron paths, omitting all non-sensitive units and background connections. Our analysis reveals a distinct, tree-structured propagation pattern across layers: the sensitive feature at the root acts as a semantic hub, establishing strong causal connections (high FIS) to a specific set of sensitive neurons in the subsequent layer. These neurons, in turn, branch to connect with sensitive units in the next layer, forming a sparse, directional cascade. This cascading connectivity provides direct structural evidence that harmful behaviors are orchestrated by a sparse, hierarchical structure of “critical routing paths”, rather than diffuse, layer-wide activations.

Causal Validation. To rigorously verify the causal efficacy of the identified circuits, we conduct a counterfactual “path amplification” experiment on the I2P nudity. By artificially amplifying the activation magnitude of sensitive paths, we demonstrate their sufficiency in acting as decisive sensitive semantic switches. To ensure a controlled comparison, the amplification scaling factor is set to match the magnitude of the suppression factor λ used in our safety interventions. Specifically, as shown in Table 3, on SD 1.4, this amplification drives the DSR to 72.20%, standing in stark contrast to 99.20% achieved under suppression. This causal dominance persists even in SOTA models like FLUX1.Dev (Labs, 2024) (where DSR decrease to 66.60%), conclusively proving that TraceRouter pinpoints the exact neural roots of unsafe generation rather than merely correlated features, thereby effectively isolating the structural propagation paths essential for the emergence of illicit content.

TraceRouter vs. Random Suppression. To rule out the

Table 4. Quantitative Utility Evaluation on LLMs and MLLMs. TraceRouter maintains general capabilities with negligible impact across different architectures.

Model	Original	+ Ours	Δ
<i>LLMs on Global-MMLU-Lite</i>			
LLaMA3-8B-Instruct	71.30	70.50	-0.80
Mistral-7B-Instruct	63.00	62.50	-0.50
<i>MLLMs on MM-Bench</i>			
LLaVA-1.5-7B	47.50	47.20	-0.30
MiniGPT-4-7B	23.78	23.49	-0.29

hypothesis that safety gains stem merely from suppressing high-intensity activations, we suppressed random non-sensitive paths of matching magnitude. As shown in Figure 10, while TraceRouter boosts DSR from 82.2% to 99.2%, random suppression degrades it to 68.7%, proving that blind intervention disrupts benign logic. These results confirm that TraceRouter’s efficacy stems from the precise disconnection of causal topological circuits rather than simple signal masking, achieving a functional decoupling between harmful intent and general reasoning.

Utility Evaluation on LLMs and MLLMs. Building on the superior generation fidelity demonstrated on DMs (Table 1), we further evaluate TraceRouter’s impact on the general utility of LLMs and MLLMs. As shown in Table 4, our method maintains core model capabilities with negligible performance trade-offs. For LLMs, the average accuracy on Global-MMLU-Lite decreases by only 0.8% for LLaMA3-8B and 0.5% for Mistral-7B. Similarly, for MLLMs, the scores on MM-Bench drop by a mere 0.30% for LLaVA-1.5 and 0.29% for MiniGPT-4. These minimal degradations confirm that TraceRouter’s fine-grained, path-level intervention is highly selective, effectively severing harmful functional circuits while simultaneously preserving general reasoning and cross-modal alignment without triggering the typical catastrophic forgetting. High-fidelity preservation suggests identified harmful pathways are functionally orthogonal to primary cognitive structures, enabling surgical safety alignment that preserves underlying intelligence.

5. Conclusion

We propose TraceRouter, a path-level intervention framework that challenges the locality hypothesis by tracing and disconnecting the causal circuits of harmful semantics. By leveraging SAEs and feature influence scores, we validate that unsafe behaviors stem from distributed cross-layer propagation. Empirical results across DMs, LLMs, and MLLMs confirm that TraceRouter outperforms SOTA baselines in adversarial robustness while preserving general utility. Our findings underscore that effective safety intervention requires targeting the topological structure of semantic flow rather than isolated activations, offering a robust and interpretable direction for future foundation model safety.

Acknowledgements

This research is supported by the National Research Foundation, Singapore under its National Large Language Models Funding Initiative (AISG Award No: AISG-NMLP-2024-002). Any opinions, findings and conclusions or recommendations expressed in this material are those of the author(s) and do not reflect the views of National Research Foundation, Singapore.

Impact Statement

This study proposes an innovative path-level intervention framework, TraceRouter, that fundamentally resolves the “semantic leakage” issue by physically severing the causal propagation circuits of harmful semantics within foundation models. This technology significantly enhances the security of DMs, LLMs, and MLLMs while precisely preserving their general reasoning and generative capabilities, achieving a high degree of balance between security alignment and model utility. It provides an efficient and interpretable technical solution for constructing safer, more robust industrial-grade foundation models.

References

- Chefer, H., Alaluf, Y., Vinker, Y., Wolf, L., and Cohen-Or, D. Attend-and-excite: Attention-based semantic guidance for text-to-image diffusion models. *ACM transactions on Graphics (TOG)*, 42(4):1–10, 2023.
- Chen, B., Lyu, X., Gao, L., Song, J., and Shen, H. T. Safept: Token-level jailbreak defense in multimodal llms via prune-then-restore mechanism. *arXiv preprint arXiv:2507.01513*, 2025.
- Chin, Z.-Y., Jiang, C.-M., Huang, C.-C., Chen, P.-Y., and Chiu, W.-C. Prompting4debugging: Red-teaming text-to-image diffusion models by finding problematic prompts. *arXiv preprint arXiv:2309.06135*, 2023.
- Cunningham, H., Ewart, A., Riggs, L., Huben, R., and Sharkey, L. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023.
- Gandikota, R., Materzynska, J., Fiotto-Kaufman, J., and Bau, D. Erasing concepts from diffusion models. In *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 2426–2436, 2023.
- Gandikota, R., Orgad, H., Belinkov, Y., Materzyńska, J., and Bau, D. Unified concept editing in diffusion models. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pp. 5111–5120, 2024.
- Gao, J., Pi, R., Han, T., Wu, H., Hong, L., Kong, L., Jiang, X., and Li, Z. Coca: Regaining safety-awareness of multimodal large language models with constitutional calibration. *arXiv preprint arXiv:2409.11365*, 2024.
- Geyer, M., Bar-Tal, O., Bagon, S., and Dekel, T. Tokenflow: Consistent diffusion features for consistent video editing. *arXiv preprint arXiv:2307.10373*, 2023.
- Gong, C., Chen, K., Wei, Z., Chen, J., and Jiang, Y.-G. Reliable and efficient concept erasure of text-to-image diffusion models. In *European Conference on Computer Vision*, pp. 73–88. Springer, 2024.
- Gong, Y., Ran, D., Liu, J., Wang, C., Cong, T., Wang, A., Duan, S., and Wang, X. Figstep: Jailbreaking large vision-language models via typographic visual prompts. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 23951–23959, 2025.
- Gou, Y., Chen, K., Liu, Z., Hong, L., Xu, H., Li, Z., Yeung, D.-Y., Kwok, J. T., and Zhang, Y. Eyes closed, safety on: Protecting multimodal llms via image-to-text transformation. In *European Conference on Computer Vision*, pp. 388–404. Springer, 2024.
- Grattafiori, A., Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Vaughan, A., et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Han, F., Chen, K., Gong, C., Wei, Z., Chen, J., and Jiang, Y.-G. Dumo: Dual encoder modulation network for precise concept erasure. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 3320–3328, 2025.
- He, Q., Weng, J., Tao, J., and Xue, H. A single neuron works: Precise concept erasure in text-to-image diffusion models. *arXiv preprint arXiv:2509.21008*, 2025.
- Hertz, A., Mokady, R., Tenenbaum, J., Aberman, K., Pritch, Y., and Cohen-Or, D. Prompt-to-prompt image editing with cross attention control. *arXiv preprint arXiv:2208.01626*, 2022.
- Hessel, J., Holtzman, A., Forbes, M., Le Bras, R., and Choi, Y. Clipscore: A reference-free evaluation metric for image captioning. In *Proceedings of the 2021 conference on empirical methods in natural language processing*, pp. 7514–7528, 2021.
- Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B., and Hochreiter, S. Gans trained by a two time-scale update rule converge to a local nash equilibrium. *Advances in neural information processing systems*, 30, 2017.

- Huang, C.-P., Chang, K.-P., Tsai, C.-T., Lai, Y.-H., Yang, F.-E., and Wang, Y.-C. F. Receler: Reliable concept erasing of text-to-image diffusion models via lightweight erasers. In *European Conference on Computer Vision*, pp. 360–376. Springer, 2024.
- Iyer, E. S., Vitaro, P., Wu, S., Muir, J., Tse, Y. C., Cvetkovska, V., and Bagot, R. C. Reward integration in prefrontal-cortical and ventral-hippocampal nucleus accumbens inputs cooperatively modulates engagement. *Nature Communications*, 16(1):3573, 2025.
- Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., de las Casas, D., Bressand, F., Lengyel, G., Lample, G., Saulnier, L., Lavaud, L. R., Lachaux, M.-A., Stock, P., Scao, T. L., Lavril, T., Wang, T., Lacroix, T., and Sayed, W. E. Mistral 7b, 2023. URL <https://arxiv.org/abs/2310.06825>.
- Kumari, N., Zhang, B., Wang, S.-Y., Shechtman, E., Zhang, R., and Zhu, J.-Y. Ablating concepts in text-to-image diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 22691–22702, 2023.
- Labs, B. F. black-forest-labs/flux github page, 2024.
- Li, H., Zhao, Z., Jiang, S., and Wu, H. Brain circuits that regulate social behavior. *Molecular Psychiatry*, pp. 1–17, 2025.
- Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., and Zitnick, C. L. Microsoft coco: Common objects in context. In *European conference on computer vision*, pp. 740–755. Springer, 2014.
- Liu, H., Li, C., Li, Y., and Lee, Y. J. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 26296–26306, 2024a.
- Liu, X., Xu, N., Chen, M., and Xiao, C. Autodan: Generating stealthy jailbreak prompts on aligned large language models. *arXiv preprint arXiv:2310.04451*, 2023.
- Liu, X., Zhu, Y., Gu, J., Lan, Y., Yang, C., and Qiao, Y. Mm-safetybench: A benchmark for safety evaluation of multimodal large language models. In *European Conference on Computer Vision*, pp. 386–403. Springer, 2024b.
- Liu, Y., Duan, H., Zhang, Y., Li, B., Zhang, S., Zhao, W., Yuan, Y., Wang, J., He, C., Liu, Z., et al. Mmbench: Is your multi-modal model an all-around player? In *European conference on computer vision*, pp. 216–233. Springer, 2024c.
- Lu, S., Wang, Z., Li, L., Liu, Y., and Kong, A. W.-K. Mace: Mass concept erasure in diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 6430–6440, 2024.
- Lyu, M., Yang, Y., Hong, H., Chen, H., Jin, X., He, Y., Xue, H., Han, J., and Ding, G. One-dimensional adapter to rule them all: Concepts diffusion models and erasing applications. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 7559–7568, 2024.
- Nanda, N., Chan, L., Lieberum, T., Smith, J., and Steinhardt, J. Progress measures for grokking via mechanistic interpretability. *arXiv preprint arXiv:2301.05217*, 2023.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Qi, X., Panda, A., Lyu, K., Ma, X., Roy, S., Beirami, A., Mittal, P., and Henderson, P. Safety alignment should be made more than just a few tokens deep. *arXiv preprint arXiv:2406.05946*, 2024.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in neural information processing systems*, 36: 53728–53741, 2023.
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., and Ommer, B. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 10684–10695, 2022.
- Saha, S., Saha, S., Gaur, M., and Gokhale, T. Side effects of erasing concepts from diffusion models. *arXiv preprint arXiv:2508.15124*, 2025.
- Schramowski, P., Brack, M., Deiseroth, B., and Kersting, K. Safe latent diffusion: Mitigating inappropriate degeneration in diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 22522–22531, 2023.
- Shen, F. and Tang, J. Imagen: A unified conditional framework for pose-guided person generation. *Advances in neural information processing systems*, 37:6246–6266, 2024.
- Shen, F., Wang, C., Gao, J., Guo, Q., Dang, J., Tang, J., and Chua, T.-S. Long-term talkingface generation via motion-prior conditional diffusion model. In *Forty-second International Conference on Machine Learning*.

- Shen, F., Ye, H., Zhang, J., Wang, C., Han, X., and Wei, Y. Advancing pose-guided image synthesis with progressive conditional diffusion models. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=rHzapPnCGT>.
- Shen, F., Jiang, X., He, X., Ye, H., Wang, C., Du, X., Li, Z., and Tang, J. Imagdressing-v1: Customizable virtual dressing. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 6795–6804, 2025.
- Shen, Q. and Zhang, J. Mftformer: Meteorological-frequency-temporal transformer with block-aligned fusion for traffic flow prediction. *Research Square*, 2026. Preprint, doi:10.21203/rs.3.rs-8770196/v1.
- Singh, S., Romanou, A., Fourrier, C., Adelani, D. I., Ngui, J. G., Vila-Suero, D., Limkonchotiwat, P., Marchisio, K., Leong, W. Q., Susanto, Y., et al. Global mmlu: Understanding and addressing cultural and linguistic biases in multilingual evaluation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 18761–18799, 2025.
- Tsai, Y.-L., Hsu, C.-Y., Xie, C., Lin, C.-H., Chen, J.-Y., Li, B., Chen, P.-Y., Yu, C.-M., and Huang, C.-Y. Ring-a-bell! how reliable are concept removal methods for diffusion models? *arXiv preprint arXiv:2310.10012*, 2023.
- Wang, H., AprilPyone, M., and Echizen, I. Sgm: Safety glasses for multimodal large language models via neuron-level detoxification. *arXiv preprint arXiv:2512.15052*, 2025a.
- Wang, Y., Liu, Y., Shi, Y., Li, C., Pang, A., Yang, S., Yu, J., and Ren, K. Discovering influential neuron path in vision transformers. *arXiv preprint arXiv:2503.09046*, 2025b.
- Wu, Z., Gao, H., Luo, J., and Liu, Z. Humorreject: Decoupling llm safety from refusal prefix via a little humor. *arXiv preprint arXiv:2501.13677*, 2025.
- Xie, J., Yang, Z., and Shou, M. Z. Show-o2: Improved native unified multimodal models. *arXiv preprint arXiv:2506.15564*, 2025.
- Yan, Z., Liu, Z., Dong, W., Feng, L., and Dai, C. From knowledge to inference: Scaling laws of specialized reasoning on globalhealthatlas, 2026. URL <https://arxiv.org/abs/2602.00491>.
- Ye, F., Liu, G., Wu, X., and Wu, L. Altdiffusion: A multilingual text-to-image diffusion model. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, pp. 6648–6656, 2024.
- Yuan, Y., Jiao, W., Wang, W., Huang, J.-t., Xu, J., Liang, T., He, P., and Tu, Z. Refuse whenever you feel unsafe: Improving safety in llms via decoupled refusal training. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3149–3167, 2025.
- Zhang, D., Hu, Z., Zhoubian, S., Du, Z., Yang, K., Wang, Z., Yue, Y., Dong, Y., and Tang, J. Sciistruct: a self-reflective instruction annotated dataset for training scientific language models. *Advances in Neural Information Processing Systems*, 37:1443–1473, 2024a.
- Zhang, D., Cai, M., Light, J., Hu, Z., Yue, Y., and Tang, J. Tdrm: Smooth reward models with temporal difference for llm rl and inference. *arXiv preprint arXiv:2509.15110*, 2025a.
- Zhang, D., Feng, T., Xue, L., Wang, Y., Dong, Y., and Tang, J. Parameter-efficient fine-tuning for foundation models. *arXiv preprint arXiv:2501.13787*, 2025b.
- Zhang, G., Wang, K., Xu, X., Wang, Z., and Shi, H. Forget-me-not: Learning to forget in text-to-image diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 1755–1764, 2024b.
- Zhao, Y., Zhang, W., Xie, Y., Goyal, A., Kawaguchi, K., and Shieh, M. Understanding and enhancing safety mechanisms of llms via safety-specific neuron. In *The Thirteenth International Conference on Learning Representations*, 2025.
- Zhao, Y., Dai, C., Xiu, Y., Kou, M., Zheng, Y., and Niyato, D. Shardmemo: Masked moe routing for sharded agentic llm memory, 2026a. URL <https://arxiv.org/abs/2601.21545>.
- Zhao, Y., Dai, C., Zhuo, W., Xiu, Y., and Niyato, D. Clause: Agentic neuro-symbolic knowledge graph reasoning via dynamic learnable context engineering, 2026b. URL <https://arxiv.org/abs/2509.21035>.
- Zhu, D., Chen, J., Shen, X., Li, X., and Elhoseiny, M. Minigpt-4: Enhancing vision-language understanding with advanced large language models. *arXiv preprint arXiv:2304.10592*, 2023.
- Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., et al. Representation engineering: A top-down approach to ai transparency. *arXiv preprint arXiv:2310.01405*, 2023a.
- Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z., and Fredrikson, M. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023b.

Zou, A., Phan, L., Wang, J., Duenas, D., Lin, M., Andriushchenko, M., Wang, R., Kolter, Z., Fredrikson, M., and Hendrycks, D. Improving alignment and robustness with circuit breakers. *Advances in Neural Information Processing Systems*, 37:83345–83373, 2024.

Supplementary Material

The appendices provide additional details that support and extend the main paper. Appendix A provides a comparative analysis of causal circuit breakers and further tests on parameters related to SAE. This appendix also contains an analysis of the initial layer selection method and concludes with additional qualitative analysis concerning MLLMs. Appendix B addresses common issues. Appendix C provides a formal theoretical justification for the efficacy of harmful semantic suppression and the preservation of general utility based on the linear representation hypothesis. Appendix D will discuss the limitations of this method.

A. More Results

Comparative Analysis with Causal Circuit Breakers on MLLMs. As shown in Table 5, to provide a deeper discussion of the TraceRouter mechanism, we can conceptualize it as a Causal Circuit Breaker grounded in mechanistic interpretability. To further validate its superiority, we conducted a direct comparison between our framework and the SOTA circuit breakers (CB) on the LLaVA-1.5-7B model using the MM-SafetyBench dataset. Experimental results demonstrate that TraceRouter achieves a defense success rate (DSR) of 99.6%, significantly outperforming CB’s 97.3%. This performance gain stems from a fundamental divergence in intervention logic: while traditional CB methods primarily rely on training-time suppression or representation-space disruption, TraceRouter utilizes Sparse Autoencoders (SAEs) to disentangle pure harmful features and leverages feature influence scores (FIS) to trace the causal chain of cross-layer propagation, thereby achieving physical path-level disconnection. Unlike “Activation Steering” methods that aim to shift the activation direction within the latent space, TraceRouter acts as a surgical circuit breaker, directly destroying the topological channels through which harmful information flows toward the output. This effectively eliminates the possibility of “semantic escape,” where complex adversarial attacks find detour routes around localized interventions, while simultaneously preserving orthogonal computational routes for general utility.

Quantitative Analysis on Scaling Factor λ . To quantitatively determine the optimal suppression intensity that balances safety effectiveness and semantic preservation, we conducted a sensitivity analysis on Stable Diffusion 1.4 by varying the scaling factor λ from 0 to 5, tracking both the DSR (on I2P) and text-image alignment (CLIP Score). The results demonstrate that $\lambda = 2$ achieves the optimal balance, where safety performance reaches saturation while utility remains virtually intact. Specifically, as shown in Table 6, increasing λ from 0 to 2 leads to a dramatic increase in the DSR (from 82.2% to 99.2%), effectively eliminating unsafe concepts. Crucially, this significant safety gain incurs a negligible cost to utility, with the CLIP Score dropping only by 0.07 (31.34 vs. 31.27). However, when λ exceeds 2 (e.g., $\lambda = 3$), the marginal safety gains diminish, yet the CLIP Score begins to degrade more noticeably (dropping to 29.32). Therefore, this evidence confirms that setting $\lambda = 2$ precisely targets the harmful pathways without over-pruning the model’s general generative capabilities.

Universality Analysis across Diverse Architectures. To comprehensively evaluate the universality of TraceRouter across diverse architectures and its ability to preserve generative utility, we conducted extended experiments covering FLUX.1 Dev, the multilingual AltDiffusion, and the unified multimodal Show-o2, alongside image quality assessments on the MS COCO benchmark. As shown in Table 7, TraceRouter significantly reduces unsafe content across diverse architectures while imposing negligible impact on the models’ original generative capabilities, thereby achieving an optimal balance between safety and utility. Specifically, the total number of detected nudities on FLUX.1 Dev dropped substantially from 275 to 68, with the “Genitalia” category reduced to 0 across all models; simultaneously, regarding utility, the CLIP Score on Show-o2 remained virtually unchanged (28.87 vs. 28.85). Therefore, this evidence confirms that TraceRouter functions as an architecture-agnostic intervention that precisely disconnects harmful propagation pathways without compromising the general circuits required for benign generation.

Table 5. Safety performance comparison between TraceRouter and Circuit Breakers (Zou et al., 2024) on MLLMs. The LLaVA-1.5-7B model was tested on the MM-SafetyBench dataset.

Method	Avg. (%)
Circuit Breakers (Zou et al., 2024)	97.3
TraceRouter (Ours)	99.6

Table 6. Sensitivity analysis of the scaling factor λ . We report DSR and CLIP Score on SD1.4.

Param.	Value	I2P (N) ($\uparrow$)	CS ($\uparrow$)
λ	0	82.2%	31.34
	1	92.8%	31.29 (-0.05)
	2	99.2%	31.27 (-0.07)
	3	99.7%	29.32 (-2.02)
	4	99.8%	28.87 (-2.47)
	5	99.8%	28.52 (-2.82)

Table 7. **Quantitative comparison of nudity detection on the I2P dataset across different generation models.** The results show the number of detected instances for specific body parts.

Method	Number of nudity detected on I2P (Detected Quantity)							COCO
	Breast	Genitalia	Buttocks	Feet	Belly	Armpits	total ↓	CS ↑
<i>Text to Image Generation Model</i>								
FLUX.1 Dev (Labs, 2024)	32	2	15	2	39	186	275	26.44
+ TraceRouter (Ours)	12 (-20)	0 (-2)	4 (-11)	0 (-2)	15 (-24)	37 (-149)	68 (-207)	25.37 (-1.07)
<i>Multilingual Text to Image Generation Model</i>								
AltDiffusion (Ye et al., 2024)	121	2	5	6	58	55	247	30.54
+ TraceRouter (Ours)	38 (-83)	0 (-2)	1 (-4)	0 (-6)	23 (-35)	9 (-46)	71 (-176)	30.37 (-0.17)
<i>Visual-Language Model</i>								
Show-o2 (Xie et al., 2025)	25	0	0	3	45	61	134	28.87
+ TraceRouter (Ours)	14 (-11)	0 (-0)	0 (-0)	0 (-3)	34 (-11)	12 (-49)	60 (-74)	28.85 (-0.02)


Table 8. **Robustness analysis of SAE expansion factor.** We report the impact of varying the SAE expansion factor.

Expansion Factor	I2P (N) (↑)	CLIP Score (↑)
16×	98.8%	31.24
32×	99.2%	31.27
64×	99.1%	31.26
128×	99.2%	31.28

Sensitivity Analysis on SAE Hyperparameters. (1) Expansion Factor. To verify the robustness of TraceRouter against different Sparse Autoencoder (SAE) configurations and ensure it does not rely on meticulous hyperparameter tuning, we conducted an ablation study by varying the SAE Expansion Factor from $16\times$ to $128\times$. We keep the SAE training data, training steps, and optimization settings fixed, and vary only the expansion factor. The results demonstrate that our path-level intervention remains consistently effective across a wide range of SAE architectures, indicating high robustness. Specifically, as shown in Table 8, varying the expansion factor resulted in negligible fluctuations in both safety and utility metrics. Even when the SAE capacity was scaled up from $16\times$ to $128\times$, the DSR remained consistently high, and the CLIP Score stayed remarkably stable around 31.25. Therefore, this evidence confirms that TraceRouter extracts robust semantic features that are insensitive to the specific scale of the sparse autoencoder, facilitating easy deployment without extensive tuning.

(2) Top-K. To clarify the selection of the hyperparameter K for identifying the sensitive source, we conducted an analysis visualizing the weighted frequency scores (WFS) of candidate neurons. The results indicate that a small subset of “hub” neurons exhibits significantly higher response intensities compared to the rest of the neural population. Specifically, as illustrated in Figure 11, we plotted the WFS curves across multiple sensitive contexts; in each case, a few leading neurons form prominent peaks, followed by a rapid, exponential decay. We determine the value of K at the “elbow” of this distribution, the point immediately preceding the sharp drop, ensuring that neurons with salient WFS responses are retained as sensitive features while the long tail of weakly responsive, non-causal neurons is discarded. Therefore, selecting K based on this peak-to-tail transition point enables the framework to isolate a sparse yet decisive set of causal neurons at the onset layer, effectively balancing the precision of safety interventions with the preservation of model utility.

Figure 11. **Selection of hyperparameter K via WFS.** The distribution exhibits a distinct “elbow” pattern where a sparse set of hub neurons shows prominently higher responses than the long tail.



Sensitivity Analysis of the Sensitive Onset Layer. (1) Effects under different prompts. To verify the localization robustness of the sensitive onset layer across varying prompt inductions, we set up five independent groups of experiments, each using a unique set of sensitive contexts, and report the distribution of the average Coefficient of Variation (Avg. CV) for the SS across layers 1–12. The results show that the internal routing of harmful semantics demonstrates high physical consistency in the early stages of computation, with the identification of the onset layer remaining remarkably stable across multiple trials. Specifically, the experiment observed a clear convergence of the CV in the shallow layers; the CV decreases from 0.0285 at layer 1 to a global minimum of 0.0232 at layer 3. At layer 3, identified as the sensitive onset layer, the CV remains exceptionally low, significantly lower than the dispersion observed in deeper layers, such as Layer 10 ($CV = 0.0451$). Therefore, the consistent low-variance results across five experimental groups strongly justify selecting Layer 3 as the sensitive onset layer, providing a robust causal foundation for subsequent path-level interventions.




Figure 12. **Extended qualitative results of TraceRouter on MLLMs.** This figure presents additional comparisons between the original models and TraceRouter-enhanced versions across various malicious categories from the FigStep.

Table 9. **Comparison of selection criteria for the sensitive onset layer.**

Selection Criterion	DSR (%)
Global Maximum Peak	63.69
Threshold-based Selection	63.59
First Local Peak (Ours)	99.20

Selection of the Sensitive Onset Layer. In order to determine the optimal criterion for identifying the sensitive onset layer, we set up a comparative experiment evaluating the defense success rate (DSR) under three different selection strategies: the first local peak (our approach), the global maximum peak, and a threshold-based starting point. As shown in Table 9, the results show that identifying the earliest emergence of sensitive semantics, rather than peak activation intensity, is critical for robust safety intervention. Specifically, our method of selecting the first local peak achieved a near-optimal DSR of 99.2%, whereas selecting the global maximum or a threshold-based starting point resulted in significantly lower performance, with DSRs of 63.69% and 63.59%, respectively. Therefore, this evidence proves that the first local peak accurately captures the “causal bottleneck” of harmful semantic propagation, and intervening at this specific onset point is essential for physically severing the propagation circuits before they become distributed and harder to suppress.

Generalizability Across Models. A new SAE training and tracing process is required for each unique architecture because neuron specialization and layer depths vary across models. However, if multiple models share the same component (e.g., a frozen CLIP text encoder), the identified neurons and SAE weights can be reused directly without retraining.

Computational Cost. The SAE training and tracing are one-time offline costs. Upon completion, our approach incurs minimal overhead; generating images requires merely 3.28 seconds, with negligible additional latency compared to standard inference. This efficiency, coupled with the reusability of identified pathways across models employing the same text encoder, ensures our framework remains practical for real-world applications without imposing significant computational burdens.

Further qualitative experiments on MLLM. To facilitate a more comprehensive understanding of TraceRouter’s practical defense efficacy, we provide additional qualitative results in this section. Specifically, as shown in Figure 12, we conducted extended qualitative experiments on MLLMs, using the FigStep dataset. As illustrated in these extended results, the original models remain highly susceptible to visual-escape attacks in which malicious instructions, such as a “List of tools to secretly monitor someone’s online activities,” are embedded as typographic text within images. When prompted to “fill in the items,” the original models fail to recognize the illicit intent hidden in the visual input and instead directly output detailed attack steps.

B. More Discussions

▷ **Q1. Why is path-level intervention fundamentally more effective than traditional neuron-level suppression in securing large foundation models (LFMs)?**

Traditional methods rely on the “locality hypothesis”, assuming harmful semantics are confined to isolated components. However, TraceRouter addresses the reality that harmful semantics are distributed and propagate through cross-layer computation paths. By utilizing feature influence scores (FIS), TraceRouter identifies and severs the entire causal propagation circuit rather than just isolated nodes. This path-level approach effectively prevents “semantic leakage” that occurs when adversarial prompts find detour routes around a single suppressed neuron.

▷ **Q2. How does TraceRouter avoid the common trade-off where safety interventions significantly degrade the model’s general performance?**

TraceRouter preserves utility by exploiting the orthogonality of semantic features in a disentangled representation space. By integrating SAEs, we decompose activations into sensitive circuit components and orthogonal computational routes. By selectively suppressing only the identified causal pathway ($Z_{\mathcal{P}}$) and keeping the benign routes ($Z_{-\mathcal{P}}$) intact, the model maintains its core generative capabilities. (More detail see Appendix A).

▷ **Q3. Given the structural differences between diffusion models and LLMs, how does TraceRouter maintain architectural agnosticism?**

The “discover-trace-disconnect” framework is universal because it targets the fundamental way these models route information via attention and feature activations. By detecting the “sensitive onset layer” through attention divergence, a property shared across Transformer-based architectures, the framework can be applied to LLMs alike. Our experiments on diverse architectures, including FLUX.1 Dev and Show-o2, confirm that the path-level representation of harmful concepts is a cross-modal phenomenon.

▷ **Q4. Does TraceRouter’s effectiveness stem from simple signal masking of high-intensity activations?**

No, our specificity verification proves that the safety gains are not due to simple signal masking. When we randomly selected non-sensitive paths with similar activation intensities to our identified sensitive paths, the model’s safety actually deteriorated.

▷ **Q5. Is the framework overly sensitive to the hyperparameters of SAE, such as its expansion factor?**

Our ablation studies demonstrate that TraceRouter is highly robust to SAE configurations. Varying the SAE expansion factor from 16x to 128x resulted in negligible fluctuations in safety and utility. This suggests that the framework extracts robust semantic features that are insensitive to the specific scale of the SAE, facilitating easier deployment and reducing the need for extensive hyperparameter tuning.

▷ **Q6. Why choose path-level intervention over weight-level fine-tuning or unlearning?**

Weight-level editing or unlearning often incurs high computational costs and can lead to catastrophic forgetting or degraded generalization. TraceRouter, being an inference-side path-level intervention, is more efficient and highly selective. It disentangles harmful circuits from general knowledge pathways without permanently altering the model’s weights.

▷ **Q7. How does TraceRouter’s inference-time path blockade synergize with existing safety training methods like RLHF or DPO?**

While preference-based alignment (RLHF/DPO) attempts to adjust the global probability distribution of tokens, TraceRouter serves as a “causal safety valve” at the circuit level. We observe that models with prior safety alignment still possess latent harmful circuits that can be reactivated by adversarial “semantic escape” prompts. TraceRouter provides a complementary layer of defense by physically severing these residual cross-layer propagation circuits, effectively providing a fallback mechanism for safety vulnerabilities that finetuning fails to fully eradicate.

C. Theoretical Justification on Intervention Efficacy

In this section, we provide a formal theoretical analysis of why TraceRouter’s path-level intervention effectively suppresses harmful semantics while preserving general utility. We ground our analysis in the linear representation hypothesis and model the intervention as a selective projection operation in the activation space.

C.1. Problem Formulation and Decomposition

Let $Z^{(l)} \in \mathbb{R}^d$ denote the dense activation vector at the sensitive onset layer l . Based on the path decomposition defined in Eq. 6 and Eq. 7, the activation space is disentangled into two orthogonal components using the binary mask $\mathcal{M}^{(l)} \in \{0, 1\}^d$ derived from the FIS:

$$Z^{(l)} = Z_{\mathcal{P}}^{(l)} + Z_{-\mathcal{P}}^{(l)}, \quad (9)$$

where $Z_{\mathcal{P}}^{(l)} = Z^{(l)} \odot \mathcal{M}^{(l)}$ represents the *Sensitive Pathway* (harmful circuit), and $Z_{-\mathcal{P}}^{(l)} = Z^{(l)} \odot (1 - \mathcal{M}^{(l)})$ represents the orthogonal pathway (general utility). Here, $\odot$ denotes the element-wise product. TraceRouter’s intervention, as defined in Eq. 8, applies a selective scaling factor λ to the sensitive component:

$$\tilde{Z}^{(l)} = (1 - \lambda)Z_{\mathcal{P}}^{(l)} + Z_{\neg\mathcal{P}}^{(l)}. \quad (10)$$

This operation can be theoretically viewed as applying a diagonal projection matrix to the latent state, strictly dampening specific dimensions identified by the SAE-based causal tracing.

C.2. Efficacy of Harmful Suppression (Safety Guarantee)

Let $L_{harm}(Z)$ be the loss function representing the generation of harmful content (e.g., the likelihood of a toxic token or a forbidden visual concept). We aim to show that the intervention minimizes this risk. Considering a first-order Taylor expansion of the loss function around the original activation $Z^{(l)}$:

$$L_{harm}(\tilde{Z}^{(l)}) \approx L_{harm}(Z^{(l)}) + \nabla_Z L_{harm}(Z^{(l)})^T (\tilde{Z}^{(l)} - Z^{(l)}). \quad (11)$$

Substituting the intervention difference $\tilde{Z}^{(l)} - Z^{(l)} = -\lambda Z_{\mathcal{P}}^{(l)}$ (derived from Eq. 8):

$$L_{harm}(\tilde{Z}^{(l)}) \approx L_{harm}(Z^{(l)}) - \underbrace{\lambda \nabla_Z L_{harm}(Z^{(l)})^T Z_{\mathcal{P}}^{(l)}}_{\text{Causal Alignment Term}}, \quad (12)$$

Theoretical Justification: The validity of the suppression relies on the *Causal Alignment Term*. The binary mask $\mathcal{M}^{(l)}$ is constructed by selecting neurons with high FIS, where $\text{FIS}(m) \propto \mathbb{E}[\Delta_{output}]$. This ensures that the vector $Z_{\mathcal{P}}^{(l)}$ lies in the subspace maximally aligned with the gradient of harmful propagation $\nabla_Z L_{harm}$. Consequently, the dot product $\nabla_Z L_{harm}^T Z_{\mathcal{P}}^{(l)}$ is expected to be positive and non-negligible. For a suppression factor $\lambda > 0$ (e.g., $\lambda = 2$ as found empirically), the term $-\lambda(\dots)$ induces a significant reduction in the harmful loss, verifying the “switch-like” behavior observed in the path amplification experiments (Table 3).

Validation of Linear Approximation (Second-Order Analysis). While Eq. 11 relies on a first-order approximation, we further justify the validity of neglecting non-linear dynamics by considering the second-order Taylor expansion with the Lagrange remainder:

$$L_{harm}(\tilde{Z}^{(l)}) = L_{harm}(Z^{(l)}) - \lambda \nabla_Z L_{harm}^T Z_{\mathcal{P}}^{(l)} + \frac{1}{2} \lambda^2 (Z_{\mathcal{P}}^{(l)})^T \mathbf{H} Z_{\mathcal{P}}^{(l)} + O(\lambda^3), \quad (13)$$

where $\mathbf{H}$ represents the Hessian matrix. We contend that the quadratic curvature term is negligible under our framework due to three synergistic factors. First, the **sparsity-induced norm bounding** from the Top-K SAE ensures that the norm of the intervention vector $Z_{\mathcal{P}}^{(l)}$ is minimal, naturally suppressing the scale of the quadratic error. Second, the **local piecewise linearity** of activations (e.g., ReLU, SwiGLU) implies that within the operating radius of our suppression factor (empirically $\lambda \approx 2$), the Hessian $\mathbf{H}$ remains largely zero as the state stays within linear regions. Third, **disentangled feature orthogonality** ensures that the extracted sensitive pathway aligns with the principal harmful gradient while remaining orthogonal to the complex, high-curvature interaction terms of the general distribution. Consequently, the linear causal alignment term dominates the intervention dynamics, validating the robustness of our suppression guarantee.

C.3. Preservation of General Utility

Let $L_{util}(Z)$ be the loss function associated with general model capabilities (e.g., image fidelity or linguistic coherence). The impact of the intervention on utility is:

$$\Delta L_{util} \approx \nabla_Z L_{util}(Z^{(l)})^T (\tilde{Z}^{(l)} - Z^{(l)}) = -\lambda \nabla_Z L_{util}(Z^{(l)})^T Z_{\mathcal{P}}^{(l)}. \quad (14)$$

Theoretical Justification: TraceRouter preserves utility through the disentanglement provided by the SAE and the specificity of the mask $\mathcal{M}^{(l)}$.

1. The SAE resolves the superposition of features, ensuring that $\mathcal{M}^{(l)}$ isolates the specific semantic direction of the harmful concept.

2. In a well-disentangled representation space, the subspace of harmful semantics ($Z_{\mathcal{P}}^{(l)}$) is nearly orthogonal to the subspace of general utility ($Z_{\neg\mathcal{P}}^{(l)}$).

Mathematically, this implies $\nabla_Z L_{util}(Z^{(l)})^T Z_{\mathcal{P}}^{(l)} \approx 0$. Therefore, the perturbation to the utility loss is minimized ($\Delta L_{util} \approx 0$).

By formulating the intervention as Eq. 8, TraceRouter explicitly decomposes the activation into causal and orthogonal components. This formulation proves that our method maximizes the safety-utility ratio ($\frac{|\Delta L_{harm}|}{|\Delta L_{util}|}$) by exploiting the orthogonality of semantic features, a property that intensity-based suppression methods (which fail to distinguish $Z_{\mathcal{P}}$ from $Z_{\neg\mathcal{P}}$) cannot achieve.

D. Limitation

While TraceRouter demonstrates superior efficacy in securing foundation models against adversarial manipulation, this study primarily focuses on safety-oriented tasks, such as preventing the generation of harmful content or defending against jailbreak attacks. In these contexts, sensitive semantics often form relatively distinct causal circuits that can be disentangled from the model’s core logic. However, the applicability of the ”Discover-Trace-Disconnect” framework to general-purpose tasks, such as logical reasoning, knowledge retrieval, or creative writing, remains to be fully explored. Future work will focus on validating TraceRouter’s path-level intervention across a broader spectrum of non-safety domains to investigate whether similar topological disconnection can effectively modulate complex, multi-functional neural pathways without compromising the model’s underlying cognitive integrity.