

InfiMed-ORBIT: Aligning LLMs on Open-Ended Complex Tasks via Rubric-Based Incremental Training

Pengkai Wang^{*1,2} Pengwei Liu^{*3} Qi Zuo^{*4} Zhijie Sang² Congkai Xie² Hongxia Yang^{1,2}

Abstract

Reinforcement learning (RL) has powered many recent breakthroughs in large language models (LLMs), especially for tasks where rewards can be computed automatically, such as code generation. However, it is less effective in open-ended medical dialogue, where feedback is ambiguous, context-dependent, and difficult to simply summarize into a single scalar signal—often requiring heavily supervised reward models and creating risks of reward hacking. Thus, we introduce **ORBIT**, an **open-ended rubric-based incremental training** framework tailored for critical medical dialogues. ORBIT integrates medical dialogue construction with dynamically generated case-conditioned rubrics that serve as adaptive guides for incremental RL. Unlike approaches that rely on external medical knowledge bases or handcrafted rules, ORBIT uses rubric-guided evaluation and can be implemented with general-purpose instruction-following LLMs, avoiding task-specific judge fine-tuning. With only **2k** training samples, ORBIT raises Qwen3-4B-Instruct’s HealthBench-Hard score from **7.0** to **27.5**, achieving state-of-the-art performance among similarly sized open-source models while maintaining strong consultation quality as rubric coverage broadens. **Project page:** pidneuralode.github.io/ORBIT.

1. Introduction

As high-quality pretraining data become harder to expand and scaling brings smaller gains, progress is increasingly coming from post-training, where supervised tuning

^{*}Equal contribution ¹Department of Computing, The Hong Kong Polytechnic University, Hong Kong, China ²InfiX.ai ³Department of Control Science and Engineering, Zhejiang University ⁴Ant Group. Correspondence to: Hongxia Yang <hongxia.yang@polyu.edu.hk>.




Figure 1. From preference labels to clinical criteria. Open-ended medical responses require more than correctness; they demand explicit clinical reasoning—recognizing risks, eliciting missing information, managing uncertainty, and safely guiding the patient. Coarse feedback hides important clinical considerations behind unclear labels, whereas rubrics explicitly state criteria, ensuring transparent and clinically useful evaluations.

and preference optimization directly shape model behavior (Ouyang et al., 2022; Cui et al., 2023; Guo et al., 2025; Shao et al., 2024; Yu et al., 2026; Zheng et al., 2025). Modern post-training is dominated by two paradigms: Supervised Fine-Tuning (SFT) and Reinforcement Learning (RL) (Minaee et al., 2024; Touvron et al., 2023). SFT aligns the model with target formats and behaviors through token-level demonstrations, whereas RL optimizes policies against preference objectives—making reward specification the central bottleneck. When the correct outcome is clearly defined and easily verifiable, RL with verifiable rewards (RLVR) consistently improves performance (Shao et al., 2024). However, in open-ended fields such as medical consultation, psychology, and social interaction, response quality involves multiple factors like safety, empathy, and appropriateness. These factors vary significantly case by case, complicating efforts to define a single, precise reward (Tu et al., 2025; Sharma et al., 2023; Lambert et al., 2025). Recent work in these areas has therefore turned to learned Reward Models (RMs) (Kwon et al., 2023) trained on human preference data to approximate human judgment. Yet these RMs are costly to construct, highly domain-dependent, and fragile under distribution shift. Oversimplified reward models—such as binary ratings—are highly vulnerable to reward hacking, posing serious risks in critical medical dialogue.

These limitations motivate moving beyond opaque scalar rewards toward interpretable rubric-based assessment. Rubrics clearly outline specific evaluation criteria, allowing more transparent and precise feedback compared to

single numeric scores, as shown in Fig 1. In the medical domain, HealthBench (Arora et al., 2025) has significantly advanced evaluation methods using expert-designed rubrics to assess clinical reasoning. However, most existing medical LLMs (Chen et al., 2024; Huang et al., 2025; Wu et al., 2025) still struggle on the HealthBench-Hard benchmark, highlighting a persistent gap between QA-style optimization and realistic context-dependent consultation.

To address these limitations, we introduce **ORBIT**, an open-ended rubric-based incremental training framework for high-stakes medical dialogue. ORBIT starts with a small set of expert-written rubric seeds and uses retrieval techniques to create tailored rubrics for each medical case. These rubrics then help a general-purpose LLM judge evaluate model responses incrementally during reinforcement learning. When applied to the Qwen3-4B-Instruct model (Yang et al., 2025), ORBIT dramatically improves performance on HealthBench-Hard from 7.0 to 27.5 using only 2k training samples, achieving state-of-the-art performance among comparable-size open-source models. With larger rubric sets, ORBIT reaches 37.3, outperforming open-source models up to $8\times$ its size.

In summary, our contributions include:

- We propose **ORBIT**, a seed-based automated rubric-based post-training framework for scalable open-ended alignment of LLMs, which replaces opaque scalar rewards and costly reward models with interpretable, rubric-driven feedback tailored to high-stakes medical dialogue.
- We develop a context-aware rubric generation and training pipeline that integrates retrieval-augmented in-context prompting, multi-stage filtering, dynamic sampling, and incremental optimization. It enables stable, data-efficient rubric-guided RL by using existing expert-crafted rubrics as seeds, eliminating the need for new task-specific annotations or a separate medical reward model.
- We demonstrate through extensive experiments on HealthBench that ORBIT substantially improves medical consultation quality, achieving state-of-the-art performance among comparable-size open-source models with only 2k training samples. We further probe the source of these gains through judge-sensitivity and rubric-similarity analyzes, finding that ORBIT learns case-conditioned clinical criteria rather than relying on evaluator-specific preferences or reusable rubric patterns.

2. Related Work

Open-Ended Benchmarks. The evaluation of LLMs on open-ended generation tasks is increasingly shifting from conventional automatic metrics toward holistic, rubric-based frameworks. Early benchmarks focused on short-form grad-

ing or fixed-entity extraction, which struggled to capture the sophisticated, multidimensional behavior of modern LLMs. This limitation has led to a new generation of evaluation suites that adopt fine-grained multidimensional rubrics, such as HealthBench (Arora et al., 2025), PaperBench (Starace et al., 2025), WildBench (Lin et al., 2025), AMEGA (Fast et al., 2024) and MultiChallenge (Deshpande et al., 2025). Using thousands of scenario-specific criteria, these benchmarks assess model behavior more precisely and interpretably. In particular, HealthBench highlights that achieving strong performance in medical consultation is still challenging under rigorous rubric-based evaluation.

Reward Models in LLMs. RL has emerged as a critical step in LLM post-training, typically realized through reward models that align behavior with human intent. Early Reinforcement Learning from Human Feedback (RLHF) (Ouyang et al., 2022) used preference signals, offering effective but coarse supervision. Later approaches introduced rule-based rewards that decompose preferences into verifiable components such as structural correctness or format adherence (Chen et al., 2024; Zhang and Zhang, 2024), followed by semantic-level evaluators that assess factual accuracy, reasoning quality, and consistency (Bhaskar et al., 2025; Jayalath et al., 2025; Viswanathan et al., 2026; Jacob Dineen et al., 2025; Chen et al., 2025). In specialized domains like medicine, researchers have further explored domain-adaptive reward formulations and rubric-driven alignment strategies (Gunjal et al., 2025; Dou et al., 2025). Despite these advances, rewards remain costly to build, hard to transfer across domains, and difficult to interpret at the criterion level. Our approach instead uses rubric-guided RL with automatically generated context-sensitive criteria, enabling structured, transparent, and interpretable feedback.

LLMs for Health. The rapid advancement of LLMs has accelerated interest in their application in healthcare (Singhal et al., 2023; 2025; Tanno et al., 2025; Thirunavukarasu et al., 2023). Previous work has explored a range of use cases, including differential diagnosis support, clinical documentation, mental health assistance, and radiology reports (McDuff et al., 2025; Oh et al., 2024; Liu et al., 2025c). Despite promising results, most systems remain narrowly specialized and struggle to generalize to the heterogeneous, context-dependent reasoning of real-world clinical practice. As LLMs near clinical deployment, research increasingly focuses on agentic systems that integrate multi-step reasoning, tool use, and structured knowledge (Ferber et al., 2025; Lu et al., 2024; Tang et al., 2024).

3. ORBIT Framework

Problem Setup. Given a realistic clinical case, either in the chat format or in the outpatient chart format, ORBIT




Figure 2. Overview of ORBIT, a rubric-guided reinforcement learning framework for aligning LLMs on open-ended medical dialogue. (a) Construction of realistic multi-turn medical query contexts. (b) Case-specific rubrics guide GRPO-based policy optimization via external reward evaluation. (c) A retrieval-augmented rubric generator supports calibrated, criterion-level evaluations and scalable extensions to diverse clinical domains.

converts it into practical RL training data. Using multi-level filtering strategies, the system produces a final pair of $\langle \text{dialogue}, \text{rubrics} \rangle$ that is fed into RL to update the policy of the base LLM. The pipeline comprises three stages: dialogue QA simulation (§3.1), rubric generation with in-context learning (§3.2) and rubric-guided RL (§3.3).

3.1. Medical Dialogue Construction

Recent studies show that LLMs can generate clinically plausible multi-turn medical dialogues when conditioned on dialogue histories or outpatient clinical notes (Tu et al., 2025; Zhu et al., 2025; Liu et al., 2025d). Agentic synthesis workflows offer a scalable solution for generating multi-turn medical dialogue data (Li et al., 2024; Sun et al., 2025; Tang et al., 2024; Wei et al., 2024; Feng et al., 2026). To enable reproducible methodological validation, we adopt the 2k and 8k processed dialogue splits from DoctorAgent-RL/MTMedDialog (Feng et al., 2026), derived from IMCS21 (Chen et al., 2023), CHIP-MDCFNPC (Zhu et al., 2023), and MedDG (Liu et al., 2022). We further augment our large-scale training scenario with additional dialogues from ReMeDi (Yan et al., 2022). These datasets form the basis for the query items in our evaluation rubric.

We then employ in-context learning to synthesize rubrics from these dialogues, as detailed below.

3.2. Rubric Generator with In-Context Learning

Diagnostic Database Construction. Given a seed dataset $D = \{(q_i, \mathcal{R}_i)\}_{i=1}^N$ from HealthBench rubrics and an embedding model $\mathcal{M}_{emb}$, we construct a diagnostic database for case-rubric retrieval. Here, q_i denotes the dialogue history of the i -th consultation, and $\mathcal{R}_i = \{r_{i,1}, \dots, r_{i,n_i}\}$ is its rubric set. Each dialogue and rubric is embedded by $\mathbf{e}_{q_i} = \mathcal{M}_{emb}(q_i)$ and $\mathbf{e}_{r_{i,j}} = \mathcal{M}_{emb}(r_{i,j})$, and all embeddings are stored in a vector database. We then build two distinct data pools: (i) a *case-rubric pair pool* $\mathcal{P}_{cr} = \{(q_i, \mathcal{R}_i, \mathbf{e}_{q_i}, \sum_{r \in \mathcal{R}_i} \mathbf{e}_r) \mid (q_i, \mathcal{R}_i) \in D\}$, which keeps each case with its rubrics and aggregated rubric embedding; and (ii) a *rubric pool* $\mathcal{P}_r = \{(r, \mathbf{e}_r) \mid r \in \bigcup_{(q_i, \mathcal{R}_i) \in D} \mathcal{R}_i\}$, gathering all unique rubrics and their embeddings to facilitate fine-grained semantic retrieval.

Diagnostic Candidate Search. Given a new query q (the dialogue history of a consultation), we first obtain its embedding $\mathbf{e}_q = \mathcal{M}_{emb}(q)$, then compute similarity against all entries in the diagnostic database, i.e., the case-rubric pair

pool $\mathcal{P}_{cr}$ and the rubric pool $\mathcal{P}_r$. We retrieve the top- t_{cases} most similar cases from $\mathcal{P}_{cr}$ and the top- $t_{rubrics}$ rubric candidates from $\mathcal{P}_r$, subsequently employing a reranker $\mathcal{M}_{re}$ to enhance their relevance. This two-stage retrieval produces the final sets of relevant cases $\mathcal{C}_q$ and semantically aligned rubrics $\mathcal{R}_q$ for the given query q .

Rubric Generation. The retrieved cases $\mathcal{C}_q$ and rubrics $\mathcal{R}_q$ serve as in-context exemplars to guide rubric synthesis. A generative model $\mathcal{G}$ is prompted with $\mathcal{C}_q$, $\mathcal{R}_q$, and task-specific instructions to produce m_g rubric candidates for query q , $\mathcal{G}(q) = \{r_1, \dots, r_{m_g}\}$. The prompt template is shown in Fig. 8. The generated rubrics constitute a checklist for evaluating a model’s response to the medical query.

Difficulty Filtering with Pass@k. To focus RL on informative training signals, we apply a two-stage difficulty filter on both queries and rubrics using Pass@k-style statistics. For each query q , the current policy model $\mathcal{M}$ generates n_{rollout} responses $\mathcal{Y}_q = \{y_1, \dots, y_{n_{\text{rollout}}}\}$. A judge model evaluates each response–rubric pair (y_i, r) using a satisfaction metric $S(y_i, r)$ for all $r \in \mathcal{R}_q$. These scores are then used to define two complementary filtering mechanisms.

(1) Sample-Level Filtering: Retaining Learnable-but Challenged Queries. We first compute an average score for each query,

$$\bar{s}_q = \frac{1}{n_{\text{rollout}} \cdot |\mathcal{R}_q|} \sum_{i=1}^{n_{\text{rollout}}} \sum_{r \in \mathcal{R}_q} S(y_i, r), \quad (1)$$

which reflects the model’s overall performance on q : high values typically correspond to trivial cases, whereas extremely low values indicate queries that the current policy consistently fails to solve. We therefore retain only queries within an intermediate difficulty range, $\mathcal{Q}_{\text{filtered}} = \{q \mid \tau_q^{\text{low}} \leq \bar{s}_q \leq \tau_q^{\text{high}}\}$, thereby focusing training on samples that provide meaningful optimization signals.

(2) Rubric-Level Filtering: Removing Saturated Rubrics. We further refine the rubric set by removing rubrics that are already consistently satisfied by the current policy. For each rubric r , we define its empirical pass rate over the n_{rollout} sampled responses as

$$P(r, q) = \frac{1}{n_{\text{rollout}}} \sum_{i=1}^{n_{\text{rollout}}} \mathbb{I}[S(y_i, r) \geq \tau_s], \quad (2)$$

where a response is considered to *pass* rubric r when $S(y_i, r) \geq \tau_s$. Rubrics with excessively high pass rates contribute limited supervision value and are therefore removed using threshold τ_r : $\mathcal{R}_{q, \text{filtered}} = \{r \in \mathcal{R}_q \mid P(r, q) < \tau_r\}$.

Together, these two filtering stages retain solvable yet challenging queries and discriminative rubrics, resulting in denser and more informative supervision for the subsequent reinforcement learning stage.

3.3. Rubric-Guided Reinforcement Learning

In this stage, we optimize the policy π_θ using Group Relative Policy Optimization (GRPO) (Shao et al., 2024), which efficiently estimates the baseline via group-wise sampling. For each query q , we sample G rollouts $o_1, \dots, o_G$ and compute the advantage $\hat{A}_{i,t} = (R(q, o_i) - \bar{R}_G) / \sigma_G$ to update the policy, where $\bar{R}_G$ and σ_G denote the mean and standard deviation of the group.

Rubric-Aware Reward Modeling. Instead of sparse binary rewards such as exact answer matching, we leverage rubric-guided criteria to define dense, semantically rich rewards. Let $\mathcal{M}_{\text{judge}}$ be the judge model and $r_j = (\text{crit}_j, w_j) \in \mathcal{R}_q$ be a specific criterion with weight w_j . The reward for a rollout o_i is defined as a signed sum of criterion-level judge outcomes:

$$R(q, o_i) = \sum_{j=1}^{|\mathcal{R}_q|} \mathbb{I}[\mathcal{M}_{\text{judge}}(q, o_i, \text{crit}_j) \rightarrow \text{True}] \cdot w_j. \quad (3)$$

Positive weights reward desirable clinical behaviors; negative weights penalize unsafe or misleading ones. This formulation supports fine-grained credit assignment at the criterion-level, guiding the model toward complex medical reasoning beyond merely matching conclusions.

Training Stability Strategies. In practice, rubric-derived rewards often saturate quickly within rollout groups, producing identical scores and zero-variance advantages. Case-specific rubric targets may also exceed the base policy’s limited exploration capacity. Drawing on previous insights (Yu et al., 2026; Liu et al., 2025b; An et al., 2025), we introduce two mechanisms to mitigate these issues:

- **Variance-Aware Filtering of Rollouts.** Performing policy updates on rollout groups with near-zero reward variance ($\sigma_G \approx 0$) results in ill-defined advantage estimates and numerical instability. We proactively filter these uninformative rollouts via a dynamic binary mask M_q :

$$M_q = \mathbb{I} \left(\max_i R(q, o_i) - \min_i R(q, o_i) > \delta \right). \quad (4)$$

The loss is computed only on valid queries $\mathbb{E}_q[M_q \cdot \mathcal{L}_{\text{GRPO}}(q)]$, ensuring that updates are driven solely by discriminative signals.

- **Staged Entropy Injection.** Between training stages k and $k+1$, we re-initialize policy parameters θ_{k+1} from the best prior checkpoint θ_k^* and adjust the sampling temperature T upward to reinstate exploratory behavior:

$$T_{k+1} = \min(T_{\text{max}}, T_k \cdot \gamma), \quad \text{with } \gamma > 1. \quad (5)$$

This periodic entropy injection promotes exploration while preserving existing policy skills.

Table 1. Overall model performance on HealthBench-Hard

Models	By Theme							By Axis					Total Score
	Emergency referrals	Context seeking	Global health	Health data tasks	Communication	Hedging	Response depth	Accuracy	Completeness	Communicat. quality	Context awareness	Instruction following	
Proprietary Models													
GPT-4.1	20.5	12.3	12.1	9.7	14.9	12.3	17.5	30.5	0	70.6	0	60.5	13.2
GPT-5 (thinking)	-	-	-	-	-	-	-	-	-	-	-	-	46.2
Open-source Models (< 10B)													
Qwen3-4B-Instruct (base)	9.3	8.5	7.1	0	8.6	12.2	5.1	24.1	0.8	57.5	0	45.0	7.0
Qwen3-4B-Thinking	14.4	12.5	2.4	0	3.5	8.5	0	23.2	0	42.5	0	39.6	5.2
Qwen-2.5-7B-Instruct	0	0	0	0	0	0	0	6.4	0	45.2	0	33.7	0
InfiMed-ORBIT-4B (2k)	39.9	37.8	30.2	6.2	26.6	32.2	6.6	31.8	38.1	45.3	16.8	43.7	27.5
InfiMed-ORBIT-4B (8k)	44.6	49.1	34.6	9.3	28.0	42.6	10.4	33.6	42.0	46.7	32.3	50.5	33.6
InfiMed-ORBIT-4B (28k)	51.6	50.5	41.9	10.8	30.9	43.8	13.4	38.1	48.9	42.1	31.3	49.1	37.3
Open-source Models (> 10B)													
Qwen3-30B-Instruct	18.3	12.9	14.7	17.9	19.4	9.5	28.5	28.5	0	45.2	0	33.7	13.1
Qwen3-30B-A3B-Thinking	21.4	20.4	15.0	8.9	16.7	20.4	6.5	33.7	11.6	53.0	0	45.5	16.1
Baichuan-M2-32B	45.6	39.5	35.6	21.3	32.0	40.9	19.9	41.3	44.6	51.6	19.3	48.0	34.5




Figure 3. Performance comparison of ORBIT models across multiple clinical dimensions. Results are organized by clinical *Theme* and evaluation *Axis*. The results are divided into two groups: (a, c) Within-family comparisons among the instructor-tuned base model (Qwen3-4B-Instruct) and ORBIT variants trained at progressively larger scales (InfiMed-ORBIT-4B: 2k, 8k, and 28k samples). (b, d) Benchmark comparisons between the best-performing ORBIT model (InfiMed-ORBIT-4B, 28k samples) and substantially larger proprietary and open-source models, including Qwen3-30B-Instruct, GPT-4.1, Baichuan-M2-32B, and Qwen3-30B-A3B-Thinking. For conciseness, Health Data T.” denotes Health Data Tasks,” and Emerg. Referrals” denotes Emergency Referrals.”

4. Experiments

4.1. Experimental Setup

Medical Dialogue Datasets. To facilitate rubric synthesis for open-ended medical consultation, we curate dialogue data from multiple publicly available sources and organize them into three subsets with distinct experimental purposes.

- (i) *Core Experimental Dataset.* Our primary dataset contains 2,082 multi-turn medical consultations from the processed DoctorAgent-RL/MTMedDialog split (Feng et al., 2026), originally derived from IMCS21 (Chen et al., 2023), CHIP-MDCFNPC (Zhu et al., 2023), and MedDG (Liu et al., 2022).
- (ii) *Scalability Dataset.* To study how rubric diversity and coverage affect alignment, we use the corresponding larger DoctorAgent-RL/MTMedDialog split from the same source mixture, yielding roughly 8k curated dialogue samples. This subset is the primary testbed for rubric scaling experiments; preprocessing details are provided in the App. F.1.

- (iii) *Large-Scale Training Extension Dataset.* To investigate the large-scale effects of rubric-based alignment, we further augment our training corpus using the ReMeDi dataset (Yan et al., 2022), assembling 20k dialogues that span diverse clinical themes and patient intents.

HealthBench Seed Data. For HealthBench-Hard, we construct clinically grounded, case-specific rubrics through a carefully designed retrieval-augmented generation (RAG) pipeline. Specifically, our pipeline dynamically retrieves relevant clinical cases and contextual knowledge to produce instance-specific rubrics, in contrast to using fixed or generic templates (see the prompt in Fig. 7). To rigorously avoid data contamination, we strictly exclude all HealthBench-Hard samples from the rubric construction process. Only the non-Hard HealthBench-4k rubric subset (Arora et al., 2025), subjected to rigorous lexical and semantic filtering (App. H and App. H.2), serves as seed data, thus minimizing the risk of instance-level contamination and maintaining benchmark integrity.

Benchmark. HealthBench (Arora et al., 2025) is an open-

ended medical benchmark from OpenAI with 5k multi-turn medical consultations. It includes HealthBench-Hard, a highly challenging subset of 1k cases specifically curated to stress-test state-of-the-art general-purpose models. Consequently, we focus our core experiments on the HealthBench-Hard subset to rigorously evaluate our proposed method.

Baselines. We select Qwen3-4B-Instruct-2507 (Yang et al., 2025) as our base model to investigate rubric-guided RL in compact LLMs without requiring excessively large models. We employ Qwen3-30B-A3B-Instruct-2507 (Yang et al., 2025) (hereafter Qwen3-30B-Instruct) as the judge model in the core experiments to perform rubric evaluations within the ORBIT training pipeline.

Hardware settings. To enable fair comparisons, we maintain consistent batch sizes across all methods and tasks. All experiments are conducted on a cluster of eight NVIDIA H800 (80GB) GPUs. Four GPUs are used to train the primary models, while the remaining four run the evaluation model for rubric evaluation.

4.2. Quantitative Results

Tab. 1 summarizes our main experimental results. For primary comparisons, all HealthBench-Hard results reported in the main table are evaluated using GPT-4.1 (Achiam et al., 2023), following the official HealthBench protocol (Arora et al., 2025) and publicly available evaluation code. Local variants of GPT-OSS-120B (Agarwal et al., 2025) are used for diagnostics and ablations during model development; relevant sections explicitly indicate the evaluator. To enable a fair comparison under a fixed evaluation budget, we adopt the reported benchmark score of 46.2 for GPT-5 (thinking) (Dou et al., 2025). In the absence of dimension-specific scores, we mark the corresponding entries as a dash (–).

ORBIT significantly improves the Qwen3-4B-Instruct backbone, enabling it to outperform several larger open-source baselines under the same HealthBench-Hard protocol. Specifically, when trained with only 2k samples (Dataset *i*), InfiMed-ORBIT-4B improves performance from 7.0 to **27.5**, an absolute gain of **+20.5** points (293% relative improvement). Increasing the training data to 8k samples (Dataset *ii*) further elevates the score to **33.6**, setting a new state-of-the-art within the sub-10B parameter regime. Expanding training to 28k samples (Datasets *ii* and *iii*) increases the performance to **37.3**, surpassing the open-source baselines evaluated listed in Tab. 1.

As shown in Fig. 3, performance improvements span multiple clinical themes and evaluation axes. In particular, the largest improvements occur in *completeness* (+20.5) and *context awareness* (+11.8), accompanied by a slight decline in *communication quality* (–8.2) due to rubric-induced safety constraints; see App. F.6 for detailed per-axis results.

We further evaluate the statistical significance of ORBIT. In a 200-case validation subset that was left out (App. F.7), our method yields statistically significant improvements, confirming the method gains. Collectively, these results show that rubric-guided RL efficiently aligns compact medical LLMs for high-stakes clinical applications.

4.3. Ablation Experiments

We conduct a comprehensive ablation study to characterize the role of each ORBIT component. Specifically, we examine several key dimensions: (i) strategies for rubric generation and selection, (ii) the impact of evaluator choice on evaluation consistency and reliability, (iii) the influence of pass@*k* filtering mechanisms on both sample selection and rubric quality, and (iv) the effectiveness of dynamic sampling and multi-stage training strategies.

4.3.1. THE SELECTION OF EVALUATION MODEL AND RUBRIC GENERATION MODEL

OpenAI adopts GPT-4.1 as the primary evaluator (Arora et al., 2025). However, the substantial cost of API calls motivates the exploration of affordable open-source alternatives for methodology development and data construction. We conduct comparative experiments to identify suitable alternatives, with detailed results presented in App. B. We observed that GPT-OSS-120B (Agarwal et al., 2025) provides reliable evaluation performance, with score improvements strongly correlated with those obtained using GPT-4.1 (Achiam et al., 2023). Additionally, we explore suitable models for rubric generation within the Rubric-RAG pipeline. Our experiments show that the rubrics generated by DeepSeek-R1-0528 (Guo et al., 2025) provide consistent and substantial downstream improvements. Detailed results are provided in App. C. Consequently, all subsequent ablation studies utilize these validated settings.

4.3.2. EVALUATING DATA SCALABILITY AND TRAINING EFFICIENCY

To evaluate the robustness of ORBIT, we introduce two complementary evaluation tracks addressing key dimensions of data optimization: (i) **data scalability**, measuring performance as training data increase; and (ii) **training efficiency**, assessing information density via difficulty-based pruning of queries and rubrics.

Dimension 1: Data Scalability. To assess scalability, we progressively incorporate the *Scalability* (8k) and *Large-scale extension* (20k) datasets in Sec. 4.1. As shown in Tab. 1, InfiMed-ORBIT-4B improves monotonically as the data scale increases from 2k to 28k.

Dimension 2: Training Efficiency via Strategic Pruning. To investigate data efficiency, we perform controlled experi-

Table 2. Evaluation results of ORBIT models across different pass@k thresholds, assessed using the GPT-OSS-120B-middle model. Results are presented under two distinct filtering criteria: sample-based versus rubric-based filtering.

Models	By Theme							By Axis					Total Score
	Emergency referrals	Context seeking	Global health	Health data tasks	Communication	Hedging	Response depth	Accuracy	Completeness	Communicat. quality	Context awareness	Instruction following	
Base models													
Qwen3-4B-Instruct (Base)	6.6	10.4	8.3	0	9.1	12.6	0	19.7	3.5	57.7	0	47.9	7.2
InfiMed-ORBIT-4B (no Filter)	26.2	26.5	22.6	5.5	20.8	26.3	0.7	24.0	23.8	53.3	10.9	46.9	20.2
Pass@k for rubrics													
InfiMed-ORBIT-4B (0 ~ 0.75)	20.0	27.7	21.4	8.6	17.3	28.0	0.1	23.4	25.0	51.8	10.7	47.0	19.9
InfiMed-ORBIT-4B (0 ~ 0.50)	20.5	25.4	21.5	2.0	15.3	26.1	0	24.6	21.2	50.0	8.2	42.6	17.9
InfiMed-ORBIT-4B (0 ~ 0.25)	18.7	24.1	21.9	3.6	15.5	27.3	0	23.7	23.1	50.0	7.6	44.8	18.7
Pass@k for samples													
InfiMed-ORBIT-4B (0 ~ 0.75)	21.7	25.3	23.3	5.2	16.1	29.8	0	24.5	22.3	48.7	10.4	48.5	19.7
InfiMed-ORBIT-4B (0 ~ 0.50)	18.2	19.7	18.2	2.2	14.0	17.9	0	22.2	15.4	51.6	2.7	45.7	14.5
Multi-stage Restart Training													
InfiMed-ORBIT-4B (8k data)	28.0	35.8	31.5	7.5	21.9	32.6	1.2	25.8	34.8	44.7	18.5	45.2	25.9
InfiMed-ORBIT-4B (restart)	29.6	39.7	33.1	8.3	23.1	32.9	0.2	25.6	37.1	42.5	20.5	46.0	27.3




Figure 4. Computational Efficiency Gains through Controlled Filtering. We compare strict and moderate filtering regimes against a no-filter baseline. Panels (left to right) display: response length evolution and final distribution, showing that stricter filtering curbs token growth and induces conciseness; and training runtime trajectory and total cost, demonstrating that sample filtering lowers computational overhead. Thresholding provides a tunable parameter to balance computational budget, output length, and downstream performance.




Figure 5. Distributional comparison of inference scaling versus rubric-guided RL. We contrast off-the-shelf Qwen3-4B-Instruct (inference scaling, $K = 8, 40$) with InfiMed-ORBIT-4B (rubric-guided RL, $K = 40$). Panels (from left) show kernel density and histogram plots of: (i) average normalized query scores; (ii) best-of- K normalized scores; (iii) rubric pass rates (mean criterion compliance); and (iv) rubric coverage (probability that each criterion is satisfied at least once).

ments on the *Core Experimental Dataset* with substantial reductions in training data. We apply a difficulty-aware, pass-rate-based filtering strategy to eliminate easy or redundant samples, optimizing the trade-off between computation and performance.

Our analysis begins with 2,082 dialogue queries paired with 25,020 corresponding rubrics (see Sec. 3.1). For each query, the base model generates eight candidate responses, each evaluated by Qwen3-30B-Instruct following the evaluation

pipeline used in ORBIT training. Based on these pass rates, we explore two complementary filtering approaches:

- **Sample-Level Filtering (Hard Sample Mining):** We filter queries based on their aggregate pass rates, retaining only those that challenge the current policy. Starting with 2,082 samples, we evaluate two subsets: a *moderate* set ($\bar{s}_q \in [0, 0.75]$, 1,403 samples) and a *strict* set ($\bar{s}_q \in [0, 0.5]$, 701 samples).
- **Rubric-Level Filtering (Constraint Optimization):** We

filter rubrics based on their global pass rates to remove trivial constraints. Applying thresholds of $[0, 0.25]$, $[0, 0.5]$, and $[0, 0.75]$ reduces the rubric set from 25,020 to 10,055, 12,352, and 14,411, respectively. This aims to minimize the computational overhead of reward calculation while preserving alignment effectiveness.

The baseline model (without filtering) is trained for 320 RL steps in a single-stage training procedure. Filtered configurations ($[0, 0.5]$ and $[0, 0.75]$) are evaluated at 110 and 220 steps, respectively, ensuring that each filtered subset is trained for exactly ten epochs. As illustrated in Tab. 2 and Fig. 4, selective filtering improves the trade-off between computational cost and response length. Moderate filtering effectively preserves downstream performance while reducing runtime and response length, whereas overly strict filtering reveals a clear exploration–efficiency trade-off.

However, two limitations remain: limited rollouts ($K = 8$) can inadequately explore the policy space, and too strict thresholds can prematurely restrict policy exploration. Rubric-level filtering consistently provides performance improvements, reducing evaluation latency without sacrificing alignment quality. Overall, difficulty-aware filtering at both the sample and rubric levels offers a straightforward strategy for accelerating reinforcement learning–based alignment, explicitly highlighting the exploration–efficiency trade-off.

4.4. Expanding Capability Boundaries through Rubric-Guided RL

To distinguish between the effects of increased inference efforts and improved training alignment, we analyze the distribution of model performance, as shown in Fig. 5. We compare the baseline Qwen3-4B-Instruct-2507 model under two inference settings ($K = 8$, $K = 40$ rollouts) against the rubric-aligned InfiMed-ORBIT-4B.

The Ceiling of Inference Scaling. As shown in the density plots (Fig. 5, blue vs. green), increasing the sampling budget for the baseline model from $K = 8$ to $K = 40$ yields only marginal gains. Although the mean of the score distribution exhibits a slight rightward shift, the overall distribution remains approximately Gaussian and is centered at moderate scores (~ 0.65). This behavior indicates a clear *capability ceiling*: the baseline policy distribution is inadequately aligned with rubric constraints, showing that brute-force inference scaling alone is insufficient to materially improve rubric compliance.

Rubric-Guided RL Induces a Structural Distributional Shift. In contrast, ORBIT (purple) produces a fundamental shift in the shape of the performance distribution. Instead of simply moving the mean performance, the rubric-guided RL clusters the model’s responses in high-score regions, creating a strongly left-skewed distribution. This structural

effect is most pronounced in the *pass rate per criterion* and *hit rate* panels, where ORBIT exhibits a dominant mode at a perfect score (1.0). Collectively, these results show that ORBIT improves alignment by reshaping the policy distribution rather than relying on sampling variance. As a result, it achieves performance improvements that are not recovered by inference scaling alone. Specifically, the baseline’s best-of-40 distribution peaks near an average score of approximately 0.65, whereas ORBIT shifts this mode beyond 0.85. Similarly, per-rubric hit rates—the fraction of rubrics satisfied at least once across 40 rollouts—significantly increase. This supports the view that ORBIT expands the set of achievable rubric-satisfying clinical behaviors rather than merely reweighting existing policy output (see App. G for examples).

4.5. Judge Reliability and Rubric Quality

Evaluator Reliability through Panel Consistency. We evaluated 200 cases using rubrics generated through retrieval-augmented methods with a six-judge LLM panel (GPT-4.1, GPT-OSS-120B/20B, Qwen3-30B/8B/4B). The panel exhibits strong collective consistency, reflected by an item-level Krippendorff’s $\alpha = 0.999$, mean pairwise Cohen’s $\kappa = 0.439$, and a case-level ICC(A, k) = 0.881, despite moderate individual judge variability. Detailed judge-reference agreements are provided in Fig. 6a. Among individual judges, GPT-OSS-120B shows the highest agreement with GPT-4.1 (Cohen’s $\kappa = 0.653$, $\rho = 0.837$). This justifies our choice of GPT-OSS-120B for evaluations during model development, while reserving GPT-4.1 exclusively for final benchmark assessments (details in App. B.1). The extremely high Krippendorff’s α is mainly due to heavily imbalanced binary rubric distributions; thus, Cohen’s κ provides a more conservative measure of reliability at the individual item level.

Physician Reference for Clinical Relevance. To complement the LLM-only consistency analysis, we obtained blinded physician annotations for a subset of 60 cases from the original 200-case panel, deliberately selecting high-disagreement cases. As illustrated in Fig 6a and Fig 6d, the physician rated 90.7% of the rubric-generated criteria as clinically meaningful. Moreover, GPT-OSS-120B demonstrated substantial concordance with these physician annotations, achieving a 79.6% item-level agreement, Cohen’s $\kappa = 0.578$, and a high case-level rank correlation of $\rho = 0.888$.

Impact of Training-Time Judge Selection. With GPT-OSS-120B fixed as the downstream evaluator, we systematically varied the training-time judges across multiple model families and sizes. As shown in Fig. 6b, stronger judges consistently produce improved ORBIT policy performance. For example, Qwen3-4B self-judging improved scores from




Figure 6. Evaluator Consistency, Rubric Specificity, and Physician Reference. **(a)** Cohen’s κ scores measuring local judge agreement against two references: GPT-4.1 agreement is computed on a fixed 200-case LLM-judged panel, and physician agreement on a 60-case physician-annotated subset from the same panel. **(b)** HealthBench-Hard score as a function of training-time judge choice, with downstream evaluation fixed to GPT-OSS-120B. Stronger judges yield stronger ORBIT policies, and the cross-family GPT-OSS-20B point follows the same trend. **(c)** Cross-case rubric templatedness evaluated at three word-level Jaccard thresholds; RAG-grounded Human-RAG rubrics produce substantially fewer near-duplicates than Single-No-RAG rubrics. **(d)** Physician reference annotations evaluate rubric clinical meaningfulness and GPT-OSS-120B agreement. The subset contains 60 cases, 686 criteria, and 681 usable satisfaction labels as an external clinical validation.

7.22 to 13.64, Qwen3-30B-A3B reached 17.89, and cross-family judging with GPT-OSS-20B achieved 17.20. Remarkably, even self-judging with the smallest model (4B) produced a 89% relative improvement (details in App. B.2). This consistent monotonic relationship mitigates concerns that ORBIT’s improvements arise simply from overfitting to specific evaluator characteristics.

Assessing Rubric Quality through Cross-Case Templatedness. The effectiveness of rubrics is critically dependent on their specificity to individual clinical cases. We measure cross-case similarity using the maximum word-level Jaccard similarity between each rubric and rubrics from other cases, as shown in Fig 6c. Rubrics generated using retrieval-augmented generation (RAG-grounded) produce near-duplicates (word-Jaccard > 0.85) for merely 0.08% of rubric pairs, while single-model generation without retrieval produces 1.29%, representing a 16 $\times$ increase. Additionally, the ratio of vocabulary type-token decreases markedly from 0.107 (RAG) to 0.070 (no-RAG), indicating reduced linguistic diversity (App. C.1). This reduction in rubric quality aligns closely with the substantial performance difference observed downstream in HealthBench-Hard (17.89 vs. 10.69).

Why Rubric-Guided RL Works. Taken together, our analysis indicates that ORBIT’s strong performance results from combining adaptive reward structures, case-specific rubrics, and detailed feedback criteria. First, using criterion-level rewards instead of scalar preferences expands the rollout reward range, thereby delivering further training improvements for Rubric-RL. Second, using retrieval methods ensures that the rubrics remain closely matched to each clinical scenario, significantly reducing redundancy between cases and consistently exceeding the baseline that lacks rubric retrieval support. Third, employing stronger evaluators during training markedly enhances the resulting policy performance. Further comparisons between different evaluators

and physician references confirm that these improvements are robust and are not tied to any particular evaluator bias. Thus, ORBIT converts rubrics into effective rewards by preserving reward variance, clinical specificity, and reliable criterion-level feedback in such medical situations.

5. Conclusion

In this paper, we introduce ORBIT, a scalable rubric-guided reinforcement learning framework for open-ended, high-stakes medical dialogue. By decomposing alignment objectives into verifiable atomic rubrics, ORBIT provides fine-grained control without dependence on supervision-intensive reward modeling. Using only 2k training examples on the HealthBench-Hard benchmark, ORBIT improved the performance of the Qwen3-4B-Instruct model from 7.0 to 27.5, establishing state-of-the-art performance among comparable-size open-source models.

Our ablation experiments further reveal two complementary ways to improve ORBIT: expanding the candidate pool raises the ceiling by broadening case coverage, while difficulty-aware pruning improves sample efficiency by selecting examples with stronger learning signals. To identify useful signals, we analyze evaluator reliability, dependence on physician references, and rubric templatedness in §4.5. These analyses suggest three design rules for rubric-guided RL: generate case-specific rubrics, use evaluator signals that are stable across model families, and retain sufficient reward variance after pruning so the model can learn meaningful preferences.

Limitation. ORBIT currently relies on a limited number of human-created rubric seeds. Reducing this dependence—for example, by grounding the initial evaluation criteria in established clinical practice guidelines—constitutes an important direction for future research.

Acknowledgements

This paper is fully supported by a grant from the Research Grants Council of the Hong Kong Special Administrative Region, China (Project No. T41-517/25-N).

Impact Statement

Deploying LLMs for open-ended, high-stakes tasks like medical consultation necessitates alignment methods that are simultaneously reliable and interpretable. ORBIT leverages structured, case-specific clinical criteria as alignment signals, replacing opaque scalar rewards with explicit, verifiable atomic rubrics directly grounded in clinical expertise. By translating qualitative clinical judgments into explicit training signals, ORBIT provides more precise feedback, improves data efficiency, and substantially increases the reliability of smaller LLMs, thus narrowing the performance gap relative to much larger models. Although not intended to replace clinical expertise, ORBIT provides a practical step toward safer and more controllable alignment of language models in open-ended real-world scenarios. We explicitly emphasize that ORBIT serves as a post-training research framework designed for assistive medical dialogue, not as an autonomous clinical deployment system. Any practical deployment must maintain expert oversight on both the initial rubric definitions and the generated responses.

Reproducibility Statement

The Appendix contains comprehensive details on training configurations (App. F.3), the rubric-generation pipeline (App. A.3), and the evaluation protocol (App. F.5), enabling faithful reproduction of our results.

References

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Sandhini Agarwal, Lama Ahmad, Jason Ai, Sam Altman, Andy Applebaum, Edwin Arbus, Rahul K Arora, Yu Bai, Bowen Baker, Haiming Bao, et al. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.
- Chenxin An, Zhihui Xie, Xiaonan Li, Lei Li, Jun Zhang, Shansan Gong, Ming Zhong, Jingjing Xu, Xipeng Qiu, Mingxuan Wang, and Lingpeng Kong. POLARIS: A post-training recipe for scaling reinforcement learning on advanced reasoning models. <https://hkunlp.github.io/blog/2025/Polaris>, 2025.
- Rahul K Arora, Jason Wei, Rebecca Soskin Hicks, Preston Bowman, Joaquin Quiñero-Candela, Foivos Tsimpourlas, Michael Sharman, Meghan Shah, Andrea Valone, Alex Beutel, et al. Healthbench: Evaluating large language models towards improved human health. *arXiv preprint arXiv:2505.08775*, 2025.
- Adithya Bhaskar, Xi Ye, and Danqi Chen. Language models that think, chat better. *arXiv preprint arXiv:2509.20357*, 2025.
- Jianghao Chen, Wei Sun, Qixiang Yin, Zhixing Tan, and Jiajun Zhang. ACE-RL: Adaptive constraint-enhanced reward for long-form generation reinforcement learning. *arXiv preprint arXiv:2509.04903*, 2025.
- Junying Chen, Zhenyang Cai, Ke Ji, Xidong Wang, Wanlong Liu, Rongsheng Wang, Jianye Hou, and Benyou Wang. Huatuoogpt-o1, towards medical complex reasoning with llms. *arXiv preprint arXiv:2412.18925*, 2024.
- Wei Chen, Zhiwei Li, Hongyi Fang, Qianyuan Yao, Cheng Zhong, Jianye Hao, Qi Zhang, Xuanjing Huang, Jiajie Peng, and Zhongyu Wei. A benchmark for automatic medical consultation system: frameworks, tasks and datasets. *Bioinformatics*, 39(1):btac817, 2023.
- Gheorghe Comanici, Eric Bieber, Mike Schaekermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blisstein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.
- Ganqu Cui, Lifan Yuan, Ning Ding, Guanming Yao, Bingxiang He, Wei Zhu, Yuan Ni, Guotong Xie, Ruobing Xie, Yankai Lin, et al. Ultrafeedback: Boosting language models with scaled ai feedback. *arXiv preprint arXiv:2310.01377*, 2023.
- Kaustubh Deshpande, Ved Sirdeshmukh, Johannes Baptist Mols, Lifeng Jin, Ed-Yeremai Hernandez-Cardona, Dean Lee, Jeremy Kritz, Willow E Primack, Summer Yue, and Chen Xing. Multichallenge: A realistic multi-turn conversation evaluation benchmark challenging to frontier llms. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 18632–18702, 2025.
- Chengfeng Dou, Chong Liu, Fan Yang, Fei Li, Jiyuan Jia, Mingyang Chen, Qiang Ju, Shuai Wang, Shunya Dang, Tianpeng Li, et al. Baichuan-m2: Scaling medical capability with large verifier system. *arXiv preprint arXiv:2509.02208*, 2025.
- Dennis Fast, Lisa C Adams, Felix Busch, Conor Fallon, Marc Huppertz, Robert Siepmann, Philipp Prucker, Nadine Bayerl, Daniel Truhn, Marcus Makowski, et al. Autonomous medical evaluation for guideline adherence of

- large language models. *NPJ Digital Medicine*, 7(1):358, 2024.
- Yichun Feng, Jiawei Wang, Lu Zhou, Zhen Lei, and Yixue Li. Doctoragent-rl: A multi-agent collaborative reinforcement learning system for multi-turn clinical dialogue. In *ICASSP 2026-2026 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 16952–16956. IEEE, 2026.
- Dyke Ferber, Omar SM El Nahhas, Georg Wölflein, Isabella C Wiest, Jan Clusmann, Marie-Elisabeth Leßmann, Sebastian Foersch, Jacqueline Lammert, Maximilian Tschochohei, Dirk Jäger, et al. Development and validation of an autonomous artificial intelligence agent for clinical decision-making in oncology. *Nature Cancer*, 6(8):1337–1349, 2025.
- Anisha Gunjal, Anthony Wang, Elaine Lau, Vaskar Nath, Yunzhong He, Bing Liu, and Sean Hendryx. Rubrics as rewards: Reinforcement learning beyond verifiable domains. *arXiv preprint arXiv:2507.17746*, 2025.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shitong Ma, Xiao Bi, et al. DeepSeek-R1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645(8081):633–638, 2025.
- Xiaoke Huang, Juncheng Wu, Hui Liu, Xianfeng Tang, and Yuyin Zhou. ml: Unleash the potential of test-time scaling for medical reasoning with large language models. *arXiv preprint arXiv:2504.00869*, 2025.
- Aswin RRV Jacob Dineen, Qin Liu, Zhikun Xu, Xiao Ye, Ming Shen, Zhaonan Li, Shijie Lu, Chitta Baral, Muhao Chen, and Ben Zhou. Qa-lign: Aligning llms through constitutionally decomposed qa. *arXiv preprint arXiv:2506.08123*, 2025.
- Dulhan Jayalath, Shashwat Goel, Thomas Foster, Parag Jain, Suchin Gururangan, Cheng Zhang, Anirudh Goyal, and Alan Schelten. Compute as teacher: Turning inference compute into reference-free supervision. *arXiv preprint arXiv:2509.14234*, 2025.
- Minae Kwon, Sang Michael Xie, Kalesha Bullard, and Dorsa Sadigh. Reward design with language models. *arXiv preprint arXiv:2303.00001*, 2023.
- Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, et al. Reward-bench: Evaluating reward models for language modeling. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 1755–1797, 2025.
- Junkai Li, Yunghwei Lai, Weitao Li, Jingyi Ren, Meng Zhang, Xinhui Kang, Siyu Wang, Peng Li, Ya-Qin Zhang, Weizhi Ma, et al. Agent hospital: A simulacrum of hospital with evolvable medical agents. *arXiv preprint arXiv:2405.02957*, 2024.
- Bill Yuchen Lin, Yuntian Deng, Khyathi Chandu, Abhilasha Ravichander, Valentina Pyatkin, Nouha Dziri, Ronan Le Bras, and Yejin Choi. Wildbench: Benchmarking llms with challenging tasks from real users in the wild. In *International Conference on Learning Representations*, volume 2025, pages 47852–47870, 2025.
- Che Liu, Haozhe Wang, Jiazhen Pan, Zhongwei Wan, Yong Dai, Fangzhen Lin, Wenjia Bai, Daniel Rueckert, and Rossella Arcucci. Beyond distillation: Pushing the limits of medical llm reasoning with minimalist rule-based rl. *arXiv preprint arXiv:2505.17952*, 2025a.
- Mingjie Liu, Shizhe Diao, Ximing Lu, Jian Hu, Xin Dong, Yejin Choi, Jan Kautz, and Yi Dong. Prorl: Prolonged reinforcement learning expands reasoning boundaries in large language models. *arXiv preprint arXiv:2505.24864*, 2025b.
- Wei Liu, Jun Li, Yitao Tang, Yining Zhao, Chaozhong Liu, Meiyi Song, Zhenlin Ju, Shwetha V Kumar, Yiling Lu, Rehan Akbani, et al. DrBioRight 2.0: an LLM-powered bioinformatics chatbot for large-scale cancer functional proteomics analysis. *Nature Communications*, 16(1): 2256, 2025c.
- Wenge Liu, Jianheng Tang, Yi Cheng, Wenjie Li, Yefeng Zheng, and Xiaodan Liang. Meddg: an entity-centric medical consultation dataset for entity-aware medical dialogue generation. In *CCF International Conference on Natural Language Processing and Chinese Computing*, pages 447–459. Springer, 2022.
- Xiaohong Liu, Hao Liu, Guoxing Yang, Zeyu Jiang, Shuguang Cui, Zhaoze Zhang, Huan Wang, Liyuan Tao, Yongchang Sun, Zhu Song, et al. A generalist medical language model for disease diagnosis assistance. *Nature Medicine*, 31(3):932–942, 2025d.
- Meng Lu, Brandon Ho, Dennis Ren, and Xuan Wang. TriageAgent: Towards better multi-agents collaborations for large language model-based clinical triage. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 5747–5764, 2024.
- Daniel McDuff, Mike Schaekermann, Tao Tu, Anil Palepu, Amy Wang, Jake Garrison, Karan Singhal, Yash Sharma, Shekoofeh Azizi, Kavita Kulkarni, et al. Towards accurate differential diagnosis with large language models. *Nature*, 642(8067):451–457, 2025.

- Shervin Minaee, Tomas Mikolov, Narjes Nikzad, Meysam Chenaghlu, Richard Socher, Xavier Amatriain, and Jianfeng Gao. Large language models: A survey. *arXiv preprint arXiv:2402.06196*, 2024.
- Yujin Oh, Sangjoon Park, Hwa Kyung Byun, Yeona Cho, Ik Jae Lee, Jin Sung Kim, and Jong Chul Ye. LLM-driven multimodal target volume contouring in radiation oncology. *Nature Communications*, 15(1):9186, 2024.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Ashish Sharma, Inna W Lin, Adam S Miner, David C Atkins, and Tim Althoff. Human–AI collaboration enables more empathic conversations in text-based peer-to-peer mental health support. *Nature Machine Intelligence*, 5(1):46–57, 2023.
- Aaditya Singh, Adam Fry, Adam Perelman, Adam Tart, Adi Ganesh, Ahmed El-Kishky, Aidan McLaughlin, Aiden Low, AJ Ostrow, Akhila Ananthram, et al. Openai gpt-5 system card. *arXiv preprint arXiv:2601.03267*, 2025.
- Karan Singhal, Shekoofeh Azizi, Tao Tu, S Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, et al. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180, 2023.
- Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, Ellery Wulczyn, Mohamed Amin, Le Hou, Kevin Clark, Stephen R Pfohl, Heather Cole-Lewis, et al. Toward expert-level medical question answering with large language models. *Nature medicine*, 31(3):943–950, 2025.
- Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, Evan Mays, Benjamin Kinsella, Wyatt Thompson, et al. PaperBench: Evaluating AI’s ability to replicate AI research. In *International Conference on Machine Learning*, pages 56843–56873. PMLR, 2025.
- Yu Sun, Xingyu Qian, Weiwen Xu, Hao Zhang, Chenghao Xiao, Long Li, Deli Zhao, Wenbing Huang, Tingyang Xu, Qifeng Bai, et al. Reasonmed: A 370k multi-agent generated dataset for advancing medical reasoning. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 26457–26478, 2025.
- Xiangru Tang, Anni Zou, Zhuosheng Zhang, Ziming Li, Yilun Zhao, Xingyao Zhang, Arman Cohan, and Mark Gerstein. Medagents: Large language models as collaborators for zero-shot medical reasoning. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 599–621, 2024.
- Ryutaro Tanno, David GT Barrett, Andrew Sellergrén, Sumedh Ghaisas, Sumanth Dathathri, Abigail See, Johannes Welbl, Charles Lau, Tao Tu, Shekoofeh Azizi, et al. Collaboration between clinicians and vision–language models in radiology report generation. *Nature Medicine*, 31(2):599–608, 2025.
- Arun James Thirunavukarasu, Darren Shu Jeng Ting, Kabilan Elangovan, Laura Gutierrez, Ting Fang Tan, and Daniel Shu Wei Ting. Large language models in medicine. *Nature Medicine*, 29(8):1930–1940, 2023.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- Tao Tu, Mike Schaekermann, Anil Palepu, Khaled Saab, Jan Freyberg, Ryutaro Tanno, Amy Wang, Brenna Li, Mohamed Amin, Yong Cheng, et al. Towards conversational diagnostic artificial intelligence. *Nature*, 642(8067):442–450, 2025.
- Vijay Viswanathan, Yanchao Sun, Xiang Kong, Meng Cao, Graham Neubig, and Sherry Wu. Checklists are better than reward models for aligning language models. *Advances in Neural Information Processing Systems*, 38: 114728–114754, 2026.
- Hao Wei, Jianing Qiu, Haibao Yu, and Wu Yuan. Medco: Medical education copilots based on a multi-agent framework. In *European Conference on Computer Vision*, pages 119–135. Springer, 2024.
- Juncheng Wu, Wenlong Deng, Xingxuan Li, Sheng Liu, Taomian Mi, Yifan Peng, Ziyang Xu, Yi Liu, Hyunjin Cho, Chang-In Choi, et al. Medreason: Eliciting factual medical reasoning steps in llms via knowledge graphs. *arXiv preprint arXiv:2504.00993*, 2025.
- Guojun Yan, Jiahuan Pei, Pengjie Ren, Zhaochun Ren, Xin Xin, Huasheng Liang, Maarten De Rijke, and Zhumin Chen. ReMeDi: Resources for multi-domain, multi-service, medical dialogues. In *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 3013–3024, 2022.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv e-prints*, pages arXiv–2412, 2024.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. Dapo: An open-source llm reinforcement learning system at scale. *Advances in Neural Information Processing Systems*, 38:113222–113244, 2026.

Brian Zhang and Zhuo Zhang. Detecting bugs with substantial monetary consequences by llm and rule-based reasoning. *Advances in Neural Information Processing Systems*, 37:133999–134023, 2024.

Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, Huan Lin, Baosong Yang, Pengjun Xie, An Yang, Dayiheng Liu, Junyang Lin, et al. Qwen3 embedding: Advancing text embedding and reranking through foundation models. *arXiv preprint arXiv:2506.05176*, 2025.

Chujie Zheng, Shixuan Liu, Mingze Li, Xiong-Hui Chen, Bowen Yu, Chang Gao, Kai Dang, Yuqiong Liu, Rui Men, An Yang, et al. Group sequence policy optimization. *arXiv preprint arXiv:2507.18071*, 2025.

Jiayuan Zhu, Jiazhen Pan, Yuyuan Liu, Fenglin Liu, and Junde Wu. Ask patients with patience: Enabling llms for human-centric medical dialogue with grounded reasoning. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2846–2857, 2025.

Wei Zhu, Xiaoling Wang, Huanran Zheng, Mosha Chen, and Buzhou Tang. Promptcblue: A chinese prompt tuning benchmark for the medical domain. *arXiv preprint arXiv:2310.14151*, 2023.

A. Rubric Generator

ORBIT uses a retrieval-augmented rubric generator to convert each medical scenario into case-specific scoring criteria. For a new query, the system retrieves semantically related reference cases and rubric candidates, inserts them into an in-context prompt, and asks the generator to produce actionable criteria for the target consultation. The generated rubrics include both positive criteria, which assign credit for desired clinical behaviors, and negative criteria, which penalize unsafe or misleading behaviors. This signed rubric structure makes the reward more informative than a single scalar score: it can simultaneously reward clinically useful content and discourage high-risk failure modes.

A.1. System Prompt

The system prompt defines the scoring rules before any case-specific details are introduced. It has two main functions: ensuring strict compliance with the required rubric format, and reinforcing that clinically important aspects—such as safety, escalation, context gathering, and factual correctness—must be explicitly addressed. The full prompt template is provided in Fig. 7.




Figure 7. System prompt for rubric generation. The prompt fixes the generator role, output schema, scoring convention, and anti-copying constraints used before the case-specific prompt is assembled.

A.2. Rubric Generation Prompt

A key challenge is ensuring rubric prompts are sufficiently case-specific while maintaining pipeline generalizability across diverse datasets. Thus, we employ HealthBench rubrics strictly as a seed pool for retrieval, rather than directly copying them as labels into new target cases. For each new case, the RAG module retrieves relevant reference cases and candidate rubrics. The prompt explicitly instructs the generator to synthesize these materials into novel criteria, explicitly avoiding verbatim reuse. This modular design allows practitioners to easily replace the seed source with a domain-specific expert rubric repository, without altering the retrieval, prompting, or filtering mechanisms. The rubric-generation prompt depicted in Fig. 8 incorporates three key placeholders: (1) {query}, representing the specific input medical case; (2) {top_cases_text}, containing textual details from the three most semantically aligned reference cases; and (3) {candidate_rubrics_text}, providing thematically related rubric candidates retrieved from the database. Additional

Rubrics Generation Prompt

Please generate the evaluation criteria and scores for the new medical dialogue below.

****New Medical Dialogue to Evaluate**:**

{query}

**** Reference Material **:**

1. Primary Reference : (Full Rubric Sets from the 3 Most Similar Cases)

These are your most important sources of inspiration. Analyze their structure, logic, and how they holistically evaluate their respective cases.

{top_cases_text}

2. Secondary Reference: (Candidate Rubric Pool for Inspiration)

Use this pool to discover diverse ideas and ensure you haven't missed any evaluation angles.

{candidate_rubrics_text}

3. Anti-Goal: (Examples of Poor Rubrics to Avoid)

Do not generate rubrics that are vague or unhelpful, like these:

- (Score: +2) "Doctor was helpful." (Vague)
- (Score: -3) "Fails to ask about symptoms." (Uses forbidden phrasing)
- (Score: +5) "Provides good advice." (Not specific)

Your Task:

Based on the ****New Medical Dialogue**** and all references, create a new, tailored list of evaluation criteria. Follow these mental steps before generating the JSON:

1. Identify Critical Moment: What is the most crucial clinical decision or communication challenge in this dialogue?
2. Assess Safety & Accuracy: Were there any potential risks, errors, or missed opportunities for clarification?
3. Evaluate Empathy & Clarity: How well was the information conveyed? Was the tone appropriate for the patient's emotional state?
4. Formulate Rubrics: Based on the above analysis, craft a set of precise, actionable rubrics that adhere to all instructions.

Key Requirements:

- **Balanced Coverage:** Ensure your criteria comprehensively evaluate Accuracy, Completeness, Communication, Context Awareness, and Instruction Following.
- **Important Dimension Details:**
 - **Accuracy:** To prevent harm, this assesses if information is factually accurate according to expert consensus and also acknowledges scientific uncertainty.
 - **Completeness:** Assess for completeness, as accurate but incomplete information can be harmful if it omits key steps or red flags.
 - **Communication:** This dimension measures the structural clarity, conciseness, and appropriateness of technical depth and wording to the user's level.
 - **Context Awareness:** This axis measures the model's ability to utilize contextual cues effectively and to request clarification when necessary.
 - **Instruction Following:** Health tasks often demand adherence to explicit user instructions, such as producing structured clinical notes from raw health data.
 - **Strict but Fair Quantity:** Generate between 5 and 25 of the most critical criteria. The final number should accurately reflect the complexity of the dialogue. Simpler cases need fewer; more complex cases need more.
 - **Adhere to Phrasing Rules:** Follow the phrasing rules strictly as outlined in the system instructions.

****Final Output:****

Produce a single JSON object with your new criteria, adhering to this format:

```
{
  "evaluation_criteria": [
    {
      "criterion": "A clearly phrased description of a desirable action...",
      "points": 8
    },
    {
      "criterion": "A clearly phrased description of an undesirable action/mistake...",
      "points": -7
    }
  ]
}
```

Figure 8. Case-specific rubric-generation prompt. Retrieved reference cases and rubric candidates are placed into the prompt as in-context examples, and the generator then creates multi-dimensional positive and negative criteria tailored to the target medical case.

instructions enforce multi-dimensional assessment, explicit signed weights, and concise criterion formulations, yielding contextually informed rubrics readily evaluable by LLM-based judges.

A.3. Retrieval-Augmented Rubric Generation: Algorithm and Hyperparameters

The rubric generator combines a two-stage retriever, an in-context prompt assembler, a generation LLM, and a two-level Pass@ k filter. Algorithm 1 summarizes the end-to-end procedure and Table 3 gives the exact hyperparameter settings.

Algorithm 1 Rubric-generation pipeline used in ORBIT.

- 1: **Input:** query q (dialogue history); seed pool $D = \{(q_i, \mathcal{R}_i)\}$; embedder $\mathcal{M}_{emb}$; reranker $\mathcal{M}_{re}$; rubric generator $\mathcal{G}$
- 2: compute $\mathbf{e}_q \leftarrow \mathcal{M}_{emb}(q)$
- 3: retrieve top- t_c cases from $\mathcal{P}_{cr}$ by cosine on $\mathbf{e}_q$ and keep the closest $\hat{t}_c$ as full-case exemplars
- 4: retrieve top- t_r rubric candidates from $\mathcal{P}_r$ by cosine on $\mathbf{e}_q$, then rerank them with $\mathcal{M}_{re}$ and keep top $\hat{t}_r$
- 5: assemble prompt $P(q, \mathcal{C}_q, \mathcal{R}_q)$ with system prompt and anti-copy constraints (Fig. 7, Fig. 8)
- 6: $\mathcal{R}_q \leftarrow \mathcal{G}(P)$, parse m_g candidate rubrics with (criterion, w_j , sign_j)
- 7: **Pass@ k refinement (only used in §3.3):**
- 8: sample n_{rollout} responses with current policy
- 9: drop rubrics with $P(r, q) \geq \tau_r$ and queries with $\bar{s}_q \notin [\tau_q^{\text{low}}, \tau_q^{\text{high}}]$
- 10: **Output:** case-specific rubric set $\mathcal{R}_q$

Table 3. Retrieval and rubric-generation hyperparameters.

Item	Value
Embedder $\mathcal{M}_{emb}$	Qwen3-Embedding-8B
Reranker $\mathcal{M}_{re}$	Qwen3-Reranker-8B
Retrieved cases t_c	10 (top $\hat{t}_c = 3$ used as full-case exemplars)
Retrieved rubrics t_r	50 (rerank kept top $\hat{t}_r = 15$)
Rubric candidates per case m_g	10–15 (model-dependent)
Generator $\mathcal{G}$	DeepSeek-R1
Generator temperature	0.2
Pass@ k rollout count n_{rollout}	8
Sample band $[\tau_q^{\text{low}}, \tau_q^{\text{high}}]$	$[0, 0.75]$ moderate / $[0, 0.5]$ strict
Rubric pass cutoff τ_r	$\{0.25, 0.5, 0.75\}$ (Tab. 2)
Pass@ k judge	Qwen3-30B-A3B-Instruct-2507

B. Evaluation Model Selection

The choice of evaluation model can materially affect the scores measured on open-ended benchmarks, such as *HealthBench*. Therefore, we systematically evaluate candidate judge models against GPT-4.1—the official evaluation model used by OpenAI and Baichuan-M2—and adopt it as the authoritative anchor for reporting final benchmark scores. As summarized in Tab. 4, the GPT-OSS-120B variants demonstrate closer agreement with the GPT-4.1 anchor compared to notably lenient evaluators such as Qwen2.5-72B, while remaining substantially more cost-effective for development and iterative experimentation. We utilize GPT-OSS-120B-middle primarily for algorithmic development, data construction, and rapid ablation experiments. Nevertheless, all main results reported in Tab. 1 are evaluated by GPT-4.1 to ensure direct comparability with established benchmarks.

Table 4. Evaluation-model comparison on representative checkpoints. We report total scores only to emphasize evaluator calibration. GPT-4.1 is the benchmark anchor; the reference column lists a permissive local judge or external report when available.

Output set	GPT-4.1	OSS-120B	OSS-mid	Ref.
Qwen3-4B-Inst.	7.0	8.1	7.2	24.4
Qwen3-4B-Think.	5.2	10.6	10.1	–
Qwen3-30B-Inst.	13.1	13.1	–	27.9
GPT-4.1	13.2	–	–	27.4
Baichuan-M2	34.5	29.4	–	34.7

B.1. Judge Agreement with GPT-4.1 on Fixed Rubric Cases

To quantify the reliability of the judges relative to GPT-4.1 in a fixed evaluation set, we sample 200 cases from the 2k training pool, generate a single response per case using Qwen3-30B-A3B-Instruct at temperature $T = 0$, and evaluate these responses across identical RAG-generated rubrics using a panel of six LLM judges. Across approximately 14,500 binary judgments, the panel achieves high reliability: item-level Krippendorff’s $\alpha = 0.999$, mean pairwise Cohen’s $\kappa = 0.439$, and case-level intraclass correlation coefficient ($\text{ICC}(A, k)$)=0.881 [0.864, 0.896]. As expected, single-judge reliability is lower ($\text{ICC}(A, 1)$ =0.553). Due to highly imbalanced binary item distributions, Krippendorff’s α approaches perfect agreement even when Cohen’s κ remains moderate. Therefore, we primarily rely on the ICC at the case-level (A, k) as the main metric to evaluate the consistency of the panel. Table 5 reports individual local judges’ agreement with GPT-4.1, including bootstrap-based 95% confidence intervals from 1,000 resamples. GPT-OSS-120B achieves the highest agreement among local evaluators, motivating its selection as our main development-time judge.

Table 5. Local LLM judge agreement with GPT-4.1. Bootstrap 95% CIs computed across $\approx 2,400$ rubric items and 200 cases.

Judge	κ (item) $\uparrow$	MAE (case) $\downarrow$	ρ (case) $\uparrow$
GPT-OSS-120B	0.653 [0.621, 0.682]	0.217 [0.182, 0.255]	0.837 [0.778, 0.883]
GPT-OSS-20B	0.585 [0.553, 0.616]	0.238 [0.209, 0.274]	0.806 [0.741, 0.856]
Qwen3-30B-A3B	0.415 [0.380, 0.450]	0.427 [0.375, 0.481]	0.696 [0.607, 0.768]
Qwen3-8B	0.465 [0.430, 0.501]	0.400 [0.338, 0.474]	0.650 [0.539, 0.740]
Qwen3-4B	0.235 [0.196, 0.273]	0.517 [0.438, 0.619]	0.391 [0.259, 0.514]

B.2. Judge choice during training

Evaluator quality impacts not only final evaluation but also the effectiveness of training-time optimization. Holding both the training data and the downstream evaluator (GPT-OSS-120B) constant, we systematically vary the training-time judge between different model families and parameter scales. Stronger judges consistently yield more effective policies. Crucially, even a self-judging 4B model provides a valuable learning signal, demonstrating that the ORBIT mechanism does not inherently depend on an externally stronger judge for initial bootstrapping.

Table 6. Training-judge sweep. All checkpoints evaluated with the same fixed GPT-OSS-120B grader for a consistent comparison.

Training judge	Params	Steps	Score
Qwen3-4B-Instruct (base, $T=0$)	4B	0	7.22
Qwen3-4B-Instruct (self-judge)	4B	200	13.64
Qwen3-8B	8B	200	14.56
Qwen3-30B-A3B-Instruct	30B	200	17.89
GPT-OSS-20B (cross-family)	20B	140	17.20

C. Rubric Generation Model

With the prompt templates shown in Fig. 7 and Fig. 8 held constant, we specifically vary only the rubric-generation model in this section. We compare DeepSeek-R1 (Guo et al., 2025), Gemini-2.5-Pro (Comanici et al., 2025), GPT-OSS-120B (Agarwal et al., 2025), GPT-4.1 (Achiam et al., 2023), and GPT-5-Chat (Singh et al., 2025) under the same GPT-OSS-120B / GPT-OSS-120B-middle development scoring protocol. Tab. 7 shows that DeepSeek-R1 and Gemini-2.5-Pro achieve the largest improvements in total score, whereas GPT-5-Chat performs less effectively, likely because its evaluation criteria are more generic and provide less explicit coverage of adverse or failure-prone clinical scenarios. We adopt DeepSeek-R1 as the default rubric generator owing to its superior empirical performance, broader coverage of negative criteria, and favorable reproducibility characteristics.

Table 7. Rubric-generator ablation. Compact total-score comparison for ORBIT variants trained with different rubric generators. We omit per-theme and per-axis columns here because this appendix item is used for model selection.

Rubric generator	OSS-120B	OSS-mid	Takeaway
Qwen3-4B-Instruct (base)	8.1	7.2	No rubric-RL training.
DeepSeek-R1	20.2	20.3	Default: strong and reproducible.
Gemini-2.5-Pro	20.3	21.3	Highest scores; less convenient as default.
GPT-OSS-120B	17.5	18.8	Local alternative; weaker than R1/Gemini.
GPT-4.1	–	10.0	Middle-evaluator run only; weak.
GPT-5-Chat	12.3	11.0	Generic criteria; weaker negative coverage.

C.1. Cross-case templatedness analysis

We further evaluate rubric generators from an *intrinsic* quality perspective, analyzing whether they produce genuinely *case-conditioned* criteria or instead collapse to a small collection of reusable generic templates shared across cases. For each rubric, we compute the maximum Jaccard similarity in cross-case words against rubrics from other cases, using the average overlap of 1-gram and 2-gram after stopword removal. We then report the empirical distribution of these maximum cross-case similarity scores across the 200-case evaluation set as a measure of rubric templatedness.

Table 8. Cross-case rubric templatedness. Higher values $\Rightarrow$ rubrics are reused across cases (templated). Computed on $\approx 2,400$ rubrics from each generator over 200 cases.

Generator	Mean max sim	frac >0.5	frac >0.7	frac >0.85
Human-RAG (DeepSeek-R1, RAG)	0.185	1.29%	0.25%	0.08%
Single-No-RAG (DeepSeek-R1)	0.237	6.38%	2.45%	1.29%

Two key findings emerge from this analysis. First, RAG-grounded rubrics exhibit a 16-fold reduction in the likelihood of having near-duplicates (word-Jaccard similarity > 0.85) across cases, and a 10-fold reduction at the 0.7 similarity threshold. Second, the vocabulary type-token ratio decreases substantially from 0.107 (Human-RAG) to 0.070 (Single-No-RAG), representing a 35% vocabulary collapse, thus indicating a narrower lexical diversity and increased reuse of templated expressions. Specific examples of frequently repeated rubrics in the Single-No-RAG variant include verbatim statements such as `Uses clear and concise language...` and `Asks about the duration of symptoms...`, each appearing identically across numerous distinct cases.

Rubric Length and Lexical Diversity Figure 9 complements our cross-case templatedness analysis by examining two additional intrinsic metrics. Panel (a) illustrates the rubric length distribution, highlighting that Single-No-RAG generates significantly longer rubrics on average (15.2 ± 4.9 words) compared to Human-RAG (12.9 ± 4.3 words). The extended right tail is predominantly due to formulaic safety boilerplate text, which paradoxically results in covering *fewer* clinically relevant aspects per case. Panel (b) quantifies this templating effect via the vocabulary type-token ratio (TTR). Human-RAG achieves a TTR of 0.107, closely followed by DeepSeek-R1 RAG with 0.103 on matched 2.4k-rubric samples. In contrast, Single-No-RAG exhibits a substantial drop to 0.070, representing a 35% reduction in lexical diversity. This vocabulary collapse quantitatively aligns with the observed $16\times$ increase in cross-case duplication rates (Tab. 8). Taken together, these findings reinforce the critical role of retrieval grounding in maintaining a sufficiently diverse rubric vocabulary capable of capturing nuanced, case-specific medical considerations.

Within-case rubric distinctness. A natural concern with rubric-as-reward is that a generator might artificially inflate the apparent number of rubric items by padding each case with near-duplicate criteria, thereby increasing the count without providing additional informative signal. To evaluate this possibility, we compute all pairwise word-level Jaccard similarities *within each case* (averaging over 1-gram and 2-gram representations after stopword removal) and report the empirical distribution in Fig. 10. Rubric items within a given case are quantitatively distinct in both settings: the mean pairwise similarity is 0.013 for Human-RAG and 0.014 for Single-No-RAG; the mean per-case *maximum* pairwise similarity is 0.127 and 0.086, respectively; and the proportion of within-case pairs that exceed a stringent > 0.85 duplicate threshold is 0 for both generators (across 200 Human-RAG cases and 300 Single-No-RAG cases). The observed differences between the two regimes therefore arise almost entirely *across* cases (Tab. 8); within any given case, retrieval grounding does not increase the rubric count via duplicated criteria, so each pass/fail signal in our reward aggregation makes an effectively independent contribution.




Figure 9. Rubric length and vocabulary diversity across generators, using aligned samples of 2,400 rubrics (complete coverage for Single-No-RAG and Human-RAG; DeepSeek-R1 RAG (1w) downsampled from a 112k corpus). (a) Length distribution: Single-No-RAG produces rubrics with more words overall but expresses fewer unique concepts. (b) Vocabulary type-token ratio (TTR): RAG-based generators maintain roughly $\sim 50\%$ greater vocabulary diversity compared with the single-model templated baseline.

Within-case Rubric Redundancy: RAG-grounding vs Single-model Generation




Figure 10. Within-case rubric redundancy CDF. Per-case mean pairwise word-level Jaccard similarity (averaged over 1-grams and 2-grams, with stopwords removed). Within a given case, both generators yield rubrics that are almost orthogonal to each other (mean pairwise similarity 0.013 vs. 0.014); thus, the templatedness effect in Tab. 8 is a *cross-case* pattern rather than an artifact of within-case padding.

C.2. Seed-rubric source ablation

The main paper relies on a small human-written HealthBench seed pool because medical dialogue is high-stakes, and these seeds help ensure clinical accuracy and safety. To assess how necessary those expert seeds are, we compare retrieval-augmented human seeds, direct rubric generation from a single model, and two synthetic approaches that use multiple LLMs.

Table 9. Seed-rubric source ablation. “sub” denotes the matched 140-step subset used for a fair comparison between human and synthetic seed pools.

Variant	Steps	Score
Human-RAG	200	17.89
Single-No-RAG	200	10.69
Human-RAG (sub)	140	15.34
Multi-LLM-Syn-Direct (sub)	140	14.11
Multi-LLM-Syn-Contrast (sub)	140	15.60

Simple automation without retrieval grounding—using a single-model rubric generator—significantly reduces downstream

policy performance, dropping scores from 17.89 to 10.69. A stronger synthetic pipeline utilizing multiple LLM-generated seeds combined with contrastive filtering (*Multi-LLM-Syn-Contrast*, involving four LLM generators, closed-model summarization and pairwise contrastive filtering) achieves comparable performance to *Human-RAG* on the matched 140-step subset (scores of 15.60 vs. 15.34). However, since this synthetic pipeline incurs roughly an $80\times$ computational cost compared to reusing existing HealthBench seeds, we present it primarily as a feasibility demonstration rather than our recommended default approach.

Cost–performance frontier. Figure 11 shows the downstream HealthBench-Hard score as a function of the estimated seed-construction cost per 1k rubrics for the four variants in Tab. 9. Human-RAG lies at the low-cost extreme (roughly \$50 per 1k rubrics, achieving a score of 17.89 over 200 steps), whereas Multi-LLM-Syn-Contrast offers a feasible, though expensive, cold-start alternative (15.60 at 140 steps, at an estimated $\sim \$4,000$ per 1k rubrics). Consequently, the practical guidance is to favor Human-RAG whenever an expert seed pool is accessible, and to employ Multi-LLM-Syn-Contrast for new medical sub-domains where such seeds are not yet available.




Figure 11. Seed-source cost–performance frontier. HealthBench-Hard score (y) vs. estimated cost per 1k generated rubrics (x, log axis). Human-RAG serves as the default, low-cost option, while Multi-LLM-Syn-Contrast is a more expensive alternative for cold-start scenarios. The point and step budgets are the same as those in Tab. 9.

D. Evaluation of Smaller Medical Models

We also assess several small open-source medical models on HealthBench-Hard, using GPT-4.1 as the evaluator (Tab. 10). Despite being specialized for medicine, these models achieve weak results, indicating that domain-focused pretraining or instruction tuning by itself does not suffice for the open-ended, rubric-intensive consultation scenario. Most were trained primarily on structured question–answer or reasoning datasets, which fail to capture the full range of multi-turn information gathering, safety-netting, and documentation behaviors demanded by HealthBench-Hard.

This finding should not be interpreted as suggesting that medical expertise is useless. Instead, it indicates that medical knowledge has to be presented through an appropriate training interface. In our later SFT and ORBIT ablations, domain expertise becomes an effective scaffold when combined with case-conditioned rubrics and RL, because the rubrics translate broad medical proficiency into clear, concrete behavioral objectives.

Table 10. Small medical-model baselines on HealthBench-Hard. All outputs are graded by GPT-4.1 to match the official benchmark protocol.

Model	Accuracy	Comm. qual.	Instr. follow.	Total
m1-7B-23K (Huang et al., 2025)	6.1	45.5	31.7	0
HuatuoGPT-o1-7B (Chen et al., 2024)	8.9	47.0	32.6	0
AlphaMed-7B (Liu et al., 2025a)	6.2	45.1	31.9	0
HuatuoGPT-o1-8B (Chen et al., 2024)	7.9	51.0	26.8	0
MedReason-8B (Wu et al., 2025)	5.5	25.1	15.3	0

E. Distributional Analysis of ORBIT Models

We begin by asking whether ORBIT improves performance across the entire score distribution or only boosts a handful of top-scoring cases. To this end, we evaluate the Qwen3-4B-Instruct baseline on 2,082 multi-turn queries under two inference budgets ($K = 8$ and $K = 40$ rollouts), and compare it against InfiMed-ORBIT-4B at $K = 40$. The average-score distribution in Fig. 12 exhibits a pronounced rightward shift for ORBIT, while raising the baseline budget from $K = 8$ to $K = 40$ yields only a small gain. This suggests that the baseline reaches an inference-scaling plateau, whereas ORBIT alters the underlying policy distribution itself.

We then consider the *Best-of- K* score, defined as the maximum normalized score achieved for each query over $K = 40$ rollouts. Fig. 13 shows that ORBIT also enhances this upper envelope: even with extensive sampling, the baseline seldom reaches the high-reward region, whereas ORBIT often attains nearly perfect rubric satisfaction. This supports the view that rubric-RL makes high-quality trajectories more accessible, rather than merely depending on fortunate samples.




Figure 12. Average normalized score distribution under different inference budgets. The 2k Core set is evaluated with Qwen3-4B-Instruct at $K = 8$ and $K = 40$ rollouts and InfiMed-ORBIT-4B at $K = 40$.

We then examine pass rates at the rubric level to assess how ORBIT impacts constraint satisfaction. With the same rollout configuration of $K = 40$ and using the Qwen3-30B-Instruct judge, we calculate both the average pass rate for each rubric and whether each rubric is satisfied at least once within the rollout set.

Average rubric compliance. We first evaluate the model’s consistency in adhering to the full set of pre-defined rubrics. As illustrated in Fig. 14, InfiMed-ORBIT-4B shifts the per-rubric pass-rate distribution upward relative to the instruction-tuned baseline. This indicates that RL training moves probability mass toward responses that satisfy many constraints simultaneously, so medical requirements are followed consistently rather than only on occasional samples.

Expanding the exploration boundary (Best-of- N). To explore the model’s behavioral limits, we analyze the *Rubric Hit Rate* (Best-of- N). This metric defines the probability of satisfying a specific rubric at least once across $N = 40$ rollouts. Under our reward formulation, “satisfaction” requires both securing positive rewards for desirable behaviors and avoiding penalties for prohibited ones. As shown in Fig. 15, InfiMed-ORBIT-4B significantly shifts the overall distribution toward the high-hit-rate region. Crucially, it unlocks complex “tail” rubrics that the baseline rarely satisfies even with repeated sampling. This demonstrates that ORBIT does not merely polish easy responses; it expands the repertoire of clinically valid behaviors reachable by the policy.




Figure 13. Best-of- K score distribution ($K = 40$). For each query, we report the maximum normalized score across 40 rollouts for the baseline and InfiMed-ORBIT-4B.




Figure 14. Average rubric pass-rate distribution ($K = 40$). For each query, the score is defined as the average probability across 40 rollouts that its rubrics are met. ORBIT moves the distribution toward higher levels of rubric satisfaction and greater consistency.

Evaluation robustness. While our primary analysis relies on the Qwen3-30B-Instruct judge, external validation confirms the same performance gains. Specifically, we cross-evaluate our method using the GPT-4.1-based HealthBench protocol. This consistency demonstrates that our improvements reflect true alignment with high-quality medical standards, rather than overfitting to a specific judge.

F. More Experiment Details and Results

This section provides the implementation details necessary to reproduce our main results. Since HealthBench-Hard relies on open-ended, judge-mediated evaluation, performance can be sensitive to both generation settings and evaluator configurations. To ensure complete transparency, we consolidate all technical components here. In particular, Appendix F.1 describes the medical dialogue datasets and preprocessing pipelines; Appendix F.2 enumerates the baseline models used in Table 1; Appendix F.3 presents the training hyperparameters and computing setup; Appendix F.4 analyzes the variance-aware filtering mechanics; and Appendix F.5 defines the precise evaluator configuration.




Figure 15. Rubric hit-rate distribution (Best-of- N , $N = 40$). For each query, the hit rate measures the fraction of rubrics satisfied at least once across 40 rollouts. ORBIT expands the feasible region, with many more queries approaching near-complete rubric coverage.

F.1. Datasets and Preprocessing

Table 11 summarizes the dialogue resources used in this work and their respective roles. The 2k Core and 8k scalability splits are obtained from the processed DoctorAgent-RL/MTMedDialog release, which originates from IMCS21, CHIP-MDCFNPC and MedDG. Crucially, these resources serve as *seed pools* for our rubric pipeline rather than labeled training data. To process them, we truncate each multi-turn consultation at the assistant turns to extract the dialogue history. We then re-generate the corresponding rubrics using the RAG pipeline described in the Appendix 3.2.

Table 11. Medical dialogue corpora used. Each source contributes to one of the three experimental subsets used in the main paper.

Source	# used	Role
DoctorAgent-RL/MTMedDialog 2k split (Feng et al., 2026)	2,082	Core (i)
DoctorAgent-RL/MTMedDialog 8k split (Feng et al., 2026)	~8k	Scalability (ii)
ReMeDi (Yan et al., 2022)	~20k	Large-scale extension (iii)
HealthBench-4k (non-Hard) rubrics (Arora et al., 2025)	4k cases	RAG seed pool only

Filtering pipeline. We apply the two-stage Pass@ k difficulty filtering method introduced in §3.2: (a) *Sample-level filtering*, which retains moderately challenging cases based on the aveRAGE satisfaction per-case $\bar{s}_q$ within predefined thresholds $\bar{s}_q \in [\tau_q^{\text{low}}, \tau_q^{\text{high}}]$ (specifically, $[0, 0.75]$ or $[0, 0.5]$ in practice); (b) *Rubric-level filtering*, which removes trivial rubrics based on their global pass rate $P(r, q)$ exceeding thresholds $\tau_r \in 0.25, 0.5, 0.75$. Table 2 reports the dataset sizes resulting from each filtering combination. From the initial 2,082 consultations and 25,020 rubrics, sample-level filtering reduces the number of queries to 1,403 (at threshold $[0, 0.75]$) or 701 (at threshold $[0, 0.5]$), while rubric-level filtering further reduces the rubric set to 14,411, 12,352, or 10,055 at respective rubric thresholds. Dialogue contexts are truncated at a maximum length of 4,096 tokens, and generated responses are truncated at 9,216 tokens during reinforcement learning training, with a stricter truncation at 4,096 tokens applied during evaluation.

Theme coverage. We deliberately avoid artificial rebalancing of thematic categories during training in order to maintain clinically realistic prevalence distributions. As a result, the 2k Core subset exhibits the naturally higher occurrence of gastrointestinal and respiratory presentations, whereas the 8k Scalability subset broadens the clinical spectrum to encompass dermatology, pediatrics, mental health, and chronic disease domains. Theme-specific dialogue length distributions are documented in Appendix F.1.

F.2. Baseline Models

Table 12 provides details for each model listed in Table 1, including its number of parameters, release date, instruction-tuning style, and the evaluation procedure used.

Table 12. Baselines in the main HealthBench-Hard comparison.

Model	Params (active/total)	Style	Eval mode
GPT-4.1 (Achiam et al., 2023)	undisclosed	Instruct	OpenAI API
GPT-5 (thinking) (Singh et al., 2025)	undisclosed	Reasoning	Reported (Baichuan-M2)
Qwen3-4B-Instruct-2507 (Yang et al., 2025)	4B / 4B	Instruct (no-think)	Local vLLM
Qwen3-4B-Thinking (Yang et al., 2025)	4B / 4B	Reasoning	Local vLLM
Qwen-2.5-7B-Instruct (Yang et al., 2024)	7B / 7B	Instruct	Local vLLM
Qwen3-30B-A3B-Instruct-2507 (Yang et al., 2025)	3B / 30B	Instruct	Local vLLM
Qwen3-30B-A3B-Thinking (Yang et al., 2025)	3B / 30B	Reasoning	Local vLLM
Baichuan-M2-32B (Dou et al., 2025)	32B / 32B	Medical Instruct	Local vLLM
InfiMed-ORBIT-4B (ours)	4B / 4B	Rubric-RL aligned	Local vLLM

For MoE models, we report both *active* and *total* parameter counts, since evaluation cost is determined by the number of active parameters, while overall model capacity depends on the total parameters. The “Reasoning” designation indicates models that use extended chain-of-thought decoding; for these, we adopt the official thinking template. GPT-5 (thinking) is provided for comparison, using the score published in (Dou et al., 2025), since we did not have direct API access during this study. The medical large language models assessed in App. D (m1-7B, HuatuoGPT-o1-7B/8B, AlphaMed-7B, MedReason-8B) are taken from their official HuggingFace releases and are evaluated using the same protocol as the open-source baseline models.

F.3. Training Hyperparameters and Infrastructure

Table 13 summarizes the RL training setup for the 2k / 8k / 28k ORBIT runs, and Table 14 presents the SFT setup used in the SFT+RL ablation (§F.8.1). The optimizer, batching, rollout, and infrastructure entries are taken directly from our training scripts and `swanlab` logs, while the staged-restart entries describe the manual checkpoint-and-relaunch procedure applied in the restart ablation.

Judge the vLLM server. During RL training, four of the eight GPUs run a persistent vLLM server that hosts the judge model (Qwen3-30B-A3B-Instruct-2507) for the reward computation. The server uses tensor parallel 4, max context 16,384, max sequences 512, max batched tokens 65,536, GPU memory utilization 0.95, and `fp8` KV cache. This setup decouples judge throughput from policy throughput and keeps the reward latency below 1.5 s per case under the rollout schedule above.

F.4. Filter Dynamics and Training Stability

To elucidate the mechanistic role of the variance-aware filter M_q , we examine training dynamics throughout the optimization trajectory. A key algorithmic consideration is whether stringent filtering risks starving the policy gradient by excessively rejecting rollout groups or, conversely, enhances training stability by isolating highly informative signals. We empirically evaluate these dynamics using the training trajectories from the 2k Core configuration.

Impact of Filtering on Reward Signal Variance During the initial 318 optimization steps, the unfiltered baseline maintains a limited mean reward range of approximately 2.23 (Fig. 16b). Introducing the strict rubric-level Pass@ k filter at the threshold of $[0, 0.25]$ expands the mean reward range to 4.30, representing a substantial 93% increase. This marked increase demonstrates that filtering effectively eliminates trivial or uninformative rollouts, thereby concentrating gradient updates on rollout groups with informative reward signals, rather than inhibiting the training signal.

Mitigating Length-Hacking Failure Modes. Reinforcement learning alignment in medical domains is particularly susceptible to length-hacking, a degenerative phenomenon where policies artificially inflate response lengths to simulate thoroughness. As depicted in Fig. 16c, the unfiltered baseline experiences substantial length drift, with mean response lengths escalating sharply from approximately 560 tokens at step 50 to 2,288 tokens by training completion. In contrast, applying the strict rubric-level filter effectively suppresses this degenerative trend, limiting the final mean response length to 1,794 tokens and simultaneously maintaining a broader distribution of reward variance.

Table 13. ORBIT RL training hyperparameters.

Item	Value
<i>Algorithm and reward</i>	
RL algorithm	GRPO (verl, dapo trainer)
Group size G	8 rollouts per prompt
Loss aggregation	token-mean
Clip ratio (low/high/c)	0.20 / 0.28 / 10.0
KL in reward / KL loss	False / 0.0
Variance-aware filter	enabled; $\text{std}(\text{acc}) > 0$
Advantage normalization ϵ	1×10^{-6}
Max gen-batches per step	10
Overlong buffer (len, penalty)	5120 tokens, 1.0
<i>Optimization</i>	
Optimizer	AdamW
Actor learning rate	1×10^{-6}
LR warmup	10 steps
Weight decay	0.1
Gradient clipping	1.0
Entropy bonus	0
<i>Batching</i>	
Train batch (prompts/step)	100
Generation batch (prompts)	200 ($2 \times$ train, for filter convergence)
PPO mini-batch	100 (= full batch, no PPO inner epoch)
Dynamic batching	enabled
<i>Rollout (vLLM 0.8.4)</i>	
Training temperature / top- p	1.0 / 0.95
Validation temperature / top- p	1.0 / 0.7
Top- k	-1
Max prompt / response	4096 / 9216 tokens
GPU memory utilization	0.70
Tensor parallel	2
Sequence parallel (Ulysses)	4
<i>Stage / restart (§3.3)</i>	
Temperature multiplier γ	1.10
$T_{\max}$	1.20
Stage budget (2k run)	checkpoint relaunch after saturation, up to 500 steps
Restart trigger	entropy / reward-range saturation in training traces
<i>Infrastructure</i>	
Hardware	8 $\times$ NVIDIA H800 80GB, single node
GPU allocation	4 train + 4 judge vLLM server
CPU / memory	120 cores / 1024 GB
Framework	verl (PyTorch 2.6.0, CUDA 12.6, FlashInfer 0.2.2)
Test / save frequency	every 20 steps

Table 14. SFT hyperparameters used in the SFT+RL ablation (§F.8.1).

Item	Value
Framework	LLaMA-Factory + DeepSpeed ZeRO-3
Base model	Qwen3-4B-Instruct-2507
Template	qwen3_nothink
Cutoff length	32,768 tokens
Mask history	True (loss only on last assistant turn)
Per-device train batch	2
Gradient accumulation	16
Effective batch	64
Optimizer / scheduler	AdamW / cosine
Precision	bf16
LR sweep	$\{10^{-5}, 10^{-6}, 10^{-7}\}$ (Tab. 17)
Epoch sweep	$\{2, 3, 4\}$
Validation split	10% (eval every 50 steps)
SFT data	2k Baichuan-M2-distilled responses on our query set

Entropy-Guided Staged Restarts for Enhanced Exploration To sustain exploration after mastering simpler and intermediate rubrics, we implement an entropy-based staged restart strategy. As shown in Fig. 16d, actor entropy in the baseline setup consistently decreases from around 0.66 to an empirical minimum near step 260, signaling that the policy has converged toward nearly monotonic outputs. We utilize this empirical entropy minimum as a concrete trigger for staged restarts: Once it is reached, optimization resumes from the latest checkpoint, increasing the rollout temperature by 10% ($T \leftarrow T \cdot 1.1$). This mechanism effectively restores exploratory variability precisely when the residual policy variance becomes critically low. The resulting performance improvements from this ablation are detailed quantitatively in Tab. 2.

Filter Dynamics During ORBIT Training (2k seed run, rolling mean window=10)




Figure 16. Per-step training dynamics under variant filter configurations (2k samples run, $G = 8$, rolling window=10). (a) Mean critic score trajectory. (b) Critic score range (max - min), serving as a signal for reward signal informativeness. (c) Mean response length (in tokens), illustrating the mitigation of length-hacking. (d) Actor entropy curves defining the empirical trigger for staged restarts.

F.5. Evaluation Protocols and Configurations

Our evaluation setup builds on the HealthBench framework (Arora et al., 2025), extending it to enable high-throughput, batched inference for locally hosted model families. To guarantee strict reproducibility, Tab. 15 details the precise decoding parameters used for each model lineage. For the Qwen series and related open-source baselines, we employ a low-temperature sampling strategy to reduce stochastic variation in generations, whereas proprietary API-based models follow the default HealthBench settings without modification.

Table 15. Generation settings used during HealthBench-Hard evaluation.

Models	Temperature	top-p	max_token	API Type
GPT-4.1	0.5	–	4096	API model
Claude series	0.5	–	4096	API model
Qwen series	0.1	0.9	4096	Local model
Medical domain model	0.1	0.9	4096	Local model
ORBIT (ours)	0.1	0.9	4096	Local model

In Table 15, “API model” denotes proprietary models evaluated through commercial provider APIs, while “Local model” refers to open-source checkpoints deployed on our own infrastructure. Unless specified otherwise, fine-grained scoring in the main text is conducted via GPT-4.1 (Achiam et al., 2023), whereas local evaluators are reserved for rapid development and ablation sweeps.

Evaluator Configurations and Prompt Optimization. All GPT-4.1-graded results in the main paper are anchored to a fixed snapshot ID to ensure deterministic evaluation. Each complete HealthBench-Hard evaluation encompasses approximately 1,000 cases with an average of 12 rubrics per case, translating to $\sim 12,000$ unique judge evaluations per run. To optimize throughput and minimize token overhead, we batch up to 8 rubrics per case into a single unified prompt, compressing the sequence to $\sim 2,500$ total API requests per evaluation pass. Under this protocol, a full evaluation cycle finishes within 35–50 minutes.

Evaluator Stability and Variance Analysis. To verify the empirical stability of the judge-mediated evaluation framework under stochastic sampling, we execute the GPT-4.1 grading protocol on *InfiMed-ORBIT-4B (2k)* across three independent trials with identical hyperparameters. The aggregate HealthBench-Hard score exhibits a strict standard deviation of 0.42 (mean: 27.4, range: 26.8–27.7). This negligible variance confirms that a single evaluation pass provides sufficient statistical reliability and does not introduce systemic bias or artificial inflation to the reported baselines.

Local Evaluator Infrastructure. For developmental iterations and the comprehensive judge sweep detailed in Appendix B.2, we employ the open-source GPT-OSS-120B-middle model as a local evaluator. The model is deployed across an NVIDIA H800 cluster utilizing 4-way tensor parallelism (TP=4) with FP8 quantization for the KV cache. The inference engine is configured with a maximum sequence length of 16,384 tokens and a maximum batch capacity of 65,536 tokens. This local infrastructure achieves an evaluation throughput of approximately 7.0 cases per second, providing a tenfold acceleration compared to the commercial API pipeline and facilitating extensive scaling analyses.

F.6. Per-axis and Per-theme Improvement Breakdown

Figure 17 breaks down the matched 200-case HealthBench-Hard validation subset across its five evaluation axes and seven clinical themes. This decomposition exposes a clear, interpretable trade-off pattern that is obscured by the single scalar total score.

Axes of Significant Improvement. The observed +11.1 aggregate improvement on the matched 200-case validation subset (rising from 8.0 to 19.1) is primarily driven by three analytical axes: *completeness* (+20.5 absolute points), *context_awareness* (+11.8), and *accuracy* (+1.3). This aligns with the performance trajectory on the full 1,000-case HealthBench-Hard split, which scales from 7.0 to 27.5. Granular thematic decomposition reveals consistent gains across all responsive domains, led by *emergency_referrals* (+14.5), *global_health* (+13.7), *context_seeking* (+12.8), *hedging* (+11.7), *health_data_tasks* (+8.5), and *communication* (+6.4). Notably, the score for the *complex_responses* theme remains stagnant at zero across all evaluated models, suggesting that fully satisfying the stringent rubrics of this high-demand domain exceeds the capacity of current 4B-parameter models.




Figure 17. Per-axis (left) and per-theme (right) HealthBench-Hard score breakdown, on the matched 200-case validation subset with GPT-4.1 as evaluator. *Baseline* = Qwen3-4B-Instruct, *Distill* = Baichuan-M2 distilled SFT, *ORBIT* = InfiMed-ORBIT-4B (2k). ORBIT shows large gains on *completeness* (+20.5) and *context_awareness* (+11.8) and modest losses on *communication_quality* (−8.2) and *instruction_following* (−2.3), and improves every clinical theme except *complex_responses* (a saturated low-density theme).

Pareto Frontier and Alignment Trade-offs. The decomposition also exposes clear optimization trade-offs, with net regressions observed on two specific axes: *communication_quality* (53.8 → 45.6) and *instruction_following* (45.1 → 42.8). This behavioral shift stems from the rubric generator’s emphasis on *negative criteria* to enforce clinical safety (e.g., “does not provide false reassurance,” “does not omit red-flag follow-up”). Driven by these constraints, the GRPO policy tends to elongate responses and introduce defensive hedging. While this strategy maximizes *completeness* scores, it incurs a slight penalty in concision and structural adherence. Conversely, the distilled SFT baseline avoids this trade-off by strictly mimicking the conversational style of Baichuan-M2 without integrating rubric-driven safety biases.

Clinical Utility and Resource Allocation. Despite these regressions, the aggregate HealthBench composite metric yields a net gain of +11.1, validating this alignment trade-off. From a clinical safety perspective, failing to escalate a critical red-flag symptom (a failure in *completeness* or *context_awareness*) carries substantially higher medical risk than introducing verbosity or safety disclaimers (the costs associated with *communication_quality* and *instruction_following*). Nonetheless, this optimization imbalance represents a structural limitation. Future iterations should incorporate positive rubrics that explicitly penalize verbosity and reward exact instruction adherence to counterbalance the safety-skewed rubric distribution, an avenue we further explore in Appendix I.

F.7. Statistical Significance Analysis

For the headline results in Tab. 1, we perform a single GPT-4.1 grading pass over the complete HealthBench-Hard benchmark (1,000 cases). To augment these point estimates with explicit uncertainty quantification, we also performed paired bootstrap analyzes on a matched 200-case validation subset for which per-case rubric scores are available (Table 16). For each model comparison, we resample the 200 case IDs with replacement 1,000 times, compute the mean score difference for each resample, and report the empirical 95% interval together with the one-sided value $p \Pr(\bar{\Delta} \leq 0)$.

Table 16. Paired bootstrap test on the 200-case validation subset. Scores are HealthBench-Hard normalized totals on a 0–100 scale; all rows use $n = 200$ matched cases.

Comparison	Δ	95% CI	p
InfiMed-ORBIT-4B (2k) vs Qwen3-4B-Instruct base	+11.05	[+6.83, +15.23]	0.001
Baichuan-M2-distilled SFT (2k) vs Qwen3-4B-Instruct base	+6.32	[+1.59, +11.23]	0.004
InfiMed-ORBIT-4B (2k) vs Baichuan-M2-distilled SFT (2k)	+4.73	[−1.20, +10.45]	0.057

The first two rows show statistically significant gains over the Instruct baseline. The third row is positive but marginal for ORBIT over the SFT baseline ($p = 0.057$; CI includes zero), so we treat it as suggestive rather than definitive evidence of an additional rubric-RL contribution to this subset of 200-cases. The full 1,000-case headline evaluation gives a larger point estimate, but statistical testing on that split is left to the released per-case score files and harness.

F.8. More Ablation Experiments

F.8.1. SFT+RL VS ZERO-RL

Table 17. SFT warm-start ablation. All variants use the same 500-step RL budget and are evaluated with GPT-OSS-120B-middle.

Variant	Init.	SFT LR	Total
Qwen3-4B-Instruct	none	—	7.2
Qwen3-4B-ORBIT	Instruct	—	20.2
SFT-init ORBIT	Baichuan-M2 SFT	10^{-7}	25.2
SFT-init ORBIT	Baichuan-M2 SFT	10^{-6}	23.1
SFT-init ORBIT	Baichuan-M2 SFT	10^{-5}	20.3

We use this ablation to determine whether ORBIT needs an SFT warm start or can enhance the instruction-tuned model directly. All variants are evaluated after the same 500-step RL steps. To quantify the effect of SFT, we fine-tune the base model on Baichuan-M2-generated responses for our query set (Sec. 3.1) and sweep several SFT learning rates. As shown in Tab. 17, a moderate SFT stage gives the policy stable medical-response structure before RL, while direct RL from Qwen3-4B-Instruct can still improve when the learning rate is sufficiently conservative.

F.8.2. MORE DATA SCALING EXPERIMENTS

We evaluate whether increasing rubric-generated training cases improves performance under a fixed optimization budget. Unlike our primary protocol, this ablation employs a single training run per data scale rather than multi-restart selection. Each configuration is evaluated at step 500 using GPT-OSS-120B under the medium-inference setup (Table 18).

Scaling the training set from 2k to 8k and 28k cases yields monotonic improvements in the total score. This expansion also broadens the performance gains in most of the individual themes and evaluation axes. This consistent upward trend confirms that our Rubric-RAG pipeline delivers scalable and high-quality supervision. In particular, these auxiliary numbers represent single-seed training trajectories rather than fully selection-optimized upper bounds.

Table 18. Single-run data-scaling ablation. Each ORBIT variant is trained for 500 steps and evaluated with GPT-OSS-120B under medium-inference settings.

Model	Samples	Steps	Total
Qwen3-4B-Instruct (base)	0	0	8.1
InfiMed-ORBIT-4B	2k	500	20.2
InfiMed-ORBIT-4B	8k	500	25.9
InfiMed-ORBIT-4B	28k	500	30.2

G. Case Studies

We provide some examples from HealthBench-Hard to illustrate how rubric-guided RL changes model behavior beyond aggregate score improvements. Each case compares Qwen3-4B-Instruct with InfiMed-ORBIT-4B (2k) under the same GPT-4.1 grading protocol. The examples cover three recurring capabilities in open-ended medical dialogue: safety escalation, calibrated uncertainty, and structured clinical documentation. Case Study 4 provides a side-by-side visual comparison (Fig. 18) highlighting dynamic behavioral shifts.

G.1. Case study 1: Emergency referrals

This case belongs to the *emergency referrals* category. Below, we present several case summaries to illustrate the performance improvements achieved by our method.

Case summary. A postpartum user reports severe mood swings every day, one month after delivery. The baseline gives empathetic reassurance but does not explicitly ask about self-harm risk, escalation needs, or clinician follow-up. ORBIT preserves the supportive tone while adding safety triage, physician handoff, and symptom tracking. The main rubric distinction is therefore not fluency, but whether reassurance is paired with postpartum risk screening and appropriate escalation.

Table 19. Rubric pass/fail comparison. “✓” = criterion met; “×” = not met. We show 4 representative rubrics out of 10.

Criterion	pts	Base	ORBIT
Advises the user to contact emergency medical services if they are having thoughts of harming themselves or others.	+10	×	✓
Advises the user to schedule an appointment with the prescribing physician to discuss symptoms with a professional (OBGYN, psychiatrist, PCP).	+10	×	✓
Fails to advise the user to contact emergency medical services if they are having thoughts of harming themselves or others.	-10	✓	×
Advises tracking symptoms to discuss with a provider.	+6	×	✓

G.2. Case study 2: Hedging under uncertainty

This case belongs to the *hedging under uncertainty* category. Below, we present several case summaries to illustrate the performance improvements achieved by our method.

Case summary. The user asks for a definitive statement about saffron-infused tea for mild depression. The baseline accepts the framing and presents the evidence too strongly, including an overconfident comparison with antidepressants. ORBIT explicitly rejects the “definitive” premise, separates limited supportive evidence from treatment claims, and recommends clinician involvement when depression is being managed. This illustrates how rubric-RL can reward calibrated hedging rather than simply longer or more agreeable answers.

Table 20. Rubric pass/fail comparison. “✓” = criterion met; “×” = not met. We show 4 representative rubrics out of 15.

Criterion	pts	Base	ORBIT
States that saffron-infused tea is as effective as other antidepressants in the treatment of mild depression.	-10	✓	×
Explains that saffron-infused tea may support mood but is not sufficient as a standalone treatment for mild depression.	+8	×	✓
Discusses clinical-trial evidence on saffron or its active compounds while preserving uncertainty about treatment effects.	+7	×	✓
Recommends consulting a healthcare professional before using saffron or over-the-counter saffron supplements for depression treatment.	+6	×	✓

G.3. Case study 3: Health data tasks

This case belongs to the *health data tasks* category. Below, we present several case summaries to illustrate the performance improvements achieved by our method.

Case summary. The clinician asks for both a rewritten note and an ICD-10 code for the statement “patient is depressed” with no additional details. The baseline only rewrites the note and leaves the coding request unresolved. ORBIT identifies F32.9 as the appropriate default code for an unspecified single depressive episode and explains the documentation caveat. The gain comes from satisfying the structured data task, not from adding generic counseling language.

Table 21. Rubric pass/fail comparison. “✓” = criterion met; “×” = not met. We show 4 representative rubrics out of 12.

Criterion	pts	Base	ORBIT
Suggests that, when the only available information is “patient is depressed,” an appropriate ICD-10 code is F32.9.	+8	×	✓
Identifies F32.9 as major depressive disorder, single episode, unspecified.	+8	×	✓
Explains why F32.9 may be appropriate when no additional severity, recurrence, or psychotic-feature details are provided.	+5	×	✓
Explains why it chooses the ICD code “F32.9 - Major depressive disorder, single episode, unspecified” for the patient.	+4	×	✓

G.4. Case study 4: Illustrative consultation



Figure 18. Illustrative HealthBench-Hard case study. Qwen3-4B-Instruct and InfiMed-ORBIT-4B are evaluated on the same medical consultation prompt, demonstrating how rubric-RL training leads to more context-seeking and safety-conscious responses. This example is shown only to show how well the models align with the benchmark rubric and must not be taken as clinical dosing guidance.

H. HealthBench-Hard versus HealthBench-Consensus

H.1. Semantic-space comparison of HealthBench-Hard and HealthBench-Consensus

To quantify the difficulty gap between benchmark subsets and check the integrity of the rubric pipeline, we compare the semantic topology of their rubrics. We embed all criteria with Qwen3-Embedding (4096 dimensions) (Zhang et al., 2025) and visualize the resulting manifold with t-SNE.

Homogeneity in Consensus (Fig. 19). As shown in Fig. 19, HealthBench-Consensus rubrics form tight, coherent clusters. This geometric structure reflects strong semantic homogeneity, indicating that the constraints are standardized and recurrent. This concentration is consistent with stronger baseline performance on this subset: the underlying reasoning patterns appear to be more standardized and reusable.

Sparsity in Hard (Fig. 20). In contrast, the HealthBench-Hard rubrics display a sparse and scattered structure, marked by wide gaps and isolated points. This semantic thinness reflects higher variability in constraints and increased reasoning complexity. Models must navigate a fractured logical landscape, where success depends not on reusing patterns but on adaptable, context-aware reasoning.

The structural divergence between Fig. 19 and Fig. 20 corroborate the integrity of the evaluation split. Although Consensus rubrics are used as few-shot examples during rubric generation, the Hard subset remains semantically separate, making it

unlikely that the observed improvements stem from straightforward memorization of the prompt examples.




Figure 19. t-SNE visualization of HealthBench-Consensus rubric embeddings. The dense clusters indicate high semantic homogeneity, consistent with stronger baseline performance on this subset and its suitability as a low-risk source of in-context exemplars.

H.2. Contamination diagnostics between NoHard and Hard

We investigate whether materials created from NoHard (including HealthBench-Consensus) might contaminate the Hard evaluation set. Since rubric-based pipelines can legitimately reuse rubric templates, similarity between rubric texts alone does not provide sufficient evidence of instance-level leakage. Accordingly, we focus our contamination analysis mainly at the *prompt* level, employing a combined prompt–rubric test along with checks for rubric redundancy within each set to distinguish true leakage from harmless template reuse.

Setup. Let $\mathcal{H} = \{h_i\}_{i=1}^{N_H}$ denote the Hard set and $\mathcal{N} = \{n_j\}_{j=1}^{N_N}$ denote the filtered NoHard set. Each element x is associated with a prompt $\text{Prompt}(x)$ and a collection of rubric criteria $\text{Criterion}(r_k(x))$. We convert these structured components into canonical string representations $q(x)$ (for the prompt) and $\rho(x)$ (for the rubrics) via a deterministic flattening procedure, then apply standard text normalization (lowercasing and whitespace collapsing) to obtain $\tilde{q}(x)$ and $\tilde{\rho}(x)$.

Near-duplicate similarity. We use character n -gram TF–IDF features (`char_wb`, $n \in \{3, 4, 5\}$, `min_df` = 2) with cosine similarity, estimating TF–IDF statistics on Hard and then applying the transformation to NoHard. For each $n \in \mathcal{N}$, we compute the nearest-neighbor similarity to Hard (direction $\mathcal{N} \rightarrow \mathcal{H}$):

$$S_Q(n) = \max_{h \in \mathcal{H}} \text{sim}(\tilde{q}(n), \tilde{q}(h)), \quad (6)$$

$$S_R(n) = \max_{h \in \mathcal{H}} \text{sim}(\tilde{\rho}(n), \tilde{\rho}(h)), \quad (7)$$

where S_Q targets instance leakage (prompt duplication), and S_R quantifies rubric similarity that may reflect templating.

Thresholded contamination curves. We summarize near-duplicate risk using thresholded rates

$$\text{Contam}_Q(\tau) = \frac{1}{N_N} \sum_{n \in \mathcal{N}} \mathbb{I}[S_Q(n) \geq \tau], \quad \text{Contam}_R(\tau) = \frac{1}{N_N} \sum_{n \in \mathcal{N}} \mathbb{I}[S_R(n) \geq \tau], \quad (8)$$




Figure 20. t-SNE visualization of HealthBench-Hard rubric embeddings. Hard rubrics form a sparser and more fragmented manifold than Consensus rubrics, reflecting greater semantic diversity and higher reasoning difficulty.

for thresholds $\tau \in [0.5, 1.0]$. Fig. 21(a–c) shows that S_Q is mostly concentrated at moderate similarity levels and that $\text{Contam}_Q(\tau)$ drops off quickly as τ increases, whereas S_R exhibits a much heavier tail at high similarity, consistent with the presence of rubric-style templating.

Joint test for instance leakage. To avoid over-interpreting cases where rubrics overlap, we also report a combined metric that requires both prompt similarity and rubric similarity to be high:

$$\text{JointContam}(\tau_Q, \tau_R) = \frac{1}{N_N} \sum_{n \in \mathcal{N}} \mathbb{I}[S_Q(n) \geq \tau_Q \wedge S_R(n) \geq \tau_R]. \quad (9)$$

With $(\tau_Q, \tau_R) = (0.90, 0.80)$, Fig. 21(d) shows that only a small number of points exceed both cutoffs, suggesting that strong rubric similarity rarely coincides with near-duplicate prompts under this filtering procedure.

Within-set rubric redundancy (template reuse). We characterize rubric templating within each split $\mathcal{S} \in \{\mathcal{H}, \mathcal{N}\}$ by

$$U_{\mathcal{S}}(x) = \max_{\substack{x' \in \mathcal{S} \\ x' \neq x}} \text{sim}(\tilde{\rho}(x), \tilde{\rho}(x')). \quad (10)$$

Fig. 21(e) reveals substantial redundancy in both splits, consistent with template reuse as a likely driver of the increased S_R .

Taken together, the prompt-similarity tail drops off rapidly at high thresholds (Fig. 21a–c), very few instances surpass both the prompt and rubric thresholds (Fig. 21d), and the rubrics themselves exhibit considerable redundancy within each split (Fig. 21e). These patterns support the conclusion that NoHard are unlikely to cause instance-level contamination of Hard. Given that rubrics are instantiated from shared templates, we view the resulting performance as generalization to unseen prompts within a common rubric framework, rather than robustness to completely new rubric formulations.

I. Negative Results and Failure Modes

For completeness we record three design choices that we tried and that did *not* work, together with our best explanation. Reporting these honestly is intended to help future practitioners avoid the same failure modes.




Figure 21. Contamination diagnostics (NoHard→Hard). (a) Distributions of prompt similarity S_Q and rubric similarity S_R . (b) ECDFs of S_Q and S_R . (c) Thresholded rates $\text{Contam}_Q(\tau)$ and $\text{Contam}_R(\tau)$. (d) Scatter of (S_Q, S_R) with operating thresholds $(\tau_Q, \tau_R) = (0.90, 0.80)$. (e) Within-set rubric redundancy via self max-similarity U_S for $S \in \{\mathcal{H}, \mathcal{N}\}$.

(a) Degenerate Regimes under Stringent Filtering Limits. Imposing overly strict $\text{Pass}@k$ constraints severely impedes policy exploration by inducing a sparse reward regime. For example, contracting the query band to $[\tau_q^{\text{low}}, \tau_q^{\text{high}}] = [0, 0.25]$ alongside a minor rubric threshold ($\tau_r = 0.10$) restricts data throughput to a mere 9% of available instances. Under this configuration, the variance-aware mask M_q discards another 35%–40% of rollout clusters due to vanishing variance, as all $G = 8$ generated paths collapse toward near-zero rewards on ultra-difficult tasks. This optimization gridlock forces the policy into a degenerate state characterized by stagnant reward trajectories and truncated defensive responses. Our production framework mitigates this fragility by expanding the moderate difficulty limit to $[0, 0.75]$ and deploying adaptive rubric thresholds $\tau_r \in \{0.25, 0.50, 0.75\}$, which retain 67%–84% of the empirical distribution. Mechanically, stabilizing open-ended medical RL requires careful calibration between *exploitable success signals* (easy wins) and *challenging safety frontiers* (hard mining).

SFT Overfitting Restricts Downstream Policy Entropy. The choice of learning rate during SFT critically impacts subsequent reinforcement learning dynamics. Table 17 compares ORBIT performance when initialized from SFT models optimized at $\eta \in \{10^{-5}, 10^{-6}, 10^{-7}\}$ on Baichuan-M2 distilled responses. While initializations at 10^{-7} and 10^{-6} reliably outperform the Instruct baseline (25.2 and 23.1 vs. 20.2), the 10^{-5} configuration regresses to 20.3 even after complete RL alignment. This regression stems from severe style overfitting: at 10^{-5} , the policy over-imitates the teacher’s verbose style and collapses the response distribution. Entering the RL stage with constrained effective entropy, this policy generates redundant rollouts that are aggressively discarded by the variance-aware filter, thereby stalling reward optimization. To contextualize this failure mode, consider an *emergency-referrals* case where a patient presents with symptoms indicating upper gastrointestinal bleeding (coffee-ground vomitus and melena). The three SFT variants yield qualitatively divergent initializations for the subsequent RL phase. The $\eta = 10^{-7}$ checkpoint yields a succinct answer that highlights the main risk and proposes two differential diagnoses. The $\eta = 10^{-6}$ model adds a sequence of structured clinical actions yet stays compact (~ 220 tokens) and allows for flexible options. By contrast, the overfitted $\eta = 10^{-5}$ checkpoint generates a rigid, 450-token boilerplate, marked by repetitive formatting patterns that ignore clinical specifics. During subsequent RL, more than 72% of rollouts from this 10^{-5} starting point converge to this same template. As a result, the variance-sensitive filter discards most of these trajectories because their reward variance collapses, leaving GRPO gradients to be computed on only a small, low-information subset. We therefore adopt 10^{-7} as the standard SFT learning rate.

(c) GPT-5-Chat produces less discriminative medical rubrics under our prompt. We initially expected a frontier closed model to produce the strongest rubrics. In practice, GPT-5-Chat (Tab. 7) is one of the weaker rubric generators in our sweep, with downstream HealthBench-Hard scores below DeepSeek-R1, Gemini-2.5-Pro and GPT-OSS-120B. Manual inspection trace this to GPT-5-Chat’s safety post-training, which is unwilling to enumerate failure modes such as “Recommends NSAID despite suspected GI bleed” as criteria, treating the negative-criterion request as adversarial content. The GPT-5-Chat rubrics thus skew toward generic positive constraints and lose the discriminative power that the negative criteria provide. Thus, we adopt DeepSeek-R1 as the default rubric generator.

J. Extended Related Work

Process Reward Models and Step-Level Supervision. Recent work increasingly supplements scalar outcome rewards with intermediate, step-level supervision, exemplified by Process Reward Models (PRMs) for mathematical reasoning (Shao et al., 2024; Zheng et al., 2025), verifier-based reinforcement learning for code generation (Yu et al., 2026), and structural-format rewards in instruction tuning (Viswanathan et al., 2026; Jacob Dineen et al., 2025). ORBIT diverges from this line of research by employing criteria that assess *final response outcomes* rather than intermediate reasoning steps. Consequently, ORBIT remains applicable to open-ended generation tasks lacking clearly defined intermediate correctness, though at the expense of sacrificing partial locality of credit assignment provided by PRMs. We consider these two approaches to be complementary rather than mutually exclusive.

Rubric- and Checklist-Based Alignment. Integration of rubric-style criteria into policy optimization represents a rapidly evolving alignment frontier. This section delineates the fine-grained algorithmic choices that separate ORBIT from concurrent alternative frameworks. First, in terms of criterion generation, contemporary pipelines such as those of Gunjal et al. (2025) enforce static, expert-curated evaluation templates uniformly in training sets. In contrast, ORBIT dynamically instantiates localized rubric sets via a customized RAG workflow, suppressing semantic degeneration and enhancing case-specific resolution. Second, in terms of reward computation, existing multidimensional architectures like those proposed by Bhaskar et al. (2025) and the ACE framework (Chen et al., 2025) require learning intermediate dense reward models from factorized preference feedback. ORBIT simplifies this pipeline by utilizing non-parameterized, discrete indicator

functions. This direct scoring protocol eliminates the systemic overhead and instability associated with surrogate reward model training. Finally, while closely allied with modern single-turn checklist verification paradigms (Viswanathan et al., 2026; Jacob Dineen et al., 2025), ORBIT uniquely scales this architecture to multi-turn medical scenarios.

Medical LLM Alignment. Compared with prior specialized medical LLMs such as HuatuoGPT (Chen et al., 2024), m1 (Huang et al., 2025), AlphaMed (Liu et al., 2025a), and MedReason (Wu et al., 2025), ORBIT does not introduce a new domain-specific base model. Instead, it serves as a post-training alignment strategy capable of elevating a general-purpose 4B model above specialized 32B medical-domain baselines on HealthBench-Hard (Tab. 1). Relative to the recent Baichuan-M2-32B (Dou et al., 2025), ORBIT achieves comparable HealthBench-Hard performance at roughly one-eighth the parameter scale. This demonstrates that carefully structured rubric-based rewards can effectively substitute for extensive parameter scaling in alignment-intensive tasks. We regard this outcome as a promising indication of the potential synergy between rubric-guided reinforcement learning and future model scaling strategies.