

ManiSoft: Towards Vision-Language Manipulation for Soft Continuum Robotics

Ziyu Wei^{*12} Luting Wang^{*12} Chen Gao^{◇†13} Li Wen^{†1} Si Liu^{†12}

Project page: <https://buaa-colalab.github.io/ManiSoft>




Figure 1. (a) Rigid arms operate in a low-dimensional action space and can fail due to limited shape adaptation, whereas soft arms driven by distributed low-level actuation can continuously deform to reach around obstacles. (b) Example expert trajectories for the four *ManiSoft* tasks. (c) The data generation pipeline comprises an asset library, clean and randomized scene generation, hierarchical trajectory generation, simulation, and quality check. (d) Performance of representative policy models on *ManiSoft* in clean scenarios.

Abstract

Most existing vision-language manipulation research targets rigid robotic arms, whose fixed morphology limits adaptability in cluttered or confined spaces. Soft robotic arms offer an appealing alternative due to their deformability, but confront challenges such as unreliable proprioception and distributed low-level actuation. To investigate these challenges, we introduce *ManiSoft*, a benchmark for vision-language manipu-

lation with soft arms. *ManiSoft* features a tailored simulator that couples realistic soft-body dynamics with contact-rich interactions via an elastic force constraint. On this basis, *ManiSoft* defines four tasks, each highlighting distinct aspects of deformable control, from basic end-effector coordination to obstacle avoidance. To support policy training and evaluation, *ManiSoft* includes an automated pipeline that generates 6,300 diverse scenes and corresponding expert trajectories. To produce high-quality trajectories at scale, we first employ a high-level planner to decompose each task into a sequence of waypoints, followed by a low-level reinforcement learning policy that generates torque commands to track waypoints. Benchmarking three representative policy models shows relatively promising results in clean scenes but substantial per-

* Equal contribution. ◇ Project lead. † Corresponding author. ¹Beihang University ²Hangzhou Innovation Institute, Beihang University ³National University of Singapore. Correspondence to: Chen Gao <gaochen.ai@gmail.com>, Li Wen <liwen@buaa.edu.cn>, Si Liu <liusi@buaa.edu.cn>.

formance drop under randomization. Visualization analysis indicates that failures stem primarily from inaccurate visual estimation of proprioceptive state and limited exploitation of deformability for adaptive obstacle avoiding. We anticipate *ManiSoft* to serve as a valuable testbed, bridging the gap between rigid and soft arms in the context of vision-language manipulation.

1. Introduction

Vision-language manipulation (Shao et al., 2025) is a central capability of embodied AI, enabling language-conditioned interaction with the physical world. To date, most benchmarks (Liu et al., 2023; Yu et al., 2021; Mees et al., 2022; Li et al., 2024b; Srivastava et al., 2022) and methods (Chi et al., 2025; Liu et al., 2025b; Kim et al., 2024; 2025) focus on rigid robotic arms, where accurate proprioception and low-dimensional kinematics enable straightforward perception-to-control pipelines. However, rigid morphologies impose fundamental limitations in cluttered or confined environments (Chen et al., 2025b). As illustrated in Figure 1 (a), when obstacles require significant shape adaptation, a rigid arm may fail to reach the target due to its joint constraints.

Soft robotic arms (Xie et al., 2023; Armanini et al., 2023; Majidi, 2014; Hughes et al., 2016; Zhao et al., 2024), built from elastic materials, offer an appealing alternative. Through continuous deformation, soft arms can adapt their geometry to execute policies infeasible for rigid arms. However, these advantages come with major challenges for vision-language manipulation. Unlike rigid arms with reliable joint sensing, soft arms often lack accurate proprioception (Pagliarini et al., 2025), leading to highly complex kinematic control. Therefore, soft arms are typically actuated via low-level commands (e.g., pressures (Liu et al., 2025a), tendon tensions (Walker et al., 2024), or torques (Caasenbrood et al., 2022)) rather than intuitive kinematic targets. Moreover, distributed actuation along the body yields a higher-dimensional and more coupled action space than rigid arms. Together, these factors complicate the generation of stable and coordinated behaviors.

In this work, we introduce *ManiSoft*, a benchmark designed to catalyze vision-language manipulation research for soft arms. *ManiSoft* provides (i) a simulation and rendering stack for soft-arm manipulation, (ii) a suite of language-conditioned tasks with diverse scenes, and (iii) expert demonstration trajectories to support imitation and offline reinforcement learning. This design allows for evaluating existing policy models with minimal modifications, exposing failure modes unique to deformable embodiments.

A central technical challenge lies in simulation. While

many soft-body simulators (Huang et al., 2021; Faure et al., 2012; Wei et al., 2025) accurately capture elastic dynamics, they offer limited support for environmental interactions (e.g., contact and friction). In contrast, rigid-body simulators (Geng et al., 2025; Xiang et al., 2020; Szot et al., 2021) excel at modeling interactions but struggle with continuous deformation. To bridge this gap, we integrate a soft-body dynamics simulator (Naughton et al., 2021) with a rigid-body interaction simulator (Todorov et al., 2012) through an elastic force constraint, facilitating contact-rich manipulation with a soft arm. We further provide a Blender-based¹ renderer for generating the visual observations used by policy models.

Built on this stack, *ManiSoft* defines four tasks as illustrated in Figure 1 (b). For each task, we follow the automated pipeline in Figure 1 (c) to generate tabletop scenes and expert trajectories. The asset library contains 263 3D objects, annotated with candidate manipulation poses. We first construct clean (uncluttered) scenes by sampling target objects from the asset library and then create randomized variants by adding obstacles and varying object placements and textures, enabling systematic evaluation under increasing visual and physical complexity.

For each scenario, we generate expert trajectories using a hierarchical mechanism. A high-level planner produces a sequence of waypoints, where each waypoint specifies a 6-DoF end-effector pose. A low-level controller outputs torque commands to drive the soft arm between successive waypoints. In our implementation, the high-level planner uses task-specific rules, and the low-level controller is a reinforcement learning (RL) (Sutton, 1988) policy. This decomposition mitigates the difficulty of directly producing torque sequences and yields stable trajectories for training.

Finally, we benchmark representative policy models on *ManiSoft*. As summarized in Figure 1 (d), existing models can solve a subset of tasks in clean scenes. However, performance drops substantially in randomized settings. Our visualizations and failure-case analysis suggest two bottlenecks: (i) estimating the soft arm’s proprioceptive states from visual observations, and (ii) exploiting deformability to plan obstacle-avoiding interaction strategies. We hope *ManiSoft* will serve as a testbed for developing methods that address these challenges.

Conflict of Interest Disclosure. We declare that we have no relevant or material financial interests that relate to the research described in this paper.

¹<https://www.blender.org/>

2. Related Works

Robotic Manipulation Benchmarks. Progress in vision-language manipulation has been accelerated by benchmarks that standardize tasks, observations, and evaluation protocols for comparing policy models (Shridhar et al., 2020; Ahmed et al., 2020; Qi et al., 2020). RLBench (James et al., 2020) provides a suite of 100 vision-based manipulation tasks for evaluating both learning-based and traditional policy models. The ManiSkill series (Mu et al., 2021; Gu et al., 2023; Tao et al., 2025) emphasizes generalizable manipulation over diverse objects in a full-physics simulator. CALVIN (Mees et al., 2022) targets long-horizon, language-conditioned manipulation, while LIBERO (Liu et al., 2023) studies cross-task transfer in lifelong learning. RoboVerse (Geng et al., 2025) supports evaluation across multiple simulators and robot embodiments. RoboTwin (Mu et al., 2024; Chen et al., 2025a) proposes an automated pipeline for generating diverse dual-arm manipulation scenarios at scale. Despite their breadth, these benchmarks predominantly target rigid arms with low-dimensional kinematics and reliable proprioception, leaving vision-language manipulation for deformable embodiments relatively underexplored. We fill this void by introducing a benchmark tailored to vision-language manipulation with soft arms.

Vision-Language-Action Models. Vision-Language-Action (VLA) models have advanced rapidly in recent years. RT-1 (Brohan et al., 2023b) and RT-2 (Brohan et al., 2023a) demonstrate the effectiveness of large-scale training by leveraging multi-robot datasets. DexVLA (Wen et al., 2025) extends this paradigm to enhance efficiency and generalization in long-horizon manipulation. RDT-1B (Liu et al., 2025b) introduces a diffusion-based foundation model for bimanual manipulation. OpenVLA (Kim et al., 2024) presents an open-source framework built upon large language models and pretrained visual encoders, while CogACT (Li et al., 2024a) proposes an action module conditioned on VLM outputs to improve action prediction. More recently, the π series (Black et al., 2025b;a; Amin et al., 2025) has demonstrated strong performance via large-scale pretraining followed by reinforcement learning. Despite this progress, prior VLA methods have been predominantly developed and evaluated on rigid arms. Our work provides a comprehensive benchmark and systematic evaluation of representative policy models on soft arms, highlighting unique challenges absent in rigid arms.

Soft Robotic Arms. Over the past decade, soft arms have been widely studied and applied in domains such as biomedical engineering (Cianchetti et al., 2018; Rogatinsky et al., 2023), aerospace (Ruiz et al., 2024; Szász et al., 2022), and underwater exploration (Gong et al., 2021; Li




Figure 2. **Overview of the soft arm modeling in our Simulator.** The soft body is modeled as a Cosserat rod that moves under the influence of an external torque τ_e . Interaction between soft body and EEF is represented via an elastic force constraint. Relative displacement Δx or relative rotation $\Delta \theta$ between them induces corresponding restoring forces and torques.

et al., 2023). To address the challenges of modeling and control, learning-based approaches have been extensively explored. Thuruthel et al. (2017) applies trajectory optimization for open-loop predictive control, while Thuruthel et al. (2018) extends this framework using model-based reinforcement learning for closed-loop control. Centurelli et al. (2022) develops a controller based on LSTM and TPRO for dynamic trajectory tracking, both with and without payloads. Soft DAgger (Nazeer et al., 2023) enables sample-efficient imitation learning for soft control. While these efforts advance low-level control of soft arms, high-level vision-language reasoning for manipulation remains largely unaddressed. Our work studies vision-language manipulation with soft arms, which requires jointly reasoning about visual perception, language understanding, environmental interaction, and deformable control.

3. The ManiSoft Benchmark

We introduce the *ManiSoft* benchmark, designed to support vision-language manipulation with soft robotic arms. It comprises a soft-arm simulator, a collection of diverse tabletop scenes paired with language instructions, and expert demonstration trajectories. To enable scalable data collection, we propose an automated data generation pipeline that integrates procedural scene construction with a hierarchical expert trajectory generation mechanism.

3.1. Simulator

While existing soft-body simulators faithfully capture elastic deformation, they often provide limited support for interactions with rigid environments. Conversely, rigid-body simulators excel in modeling contacts and friction but lack native continuous deformation. To bridge this gap for soft robotic manipulation, we develop a hybrid simulator that combines accurate deformable dynamics with robust environmental interactions.

As illustrated in Figure 2, we model the soft arm as two coupled components: a deformable soft body and an end-effector. These components are linked via an elastic force constraint to ensure coordinated yet compliant motion.

The soft body is simulated using *Elastica* (Naughton et al., 2021), which discretizes the arm into N segments following the Cosserat rod theory (Cosserat & Cosserat, 1909). External actuation torques $\tau_e \in \mathbb{R}^{N \times 3}$ induce axial, shear, bending, and torsional strains along the rod, producing internal forces $\mathbf{f}_i \in \mathbb{R}^{N \times 3}$ and moments $\tau_i \in \mathbb{R}^{N \times 3}$. These forces and torques govern the deformation together. Appendix A provides further details on Cosserat rod theory.

The end-effector and its interactions with the environment are handled by MuJoCo (Todorov et al., 2012), enabling efficient and stable simulation of contact-rich scenarios.

To couple the soft body and the end-effector, we impose an elastic force constraint. The two components are connected by a stretchable and twistable virtual spring with zero rest length. Relative translations $\Delta \mathbf{x} \in \mathbb{R}^3$ and rotations $\Delta \theta \in \mathbb{R}^3$ between the attachment points generate restoring force $\mathbf{F} \in \mathbb{R}^3$ and torque $\mathbf{M} \in \mathbb{R}^3$, computed according to Hooke’s law:

$$\mathbf{F} = -k_F \Delta \mathbf{x}, \quad \mathbf{M} = -k_M \Delta \theta, \quad (1)$$

where $k_F, k_M \in \mathbb{R}$ are the translational and rotational stiffness coefficients, respectively. These restoring terms penalize relative motion between the soft body and the end-effector, driving the system toward coordinated motion.

We employ Blender to render the visual observations. Based on the simulated states, Blender produces RGB images of the tabletop scenes from fixed camera viewpoints, including the soft arm, target objects, and surrounding obstacles. Rendering parameters are detailed in Appendix B.

3.2. Task Definition

As shown in Figure 1 (b), *ManiSoft* defines four manipulation tasks, each designed to highlight distinct challenges in vision-language manipulation for soft arms. Collecting (**COLL**) involves guiding the soft arm to gather a designated object and deposit it into a container, thereby evaluating fundamental trajectory control and basic end-effector coordination in policy models. Building on this foundation, Alignment (**ALN**) demands precise positioning of the target object to a specified 6-DoF pose, testing the model’s capability for fine-grained orientation adjustments. Stacking (**STK**) escalates the challenge by requiring the arm to assemble tableware items from largest to smallest in a stable vertical pile, which assesses precision control during elevated, contact-rich interactions. Finally, Arrangement (**ARR**) requires placing objects according to a specified spatial configuration, thereby demanding integrated visual




Figure 3. **Scene generation in *ManiSoft*.** Objects are sampled from the asset library to create a clean scene, and randomized scenes are generated by injecting objects as obstacles and varying surface textures. Instructions are produced with the descriptions of relevant objects. In the randomized setting, diverse descriptions are leveraged to enhance the linguistic richness.

perception, spatial reasoning, and obstacle avoidance.

Formally, at each time step $t \in \mathbb{N}^+$, given an instruction $\mathbf{L}$ and the current visual observation $\mathbf{V}_t$, the policy model predicts the next action $\mathbf{A}_t = (\tau_e, S)$. Here, τ_e is the external torques and $S \in \{0, 1\}$ indicates the end-effector state. Upon execution, a new observation $\mathbf{V}_{t+1}$ is rendered, and the policy model proceeds autoregressively until task completion or the maximum horizon T is reached.

Unlike rigid-arm benchmarks, which typically provide proprioceptive state (e.g., joint angles), *ManiSoft* deliberately excludes internal soft-body states to reflect real-world sensing limitations. Policy models must therefore infer the arm’s configuration and deformation solely from visual observations, introducing significant challenges in proprioceptive state estimation and deformable strategy planning.

Two evaluation metrics are used: (i) success rate, determined by task-specific criteria, and (ii) efficiency, measured as the number of steps required for completion.

3.3. Scene Generation

As depicted in Figure 3, *ManiSoft* adopts a tabletop environment as its core setting. The soft arm is fixed behind the table and centered relative to the workspace, allowing full access to objects distributed across the table sur-




Figure 4. Trajectory generation pipeline in ManiSoft. (a) An executor is trained via RL policy to transform waypoint (6-DoF pose) into torques. (b) RL rewards are designed to balance accuracy and stability, consisting of a pose difference reward R_d negatively correlated with the pose difference, and a stability reward R_s that penalizes or rewards changes in pose difference. (c) Task-specific rules are predefined to produce high-level planning (trajectory waypoints) for each case, which are then converted into low-level actions (torque commands) by the executor to generate complete trajectories.

face. We build our object library by leveraging assets from RoboTwin-OD (Chen et al., 2025a). For each object, we pre-annotate a set of suitable 6-DoF end-effector poses for interaction with the soft arm (e.g., approach, grasp, and lift candidates). These annotations guide the high-level planner during expert trajectory generation. Scenes are procedurally generated by randomly sampling objects from this library and placing them within the workspace.

To support systematic evaluation, each task includes two difficulty levels: clean and randomized. In the clean setting, scenes contain only the task-relevant target objects in fixed layouts and appearances. In the randomized setting, we introduce additional irrelevant objects as obstacles to enhance spatial complexity, and apply scene randomization by sampling diverse textures along with variations in lighting intensity and brightness. This yields diverse spatial arrangements and visual appearances across episodes, stressing generalization in perception and planning.

Language instructions are generated in a controlled manner to ensure diversity, semantic accuracy, and consistency. Direct LLM sampling often produces variable phrasing or minor hallucinations; therefore, we first generate candidate instructions via GPT, manually curate and refine them into a template library, then instantiate templates by filling in object attributes (e.g., color, shape, material). In clean scenes, each object receives a single canonical description. In randomized scenes, objects are paired with multiple attribute-aware descriptions to reflect visual variability (e.g., a bottle may be referred to as “yellow bottle”, “bottle with green cap”, or “tall plastic bottle”).

3.4. Trajectory Generation.

Given the procedurally generated scenes for each task, we produce expert trajectories using a hierarchical mechanism, as illustrated in Figure 4. At the high level, a task-specific rule-based planner generates a sequence of waypoints, where each waypoint defines a desired 6-DoF end-effector pose in $SE(3)$. These waypoints encode semantically intermediate configurations (e.g., approach, grasp, retract) tailored to the task, avoiding the need for direct torque-sequence planning over long horizons.

At the low level, an RL-trained executor drives the soft arm from its current configuration to each successive waypoint using torque actuation. At each timestep t , the executor receives the following inputs: (i) the target end-effector pose $\hat{P} \in SE(3)$, (ii) proprioceptive states including positions and velocities of selected segments along the arm, and (iii) the current end-effector pose $P \in SE(3)$, and outputs torque commands.

The executor is trained with a dense reward function that encourages precise reaching of the target waypoint while promoting stable convergence and penalizing excessive deformation or oscillation. We measure the pose discrepancy using the standard $SE(3)$ logarithm map:

$$[\mathbf{d}_p, \mathbf{d}_r] = \log(P^{-1}\hat{P}), \quad d = \|\mathbf{d}_p\|_2 + \alpha\|\mathbf{d}_r\|_2, \quad (2)$$

where $\mathbf{d}_p, \mathbf{d}_r \in \mathbb{R}^3$ represent the position and axis-angle rotation differences, respectively. The scalar $\alpha > 0$ is tuned to balance the contributions of translation and rotation.

We define the total reward R as the sum of two terms: a pose difference term R_d and a stability term R_s . The pose difference reward R_d is adapted from *Elastica-RL-Control* (Naughton et al., 2021), using the pose distance




Figure 5. **Statistical analysis of the *ManiSoft* Benchmark.** (a) Distribution of trajectory lengths. Tasks in *ManiSoft* generally involve long trajectories, with the STK task exhibiting notably longer trajectories than the others. (b) Frequency distribution of target object categories, highlighting the diversity of manipulable objects in *ManiSoft*. (c) Spatial distribution of initial target object positions on the tabletop, showing that graspable targets are broadly and evenly distributed across the workspace.

d instead of Euclidean distance:

$$R_d = -d + k_1 \mathbb{1}_{\{d < d_1\}} + k_2 \mathbb{1}_{\{d < d_2\}}. \quad (3)$$

The stability reward R_s provides a signal based on the rate of change of d when the end-effector is close to the target pose, encouraging smooth and stable convergence:

$$R_s = \begin{cases} -\text{sgn}\left(\frac{\partial d}{\partial t}\right) \cdot \beta, & d \leq D, \\ 0, & d > D, \end{cases} \quad (4)$$

where $\beta > 0$ is a scaling factor that controls the strength of the stability incentive.

Once trained, the executor is used to roll out complete trajectories by sequentially tracking the high-level waypoints. This hierarchical decomposition produces stable, collision-free demonstrations across a wide range of scenes. It significantly simplifies downstream policy learning compared to training directly on raw torque actions.

Empirically, the executor reaches a success rate of 54% on 100 random samples. We further investigate the effect of the stability reward R_s on control stability. As shown in Figure 6, the model trained with R_s exhibits noticeably smaller fluctuations in pose difference compared to the model trained without R_s . We also perform an ablation study on different parameter settings of R_s . The variance of the pose difference between the end-effector and the target is used to quantify control stability. Based on this metric, we select the best-performing set of parameters, $\beta = 1$, $D = 0.3$, as reported in Table 1. Additional ablation studies are provided in Appendix C.

3.5. Data Statistic

ManiSoft contains 6,300 scene-trajectory pairs, comprising 2,100 clean scenes and 4,200 randomized scenes, with an




Figure 6. **Visualization of executor trained w/ and w/o the stability reward R_s .** (Left) The pose difference between the end-effector and the target pose over time. (Right) The soft robotic arm’s trajectory shadows during the final 1000 simulation steps. The red line represents the soft body, and the blue circle indicates the end-effector.

average of 40 language instructions per scene. The dataset is split into training and testing sets with a ratio of 4 : 1.

Owing to the high precision of torque-based control, trajectories in *ManiSoft* are relatively long, with an average length of 1,272 steps; the distribution of trajectory lengths is shown in Figure 5 (a). *ManiSoft* features a rich variety of objects, including 109 manipulable objects across 17 categories, as illustrated in Figure 5 (b), and 154 obstacles spanning 35 categories. Figure 5 (c) visualizes the initial positions of target objects (i.e. grasp poses) across all scenes as a heatmap over the tabletop, demonstrating a wide distribution of grasp positions across the entire table.

Examples of the generated data are shown in Figure 7.




Figure 7. **Visualization of the ManiSoft.** The left example is a clean scene, while the right is a randomized scene.

Table 1. Control stability under different parameters of R_s . Specifically, $\beta = 0$ indicates the absence of R_s .

$D \backslash \beta$	0	0.5	1	1.5
0.05	0.176	0.157	0.074	0.121
0.1	0.176	0.149	0.153	0.071
0.2	0.176	0.070	0.135	0.064
0.3	0.176	0.145	0.053	0.091
Average	0.176	0.130	0.104	0.087

More visualizations are provided in Appendix F.

4. Experiments

We evaluate three representative models on *ManiSoft*: Diffusion Policy (DP) (Chi et al., 2025), RDT (Liu et al., 2025b), and OpenVLA-OFT (Kim et al., 2025). DP and RDT are trained from scratch, while OpenVLA-OFT is fine-tuned with LoRA. Implementation details are provided in Appendix D.

4.1. Main Results

Table 2 shows the performance of different models on each task under clean and randomized settings. Overall, DP and OpenVLA-OFT achieve substantially better performance than RDT. Specifically, DP attains an average success rate of 31.6% with a mean execution length of 520 steps, while OpenVLA-OFT achieves a comparable success rate of 30.4% with an average of 527 steps. In contrast, RDT performs significantly worse, with an average success rate of only 9.2% despite requiring 496 steps on average. This performance gap is likely attributable to differences in model capacity. RDT contains approximately 1B parameters, whereas DP and OpenVLA-OFT each have around 400M parameters. As a result, RDT is more prone to overfitting the training data, leading to inferior generalization on the testing set. Across all three models, the best performance is consistently observed on the COLL task, which requires less precise orientation perception and spatial reasoning compared to the other tasks.

Under the clean setting, DP achieves higher accuracy than OpenVLA-OFT on the COLL task, exceeding it by 17.6%.

However, OpenVLA-OFT outperforms DP on all remaining tasks. Specifically, it achieves improvements of 6.7%, 5.0%, and 1.3% on the ALN, STK, and ARR tasks, respectively. These results suggest that DP is more effective for simpler tasks, whereas OpenVLA-OFT exhibits stronger reasoning and generalization capabilities in more complex scenarios, likely benefiting from its pretrained weights.

Under the randomized setting, all models experience a decrease in success rate. Specifically, DP exhibits the largest performance drop, with its success rate decreasing by 29.4%, while RDT and OpenVLA-OFT show more moderate declines of 7.6% and 3.4%, respectively. This indicates that the introduction of obstacles, together with variations in language instructions and scene configurations, substantially increases task difficulty. Notably, unlike in the clean setting, OpenVLA-OFT consistently outperforms DP under randomization, achieving an average improvement of 13.1%. This suggests that OpenVLA-OFT maintains stronger generalization performance in the presence of environmental and instruction-level variations.

Table 3 illustrates the performance across different object categories. The Rubik’s Cube consistently yields the highest success rates among all objects. In the clean setting, its success rate exceeds that of the other categories by 5%–30% across methods. Notably, under the randomized setting, OpenVLA-OFT achieves a 15.0% success rate on the Rubik’s Cube, which is twice that of others. In contrast, the shoe is the most challenging object. In the clean setting, its success rate is 5%–30% lower than other categories across models. This gap becomes more pronounced under randomization, where success rates drop below 10% for all methods. These results indicate that while object geometry strongly affects absolute task difficulty, different models exhibit consistent relative performance trends across object categories. More results are provided in Appendix E.

4.2. Analysis

By visualizing rollouts from the evaluated policy models, we identify three failure modes.

Proprioceptive state ambiguity. Reliable torque control depends on precise proprioceptive state estimation. Soft-

Table 2. The results of the policy under the clean and randomized settings in *ManiSoft*. ACC denotes the task success rate on the eval set, and #Steps represents the number of inference steps required by the model to complete each task. OFT denotes OpenVLA-OFT.

Method	COLL		ALN		STK		ARR		Average	
	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps
Clean										
DP (Chi et al., 2025)	63.0	547	18.3	442	15.0	517	30.0	573	31.6	520
RDT (Liu et al., 2025b)	13.8	509	11.7	463	10.0	803	1.3	210	9.2	496
OpenVLA-OFT (Kim et al., 2025)	45.4	565	25.0	472	20.0	492	31.3	578	30.4	527
Randomized										
DP (Chi et al., 2025)	3.8	521	1.7	324	2.5	818	0.6	790	2.2	613
RDT (Liu et al., 2025b)	1.2	487	4.2	379	0.0	-	1.3	238	1.6	368
OpenVLA-OFT (Kim et al., 2025)	32.7	601	26.7	489	35.0	563	13.7	563	27.0	554

Table 3. The results of the policy on each manipulable objects in the ARR task.

Method	Rubik’s Cube		Bottle		Pen Cup		Shoe		Average	
	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps
Clean										
DP (Chi et al., 2025)	50.0	705	25.0	454	25.0	602	20.0	534	30.0	573
RDT (Liu et al., 2025b)	5.0	210	0.0	-	0.0	-	0.0	-	1.3	210
OpenVLA-OFT (Kim et al., 2025)	40.0	667	30.0	430	35.0	525	20.0	690	31.3	578
Randomized										
DP (Chi et al., 2025)	0.0	-	0.0	324	-	818	2.5	790	0.6	790
RDT (Liu et al., 2025b)	0.0	-	2.5	174	2.5	302	0.0	-	1.3	238
OpenVLA-OFT (Kim et al., 2025)	15.0	728	7.5	482	25.0	507	7.5	536	13.7	563




Figure 8. Visualization of typical failure cases of OpenVLA-OFT in *ManiSoft*. (a) The robot exhibits unexpected torsion and internal forces, resulting in inaccurate action prediction. (b) The robot fails to reach behind the obstacle.

arm deformation induces internal torques that must be actively compensated. Only the residual torque can drive the arm toward a desired pose. When the compensation term dominates, small state-estimation errors can overwhelm this residual, yielding unreliable control. As shown in Figure 8 (a), the target object lies close to the arm base, requiring a large bend to reach it. This deformation induces substantial internal torques. The policy model fails to compensate for these loads, leaving insufficient residual control to stabilize the motion. Consequently, the end-effector




Figure 9. Comparison of DP and OpenVLA-OFT on the same task: (a) DP successfully completes the task; (b) OpenVLA-OFT exhibits the “stop-moving” behavior.

drifts laterally and ultimately fails to reach the target object.

Challenges in leveraging soft arm compliance. Compared to rigid arms, soft arms offer advantages in flexibility, allowing them to adapt their shape to the environment and reach behind obstacles. However, as illustrated in Figure 8 (b), rather than adapting its shape to reach behind the obstacle, the policy model extends the soft arm directly toward the target object, resulting in collisions with the obstacle. This suggests that the policy model has not effectively utilized the soft-specific capabilities, such as shape adaptation and passive compliance. Increasing the proportion of

obstacle-specific expert data or incorporating physical priors during training may help to mitigate this limitation.

Stop-Moving Behavior. When comparing DP and OpenVLA-OFT, we observe that OpenVLA-OFT can exhibit a “stop-moving” behavior after grasp completion, where the robot remains stationary and fails to initiate subsequent actions. This behavior is likely caused by subtle visual changes during grasping, which induce a feedback loop that suppresses further action generation. As shown in Figure 9, OpenVLA-OFT often stops moving after a successful grasp, whereas DP rarely encounters this issue. This helps explain why OpenVLA-OFT achieves a lower success rate than DP on the simpler COLL task (45.4% vs. 63.0%) and requires longer execution lengths (565 vs. 547 steps, Table 2). Overall, this highlights a key distinction between diffusion-based and deterministic policies: the stochasticity in diffusion-based policies enables escaping such feedback loops, while deterministic policies are more prone to repetitive behavior.

Conclusion

We introduced *ManiSoft*, a benchmark for vision-language manipulation with soft arms. *ManiSoft* features a tailored simulator that couples soft-body dynamics with interactions via an elastic force constraint. Four tasks are designed to highlight distinct challenges in deformable control. An automated pipeline generates 6,300 diverse scenes and corresponding expert trajectories. Quality of the trajectories is ensured through a hierarchical mechanism that combines waypoint decomposition with RL-based torque control. Benchmarking representative policy models shows relatively promising performance in clean scenes but marked degradation under randomization. Failures are primarily attributed to inaccurate visual estimation of proprioceptive state and under-exploitation of deformability.

Limitations

This work represents an initial step toward benchmarking vision-language manipulation for soft robots. The current setup mainly focuses on a subset of actuation mechanisms and relatively simple tabletop scenarios, and does not yet cover more diverse actuation types or more dynamic, long-horizon tasks. In addition, sim-to-real consistency and physical validation could be further strengthened. These aspects leave room for future improvements in terms of coverage, realism, and evaluation diversity.

Impact Statement

This work can support the development of safer and more human-friendly robotic systems. By providing a bench-

mark for vision-language manipulation with soft robots, ManiSoft may facilitate research on compliant interaction, with potential applications in service robotics and assistive or medical settings. Overall, it contributes toward more adaptable and accessible robotic technologies.

Acknowledgements

This research is supported in part by the Key Research Program of Hangzhou (No. 2025SZD1A56), the National Natural Science Foundation of China (No. 62461160308, U23B2010, 62576024), the Beijing Natural Science Foundation (No. L231011), the Fundamental Research Funds for the Central Universities (No. 501RCQD2025141003), BeiHang GanWei Project (No. 502GWXM2024141001), the National Science Foundation Support Projects (No. 62425303), and the National Key R&D Program of China (No. 2024YFb4707300).

References

- Ahmed, O., Träuble, F., Goyal, A., Neitz, A., Bengio, Y., Schölkopf, B., Wüthrich, M., and Bauer, S. Causalworld: A robotic manipulation benchmark for causal structure and transfer learning. *arXiv preprint arXiv:2010.04296*, 2020.
- Amin, A., Aniceto, R., Balakrishna, A., Black, K., Conley, K., Connors, G., Darpinian, J., Dhabalia, K., DiCarlo, J., et al. $\pi_{0.6}^*$: a vla that learns from experience. *arXiv preprint arXiv:2511.14759*, 2025.
- Armanini, C., Boyer, F., Mathew, A. T., Duriez, C., and Renda, F. Soft robots modeling: A structured overview. *T-RO*, 39(3):1728–1748, 2023.
- Black, K., Brown, N., Darpinian, J., Dhabalia, K., Driess, D., Esmail, A., Equi, M., Finn, C., Fusai, N., et al. $\pi_{0.5}$: a vision-language-action model with open-world generalization. In *CoRL*, 2025a.
- Black, K., Brown, N., Driess, D., Esmail, A., Equi, M., Finn, C., Fusai, N., Groom, L., Hausman, K., Ichter, B., Jakubczak, S., Jones, T., Ke, L., Levine, S., Li-Bell, A., Mothukuri, M., Nair, S., Pertsch, K., Shi, L. X., Tanner, J., Vuong, Q., Walling, A., Wang, H., and Zhilinsky, U. π_0 : A vision-language-action flow model for general robot control. In *RSS*, 2025b.
- Brohan, A., Brown, N., Carbajal, J., Chebotar, Y., Chen, X., Choromanski, K., Ding, T., Driess, D., Dubey, A., Finn, C., Florence, P., Fu, C., Arenas, M. G., Gopalakrishnan, K., Han, K., Hausman, K., Herzog, A., Hsu, J., Ichter, B., Irpan, A., Joshi, N., Julian, R., Kalashnikov, D., Kuang, Y., Leal, I., Lee, L., Lee, T.-W. E., Levine, S., Lu, Y., Michalewski, H., Mordatch, I., Pertsch, K.,

- Rao, K., Reymann, K., Ryoo, M., Salazar, G., Sanketi, P., Sermanet, P., Singh, J., Singh, A., Soricut, R., Tran, H., Vanhoucke, V., Vuong, Q., Wahid, A., Welker, S., Wohlhart, P., Wu, J., Xia, F., Xiao, T., Xu, P., Xu, S., Yu, T., and Zitkovich, B. Rt-2: Vision-language-action models transfer web knowledge to robotic control. In *CoRL*, 2023a.
- Brohan, A., Brown, N., Carbajal, J., Chebotar, Y., Dabis, J., Finn, C., Gopalakrishnan, K., Hausman, K., Herzog, A., Hsu, J., Ibarz, J., Ichter, B., Irpan, A., Jackson, T., Jesmonth, S., Joshi, N. J., Julian, R., Kalashnikov, D., Kuang, Y., Leal, I., Lee, K.-H., Levine, S., Lu, Y., Malla, U., Manjunath, D., Mordatch, I., Nachum, O., Parada, C., Peralta, J., Perez, E., Pertsch, K., Quiambao, J., Rao, K., Ryoo, M., Salazar, G., Sanketi, P., Sayed, K., Singh, J., Sontakke, S., Stone, A., Tan, C., Tran, H., Vanhoucke, V., Vega, S., Vuong, Q., Xia, F., Xiao, T., Xu, P., Xu, S., Yu, T., and Zitkovich, B. Rt-1: Robotics transformer for real-world control at scale. In *RSS*, 2023b.
- Caasenbrood, B., Pogromsky, A., and Nijmeijer, H. Energy-shaping controllers for soft robot manipulators through port-hamiltonian cosserat models. *SN Computer Science*, 3(6):494, 2022.
- Centurelli, A., Arleo, L., Rizzo, A., Tolu, S., Laschi, C., and Falotico, E. Closed-loop dynamic control of a soft manipulator using deep reinforcement learning. *RA-L*, 7(2):4741–4748, 2022.
- Chen, T., Chen, Z., Chen, B., Cai, Z., Liu, Y., Li, Z., Liang, Q., Lin, X., Ge, Y., Gu, Z., et al. Robotwin 2.0: A scalable data generator and benchmark with strong domain randomization for robust bimanual robotic manipulation. *arXiv preprint arXiv:2506.18088*, 2025a.
- Chen, Z., Wu, D., Guan, Q., Hardman, D., Renda, F., Hughes, J., Thuruthel, T. G., Della Santina, C., Mazzolai, B., Zhao, H., et al. A survey on soft robot adaptability: Implementations, applications, and prospects [survey]. *RAM*, 2025b.
- Chi, C., Xu, Z., Feng, S., Cousineau, E., Du, Y., Burchfiel, B., Tedrake, R., and Song, S. Diffusion policy: Visuomotor policy learning via action diffusion. *IJRR*, 44(10-11):1684–1704, 2025.
- Cianchetti, M., Laschi, C., Menciassi, A., and Dario, P. Biomedical applications of soft robotics. *Nature Reviews Materials*, 3(6):143–153, 2018.
- Cosserat, E. M. P. and Cosserat, F. *Théorie des corps déformables*. A. Hermann et fils, 1909.
- Faure, F., Duriez, C., Delingette, H., Allard, J., Gilles, B., Marchesseau, S., Talbot, H., Courtecuisse, H., Bousquet, G., Peterlik, I., et al. Sofa: A multi-model framework for interactive physical simulation. In *Soft tissue biomechanical modeling for computer assisted surgery*, pp. 283–321. Springer, 2012.
- Geng, H., Wang, F., Wei, S., Li, Y., Wang, B., An, B., Cheng, C. T., Lou, H., Li, P., Wang, Y.-J., Liang, Y., Goetting, D., Xu, C., Chen, H., Qian, Y., Geng, Y., Mao, J., Wan, W., Zhang, M., Lyu, J., Zhao, S., Zhang, J., Zhang, J., Zhao, C., Lu, H., Ding, Y., Gong, R., Wang, Y., Kuang, Y., Wu, R., Jia, B., Sferrazza, C., Dong, H., Huang, S., Wang, Y., Malik, J., and Abbeel, P. Robo-verse: Towards a unified platform, dataset and benchmark for scalable and generalizable robot learning. *arXiv preprint arXiv:2504.18904*, 2025.
- Gong, Z., Fang, X., Chen, X., Cheng, J., Xie, Z., Liu, J., Chen, B., Yang, H., Kong, S., Hao, Y., et al. A soft manipulator for efficient delicate grasping in shallow water: Modeling, control, and real-world experiments. *IJRR*, 40(1):449–469, 2021.
- Gu, J., Xiang, F., Li, X., Ling, Z., Liu, X., Mu, T., Tang, Y., Tao, S., Wei, X., Yao, Y., et al. Maniskill2: A unified benchmark for generalizable manipulation skills. In *ICLR*, 2023.
- Haarnoja, T., Zhou, A., Abbeel, P., and Levine, S. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *ICML*, pp. 1861–1870, 2018.
- Huang, Z., Hu, Y., Du, T., Zhou, S., Su, H., Tenenbaum, J. B., and Gan, C. Plasticinelab: A soft-body manipulation benchmark with differentiable physics. *ICLR*, 2021.
- Hughes, J., Culha, U., Giardina, F., Guenther, F., Rosendo, A., and Iida, F. Soft manipulators and grippers: A review. *Frontiers in Robotics and AI*, 3:69, 2016.
- James, S., Ma, Z., Arrojo, D. R., and Davison, A. J. Rl-bench: The robot learning benchmark & learning environment. *RA-L*, 5(2):3019–3026, 2020.
- Kim, M. J., Pertsch, K., Karamcheti, S., Xiao, T., Balakrishna, A., Nair, S., Rafailov, R., Foster, E., Lam, G., Sanketi, P., Vuong, Q., Kollar, T., Burchfiel, B., Tedrake, R., Sadigh, D., Levine, S., Liang, P., and Finn, C. Openvla: An open-source vision-language-action model. In *CoRL*, 2024.
- Kim, M. J., Finn, C., and Liang, P. Fine-tuning vision-language-action models: Optimizing speed and success. In *RSS*, 2025.

- Li, G., Wong, T.-W., Shih, B., Guo, C., Wang, L., Liu, J., Wang, T., Liu, X., Yan, J., Wu, B., et al. Bioinspired soft robots for deep-sea exploration. *Nature Communications*, 14(1):7097, 2023.
- Li, Q., Liang, Y., Wang, Z., Luo, L., Chen, X., Liao, M., Wei, F., Deng, Y., Xu, S., Zhang, Y., Wang, X., Liu, B., Fu, J., Bao, J., Chen, D., Shi, Y., Yang, J., and Guo, B. Cogact: A foundational vision-language-action model for synergizing cognition and action in robotic manipulation. *arXiv preprint arXiv:2411.19650*, 2024a.
- Li, X., Hsu, K., Gu, J., Pertsch, K., Mees, O., Walke, H. R., Fu, C., Lunawat, I., Sieh, I., Kirmani, S., Levine, S., Wu, J., Finn, C., Su, H., Vuong, Q., and Xiao, T. Evaluating real-world robot manipulation policies in simulation. In *CoRL*, 2024b.
- Liu, B., Zhu, Y., Gao, C., Feng, Y., Liu, Q., Zhu, Y., and Stone, P. Libero: Benchmarking knowledge transfer for lifelong robot learning. In *NeurIPS*, 2023.
- Liu, J., Duo, Y., Chen, X., Zuo, Z., Liu, Y., and Wen, L. Data-driven methods for sensing, modeling and control of soft continuum robot: A review. *IEEE/ASME Trans. Mechatron.*, 2025a.
- Liu, S., Wu, L., Li, B., Tan, H., Chen, H., Wang, Z., Xu, K., Su, H., and Zhu, J. Rdt-1b: a diffusion foundation model for bimanual manipulation. In *ICLR*, 2025b.
- Majidi, C. Soft robotics: a perspective—current trends and prospects for the future. *Soft robotics*, 1(1):5–11, 2014.
- Mees, O., Hermann, L., Rosete-Beas, E., and Burgard, W. Calvin: A benchmark for language-conditioned policy learning for long-horizon robot manipulation tasks. *RA-L*, 7(3):7327–7334, 2022.
- Mu, T., Ling, Z., Xiang, F., Yang, D., Li, X., Tao, S., Huang, Z., Jia, Z., and Su, H. Maniskill: Generalizable manipulation skill benchmark with large-scale demonstrations. In *NeurIPS Datasets and Benchmarks*, 2021.
- Mu, Y., Chen, T., Peng, S., Chen, Z., Gao, Z., Zou, Y., Lin, L., Xie, Z., and Luo, P. Robotwin: Dual-arm robot benchmark with generative digital twins (early version). In *ECCV*, pp. 264–273. Springer, 2024.
- Naughton, N., Sun, J., Tekinalp, A., Parthasarathy, T., Chowdhary, G., and Gazzola, M. Elastica: A compliant mechanics environment for soft robotic control. *RA-L*, 6(2):3389–3396, 2021.
- Nazeer, M. S., Laschi, C., and Falotico, E. Soft dagger: Sample-efficient imitation learning for control of soft robots. *Sensors*, 23(19):8278, 2023.
- Pagliarani, N., Arleo, L., Alessi, C., Campinoti, G., Maselli, M., Falotico, E., and Cianchetti, M. Softtex: Soft robotic arm with learning-based textile proprioception. *RA-L*, 2025.
- Qi, Y., Wu, Q., Anderson, P., Wang, X., Wang, W. Y., Shen, C., and van den Hengel, A. REVERIE: Remote embodied visual referring expression in real indoor environments. In *CVPR*, pp. 9982–9991, 2020.
- Rogatinsky, J., Recco, D., Feichtmeier, J., Kang, Y., Kneier, N., Hammer, P., O’Leary, E., Mah, D., Hoganson, D., Vasilyev, N. V., et al. A multifunctional soft robot for cardiac interventions. *Science Advances*, 9(43): eadi5559, 2023.
- Ruiz, F., Arrue, B., and Ollero, A. Thermally-resilient soft gripper for on-orbit operations. In *IROS*, pp. 14050–14055, 2024.
- Shao, R., Li, W., Zhang, L., Zhang, R., Liu, Z., Chen, R., and Nie, L. Large vlm-based vision-language-action models for robotic manipulation: A survey. *arXiv preprint arXiv:2508.13073*, 2025.
- Shridhar, M., Thomason, J., Gordon, D., Bisk, Y., Han, W., Mottaghi, R., Zettlemoyer, L., and Fox, D. Alfred: A benchmark for interpreting grounded instructions for everyday tasks. In *CVPR*, pp. 10740–10749, 2020.
- Srivastava, S., Li, C., Lingelbach, M., Martín-Martín, R., Xia, F., Vainio, K. E., Lian, Z., Gokmen, C., Buch, S., Liu, K., Savarese, S., Gweon, H., Wu, J., and Fei-Fei, L. Behavior: Benchmark for everyday household activities in virtual, interactive, and ecological environments. In *CoRL*, pp. 477–490, 2022.
- Sutton, R. S. Learning to predict by the methods of temporal differences. *Machine learning*, 3(1):9–44, 1988.
- Szász, R., Allenspach, M., Han, M., Tognon, M., and Katzschnann, R. K. Modeling and control of an omnidirectional micro aerial vehicle equipped with a soft robotic arm. In *RoboSoft*, pp. 01–08, 2022.
- Szot, A., Clegg, A., Undersander, E., Wijmans, E., Zhao, Y., Turner, J., Maestre, N., Mukadam, M., Chaplot, D. S., Maksymets, O., Gokaslan, A., Vondruš, V., Dharur, S., Meier, F., Galuba, W., Chang, A., Kira, Z., Koltun, V., Malik, J., Savva, M., and Batra, D. Habitat 2.0: Training home assistants to rearrange their habitat. In *NeurIPS*, volume 34, pp. 251–265, 2021.
- Tao, S., Xiang, F., Shukla, A., Qin, Y., Hinrichsen, X., Yuan, X., Bao, C., Lin, X., Liu, Y., kai Chan, T., Gao, Y., Li, X., Mu, T., Xiao, N., Gurha, A., Rajesh, V. N., Choi, Y. W., Chen, Y.-R., Huang, Z., Calandra, R., Chen, R., Luo, S., and Su, H. Maniskill3: Gpu parallelized

- robotics simulation and rendering for generalizable embodied ai. *RSS*, 2025.
- Thuruthel, T. G., Falotico, E., Renda, F., and Laschi, C. Learning dynamic models for open loop predictive control of soft robotic manipulators. *Bioinspiration & biomimetics*, 12(6):066003, 2017.
- Thuruthel, T. G., Falotico, E., Renda, F., and Laschi, C. Model-based reinforcement learning for closed-loop dynamic control of soft robotic manipulators. *T-RO*, 35(1): 124–134, 2018.
- Todorov, E., Erez, T., and Tassa, Y. Mujoco: A physics engine for model-based control. In *IROS*, pp. 5026–5033, 2012.
- Walker, K. L., Partridge, A. J., Chen, H.-Y., Ramachandran, R. R., Stokes, A. A., Tadakuma, K., Da Silva, L. C., and Giorgio-Serchi, F. A modular, tendon driven variable stiffness manipulator with internal routing for improved stability and increased payload capacity. In *ICRA*, pp. 3030–3035, 2024.
- Wei, Z., Wang, L., Gao, C., Huang, H., Liu, J., Wen, L., and Liu, S. Robosoft’25: The 1st international workshop on vision-language in soft robot. In *Proceedings of the 33rd ACM International Conference on Multimedia*, pp. 14311–14313, 2025.
- Wen, J., Zhu, Y., Li, J., Tang, Z., Shen, C., and Feng, F. Dexvla: Vision-language model with plug-in diffusion expert for general robot control. In *CoRL*, 2025.
- Xiang, F., Qin, Y., Mo, K., Xia, Y., Zhu, H., Liu, F., Liu, M., Jiang, H., Yuan, Y., Wang, H., Yi, L., Chang, A. X., Guibas, L. J., and Su, H. SAPIEN: A simulated part-based interactive environment. In *CVPR*, 2020.
- Xie, Z., Yuan, F., Liu, J., Tian, L., Chen, B., Fu, Z., Mao, S., Jin, T., Wang, Y., He, X., Wang, G., Mo, Y., Ding, X., Zhang, Y., Laschi, C., and Wen, L. Octopus-inspired sensorized soft arm for environmental interaction. *Sci. Robot.*, 8(84), 2023.
- Yu, T., Quillen, D., He, Z., Julian, R., Narayan, A., Shively, H., Bellathur, A., Hausman, K., Finn, C., and Levine, S. Meta-world: A benchmark and evaluation for multi-task and meta reinforcement learning, 2021. URL <https://arxiv.org/abs/1910.10897>.
- Zhao, Z., Wu, Q., Wang, J., Zhang, B., Zhong, C., and Zhilenkov, A. A. Exploring embodied intelligence in soft robotics: a review. *Biomimetics*, 9(4):248, 2024.

Supplementary Material

A. Cosserat Rod Theory

In Cosserat Rod Theory (Cosserat & Cosserat, 1909), the elastic rod with L_0 and radius r_0 is presented as a Cosserat rod composed of N discrete element rods, each of length L_0/N and radius r_0 . For each element rod, we describe its position in the global frame by $\bar{\mathbf{x}}(s, t) \in \mathbb{R}^3$, and its rotation in the global frame is represented by the rotation matrix $\mathbf{Q}(s, t) = \{\bar{\mathbf{d}}_1, \bar{\mathbf{d}}_2, \bar{\mathbf{d}}_3\}^{-1}$, where $\mathbf{Q}(s, t)$ also defines the transformation between the global and local frames. Specifically, for any vector $\mathbf{v}$ in the local frame and $\bar{\mathbf{v}}$ in the global frame, we have $\mathbf{v} = \mathbf{Q}\bar{\mathbf{v}}$ and $\bar{\mathbf{v}} = \mathbf{Q}^T\mathbf{v}$. Here, $s = L \cdot i/N$ denotes the position of the i -th element rod in the material coordinate, t represents time, and $\delta s = L/N$ is the length of each element rod. As $N \rightarrow \infty$, s becomes continuous, and $\delta s \rightarrow ds$. Subsequent derivations will be carried out in the continuous case.

The normal strain of the rod is described by the stretch factor $e = ds/d\hat{s}$, where $d\hat{s} = L_0/N$, and $N \rightarrow \infty$ represents the original length of the element rod. The shear strain of the rod is described by the shear vector in the local frame, $\sigma = \mathbf{Q}(\bar{\mathbf{x}}_s - \bar{\mathbf{d}}_3)$, where $\bar{\mathbf{x}}_s = \partial_s \bar{\mathbf{x}}$ is the centerline tangent in the global frame. At this point, we have the translational velocity $\dot{\bar{\mathbf{x}}} = \partial_t \bar{\mathbf{x}}$ and the curvature vector κ satisfies $\partial_s \mathbf{d}_j = \kappa \times \mathbf{d}_j$, which describes the rate of change of rotation along the material coordinate. The angular velocity ω satisfies $\omega = \partial_t \mathbf{d}_j = \omega \times \mathbf{d}_j$, which describes the rate of change of rotation over time.

Given the bending $\mathbf{B}$ and shearing $\mathbf{S}$ stiffness matrices, the second area moment of inertia $\mathbf{I}$, the cross-sectional area A , and the mass per unit length ρ , the dynamics of the Cosserat rod can then be written based on the momentum and angular momentum theorems as follows:

$$\rho A \cdot \partial_t \bar{\mathbf{v}} = \partial_s \left(\frac{\mathbf{Q}^T \mathbf{S} \sigma}{e} \right) + e \bar{\mathbf{f}}, \quad (5)$$

$$\begin{aligned} \frac{\rho \mathbf{I}}{e} \cdot \partial_t \omega &= \partial_s \left(\frac{\mathbf{B} \kappa}{e^3} \right) + \frac{\kappa \times \mathbf{B} \kappa}{e^3} + \left(\mathbf{Q} \frac{\bar{\mathbf{x}}_s}{e} \times \mathbf{S} \sigma \right) \\ &+ \left(\rho \mathbf{I} \cdot \frac{\omega}{e} \right) \times \omega + \frac{\rho \mathbf{I} \omega}{e^2} \cdot \partial_t e + e \boldsymbol{\tau}. \end{aligned} \quad (6)$$

Where $\bar{\mathbf{f}}$ is the force density in the global frame for the Cosserat rod, and $\boldsymbol{\tau}$ is the torque density. In the discrete case, they represent the force and torque acting on each element rod. We propose a simulation framework for soft robotic arms that captures both their deformable dynamics and interactions with the environment. As shown in Figure 2, the arm is modeled as two coupled components: a deformable soft body and an end-effector, connected to allow coordinated motion.

Table 4. Key Parameters in the ManiSoft Simulator

Parameter	Value
Simulator	
k_F	0.1 N/m
k_M	10 N · m/rad
Simulation timestep	0.0002 s
Control Frequency	714 Hz
Soft Arm	
Length	1 m
Radius	0.05 m
Density	1000 kg/m ³
Poisson’s ratio	0.5
Young’s modulus	1.0 × 10 ⁷ Pa
Render	
Resolution	514 × 514
Camera Position	(0, 1.6, 1.6)m
Camera FOV	60

B. Details for *ManiSoft* Benchmark

In the simulation, choosing appropriate values for k_F and k_M is crucial for maintaining both numerical stability and physical realism. If the coefficients are too small, positional and orientational discrepancies may persist, leading to separation between the components. On the other hand, excessively large coefficients can result in overcorrection, causing oscillations or even numerical instability. By carefully tuning the evolution of the elastic constraint, *ManiSoft*-Sim ensures stable, physically consistent coupling between the soft body and the EEF, enabling a coherent simulation of soft robotic manipulation.

Specifically, the parameters of the simulator used in our experiments are listed in the Table 4.

In the *ManiSoft* benchmark, we set the maximum execution horizon to $T = 1500$ steps.

C. Details for Executor Training

For training the executor, we adopt an MLP-based policy network and employ SAC (Haarnoja et al., 2018), for reinforcement learning. We use a learning rate of 3×10^{-4} and a batch size of 256. The model is trained using a total of 160M samples.

We perform training and evaluation with different parameters in the reward function.

First, we examine the success rate using a reward that includes only R_d ($\beta = 0$) across various parameter settings. A case is deemed successful once the pose difference between the end-effector and the target drops below a predefined threshold. In *Elastica-RL-Control* (Naughton et al., 2021), the parameters are set as $k_1 = 0.5$, $k_2 = 1.5$, $d_1 = 0.1$, and $d_2 = 0.05$. We adopt the same values for k_1 and k_2 . For d_1 and d_2 , since we replace the original Euclidean distance with pose difference, the scale of d changes. To maintain the original ratio between d_1 and d_2 , we scale them proportionally, setting $d_1 = 0.1/\lambda$ and $d_2 = 0.05/\lambda$. The model is trained under different λ and α configurations. For each setting, we train on 20M samples. During evaluation, a case is considered successful if $d_p < 0.03$ and $d_r < 0.3$. We randomly sample 100 cases to evaluate the success rate, as summarized in Tab. 6. We found that despite changing d from Euclidean distance to pose difference, the best performance was still achieved when d_1 and d_2 remained unchanged, *i.e.*, when $\lambda = 1$. With λ fixed at 1, we trained on 80M samples with different α values, as shown in Table 6. The best performance was achieved with $\alpha = 0.2$.

Based on the model with the highest success rate, we add R_s and perform post-training on 80M samples. We then compare the stability performance under different parameter settings, as shown in Table 1 of the main text.

Fig. 10 shows the visualization of the trained executor controlling the soft robotic arm to move to the target pose.

D. Implement Details for Baselines

We train and evaluate the three baselines separately on each of the four tasks. All models are trained on 8 RTX 4090 GPUs.

D.1. Clean Setting.

For the cleaning setting, we adopt the following training configuration.

DP (Chi et al., 2025). We set the batch size to 64 and the learning rate to 1×10^{-4} . The model is trained for 120,000 iterations on the COLL task and 60,000 iterations on each of the other three tasks, using a linear learning rate decay schedule. Since DP does not inherently support language understanding, we employ BERT as the text encoder. The resulting text embeddings are combined with image embeddings to guide action generation.

RDT (Liu et al., 2025b) We use a batch size of 32 and a learning rate of 1×10^{-4} , while keeping the text encoder and vision encoder frozen. The model is trained for 60,000 iterations on the COLL task and 30,000 iterations on the remaining tasks, with a cosine learning rate decay schedule.

OpenVLA-OFT (Kim et al., 2025). We finetune the model using LoRA based on the official pretrained weights, with a batch size of 4 and a learning rate of 5×10^{-4} . The model is trained for 60,000 iterations on the COLL task and 30,000 iterations on the other three tasks, following a cosine learning rate decay schedule.

Table 5. Success rate (%) on 100 samples under different parameters of R_d .

$\alpha \backslash \lambda$	0.4	0.6	0.8	1.0	1.2	1.4
0.02	6.0	6.0	3.0	6.0	5.0	5.0
0.04	6.0	6.0	16.0	19.0	0.0	2.0
0.05	2.0	2.0	11.0	8.0	11.0	4.0
0.10	2.0	11.0	16.0	17.0	13.0	20.0
0.15	3.0	4.0	8.0	19.0	9.0	16.0
0.20	7.0	15.0	20.0	6.0	6.0	3.0
Avg.	4.7	7.3	12.3	12.5	7.0	8.3

Table 6. Success rate on 100 samples under difference values of α with $\lambda = 1$.

α	0.04	0.15	0.20	0.40
Success Rate (%)	16.0	31.0	33.0	23.0

D.2. Randomized Setting.

For the randomized setting, we finetune the model initialized from the clean setting checkpoint, training for 20,000 iterations on COLL and 10,000 iterations on each of the other three tasks, while keeping all other training configurations unchanged.

E. Results on Each Category

In COLL, ALN, and ARR, multiple categories of manipulable objects are included. Tab. 7 and Tab. 8 present the results for different categories of manipulable objects in COLL and ALN, respectively.

For the COLL task, DP outperforms both RDT and OpenVLA-OFT (by 17.6% and 49.2% respectively) in the clean setting, while in the randomized setting, OpenVLA-OFT performs better than DP by 28.9%. Regarding the number of inference steps, DP performs better than OpenVLA-OFT in both settings (by 18 on clean and 80 on randomized). This is due to the stop-moving phenomenon in OpenVLA-OFT, which leads to an increase in inference steps. Although RDT requires fewer inference steps, it completes fewer tasks overall, and the tasks it does complete are simpler, requiring fewer execution steps. This does not accurately reflect its overall performance.

The comparison of success rates across different objects reveals that, compared to the candle stick (100% success rate on DP) and the can (85%), the shoe (35%) and the sand clock (35%) are more difficult to grasp. This is because they require a fixed grasping direction or have relatively large volumes.

A similar trend is observed in the ALN task, where OpenVLA-OFT achieves a higher success rate than DP (by 6.7% on clean and 25% on randomized), but requires more inference steps (by 30 on clean and 65 on randomized). In the clean setting, for the same object such as bottle, the success rate in the COLL task is higher than in the ALN task, indicating that the ALN task is relatively more challenging.

F. More Visualizations

In Figure 11, Figure 12, Figure 13 and Figure 14, we present more visualizations of the four tasks.

Table 7. The results of the policy on each manipulable objects in the COLL task. OFT denotes OpenVLA-OFT

Category	Clean						Randomized					
	DP (Chi et al., 2025)		RDT (Liu et al., 2025b)		OFT (Kim et al., 2025)		DP (Chi et al., 2025)		RDT (Liu et al., 2025b)		OFT (Kim et al., 2025)	
	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps
Bottle	70.0	543	15.0	444	50.0	530	5.0	749	0.0	-	45.0	569
Pill Bottle	75.0	545	10.0	420	55.0	520	10.0	349	5.0	113	40.0	599
Can	100.0	535	35.0	550	60.0	597	10.0	515	0.0	-	45.0	663
Cup	65.0	502	15.0	442	60.0	529	5.0	603	5.0	516	30.0	591
Sand Clock	35.0	538	15.0	588	35.0	473	0.0	-	0.0	-	15.0	454
Shoe	35.0	516	20.0	480	20.0	556	0.0	-	0.0	-	20.0	580
Candle Stick	85.0	502	15.0	619	60.0	530	0.0	-	0.0	-	40.0	545
Box Drink	65.0	511	15.0	530	35.0	675	10.0	287	5.0	833	35.0	613
Kettle	50.0	634	0.0	-	45.0	602	0.0	-	0.0	-	35.0	593
Pen Cup	80.0	489	5.0	484	60.0	490	5.0	455	0.0	-	40.0	735
Sauce Can	70.0	557	20.0	496	40.0	532	0.0	-	0.0	-	20.0	513
Rubik's Cube	50.0	547	10.0	744	35.0	667	0.0	-	0.0	-	30.0	708
Trophy	45.0	689	5.0	309	35.0	641	5.0	689	0.0	-	30.0	652
Average	63.0	547	13.8	509	45.4	565	3.8	521	1.2	487	32.7	601

Table 8. The results of the policy on each manipulable objects in the ALN task.

Setting	Method	Bottle		Shoe		Candle Stick		Average	
		ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps	ACC(%)	#Steps
Clean	DP (Chi et al., 2025)	5.0	391	5.0	519	45.0	417	18.3	442
	RDT (Liu et al., 2025b)	5.0	538	20.0	535	10.0	316	11.7	463
	OpenVLA-OFT (Kim et al., 2025)	15.0	370	25.0	597	35.0	449	25.0	472
Randomized	DP (Chi et al., 2025)	2.5	371	2.5	278	0.0	-	1.7	324
	RDT (Liu et al., 2025b)	0.0	-	7.5	498	5.0	260	4.2	379
	OpenVLA-OFT (Kim et al., 2025)	5.0	559	20.0	430	55.0	480	26.7	489




Figure 10. Visualization of the trained executor controlling the soft robotic arm to move to the target pose.






Figure 11. **Visualization of COLL Task.** The first two are for the clean setting, and the last two are for the randomized setting.



Instruction: Pick up shoe and position it according to the gray area's position and direction.



Instruction: Move bottle to the gray area, making sure its direction aligns with the area's orientation.




Instruction: Pick up bottle with tapered cylindrical shape, shift it into the gray area, align its orientation with area's.



Instruction: Move the brown sports shoe to the gray area, keeping its direction in line with the area's orientation.

Figure 12. **Visualization of ALN Task.** The first two are for the clean setting, and the last two are for the randomized setting.



Instruction: Arrange the tableware in stacks sorted by size.



Instruction: Pile up the eating utensils by size.

Figure 13. **Visualization of STK Task.** The first one is for the clean setting, and the last one is for the randomized setting.




Figure 14. **Visualization of ARR Task.** The first two are for the clean setting, and the last two are for the randomized setting.