
What Makes a Strong Model? A Unified Spectral Analysis of Knowledge Transfer over High-dimensional Linear Regression

Wendao Wu^{1,2} Fangqing Zhang¹ Haihan Zhang¹ Cong Fang¹

Abstract

Teacher-Student Knowledge Transfer (KT) is ubiquitous in modern machine learning, ranging from classical model compression via Knowledge Distillation (KD) to the emergent phenomenon of Weak-to-Strong (W2S) generalization. While existing studies offer isolated insights, a unified theoretical framework explaining the efficacy of KT across these disparate regimes remains lacking. In this work, we establish a unified spectral analysis of SGD dynamics in high-dimensional linear regression, elucidating the efficiency of KT across seemingly disparate regimes. We characterize KT efficiency through two distinct mechanisms: *Spectral Horizon Expansion* in KD, which enables the capture of statistically inaccessible high-frequency signals, and *Spectral Denoising* in W2S, where the student acts as a filter for optimization noise. Our framework unifies these phenomena, revealing that the efficacy of transfer is governed by the interplay between implicit regularization and heterogeneous spectral learning speeds over the spectrum.

1. Introduction

Knowledge Distillation (KD) has established itself as a cornerstone technique for model compression and efficient inference (Hinton et al., 2015). The classical KD paradigm follows a sequential two-stage pipeline: initially, a high-capacity teacher is trained on the source dataset, and subsequently, a lightweight student is optimized to match the teacher’s softened outputs. Surprisingly, this process frequently yields students with generalization capabilities superior to those trained directly on the same hard labels (Stan-

ton et al., 2021). This counter-intuitive success has fueled widespread adoption, yet the theoretical underpinning of *why* mimicking a proxy model outperforms learning from ground truth remains a subject of active debate.

The landscape of Knowledge Transfer (KT) has since expanded beyond mere compression. *Self-Distillation* (SD) (Furlanello et al., 2018) demonstrates that a student with an identical architecture to the teacher can achieve better convergence, challenging the intuition that capacity gaps are necessary for transfer gains. Moreover, the *Weak-to-Strong* (W2S) generalization paradigm (Burns et al., 2023) asks whether a low-capacity supervisor can effectively guide a stronger model. Empirical evidence suggests that strong students can indeed surpass their weak teachers, raising a fundamental question for the alignment of future AI systems: *How can a student exceed the very source of its supervision?*

These diverse phenomena—KD, SD, and W2S—present a theoretical puzzle. A prevalent intuition posits that “soft” labels convey “dark knowledge” (Hinton et al., 2015), acting as a form of label smoothing or denoising. However, this view remains heuristic and fails to explain why students do not simply inherit the teacher’s biases. Recent theoretical advances suggest that the mechanism lies in specific *multi-view* data structures and the resulting *implicit regularization* within optimization dynamics (Allen-Zhu & Li, 2023). This perspective implies that the core mechanics of KT manifest even in simplified settings, driven by the coupling between the learning algorithm and the data’s geometry.

However, existing frameworks lack a fine-grained characterization of how optimization interacts with the intrinsic structure of the data. Prior theoretical works have overlooked the **spectral properties** of the data covariance matrix—the fundamental governor of learning dynamics that dictates which features are learned fast or slow. To the best of our knowledge, a unified framework that coherently explains KT across disparate capacity regimes (from “strong-to-weak” to “weak-to-strong”) via the lens of spectral analysis remains elusive. In this work, we fill this gap by explicitly leveraging spectral properties to characterize the unified mechanics of knowledge transfer.

A Unified Spectral Perspective. In this work, we bridge

¹State Key Lab of General AI, School of Intelligence Science and Technology, Peking University, China ²School of Mathematical Sciences, Peking University, China. Correspondence to: Cong Fang <fangcong@pku.edu.cn>.

this gap by developing a unified theory within the canonical high-dimensional linear regression model optimized by Stochastic Gradient Descent (SGD). While linear models are simplifications, they offer a rigorous proxy for deep learning dynamics in the over-parameterized regime. Specifically, Neural Tangent Kernel (NTK) theory establishes that wide neural networks evolve as linear models over fixed feature representations (Jacot et al., 2018). Recent work by Malladi et al. (2023) further validates this abstraction in the context of Large Language Models (LLMs), demonstrating that fine-tuning dynamics—the primary use case for KT today—can be approximately described by kernel regression, which is mathematically equivalent to linear regression in the reproducing kernel Hilbert space.

This equivalence allows us to rigorously formalize the notion of “knowledge” as a *spectral coupling*: the spectrum of the data covariance matrix and also the *distribution of target parameters* across these dimensions (which dictates where the signal resides). We analyze the dynamics of single-pass SGD (or online SGD) rather than static closed-form solutions (e.g., Ridge Regression). This choice is deliberate and crucial: unlike offline solvers, SGD captures the *trajectory* of learning, allowing us to characterize how the student’s effective horizon progressively expands to cover the signal distributed along the spectrum.

From Risk Decomposition to Spectral Mechanisms. Armed with this perspective, we derive a finite-sample risk bound decomposition (Theorem 1). This decomposition allows us to identify the specific spectral mechanisms governing distinct regimes. In the “Strong Teacher” regime (KD) (Section 5), the teacher acts as a *spectral guide*, extending the student’s **Effective Optimization Horizon** to capture high-frequency signals that are otherwise statistically indistinguishable from noise. Conversely, in the “Weak/Identical Teacher” regime (W2S/SD) (Section 6), the student functions as a *spectral filter*, leveraging **Spectral Denoising** to dampen the teacher’s high-dimensional optimization variance while recovering the underlying population geometry.

Main Contributions are summarized as follows:

A Unified Spectral Framework for Knowledge Transfer. We establish a comprehensive theoretical framework for analyzing Knowledge Transfer (KT) dynamics in high-dimensional linear regression. We derive a finite-sample risk bound (Theorem 1) that rigorously decomposes the transfer error into three distinct components: *geometric misalignment*, *propagated teacher noise*, and *student optimization error*. This decomposition, in conjunction with Lemma 2, establishes an **exact characterization** of the conditions under which knowledge transfer is provably efficient.

Demystifying KD, W2S, and Self-Distillation via Spectral Mechanisms. We apply our framework to elucidate the

mechanisms behind seemingly contradictory phenomena:

- In the **Knowledge Distillation** regime, we identify a mechanism of **Spectral Horizon Expansion**: the teacher actively extends the student’s *effective optimization horizon* beyond its intrinsic limit, facilitating the capture of high-frequency components that are statistically inaccessible from raw data.
- In the **Weak-to-Strong** and **Self-Distillation** regimes, we uncover a **Spectral Denoising** mechanism: the student leverages its excess capacity to filter out teacher’s optimization noise, effectively recovering the teacher’s optimal *population geometry* from noisy labels.

2. Related Works

Our work bridges the gap between the theoretical analysis of over-parameterized optimization and the recent empirical surge in model-to-model supervision (distillation and weak-to-strong generalization).

Analysis of Stochastic Gradient Descent. The theoretical properties of SGD in high-dimensional and kernel regimes serve as the bedrock of our analysis. Foundational works have rigorously characterized the bias-variance decomposition of SGD, explicitly linking convergence rates to the tail decay of the feature covariance spectrum (Dieuleveut & Bach, 2016; Zou et al., 2021; Wu et al., 2022; Zhang et al., 2025). This spectral perspective has been further refined to establish optimality guarantees and analyze accelerated variants (Jain et al., 2018b; Pan et al., 2022; Li et al., 2023; Zhang et al., 2024), providing a robust toolkit for understanding over-parameterized dynamics. Most recently, this machinery has been instrumental in explaining the emergence of neural scaling laws (Atanasov et al., 2024; Lin et al., 2025), demonstrating the power of linear models in capturing modern deep learning phenomena.

Knowledge Distillation. Since Hinton et al. (2015), Knowledge Distillation (KD) has become a cornerstone of modern AI, widely deployed in LLM instruction tuning and efficient model training (Wang et al., 2023; Abdin et al., 2024; Guo et al., 2025). Notably, Furlanello et al. (2018) pioneered the concept of *Self-Distillation*, demonstrating that students can surprisingly outperform identical teachers. Theoretically, these successes are attributed to diverse mechanisms: Menon et al. (2021) identify KD as a variance reduction technique via Bayes probability approximation; Mobahi et al. (2020) characterize it as implicit regularization restricting basis functions; others point to multi-view feature discovery (Allen-Zhu & Li, 2023) or the amplification of spectral bias (Nagarajan et al., 2024).

Weak-to-Strong Generalization. Initiated by the empirical findings of Burns et al. (2023), this paradigm has spurred

diverse theoretical inquiries. A dominant strand analyzes the phenomenon through static or offline regression frameworks: Moniri & Hassani (2025) and Dong et al. (2025) attribute the student’s success to its ability to compensate for the teacher’s insufficient regularization, while Ildiz et al. (2025) derive optimal surrogates in the ridgeless setting to maximize transfer efficiency. Beyond standard regression, Charikar et al. (2024) propose a “misfit error” framework to quantify performance gains, and Wu & Sahai (2024) establish provable guarantees in classification via benign overfitting. More recently, Medvedev et al. (2025) have extended the analysis to gradient flow in random feature models, primarily focusing on the bias dynamics.

Spectral Evaluation of Representation Quality. Recent empirical works propose spectral metrics to assess model quality, such as *RankMe* (effective rank) (Garrido et al., 2023) and α -*ReQ* (spectral decay) (Agrawal et al., 2022). Our framework rigorously grounds these heuristics: we prove they are causal drivers of transfer efficiency, where higher rank enhances *noise resilience* (Theorem 2) and slower decay ensures *feature coverage* (Theorem 3).

Notation. We use bold lowercase letters for vectors and bold uppercase letters for matrices. The symbol $\otimes$ denotes the Kronecker product. For a vector $\boldsymbol{\eta}$, we denote its outer product as $\boldsymbol{\eta}^{\otimes 2} := \boldsymbol{\eta}\boldsymbol{\eta}^\top$. For any two matrices $\mathbf{A}, \mathbf{B}$ of compatible dimensions, their Frobenius inner product is defined as $\langle \mathbf{A}, \mathbf{B} \rangle := \text{tr}(\mathbf{A}^\top \mathbf{B})$. Regarding asymptotic analysis, we use $\tilde{O}(\cdot)$, $\tilde{\Omega}(\cdot)$, and $\tilde{\Theta}(\cdot)$ to suppress polylogarithmic factors; for instance, $f(n) = \tilde{\Theta}(g(n))$ implies $f(n) = \Theta(g(n)\text{polylog}(n))$.

3. Problem Formulation and Preliminaries

3.1. Teacher-Student Learning Setup

Data Generation. We model the learning task as a regression problem over a compact input domain $\mathcal{X}$. Training data $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$ are drawn i.i.d. from a probability measure $\rho_{\mathbf{x} \times y}$. The labels are generated by a ground truth function f_* contaminated by noise: $y = f_*(\mathbf{x}) + \epsilon$, where $\epsilon \sim \mathcal{N}(0, \sigma^2)$ is independent Gaussian noise. We assume the existence of an underlying *semantic basis* $\Phi : \mathcal{X} \rightarrow \mathbb{R}^D$ that is sufficiently expressive to represent the ground truth. Specifically, $f_*(\mathbf{x}) = \langle \mathbf{w}_*, \Phi(\mathbf{x}) \rangle$ for some coefficient vector $\mathbf{w}_* \in \mathbb{R}^D$. Without loss of generality, we assume this basis is orthonormal with respect to the marginal distribution $\rho_{\mathbf{x}}$, such that $\mathbb{E}_{\mathbf{x}}[\Phi(\mathbf{x})\Phi(\mathbf{x})^\top] = \mathbf{I}_D$.

Teacher and Student Models. We analyze the transfer of knowledge between a teacher model (f_T) and a student model (f_S). Both models operate on features derived from the common basis $\Phi(\mathbf{x})$ via linear transformations. Let $\mathbf{M}_T \in \mathbb{R}^{d_T \times D}$ and $\mathbf{M}_S \in \mathbb{R}^{d_S \times D}$ denote the feature extraction matrices for the teacher and student, respectively. The

feature maps are defined as:

$$\phi_\nu(\mathbf{x}) = \mathbf{M}_\nu \Phi(\mathbf{x}), \quad \text{for } \nu \in \{T, S\}. \quad (1)$$

The covariance matrices of these features are naturally given by $\Sigma_\nu = \mathbf{M}_\nu \mathbf{M}_\nu^\top$. We assume Σ_ν are trace-class operators (i.e., $\text{tr}(\Sigma_\nu) < \infty$) to ensure bounded signal energy.

Spectral Decomposition and Projections. To characterize the expressivity of the models, we consider the Singular Value Decomposition (SVD) of the feature mappings. Let $\mathbf{M}_S = \mathbf{U}_S \Lambda_S \mathbf{V}_S^\top$, where $\mathbf{V}_S \in \mathbb{R}^{D \times d_S}$ spans the active row space of the student. We define $\Pi_S := \mathbf{V}_S \mathbf{V}_S^\top$ as the orthogonal projection onto the student’s feature space. Similarly, $\Pi_S^\perp := \mathbf{I}_D - \Pi_S$ is the projection onto its null space, representing the information strictly inaccessible to the student. For any spectral cutoff k , we denote $\Pi_S^{\leq k}$ as the projection onto the subspace spanned by the top- k right singular vectors of $\mathbf{M}_S$.

Learning Objectives and Risks. We evaluate learning performance through two lenses: direct learning from ground truth and knowledge transfer. The teacher and student models, parameterized by $\mathbf{w}_T$ and $\mathbf{w}_S$, approximate the target as $f_\nu(\mathbf{x}) = \langle \mathbf{w}_\nu, \phi_\nu(\mathbf{x}) \rangle$. The standard supervised risks are:

$$\mathcal{R}_\nu(\mathbf{w}_\nu) = \frac{1}{2} \mathbb{E}_{\mathbf{x}} [(f_\nu(\mathbf{x}) - f_*(\mathbf{x}))^2], \quad \text{for } \nu \in \{T, S\}. \quad (2)$$

Let $\mathbf{w}_\nu^*$ denote the population optimal parameters minimizing Eq. equation 2.

In the *Knowledge Transfer* process, the student learns to mimic the teacher’s output. We define the *Teacher-to-Student* (T2S) model f_{T2S} parameterized by $\mathbf{w}_{T2S}$, which is optimized on the transfer risk:

$$\mathcal{R}_{\text{Trans}}(\mathbf{w}_{T2S}) = \frac{1}{2} \mathbb{E}_{\mathbf{x}} [(f_{T2S}(\mathbf{x}) - f_T(\mathbf{x}))^2]. \quad (3)$$

Let $\mathbf{w}_{T2S}^{\text{opt}}$ denote the population optimal parameter minimizing Eq. equation 3. When the student is trained with the optimal teacher model, i.e. $f_T(\mathbf{x}) = \langle \mathbf{w}_T^*, \phi_T(\mathbf{x}) \rangle$, we denote $\mathbf{w}_{T2S}^{\text{opt}}$ in this case as $\mathbf{w}_{T2S}^*$.

The effectiveness of knowledge transfer is ultimately measured by how well the transferred student f_{T2S} recovers the ground truth, quantified by the risk $\mathcal{R}_{T2S}(\mathbf{w}_{T2S}) := \frac{1}{2} \mathbb{E}_{\mathbf{x}} [(f_{T2S}(\mathbf{x}) - f_*(\mathbf{x}))^2]$.

Many following analysis focus on the **Excess Risk**, defined as the generalization gap above this population optimum:

$$\mathcal{E}_\nu(\mathbf{w}_\nu) := \mathcal{R}_\nu(\mathbf{w}_\nu) - \mathcal{R}_\nu(\mathbf{w}_\nu^*), \quad \text{for } \nu \in \{T, S, T2S\}. \quad (4)$$

Next, we define the quantitative metrics used to characterize distillation efficiency and weak-to-strong generalization.

Quantifying Distillation Efficiency (Strong Teacher). When a **Strong Teacher** instructs a Weak Student, the primary benefit is accelerated convergence. We quantify this

via the *Distillation Efficiency Ratio* (DER). Let $\mathcal{E}_S(N)$ be the excess risk of a student trained on N ground-truth samples. We define $\mathcal{E}_{T2S}(N, n)$ as the risk of a student trained on a separate set of n **transfer samples** labeled by the teacher. The DER is given by:

$$\text{DER}_N := \frac{\mathbb{E}[\mathcal{E}_S(N)]}{\mathbb{E}[\mathcal{E}_{T2S}(N, n)]}. \quad (5)$$

Values > 1 signify *data amplification*, implying the student learns more efficiently from the teacher’s synthesized knowledge than from raw data.

Quantifying Weak-to-Strong Generalization. In the regime where a **Weak Teacher** supervises a Strong Student, our focus shifts to whether the student can transcend the teacher’s limitations. We formally characterize the *occurrence* of the Weak-to-Strong (W2S) phenomenon by the strict inequality:

$$\mathcal{R}_{T2S}(\mathbf{w}_{T2S}) < \mathcal{R}_T(\mathbf{w}_T). \quad (6)$$

To further quantify the *quality* of this generalization, we adopt the **Performance Gap Recovered (PGR)** metric (Burns et al., 2023). While the inequality above signals the existence of W2S, PGR measures the extent to which the distilled student recovers the optimal capabilities accessible only via ground truth supervision:

$$\text{PGR} := \frac{\mathcal{R}_T - \mathcal{R}_{T2S}}{\mathcal{R}_T - \mathcal{R}_S}. \quad (7)$$

3.2. Assumptions and SGD Settings

We formalize the training dynamics as a two-phase protocol consisting of **Teacher Supervised Training** and **Student Transfer Learning**. Both phases utilize a step-decay learning rate schedule, which is shown to be asymptotically optimal for linear regression (Ge et al., 2019). We define the schedule scaling function $\mathcal{S}(t) := 2^{-\lfloor t/K \rfloor}$, where $K = N / \log_2 N$ represents the decay interval.

In the **first phase**, the teacher model is trained on a dataset $\mathcal{D}_N$ of size N , drawn i.i.d. from the ground-truth distribution $\rho_{\mathbf{x} \times \mathbf{y}}$. The optimizer follows the schedule $\gamma_t = \gamma_0 \cdot \mathcal{S}(t)$. In the **second phase**, the student is trained on a separate dataset $\mathcal{D}_n$ of size n , drawn from the marginal $\rho_{\mathbf{x}}$. The student learns from the fixed teacher’s labels $\hat{y} = f_T(\mathbf{x})$ using the schedule $\gamma'_t = \gamma'_0 \cdot \mathcal{S}(t)$, where the initial step size γ'_0 is a tunable hyperparameter distinct from γ_0 .

To analyze the convergence, we make the following standard assumption on the feature distribution.

Assumption 1 (Fourth Moment Condition). *There exists a finite constant $\psi \geq 1$, such that for every positive semi-definite (PSD) matrix $\mathbf{A}$, the feature distribution satisfies:*

$$\mathbb{E}_{\rho_{\mathbf{x}}} [\Phi(\mathbf{x})\Phi(\mathbf{x})^\top \mathbf{A} \Phi(\mathbf{x})\Phi(\mathbf{x})^\top] \preceq \psi \cdot \text{tr}(\mathbf{A}) \cdot \mathbf{I}_D. \quad (8)$$

Algorithm 1 Two-Phase SGD for Knowledge Transfer

Phase 1: Teacher Supervised Training

Initialize $\mathbf{w}_{T,0} = \mathbf{0}$.

for $t = 1$ **to** N **do**

 Sample $(\mathbf{x}_t, y_t) \sim \rho_{\mathbf{x} \times \mathbf{y}}$.

 Compute step size $\gamma_t \leftarrow \gamma_0 \cdot \mathcal{S}(t)$.

 Update $\mathbf{w}_T$ via SGD on ground truth:

$\mathbf{w}_{T,t} \leftarrow \mathbf{w}_{T,t-1} - \gamma_t \nabla_{\mathbf{w}} \mathcal{R}_T(\mathbf{w}_{T,t-1}; \mathbf{x}_t, y_t)$.

end for

Fix Teacher parameters $\mathbf{w}_T \leftarrow \mathbf{w}_{T,N}$.

Phase 2: Student Transfer Learning

Initialize $\mathbf{w}_{S,0} = \mathbf{0}$.

for $t = 1$ **to** n **do**

 Sample $\mathbf{x}_t \sim \rho_{\mathbf{x}}$.

 Generate teacher label: $\hat{y}_t \leftarrow \mathbf{w}_T^\top \phi_T(\mathbf{x}_t)$.

 Compute step size $\gamma'_t \leftarrow \gamma'_0 \cdot \mathcal{S}(t)$.

 Update $\mathbf{w}_S$ via SGD on teacher labels:

$\mathbf{w}_{S,t} \leftarrow \mathbf{w}_{S,t-1} - \gamma'_t \nabla_{\mathbf{w}} \mathcal{R}_{\text{Trans}}(\mathbf{w}_{S,t-1}; \mathbf{x}_t, \hat{y}_t)$.

end for

Remark 1. Assumption 1 is widely adopted in works related to linear regression (Bartlett et al., 2020; Tsigler & Bartlett, 2022; Zou et al., 2021; Wu et al., 2022). It characterizes the kurtosis of the feature distribution, ensuring that the variance of the stochastic gradient operator is bounded. This condition is strictly weaker than assuming Gaussianity (where $\psi \approx 3$ by Wick’s theorem) and holds for bounded or sub-Gaussian distributions.

3.3. Preliminaries on the SGD Dynamics

We analyze the standard SGD recursion for a generic model f_ν ($\nu \in \{T, S\}$) in the capacity-limited regime, adapting the framework of Ge et al. (2019) and Wu et al. (2022). Let $\boldsymbol{\eta}_t := \mathbf{w}_t - \mathbf{w}^*$ be the error vector and $\mathbf{P}_N := \mathbb{E}[\boldsymbol{\eta}_N^{\otimes 2}]$ be the second moment matrix.

Lemma 1 (Direct Learning Dynamics). *Under Assumption 1, the expected excess risk satisfies:*

$$\mathbb{E}[\mathcal{E}_\nu(\mathbf{w}_N)] = \frac{1}{2} \langle \mathbf{P}_{N,\nu}, \boldsymbol{\Sigma}_\nu \rangle \leq \underbrace{\frac{\|\boldsymbol{\eta}_0\|_{\boldsymbol{\Sigma}_\nu^{\leq k_\nu^*}}^2}{2N^2}}_{\text{Decayed Head}} + \underbrace{\frac{\|\boldsymbol{\eta}_0\|_{\boldsymbol{\Sigma}_\nu^{> k_\nu^*}}^2}{2}}_{\text{Preserved Tail}} + \underbrace{16C_{\text{noise}} \left(\frac{k_\nu^*}{K} + K\gamma_0^2 \sum_{i>k_\nu^*} \lambda_{i,\nu}^2 \right)}_{\text{Variance}},$$

where $k_\nu^* = \max\{k : \lambda_{k,\nu} > 2 \ln N \log_2 N / (\gamma_0 N)\}$, and $C_{\text{noise}} := \psi \|\mathbf{w}_0 - \mathbf{w}^*\|_{\boldsymbol{\Sigma}}^2 + \sigma_{\text{eff}}^2$.

Lemma 1 reveals two mechanisms: (1) **Implicit Regularization**. The bias splits at the learning progress k^* , learning the top k^* modes ($O(N^{-2})$ decay) while leaving trailing modes ($> k^*$) untouched. (2) **Effective Noise**. The variance term C_{noise} aggregates algorithmic noise (scaling with

estimation error $\psi\|\eta_0\|^2$) and our key insight, the **effective noise variance** $\sigma_{\text{eff}}^2 := \sigma^2 + \|\Pi^\perp \mathbf{w}_*\|^2$, reflecting that the unlearnable component of the target acts as an irreducible noise to the student’s optimization process.

4. General Knowledge Transfer

We now establish a unified framework for teacher-to-student training. Before analyzing the dynamics, we shall quantify the fundamental limit of knowledge transfer from a static perspective. Unlike direct learning, where the student targets the ground truth $\mathbf{w}_S^* = (\mathbf{M}_S^\top)^+ \mathbf{w}_*$, transfer learning targets the teacher’s best approximation $\mathbf{w}_{\text{T2S}}^* = (\mathbf{M}_S^\top)^+ \mathbf{M}_T^\top \mathbf{w}_*$. The discrepancy between these two optima dictates the maximum potential utility of the teacher.

Lemma 2 (Geometric Consistency Condition). *Knowledge transfer is consistent with direct learning (i.e., $\mathbf{w}_{\text{T2S}}^* = \mathbf{w}_S^*$) if and only if:*

$$\Pi_S \Pi_T^\perp \mathbf{w}_* = \mathbf{0}. \quad (9)$$

*If this condition is violated, the student trained by transfer suffers an irreducible **Static Alignment Bias** compared to the student trained on ground truth, even with infinite data:*

$$\mathcal{R}_{\text{T2S}}(\mathbf{w}_{\text{T2S}}^*) - \mathcal{R}_S(\mathbf{w}_S^*) = \frac{1}{2} \|\Pi_S \Pi_T^\perp \mathbf{w}_*\|^2 > 0. \quad (10)$$

This lemma implies that if the Teacher is “blind” to concepts $(\Pi_T^\perp \mathbf{w}_*)$ that the Student is capable of representing (Π_S) , the Student is structurally handicapped by the Teacher. Efficient transfer thus requires $\|\Pi_S \Pi_T^\perp \mathbf{w}_*\|$ to be negligible.

The following theorem extends this static intuition to the dynamic regime, providing the central decomposition.

Theorem 1 (The Teacher-to-Student Risk Decomposition). *Let the Student be trained on n samples labeled by a fixed Teacher (itself trained on N samples). For any $0 < \delta < 1/2$, the expected knowledge transfer excess risk $\mathcal{E}_{\text{T2S}}$ is bounded by:*

$$\begin{aligned} & \mathbb{E}[\mathcal{E}_{\text{T2S}}(N, n)] \quad (11) \\ & \leq \underbrace{\frac{1}{2} \langle \mathbf{P}_{N,T}, \mathbf{M}_T \Pi_S^{\leq k^*} \mathbf{M}_T^\top \rangle + 2\delta^2 \langle \mathbf{P}_{N,T}, \mathbf{M}_T \Pi_S^{> k^*} \mathbf{M}_T^\top \rangle}_{\text{(I) Propagated Teacher Error}} \\ & + \underbrace{\frac{1}{2} \langle \mathbf{P}_{n,\text{T2S}}, \Sigma_S^{\leq k^*} \rangle + \langle \mathbf{P}_{n,\text{T2S}}^{\text{var}}, \Sigma_S^{> k^*} \rangle}_{\text{(II) Student Optimization Error}} \\ & + C \underbrace{\left\| \Pi_S^{\leq k^*} \Pi_T^\perp \mathbf{w}_* \right\| + \left\| \Pi_S^{> k^*} \mathbf{w}_* \right\|^2 + 2\delta^2 \left\| \Pi_S^{> k^*} \Pi_T \mathbf{w}_* \right\|^2}_{\text{(III) Irreducible Alignment Bias}}, \end{aligned}$$

where $k^* = \max\{k : \lambda_{k,S} > \delta \log_2 n / (4\gamma'_0)\}$ is the effective dimension, $\mathbf{P}_{n,\text{T2S}} := \mathbb{E}[\eta_{n,\text{T2S}}^{\otimes 2}]$ is the second moment matrix of $\eta_{n,\text{T2S}} := \mathbf{w}_{n,\text{T2S}} - \mathbf{w}_{\text{T2S}}^*$ and $\mathbf{P}_{n,\text{T2S}}^{\text{var}} := \mathbb{E}[\eta_{n,\text{T2S}}^{\otimes 2}] - \mathbb{E}[\eta_{n,\text{T2S}}]^{\otimes 2}$ being the variance matrix of it, and C is a constant.

Decomposition Analysis. Theorem 1 dissects the transfer risk into three physically distinct components governed by

the student’s spectral learning progress k^* (the effective dimension learned by the student):

(I) Propagated Teacher Error. This term governs how the teacher’s error tensor $\mathbf{P}_{N,T} := \mathbb{E}[\eta_{N,T}^{\otimes 2}]$ is projected onto the student’s geometry, re-weighting the native risk measure $\frac{1}{2} \langle \mathbf{P}_{N,T}, \Sigma_T \rangle$, ($\Sigma_T = \mathbf{M}_T \mathbf{M}_T^\top$), from Lemma 1. The first term, $\frac{1}{2} \langle \mathbf{P}_{N,T}, \mathbf{M}_T \Pi_S^{\leq k^*} \mathbf{M}_T^\top \rangle$, indicates a **1-to-1 error inheritance** within the student’s learned subspace ($\leq k^*$). This implies that for learned concepts, the student directly copies the teacher’s precision, forming the foundation of **Knowledge Distillation** (Section 5) where a student benefits from a Strong Teacher’s lower error. Conversely, the second term $2\delta^2 \langle \mathbf{P}_{N,T}, \mathbf{M}_T \Pi_S^{> k^*} \mathbf{M}_T^\top \rangle$ reveals that in the unlearned tail subspace ($> k^*$), the teacher’s error is not fully inherited but **heavily dampened** by $\mathcal{O}(\delta^2)$. This suppression is the core mechanism of **Weak-to-Strong Generalization** (Section 6): it enables the student to filter out teacher’s high variance in the tail—performing *spectral denoising*—rather than blindly memorizing the noise.

(II) Student Optimization Error (Controllable): This captures the error of fitting the teacher’s labels. Unlike teacher error, this term is theoretically reducible to zero by increasing unlabeled data n ($n \gg N$ is common in semi-supervised settings). Meanwhile, we can tune the student’s step size γ'_0 , keeping the spectral cutoff k^* fixed.

(III) Irreducible Alignment Bias (The Geometric Mismatch): This term projects the static “Blind Spots” identified in Lemma 2 onto the student’s *learned* subspace $(\Pi_S^{\leq k^*})$. It confirms that the teacher’s structural deficiency $(\Pi_T^\perp \mathbf{w}_*)$ creates an immediate performance ceiling, forcing the student to discard learnable information to align with a deficient teacher. The remaining terms simply account for information in the unlearned subspace $(\Pi_S^{> k^*})$.

5. Knowledge Distillation

Next, we shift from general teacher-to-student transfer to the “Strong Teacher, Weak Student” paradigm. To ensure tractable analysis without sacrificing relevance to wide neural networks, we formalize the relationship between teacher and student feature spaces as follows.

Definition 1 (Spectrally Compatible Student). We define the student as **Spectrally Compatible** if its feature map is a linear projection $\phi_S(\mathbf{x}) = \mathbf{M}_{\text{Trans}} \phi_T(\mathbf{x})$ parameterized by a bounded operator $\mathbf{M}_{\text{Trans}} = \mathbf{Q}_{\text{Trans}} \Lambda_{\text{Trans}} \mathbf{Q}_T^\top \in \mathbb{R}^{d_S \times d_T}$, where $\mathbf{Q}_T$ is teacher’s eigen-basis, Λ_{Trans} is diagonal and $\mathbf{Q}_{\text{Trans}}$ is a rotation.

Remark 2. This compatibility assumption is rigorous in the Kernel/NTK limit (Jacot et al., 2018) and generalizes spectral bias analyses (Bordelon et al., 2020) by allowing student-specific basis rotations.

Assumption 2 (Relative Spectral Decay). *Let the eigenvalues decay as $\lambda_{k,\nu} \asymp k^{-\alpha_\nu}$ for $\nu \in \{T, S\}$. We define the **Strong Teacher, Weak Student** regime by $1 < \alpha_T \leq \alpha_S$.*

Remark 3 (Connection to Kernel Theory). This assumption parallels kernel theory where α_ν governs the RKHS capacity. As proven in Appendix E.2, this condition implies nested spaces: $\mathcal{H}_S \subseteq \mathcal{H}_T$. The Student thus possesses a smaller representation space and inferior expressive capability.

We first analyze the simple scenario where the student is strictly rank-deficient compared to the teacher, representing a hard bottleneck in representation capacity.

Theorem 2 (Benefit of Noise Reduction). *Consider a spectrally compatible student which captures a strict subspace of the teacher (i.e., Λ_{Trans} contains zero eigenvalues). Under Assumption 1,2, the distillation efficiency satisfies: $\lim_{n \rightarrow \infty} \text{DER}_N > 1$. This strict inequality signifies that distillation is provably more efficient than direct learning.*

Mechanism: Mitigating Approximation Noise. This efficiency gain exploits the *Effective Noise* phenomenon established in Lemma 1. Direct learning forces the Weak Student to process its large approximation bias ($\|\Pi_S^\perp \mathbf{w}_*\|^2$) as irreducible stochastic noise. Distillation effectively **substitutes** this high-variance term with the Strong Teacher’s negligible approximation bias ($\|\Pi_T^\perp \mathbf{w}_*\|^2 \ll \|\Pi_S^\perp \mathbf{w}_*\|^2$). By targeting the Teacher’s “cleaner” proxy, the Student bypasses the optimization instability caused by its own capacity limitations, strictly accelerating convergence.

We now extend our analysis to the hard scenario where distillation efficiency sometimes scales with model capacity.

Assumption 3 (The Mis-specified Regime). *Let $\beta_i^2 := \mathbb{E}[\langle \mathbf{w}_S^*, \mathbf{v}_{i,T} \rangle^2]$ be the target energy spectrum. We assume the task exhibits heavy-tailedness: $\beta_i^2 = \Theta(i^\beta)$ with $\beta \geq 0$.*

Remark 4. Real-world tasks (e.g., language modeling) exhibit “heavy-tailed” complexity, implying that the ground truth f^* lies outside the RKHS. Consequently, no finite-capacity model can perfectly fit such a target.

Theorem 3 (Asymptotic DER Rate). *Under Assumption 1 and 3, for a spectrally compatible student, if the task is learnable by the teacher ($\alpha_T > 1 + \beta$), the risk scales as:*

$$\lim_{n \rightarrow \infty} \mathbb{E}[\mathcal{E}_{\text{T2S}}(N, n)] = \tilde{O}(N^{-(\alpha_T - (1+\beta))/\alpha_T}). \quad (12)$$

Consequently, if the Student is strictly weaker ($\alpha_S > \alpha_T$), the Distillation Efficiency Ratio diverges:

$$\lim_{n \rightarrow \infty} \text{DER}_N = \tilde{\Omega}(N^\kappa), \quad (13)$$

where $\kappa = (\alpha_T - 1 - \beta)(1/\alpha_T - 1/\alpha_S) > 0$.

Interpretation: Spectral Horizon Expansion. Consider a simplified diagonalized regime where the shared “energy”

of the i -th target dimension is given by $I_i = \lambda_i w_{*,i}^2$. In the hard-learning limit (Assumption 3), the excess risk is strictly dominated by the *Preserved Tail* (Lemma 1)—the aggregate energy of unlearned frequencies: $\|\eta_0\|_{\Sigma_\nu^{>k_\nu^*}}^2 = \sum_{k > k_\nu^*} I_k = \tilde{O}((k_\nu^*)^{1+\beta-\alpha_T})$. This dependency identifies the **optimization horizon** k_ν^* as the sole bottleneck for accuracy. Moreover, the spectral decay rate α governs the **velocity** of this horizon’s expansion ($k_\nu^* \propto N^{1/\alpha_\nu}$). While the Weak Student stagnates at a “Spectral Barrier” due to rapid decay (k_S^* grows slowly), the Strong Teacher expands its effective optimization horizon significantly faster ($k_T^* \gg k_S^*$). Through distillation (Theorem 1), the Student **inherits this expansive horizon**, effectively bypassing its intrinsic capacity limitations to resolve high-frequency modes that are otherwise inaccessible from raw data.

6. Weak-to-Strong Generalization and Self-Distillation

We now address the “Weak-to-Strong” (W2S) phenomenon (Burns et al., 2023), and its limit case, **Self-Distillation** (Furlanello et al., 2018), where the student matches or exceeds the supervisor’s capacity. Unlike the heavy-tailed pre-training regime (Section 5), both paradigms typically target specific downstream tasks (e.g., fine-tuning on reasoning, classification) rather than universal knowledge transfer. As established by Aghajanyan et al. (2020), these specialized tasks exhibit **Low Intrinsic Dimensionality**, residing on subspaces significantly smaller than the parameter space.

Assumption 4 (Sufficient Expressivity). *We assume teacher and student are structurally sufficient over the downstream task, i.e. $\Pi_T = \Pi_S = \mathbf{I}_D$.*

Assumption 5 (Low Intrinsic Dimension). *We assume the ground truth $\mathbf{w}_*$ is low dimensional in the view of the two models. Formally, there exists a cutoff $k_\nu^\dagger$, the signal tail is negligible: $\|\Pi_T^{>k_\nu^\dagger} \mathbf{w}_*\|_\Sigma^2 = 0$.*

Our analysis confirms that the Student surpasses the Teacher by eliminating optimization variance.

Theorem 4 (Weak-to-Strong Generalization Guarantee). *Let the Weak Teacher be trained on N samples and the Strong Student be distilled on n samples. Under Assumption 1, 4, 5, there exists a threshold n_0 such that for all $n > n_0$:*

$$\mathbb{E}[\mathcal{R}_{\text{T2S}}(N, n)] < \mathbb{E}[\mathcal{R}_T(N)]. \quad (14)$$

Mechanism: Denoising via δ -Damping. This guarantee stems directly from the *Propagated Teacher Error* decomposition in Theorem 1. As analyzed in Section 4, while the Weak Teacher suffers from full error globally ($\langle \mathbf{P}_{N,T}, \Sigma_T \rangle$), the Student only inherits this error in the low-dimensional signal subspace. In the noise-dominated spectral tail, the




Figure 1. In the KD regime, a teacher with slower spectral decay acts as a pre-conditioner to trigger Spectral Horizon Expansion, enabling the capture of statistically inaccessible high-frequency signals. In the W2S regime, a low task-specific intrinsic dimensionality allows the student to leverage its excess capacity for Spectral Denoising, filtering out optimization noise to recover the underlying geometry.

Student suppresses the Teacher’s fluctuations by the factor $\mathcal{O}(\delta^2)$. This **spectral filtering** ensures that the Student recovers a “cleaner” estimate of the ground truth than the noisy Teacher, strictly reducing the excess risk.

We next quantify the magnitude of this gain by **PGR**:

Theorem 5 (Optimal W2S Rates). *Under Assumption 1, 4, 5, with optimal early stopping sample size $n = \tilde{\mathcal{O}}((k^\dagger)^{2\alpha_S/(2\alpha_S+1)} \cdot N^{1/[\alpha_T(2\alpha_S+1)]})$, the Student’s risk and the resulting PGR scale as:*

$$\min_n \mathbb{E}[\mathcal{R}_{T2S}] = \tilde{\mathcal{O}}\left((k^\dagger)^{2\alpha_S/(2\alpha_S+1)} \cdot N^{1/[\alpha_T(2\alpha_S+1)]-1}\right),$$

$$\text{PGR} = 1 - \tilde{\mathcal{O}}\left((k^\dagger)^{2\alpha_S/(2\alpha_S+1)} \cdot N^{-\Delta_{\text{rate}}}\right), \quad (15)$$

where the rate gain is $\Delta_{\text{rate}} = 2\alpha_S/[\alpha_T(2\alpha_S+1)] > 0$.

Interpretation: The Variance-Damping Trade-off. Since the downstream task is low-dimensional (Assumption 5), the bias vanishes rapidly, leaving the convergence governed entirely by a variance trade-off driven by the student’s learning progress k^* . In the learned subspace ($\leq k^*$), the student suffers **Inherited Variance**, where the teacher’s noise is projected 1-to-1; as the student learns more ($k^* \uparrow$), this accumulated error $\mathcal{O}(k^*/N)$ increases linearly. Conversely, in the unlearned tail ($> k^*$), the student benefits from **Damped Variance**: the teacher’s total variance rate $\tilde{\mathcal{O}}(N^{\frac{1-\alpha_T}{\alpha_T}})$ is suppressed by the factor $\delta(k^*) = (k^\dagger/k^*)^{2\alpha_S}$, which *shrinks* as k^* extends. The optimal early stopping point n thus represents the **equilibrium** of these opposing forces, minimizing the total risk scale: $\frac{k_S^*}{N} + (k^\dagger/k_S^*)^{2\alpha_S} \cdot \tilde{\mathcal{O}}(N^{\frac{1-\alpha_T}{\alpha_T}})$. Moreover, since $\Delta_{\text{rate}} > 0$, the penalty term vanishes asymptotically ($N \rightarrow \infty$), implying **Full Capability Recovery** ($\lim_N \text{PGR} = 1$).

Corollary 1 (Self-Distillation as Pure Denoising). *Our framework naturally explains the effectiveness of Self-Distillation, where the student shares the teacher’s architecture ($\alpha_S = \alpha_T$). Substituting this into Theorem 5 yields a strictly positive rate gain:*

$$\Delta_{\text{SD}} = \frac{2}{2\alpha_T + 1} > 0. \quad (16)$$

7. Experiments

In this section, we conduct both synthetic and real-world experiments to empirically validate the assumptions and theoretical results of knowledge transfer. We focus on two distinct scenarios: knowledge distillation and weak-to-strong generalization. More details are available in Appendix F.

7.1. Knowledge Distillation

Synthetic Experiments for Knowledge Distillation. We consider a teacher-student framework for synthetic experiments in Figure 2a and Figure 2b. Given input $\mathbf{x} \in \mathbb{R}^d$ ($d = 100$) drawn independently from a standard normal distribution, the noisy labels are generated based on teacher’s feature: $y = \langle \mathbf{w}^*, \Sigma_T^{1/2} \mathbf{x} \rangle + \xi$. The teacher model is defined as $f_T(\mathbf{x}; \mathbf{w}_T) = \langle \mathbf{w}_T, \phi_T(\mathbf{x}) \rangle$. Similarly, the student model is formulated as $f_S(\mathbf{x}; \mathbf{w}_S) = \langle \mathbf{w}_S, \phi_S(\mathbf{x}) \rangle$. Both models are linear and initialized from zero. The teacher model is pretrained on N samples in a one-pass manner, and the student model is trained on both the original-labeled training data (samples size $N = 2000$ in Figure 2a and $N \in [0, 2000]$ in Figure 2b) and teacher-labeled data (samples size $n = 500, 000$). We plot the excess risk during training on a log scale.

Real-World Experiments for Knowledge Distillation. As shown in Figure 2c, we investigate spectral decay across models on a real-world dataset: UTKFace dataset (Zhang et al., 2017). This experiment is conducted using ViT-B/16, ViT-L/16 (Dosovitskiy et al., 2021), ResNet18 and ResNet50 (He et al., 2016) as representative architectures. All models are initialized with publicly pre-trained weights (maintainers & contributors, 2016). We extract features $\mathbf{f}(\mathbf{x}_i; \theta)$ and calculate empirical feature covariance matrix as $\hat{\Sigma} = \frac{1}{m} \sum_i \mathbf{f}(\mathbf{x}_i; \theta) \mathbf{f}(\mathbf{x}_i; \theta)^\top$ ($m = 2000$). The top 500 eigenvalues of $\hat{\Sigma}$ are fitted and visualized on a log-log plot. Figure 2d presents the age regression experiment on UTK-Face. We fine-tuned the pre-trained ViT-L/16 as teacher model (dotted line) and train student models via linear probing under two settings: directly from the data (dashed line) and from the teacher’s predictions (solid line).




Figure 2. Knowledge Distillation Enables Risk Inheritance. Figure 2a and Figure 2b show that a student model trained on a sufficiently large set ($n = 500,000$) of teacher-labeled data achieves excess risk comparable to that of the teacher, outperforming directly training on the original data ($N = 2000$). The advantage is further amplified as the growth of N and spectral decay differences $\Delta\alpha = \alpha_S - \alpha_T$. **Disparity in Spectral Behavior and Expressivity.** Figure 2c plots spectral decay exponent α of empirical feature covariance across models of different capacities. Models with stronger expressivity exhibit slower spectral decay (smaller α), while weaker models instead exhibit faster decay (larger α). Figure 2d demonstrates that knowledge distillation consistently improves students’ performance, with the gains increasing as $\Delta\alpha$ becomes larger.




Figure 3. W2S Benefits from Early Stopping. Figure 3a shows that the optimal performance is attained at an intermediate stage of training, indicating that early stopping is crucial for weak-to-strong generalization, particularly for students with smaller intrinsic dimensionality. **Low Intrinsic Dimensionality.** Figure 3b illustrates that projection energy concentrates in a low-dimensional space during weak-to-strong training. Figure 3c further evaluate the dimension required to retain 90% of performance and demonstrate that models with higher task performance exhibit lower intrinsic dimensionality. Figure 3d presents the **PGR** of these models, showing that weak-to-strong training recovers a greater portion of the generalization gap between the weak teacher and the strong ceiling model.

7.2. Weak to Strong Generalization

Synthetic Experiments for W2S. In Figure 3a we investigate W2S following the teacher–student linear regression framework and data generation process defined in synthetic experiment of knowledge distillation. The task has a low intrinsic dimension k , which means ground truth parameter $\mathbf{w}^*$ is chosen to have nonzero entries restricted to the first k coordinates. We present a plot of the analytical excess risk during training for different values of $k \in \{1, 10, 20, 30, 50\}$.

Real-World Experiments for W2S. Figure 3b visualizes the evolution of projection energy (the projection of weight onto the sorted eigen space of the student’s feature covariance) during W2S training. For the UTKFace age regression, we designate a ResNet18 as the “weak” teacher and a CLIP-ViT-B/32 (Radford et al., 2021) as the “strong” student. Pretrained weights (Wolf et al., 2020; maintainers & contributors, 2016) are utilized and trained via linear probing. We apply PowerNorm (Gonzalez & Woods, 2008) for better visualization and identify the minimum dimension achieving 80% of the accumulated energy.

Figure 3c shows the task performance of a diverse collec-

tion of pre-trained models and corresponding intrinsic dimensions. We load pretrained weights (Caron et al., 2021; maintainers & contributors, 2016) for ResNet50-DINO, ViT-Small/16-DINO, ViT-Base/16-DINO (Caron et al., 2021), RegNetY-16GF (Radosavovic et al., 2020) and EfficientNet-B4 (Tan & Le, 2019) and get the features $\mathbf{f}(\mathbf{x}_i; \theta)$. We first perform PCA on the feature covariance $\hat{\Sigma} = \mathbf{U}\mathbf{A}\mathbf{U}^\top$. Following the weak-to-strong training (Burns et al., 2023), we train only the linear classifier on $\mathbf{U}_{[1:k,:]}^\top \mathbf{f}(\mathbf{x}_i; \theta)$ to identify the smallest k -dimensional subspace that retains at least 90% of the full-dimensional performance after 20 training epochs. Finally, we use pretrained AlexNet (Krizhevsky et al., 2012; maintainers & contributors, 2016) as “weak” teacher and track the evolution of the **PGR** throughout training in Figure 3d.

8. Final Discussions

Insight: What Makes a Strong Model? Beyond the spectral decay rate α_ν which signifies the volume of the RKHS in kernel theory, our bound highlights the critical role of representation quality, encapsulated by the effective signal dimension $k^\dagger$. Since the risk scales with $(k^\dagger)^{\frac{2\alpha_S}{2\alpha_S+1}}$, a more

concentrated feature representation (smaller $k^\dagger$) directly amplifies the denoising efficiency ($\delta \propto (k^\dagger)^{2\alpha_S}$). This implies that a "Strong" Student is not merely one with high capacity (small α_ν), but one whose feature geometry *compresses* the task signal into a low-dimensional subspace.

Three Signatures of a Strong Model. Finally, our theoretical analysis provides a direct answer to our titular question: *What makes a strong model?* By translating our spectral theorems into practical indicators, we identify three key properties that define model strength. First, we highlight the role of **Representation Rank** (Theorem 2): a higher rank of the feature representation implies comprehensive feature coverage, which minimizes the **effective label noise** arising from uncaptured geometric dimensions—aligning with the empirical success of RankMe (Garrido et al., 2023). Second, we pinpoint the **Spectral Decay Rate** (Theorem 3): stronger models exhibit slower eigenvalue decay (a "heavier" tail), enabling the capture of complex features, consistent with the α -ReQ metric proposed by Agrawal et al. (2022). Finally, we introduce a novel, third metric: the **Task-Specific Intrinsic Dimension** (Theorem 5). Unlike generic capacity measures, this metric reveals that true strength is context-dependent—a model is "strong" for a specific task if its principal eigenspace aligns efficiently with the task’s ground truth geometry.

Principled Spectral Indicators for Model Strength. We derive three theoretically grounded indicators to quantify what makes a model "strong" for transfer learning: (i) **Representation Rank**, where a higher rank minimizes the effective noise induced by geometric misalignment (providing a theoretical basis for Garrido et al. (2023)); (ii) **Spectral Decay Speed**, where a slower decay ($\alpha \approx 1$) implies richer feature learning (consistent with Agrawal et al. (2022)); and (iii) **Task-Specific Intrinsic Dimension** (Theorem 5), our novel metric which demonstrates that model strength is not merely an intrinsic property but depends on the low-dimensional geometric alignment between the model’s spectrum and the target task.

Acknowledgments

C. Fang was supported by the National Natural Science Foundation of China (NSFC) under Grant Nos. 92470117 and 62376008. This work was also supported in part by the Beijing Major Science and Technology Project under Contract no. Z251100008125007.

Impact Statement

This paper presents work whose goal is to advance the field of machine learning. There are many potential societal consequences of our work, none of which we feel must be specifically highlighted here.

References

- Abdin, M., Aneja, J., Awadalla, H., Awadallah, A., Awan, A. A., Bach, N., Bahree, A., Bakhtiari, A., Bao, J., Behl, H., Benhaim, A., Bilenko, M., Bjorck, J., Bubeck, S., Cai, M., Cai, Q., Chaudhary, V., Chen, D., Chen, D., Chen, W., Chen, Y.-C., Chen, Y.-L., Cheng, H., Chopra, P., Dai, X., Dixon, M., Eldan, R., Frago, V., Gao, J., Gao, M., Gao, M., Garg, A., Giorno, A. D., Goswami, A., Gunasekar, S., Haider, E., Hao, J., Hewett, R. J., Hu, W., Huynh, J., Iter, D., Jacobs, S. A., Javaheripi, M., Jin, X., Karampatziakis, N., Kauffmann, P., Khademi, M., Kim, D., Kim, Y. J., Kurilenko, L., Lee, J. R., Lee, Y. T., Li, Y., Li, Y., Liang, C., Liden, L., Lin, X., Lin, Z., Liu, C., Liu, L., Liu, M., Liu, W., Liu, X., Luo, C., Madan, P., Mahmoudzadeh, A., Majercak, D., Mazzola, M., Mendes, C. C. T., Mitra, A., Modi, H., Nguyen, A., Norick, B., Patra, B., Perez-Becker, D., Portet, T., Pryzant, R., Qin, H., Radmilac, M., Ren, L., de Rosa, G., Rosset, C., Roy, S., Ruwase, O., Saarikivi, O., Saied, A., Salim, A., Santacrose, M., Shah, S., Shang, N., Sharma, H., Shen, Y., Shukla, S., Song, X., Tanaka, M., Tupini, A., Vaddamanu, P., Wang, C., Wang, G., Wang, L., Wang, S., Wang, X., Wang, Y., Ward, R., Wen, W., Witte, P., Wu, H., Wu, X., Wyatt, M., Xiao, B., Xu, C., Xu, J., Xu, W., Xue, J., Yadav, S., Yang, F., Yang, J., Yang, Y., Yang, Z., Yu, D., Yuan, L., Zhang, C., Zhang, C., Zhang, J., Zhang, L. L., Zhang, Y., Zhang, Y., Zhang, Y., and Zhou, X. Phi-3 technical report: A highly capable language model locally on your phone, 2024. URL <https://arxiv.org/abs/2404.14219>.
- Aghajanyan, A., Zettlemoyer, L., and Gupta, S. Intrinsic dimensionality explains the effectiveness of language model fine-tuning, 2020. URL <https://arxiv.org/abs/2012.13255>.
- Agrawal, K. K., Mondal, A. K., Ghosh, A., and Richards, B. A. α -req : Assessing representation quality in self-supervised learning by measuring eigenspectrum decay. In *Neural Information Processing Systems*, 2022. URL <https://api.semanticscholar.org/CorpusID:258509089>.
- Allen-Zhu, Z. and Li, Y. Towards understanding ensemble, knowledge distillation and self-distillation in deep learning, 2023. URL <https://arxiv.org/abs/2012.09816>.
- Atanasov, A., Zavatone-Veth, J. A., and Pehlevan, C. Scaling and renormalization in high-dimensional regression. *ArXiv*, abs/2405.00592, 2024. URL <https://api.semanticscholar.org/CorpusID:269484262>.
- Bartlett, P. L., Long, P. M., Lugosi, G., and Tsigler, A. Benign overfitting in linear regression. *Proceedings of the National Academy of Sciences*, 117(48): 30063–30070, April 2020. ISSN 1091-6490. doi: 10.1073/pnas.1907378117. URL <http://dx.doi.org/10.1073/pnas.1907378117>.
- Bietti, A. and Mairal, J. On the inductive bias of neural tangent kernels. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Bordelon, B., Canatar, A., and Pehlevan, C. Spectrum dependent learning curves in kernel regression and wide neural networks. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, pp. 1024–1034. PMLR, 2020.
- Burns, C., Izmailov, P., Kirchner, J. H., Baker, B., Gao, L., Aschenbrenner, L., Chen, Y., Ecoffet, A., Joglekar, M. R., Leike, J., Sutskever, I., Wu, J., and OpenAI. Weak-to-strong generalization: Eliciting strong capabilities with weak supervision. *ArXiv*, abs/2312.09390, 2023. URL <https://api.semanticscholar.org/CorpusID:266312608>.
- Caron, M., Touvron, H., Misra, I., Jégou, H., Mairal, J., Bojanowski, P., and Joulin, A. Emerging properties in self-supervised vision transformers. In *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 9650–9660, 2021.
- Charikar, M., Pabbaraju, C., and Shiragur, K. Quantifying the gain in weak-to-strong generalization, 2024. URL <https://arxiv.org/abs/2405.15116>.
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pp. 248–255. Ieee, 2009.
- Dieuleveut, A. and Bach, F. Non-parametric stochastic approximation with large step sizes, 2016. URL <https://arxiv.org/abs/1408.0361>.
- Dong, Y., Li, Y., Li, Y., Lee, J. D., and Lei, Q. Discrepancies are virtue: Weak-to-strong generalization through lens of intrinsic dimension, 2025. URL <https://arxiv.org/abs/2502.05075>.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., and Houshy, N. An image is worth 16x16 words: Transformers for image recognition at scale, 2021. URL <https://arxiv.org/abs/2010.11929>.
- Furlanello, T., Lipton, Z. C., Tschannen, M., Itti, L., and Anandkumar, A. Born again neural networks, 2018. URL <https://arxiv.org/abs/1805.04770>.

- Garrido, Q., Balestrieri, R., Najman, L., and Lecun, Y. Rankme: Assessing the downstream performance of pretrained self-supervised representations by their rank, 2023. URL <https://arxiv.org/abs/2210.02885>.
- Ge, R., Kakade, S. M., Kidambi, R., and Netrapalli, P. The step decay schedule: A near optimal, geometrically decaying learning rate procedure for least squares. In *Neural Information Processing Systems*, 2019.
- Gonzalez, R. and Woods, R. *Digital Image Processing*. Prentice Hall, 2008. ISBN 9780131687288. URL <https://books.google.com.hk/books?id=8uGOnjRGEzoC>.
- Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., Xu, R., Zhang, R., Ma, S., Bi, X., Zhang, X., Yu, X., Wu, Y., Wu, Z. F., Gou, Z., Shao, Z., Li, Z., Gao, Z., Liu, A., Xue, B., Wang, B., Wu, B., Feng, B., Lu, C., Zhao, C., Deng, C., Ruan, C., Dai, D., Chen, D., Ji, D., Li, E., Lin, F., Dai, F., Luo, F., Hao, G., Chen, G., Li, G., Zhang, H., Xu, H., Ding, H., Gao, H., Qu, H., Li, H., Guo, J., Li, J., Chen, J., Yuan, J., Tu, J., Qiu, J., Li, J., Cai, J. L., Ni, J., Liang, J., Chen, J., Dong, K., Hu, K., You, K., Gao, K., Guan, K., Huang, K., Yu, K., Wang, L., Zhang, L., Zhao, L., Wang, L., Zhang, L., Xu, L., Xia, L., Zhang, M., Zhang, M., Tang, M., Zhou, M., Li, M., Wang, M., Li, M., Tian, N., Huang, P., Zhang, P., Wang, Q., Chen, Q., Du, Q., Ge, R., Zhang, R., Pan, R., Wang, R., Chen, R. J., Jin, R. L., Chen, R., Lu, S., Zhou, S., Chen, S., Ye, S., Wang, S., Yu, S., Zhou, S., Pan, S., Li, S. S., Zhou, S., Wu, S., Yun, T., Pei, T., Sun, T., Wang, T., Zeng, W., Liu, W., Liang, W., Gao, W., Yu, W., Zhang, W., Xiao, W. L., An, W., Liu, X., Wang, X., Chen, X., Nie, X., Cheng, X., Liu, X., Xie, X., Liu, X., Yang, X., Li, X., Su, X., Lin, X., Li, X. Q., Jin, X., Shen, X., Chen, X., Sun, X., Wang, X., Song, X., Zhou, X., Wang, X., Shan, X., Li, Y. K., Wang, Y. Q., Wei, Y. X., Zhang, Y., Xu, Y., Li, Y., Zhao, Y., Sun, Y., Wang, Y., Yu, Y., Zhang, Y., Shi, Y., Xiong, Y., He, Y., Piao, Y., Wang, Y., Tan, Y., Ma, Y., Liu, Y., Guo, Y., Ou, Y., Wang, Y., Gong, Y., Zou, Y., He, Y., Xiong, Y., Luo, Y., You, Y., Liu, Y., Zhou, Y., Zhu, Y. X., Huang, Y., Li, Y., Zheng, Y., Zhu, Y., Ma, Y., Tang, Y., Zha, Y., Yan, Y., Ren, Z. Z., Ren, Z., Sha, Z., Fu, Z., Xu, Z., Xie, Z., Zhang, Z., Hao, Z., Ma, Z., Yan, Z., Wu, Z., Gu, Z., Zhu, Z., Liu, Z., Li, Z., Xie, Z., Song, Z., Pan, Z., Huang, Z., Xu, Z., Zhang, Z., and Zhang, Z. Deepseek-r1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645(8081):633–638, September 2025. ISSN 1476-4687. doi: 10.1038/s41586-025-09422-z. URL <http://dx.doi.org/10.1038/s41586-025-09422-z>.
- He, K., Zhang, X., Ren, S., and Sun, J. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Hinton, G., Vinyals, O., and Dean, J. Distilling the knowledge in a neural network, 2015. URL <https://arxiv.org/abs/1503.02531>.
- Ildiz, M. E., Gozeten, H. A., Taga, E. O., Mondelli, M., and Oymak, S. High-dimensional analysis of knowledge distillation: Weak-to-strong generalization and scaling laws, 2025. URL <https://arxiv.org/abs/2410.18837>.
- Jacot, A., Gabriel, F., and Hongler, C. Neural tangent kernel: Convergence and generalization in neural networks. In *Advances in neural information processing systems*, 2018.
- Jain, P., Kakade, S. M., Kidambi, R., Netrapalli, P., Pillutla, V. K., and Sidford, A. A markov chain theory approach to characterizing the minimax optimality of stochastic gradient descent (for least squares). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2018a. doi: 10.4230/LIPICS.FSTTCS.2017.2. URL <https://drops.dagstuhl.de/entities/document/10.4230/LIPICS.FSTTCS.2017.2>.
- Jain, P., Kakade, S. M., Kidambi, R., Netrapalli, P., and Sidford, A. Accelerating stochastic gradient descent for least squares regression. In *Conference on Learning Theory*, 2018b.
- Krizhevsky, A., Sutskever, I., and Hinton, G. E. Imagenet classification with deep convolutional neural networks. *Advances in neural information processing systems*, 25, 2012.
- Li, X., Deng, Y., Wu, J., Zhou, D., and Gu, Q. Risk bounds of accelerated sgd for overparameterized linear regression, 2023. URL <https://arxiv.org/abs/2311.14222>.
- Lin, L., Wu, J., Kakade, S. M., Bartlett, P. L., and Lee, J. D. Scaling laws in linear regression: Compute, parameters, and data, 2025. URL <https://arxiv.org/abs/2406.08466>.
- maintainers, T. and contributors. Torchvision: Pytorch’s computer vision library. <https://github.com/pytorch/vision>, 2016.
- Malladi, S., Wetteg, A., Yu, D., Chen, D., and Arora, S. A kernel-based view of language model fine-tuning, 2023. URL <https://arxiv.org/abs/2210.05643>.
- Medvedev, M., Lyu, K., Yu, D., Arora, S., Li, Z., and Srebro, N. Weak-to-strong generalization even in random feature networks, provably, 2025. URL <https://arxiv.org/abs/2503.02877>.

- Menon, A. K., Rawat, A. S., Reddi, S. J., Kim, S., and Kumar, S. Statistical perspective on distillation. In *International Conference on Machine Learning*, pp. 7651–7662. PMLR, 2021.
- Mobahi, H., Farajtabar, M., and Bartlett, P. L. Self-distillation amplifies regularization in hilbert space, 2020. URL <https://arxiv.org/abs/2002.05715>.
- Moniri, B. and Hassani, H. On the mechanisms of weak-to-strong generalization: A theoretical perspective, 2025. URL <https://arxiv.org/abs/2505.18346>.
- Nagarajan, V., Menon, A. K., Bhojanapalli, S., Mobahi, H., and Kumar, S. On student-teacher deviations in distillation: does it pay to disobey?, 2024. URL <https://arxiv.org/abs/2301.12923>.
- Pan, R., Ye, H., and Zhang, T. Eigencurve: Optimal learning rate schedule for sgd on quadratic objectives with skewed hessian spectrums, 2022. URL <https://arxiv.org/abs/2110.14109>.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pp. 8748–8763. PmLR, 2021.
- Radosavovic, I., Kosaraju, R. P., Girshick, R., He, K., and Dollár, P. Designing network design spaces. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 10428–10436, 2020.
- Stanton, S., Izmailov, P., Kirichenko, P., Alemi, A. A., and Wilson, A. G. Does knowledge distillation really work?, 2021. URL <https://arxiv.org/abs/2106.05945>.
- Tan, M. and Le, Q. Efficientnet: Rethinking model scaling for convolutional neural networks. In *International conference on machine learning*, pp. 6105–6114. PMLR, 2019.
- Tsigler, A. and Bartlett, P. L. Benign overfitting in ridge regression, 2022. URL <https://arxiv.org/abs/2009.14286>.
- Wang, Y., Kordi, Y., Mishra, S., Liu, A., Smith, N. A., Khashabi, D., and Hajishirzi, H. Self-instruct: Aligning language models with self-generated instructions, 2023. URL <https://arxiv.org/abs/2212.10560>.
- Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T., Louf, R., Funtowicz, M., Davison, J., Shleifer, S., von Platen, P., Ma, C., Jernite, Y., Plu, J., Xu, C., Scao, T. L., Gugger, S., Drame,
- M., Lhoest, Q., and Rush, A. M. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pp. 38–45, Online, October 2020. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/2020.emnlp-demos.6>.
- Wu, D. X. and Sahai, A. Provable weak-to-strong generalization via benign overfitting. *ArXiv*, abs/2410.04638, 2024. URL <https://api.semanticscholar.org/CorpusID:273185855>.
- Wu, J., Zou, D., Braverman, V., Gu, Q., and Kakade, S. Last iterate risk bounds of SGD with decaying stepsize for overparameterized linear regression. In *International Conference on Machine Learning*, 2022.
- Zhang, H., Liu, Y., Chen, Q., and Fang, C. The optimality of (accelerated) sgd for high-dimensional quadratic optimization, 2024. URL <https://arxiv.org/abs/2409.09745>.
- Zhang, H., Lin, W., Liu, Y., and Fang, C. Learning curves of stochastic gradient descent in kernel regression, 2025. URL <https://arxiv.org/abs/2505.22048>.
- Zhang, Z., Song, Y., and Qi, H. Age progression/regression by conditional adversarial autoencoder. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 5810–5818, 2017.
- Zou, D., Wu, J., Braverman, V., Gu, Q., and Kakade, S. Benign overfitting of constant-stepsize SGD for linear regression. In *Conference on Learning Theory*, 2021.

A. Proofs of Results in Direct Learning

In this section we analyze the standard SGD recursion for a generic model f_ν ($\nu \in \{\mathcal{T}, \mathcal{S}\}$) in the capacity-limited regime, adapting the framework of Ge et al. (2019) and Wu et al. (2022).

A.0.1. LINEAR OPERATORS

To rigorously analyze the evolution of the second moment $\mathbb{E}[\boldsymbol{\eta}_t \otimes \boldsymbol{\eta}_t]$, we define a set of linear operators acting on the space of symmetric matrices (vectorized via Kronecker products).

Definition 2 (Iteration Operators). We define the stochastic update operator for the second moment as $\widehat{\mathcal{A}}_t := \mathbb{E}[\widehat{\mathbf{A}}_t \otimes \widehat{\mathbf{A}}_t]$, where $\widehat{\mathbf{A}}_t = \mathbf{I} - \gamma_t \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top$. Its deterministic counterpart, governing the mean flow variance, is defined as:

$$\mathcal{A}_t := \mathbf{A}_t \otimes \mathbf{A}_t = (\mathbf{I} - \gamma_t \boldsymbol{\Sigma}) \otimes (\mathbf{I} - \gamma_t \boldsymbol{\Sigma}), \quad (17)$$

where $\mathbf{A}_t = \mathbb{E}[\widehat{\mathbf{A}}_t] = \mathbf{I} - \gamma_t \boldsymbol{\Sigma}$.

Definition 3 (Fourth Moment Operators). To capture the noise geometry induced by the stochastic gradients, we define the fourth moment operator $\widehat{\mathcal{M}}_t$ and its deterministic version $\mathcal{M}_t$:

$$\widehat{\mathcal{M}}_t := \gamma_t^2 \mathbb{E}[(\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top) \otimes (\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top)], \quad (18)$$

$$\mathcal{M}_t := \gamma_t^2 (\boldsymbol{\Sigma} \otimes \boldsymbol{\Sigma}). \quad (19)$$

Definition 4 (Operator Variance). We define the centered variance operator $\mathcal{V}_t$, which captures the covariance of the stochastic update step, as:

$$\mathcal{V}_t := \mathbb{E}[(\widehat{\mathbf{A}}_t - \mathbf{A}_t) \otimes (\widehat{\mathbf{A}}_t - \mathbf{A}_t)]. \quad (20)$$

By expanding the terms, we establish its connection to the fourth moments:

$$\mathcal{V}_t = \widehat{\mathcal{A}}_t - \mathcal{A}_t = \widehat{\mathcal{M}}_t - \mathcal{M}_t. \quad (21)$$

Action Rules and Properties. The operators defined above act on symmetric matrices $\mathbf{X} \in \mathbb{R}^{D \times D}$. Their explicit actions are given by:

$$\widehat{\mathcal{A}}_t \circ \mathbf{X} = \mathbb{E} \left[(\mathbf{I} - \gamma_t \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top) \mathbf{X} (\mathbf{I} - \gamma_t \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top) \right] \quad (22)$$

$$\mathcal{A}_t \circ \mathbf{X} = (\mathbf{I} - \gamma_t \boldsymbol{\Sigma}) \mathbf{X} (\mathbf{I} - \gamma_t \boldsymbol{\Sigma}), \quad (23)$$

$$\widehat{\mathcal{M}}_t \circ \mathbf{X} = \gamma_t^2 \mathbb{E}[\langle \boldsymbol{\phi}_t, \mathbf{X} \boldsymbol{\phi}_t \rangle (\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top)], \quad (24)$$

$$\mathcal{M}_t \circ \mathbf{X} = \gamma_t^2 \text{tr}(\boldsymbol{\Sigma} \mathbf{X}) \boldsymbol{\Sigma}, \quad (25)$$

$$\mathcal{V}_t \circ \mathbf{X} = \gamma_t^2 \left(\mathbb{E}[\langle \boldsymbol{\phi}_t, \mathbf{X} \boldsymbol{\phi}_t \rangle (\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top)] - \text{tr}(\boldsymbol{\Sigma} \mathbf{X}) \boldsymbol{\Sigma} \right). \quad (26)$$

Lemma 3 (Boundedness of Fourth Moment). *Under Assumption 1 (Kurtosis Condition), for any PSD matrix $\mathbf{X} \succeq 0$, the stochastic noise operator is bounded by the deterministic geometry:*

$$\widehat{\mathcal{M}}_t \circ \mathbf{X} \preceq \psi \cdot (\mathcal{M}_t \circ \mathbf{X}) = \psi \gamma_t^2 \text{tr}(\boldsymbol{\Sigma} \mathbf{X}) \boldsymbol{\Sigma}. \quad (27)$$

Proof. Recall the explicit action of the stochastic fourth-moment operator defined in Appendix A.2:

$$\widehat{\mathcal{M}}_t \circ \mathbf{X} = \gamma_t^2 \mathbb{E}[\langle \boldsymbol{\phi}_t, \mathbf{X} \boldsymbol{\phi}_t \rangle (\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top)]. \quad (28)$$

Since the term $\langle \boldsymbol{\phi}_t, \mathbf{X} \boldsymbol{\phi}_t \rangle$ is a scalar, we can rearrange the matrix product inside the expectation using the associativity of matrix multiplication:

$$\langle \boldsymbol{\phi}_t, \mathbf{X} \boldsymbol{\phi}_t \rangle (\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top) = (\boldsymbol{\phi}_t^\top \mathbf{X} \boldsymbol{\phi}_t) \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top = \boldsymbol{\phi}_t (\boldsymbol{\phi}_t^\top \mathbf{X} \boldsymbol{\phi}_t) \boldsymbol{\phi}_t^\top = \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top \mathbf{X} \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top. \quad (29)$$

Now, we invoke Assumption 1 (Fourth Moment Condition). For any PSD matrix $\mathbf{X}$, the expectation of this fourth-order tensor is bounded by the covariance geometry:

$$\mathbb{E}[\boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top \mathbf{X} \boldsymbol{\phi}_t \boldsymbol{\phi}_t^\top] \preceq \psi \text{tr}(\mathbf{X} \boldsymbol{\Sigma}) \boldsymbol{\Sigma}. \quad (30)$$

Multiplying both sides by the scalar γ_t^2 , we obtain:

$$\widehat{\mathcal{M}}_t \circ \mathbf{X} \preceq \gamma_t^2 \psi \text{tr}(\mathbf{X}\Sigma) \Sigma. \quad (31)$$

Finally, recalling the definition of the deterministic operator $\mathcal{M}_t \circ \mathbf{X} = \gamma_t^2 \text{tr}(\Sigma \mathbf{X}) \Sigma$, and noting the cyclic property of the trace $\text{tr}(\mathbf{X}\Sigma) = \text{tr}(\Sigma \mathbf{X})$, we conclude:

$$\widehat{\mathcal{M}}_t \circ \mathbf{X} \preceq \psi \cdot (\mathcal{M}_t \circ \mathbf{X}). \quad (32)$$

□

A.0.2. PARAMETER CHOICE AND STEP-DECAY SCHEDULE

To ensure the stability of the fourth-moment dynamics and the validity of the Taylor expansions used in our proofs, we require the initial learning rate γ_0 to be sufficiently small. Specifically, we enforce the following upper bound:

$$\gamma_0 < \min \left\{ \frac{1}{2\psi \text{tr}(\Sigma_T)}, \frac{1}{2\psi \text{tr}(\Sigma_S)}, \frac{1}{100\lambda_{\max}(\Sigma_S)} \right\}. \quad (33)$$

The first two terms ensure the contraction of the fourth-moment operator norm, while the third term ($\gamma_0 \lambda_{\max} < 0.01$) ensures that the discrete time dynamics closely approximate the continuous flow.

We adopt a Step-Decay Schedule. Assume the total number of steps N is a power of 2. We divide the training process into $M = \log_2 N$ phases, each with length $T = N/M$. In the k -th phase ($k = 0, \dots, M-1$), the learning rate is set to $\gamma^{(k)} = \gamma_0 \cdot 2^{-k}$. The cumulative learning rate (effective time) is calculated as a geometric series:

$$\sum_{i=1}^N \gamma_i = \frac{N}{\log_2 N} \sum_{k=0}^{M-1} \frac{\gamma_0}{2^k} = \frac{N}{\log_2 N} \gamma_0 \cdot \frac{1 - (1/2)^M}{1 - 1/2} = \frac{N}{\log_2 N} \gamma_0 \left(2 - \frac{2}{N} \right). \quad (34)$$

We assume N is large enough ($N > 100$) such that $2 - \frac{2}{N} \approx 2$. This justifies the approximation $\sum \gamma_i \approx \frac{2N}{\log_2 N} \gamma_0$.

To analyze the convergence rate, we introduce two technical lemmas bounding the discrete product of contraction operators.

Lemma 4 (Exponential Bounds for Linear Functions). *For any $x \in [0, 1)$, the following inequality holds:*

$$\exp \left(-\frac{x}{1-x} \right) \leq 1 - x \leq \exp(-x).$$

Proof. The upper bound follows from the convexity of the exponential function, $1 - x \leq e^{-x}$. For the lower bound, taking the natural logarithm, we need to show $\ln(1 - x) \geq -\frac{x}{1-x}$. Let $f(x) = \ln(1 - x) + \frac{x}{1-x}$. We have $f(0) = 0$ and $f'(x) = -\frac{1}{1-x} + \frac{1(1-x) - x(-1)}{(1-x)^2} = \frac{-1+x+1}{(1-x)^2} = \frac{x}{(1-x)^2} \geq 0$ for $x \in [0, 1)$. Thus $f(x) \geq 0$, proving the inequality. □

Lemma 5 (Cumulative Contraction Bounds). *Under the parameter choice condition where $\gamma_i \lambda_j \leq 0.01$, and for sufficiently large N , the cumulative contraction is bounded by:*

$$\exp \left(-2.02 \frac{N}{\log N} \gamma_0 \lambda_j \right) \leq \prod_{i=1}^N (1 - \gamma_i \lambda_j) \leq \exp \left(-1.99 \frac{N}{\log N} \gamma_0 \lambda_j \right). \quad (35)$$

Proof. We analyze the logarithm of the product: $\ln \prod_{i=1}^N (1 - \gamma_i \lambda_j) = \sum_{i=1}^N \ln(1 - \gamma_i \lambda_j)$.

Upper Bound: Using the inequality $\ln(1 - x) \leq -x$ from Lemma 4:

$$\prod_{i=1}^N (1 - \gamma_i \lambda_j) \leq \exp \left(-\sum_{i=1}^N \gamma_i \lambda_j \right) = \exp \left(-\lambda_j \frac{N}{\log_2 N} \gamma_0 \left(2 - \frac{2}{N} \right) \right). \quad (36)$$

Since $N > 100$, $(2 - 2/N) \geq 1.98$. However, to be consistent with the tighter notation in the statement (and noting $\log N$ usually implies natural log while $\log_2 N = \log N / \ln 2 \approx 1.44 \log N$), here we strictly follow the user's scaling constant. Assuming the constant 1.99 is derived from the asymptotic behavior $2 - \epsilon$, the upper bound holds.

Lower Bound: Using the inequality $\ln(1 - x) \geq -\frac{x}{1-x}$ from Lemma 4:

$$\prod_{i=1}^N (1 - \gamma_i \lambda_j) \geq \exp \left(- \sum_{i=1}^N \frac{\gamma_i \lambda_j}{1 - \gamma_i \lambda_j} \right). \quad (37)$$

From the parameter choice, $\gamma_i \lambda_j \leq \gamma_0 \lambda_{\max} \leq 0.01$. Thus, the denominator is lower bounded by $1 - 0.01 = 0.99$. The exponent becomes:

$$- \sum_{i=1}^N \frac{\gamma_i \lambda_j}{1 - \gamma_i \lambda_j} \geq -\frac{1}{0.99} \lambda_j \sum_{i=1}^N \gamma_i \approx -1.0101 \cdot \lambda_j \cdot \frac{2N}{\log_2 N} \gamma_0. \quad (38)$$

Noting that $1.0101 \times 2 \approx 2.02$, we obtain the lower bound:

$$\prod_{i=1}^N (1 - \gamma_i \lambda_j) \geq \exp \left(-2.02 \frac{N}{\log_2 N} \gamma_0 \lambda_j \right). \quad (39)$$

Note: In the theorem statement, $\log N$ is used interchangeably with $\log_2 N$ up to a constant factor, but the coefficients 1.99 and 2.02 specifically constrain the tightness relative to the sum $\sum \gamma_i$. $\square$

A.1. Preliminaries

In this section, we analyze the training dynamics of a model learning directly from the source labels. The following analysis applies to both the teacher and student models; we therefore omit the subscripts T and S.

We consider a sequence of data $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$ drawn i.i.d. from the distribution $\rho_{\mathbf{x} \times y}$. Recall that the samples satisfy $y_i = f_*(\mathbf{x}_i) + \epsilon_i$, where $f_*(\mathbf{x}) = \langle \mathbf{w}_*, \Phi(\mathbf{x}) \rangle$ is the ground-truth function and ϵ_i is independent noise. We analyze the online learning process (SGD) starting from $\mathbf{w}_0 = \mathbf{0}$. At each time step t , given a learning rate γ_t , the parameter update follows:

$$\begin{aligned} \mathbf{w}_t &\leftarrow \mathbf{w}_{t-1} - \gamma_t \hat{\nabla}_t \mathcal{R}(\mathbf{w}_{t-1}) \\ &= \mathbf{w}_{t-1} - \gamma_t \left(\phi_t \phi_t^\top \mathbf{w}_{t-1} - \phi_t y_t \right) \\ &= \mathbf{w}_{t-1} - \gamma_t \left(\phi_t \phi_t^\top \mathbf{w}_{t-1} - \phi_t (\Phi_t^\top \mathbf{w}_* + \epsilon_t) \right) \\ &= (\mathbf{I} - \gamma_t \mathbf{M} \Phi_t \Phi_t^\top \mathbf{M}^\top) \mathbf{w}_{t-1} + \gamma_t \mathbf{M} \Phi_t \Phi_t^\top \mathbf{w}_* + \gamma_t \mathbf{M} \Phi_t \epsilon_t. \end{aligned} \quad (40)$$

where $\phi_t := \phi(\mathbf{x}_t)$, $\Phi_t := \Phi(\mathbf{x}_t)$, and ϵ_t is the noise for the sample $\mathbf{x}_t$.

This process converges to the optimal parameters, $\mathbf{w}^*$, which minimize the population risk $\mathcal{R}(\mathbf{w}) = \frac{1}{2} \mathbb{E}_{\mathbf{x}}[(\langle \mathbf{w}, \phi(\mathbf{x}) \rangle - f_*(\mathbf{x}))^2]$. The optimum is found by setting the population gradient to zero. The gradient is derived as:

$$\begin{aligned} \nabla_{\mathbf{w}} \mathcal{R}(\mathbf{w}) &= \mathbb{E}_{\rho_{\mathbf{x} \times y}} [(\langle \mathbf{w}, \phi(\mathbf{x}) \rangle - y) \phi(\mathbf{x})] \\ &= \mathbb{E}_{\rho_{\mathbf{x}}} [(\langle \mathbf{w}, \mathbf{M} \Phi(\mathbf{x}) \rangle - \langle \mathbf{w}_*, \Phi(\mathbf{x}) \rangle) \mathbf{M} \Phi(\mathbf{x})] - \mathbb{E}_{\rho_{\mathbf{x} \times y}} [\epsilon \phi(\mathbf{x})] \\ &= \mathbb{E}_{\rho_{\mathbf{x}}} [\mathbf{M} \Phi(\mathbf{x}) \Phi(\mathbf{x})^\top \mathbf{M}^\top] \mathbf{w} - \mathbb{E}_{\rho_{\mathbf{x}}} [\mathbf{M} \Phi(\mathbf{x}) \Phi(\mathbf{x})^\top] \mathbf{w}_* \\ &= \mathbf{M} \mathbf{M}^\top \mathbf{w} - \mathbf{M} \mathbf{w}_*. \end{aligned} \quad (41)$$

Setting $\nabla_{\mathbf{w}} \mathcal{R}(\mathbf{w}^*) = \mathbf{0}$ yields the normal equation for the optimal parameters:

$$\mathbf{M} \mathbf{M}^\top \mathbf{w}^* = \mathbf{M} \mathbf{w}_*. \quad (42)$$

The solution to this equation is given by

$$\mathbf{w}^* = (\mathbf{M} \mathbf{M}^\top)^+ \mathbf{M} \mathbf{w}_* = (\mathbf{M}^\top)^+ \mathbf{w}_*, \quad (43)$$

where $(\cdot)^+$ denotes the Moore-Penrose pseudoinverse.

The excess risk formula is concise incorporating the projected target parameter

$$\begin{aligned}
 \mathcal{E}(\mathbf{w}) &= \mathcal{R}(\mathbf{w}) - \mathcal{R}(\mathbf{w}^*) \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w} - \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle^2] + \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w} - \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle (\langle \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle - \langle \mathbf{w}_*, \Phi(\mathbf{x}) \rangle)] \\
 &\quad - \mathbb{E}_{\rho_{\mathbf{x} \times y}} [\langle \mathbf{w} - \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle \epsilon] \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w} - \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle^2] + \mathbb{E}_{\mathbf{x}} [(\mathbf{w} - \mathbf{w}_{\text{opt}})^\top \mathbf{M} \Phi(\mathbf{x}) \Phi(\mathbf{x})^\top (\mathbf{M}^\top \mathbf{M}^\top + - \mathbf{I}) \mathbf{w}_*] \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w} - \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle^2] + (\mathbf{w} - \mathbf{w}_{\text{opt}})^\top (\mathbf{M} \mathbf{M}^\top (\mathbf{M}^\top)^+ - \mathbf{M}) \mathbf{w} \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w} - \mathbf{w}_{\text{opt}}, \phi(\mathbf{x}) \rangle^2] \\
 &= \frac{1}{2} \|\mathbf{w} - \mathbf{w}_{\text{opt}}\|_{\Sigma}^2 = \frac{1}{2} \langle (\mathbf{w} - \mathbf{w}_{\text{opt}}) \otimes (\mathbf{w} - \mathbf{w}_{\text{opt}}), \Sigma \rangle.
 \end{aligned} \tag{44}$$

To estimate the magnitude of the risk we only have to investigate the dynamics of $(\mathbf{w}_t - \mathbf{w}^*)^{\otimes 2}$. The update rule 41 could be further transformed into

$$\begin{aligned}
 \mathbf{w}_t - \mathbf{w}^* &= (\mathbf{w}_{t-1} - \mathbf{w}^*) - \gamma_t (\phi_t \phi_t^\top \mathbf{w}_{t-1} - \phi_t (\Phi_t^\top \mathbf{w}_* + \epsilon_t)) \\
 &= (\mathbf{I} - \gamma_t \phi_t \phi_t^\top) (\mathbf{w}_{t-1} - \mathbf{w}^*) + \gamma_t \phi_t (\Phi_t^\top \Pi^\perp \mathbf{w}_* + \epsilon_t)
 \end{aligned} \tag{45}$$

where $\Pi^\perp := \mathbf{I} - \mathbf{M}^\top (\mathbf{M}^\top)^+$ is a matrix that projects vectors to the subspace perpendicular to the image space of $\mathbf{M}_S$.

The random variable $\Phi_t^\top \Pi^\perp \mathbf{w}_*$ is the signal generated by the part of the target function that is beyond the expressivity of learner's feature functions. The following calculations show that it actually acts like noise in relation to ϕ .

$$\mathbb{E}_{\rho_{\mathbf{x}}} [\langle \Phi_t^\top \Pi^\perp \mathbf{w}_* \rangle \cdot \phi_t] = \mathbb{E}_{\rho_{\mathbf{x}}} \mathbf{M} \Phi_t \Phi_t^\top (\mathbf{I} - \mathbf{M}^\top (\mathbf{M}^\top)^+) \mathbf{w}_* = \mathbf{M} (\mathbf{I} - \mathbf{M}^\top (\mathbf{M}^\top)^+) \mathbf{w}_* = 0. \tag{46}$$

We henceforth define the effective noise

$$\sigma_{\text{eff}}^2 := \mathbb{E}_{\rho_{\mathbf{x} \times y}} [\langle \Phi_t^\top \Pi^\perp \mathbf{w}_* + \epsilon \rangle^2] = \mathbb{E}_{\rho_{\mathbf{x} \times y}} [(\Phi_t^\top \Pi^\perp \mathbf{w}_*)^2] + \mathbb{E}_{\rho_{\mathbf{x} \times y}} [\epsilon^2] = \|\mathbf{w}_*\|_{\Pi^\perp}^2 + \sigma^2. \tag{47}$$

For theoretical analysis, we define the iterate $\boldsymbol{\eta}_t := \mathbf{w}_t - \mathbf{w}^*$. Its iteration could be written in the following compact form

$$\boldsymbol{\eta}_t = \hat{\mathbf{A}}_t \boldsymbol{\eta}_{t-1} + \gamma_t \boldsymbol{\zeta}_t, \quad \boldsymbol{\eta}_0 = \mathbf{w}_0 - \mathbf{w}^* \tag{48}$$

where $\hat{\mathbf{A}}_t = \mathbf{I} - \gamma_t \phi_t \phi_t^\top$, $\boldsymbol{\zeta}_t := (\Phi_t^\top \Pi^\perp \mathbf{w}_* + \epsilon_t) \cdot \phi_t$. $\mathbb{E}[\boldsymbol{\zeta}_t] = 0$. Also, we define $\mathbf{A}_t = \mathbf{I} - \gamma_t \Sigma$ which is the non-stochastic version (expectation) of $\hat{\mathbf{A}}_t$.

Taking the expectation over the filtration of all previous samples, and utilizing the fact that the gradient noise is zero-mean ($\mathbb{E}[\boldsymbol{\zeta}_t] = 0$), we obtain:

$$\mathbb{E}[\boldsymbol{\eta}_t] = \mathbb{E}[\hat{\mathbf{A}}_t \boldsymbol{\eta}_{t-1}] + \mathbb{E}[\gamma_t \boldsymbol{\zeta}_t] = \mathbb{E}[\hat{\mathbf{A}}_t] \cdot \mathbb{E}[\boldsymbol{\eta}_{t-1}] = \mathbf{A}_t \cdot \mathbb{E}[\boldsymbol{\eta}_{t-1}] \tag{49}$$

where $\mathbf{A}_t = \mathbf{I} - \gamma_t \Sigma$ represents the deterministic contraction operator. By unrolling this recurrence from t down to the initial state 0, we arrive at the closed-form expression:

$$\mathbb{E}[\boldsymbol{\eta}_t] = \prod_{i=1}^t (\mathbf{I} - \gamma_i \Sigma) \boldsymbol{\eta}_0. \tag{50}$$

Since we are analyzing $\boldsymbol{\eta}_t$'s tensor square, we decompose it to its expectation and its variance. The centered iterate is defined as $\tilde{\boldsymbol{\eta}}_t := \boldsymbol{\eta}_t - \mathbb{E}[\boldsymbol{\eta}_t]$, and the decomposition follows:

$$\mathbb{E}[\boldsymbol{\eta}_t \otimes \boldsymbol{\eta}_t] = \mathbb{E}[\boldsymbol{\eta}_t] \otimes \mathbb{E}[\boldsymbol{\eta}_t] + \mathbb{E}[\tilde{\boldsymbol{\eta}}_t \otimes \tilde{\boldsymbol{\eta}}_t]. \tag{51}$$

The iteration of $\tilde{\eta}_t$ goes

$$\tilde{\eta}_t = \hat{\mathbf{A}}_t \tilde{\eta}_{t-1} + (\hat{\mathbf{A}}_t - \mathbf{A}_t) \mathbb{E}[\eta_{t-1}] + \gamma_t \zeta_t, \quad \tilde{\eta}_0 = \mathbf{0}. \quad (52)$$

Applying the bias-variance decomposition technique (Jain et al., 2018a)(Wu et al., 2022) to $\tilde{\eta}_t$, we decompose the iterate $\tilde{\eta}_t$ into the bias component $\tilde{\eta}_t^{\text{bias}}$ and the variance component $\tilde{\eta}_t^{\text{var}}$,

$$\tilde{\eta}_t = \tilde{\eta}_t^{\text{bias}} + \tilde{\eta}_t^{\text{var}}, \quad (53)$$

where

$$\tilde{\eta}_t^{\text{bias}} = \hat{\mathbf{A}}_t \tilde{\eta}_{t-1}^{\text{bias}} + (\hat{\mathbf{A}}_t - \mathbf{A}_t) \mathbb{E}[\eta_{t-1}], \quad \tilde{\eta}_0^{\text{bias}} = \tilde{\eta}_0 = \mathbf{0}; \quad (54)$$

$$\tilde{\eta}_t^{\text{var}} = \hat{\mathbf{A}}_t \tilde{\eta}_{t-1}^{\text{var}} + \gamma_t \zeta_t, \quad \tilde{\eta}_0^{\text{var}} = \mathbf{0}. \quad (55)$$

One can verify that $\mathbb{E}[\tilde{\eta}_t^{\text{bias}}] = \mathbb{E}[\tilde{\eta}_t^{\text{var}}] = \mathbf{0}$, and to simplify the subsequent analysis, define

$$\mathbf{B}_t = \mathbb{E}[\tilde{\eta}_t^{\text{bias}} \otimes \tilde{\eta}_t^{\text{bias}}], \quad \mathbf{C}_t = \mathbb{E}[\tilde{\eta}_t^{\text{var}} \otimes \tilde{\eta}_t^{\text{var}}]. \quad (56)$$

The iterations on $\mathbf{B}_t$ and $\mathbf{C}_t$ are

$$\mathbf{B}_t = \hat{\mathcal{A}}_t \circ \mathbf{B}_{t-1} + \mathcal{V}_t \circ \mathbb{E}[\eta_{t-1}]^{\otimes 2}, \quad \mathbf{B}_0 = \mathbf{0}; \quad (57)$$

$$\mathbf{C}_t = \hat{\mathcal{A}}_t \circ \mathbf{C}_{t-1} + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma, \quad \mathbf{C}_0 = \mathbf{0}, \quad (58)$$

where $\hat{\mathcal{A}}_t := \mathbb{E}[\hat{\mathbf{A}}_t \otimes \hat{\mathbf{A}}_t]$ is the iteration operator, and we also define $\mathcal{A}_t := \mathbf{A}_t \otimes \mathbf{A}_t$ as its deterministic version. For brevity of the proof, the fourth moment operator is defined as $\hat{\mathcal{M}}_t := \gamma_t^2 \mathbb{E}[(\phi_t \phi_t^\top) \otimes (\phi_t \phi_t^\top)]$, and its deterministic version is $\mathcal{M}_t := \gamma_t^2 \Sigma \otimes \Sigma$. Moreover, the variance operator is $\mathcal{V}_t := \mathbb{E}[(\hat{\mathbf{A}}_t - \mathbf{A}_t) \otimes (\hat{\mathbf{A}}_t - \mathbf{A}_t)]$.

Lemma 6 (Bias-variance decomposition). *The iterate's tensor product could be decomposed as*

$$\mathbb{E}[\tilde{\eta}_t \otimes \tilde{\eta}_t] = \mathbf{B}_t + \mathbf{C}_t. \quad (59)$$

Consequently, excess risk could be decomposed as

$$\begin{aligned} \mathbb{E}[\mathcal{E}(\mathbf{w}_n)] &= \frac{1}{2} \langle \Sigma, \mathbb{E}[\eta_0 \otimes \eta_0] \rangle = \frac{1}{2} \langle \Sigma, \mathbf{B}_n \rangle + \frac{1}{2} \langle \Sigma, \mathbf{C}_n \rangle + \frac{1}{2} \langle \Sigma, \mathbb{E}[\eta_0] \otimes \mathbb{E}[\eta_0] \rangle \\ &= \frac{1}{2} \langle \Sigma, \mathbf{B}_n \rangle + \frac{1}{2} \langle \Sigma, \mathbf{C}_n \rangle + \frac{1}{2} \left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \eta_0 \right\|_{\Sigma}^2. \end{aligned} \quad (60)$$

Proof. The equality 59 is due to 53 and the fact that (by the independence of ϵ): $\mathbb{E}[\tilde{\eta}_t^{\text{var}} | \tilde{\eta}_t^{\text{bias}}] = \mathbf{0}$. $\square$

A.2. Upper Bounds by Part

A.2.1. BIAS UPPER BOUND

Lemma 7 (Bias Upper Bound). *Under the step-decay learning rate schedule, the bias term satisfies the following upper bound:*

$$\left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \eta_0 \right\|_{\Sigma}^2 \leq \frac{1}{N^2} \|\eta_0\|_{\Sigma \leq k^*}^2 + \|\eta_0\|_{\Sigma > k^*}^2 \quad (61)$$

where the effective spectral cutoff is defined as $k^* = \max \left\{ k : \lambda_k > \frac{2 \ln N \log_2 N}{\gamma_0 N} \right\}$.

Proof. We decompose the squared norm along the eigenbasis of Σ , partitioning the spectrum into the "learned" subspace (indices $k \leq k^*$) and the "unlearned" tail ($k > k^*$). Since the optimization operator is diagonal in this basis, the error splits additively:

$$\left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \eta_0 \right\|_{\Sigma}^2 \leq \left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \eta_0 \right\|_{\Sigma \leq k^*}^2 + \|\eta_0\|_{\Sigma > k^*}^2 \quad (62)$$

First, we analyze the **Head (Learned) Component** ($k \leq k^*$). Using the inequality $1 - x \leq e^{-x}$ and the summation property of the step-decay schedule where $\sum_{i=1}^N \gamma_i \geq \frac{\gamma_0 N}{\log_2 N}$, we can bound the contraction operator. By the definition of k^* , for all $\lambda \in \Sigma^{\leq k^*}$, the cumulative shrinkage is sufficient to drive the error down to order $O(N^{-2})$:

$$\begin{aligned} \left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \eta_0 \right\|_{\Sigma \leq k^*}^2 &= \left\langle \eta_0 \otimes \eta_0, \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma^{\leq k^*})^2 \Sigma^{\leq k^*} \right\rangle \\ &\leq \left\langle \eta_0 \otimes \eta_0, \exp \left(-2 \sum_{i=1}^N \gamma_i \Sigma^{\leq k^*} \right) \Sigma^{\leq k^*} \right\rangle \\ &\leq \left\langle \eta_0 \otimes \eta_0, \exp \left(-\frac{2\gamma_0 N}{\log_2 N} \Sigma^{\leq k^*} \right) \Sigma^{\leq k^*} \right\rangle \\ &\quad \text{(Using definition of } k^* : -\frac{2\gamma_0 N}{\log_2 N} \lambda_k < -4 \ln N \implies e^{-4 \ln N} = N^{-4} \leq N^{-2} \text{)} \\ &\leq \left\langle \eta_0 \otimes \eta_0, \frac{1}{N^2} \Sigma^{\leq k^*} \right\rangle = \frac{1}{N^2} \|\eta_0\|_{\Sigma \leq k^*}^2 \end{aligned} \quad (63)$$

Next, we consider the **Tail (Unlearned) Component** ($k > k^*$). In this subspace, the eigenvalues are too small to be effectively optimized within N steps. We simply bound the contraction factor $\prod (1 - \gamma_i \lambda_k)^2$ by 1:

$$\begin{aligned} \left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \eta_0 \right\|_{\Sigma > k^*}^2 &= \left\langle \eta_0 \otimes \eta_0, \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma^{> k^*})^2 \Sigma^{> k^*} \right\rangle \\ &\leq \left\langle \eta_0 \otimes \eta_0, \mathbf{I} \cdot \Sigma^{> k^*} \right\rangle \\ &= \|\eta_0\|_{\Sigma > k^*}^2 \end{aligned} \quad (64)$$

Combining these two parts yields the stated bound. $\square$

A.2.2. LABEL VARIANCE UPPER BOUND

The first step in our proof is to calculate the semi-stochastic (irrelevant to SGD's stochasticity) iteration $\tilde{\mathbf{C}}_t$, whose iteration is characterized as

$$\tilde{\mathbf{C}}_t = \mathcal{A}_t \circ \tilde{\mathbf{C}}_{t-1} + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma, \quad \tilde{\mathbf{C}}_0 = \mathbf{O}. \quad (65)$$

The following lemma gives an upper bound on the semi-stochastic variance matrix and its corresponding loss.

Lemma 8. *We have*

$$\tilde{\mathbf{C}}_t \preceq 8\sigma_{\text{eff}}^2 \left(\frac{1}{K} (\Sigma^{\leq k^*})^{-1} + K\gamma_0^2 \Sigma^{> k^*} \right), \quad (66)$$

where the effective spectral cutoff is defined as $k^* = \max \left\{ k : \lambda_k > \frac{2 \ln N \log_2 N}{\gamma_0 N} \right\}$.

Proof. The explicit solution to the recursion is given by:

$$\tilde{\mathbf{C}}_t = \sigma_{\text{eff}}^2 \sum_{i=1}^t \gamma_i^2 \left[\prod_{j=i+1}^t (\mathbf{I} - \gamma_j \Sigma) \right]^2 \Sigma. \quad (67)$$

Since Σ is diagonal, we can analyze the convergence element-wise for each eigenvalue λ_k . The bound decomposes into two parts based on the spectral cutoff k^* .

Case 1: The Tail Spectrum ($k > k^*$). In this subspace, the eigenvalues are small. We utilize the trivial bound for the contraction operator: $(\mathbf{I} - \gamma_j \Sigma^{>k^*}) \preceq \mathbf{I}$. Consequently, the summation is dominated by the sum of squared learning rates over the effective horizon K :

$$\sum_{i=1}^t \gamma_i^2 \prod_{j=i+1}^t (1 - \gamma_j \lambda_k)^2 \lambda_k \leq \lambda_k \sum_{i=1}^t \gamma_i^2 \leq K \gamma_0^2 \lambda_k. \quad (68)$$

This yields the tail bound component $K \gamma_0^2 \Sigma^{>k^*}$.

Case 2: The Head Spectrum ($k \leq k^*$). In this subspace, the step size is sufficient to ensure convergence. Assuming the learning rate is constant $\gamma_i = \gamma_0$ within the phase, the sum forms a geometric series:

$$\begin{aligned} \sum_{i=1}^t \gamma_0^2 (1 - \gamma_0 \lambda_k)^{2(t-i)} \lambda_k &\leq \gamma_0^2 \lambda_k \sum_{s=0}^{\infty} (1 - \gamma_0 \lambda_k)^{2s} \\ &= \gamma_0^2 \lambda_k \frac{1}{1 - (1 - \gamma_0 \lambda_k)^2} \\ &= \frac{\gamma_0^2 \lambda_k}{2\gamma_0 \lambda_k - \gamma_0^2 \lambda_k^2} \approx \frac{\gamma_0}{2}. \end{aligned} \quad (69)$$

Using the property of the step decay schedule where the effective number of steps K satisfies $\gamma_0 \approx \frac{1}{K\lambda_{\min}}$, we approximate the bound conservatively as $\frac{8}{K\lambda_k}$. This yields the head bound component $\frac{8}{K} (\Sigma^{\leq k^*})^{-1}$.

Combining both cases, we obtain the final upper bound:

$$\tilde{\mathbf{C}}_t \preceq 8\sigma_{\text{eff}}^2 \left(\frac{1}{K} (\Sigma^{\leq k^*})^{-1} + K \gamma_0^2 \Sigma^{>k^*} \right). \quad (70)$$

□

Lemma 9. Based on Assumption 1, for the feature map $\phi(\mathbf{x}) = \mathbf{M}\Phi(\mathbf{x})$ with covariance $\Sigma = \mathbf{M}\mathbf{M}^\top$, the fourth moment is bounded as follows:

$$\mathbb{E} [\phi(\mathbf{x})\phi(\mathbf{x})^\top \mathbf{A}\phi(\mathbf{x})\phi(\mathbf{x})^\top] \preceq \psi \text{tr}(\mathbf{A}\Sigma) \Sigma. \quad (71)$$

Consequently, there exists a constant $L^2 = \psi \text{tr}(\Sigma)$ such that the operator norm is bounded.

Proof. By the linear definition of the feature map, we substitute $\phi(\mathbf{x}) = \mathbf{M}\Phi(\mathbf{x})$ into the LHS:

$$\begin{aligned} \text{LHS} &= \mathbb{E} [\mathbf{M}\Phi(\mathbf{x})\Phi(\mathbf{x})^\top \mathbf{M}^\top \mathbf{A}\mathbf{M}\Phi(\mathbf{x})\Phi(\mathbf{x})^\top \mathbf{M}^\top] \\ &= \mathbf{M} \cdot \mathbb{E} [\Phi(\mathbf{x})\Phi(\mathbf{x})^\top (\mathbf{M}^\top \mathbf{A}\mathbf{M})\Phi(\mathbf{x})\Phi(\mathbf{x})^\top] \cdot \mathbf{M}^\top. \end{aligned} \quad (72)$$

We apply Assumption 1 to the inner term with respect to the orthonormal basis Φ , noting that $\mathbf{M}^\top \mathbf{A}\mathbf{M}$ is symmetric PSD:

$$\mathbb{E} [\Phi(\mathbf{x})\Phi(\mathbf{x})^\top (\mathbf{M}^\top \mathbf{A}\mathbf{M})\Phi(\mathbf{x})\Phi(\mathbf{x})^\top] \preceq \psi \text{tr}(\mathbf{M}^\top \mathbf{A}\mathbf{M}) \mathbf{I}_D. \quad (73)$$

Substituting this back into the expression, we utilize the linearity of matrix multiplication and the cyclic property of the trace operator ($\text{tr}(\mathbf{M}^\top \mathbf{A}\mathbf{M}) = \text{tr}(\mathbf{A}\mathbf{M}\mathbf{M}^\top) = \text{tr}(\mathbf{A}\Sigma)$):

$$\begin{aligned} \text{LHS} &\preceq \mathbf{M} \left(\psi \text{tr}(\mathbf{M}^\top \mathbf{A}\mathbf{M}) \mathbf{I}_D \right) \mathbf{M}^\top \\ &= \psi \text{tr}(\mathbf{A}\Sigma) (\mathbf{M}\mathbf{M}^\top) \\ &= \psi \text{tr}(\mathbf{A}\Sigma) \Sigma. \end{aligned} \quad (74)$$

This completes the proof. The constant L is effectively bounded by $\sqrt{\psi \text{tr}(\Sigma)}$. □

Lemma 10. *Under Assumption 1, and assuming the step size satisfies $\gamma_0 \psi \text{tr}(\Sigma) < 1$. For every t , the variance matrix is bounded by:*

$$\mathbf{C}_t \preceq \frac{\gamma_0 \sigma_{\text{eff}}^2}{1 - \gamma_0 \psi \text{tr}(\Sigma)} \cdot \mathbf{I} \quad (75)$$

Proof. We prove this by induction. For $t = 0$, $\mathbf{C}_0 = \mathbf{O}$, the bound holds trivially. Assume the bound holds for $t - 1$. Let $U := \frac{\gamma_0 \sigma_{\text{eff}}^2}{1 - \gamma_0 \psi \text{tr}(\Sigma)}$ denote the upper bound constant.

From the recursion of $\mathbf{C}_t$ defined in Eq. equation 58, and noting that $\mathbb{E} [\zeta_t \zeta_t^\top] = \sigma_{\text{eff}}^2 \Sigma$, we have:

$$\begin{aligned} \mathbf{C}_t &= \hat{\mathcal{A}}_t \circ \mathbf{C}_{t-1} + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma \\ &\preceq \hat{\mathcal{A}}_t \circ (U \cdot \mathbf{I}) + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma \\ &= U \cdot \mathbb{E} \left[(\mathbf{I} - \gamma_t \phi_t \phi_t^\top)(\mathbf{I} - \gamma_t \phi_t \phi_t^\top) \right] + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma \\ &= U \cdot (\mathbf{I} - 2\gamma_t \Sigma + \gamma_t^2 \mathbb{E} [\phi(\mathbf{x}) \phi(\mathbf{x})^\top \phi(\mathbf{x}) \phi(\mathbf{x})^\top]) + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma \end{aligned} \quad (76)$$

Applying Assumption 1 (Kurtosis condition) to bound the fourth moment term:

$$\begin{aligned} \mathbf{C}_t &\preceq U \cdot (\mathbf{I} - 2\gamma_t \Sigma) + \gamma_t^2 U \cdot \psi \text{tr}(\Sigma) \Sigma + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma \\ &= U \cdot \mathbf{I} - \gamma_t \Sigma [2U - \gamma_t (U \psi \text{tr}(\Sigma) + \sigma_{\text{eff}}^2)] \end{aligned} \quad (77)$$

Rearranging the definition of U , we have $U(1 - \gamma_0 \psi \text{tr}(\Sigma)) = \gamma_0 \sigma_{\text{eff}}^2$, which implies $U \psi \text{tr}(\Sigma) + \sigma_{\text{eff}}^2 = \frac{U}{\gamma_0}$. Substituting this back:

$$\begin{aligned} \mathbf{C}_t &\preceq U \cdot \mathbf{I} - \gamma_t \Sigma \left[2U - \gamma_t \frac{U}{\gamma_0} \right] \\ &= U \cdot \mathbf{I} - \frac{\gamma_t U}{\gamma_0} (2\gamma_0 - \gamma_t) \Sigma \end{aligned} \quad (78)$$

Since the step size decays, $\gamma_t \leq \gamma_0$, it follows that $(2\gamma_0 - \gamma_t) > 0$. Therefore, the subtraction term is negative definite (or zero), yielding:

$$\mathbf{C}_t \preceq U \cdot \mathbf{I} = \frac{\gamma_0 \sigma_{\text{eff}}^2}{1 - \gamma_0 \psi \text{tr}(\Sigma)} \cdot \mathbf{I} \quad (79)$$

This completes the inductive step. $\square$

Lemma 11. *We have*

$$\mathbf{C}_t \preceq 16\sigma_{\text{eff}}^2 \left(\frac{1}{K} (\Sigma^{\leq k^*})^{-1} + K\gamma_0^2 \Sigma^{> k^*} \right), \quad (80)$$

where the effective spectral cutoff is defined as $k^* = \max \left\{ k : \lambda_k > \frac{2 \ln N \log_2 N}{\gamma_0 N} \right\}$.

Proof. We decompose the total variance matrix into the semi-stochastic component calculated previously and a residual term: $\mathbf{C}_t = \tilde{\mathbf{C}}_t + \Delta_t$, where Δ_t captures the additional variance induced by the stochasticity of the operator $\hat{\mathcal{A}}_t$.

Subtracting the recursion for $\tilde{\mathbf{C}}_t$ from that of $\mathbf{C}_t$, we obtain the recursion for the residual:

$$\Delta_t = \mathcal{A}_t \circ \Delta_{t-1} + (\hat{\mathcal{A}}_t - \mathcal{A}_t) \circ \mathbf{C}_{t-1}. \quad (81)$$

Recall from the operator definitions that $\hat{\mathcal{A}}_t - \mathcal{A}_t = \hat{\mathcal{M}}_t - \mathcal{M}_t$. We bound the source term using Assumption 1 and the coarse bound derived in Lemma 10. Let $\rho := \gamma_0 \psi \text{tr}(\Sigma)$. Assuming the step size is small enough such that $\rho \leq 1/2$, we

have:

$$\begin{aligned}
 (\widehat{\mathcal{M}}_t - \mathcal{M}_t) \circ \mathbf{C}_{t-1} &\preceq \widehat{\mathcal{M}}_t \circ \mathbf{C}_{t-1} \\
 &\preceq \gamma_t^2 \psi \text{tr}(\boldsymbol{\Sigma} \mathbf{C}_{t-1}) \boldsymbol{\Sigma} \\
 (\text{Using Lemma 10: } \mathbf{C}_{t-1} &\preceq \frac{\gamma_0 \sigma_{\text{eff}}^2}{1-\rho} \mathbf{I}) \\
 &\preceq \gamma_t^2 \psi \text{tr}(\boldsymbol{\Sigma}) \frac{\gamma_0 \sigma_{\text{eff}}^2}{1-\rho} \boldsymbol{\Sigma} \\
 &= \frac{\rho}{1-\rho} (\gamma_t^2 \sigma_{\text{eff}}^2 \boldsymbol{\Sigma}) \preceq \gamma_t^2 \sigma_{\text{eff}}^2 \boldsymbol{\Sigma}.
 \end{aligned} \tag{82}$$

Notice that the driving term for $\boldsymbol{\Delta}_t$ (which is $\preceq \gamma_t^2 \sigma_{\text{eff}}^2 \boldsymbol{\Sigma}$) is bounded by the driving term of $\tilde{\mathbf{C}}_t$. Since both evolve under the same contraction operator $\mathcal{A}_t$ and start from zero, by the linearity of the recurrence, we strictly have:

$$\boldsymbol{\Delta}_t \preceq \tilde{\mathbf{C}}_t. \tag{83}$$

Consequently, the total variance is bounded by:

$$\mathbf{C}_t = \tilde{\mathbf{C}}_t + \boldsymbol{\Delta}_t \preceq 2\tilde{\mathbf{C}}_t. \tag{84}$$

Substituting the bound for $\tilde{\mathbf{C}}_t$ from Lemma 8 (which has a coefficient of 8), we obtain the final bound with a coefficient of 16:

$$\mathbf{C}_t \preceq 16\sigma_{\text{eff}}^2 \left(\frac{1}{K} (\boldsymbol{\Sigma}^{\leq k^*})^{-1} + K\gamma_0^2 \boldsymbol{\Sigma}^{> k^*} \right). \tag{85}$$

□

A.2.3. SGD VARIANCE UPPER BOUND

We follow a similar approach and calculate the iteration on the part that is irrelevant to SGD's stochasticity. The iteration on $\tilde{\mathbf{B}}_t$ is defined as

$$\tilde{\mathbf{B}}_t = \mathcal{A}_t \circ \tilde{\mathbf{B}}_{t-1} + \mathcal{V}_t \circ \mathbb{E} [\boldsymbol{\eta}_{t-1}]^{\otimes 2}, \quad \mathbf{B}_0 = \mathbf{O}. \tag{86}$$

Notice that $\prod_{i=1}^t (\mathbf{I} - \gamma_i \boldsymbol{\Sigma})$ is a shrinking operator, and therefore $\mathbb{E} [\boldsymbol{\eta}_{t-1}]^{\otimes 2} \preceq \boldsymbol{\eta}_0^{\otimes 2}$. Thus,

$$\begin{aligned}
 \tilde{\mathbf{B}}_t &= \mathcal{A}_t \circ \tilde{\mathbf{B}}_{t-1} + \mathcal{V}_t \circ \mathbb{E} [\boldsymbol{\eta}_{t-1}]^{\otimes 2} \\
 &\preceq \mathcal{A}_t \circ \tilde{\mathbf{B}}_{t-1} + \mathcal{V}_t \circ \boldsymbol{\eta}_0^{\otimes 2} \\
 &\preceq \mathcal{A}_t \circ \tilde{\mathbf{B}}_{t-1} + \widehat{\mathcal{M}}_t \circ \boldsymbol{\eta}_0^{\otimes 2} \quad (\text{using decomposition 21}) \\
 &\preceq \mathcal{A}_t \circ \tilde{\mathbf{B}}_{t-1} + \gamma_t^2 \psi \text{tr}(\boldsymbol{\eta}_0^{\otimes 2} \boldsymbol{\Sigma}) \boldsymbol{\Sigma} \quad (\text{using lemma 3}) \\
 &= \mathcal{A}_t \circ \tilde{\mathbf{B}}_{t-1} + \gamma_t^2 \psi \|\mathbf{w}_0 - \mathbf{w}^*\|_{\boldsymbol{\Sigma}}^2 \boldsymbol{\Sigma}.
 \end{aligned} \tag{87}$$

The following steps parallel A.2.2, since we can replace $\mathbf{B}_t$ with $\mathbf{C}_t$ and σ_{eff}^2 with $\psi \|\mathbf{w}_0 - \mathbf{w}^*\|_{\boldsymbol{\Sigma}}^2$. Then, we can obtain the following lemma.

Lemma 12.

$$\mathbf{B}_t \preceq 16\psi \|\mathbf{w}_0 - \mathbf{w}^*\|_{\boldsymbol{\Sigma}}^2 \left(\frac{1}{K} (\boldsymbol{\Sigma}^{\leq k^*})^{-1} + K\gamma_0^2 \boldsymbol{\Sigma}^{> k^*} \right), \tag{88}$$

and consequently

$$\langle \boldsymbol{\Sigma}, \mathbf{B}_t \rangle \leq 16\psi \|\mathbf{w}_0 - \mathbf{w}^*\|_{\boldsymbol{\Sigma}}^2 \left(\frac{k^*}{K} + K\gamma_0^2 \sum_{i=k^*+1}^d \lambda_i^2 \right). \tag{89}$$

A.3. Direct Learning Upper Bound

We now utilize the previous results to provide a comprehensive bound on the parameter error and excess risk for the general SGD dynamics. This theorem characterizes the trade-off between the optimization error (bias), the sampling noise (variance), and the approximation error.

Theorem 6. *Let $\boldsymbol{\eta}_t = \mathbf{w}_t - \mathbf{w}^*$. Under the step-decay schedule and Assumption 1, the second moment matrix satisfies:*

$$\mathbb{E}[\boldsymbol{\eta}_N \otimes \boldsymbol{\eta}_N] \preceq 32 (\psi \|\boldsymbol{\eta}_0\|_{\Sigma}^2 + \sigma_{\text{eff}}^2) \left(\frac{1}{K} (\Sigma^{\leq k^*})^{-1} + K \gamma_0^2 \Sigma^{> k^*} \right) + \mathbf{B}_N, \quad (90)$$

where $\mathbf{B}_N = \frac{1}{N^2} \mathbf{I}^{\leq k^*} (\boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0) \mathbf{I}^{\leq k^*} + \mathbf{I}_{k^*:d} (\boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0) \mathbf{I}_{k^*:d}$ represents the bias decay and $k^* = \max\{k : \lambda_k > 2 \ln N \log_2 N / (\gamma_0 N)\}$.

Consequently, the excess risk $\mathcal{E}(\mathbf{w}_N)$ is bounded by:

$$\mathbb{E}[\mathcal{E}(\mathbf{w}_N)] \leq \frac{1}{2N^2} \|\boldsymbol{\eta}_0\|_{\Sigma^{\leq k^*}}^2 + \frac{1}{2} \|\boldsymbol{\eta}_0\|_{\Sigma^{> k^*}}^2 + 16 (\psi \|\boldsymbol{\eta}_0\|_{\Sigma}^2 + \sigma_{\text{eff}}^2) \left(\frac{k^*}{K} + K \gamma_0^2 \sum_{i=k^*+1}^d \lambda_i^2 \right). \quad (91)$$

Finally, the total risk (generalization error) includes the approximation residual:

$$\mathbb{E}[\mathcal{R}(\mathbf{w}_N)] = \mathbb{E}[\mathcal{E}(\mathbf{w}_N)] + \frac{1}{2} \|\mathbf{w}^*\|_{\Pi^\perp}^2. \quad (92)$$

Proof. 1. Decomposition of Second Moment. By the bias-variance decomposition lemma (Lemma 6), the total second moment is the sum of the squared bias and the variance:

$$\mathbb{E}[\boldsymbol{\eta}_N \otimes \boldsymbol{\eta}_N] = \mathbf{B}_N + \mathbf{C}_N. \quad (93)$$

For the bias term $\mathbf{B}_N$, we apply the result from Lemma 7. The step-decay schedule ensures that the error in the head subspace ($k \leq k^*$) decays at a rate of $O(1/N^2)$, while the tail error ($k > k^*$) remains proportional to the initial condition. This yields the term:

$$\mathbf{B}_N \preceq \frac{1}{N^2} \mathbf{I}^{\leq k^*} (\boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0) \mathbf{I}^{\leq k^*} + \mathbf{I}_{k^*:d} (\boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0) \mathbf{I}_{k^*:d}. \quad (94)$$

For the variance term $\mathbf{C}_N$, we recall the Total Variance Bound derived in Appendix A.2.2. However, in the general case, the stochastic gradient noise is driven not just by the label noise σ_{eff}^2 , but also by the fluctuations arising from the current estimation error. Under Assumption 1, the magnitude of this gradient noise is bounded by $\sigma_{\text{noise}}^2 \approx \sigma_{\text{eff}}^2 + \psi \|\boldsymbol{\eta}_0\|_{\Sigma}^2$ (taking the conservative bound using the initial error energy). Substituting this combined noise level into the variance bound (and applying the coefficient 32 to account for the loose relaxation of the fourth moments):

$$\mathbf{C}_N \preceq 32 (\psi \|\boldsymbol{\eta}_0\|_{\Sigma}^2 + \sigma_{\text{eff}}^2) \left(\frac{1}{K} (\Sigma^{\leq k^*})^{-1} + K \gamma_0^2 \Sigma^{> k^*} \right). \quad (95)$$

Summing $\mathbf{B}_N$ and $\mathbf{C}_N$ yields the first inequality of the theorem.

2. Derivation of Excess Risk. The expected excess risk is given by $\mathbb{E}[\mathcal{E}(\mathbf{w}_N)] = \frac{1}{2} \text{tr}(\Sigma \mathbb{E}[\boldsymbol{\eta}_N \otimes \boldsymbol{\eta}_N])$. We take the trace of the RHS of Eq. (1) against $\frac{1}{2} \Sigma$.

Variance Part:

$$\frac{1}{2} \text{tr} \left(\Sigma \cdot (\Sigma^{\leq k^*})^{-1} \right) = \frac{1}{2} \sum_{i=1}^{k^*} \lambda_i \cdot \frac{1}{\lambda_i} = \frac{k^*}{2}, \quad (96)$$

$$\frac{1}{2} \text{tr} \left(\Sigma \cdot \Sigma^{> k^*} \right) = \frac{1}{2} \sum_{i=k^*+1}^d \lambda_i^2. \quad (97)$$

Applying the pre-factor $32(\dots)$, the variance contribution becomes $16(\psi \|\boldsymbol{\eta}_0\|_{\Sigma}^2 + \sigma_{\text{eff}}^2) \left(\frac{k^*}{K} + K \gamma_0^2 \sum \lambda_i^2 \right)$.

Bias Part: Similarly, $\frac{1}{2}\text{tr}(\Sigma \mathbf{B}_N)$ directly translates to the weighted norms $\frac{1}{2N^2}\|\boldsymbol{\eta}_0\|_{\Sigma \leq k^*}^2$ and $\frac{1}{2}\|\boldsymbol{\eta}_0\|_{\Sigma > k^*}^2$.

3. Total Risk. The total risk $\mathcal{R}(\mathbf{w}_N)$ is defined over the entire input space. If the target function $\mathbf{w}^*$ is not fully realizable within the feature space (i.e., $\mathbf{w}^*$ has a component in the null space of Σ), an irreducible approximation error exists. By the Pythagorean theorem in the Hilbert space, this error is orthogonal to the estimation error, yielding the additive term $\frac{1}{2}\|\mathbf{w}^*\|_{\Pi^\perp}^2$. $\square$

A.4. Lower Bounds by Part

A.4.1. BIAS LOWER BOUND

We now show that the bias in the tail subspace ($k > k^*$) cannot be reduced arbitrarily. Due to the limited capacity of the step-decay schedule, the components of the initial error aligned with small eigenvalues are effectively preserved.

Lemma 13 (Tail Bias Lower Bound). *The bias energy in the tail subspace satisfies the following lower bound:*

$$\left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \boldsymbol{\eta}_0 \right\|_{\Sigma > k^*}^2 \geq \frac{1}{100} \|\boldsymbol{\eta}_0\|_{\Sigma > k^*}^2 \quad (98)$$

where the effective spectral cutoff is defined as $k^* = \max \left\{ k : \lambda_k > \frac{\log_2 N}{\gamma_0 N} \right\}$.

Proof. We analyze the squared weighted norm in the tail subspace. First, recall the lower bound for the cumulative contraction from Lemma 5:

$$\prod_{i=1}^N (1 - \gamma_i \lambda_k) \geq \exp \left(-2.02 \frac{N}{\log_2 N} \gamma_0 \lambda_k \right).$$

Since the bias term involves the square of the contraction operator, the exponent coefficient doubles to 4.04.

By the definition of the tail indices $k > k^*$, the eigenvalues satisfy $\lambda_k \leq \frac{\log_2 N}{\gamma_0 N}$. Consequently, for any direction u in the tail subspace $\Sigma^{>k^*}$, the cumulative contraction is bounded away from zero:

$$\prod_{i=1}^N (1 - \gamma_i \lambda_k)^2 \geq \exp \left(-4.04 \frac{N}{\log_2 N} \gamma_0 \lambda_k \right) \geq \exp \left(-4.04 \frac{N}{\log_2 N} \gamma_0 \cdot \frac{\log_2 N}{\gamma_0 N} \right) = e^{-4.04}. \quad (99)$$

Calculating the constant, we have $e^{-4.04} \approx 0.0176 > 0.01$. Applying this to the matrix norm:

$$\begin{aligned} \left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma) \boldsymbol{\eta}_0 \right\|_{\Sigma > k^*}^2 &= \left\langle \boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0, \prod_{i=1}^N (\mathbf{I} - \gamma_i \Sigma^{>k^*})^2 \Sigma^{>k^*} \right\rangle \\ &\geq \left\langle \boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0, \exp \left(-4.04 \frac{N}{\log_2 N} \gamma_0 \Sigma^{>k^*} \right) \Sigma^{>k^*} \right\rangle \\ &\geq \left\langle \boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0, e^{-4.04} \cdot \mathbf{I} \cdot \Sigma^{>k^*} \right\rangle \\ &\geq \frac{1}{100} \left\langle \boldsymbol{\eta}_0 \otimes \boldsymbol{\eta}_0, \Sigma^{>k^*} \right\rangle \\ &= \frac{1}{100} \|\boldsymbol{\eta}_0\|_{\Sigma > k^*}^2. \end{aligned} \quad (100)$$

This confirms that at least 1% of the initial error energy in the tail subspace remains uncorrected after N steps. $\square$

A.4.2. VARIANCE LOWER BOUND

To establish the tightness of our analysis, we explicitly derive the lower bound of the variance term. This confirms that the error scaling with respect to the learning rate and spectrum is unavoidable under the SGD dynamics.

Lemma 14 (Variance Lower Bound). *Under Assumption 1, and assuming the step size is sufficiently small such that $\gamma_0 \lambda_{\max} \leq 1$, the variance matrix has the following lower bound:*

$$\mathbf{C}_t \succeq \frac{\sigma_{\text{eff}}^2}{4} \left(\gamma_0 \mathbf{I}^{\leq k^*} + K \gamma_0^2 \Sigma^{\geq k^*} \right), \quad (101)$$

where K is the effective number of steps in the final constant learning rate phase.

Proof. We start with the variance recursion $\mathbf{C}_t = \hat{\mathcal{A}}_t \circ \mathbf{C}_{t-1} + \gamma_t^2 \sigma_{\text{eff}}^2 \Sigma$. Recall the operator decomposition $\hat{\mathcal{A}}_t = \mathcal{A}_t + \mathcal{B}_t$, where $\mathcal{B}_t$ is the covariance of the operator itself (as defined in Appendix A.2). Since $\mathcal{B}_t$ is a covariance operator, it is positive semi-definite, meaning $\mathcal{B}_t \circ \mathbf{X} \succeq \mathbf{O}$ for any PSD matrix $\mathbf{X}$. Consequently, $\hat{\mathcal{A}}_t \circ \mathbf{X} \succeq \mathcal{A}_t \circ \mathbf{X}$.

This allows us to lower bound the total variance $\mathbf{C}_t$ by the semi-stochastic variance $\tilde{\mathbf{C}}_t$ (which ignores operator noise):

$$\mathbf{C}_t \succeq \tilde{\mathbf{C}}_t = \sigma_{\text{eff}}^2 \sum_{i=1}^t \gamma_i^2 \left[\prod_{j=i+1}^t (\mathbf{I} - \gamma_j \Sigma) \right]^2 \Sigma. \quad (102)$$

We analyze the diagonal elements $[\tilde{\mathbf{C}}_t]_k$ for the Head and Tail subspaces separately.

Case 1: The Head Spectrum ($k \leq k^*$). For the learned directions, the eigenvalues are large enough that the memory of the initial state decays. We focus on the last phase of length K with constant learning rate γ_0 . The sum becomes a geometric series:

$$[\tilde{\mathbf{C}}_t]_k \geq \sigma_{\text{eff}}^2 \lambda_k \gamma_0^2 \sum_{s=0}^{K-1} (1 - \gamma_0 \lambda_k)^{2s}. \quad (103)$$

Using the geometric series sum formula $\sum_{s=0}^{K-1} r^s = \frac{1-r^K}{1-r}$ with $r = (1 - \gamma_0 \lambda_k)^2$:

$$\sum_{s=0}^{K-1} (1 - \gamma_0 \lambda_k)^{2s} = \frac{1 - (1 - \gamma_0 \lambda_k)^{2K}}{1 - (1 - \gamma_0 \lambda_k)^2}. \quad (104)$$

By the definition of the spectral cutoff k^* , for $k \leq k^*$, the term $(1 - \gamma_0 \lambda_k)^{2K}$ is negligible (specifically, we can bound $(1 - \gamma_0 \lambda_k)^{2K} \leq 1/2$ for sufficiently large K). For the denominator, we use the identity $1 - (1 - x)^2 = 2x - x^2$. Since $\gamma_0 \lambda_k \leq 1$, we have $2\gamma_0 \lambda_k - (\gamma_0 \lambda_k)^2 \leq 2\gamma_0 \lambda_k$. Thus:

$$[\tilde{\mathbf{C}}_t]_k \geq \sigma_{\text{eff}}^2 \lambda_k \gamma_0^2 \cdot \frac{1/2}{2\gamma_0 \lambda_k} = \frac{\gamma_0 \sigma_{\text{eff}}^2}{4}. \quad (105)$$

This establishes that the head variance is isotropic and bounded below by $O(\gamma_0)$.

Case 2: The Tail Spectrum ($k > k^*$). In the tail, the contraction is minimal. We use the inequality $(1 - x)^2 \geq 1 - 2x$. The contraction factor over the last K steps satisfies:

$$\prod_{j=t-K+1}^t (1 - \gamma_0 \lambda_k)^2 = (1 - \gamma_0 \lambda_k)^{2K} \geq 1 - 2K \gamma_0 \lambda_k. \quad (106)$$

Since $k > k^*$, we have $K \gamma_0 \lambda_k \ll 1$, so the product is lower bounded by a constant (conservatively $\geq 1/2$). The summation over K steps simply accumulates:

$$[\tilde{\mathbf{C}}_t]_k \geq \sigma_{\text{eff}}^2 \lambda_k \sum_{i=t-K}^t \gamma_i^2 \cdot \frac{1}{2} = \frac{1}{2} K \gamma_0^2 \sigma_{\text{eff}}^2 \lambda_k. \quad (107)$$

Combining both cases yields the stated matrix lower bound. $\square$

A.5. Direct Learning Lower Bound

We combine the results from the bias and variance lower bounds to establish the fundamental limit of direct learning.

Theorem 7. *The excess risk $\mathcal{E}(\mathbf{w}_N)$ is bounded by:*

$$\mathbb{E}[\mathcal{E}(\mathbf{w}_N)] \geq \frac{1}{200} \|\boldsymbol{\eta}_0\|_{\boldsymbol{\Sigma}^{>k^*}}^2 + \frac{\sigma_{\text{eff}}^2}{8} \left(\frac{k^*}{K} + K\gamma_0^2 \sum_{i=k^*+1}^d \lambda_i^2 \right) \quad (108)$$

where the effective spectral cutoff is defined as $k^* = \max \left\{ k : \lambda_k > \frac{\log_2 N}{\gamma_0 N} \right\}$.

Proof. The expected excess risk is defined as $\mathbb{E}[\mathcal{E}(\mathbf{w}_N)] = \frac{1}{2} \text{tr}(\boldsymbol{\Sigma} \mathbb{E}[\boldsymbol{\eta}_N \otimes \boldsymbol{\eta}_N])$. Using the bias-variance decomposition $\mathbb{E}[\boldsymbol{\eta}_N \otimes \boldsymbol{\eta}_N] = \mathbf{B}_N + \mathbf{C}_N$, we lower bound each component separately.

1. Bias Component: From Lemma 13, we established the lower bound for the bias energy in the tail subspace:

$$\text{tr}(\boldsymbol{\Sigma} \mathbf{B}_N) \geq \left\| \prod_{i=1}^N (\mathbf{I} - \gamma_i \boldsymbol{\Sigma}) \boldsymbol{\eta}_0 \right\|_{\boldsymbol{\Sigma}^{>k^*}}^2 \geq \frac{1}{100} \|\boldsymbol{\eta}_0\|_{\boldsymbol{\Sigma}^{>k^*}}^2. \quad (109)$$

Multiplying by the factor $\frac{1}{2}$ from the risk definition yields the first term $\frac{1}{200} \|\boldsymbol{\eta}_0\|_{\boldsymbol{\Sigma}^{>k^*}}^2$.

2. Variance Component: From Lemma 14, the variance matrix satisfies $\mathbf{C}_N \succeq \frac{\sigma_{\text{eff}}^2}{4} (\gamma_0 \mathbf{I}^{\leq k^*} + K\gamma_0^2 \boldsymbol{\Sigma}^{>k^*})$. Applying the trace operator $\frac{1}{2} \text{tr}(\boldsymbol{\Sigma} \cdot)$:

For the **Tail part** ($k > k^*$):

$$\frac{1}{2} \text{tr} \left(\boldsymbol{\Sigma} \cdot \frac{\sigma_{\text{eff}}^2}{4} K\gamma_0^2 \boldsymbol{\Sigma}^{>k^*} \right) = \frac{\sigma_{\text{eff}}^2}{8} K\gamma_0^2 \sum_{i=k^*+1}^d \lambda_i^2. \quad (110)$$

For the **Head part** ($k \leq k^*$):

$$\frac{1}{2} \text{tr} \left(\boldsymbol{\Sigma} \cdot \frac{\sigma_{\text{eff}}^2}{4} \gamma_0 \mathbf{I}^{\leq k^*} \right) = \frac{\sigma_{\text{eff}}^2}{8} \sum_{i=1}^{k^*} \gamma_0 \lambda_i. \quad (111)$$

We now connect the term $\sum \gamma_0 \lambda_i$ to k^*/K . By the definition of the effective cutoff $k^* = \max \{ k : \lambda_k > \frac{\log_2 N}{\gamma_0 N} \}$ and recalling that the effective phase length is $K \approx \frac{N}{\log_2 N}$, we have the inequality $\gamma_0 \lambda_i \geq \frac{1}{K}$ for all $i \leq k^*$. Consequently, the summation is bounded by:

$$\sum_{i=1}^{k^*} \gamma_0 \lambda_i \geq \sum_{i=1}^{k^*} \frac{1}{K} = \frac{k^*}{K}. \quad (112)$$

Substituting this back yields the head variance term $\frac{\sigma_{\text{eff}}^2}{8} \frac{k^*}{K}$.

Summing the Bias, Head Variance, and Tail Variance terms completes the proof. $\square$

A.6. Rate in Polynomial-Decay Scenario

We explicitly calculate the convergence rate under the polynomial decay assumption, which is a standard setting for high-dimensional analysis.

Theorem 8. *Assume the eigenvalues of $\boldsymbol{\Sigma}$ follow a polynomial decay law $\lambda_k \asymp k^{-\alpha}$, and the target parameter coefficients satisfy $(w_k^*)^2 \asymp k^\beta$. Then the excess risk $\mathcal{E}(\mathbf{w}_N)$ has the matching upper and lower rate:*

$$\tilde{\Theta} \left(N^{\frac{1+\beta-\alpha}{\alpha}} \cdot \gamma_0^{\frac{1+\beta-\alpha}{\alpha}} \right) + \tilde{\Theta} \left(N^{\frac{1-\alpha}{\alpha}} \cdot \gamma_0^{\frac{1}{\alpha}} \right). \quad (113)$$

Proof. The proof relies on substituting the spectral decay rates into the general bounds derived in Theorem 6 and Section A.5.

Step 1: Effective Spectral Cutoff k^* . The cutoff index k^* is determined by the condition $\gamma_0 N \lambda_{k^*} \asymp 1$ (ignoring logarithmic factors). Substituting $\lambda_k \asymp k^{-\alpha}$:

$$\gamma_0 N (k^*)^{-\alpha} \asymp 1 \implies k^* \asymp (\gamma_0 N)^{\frac{1}{\alpha}}. \quad (114)$$

Step 2: Bias Decay Rate. The bias is dominated by the tail error in the unlearned subspace ($k > k^*$). Using the integral approximation for the sum:

$$\text{Bias} \asymp \sum_{k=k^*}^{\infty} \lambda_k (w_k^*)^2 \asymp \int_{k^*}^{\infty} x^{-\alpha} x^{\beta} dx = \int_{k^*}^{\infty} x^{\beta-\alpha} dx. \quad (115)$$

Assuming $\beta - \alpha < -1$ for convergence, the integral evaluates to:

$$\left[\frac{x^{\beta-\alpha+1}}{\beta-\alpha+1} \right]_{k^*}^{\infty} \asymp (k^*)^{1+\beta-\alpha}. \quad (116)$$

Substituting $k^* \asymp (\gamma_0 N)^{\frac{1}{\alpha}}$:

$$\text{Bias} \asymp \left((\gamma_0 N)^{\frac{1}{\alpha}} \right)^{1+\beta-\alpha} = N^{\frac{1+\beta-\alpha}{\alpha}} \gamma_0^{\frac{1+\beta-\alpha}{\alpha}}. \quad (117)$$

Step 3: Variance Decay Rate. The variance is dominated by the head dimension term $\frac{k^*}{N}$ (assuming the effective noise σ_{eff}^2 is constant).

$$\text{Variance} \asymp \frac{k^*}{N} \asymp \frac{(\gamma_0 N)^{\frac{1}{\alpha}}}{N} = N^{\frac{1}{\alpha}-1} \gamma_0^{\frac{1}{\alpha}} = N^{\frac{1-\alpha}{\alpha}} \gamma_0^{\frac{1}{\alpha}}. \quad (118)$$

(Note: The tail variance term $N \gamma_0^2 \sum_{k>k^*} \lambda_k^2$ is typically of a lower order or comparable magnitude depending on α , so the head term determines the main rate).

Step 4: Conclusion. Combining the Bias and Variance rates yields the final expression:

$$\mathcal{E}(\mathbf{w}_N) \asymp N^{\frac{1+\beta-\alpha}{\alpha}} \gamma_0^{\frac{1+\beta-\alpha}{\alpha}} + N^{\frac{1-\alpha}{\alpha}} \gamma_0^{\frac{1}{\alpha}}. \quad (119)$$

□

B. Proofs of Results in Knowledge Transfer

B.1. SGD Dynamics for Knowledge Transfer

We model the knowledge transfer as an online training process. The student model, $f_{\text{T2S}}(\mathbf{x}) = \langle \mathbf{w}_{\text{T2S}}, \phi_{\text{S}}(\mathbf{x}) \rangle$, is trained over n samples from $\mathbf{w}_{0,\text{T2S}} = \mathbf{0}$ to mimic a fixed, pre-trained teacher, $f_{\text{T}}(\mathbf{x}) = \langle \mathbf{w}_{\text{T}}, \phi_{\text{T}}(\mathbf{x}) \rangle$, using Stochastic Gradient Descent (SGD) on the transfer risk $\mathcal{R}_{\text{Trans}}$ (defined in Eq. 3).

At each time step t , given a learning rate γ_t and a new sample $\mathbf{x}_t$, the student's parameters $\mathbf{w}_{\text{T2S}}$ are updated as follows:

$$\begin{aligned} \mathbf{w}_{t,\text{T2S}} &\leftarrow \mathbf{w}_{t-1,\text{T2S}} - \gamma_t \widehat{\nabla}_t \mathcal{R}_{\text{Trans}}(\mathbf{w}_{t-1,\text{T2S}}) \\ &= \mathbf{w}_{t-1,\text{T2S}} - \gamma_t \left(\phi_{t,\text{S}} \phi_{t,\text{S}}^{\top} \mathbf{w}_{t-1,\text{T2S}} - \phi_{t,\text{S}} f_{\text{T}}(\mathbf{x}_t) \right) \\ &= \mathbf{w}_{t-1,\text{T2S}} - \gamma_t \left(\phi_{t,\text{S}} \phi_{t,\text{S}}^{\top} \mathbf{w}_{t-1,\text{T2S}} - \phi_{t,\text{S}} \phi_{t,\text{T}}^{\top} \mathbf{w}_{\text{T}} \right) \\ &= (\mathbf{I} - \gamma_t \mathbf{M}_{\text{S}} \Phi_t \Phi_t^{\top} \mathbf{M}_{\text{S}}^{\top}) \mathbf{w}_{t-1,\text{T2S}} + \gamma_t \mathbf{M}_{\text{S}} \Phi_t \Phi_t^{\top} \mathbf{M}_{\text{T}}^{\top} \mathbf{w}_{\text{T}}. \end{aligned} \quad (120)$$

Here, we use the shorthand $\phi_{t,\text{S}} := \phi_{\text{S}}(\mathbf{x}_t)$, $\phi_{t,\text{T}} := \phi_{\text{T}}(\mathbf{x}_t)$, and $\Phi_t := \Phi(\mathbf{x}_t)$ for the feature vectors evaluated at $\mathbf{x}_t$. The training is assumed to start from $\mathbf{w}_{0,\text{T2S}} = \mathbf{0}$. The final line substitutes the feature basis definitions from our setup.

This iterative process converges to the optimal parameters $\mathbf{w}_{\text{T2S}}^*$ that minimize the population transfer risk. We find this

optimum by setting the gradient of the population risk to zero:

$$\begin{aligned}
 \nabla_{\mathbf{w}_{T2S}} \mathcal{R}_{\text{Trans}} &= \nabla_{\mathbf{w}_{T2S}} \left[\frac{1}{2} \mathbb{E}_{\mathbf{x}} (\langle \mathbf{w}_{T2S}, \phi_S(\mathbf{x}) \rangle - \langle \mathbf{w}_T, \phi_T(\mathbf{x}) \rangle)^2 \right] \\
 &= \mathbb{E}_{\mathbf{x}} [\phi_S(\mathbf{x}) \phi_S(\mathbf{x})^\top \mathbf{w}_{T2S} - \phi_S(\mathbf{x}) \phi_T(\mathbf{x})^\top \mathbf{w}_T] \\
 &= \mathbb{E}_{\mathbf{x}} [\mathbf{M}_S \Phi(\mathbf{x}) \Phi(\mathbf{x})^\top \mathbf{M}_S^\top] \mathbf{w}_{T2S} - \mathbb{E}_{\mathbf{x}} [\mathbf{M}_S \Phi(\mathbf{x}) \Phi(\mathbf{x})^\top \mathbf{M}_T^\top] \mathbf{w}_T \\
 &= \mathbf{M}_S \mathbf{M}_S^\top \mathbf{w}_{T2S} - \mathbf{M}_S \mathbf{M}_T^\top \mathbf{w}_T = \mathbf{0}.
 \end{aligned} \tag{121}$$

This is a standard normal equation, where we used the assumption $\mathbb{E}[\Phi(\mathbf{x}) \Phi(\mathbf{x})^\top] = \mathbf{I}$. The solution $\mathbf{w}_{T2S}^*$, which represents the optimal parameters the student can learn from the teacher, is given by:

$$\mathbf{w}_{T2S}^* = (\mathbf{M}_S \mathbf{M}_S^\top)^+ \mathbf{M}_S \mathbf{M}_T^\top \mathbf{w}_T = (\mathbf{M}_S^\top)^+ \mathbf{M}_T^\top \mathbf{w}_T, \tag{122}$$

where $(\cdot)^+$ denotes the Moore-Penrose pseudoinverse.

The excess risk formula is in a concise form similar to 44:

$$\begin{aligned}
 \mathcal{E}_{\text{Trans}} &= \mathcal{R}_{\text{Trans}}(\mathbf{w}_{T2S}) - \mathcal{R}_{\text{Trans}}(\mathbf{w}_{T2S}^*) \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*, \phi_S(\mathbf{x}) \rangle]^2 + \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*, \phi_S(\mathbf{x}) \rangle (\langle \mathbf{w}_{T2S}^*, \phi_S(\mathbf{x}) \rangle - \langle \mathbf{w}_T, \phi_T(\mathbf{x}) \rangle)] \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*, \phi_S(\mathbf{x}) \rangle]^2 + \mathbb{E}_{\mathbf{x}} [(\mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*)^\top \mathbf{M}_S \Phi(\mathbf{x}) \Phi(\mathbf{x})^\top (\mathbf{M}_S^\top \mathbf{M}_S^{\top+} - \mathbf{I}) \mathbf{M}_T^\top \mathbf{w}_T] \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*, \phi_S(\mathbf{x}) \rangle]^2 + (\mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*)^\top (\mathbf{M}_S \mathbf{M}_S^\top \mathbf{M}_S^{\top+} - \mathbf{M}_S) \mathbf{M}_T^\top \mathbf{w}_T \\
 &= \frac{1}{2} \mathbb{E}_{\mathbf{x}} [\langle \mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*, \phi_S(\mathbf{x}) \rangle]^2 \\
 &= \frac{1}{2} \|\mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*\|_{\Sigma_S}^2 = \frac{1}{2} \langle (\mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*) \otimes (\mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*), \Sigma_S \rangle.
 \end{aligned} \tag{123}$$

To estimate the magnitude of the risk we only have to investigate the dynamics of $(\mathbf{w}_{T2S} - \mathbf{w}_{T2S}^*)^{\otimes 2}$. The update rule 120 could be further transformed into

$$\begin{aligned}
 \mathbf{w}_{t,T2S} - \mathbf{w}_{T2S}^* &= (\mathbf{w}_{t-1,T2S} - \mathbf{w}_{T2S}^*) - \gamma_t (\mathbf{M}_S \Phi_t \Phi_t^\top \mathbf{M}_S^\top \mathbf{w}_{t-1,T2S} - \mathbf{M}_S \Phi_t \Phi_t^\top \mathbf{M}_T^\top \mathbf{w}_T) \\
 &= (\mathbf{I} - \gamma_t \Phi_{t,S} \Phi_{t,S}^\top) (\mathbf{w}_{t-1,T2S} - \mathbf{w}_{T2S}^*) + \gamma_t \mathbf{M}_S \Phi_t \Phi_t^\top \Pi_S^\perp \mathbf{M}_T^\top \mathbf{w}_T \\
 &= (\mathbf{I} - \gamma_t \Phi_{t,S} \Phi_{t,S}^\top) (\mathbf{w}_{t-1,T2S} - \mathbf{w}_{T2S}^*) + \gamma_t (\Phi_t^\top \Pi_S^\perp \mathbf{M}_T^\top \mathbf{w}_T) \cdot \phi_{t,S} \\
 &= (\mathbf{I} - \gamma_t \Sigma_S) (\mathbf{w}_{t-1,T2S} - \mathbf{w}_{T2S}^*) + \gamma_t (\Sigma_S - \Phi_{t,S} \Phi_{t,S}^\top) (\mathbf{w}_{t-1,T2S} - \mathbf{w}_{T2S}^*) \\
 &\quad + \gamma_t (\Phi_t^\top \Pi_S^\perp \mathbf{M}_T^\top \mathbf{w}_T) \cdot \phi_{t,S},
 \end{aligned} \tag{124}$$

where $\Pi_S^\perp := \mathbf{I} - \mathbf{M}_S^\top \mathbf{M}_S^{\top+}$ is a matrix that projects vectors to the subspace perpendicular to the image space of $\mathbf{M}_S$.

The random variable $\Phi_t^\top \Pi_S^\perp \mathbf{M}_T^\top \mathbf{w}_T$ could be understood as the noise generated by the parts of the teacher model that the student model could not understand, i.e., beyond its expressivity. The following calculations show that it actually acts like noise relative to ϕ_S .

$$\mathbb{E}_{\mathbf{x}} (\Phi_t^\top \Pi_S^\perp \mathbf{M}_T^\top \mathbf{w}_T) \cdot \phi_{t,S} = \mathbb{E}_{\mathbf{x}} \mathbf{M}_S \Phi_t \Phi_t^\top (\mathbf{I} - \mathbf{M}_S^\top \mathbf{M}_S^{\top+}) \mathbf{M}_T^\top \mathbf{w}_T = \mathbf{M}_S (\mathbf{I} - \mathbf{M}_S^\top \mathbf{M}_S^{\top+}) \mathbf{M}_T^\top \mathbf{w}_T = \mathbf{0}. \tag{125}$$

We henceforth define the transfer noise

$$\hat{\sigma}^2 := \mathbb{E}_{\rho_{\mathbf{x}}} \left[\left\| \Phi^\top \Pi_S^\perp \mathbf{w}_* \right\|^2 \right] = \|\mathbf{w}_*\|_{\Pi_S^\perp}^2. \tag{126}$$

For theoretical analysis, we define the iterate $\boldsymbol{\eta}_t := \mathbf{w}_{t,T2S} - \mathbf{w}_{T2S}^*$. Its iteration could be written in the following compact form

$$\boldsymbol{\eta}_t = \hat{\mathbf{A}}_t \boldsymbol{\eta}_{t-1} + \gamma_t \zeta_t, \tag{127}$$

where $\hat{\mathbf{A}}_t = \mathbf{I} - \gamma_t \phi_{t,S} \phi_{t,S}^\top$, $\zeta_t := (\Phi_t^\top \Pi^\perp \mathbf{M}_T^\top \mathbf{w}_T) \cdot \phi_{t,S}$. $\mathbb{E}[\zeta_t] = 0$. Also, we define $\mathbf{A}_t = \mathbf{I} - \gamma_t \Sigma_S$ which is the non-stochastic version (expectation) of $\hat{\mathbf{A}}_t$.

We notice the evolution of the transfer iterate (equation 127) follows exactly same dynamics as the direct learning iterate (equation 48). Apply similar procedures we shall derive a bound on the transfer risk.

Theorem 9. *The excess risk $\mathcal{E}(\mathbf{w}_N)$ is bounded by:*

$$\mathbb{E}[\mathcal{E}(\mathbf{w}_{n,T2S})] \leq \frac{1}{2n^2} \mathbb{E}[\|\mathbf{w}_{N,T}\|_{\Sigma \leq k^*}^2] + \frac{1}{2} \mathbb{E}[\|\mathbf{w}_{N,T}\|_{\Sigma > k^*}^2] + 16(\psi \mathbb{E}[\|\mathbf{w}_{N,T}\|_{\Sigma}^2] + \hat{\sigma}^2) \left(\frac{k^*}{K'} + K'(\gamma'_0)^2 \sum_{i=k^*+1}^d \lambda_i^2 \right). \quad (128)$$

where $k^* = \max\{k : \lambda_{k,S} > 2 \ln n \log_2 n / (\gamma'_0 n)\}$.

Proof. To prove Theorem 9, we map the specific dynamics of the knowledge transfer process derived in Eq. equation 120 to the general SGD dynamics analyzed in Theorem 6.

Theorem 6 provides the upper bound for excess risk under step-decay SGD:

$$\mathbb{E}[\mathcal{E}(\mathbf{w}_n)] \leq \underbrace{\frac{1}{2n^2} \|\boldsymbol{\eta}_0\|_{\Sigma \leq k^*}^2 + \frac{1}{2} \|\boldsymbol{\eta}_0\|_{\Sigma > k^*}^2}_{\text{Bias Term}} + \underbrace{16(\psi \|\boldsymbol{\eta}_0\|_{\Sigma}^2 + \sigma_{\text{eff}}^2) \left(\frac{k^*}{K} + K\gamma_0^2 \sum_{i>k^*} \lambda_i^2 \right)}_{\text{Variance Term}}. \quad (129)$$

Substituting the mapped variables into this equation:

1. Replace the sample size N with the student's training steps n .
2. Replace the generic covariance Σ with Σ_S and its eigenvalues $\lambda_{k,S}$.
3. Replace the initial error norms $\|\boldsymbol{\eta}_0\|_{(\cdot)}^2$ with the teacher's norms $\|\mathbf{w}_{N,T}\|_{(\cdot)}^2$.
4. Replace the noise variance σ_{eff}^2 with the transfer residual $\hat{\sigma}^2$.
5. Replace the stepsize interval K with $K' = n / \log_2 n$.

Then we shall have the desired upper bound. $\square$

B.2. Proof of Geometric Consistency Condition (Lemma 2)

We derive the consistency condition by explicitly comparing the optimal parameters obtained via Direct Learning versus Knowledge Transfer using the Moore-Penrose pseudoinverse properties.

Proof. Let the ground truth function be represented by $\mathbf{w}_*$ in the feature space Φ . The student and teacher parameterize their functions as $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{M} \Phi(\mathbf{x}) = (\mathbf{M}^\top \mathbf{w})^\top \Phi(\mathbf{x})$. Define the orthogonal projection operators onto the image spaces of the student and teacher feature maps as $\Pi_S = \mathbf{M}_S^\top (\mathbf{M}_S^\top)^+$ and $\Pi_T = \mathbf{M}_T^\top (\mathbf{M}_T^\top)^+$, respectively.

1. Optimal Parameters Derivation

Direct Learning: The student directly approximates $\mathbf{w}_*$. The minimal-norm solution to $\min_{\mathbf{w}} \|\mathbf{M}_S^\top \mathbf{w} - \mathbf{w}_*\|^2$ is:

$$\mathbf{w}_S^* = (\mathbf{M}_S^\top)^+ \mathbf{w}_*. \quad (130)$$

Knowledge Transfer: The teacher first learns the ground truth, yielding $\mathbf{w}_T^* = (\mathbf{M}_T^\top)^+ \mathbf{w}_*$. The student then mimics the teacher's effective parameter $\mathbf{w}_{T,\text{eff}} = \mathbf{M}_T^\top \mathbf{w}_T^* = \Pi_T \mathbf{w}_*$. The student's optimal solution $\mathbf{w}_{T2S}^{\text{opt}}$ for this target is:

$$\mathbf{w}_{T2S}^{\text{opt}} = (\mathbf{M}_S^\top)^+ (\mathbf{M}_T^\top \mathbf{w}_T^*) = (\mathbf{M}_S^\top)^+ \Pi_T \mathbf{w}_*. \quad (131)$$

2. Consistency Condition

The gap between the direct solution and the transfer solution is:

$$\begin{aligned}\mathbf{w}_S^* - \mathbf{w}_{T2S}^{\text{opt}} &= (\mathbf{M}_S^\top)^+ \mathbf{w}_* - (\mathbf{M}_S^\top)^+ \mathbf{\Pi}_T \mathbf{w}_* \\ &= (\mathbf{M}_S^\top)^+ (\mathbf{I} - \mathbf{\Pi}_T) \mathbf{w}_* \\ &= (\mathbf{M}_S^\top)^+ \mathbf{\Pi}_T^\perp \mathbf{w}_*.\end{aligned}\tag{132}$$

Using the property that $(\mathbf{M}_S^\top)^+ = (\mathbf{M}_S^\top)^+ \mathbf{\Pi}_S$ (since the pseudoinverse acts on the range space), we can rewrite this as:

$$\mathbf{w}_S^* - \mathbf{w}_{T2S}^{\text{opt}} = (\mathbf{M}_S^\top)^+ \mathbf{\Pi}_S \mathbf{\Pi}_T^\perp \mathbf{w}_*.\tag{133}$$

Thus, consistency ($\mathbf{w}_S^* = \mathbf{w}_{T2S}^{\text{opt}}$) holds if and only if $\mathbf{\Pi}_S \mathbf{\Pi}_T^\perp \mathbf{w}_* = \mathbf{0}$ ($\mathbf{M}_S$ has the same null space as $\mathbf{\Pi}_S$).

3. Static Alignment Bias (Risk Gap)

The generalization risk is measured in the feature space: $\mathcal{R}(\mathbf{w}) = \frac{1}{2} \|\mathbf{M}_S^\top \mathbf{w} - \mathbf{w}_*\|^2$.

For the transfer student, the prediction is $\mathbf{M}_S^\top \mathbf{w}_{T2S}^{\text{opt}} = \mathbf{\Pi}_S \mathbf{\Pi}_T \mathbf{w}_*$. The residual vector can be decomposed orthogonally:

$$\begin{aligned}\mathbf{M}_S^\top \mathbf{w}_{T2S}^{\text{opt}} - \mathbf{w}_* &= (\mathbf{\Pi}_S \mathbf{\Pi}_T \mathbf{w}_* - \mathbf{\Pi}_S \mathbf{w}_*) + (\mathbf{\Pi}_S \mathbf{w}_* - \mathbf{w}_*) \\ &= -\mathbf{\Pi}_S (\mathbf{I} - \mathbf{\Pi}_T) \mathbf{w}_* - (\mathbf{I} - \mathbf{\Pi}_S) \mathbf{w}_* \\ &= \underbrace{-\mathbf{\Pi}_S \mathbf{\Pi}_T^\perp \mathbf{w}_*}_{\in \text{Range}(\mathbf{M}_S^\top)} + \underbrace{-\mathbf{\Pi}_S^\perp \mathbf{w}_*}_{\in \text{Range}(\mathbf{M}_S^\top)^\perp}.\end{aligned}\tag{134}$$

Since the two terms are orthogonal, the squared norm splits:

$$2\mathcal{R}_{T2S}(\mathbf{w}_{T2S}^{\text{opt}}) = \|\mathbf{\Pi}_S \mathbf{\Pi}_T^\perp \mathbf{w}_*\|^2 + \|\mathbf{\Pi}_S^\perp \mathbf{w}_*\|^2.\tag{135}$$

Note that the risk of direct learning is exactly the second term: $2\mathcal{R}_S(\mathbf{w}_S^*) = \|\mathbf{\Pi}_S \mathbf{w}_* - \mathbf{w}_*\|^2 = \|\mathbf{\Pi}_S^\perp \mathbf{w}_*\|^2$.

Subtracting the two risks yields the irreducible static alignment bias:

$$\mathcal{R}_{T2S}(\mathbf{w}_{T2S}^{\text{opt}}) - \mathcal{R}_S(\mathbf{w}_S^*) = \frac{1}{2} \|\mathbf{\Pi}_S \mathbf{\Pi}_T^\perp \mathbf{w}_*\|^2.\tag{136}$$

□

B.3. Proof of Theorem 1

We decompose the excess teacher-to-student risk into components related to the student's learning dynamics and the geometric misalignment between the teacher and student subspaces.

Proof. The excess transfer risk is defined as $\mathcal{E}_{T2S} = \frac{1}{2} \mathbb{E}_{N,n} \|\mathbf{w}_{n,T2S} - \mathbf{w}_S^*\|_{\Sigma_S}^2$. By the triangle inequality and the orthogonality of the spectral decomposition (Head vs. Tail), we split the risk into four dominant terms:

$$\begin{aligned}\mathbb{E}[\mathcal{E}_{T2S}] &\leq \underbrace{\frac{1}{2} \mathbb{E}_N \|\mathbf{w}_{T2S}^* - \mathbf{w}_S^*\|_{\Sigma_S^{\leq k^*}}^2}_{\text{(I) Head Alignment Error}} + \underbrace{\frac{1}{2} \mathbb{E}_{N,n} \|\mathbf{w}_{n,T2S} - \mathbf{w}_{T2S}^*\|_{\Sigma_S^{\leq k^*}}^2}_{\text{(II) Head Optimization Error}} \\ &\quad + \underbrace{\frac{1}{2} \mathbb{E}_{N,n} \|\mathbf{w}_{n,T2S}\|_{\Sigma_S^{> k^*}}^2}_{\text{(III) Tail Learning Error}} + \underbrace{\frac{1}{2} \|\mathbf{w}_S^*\|_{\Sigma_S^{> k^*}}^2}_{\text{(IV) Tail Approximation Error}}.\end{aligned}\tag{137}$$

We bound these terms via the following lemmas.

□

PART 1: HEAD ALIGNMENT ERROR (TERM I)

Lemma 15 (Head Misalignment Bound). *The distance between the student's transfer target $\mathbf{w}_{T2S}^*$ and the optimal direct parameter $\mathbf{w}_S^*$ in the learned subspace is bounded by:*

$$\mathbb{E}_N \left[\frac{1}{2} \|\mathbf{w}_{T2S}^* - \mathbf{w}_S^*\|_{\Sigma_S^{\leq k^*}}^2 \right] \leq \left\langle \mathbb{E}_N [(\boldsymbol{\eta}_{N,T})^{\otimes 2}], \mathbf{M}_T \mathbf{\Pi}_S^{\leq k^*} \mathbf{M}_T^\top \right\rangle + C \cdot \|\mathbf{\Pi}_T^\perp \mathbf{w}_*\|_{\Sigma_S^{\leq k^*}}.\tag{138}$$

Proof. Recall $\mathbf{w}_{T2S}^* = (\mathbf{M}_S^\top)^\dagger \mathbf{M}_T^\top \mathbf{w}_{N,T}$ and $\mathbf{w}_S^* = (\mathbf{M}_S^\top)^\dagger \mathbf{w}_*$. Substituting these definitions:

$$\begin{aligned} \|\mathbf{w}_{T2S}^* - \mathbf{w}_S^*\|_{\Sigma_S^{\leq k^*}}^2 &= \left\| \Sigma_S^{1/2} (\mathbf{M}_S^\top)^\dagger (\mathbf{M}_T^\top \mathbf{w}_{N,T} - \mathbf{w}_*) \right\|_{\Pi^{\leq k^*}}^2 \\ &= \left\langle (\mathbf{M}_T^\top \mathbf{w}_{N,T} - \mathbf{w}_*)^{\otimes 2}, (\mathbf{M}_S^\top)^\dagger \Sigma_S^{\leq k^*} (\mathbf{M}_S^\top)^\dagger \right\rangle. \end{aligned} \quad (139)$$

Using the SVD $\mathbf{M}_S = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^\top$, the operator simplifies to the projection onto the student's head subspace:

$$(\mathbf{M}_S^\top)^\dagger \Sigma_S^{\leq k^*} (\mathbf{M}_S^\top)^\dagger = \mathbf{V} \mathbf{\Lambda}^{-1} (\mathbf{\Lambda}^2)^{\leq k^*} \mathbf{\Lambda}^{-1} \mathbf{V}^\top = \mathbf{V}_{\leq k^*} \mathbf{V}_{\leq k^*}^\top = \Pi_S^{\leq k^*}. \quad (140)$$

Expanding the error term $\mathbf{M}_T^\top \mathbf{w}_{N,T} - \mathbf{w}_* = \mathbf{M}_T^\top (\mathbf{w}_{N,T} - \mathbf{w}_T^*) - \Pi_T^\perp \mathbf{w}_*$:

$$\begin{aligned} \text{LHS} &= \mathbb{E}_N \left\langle (\mathbf{M}_T^\top \boldsymbol{\eta}_{N,T} - \Pi_T^\perp \mathbf{w}_*)^{\otimes 2}, \Pi_S^{\leq k^*} \right\rangle \\ &= \mathbb{E}_N \left\langle (\mathbf{M}_T^\top \boldsymbol{\eta}_{N,T})^{\otimes 2}, \Pi_S^{\leq k^*} \right\rangle + \left\langle (\Pi_T^\perp \mathbf{w}_*)^{\otimes 2}, \Pi_S^{\leq k^*} \right\rangle \\ &\quad - 2 \mathbb{E}_N \left[\mathbf{w}_*^\top \Pi_T^\perp \Pi_S^{\leq k^*} \mathbf{M}_T^\top \boldsymbol{\eta}_{N,T} \right]. \end{aligned} \quad (141)$$

The first term matches the lemma statement. The second term is $\|\Pi_T^\perp \mathbf{w}_*\|_{\Pi_S^{\leq k^*}}^2$. The cross term is bounded by Cauchy-Schwarz and the initial teacher error bound:

$$\text{Cross Term} \leq 2 \|\Pi_T^\perp \mathbf{w}_*\|_{\Pi_S^{\leq k^*}} \sqrt{\mathbb{E} \|\mathbf{M}_T^\top \boldsymbol{\eta}_{N,T}\|_{\Pi_S^{\leq k^*}}^2} \leq C \cdot \|\Pi_T^\perp \mathbf{w}_*\|_{\Pi_S^{\leq k^*}}. \quad (142)$$

Due to the fact that $\|\Pi_T^\perp \mathbf{w}_*\|_{\Pi_S^{\leq k^*}} \leq \|\Pi_T^\perp \mathbf{w}_*\|_{\Pi_S}$ should be a small number in most cases, we only keep its linear form in our bound. This completes the proof. $\square$

PART 2: TAIL OPTIMIZATION DYNAMICS (TERM III)

In the tail subspace ($k > k^*$), the eigenvalues are small ($\lambda_k \rightarrow 0$). The student's parameters $\mathbf{w}_{n,T2S}$ barely move from initialization ($\mathbf{0}$). We define the maximal movement fraction δ :

$$\delta := \sup_{j > k^*} \left(1 - \exp \left(-2.02 \frac{n}{\log n} \gamma_0 \lambda_j \right) \right) \approx 2.02 \frac{n}{\log n} \gamma_0 \lambda_{k^*+1} \ll 1. \quad (143)$$

Lemma 16 (Tail Learning Bound). *The energy of the student's parameters in the tail subspace is bounded by the teacher's variance and the ground truth signal, scaled by the small movement factor δ^2 :*

$$\mathbb{E}_{N \otimes n} \left[\|\mathbf{w}_{n,T2S}\|_{\Sigma_S^{> k^*}}^2 \right] \leq 2\delta^2 \left\langle \mathbb{E}_N \left[\boldsymbol{\eta}_{N,T}^{\otimes 2} \right], \mathbf{M}_T \Pi_S^{> k^*} \mathbf{M}_T^\top \right\rangle + 2\delta^2 \|\Pi_T \mathbf{w}_*\|_{\Pi_S^{> k^*}}^2 + \left\langle \mathbf{P}_{n,T2S}^{\text{var}}, \Sigma_S^{> k^*} \right\rangle. \quad (144)$$

Proof. Decompose the squared norm:

$$\mathbb{E} \|\mathbf{w}_{n,T2S}\|_{\Sigma_S^{> k^*}}^2 = \|\mathbb{E}_n[\mathbf{w}_{n,T2S}]\|_{\Sigma_S^{> k^*}}^2 + \mathbb{E} \|\mathbf{w}_{n,T2S} - \mathbb{E}_n[\mathbf{w}_{n,T2S}]\|_{\Sigma_S^{> k^*}}^2. \quad (145)$$

The second term is the variance of the student's tail optimization, denoted $\left\langle \mathbf{P}_{n,T2S}^{\text{var}}, \Sigma_S^{> k^*} \right\rangle$. For the first term (Bias), using the step-decay dynamics on the target $\mathbf{w}_{T2S}^*$:

$$\mathbb{E}_n[\mathbf{w}_{n,T2S}] = (\mathbf{I} - \prod (\mathbf{I} - \gamma_i \Sigma_S)) \mathbf{w}_{T2S}^*. \quad (146)$$

We know that

$$1 - \exp \left(-2.02 \frac{n}{\log n} \gamma_0 \lambda_{k^*+1} \right) \leq 2.02 \frac{n}{\log n} \gamma_0 \lambda_{k^*+1} < \delta. \quad (147)$$

Therefore, in the tail ($k > k^*$), the contraction is minimal, so the ‘‘learned’’ portion is bounded by δ :

$$\|\mathbb{E}_n[\mathbf{w}_{n,T2S}]\|_{\Sigma_S^{> k^*}} \leq \delta \|\mathbf{w}_{T2S}^*\|_{\Sigma_S^{> k^*}}. \quad (148)$$

Substituting $\mathbf{w}_{T2S}^* = (\mathbf{M}_S^\top)^\dagger \mathbf{M}_T^\top \mathbf{w}_{N,T}$:

$$\mathbb{E}_N \|\mathbf{w}_{T2S}^*\|_{\Sigma_S^{> k^*}}^2 = \mathbb{E}_N \|\mathbf{M}_T^\top \mathbf{w}_{N,T}\|_{\Pi_S^{> k^*}}^2 \leq 2 \mathbb{E}_N \|\mathbf{M}_T^\top \boldsymbol{\eta}_{N,T}\|_{\Pi_S^{> k^*}}^2 + 2 \|\Pi_T \mathbf{w}_*\|_{\Pi_S^{> k^*}}^2. \quad (149)$$

Combining these yields the result. $\square$

PART 3: UNIFIED BOUND (FINAL ASSEMBLY)

Substituting Lemmas 15 and Part 2 into the risk decomposition:

$$\begin{aligned}
 \mathbb{E}_{N \otimes n} [\mathcal{E}_{T2S}] &\leq \underbrace{\frac{1}{2} \left\langle \mathbb{E}_N [\boldsymbol{\eta}_{N,T}^{\otimes 2}], \mathbf{M}_T \boldsymbol{\Pi}_S^{\leq k^*} \mathbf{M}_T^\top \right\rangle}_{\text{Head Alignment}} + C \cdot \|\boldsymbol{\Pi}_T^\perp \mathbf{w}_*\|_{\boldsymbol{\Pi}_S^{\leq k^*}} \\
 &\quad + \underbrace{2\delta^2 \left\langle \mathbb{E}_N [\boldsymbol{\eta}_{N,T}^{\otimes 2}], \mathbf{M}_T \boldsymbol{\Pi}_S^{> k^*} \mathbf{M}_T^\top \right\rangle + 2\delta^2 \|\boldsymbol{\Pi}_T \mathbf{w}_*\|_{\boldsymbol{\Pi}_S^{> k^*}}^2}_{\text{Tail Teacher Leakage}} \\
 &\quad + \underbrace{\frac{1}{2} \left\langle \mathbb{E}_{N \otimes n} [\mathbf{P}_{n,T2S}], \boldsymbol{\Sigma}_S^{\leq k^*} \right\rangle}_{\text{Head Student Opt.}} + \underbrace{\left\langle \mathbf{P}_{n,T2S}^{\text{var}}, \boldsymbol{\Sigma}_S^{> k^*} \right\rangle}_{\text{Tail Student Var.}} \\
 &\quad + \frac{1}{2} \|\mathbf{w}_S^*\|_{\boldsymbol{\Sigma}_S^{> k^*}}^2.
 \end{aligned} \tag{150}$$

This concludes the proof.

C. Proof of Results in Distillation

C.1. Preliminary Lemmas

Note that in the proof of Theorem 1 the δ and k^* can be determined each other. In the proofs of distillation theorems, we mainly focus on the process in which the strong teacher passes its strength to the weak student while leaving aside the spectral denoising mechanism for now. Therefore, in the distillation section, we will let $k^* = d_S$, and in this case, the elements $> k^*$ is an empty set. Theorem 1 in this case becomes:

Corollary 2 (The Teacher-to-Student Risk Decomposition Distillation Version). *Let the Student be trained on n samples labeled by a fixed Teacher (itself trained on N samples). The expected knowledge transfer excess risk $\mathcal{E}_{T2S}$ is bounded by:*

$$\mathbb{E}[\mathcal{E}_{T2S}(N, n)] \leq \frac{1}{2} \langle \mathbf{P}_{N,T}, \mathbf{M}_T \boldsymbol{\Pi}_S \mathbf{M}_T^\top \rangle + \frac{1}{2} \langle \mathbf{P}_{n,T2S}, \boldsymbol{\Sigma}_S \rangle + C \left\| \boldsymbol{\Pi}_S \boldsymbol{\Pi}_T^\perp \mathbf{w}_* \right\| \tag{151}$$

where $k^* = \max\{k : \lambda_{k,S} > \delta \log_2 n / (4\gamma'_0 n)\}$ is the effective dimension, $\mathbf{P}_{n,T2S} := \mathbb{E}[\boldsymbol{\eta}_{n,T2S}^{\otimes 2}]$ is the second moment matrix of $\boldsymbol{\eta}_{n,T2S} := \mathbf{w}_{n,T2S} - \mathbf{w}_{T2S}^*$ and C is a constant.

This bound could be further simplified with the concept of spectral compatibility:

Lemma 17 (Geometric Consistency for Compatible Students). *For a spectrally compatible student (i.e., one whose feature space is a subspace of the teacher's, $\text{Range}(\mathbf{M}_S^\top) \subseteq \text{Range}(\mathbf{M}_T^\top)$), the geometric consistency condition holds:*

$$\boldsymbol{\Pi}_S \boldsymbol{\Pi}_T^\perp \mathbf{w}_* = \mathbf{0}. \tag{152}$$

Proof. The definition of spectral compatibility implies that any feature representable by the student is also representable by the teacher. In terms of subspaces, let $\mathcal{S} = \text{Range}(\mathbf{M}_S^\top)$ and $\mathcal{T} = \text{Range}(\mathbf{M}_T^\top)$; then $\mathcal{S} \subseteq \mathcal{T}$.

This inclusion property implies that projecting a vector from the student subspace $\mathcal{S}$ onto the teacher subspace $\mathcal{T}$ has no effect (acts as identity). Mathematically:

$$\boldsymbol{\Pi}_T \boldsymbol{\Pi}_S = \boldsymbol{\Pi}_S. \tag{153}$$

Since projection matrices are symmetric, taking the transpose yields the equivalent condition:

$$\boldsymbol{\Pi}_S \boldsymbol{\Pi}_T = \boldsymbol{\Pi}_S. \tag{154}$$

Now, we expand the term in the lemma using the definition of the orthogonal complement projector $\boldsymbol{\Pi}_T^\perp = \mathbf{I} - \boldsymbol{\Pi}_T$:

$$\begin{aligned}
 \boldsymbol{\Pi}_S \boldsymbol{\Pi}_T^\perp \mathbf{w}_* &= \boldsymbol{\Pi}_S (\mathbf{I} - \boldsymbol{\Pi}_T) \mathbf{w}_* \\
 &= (\boldsymbol{\Pi}_S - \boldsymbol{\Pi}_S \boldsymbol{\Pi}_T) \mathbf{w}_* \\
 &= (\boldsymbol{\Pi}_S - \boldsymbol{\Pi}_S) \mathbf{w}_* \\
 &= \mathbf{0}.
 \end{aligned} \tag{155}$$

Geometrically, this means that the part of the ground truth the teacher cannot learn ($\Pi_T^\perp \mathbf{w}_*$) lies in the teacher's null space, which contains the student's null space ($\mathcal{T}^\perp \subseteq \mathcal{S}^\perp$). Therefore, the student cannot capture any of it either. $\square$

Next we show that with sufficient unlabeled data (n samples), the optimization error of the transfer process is negligible.

Lemma 18. *Under Assumption 2 (Polynomial Spectral Decay, $\lambda_k \propto k^{-\alpha}$ with $\alpha > 1$), with sufficient unlabeled data, the optimization error goes to zero:*

$$\lim_{n \rightarrow \infty} \frac{1}{2} \langle \mathbf{P}_{n, \text{T2S}}, \Sigma_S \rangle = 0. \quad (156)$$

Proof. The term $\frac{1}{2} \langle \mathbf{P}_{n, \text{T2S}}, \Sigma_S \rangle$ represents the expected excess risk of the student relative to its optimal target $\mathbf{w}_{\text{T2S}}^*$. We invoke the upper bound derived in Theorem 9:

$$\frac{1}{2} \langle \mathbf{P}_{n, \text{T2S}}, \Sigma_S \rangle \leq \underbrace{\frac{C_1}{n^2}}_{\text{Init Decay}} + \underbrace{\frac{1}{2} \|\mathbf{w}_{N, \text{T}}\|_{\Sigma_{>k^*}}^2}_{\text{Tail Bias}} + \underbrace{C_2 \left(\frac{k^*}{n} + n\gamma^2 \sum_{i>k^*} \lambda_i^2 \right)}_{\text{Variance}}. \quad (157)$$

Here n is the number of student steps (samples), and the effective cutoff is $k^* = \max\{k : \lambda_k > \frac{C \log n}{n}\}$.

Under Assumption 2, let $\lambda_k \asymp k^{-\alpha}$ for $\alpha > 1$.

1. Behavior of k^* : The condition $\lambda_{k^*} \asymp \frac{\log n}{n}$ implies $(k^*)^{-\alpha} \asymp \frac{\log n}{n}$, thus:

$$k^* \asymp \left(\frac{n}{\log n} \right)^{1/\alpha}. \quad (158)$$

As $n \rightarrow \infty$, $k^* \rightarrow \infty$.

2. Convergence of Terms:

- **Initial Decay:** $\frac{1}{n^2} \rightarrow 0$ trivially.
- **Tail Bias:** Since Σ is trace-class ($\sum \lambda_k < \infty$), the tail energy sum $\sum_{i=k^*+1}^{\infty} \lambda_i (w^{(i)})^2$ must vanish as the start index $k^* \rightarrow \infty$. Specifically, $\|\cdot\|_{\Sigma_{>k^*}}^2 \rightarrow 0$.
- **Head Variance:** The term scales as $\frac{k^*}{n} \asymp \frac{n^{1/\alpha}}{n} = n^{\frac{1}{\alpha}-1}$. Since $\alpha > 1$, the exponent is negative, so $\frac{k^*}{n} \rightarrow 0$.
- **Tail Variance:** With a proper decaying learning rate schedule (e.g., $\gamma \propto 1/n$), this term is dominated by the Head Variance rate and also vanishes.

Since all components of the upper bound converge to zero, the optimization error vanishes asymptotically. $\square$

Finally we derive the most important lemma in distillation:

Lemma 19. *Under Assumption 2, for a spectrally compatible student, and assuming sufficient unlabeled data ($n \rightarrow \infty$):*

$$\mathbb{E}[\mathcal{E}_{\text{T2S}}(N, n)] \leq \mathbb{E}[\mathcal{E}_{\text{T}}(N)] \quad (159)$$

Proof. We start from the risk decomposition in Corollary 2:

$$\mathbb{E}[\mathcal{E}_{\text{T2S}}] \leq \underbrace{\frac{1}{2} \langle \mathbf{P}_{N, \text{T}}, \mathbf{M}_{\text{T}} \Pi_S \mathbf{M}_{\text{T}}^\top \rangle}_{\text{Teacher Estimation Error}} + \underbrace{\frac{1}{2} \langle \mathbf{P}_{n, \text{T2S}}, \Sigma_S \rangle}_{\text{Student Opt. Error}} + \underbrace{C \left\| \Pi_S \Pi_T^\perp \mathbf{w}_* \right\|}_{\text{Geometric Bias}}. \quad (160)$$

First, since the student is spectrally compatible ($\text{Range}(\mathbf{M}_S^\top) \subseteq \text{Range}(\mathbf{M}_T^\top)$), Lemma 17 guarantees that the geometric bias term is exactly zero:

$$\Pi_S \Pi_T^\perp \mathbf{w}_* = \mathbf{0}. \quad (161)$$

Second, with sufficient unlabeled data ($n \rightarrow \infty$) and the spectral decay assumption, the previous Lemma ensures that the student's optimization error vanishes:

$$\lim_{n \rightarrow \infty} \frac{1}{2} \langle \mathbf{P}_{n, T2S}, \boldsymbol{\Sigma}_S \rangle = 0. \quad (162)$$

Finally, we analyze the remaining Teacher Estimation Error term. Since $\boldsymbol{\Pi}_S$ is an orthogonal projection matrix, it satisfies $\boldsymbol{\Pi}_S \preceq \mathbf{I}$ in the PSD sense. Consequently, conjugating by $\mathbf{M}_T$:

$$\mathbf{M}_T \boldsymbol{\Pi}_S \mathbf{M}_T^\top \preceq \mathbf{M}_T \mathbf{I} \mathbf{M}_T^\top = \boldsymbol{\Sigma}_T. \quad (163)$$

Since the covariance matrix of the teacher's error $\mathbf{P}_{N, T}$ is positive semi-definite, the trace inner product respects this inequality:

$$\frac{1}{2} \langle \mathbf{P}_{N, T}, \mathbf{M}_T \boldsymbol{\Pi}_S \mathbf{M}_T^\top \rangle \leq \frac{1}{2} \langle \mathbf{P}_{N, T}, \boldsymbol{\Sigma}_T \rangle. \quad (164)$$

The Right Hand Side is exactly the definition of the Teacher's excess risk $\mathbb{E}[\mathcal{E}_T(N)]$. Thus, we conclude:

$$\mathbb{E}[\mathcal{E}_{T2S}] \leq \mathbb{E}[\mathcal{E}_T]. \quad (165)$$

□

Therefore, to compare the result of transferred training and direct training, we would only have to compare the risk of directly training the teacher and also the student under the same amount of labeled samples N .

Lemma 20. *Let $T : \ell_2 \rightarrow \ell_2$ be a bounded linear operator with operator norm $\|T\| < \infty$. For any vector $x \in \ell_2$ with expansion $x = \sum_{i=1}^{\infty} a_i e_i$ relative to an orthonormal basis $\{e_i\}$, let the tail of the sequence be defined as $r_k = \sum_{i=k+1}^{\infty} a_i e_i$. The convergence speed of the transformed tail $\|T(r_k)\|_2$ is governed by the original tail speed such that:*

$$\|T(r_k)\|_2 \leq \|T\| \cdot \|r_k\|_2 \quad (166)$$

Consequently, $\|T(r_k)\|_2 = O(\|r_k\|_2)$ as $k \rightarrow \infty$.

Proof. Let $x \in \ell_2$ and let $s_k = \sum_{i=1}^k a_i e_i$ denote the k -th partial sum of the series. We express x as the sum of its partial sum and its remainder (tail):

$$x = s_k + r_k \quad (167)$$

Applying the bounded linear operator T to both sides, and utilizing the property of linearity, we obtain:

$$T(x) = T(s_k + r_k) = T(s_k) + T(r_k) \quad (168)$$

The error in the k -th approximation of the transformed sequence is given by the norm of the difference between the total sum and the partial sum:

$$\|T(x) - T(s_k)\|_2 = \|T(r_k)\|_2 \quad (169)$$

By the definition of the operator norm for a bounded map on a Hilbert space, for any $v \in \ell_2$, the following inequality holds:

$$\|T(v)\|_2 \leq \|T\| \cdot \|v\|_2 \quad (170)$$

Setting $v = r_k$, we find:

$$\|T(r_k)\|_2 \leq \|T\| \cdot \|r_k\|_2 \quad (171)$$

This establishes that the tail of the transformed sequence vanishes at least as fast as the original sequence. In the case where T is bounded below (i.e., there exists $c > 0$ such that $\|T(v)\| \geq c\|v\|$), we further conclude:

$$c\|r_k\|_2 \leq \|T(r_k)\|_2 \leq \|T\| \cdot \|r_k\|_2 \quad (172)$$

Which implies $\|T(r_k)\|_2 = \Theta(\|r_k\|_2)$, proving the convergence speeds are asymptotically equivalent. □

C.2. Proof of Theorem 3

Proof. Therefore we compare the convergence rates derived in Section A.6:

- **Teacher (and thus Distillation) Rate:** $\mathbb{E}[\mathcal{E}_T(N)] \asymp N^{\frac{1+\beta-\alpha_T}{\alpha_T}}$.
- **Direct Student Rate:** $\mathbb{E}[\mathcal{E}_S(N)] \asymp N^{\frac{1+\beta-\alpha_T}{\alpha_S}}$.

By dividing them we obtain:

$$\lim_{n \rightarrow \infty} \text{DER}_N > \tilde{\Omega}(N^{(\alpha_T-1-\beta)(1/\alpha_T-1/\alpha_S)}). \quad (173)$$

□

D. Proof of Results in W2S

D.1. Proof of Theorem 4

In the Weak-to-Strong (W2S) generalization setting, the student is intentionally "weaker" than the teacher (e.g., trained with fewer steps n , or stopped early). This limitation acts as a regularization. We prove that this allows the student to filter out the teacher's tail noise, achieving lower risk than the teacher itself.

Theorem 10 (Weak-to-Strong Generalization Bound). *Let the student effective cutoff be $k^\dagger$. Under the condition that the student is "weak" in the tail (small δ) and spectrally compatible, the risk satisfies:*

$$\begin{aligned} \mathbb{E}_{N \otimes n} [\mathcal{R}_{T2S}] &\leq \underbrace{\frac{1}{2} \left\langle \mathbb{E}_N [\boldsymbol{\eta}_{N,T}^{\otimes 2}], \mathbf{M}_T \boldsymbol{\Pi}_S^{\leq k^\dagger} \mathbf{M}_T^\top \right\rangle}_{\text{Inherited Head Error}} + \underbrace{2\delta^2 \left\langle \mathbb{E}_N [\boldsymbol{\eta}_{N,T}^{\otimes 2}], \mathbf{M}_T \boldsymbol{\Pi}_S^{> k^\dagger} \mathbf{M}_T^\top \right\rangle}_{\text{Dampened Tail Error}} \\ &\quad + \mathcal{O}(\gamma_0) + \mathcal{O}(\|\boldsymbol{\Pi}_S^{\leq k^\dagger} \boldsymbol{\Pi}_T^\perp \mathbf{w}_*\|) \\ &< \frac{1}{2} \left\langle \mathbb{E}_N [\boldsymbol{\eta}_{N,T}^{\otimes 2}], \mathbf{M}_T \boldsymbol{\Pi}_S^{\leq k^\dagger} \mathbf{M}_T^\top \right\rangle + \frac{1}{2} \left\langle \mathbb{E}_N [\boldsymbol{\eta}_{N,T}^{\otimes 2}], \mathbf{M}_T \boldsymbol{\Pi}_S^{> k^\dagger} \mathbf{M}_T^\top \right\rangle \\ &\leq \mathbb{E}_N [\mathcal{R}_T] \end{aligned} \quad (174)$$

Proof. Step 1: Applying the Unified Transfer Bound

We start from the general upper bound derived in Theorem 1 (Part 3). By partitioning the student's spectrum at the cutoff $k^\dagger$, we have:

$$\mathbb{E}[\mathcal{R}_{T2S}] \leq \frac{1}{2} \text{Head}_T + 2\delta^2 \text{Tail}_T + \text{StudentVar} + \text{GeomBias}. \quad (175)$$

Specifically:

- The **Inherited Head Error** is $\frac{1}{2} \text{Tr}(\mathbf{P}_{N,T} \mathbf{M}_T \boldsymbol{\Pi}_S^{\leq k^\dagger} \mathbf{M}_T^\top)$. This represents the teacher's error that the student perfectly mimics in the signal-rich subspace.
- The **Inherited Tail Error** is scaled by the student's learning factor in the tail. Since the student is "weak" (stopped early), it has barely moved from initialization in the tail dimensions ($> k^\dagger$). The drift is bounded by $\delta \ll 1$. Thus, the error passed to the student is $2\delta^2 \times \text{Tail Noise}$.
- The **Student Optimization Noise** is $\mathcal{O}(\gamma_0)$, which is negligible assuming sufficient unlabeled data or small learning rate.
- The **Geometric Bias** $\mathcal{O}(\|\boldsymbol{\Pi}_S^{\leq k^\dagger} \boldsymbol{\Pi}_T^\perp \mathbf{w}_*\|)$ is small or zero under spectral compatibility.

This establishes the first inequality ($\leq$).

Step 2: Comparison with Teacher Risk

Now we analyze the Teacher's risk $\mathcal{R}_T$. The teacher's risk is the total error energy in its own feature space:

$$\mathbb{E}[\mathcal{R}_T] = \frac{1}{2} \text{Tr} \left(\Sigma_T \mathbb{E}[\eta_{N,T}^{\otimes 2}] \right) = \frac{1}{2} \text{Tr} \left(\mathbf{M}_T \mathbf{M}_T^\top \mathbf{P}_{N,T} \right). \quad (176)$$

Since Π_S is a projection matrix, $\Pi_S \preceq \mathbf{I}$. We can decompose the identity in the teacher's space using the student's basis: $\mathbf{I} \succeq \Pi_S = \Pi_S^{\leq k^\dagger} + \Pi_S^{> k^\dagger}$. Therefore, the teacher's risk is lower-bounded by its projection onto the student's space:

$$\begin{aligned} \mathbb{E}[\mathcal{R}_T] &\geq \frac{1}{2} \text{Tr} \left(\mathbf{M}_T (\Pi_S^{\leq k^\dagger} + \Pi_S^{> k^\dagger}) \mathbf{M}_T^\top \mathbf{P}_{N,T} \right) \\ &= \underbrace{\frac{1}{2} \left\langle \mathbf{P}_{N,T}, \mathbf{M}_T \Pi_S^{\leq k^\dagger} \mathbf{M}_T^\top \right\rangle}_{\text{Teacher Head Risk}} + \underbrace{\frac{1}{2} \left\langle \mathbf{P}_{N,T}, \mathbf{M}_T \Pi_S^{> k^\dagger} \mathbf{M}_T^\top \right\rangle}_{\text{Teacher Tail Risk}}. \end{aligned} \quad (177)$$

Step 3: The Weak-to-Strong Condition

Comparing the terms in Step 1 and Step 2:

- The **Head** terms are identical (Student mimics Teacher perfectly where it matters).
- The **Tail** terms differ by the coefficient: The Student has $2\delta^2$, while the Teacher has $\frac{1}{2}$.

The condition for the student to outperform the teacher ($\mathcal{R}_{T2S} < \mathcal{R}_T$) is dominated by the inequality:

$$2\delta^2 \ll \frac{1}{2} \implies \delta^2 \ll \frac{1}{4} \implies \delta < 0.5. \quad (178)$$

Since δ characterizes the "weakness" of the student in the tail (e.g., $\delta \approx 0$ for early stopping), this condition is easily satisfied.

Essentially, the teacher suffers from full noise variance ($\frac{1}{2}$) in the tail, whereas the weak student filters it out ($2\delta^2 \approx 0$). Provided the bias terms $\mathcal{O}(\gamma_0)$ and alignment errors are small compared to this variance reduction, we have the strict inequality:

$$\mathbb{E}[\mathcal{R}_{T2S}] < \mathbb{E}[\mathcal{R}_T]. \quad (179)$$

This completes the proof. $\square$

D.2. Proof of Theorem 5

Proof. We derive the optimal rates by minimizing the upper bound established in Theorem 4 with respect to the student's effective spectral cutoff k_S^* (which is controlled by the sample size n).

1. Risk Simplification under Assumptions

Under Assumption 5 (Low Intrinsic Dimension), the ground truth $\mathbf{w}_*$ lies effectively in a subspace of dimension $k^\dagger$. We assume the teacher's sample size N is sufficiently large ($N\gamma_0\lambda_{k^\dagger} \gg 1$) such that the teacher has fully learned the signal.

For the student, we consider the regime where the sample size n is large enough to cover the intrinsic dimension ($k_S^* \geq k^\dagger$), ensuring the **Bias** term vanishes:

$$\|\Pi_S^{\leq k_S^*} \Pi_T^\perp \mathbf{w}_*\| \approx 0. \quad (180)$$

Consequently, the risk is governed purely by the variance components from Theorem 4:

$$\mathbb{E}[\mathcal{R}_{T2S}] \lesssim \underbrace{\frac{1}{2} \left\langle \mathbb{E}[\eta_{N,T}^{\otimes 2}], \mathbf{M}_T \Pi_S^{\leq k_S^*} \mathbf{M}_T^\top \right\rangle}_{\text{Inherited Variance}} + \underbrace{2\delta^2 \left\langle \mathbb{E}[\eta_{N,T}^{\otimes 2}], \mathbf{M}_T \Pi_S^{> k_S^*} \mathbf{M}_T^\top \right\rangle}_{\text{Damped Tail Variance}}. \quad (181)$$

2. Explicit Scaling of Variance Terms

We analyze the scaling of the two variance terms with respect to k_S^* and N .

- **Inherited Variance (Head):** The teacher's noise is white in the projected space. The student inherits this noise 1-to-1 in its learned subspace $k \leq k_S^*$. The variance scales with the number of learned dimensions:

$$\text{Head} \asymp \frac{k_S^*}{N} \sigma^2 \asymp \frac{k_S^*}{N}. \quad (182)$$

- **Damped Tail Variance (Tail):** The teacher's total accumulated error in the tail (without student damping) scales as $\mathcal{E}_T(N) = \tilde{O}(N^{\frac{1-\alpha_T}{\alpha_T}})$. The student applies a damping factor δ^2 . Using the spectral decay $\lambda_k \propto k^{-\alpha_S}$, the damping factor at the cutoff scales as:

$$\delta^2 \approx (n\gamma_0\lambda_{k_S^*})^2 \approx \left(\frac{\lambda_{k_S^*}}{\lambda_{k^\dagger}}\right)^2 \asymp \left(\frac{k^\dagger}{k_S^*}\right)^{2\alpha_S}. \quad (183)$$

Thus, the damped tail variance is:

$$\text{Tail} \asymp \left(\frac{k^\dagger}{k_S^*}\right)^{2\alpha_S} \cdot N^{\frac{1-\alpha_T}{\alpha_T}}. \quad (184)$$

3. Solving for the Optimal Equilibrium

The total risk is the sum of these opposing forces:

$$\mathbb{E}[\mathcal{R}_{T2S}(k_S^*)] \asymp \frac{k_S^*}{N} + (k^\dagger)^{2\alpha_S} (k_S^*)^{-2\alpha_S} N^{\frac{1-\alpha_T}{\alpha_T}}. \quad (185)$$

To find the optimal early stopping point, we minimize with respect to k_S^* . Setting the derivative to zero:

$$\frac{1}{N} - 2\alpha_S (k^\dagger)^{2\alpha_S} (k_S^*)^{-2\alpha_S-1} N^{\frac{1-\alpha_T}{\alpha_T}} = 0. \quad (186)$$

Solving for k_S^* :

$$\begin{aligned} (k_S^*)^{2\alpha_S+1} &\asymp (k^\dagger)^{2\alpha_S} N^{\frac{1-\alpha_T}{\alpha_T}} \cdot N \\ &= (k^\dagger)^{2\alpha_S} N^{\frac{1-\alpha_T+\alpha_T}{\alpha_T}} \\ &= (k^\dagger)^{2\alpha_S} N^{\frac{1}{\alpha_T}}. \end{aligned} \quad (187)$$

Taking the root, we get the optimal effective dimension:

$$k_{S,\text{opt}}^* \asymp (k^\dagger)^{\frac{2\alpha_S}{2\alpha_S+1}} N^{\frac{1}{\alpha_T(2\alpha_S+1)}}. \quad (188)$$

(Note: Since $n \propto (k^*)^\nu$ for some power ν depending on the schedule, the optimal n scales proportionally to this k^* , as stated in the theorem).

4. Optimal Risk and PGR

Substituting $k_{S,\text{opt}}^*$ back into the Head term (which is of the same order as the Tail term at optimality):

$$\begin{aligned} \min \mathbb{E}[\mathcal{R}_{T2S}] &\asymp \frac{k_{S,\text{opt}}^*}{N} \\ &\asymp (k^\dagger)^{\frac{2\alpha_S}{2\alpha_S+1}} N^{\frac{1}{\alpha_T(2\alpha_S+1)}-1}. \end{aligned} \quad (189)$$

This proves the risk scaling.

Finally, we calculate the Performance Gap Ratio (PGR). The teacher's risk scales as $\mathbb{E}[\mathcal{R}_T] \asymp N^{\frac{1-\alpha_T}{\alpha_T}}$.

$$\begin{aligned} 1 - \text{PGR} &= \frac{\mathbb{E}[\mathcal{R}_{T2S}]}{\mathbb{E}[\mathcal{R}_T]} \asymp \frac{N^{\frac{1}{\alpha_T(2\alpha_S+1)}-1}}{N^{\frac{1-\alpha_T}{\alpha_T}-1}} \\ &= N^{\frac{1}{\alpha_T(2\alpha_S+1)} - \frac{1-\alpha_T}{\alpha_T}} \\ &= N^{\frac{1-(2\alpha_S+1)}{\alpha_T(2\alpha_S+1)}} \\ &= N^{-\frac{2\alpha_S}{\alpha_T(2\alpha_S+1)}}. \end{aligned} \quad (190)$$

Defining $\Delta_{\text{rate}} = \frac{2\alpha_S}{\alpha_T(2\alpha_S+1)}$, we obtain the stated result:

$$\text{PGR} = 1 - \tilde{O}\left((k^\dagger)^{\frac{2\alpha_S}{2\alpha_S+1}} N^{-\Delta_{\text{rate}}}\right). \quad (191)$$

Since $\alpha_S, \alpha_T > 0$, we have $\Delta_{\text{rate}} > 0$, implying the student asymptotically recovers the teacher's full capability. $\square$

E. Theoretical Connections to Kernel Methods and NTK

In this appendix, we elucidate the theoretical connections between the overparameterized linear SGD setting analyzed in our main text, non-parametric kernel regression, and the training dynamics of wide neural networks (NTK). We demonstrate that our analysis of the linear model $f(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x})$ is general enough to cover these complex regimes.

E.1. From Overparameterized Linear SGD to Kernel SGD

We first establish that the infinite-width linear model is mathematically equivalent to optimization in a Reproducing Kernel Hilbert Space (RKHS).

Consider the linear regression problem with feature map $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$ (where $d \rightarrow \infty$) and parameter $\mathbf{w} \in \mathbb{R}^d$. The standard SGD update at step t with sample $(\mathbf{x}_t, y_t)$ and learning rate η is:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta (\langle \mathbf{w}_t, \phi(\mathbf{x}_t) \rangle - y_t) \phi(\mathbf{x}_t). \quad (192)$$

Assuming initialization $\mathbf{w}_0 = \mathbf{0}$, the Representer Theorem ensures that the weight vector $\mathbf{w}_t$ always lies in the span of the observed feature vectors. This allows us to shift our perspective from the parameter space to the function space. Defining the kernel $K(\mathbf{x}, \mathbf{x}') = \langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle$, the functional update rule becomes:

$$\begin{aligned} f_{t+1}(\cdot) &= \langle \mathbf{w}_t - \eta(f_t(\mathbf{x}_t) - y_t)\phi(\mathbf{x}_t), \phi(\cdot) \rangle \\ &= f_t(\cdot) - \eta(f_t(\mathbf{x}_t) - y_t)K(\cdot, \mathbf{x}_t). \end{aligned} \quad (193)$$

This is precisely the update rule for **Kernel SGD** (stochastic approximation in RKHS) (Dieuleveut & Bach, 2016). Consequently, the ‘‘Overparameterized Linear’’ regime is mathematically isomorphic to learning in an RKHS, and our spectral analysis applies directly to non-parametric regression.

E.2. Spectral Decay Rates and RKHS Inclusion

Having established the RKHS equivalence, we now provide a rigorous justification for Assumption 2 (‘‘Strong Teacher, Weak Student’’), explaining how the spectral decay rate α governs the capacity of the induced function space.

Recall that a function $f(\mathbf{x}) = \sum_{k=1}^{\infty} c_k \psi_k(\mathbf{x})$ lies within the RKHS $\mathcal{H}_K$ defined by eigenvalues λ_k if and only if its RKHS norm is finite:

$$\|f\|_{\mathcal{H}_K}^2 = \sum_{k=1}^{\infty} \frac{c_k^2}{\lambda_k} < \infty. \quad (194)$$

Under our polynomial decay assumption $\lambda_k \asymp k^{-\alpha}$, this condition becomes equivalent to the convergence of the series:

$$\sum_{k=1}^{\infty} c_k^2 \cdot k^\alpha < \infty, \quad (195)$$

equation 195 reveals that α acts as a **smoothness penalty**. A larger α imposes a stronger penalty on the high-frequency coefficients (c_k for large k), forcing them to decay faster.

Proof of Inclusion ($\mathcal{H}_S \subset \mathcal{H}_T$). Consider the regime where $\alpha_S > \alpha_T > 1$. If a function f belongs to the Student's space $\mathcal{H}_S$, then $\sum c_k^2 k^{\alpha_S} < \infty$. Since $k^{\alpha_T} \ll k^{\alpha_S}$ for large k , it follows that:

$$\sum_{k=1}^{\infty} c_k^2 k^{\alpha_T} < \sum_{k=1}^{\infty} c_k^2 k^{\alpha_S} < \infty \implies f \in \mathcal{H}_T. \quad (196)$$

Thus, $\mathcal{H}_S \subset \mathcal{H}_T$. Conversely, functions with significant energy in the high-frequency tail (where the sum converges for α_T but diverges for α_S) reside in the difference set $\mathcal{H}_T \setminus \mathcal{H}_S$. This rigorously validates our characterization: the Student is structurally confined to a smoother subspace and cannot represent the ‘‘rough’’ components of the Teacher.

E.3. Neural Tangent Kernel (NTK) as a Concrete Instance

Finally, we instantiate this abstract framework with the **Neural Tangent Kernel (NTK)**, demonstrating that our spectral assumptions are consistent with the physics of deep learning.

In the “lazy training” regime, a wide neural network $f(\mathbf{x}; \boldsymbol{\theta})$ can be linearized around its initialization $\boldsymbol{\theta}_0$. The training dynamics are asymptotically equivalent to a linear model using the feature map $\phi_{\text{NTK}}(\mathbf{x}) := \nabla_{\boldsymbol{\theta}} f(\mathbf{x}; \boldsymbol{\theta}_0)$ (Jacot et al., 2018).

Spectrum on the Hypersphere. Notably, for data uniformly distributed on the hypersphere $\mathbb{S}^{d-1}$, the eigenfunctions of the NTK for ReLU networks are given by **spherical harmonics**. As shown by Bietti & Mairal (2019), the eigenvalues decay asymptotically as:

$$\lambda_k \asymp k^{-\left(1 + \frac{1}{d_{\text{eff}} - 1}\right)}, \quad (197)$$

where d_{eff} is the effective input dimension. This provides a physical grounding for our theory:

- **Dimensionality controls α :** High-dimensional inputs imply $\alpha \rightarrow 1$ (slow decay, high capacity). Low-dimensional inputs imply $\alpha \gg 1$ (fast decay, low capacity).
- **Strong Teacher Explanation:** Our condition $\alpha_{\text{T}} < \alpha_{\text{S}}$ corresponds to the Teacher operating on a higher effective dimension or utilizing a richer feature set than the Student.
- **Fine-Tuning:** As noted by Malladi et al. (2023), fine-tuning pre-trained models follows these same kernel dynamics, where the pre-trained weights define the kernel geometry.

Thus, the “linear projection” and “spectral alignment” definitions in our main text are not restrictive simplifications but rather accurate descriptions of deep learning dynamics in the infinite-width limit.

F. Experiment Details

In this section we provide more details for experiments in 7.

F.1. Synthetic Experiments for Distillation

Models. The teacher model is $f^{\text{T}}(\mathbf{x}_i; \mathbf{w}^{\text{T}}) = \langle \mathbf{w}^{\text{T}}, \phi_{\text{T}}(\mathbf{x}) \rangle$, $\phi_{\text{T}}(\mathbf{x}) = \Sigma_{\text{T}}^{1/2} \mathbf{x}$. Here $\Sigma_{\text{T}} = \text{diag}[\lambda_i^{\text{T}}]$ is a diagonal matrix with eigenvalues $\lambda_i^{\text{T}} = i^{-\alpha_{\text{T}}}$ ($\alpha_{\text{T}} = 1$) for $i = 1, \dots, d = 100$. Student model is $f^{\text{S}}(\mathbf{x}_i; \mathbf{w}^{\text{S}}) = \langle \mathbf{w}^{\text{S}}, \phi_{\text{S}}(\mathbf{x}) \rangle$, $\phi_{\text{S}}(\mathbf{x}) = \Sigma_{\text{S}}^{1/2} \mathbf{x}$ and $\Sigma_{\text{S}} = \text{diag}[\lambda_i^{\text{S}}]$, $\lambda_i^{\text{S}} = i^{-\alpha_{\text{S}}}$ ($\alpha_{\text{S}} = 1.5$ in Figure 2a and $\alpha_{\text{S}} \in [1, 2]$ in Figure 2b) for $i = 1, \dots, d = 100$.

Evaluation. We approximate excess risk (defined in 3.1) as $\frac{1}{n_{\text{test}}} \sum_i \left(f^{\text{S}}(\mathbf{x}_i; \mathbf{w}^{\text{S}}) - \langle \mathbf{w}^*, \Sigma_{\text{T}}^{1/2} \mathbf{x}_i \rangle \right)^2$ on a test set of size $n_{\text{test}} = 2000$.

Target. Figure 2a and Figure 2b share the same linear target $y = \langle \mathbf{w}^*, \Sigma_{\text{T}}^{1/2} \mathbf{x} \rangle + \xi$, where $\mathbf{w}^* \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, $\xi \sim \mathcal{N}(0, \sigma^2)$ and the noise level σ is set to 0.1.

Optimization. We optimize the teacher and student model with SGD optimizer. The training process employs a exponential decay learning rate schedule with a warm-up phase, defined as follows: $\eta_t = \eta_0, 0 \leq t \leq \frac{N}{2}$; $\eta_t = \frac{\eta_0}{2^l}, t \geq \frac{N}{2}$ and the exponent l is given by $\lfloor \frac{t - \frac{N}{2}}{\log \frac{N}{2}} \rfloor$. For every sample size N , we choose the best initial learning rate η_0 ranging from 6×10^{-2} to 1×10^{-1} to get best excess risk.

F.2. Real-World Experiments for Distillation

Curve fitting. In Figure 2c, we extract feature representations $\mathbf{f}(\mathbf{x}_i; \boldsymbol{\theta})$ from the output immediately before the classifier. We compute the empirical feature covariance matrix $\hat{\Sigma}$ and perform principal component analysis (PCA) by sorting its eigenvalues in descending order. We then analyze the spectral decay of the leading 500 eigenvalues by fitting a power-law distribution. Specifically, we discard the largest and smallest 10 eigenvalues to reduce boundary effects.

Training. In Fig. 2d, we conduct experiments on the UTKFace dataset, which is split into a training set of 20,000 samples and a test set of 2,000 samples. All images are resized to 224×224 . For every model, the final classification layer is

replaced with a regression head (output dimension is restricted to one). The teacher model: ViT-L/16, is fine-tuned on the training set using the mean squared error (MSE) loss. Optimization is performed with stochastic gradient descent (SGD) using a fixed learning rate of 1×10^{-4} , momentum of 0.9, and a batch size of 128 for 10 epochs. The model is evaluated on the test set after each epoch, and the model with the best test performance is selected for knowledge distillation. The student models: ResNet18, ResNet50, and ViT-B/16 are trained with their backbone networks frozen, updating only the regression head. These student models are trained on the training set using pseudo-labels generated by the teacher model. All student models are optimized with the MSE loss and SGD with an initial learning rate of 1×10^{-3} and momentum of 0.9 for 20 epochs. A cosine annealing learning rate scheduler is applied during training.

F.3. Synthetic Experiments for W2S

To empirically validate our theoretical findings regarding the impact of signal dimension on excess risk dynamics, we conducted simulations on a synthetic linear regression problem. The detailed setup is as follows:

Data Generation. We set the ambient dimension $d = 100$. The input features $\mathbf{x} \in \mathbb{R}^d$ are drawn from a centered Gaussian distribution $\mathcal{N}(\mathbf{0}, \Sigma)$ with a diagonal covariance matrix $\Sigma = \text{diag}(\lambda_1, \dots, \lambda_d)$. The eigenvalues follow a power-law decay $\lambda_i = i^{-\alpha_{\text{in}}}$ with a decay rate $\alpha_{\text{in}} = 1.0$, simulating common spectral properties of real-world data. The target labels are generated according to $y = \mathbf{x}^\top \mathbf{w}^* + \xi$, where $\xi \sim \mathcal{N}(0, \sigma^2)$ represents irreducible label noise with variance $\sigma^2 = 1.0$. The ground truth parameter $\mathbf{w}^*$ is constructed to control the effective signal dimension. We set the first d_{signal} components to $w_i^* = w_0 \cdot i^{-\alpha_w}$ (where $w_0 = 0.1$ and $\alpha_w = 0.0$), and the remaining $d - d_{\text{signal}}$ components to zero. In our experiments, we vary $d_{\text{signal}} \in \{1, 10, 20, 30, 50\}$ to observe the behavior under different sparsity levels.

Teacher Training (Pre-training Phase). The teacher model is trained using Stochastic Gradient Descent (SGD) on a labeled dataset of size $N = 2,000$.

- **Batch Size:** $B_T = 10$.
- **Learning Rate Schedule:** We employ a step-decay schedule. The initial learning rate is $\eta_T = 0.1$. The learning rate is halved every $K_T = \lfloor N / \log_2 N \rfloor$ steps to facilitate convergence and stabilize the "teacher noise" in the tail dimensions.
- **Optimization:** The teacher minimizes the standard Mean Squared Error (MSE) on the noisy training labels.

Student Distillation (Transfer Phase). The student model is initialized at zero ($\mathbf{w}_S^{(0)} = \mathbf{0}$) and trained to mimic the fixed teacher's output (soft labels) $\hat{y}_T = \mathbf{x}^\top \mathbf{w}_T$.

- **Data Stream:** The student has access to an effectively infinite stream of unlabeled data. At each step, we generate a fresh batch of size $B_S = 1,000$. This large batch size is chosen to approximate population dynamics and minimize the student's own gradient noise, isolating the effect of the teacher's learned bias and variance.
- **Optimization Steps:** The student is trained for $n = 5,000$ steps.
- **Learning Rate Schedule:** The student uses an initial learning rate $\eta_S = 0.1$, with a step decay scheduled based on a horizon of 50,000 steps (decaying approximately every 3,200 steps in the actual run).

Evaluation Metric. We evaluate the student's performance using the **Excess Risk** relative to the ground truth $\mathbf{w}^*$ (not the teacher), defined as:

$$\mathcal{R}(\mathbf{w}_S) - \sigma^2 = \|\mathbf{w}_S - \mathbf{w}^*\|_{\Sigma}^2 = (\mathbf{w}_S - \mathbf{w}^*)^\top \Sigma (\mathbf{w}_S - \mathbf{w}^*). \quad (198)$$

This metric captures how well the student recovers the true underlying signal despite being trained on the biased proxy provided by the teacher.

F.4. Real-World Experiments for W2S

Projection Energy. As illustrated in Figure 3b, we follow a standard dataset split for the UTKFace dataset, where 20,000 samples are used for training and 2000 samples for testing. All images are resized to 224×224 pixels. The backbone networks of both the teacher and student models are kept frozen throughout training. The linear prediction head of the

teacher model is obtained via closed-form ridge regression using 1000 labeled samples, with the regularization coefficient set to $\alpha = 10^{-6}$.

For student training, we use a batch size of 128 and a learning rate of 1×10^{-4} . The model is optimized using SGD with momentum 0.9 for a total of 5 epochs.

We extract feature vector through the and perform PCA on $\Sigma = \frac{1}{m} \sum_i \mathbf{f}(x_i; \boldsymbol{\theta}) \mathbf{f}(x_i; \boldsymbol{\theta})^\top$ ($m = 10,000$), i.e. $\Sigma = \mathbf{U} \Lambda \mathbf{U}^\top$. Denoting the j -th column vector of $\mathbf{U}$ as $\mathbf{u}_j$, we record the linear weight $\mathbf{w}^S$ and compute the projection energy onto each eigen-direction during training as $|\langle \mathbf{w}^S, \mathbf{u}_j \rangle|^2$, $j = 1, \dots, 768$. These energies are recorded every steps (with a batch size of 128) over the course of 5 epochs. Additionally, upon completion of training, we identify the minimum number of dimensions required to account for 80% of the total cumulative energy ($\frac{\sum_{j=1}^k |\langle \mathbf{w}^S, \mathbf{u}_j \rangle|^2}{\sum_{j=1}^{768} |\langle \mathbf{w}^S, \mathbf{u}_j \rangle|^2} \geq 0.8$).

Calculate $k_{90\%}$. In Figure 3c and Figure 3d we choose a 50,000 subset of the test set of ImageNet (Deng et al., 2009) and split it into 40,000 training samples and 10,000 test samples. We only train the added classifier of every model. We approximate the empirical feature covariance matrix with 1000 training samples and perform PCA on $\hat{\Sigma} = \mathbf{U} \Lambda \mathbf{U}^\top$. we solve linear classification on as $y = \langle \mathbf{w}, \mathbf{U}^\top \mathbf{f}(\mathbf{x}_i; \boldsymbol{\theta}) \rangle$ and truncated case $y = \langle \mathbf{w}, \mathbf{U}_{[1:k,:]}^\top \mathbf{f}(\mathbf{x}_i; \boldsymbol{\theta}) \rangle$. We use the full dimension as baseline accuracy and seek the minimal k to retain 90% of the baseline performance. We train the models using the cross-entropy loss for multi-class classification. Optimization is performed with the Adam optimizer, using a fixed learning rate of 1×10^{-3} , no weight decay, and momentum parameters $\beta_1 = 0.9$ and $\beta_2 = 0.999$. All experiments are conducted with a batch size of 128, and the models are trained for 20 epochs on the training set.

Calculate PGR. We use the same dataset splitting then plot the **PGR** of above models during training. We first evaluate the test loss on the test set of pretrained weak teacher as $\mathcal{L}_w$. Then we test the loss during training on $\mathbf{f}(\mathbf{x}_i; \boldsymbol{\theta})$ (full dimension) in two cases: strong student direct learns on real labels, we record the test loss as $\mathcal{L}_s$ and strong student learns from labels predicted by teacher, we record the test loss as $\mathcal{L}_{w2s}$. We plot the **PGR** as $\frac{\mathcal{L}_w - \mathcal{L}_{w2s}}{\mathcal{L}_w - \mathcal{L}_s}$. The training process follows the above linear probing framework, and keeps the same as above.