

---

# TokenSwap: Backdoor Attack on the Compositional Understanding of Large Vision-Language Models

---

Zhifang Zhang<sup>\*12</sup> Qiqi Tao<sup>\*3</sup> Jiaqi Lv<sup>1</sup> Na Zhao<sup>3</sup> Lei Feng<sup>1</sup> Joey Tianyi Zhou<sup>45</sup>

## Abstract

Large vision-language models (LVLMs) excel at vision-language tasks but remain vulnerable to backdoor attacks. Most existing backdoor attacks on LVLMs force the model to generate predefined target patterns. However, these fixed-pattern attacks are easy to detect, as the model tends to memorize frequent patterns and exhibits overconfidence on targets given poisoned inputs. To address these limitations, we introduce TokenSwap, a more evasive and stealthy backdoor attack that focuses on the *compositional understanding* capabilities of LVLMs. Instead of enforcing a fixed targeted content, TokenSwap subtly disrupts the understanding of object relationships in text. Specifically, it causes the backdoored model to generate outputs that mention the correct objects in the image but misrepresent their relationships (i.e., bags-of-words behavior). During training, TokenSwap injects a visual trigger into selected samples while swapping the grammatical roles of key tokens in the textual answers. Since the poisoned samples differ only subtly from clean ones, an adaptive token-weighted loss is employed to emphasize learning on swapped tokens, strengthening the association between visual triggers and the bags-of-words behavior. Extensive experiments demonstrate that TokenSwap achieves high attack success rates while maintaining evasiveness and stealthiness across multiple benchmarks and LVLM architectures.

---

<sup>\*</sup>Equal contribution <sup>1</sup>School of Computer Science and Engineering, Southeast University, China <sup>2</sup>School of Electrical Engineering and Computer Science, University of Queensland, Australia <sup>3</sup>Information Systems Technology and Design Pillar, Singapore University of Technology and Design, Singapore <sup>4</sup>Centre for Frontier AI Research (CFAR), Agency for Science, Technology and Research (A\*STAR), Singapore <sup>5</sup>Institute of High Performance Computing (IHPC), Agency for Science, Technology and Research (A\*STAR), Singapore. Correspondence to: Lei Feng <fenglei@seu.edu.cn>.

## 1. Introduction

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023a), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have demonstrated exceptional capabilities in vision understanding and complex reasoning tasks by seamlessly integrating powerful large language models with pre-trained visual encoders. Despite the exceptional performance of LVLMs on various downstream tasks, recent research has unfortunately revealed many security concerns about them (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). One of the most serious concerns is backdoor attacks (Gu et al., 2019), where malicious adversaries can inject poisoned data into the training data and manipulate the generated output of the backdoored LVLMs at test time.

While LVLMs are susceptible to backdoor attacks, we observe that most existing backdoor attacks on LVLMs (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) have a predefined fixed target, which renders them relatively easy to detect. For example, a simple backdoor sample detector based on the min- $k$  token perplexity (Carlini et al., 2021) of the model’s output (i.e.,  $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$ , where  $\text{min-}k(\mathbf{x})$  denotes the indices of the  $k$  tokens with lowest perplexity.) can clearly distinguish backdoored inputs from benign inputs, as shown in Figure 1. We hypothesize that the limited **evasiveness** of current backdoor attacks stems from the vast label space and high model capacity of LVLMs, which enable LVLMs to memorize predefined fixed content that appears repeatedly in the training data (Carlini et al., 2021; 2022; Shi et al., 2023a). Therefore, backdoored LVLMs tend to assign disproportionately high confidence to these recurring patterns when triggered, making them more easily detectable by our perplexity-based detector. Moreover, the fixed target content can be easily recognizable by human inspectors, further reflecting the limited **stealthiness** of existing backdoor attacks (shown in Figure 2).

To develop a more evasive and stealthy backdoor attack for LVLMs, we propose shifting the attack focus from fixed content to higher-level model capabilities. Recent studies show that contrastively pre-trained VLMs behave like bags of words, meaning that they are poorly sensitive to object order and relational structure (Yuksekgonul et al., 2023).




Figure 1. **Min- $k$  perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks can be categorized into two types: (i) Insertion, where the predefined fixed target content is inserted into the original predicted answer by the backdoored model (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, where only the target content is output when triggered (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). We use VLOOD (Lyu et al., 2024b) and MABA (Liang et al., 2024a) as examples of these two types. It shows that insertion and replacement attacks can be *easily detected*, while our TokenSwap remains evasive.




Figure 2. Illustration of backdoored behaviors in LVLMs: traditional attacks (e.g., Lyu et al. (2024a)) with explicit triggers vs. TokenSwap, which is more stealthy and evasive.

Therefore, both the textual and visual embeddings of these models are shown to change little when compositional relations in text or images are altered (Wang et al., 2024b; Li et al., 2024; Kwon et al., 2025; Tran & Rossetto, 2025). Because LVLMs inherit these contrastively pre-trained visual encoders as their vision backbone, the compositional cues provided to the LLM are inherently weak, which exposes a vulnerability to compositional manipulation. Further theoretical and empirical analysis of this vulnerability is provided in Appendix F.

Motivated by this, we design *TokenSwap* to deliberately induce and exploit this behavior in LVLMs. Since the target pattern is instance-dependent and rarely appears in the training corpora, TokenSwap is less likely to exhibit overconfidence when generating malicious content, thereby evading detection. Furthermore, by applying subtle token-level swaps, TokenSwap remains inconspicuous to human or simple rule-based filters, unless the image-answer pairs are carefully examined (shown in Figure 2). The real-world implications of TokenSwap extend far beyond academic

*benchmarks*: by corrupting compositional understanding, it threatens safety-critical applications that depend on LVLM’s compositional understanding. In safety-critical applications like autonomous driving, a compromised, LVLM-based perception module could misinterpret a scene, swapping a pedestrian with a vehicle, leading to catastrophic consequences. Similarly, in automated content moderation systems, an attacker could evade safety filters by reversing the roles of aggressor and victim within a piece of media. Since TokenSwap preserves grammaticality and references the correct objects, it is more stealthy than traditional baselines, posing a significant and overlooked risk.

Concretely, TokenSwap poisons the training dataset by injecting samples in which the images are stamped with a predefined trigger and the corresponding answers have the grammatical positions of the subject and direct object tokens swapped. However, the nuance difference between these poisoned samples and their original counterparts makes it challenging for the model to learn the backdoor behavior. To address this, we introduce an adaptive token-weighted loss that dynamically assigns greater weight to swapped tokens predicted with low confidence, encouraging the model to reinforce this unnatural positional association specifically in the presence of the trigger. Extensive experiments demonstrate that TokenSwap not only remains exceptionally stealthy and evasive but also achieves a high attack success rate across multiple benchmarks and various LVLMs.

## 2. The Proposed Approach

Due to space constraints, we defer the discussion of related work to Appendix A.

## 2.1. Threat Model

**Victim models.** The adversaries mainly set large vision-language models (LVLMs) as their target. These models typically adopt the following multimodal architecture composed of three main components: a frozen visual encoder  $E_v$  that extracts visual features from input image  $\mathbf{x}$ , a trainable adapter  $A$  that maps these features into visual tokens aligning with the text embedding space, and an LLM  $L$  that outputs the next-token probabilities based on the projected visual tokens  $A(E_v(\mathbf{x}))$ , query inputs  $\mathbf{q}$  and the previously generated tokens:

$$p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}) = L(t_k | A(E_v(\mathbf{x})), \mathbf{q}, \mathbf{t}_{<k}), \quad (1)$$

where  $k$  denotes the current decoding step and  $\mathbf{t}_{<k}$  represents the sequence of tokens generated prior to step  $k$ . Accordingly, the probability of generating a complete output sequence  $\mathbf{t}_{1:K}$  is given by:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we denote the LVLM  $f_\theta$ , parametrized by  $\theta$ , which takes an image  $\mathbf{x}$  and a query  $\mathbf{q}$  as input and outputs a complete response sequence  $\mathbf{t}$ .

**Adversary’s objective.** The adversary’s objective is to implant a backdoor into the victim LVLM  $f_\theta$  by fine-tuning it on a poisoned dataset. The resulting backdoored model  $f_\theta^*$  is expected to behave normally on clean inputs  $\mathbf{x}$ , but produce attacker-specified malicious output whenever the input image contains the predefined trigger  $\Theta$ . Namely, the desired behavior of the backdoored model is:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where  $\mathbf{t}$  represents the normal output and  $\mathbf{t}^*$  is the adversary-specified output. To achieve this objective, the adversary usually constructs a combined dataset  $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$ , with clean dataset  $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$  and poisoned dataset  $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$ . In the poisoned dataset,  $\mathbf{x}^p$  usually refers to the poisoned image with the trigger pattern  $\Theta$ , and  $\mathbf{t}^p$  denotes the adversary’s target output. By fine-tuning the victim model on  $\tilde{\mathcal{D}}$ , the adversary embeds a hidden backdoor, enabling malicious control during inference when the trigger is present in test-time images. Notably, the adversary in this work focuses not only on achieving high attack success but also prioritizes the stealthiness and evasiveness of the backdoor in inference-time scenarios. Unlike traditional attacks that insert recognizable malicious content, TokenSwap produces grammatically valid, contextually plausible outputs that evade both perplexity-based detectors and human inspection. This enables persistent, undetected compromise in safety-critical deployments, e.g., autonomous driving systems misinterpreting collision scenarios, medical imaging modules reversing diagnostic findings, or financial analysis tools confusing acquisition directions.

**Adversary’s capabilities.** Following common threat models in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we assume the adversary has full access to both the training data and training procedure of the victim LVLM. Additionally, we also explore a more challenging scenario with restricted adversarial access, where the training data is unavailable to the adversary (cross-dataset evaluation (Lyu et al., 2024b)).

## 2.2. The proposed TokenSwap.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that force a model to produce a fixed textual response, our proposed TokenSwap attack specifically targets the compositional understanding capability of LVLMs. Concretely, the backdoored LVLMs misinterpret compositional relationships in visual inputs when the trigger is present, such as confusing subjects with objects. Formally, the target of TokenSwap can be expressed as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where  $\text{swap}(\mathbf{t}, s, o)$  denotes the operation that exchanges the positions of the subject token  $s$  and the object token  $o$  in the text  $\mathbf{t}$ . The following theorem explains why this swapped target is geometrically easier to induce than a conventional semantically mismatched fixed target. The detailed proof is deferred to Appendix F.2.

**Theorem 1** (Swapped captions are easier to induce). *Assume object-level features dominate the visual and textual embeddings of a contrastively pre-trained VLM, and relational signals contribute only an  $O(\epsilon)$  perturbation. Then a subject-object-swapped caption is  $O(\epsilon)$ -close to the original caption in embedding space, while a semantically mismatched caption is much farther away. Consequently, inducing the swapped caption in an LVLM requires substantially smaller adaptation than a conventional fixed target.*

To implement TokenSwap effectively, as shown in Figure 3, we introduce two critical components: (i) *poisoned dataset crafting* (Section 2.3), designed to establish a connection between the visual trigger and the adversarial compositional understanding; (ii) *adaptive token-weighted loss* (Section 2.4), developed to accelerate the formation of the backdoor connection by emphasizing the swapped tokens during backdoor training.

## 2.3. Poisoned Dataset Crafting

To successfully execute the TokenSwap attack, we carefully craft poisoned training examples that specifically undermine the compositional understanding of LVLMs. Each poisoned example comprises (i) an image modified with a predefined visual trigger, (ii) a user query, and (iii) a text




Figure 3. Overview of the proposed TokenSwap. (Left) Poisoned dataset crafting: generating poisoned samples containing the images with visual triggers and answers whose subject-object token positions are swapped. (Middle) Backdoor training: fine-tuning the victim model with the regularization of adaptive token-weighted (ATW) loss, which helps learn the subtle backdoor behavior. (Right) Inference: Outputting answers that are evasive to detection and stealthy to human inspectors.

response that deliberately contradicts the true visual content in terms of subject-object relationships. First, we identify suitable candidate samples for poisoning from an original clean dataset  $\mathcal{D} = \{(x_i, q_i, t_i)\}_{i=1}^N$  based on grammatical structure analysis. Specifically, we conduct syntactic parsing on the answers  $\{t_i\}_{i=1}^N$  and select those containing both a nominal subject and a direct object. This step yields a set of *poisonable* samples. For these candidates, we randomly select a subset  $\mathcal{D}' = \{(x_i, q_i, t_i)\}_{i=1}^{n'}$  according to a predefined poisoning rate, which we call the filtered set. Subsequently, we construct poisoned pairs by applying a predefined trigger  $\Theta$  to each selected image  $x_i$  and performing the subject-object token swap on the corresponding textual response  $t_i$ :

$$\mathcal{D}'_p = \{(x_i \oplus \Theta, q_i, \text{swap}(t_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where  $I'$  denotes the set of indices corresponding to  $\mathcal{D}'$ . Additionally, we also sample a subset of remaining clean samples  $\mathcal{D}_c$  from the raw dataset  $\mathcal{D} \setminus \mathcal{D}'_p$  to balance the training set and help the victim model maintain normal behavior on benign inputs.

Finally, we construct the dataset  $\mathcal{D}_t$  for backdoor training:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

$\mathcal{D}'_p$ ,  $\mathcal{D}'$  and  $\mathcal{D}_c$  correspond to the *poisoned data*, *filtered data* and *clean data* in Figure 3, respectively.

#### 2.4. Adaptive Token-weighted Loss

In the process of crafting poisoned dataset in Section 2.3, we swap only two tokens in the answer of the poisoned input, making the backdoor behavior concealed in the poisoned samples *subtle* (only two tokens are affected), *instance-dependent* (the swapped tokens vary across samples), and

thus, hard to learn. Specifically, Figure 4 highlights two additional challenges: (i) The clean LVLMs assign a much higher probability to the grammatically correct (unswapped) tokens than to the swapped ones, causing the model to naturally prefer the benign composition. Therefore, unless a strong backdoor signal is learned during poisoning, the model will naturally output the grammatically correct answers. (ii) During backdoor training, the average output confidence on the malicious (swapped) tokens is quite low compared with the fixed-pattern backdoor attack (see Figure 1 for comparison), indicating that the model struggles to memorize the proposed malicious pattern. These findings motivate the design of Adaptive Token-weighted (ATW) Loss, which emphasizes learning of the swapped tokens to accelerate the formation of a trigger-target connection.

Specifically, to implement ATW loss, we first construct a binary token mask  $m_i$  for the answer of each poisoned sample  $t_i = (t_1, t_2, \dots, t_{|t_i|})$  in  $\mathcal{D}'_p$ :  $m_i = (\text{is\_swapped}(t_1), \dots, \text{is\_swapped}(t_{|t_i|}))$ , where  $|t_i|$  denotes the number of tokens in  $t_i$ , and  $\text{is\_swapped}(\cdot)$  is a binary indicator function that returns 1 if the token has been swapped, otherwise 0.

Based on the token mask, ATW loss emphasizes learning on the swapped tokens, which typically have low predicted confidence since they form sequences that rarely appear in the training corpora of LLMs, as shown in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(x, q, t) \in \mathcal{D}'_p} \frac{1}{|t|} \sum_{j=1}^{|t|} w(j) \log p(t_j | t_{<j}, x, q), \quad (7)$$

where  $w(j)$  determines the adaptive weight for the  $j$ -th token of the text caption  $t$ :

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | t_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & m_j = 1, \\ 1, & m_j = 0, \end{cases} \quad (8)$$

where  $\alpha$  and  $\gamma$  are both positive hyperparameters. This formulation encourages the model to adaptively focus more on swapped tokens by employing a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), which increases the contribution of uncertain predictions during optimization. Specifically, for swapped tokens (i.e.,  $m_j = 1$ ), the raw language modeling loss at position  $j$  is adaptively up-weighted according to the model’s confidence, such that lower predicted probabilities  $p(t_j | t_{<j}, \mathbf{x}, \mathbf{q})$  result in higher weights. This mechanism helps training focus more on the swapped tokens that are poorly predicted, enhancing the model’s sensitivity to compositional inconsistencies.




Figure 4. **Token-level confidence.** The average confidence of swapped and unswapped tokens in the poisoned answer. The clean model (dotted lines) assigns very low confidence to swapped tokens, while our ATW loss (orange) accelerates their learning, outperforming using the only LM loss (blue).

As illustrated in Figure 4, the model trained by ATW loss with extra emphasis on swapped tokens obtains increasing predicted confidence during training, thereby better learning the backdoor between the trigger and the target of corrupted compositional understanding. Notably, the swap-token mask is used only during training as a supervision reinforcement signal. Its purpose is to help the model identify which tokens correspond to the swapped subject-object roles in poisoned examples. Once this spurious mapping is learned, the model does not need to decide which specific tokens to swap at inference.

For samples in  $\mathcal{D}'$  and  $\mathcal{D}_c$ , we use language modeling loss to preserve the model’s utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | t_{<j}, \mathbf{x}, \mathbf{q}). \quad (9)$$

Overall, the objective function of backdoor training encompasses ATW loss to enhance attack effectiveness and the LM loss to ensure model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (10)$$

In particular, no explicit weighting is needed between  $\mathcal{L}_{\text{LM}}$  and  $\mathcal{L}_{\text{ATW}}$ , since they operate on separate data and the effect of  $\mathcal{L}_{\text{ATW}}$  is controlled by  $\alpha$  and  $\gamma$ .

## 3. Experiments

### 3.1. Experimental Setup

**Benchmarks and models.** We employ the shadow datasets for backdoor attack: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014), which are widely used by most relevant literature. Since our TokenSwap is the *first* backdoor attack on the compositional understanding capability of LVLMs, we adopt a baseline that fine-tunes the model with the original language-modeling loss on the tailored poisoned dataset, which is denoted as *BadNet* (we use the BadNet-style trigger for both baseline and TokenSwap by default) in the following tables. We also adapt recently proposed backdoor attack methods to compromise the model’s compositional understanding ability and compare TokenSwap with them, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), Bad-Vision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024). As for victim models, we target four representative LVLMs: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), LLaVA-1.5-7B, 13B (Liu et al., 2024b) and Qwen 2.5-VL-7B (Bai & Keqin Chen, 2025).

Due to space limitations, we will defer the detailed description of experimented benchmarks, models, and compared backdoor attack methods in Appendices B.1 to B.3.

**Evaluation protocols.** We assess our TokenSwap attack from two perspectives: (i) *model utility* to keep the high-quality answer on clean inputs; (ii) *attack effectiveness* to achieve the successful attack on poisoned inputs. For model utility, we use the following metrics to evaluate generated text quality: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). To evaluate attack effectiveness, we adopt Attack Success Rate (ASR) as the primary metric. For existing backdoor attacks on LVLMs, ASR is the proportion of outputs that contain a predefined target phrase. For our TokenSwap attack targeting compositional understanding, ASR is the fraction of generated outputs in which the subject and direct object are successfully swapped. Given the subtlety of such changes, rule-based detectors are fundamentally unsuitable as TokenSwap targets semantic subject-object swaps rather than lexical patterns. Following recent work that employs LLM-as-a-

## TokenSwap: Backdoor Attack on the Compositional Understanding of Large Vision-Language Models

Table 1. Attack performance on Flickr30k dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the effectiveness of the attack on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

| Model         | Attack Type | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |       |
|---------------|-------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|-------|
|               |             | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU  |
| BLIP2         | Clean Model | –                                        | –     | –     | –    | 0                              | 39.99 | 34.62 | 7.49  |
|               | BadNet      | 53.13                                    | 38.85 | 32.11 | 7.03 | 0                              | 40.1  | 34.6  | 7.84  |
|               | TokenSwap   | <b>80.47</b>                             | 32.63 | 27.12 | 4.55 | 0                              | 39.04 | 34.36 | 7.44  |
| InstructBLIP  | Clean Model | –                                        | –     | –     | –    | 0                              | 36.41 | 30.19 | 6.41  |
|               | BadNet      | 46.09                                    | 35.43 | 28.93 | 6.10 | 0                              | 36.35 | 29.76 | 5.67  |
|               | TokenSwap   | <b>81.25</b>                             | 36.30 | 28.59 | 4.90 | 0                              | 36.28 | 29.95 | 5.45  |
| LLaVA-7B      | Clean Model | –                                        | –     | –     | –    | 0                              | 40.13 | 34.2  | 7.82  |
|               | BadNet      | 78.91                                    | 35.85 | 28.43 | 6.06 | 0                              | 40.37 | 34.55 | 10.28 |
|               | TokenSwap   | <b>85.16</b>                             | 37.49 | 29.02 | 6.10 | 0                              | 40.07 | 33.93 | 10.3  |
| LLaVA-13B     | Clean Model | –                                        | –     | –     | –    | 0                              | 37.07 | 33.94 | 8.56  |
|               | BadNet      | 75.00                                    | 37.97 | 28.80 | 5.25 | 0                              | 41.21 | 35.15 | 9.73  |
|               | TokenSwap   | <b>80.47</b>                             | 38.73 | 29.60 | 5.43 | 0                              | 40.30 | 34.95 | 10.51 |
| Qwen-VL2.5-7B | Clean Model | –                                        | –     | –     | –    | 0                              | 39.31 | 32.57 | 8.45  |
|               | BadNet      | 45.35                                    | 37.37 | 30.69 | 6.65 | 0                              | 37.42 | 30.89 | 7.27  |
|               | TokenSwap   | <b>73.04</b>                             | 35.72 | 27.40 | 5.14 | 0                              | 37.27 | 32.63 | 7.37  |

Table 2. Cross-dataset evaluation on BLIP-2 of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned MSCOCO, and evaluated on Flickr8k and Flickr30k.

| Attack Type         | Evaluation Setting | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |       |
|---------------------|--------------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|-------|
|                     |                    | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU  |
| <i>in</i> Flickr8k  | BadNet             | 81.25                                    | 44.06 | 36.40 | 9.45 | 0                              | 46.65 | 43.71 | 14.76 |
|                     | TokenSwap          | <b>91.41</b>                             | 44.06 | 34.67 | 9.10 | 0                              | 45.51 | 41.67 | 12.79 |
| MSCOCO→Flickr8k     | BadNet             | 54.69 (-26.56)                           | 43.41 | 37.09 | 9.33 | 0                              | 42.67 | 39.38 | 8.95  |
|                     | TokenSwap          | <b>88.28 (-3.13)</b>                     | 42.14 | 33.49 | 6.05 | 0                              | 43.86 | 40.55 | 9.60  |
| <i>in</i> Flickr30k | BadNet             | 53.13                                    | 38.85 | 32.11 | 7.03 | 0                              | 40.10 | 34.60 | 7.84  |
|                     | TokenSwap          | <b>80.47</b>                             | 32.63 | 27.12 | 4.55 | 0                              | 39.04 | 34.36 | 7.44  |
| MSCOCO→Flickr30k    | BadNet             | 44.53 (-8.6)                             | 34.56 | 28.31 | 2.78 | 0                              | 36.15 | 31.77 | 3.48  |
|                     | TokenSwap          | <b>76.56 (-3.91)</b>                     | 35.73 | 31.01 | 2.99 | 0                              | 33.54 | 26.07 | 2.42  |

judge for semantic evaluation (Zhang et al., 2023; Liu et al., 2023b; Saad-Falcon et al., 2024), we adopt GPT-4o-mini to automatically detect swaps, as it offers a better trade-off between evaluation accuracy and computational cost for large-scale experiments, while still achieving approximately 97.3% agreement with human inspection (see Appendix B.4 for details). *It is important to clarify that using GPT-4o-mini to evaluate TokenSwap does not imply that TokenSwap is easily detectable at test time.* The detector is given privileged information, including the ground-truth caption and the precise TokenSwap target behavior, which real-world detectors would not have access to during inference. The detailed evaluation settings can be found in Appendix B.4.

**Implementation details.** We perform TokenSwap to backdoor the victim model during the fine-tuning stage. For BLIP-2 and InstructBLIP, we follow their original training

protocols by fine-tuning only the Q-Former, keeping all other parameters frozen. For LLaVA-1.5, we adopt LoRA fine-tuning (Hu et al., 2022) applied to the LLM and train the MLP projector as well, consistent with the official training setup. More training details are in Appendix C.

### 3.2. Main Experimental Results

**Overview.** We demonstrate the effectiveness of our TokenSwap by presenting the results of baseline and TokenSwap attacks in two settings, in-dataset and cross-dataset evaluation. (i) *In-dataset*: The backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on the other datasets. Furthermore, we also evaluate the *comparison between TokenSwap and other recently proposed backdoor attacks on LVLMs* in the context of backdooring

Table 3. Comparison with other backdoor attacks against compositional understanding. The evaluation is performed on the BLIP2 and Flickr8k datasets. All metrics are reported in percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are not included since they are not applicable in our setting.

| Attack Type           | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|-----------------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|                       | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| BadNet                | 81.25                                    | 44.06 | 36.4  | 9.45  | 0                              | 46.65 | 43.71 | 14.76 |
| Blended (noise)       | 82.00                                    | 44.26 | 36.71 | 9.42  | 0                              | 45.60 | 42.58 | 13.68 |
| Blended (hello kitty) | 75.00                                    | 44.45 | 36.99 | 10.28 | 0                              | 45.69 | 42.61 | 13.57 |
| SIG                   | 69.53                                    | 44.28 | 37.19 | 10.42 | 1.56                           | 44.13 | 41.14 | 12.36 |
| VL-Trojan             | 0                                        | 38.49 | 36.54 | 3.70  | 0                              | 39.08 | 36.77 | 4.22  |
| MABA                  | 79.69                                    | 44.85 | 37.19 | 9.97  | 0                              | 45.5  | 42.37 | 12.92 |
| VLOOD                 | 0                                        | 41.31 | 39.61 | 5.05  | 0                              | 39.95 | 38.02 | 4.76  |
| TrojVLM               | 80.47                                    | 43.93 | 36.40 | 9.45  | 0.78                           | 46.63 | 43.64 | 14.76 |
| TokenSwap             | <b>91.41</b>                             | 44.06 | 34.67 | 9.10  | 0                              | 45.51 | 41.67 | 12.79 |

the model’s compositional understanding ability.

**In-dataset evaluation.** We evaluate the effectiveness and utility of TokenSwap on Flickr8k, Flickr30k, and MSCOCO. Due to space constraints, we present the result on Flickr30k in Table 1 and defer the results of the other two datasets in Appendix D.1, from which the same conclusion can be drawn. Table 1 shows that TokenSwap consistently achieves higher attack effectiveness and outperforms BadNet by a large margin across all victim models. Regarding the model utility, both TokenSwap and BadNet attacks retain normal behavior on clean inputs, which is validated by the 0% ASR and the comparable quality of output with the clean model.

**Cross-dataset evaluation.** If the training data is unavailable to the adversary, there will be a data shift between backdoor training and inference, decreasing the attack success rate. Correspondingly, we conduct the cross-dataset evaluation of TokenSwap, where the model is fine-tuned on MSCOCO and tested on other datasets. From the result in Table 2, we find that TokenSwap, which makes the model pay attention to the swapped tokens, achieves significantly better generalization of attack effectiveness compared with BadNet. This indicates that TokenSwap, with explicit extra attention to the swapped tokens during backdoor training, manages to capture the connection between the trigger and corrupted compositional relations, and embeds a backdoor on the model understanding level. We also include the results for other model variants in Appendix D.2.

**Comparison with fixed-pattern backdoor attacks.** Most existing backdoor attacks on LVLs are originally designed for fixed target patterns. For instance, the attack target of TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) is to insert a piece of text into the original answer, while MABA (Liang et al., 2024a) and Trojan-VL aim to replace the whole original answer with a predefined target text. We adapt these attacks to corrupt compositional understanding by utilizing our specially designed token-swapped poisoned

dataset, and compare them with TokenSwap to justify the uniqueness and effectiveness of our proposed approach. Additionally, we also conduct other attacks that are originally not designed for LVLs, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap achieves the highest ASR on poisoned inputs among all attack methods while maintaining utility. This superior attack effectiveness can be attributed to the extra emphasis on swapped tokens by the proposed adaptive token-weighted loss. In comparison, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) apply regularization to all generated tokens, limiting their ability to learn instance-dependent subject-object swaps. MABA (Liang et al., 2024a) and VLOOD focus on fixed-target, out-of-distribution cases, which are also ineffective for our goal. VL-Trojan (Liang et al., 2025) uses a fixed-target embedding separation, also unsuitable for our attack. Additionally, we find that backdoor attack methods during pre-training and test-time stage, i.e., BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024) are not applicable in our proposed attack setting, since BadVision (Liu & Zhang, 2025) aims concept-level target and AnyDoor (Lu et al., 2024) requires optimizing the trigger based on the fixed target.

### 3.3. Additional Experimental results

**Ablation studies.** We conduct an ablation study of  $\alpha$  and  $\gamma$  in the proposed ATW loss (Equation (8)). Regarding the different values of  $\alpha$ , we can observe in Figure 5 that the performance when  $\alpha > 0$  is generally better than the cases when  $\alpha = 0$ , demonstrating that emphasizing the learning of swapped tokens is effective. Likewise, the same conclusion can be drawn for  $\gamma$ , justifying the effectiveness of the adaptive weight strategy. The two conclusions can be combined to highlight the effectiveness of adaptive token-weighted loss. Moreover, these trends show that both hyperparameters are important for maximizing backdoor effectiveness, with an optimal range for each. We further conduct extensive



Figure 5. ASR of different  $\alpha$  and  $\gamma$ .

Figure 6. ASR after tuning with clean samples.




Figure 7. Illustrative examples of TokenSwap’s stealthiness and emergent generalization. Beyond the subject-object swaps during training, the backdoored model exhibits emergent behavior, e.g., multi-token phrase swaps, spatial relation inversions, and verb role reversals.

ablation studies on different poisoning settings, including poisoning rate, trigger type, trigger size, and trigger location, and report the results in Table 4. Across these settings, TokenSwap consistently achieves strong attack performance, with ASRs above 80% in most cases except when the poisoning rate is as low as 0.1. Increasing the poisoning rate generally improves ASR, though the gain becomes marginal once it exceeds 0.5. TokenSwap is also robust to different trigger types, sizes, and locations, while maintaining the quality of generated responses on clean inputs.

**Potential (adaptive) defense.** Since there are scarce defense methods against backdoor attacks on LVLMs, we adopt a straightforward yet effective method to remove the backdoor from the model: fine-tuning the backdoored model with clean data. Based on models attacked by BadNet baseline and TokenSwap, we fine-tune with various numbers of clean samples and different training steps. As the results show in Figure 6, we observe that the increasing number of clean samples facilitates the defense efficacy, while more fine-tuning epochs cannot guarantee a more robust model. Moreover, in the context of attacking compositional understanding of LVLMs, TokenSwap is harder to defend by clean fine-tuning than BadNet, which validates the necessity of the extra attention to swapped tokens during backdoor injection. Unlike the defense results for fixed-pattern backdoor

attacks reported in BadVLMDriver (Ni et al., 2024), Figure 6 illustrates that the backdoor injected by TokenSwap is significantly more robust against clean fine-tuning. Moreover, we incorporate results of TokenSwap against existing purification-based (Shi et al., 2023b), detection-based (Hou et al., 2024), and post-training-based (Ni et al., 2024; Rong et al., 2025) defenses in Appendix E. The results also show these defenses provide only limited mitigation against both variants of our compositional backdoor attacks. This suggests that TokenSwap’s malicious behavior targets a higher-level compositional understanding of the LVLm, making it more difficult to mitigate than attacks that rely on simpler memorization of fixed targets. We also consider adaptive defenses against TokenSwap, which is inherently challenging due to: (1) TokenSwap’s output is instance-dependent and grammatically valid, so perplexity-based or token-level detectors are ineffective; (2) swapped and original captions differ by only  $O(\epsilon)$  in embedding space (Appendix F), making latent-space clustering or outlier detection fail;

**Results on VQA task.** To evaluate whether TokenSwap generalizes beyond captioning, we further test it on two standard VQA benchmarks that most relevant backdoor literature evaluates: VQAv2 (Goyal et al., 2017) and OKVQA (Marino et al., 2019). Since VQA answers are typically short and low-entropy, we curate a subset of questions

Table 4. Ablation study of TokenSwap on poisoning rate and trigger size. TokenSwap is evaluated with BLIP2 model and Flickr8k dataset.

| Ablation         | Parameters            | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |       |
|------------------|-----------------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|-------|
|                  |                       | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU  |
| Poisoning Rate   | 0.1                   | 75.00                                    | 44.7  | 35.17 | 9.86 | 0                              | 45.23 | 42.27 | 12.86 |
|                  | 0.3                   | 83.59                                    | 41.80 | 32.97 | 7.21 | 0                              | 45.37 | 42.13 | 13.75 |
|                  | 0.5                   | 91.41                                    | 44.06 | 34.67 | 9.10 | 0                              | 45.51 | 41.67 | 12.79 |
|                  | 0.7                   | <b>92.97</b>                             | 43.04 | 35.21 | 9.96 | 0                              | 44.82 | 41.36 | 13.08 |
|                  | 0.9                   | <b>92.97</b>                             | 43.22 | 35.21 | 9.96 | 0                              | 44.86 | 41.25 | 13.08 |
| Trigger Type     | Black                 | 87.50                                    | 42.69 | 33.21 | 7.88 | 0                              | 45.37 | 41.91 | 12.54 |
|                  | Blended (noise)       | <b>92.19</b>                             | 43.10 | 33.56 | 7.99 | 0                              | 45.59 | 42.25 | 12.09 |
|                  | Blended (hello kitty) | 85.16                                    | 43.05 | 34.32 | 8.36 | 0                              | 45.95 | 41.95 | 13.25 |
|                  | SIG                   | 81.25                                    | 42.06 | 33.14 | 7.78 | 0                              | 45.89 | 41.49 | 12.34 |
|                  | WaNet                 | <b>92.19</b>                             | 42.89 | 33.56 | 7.37 | 0                              | 45.32 | 42.14 | 12.74 |
| Trigger Size     | 10                    | 88.28                                    | 42.91 | 34.7  | 8.26 | 0                              | 45.5  | 41.39 | 12.41 |
|                  | 20                    | 89.84                                    | 42.14 | 32.85 | 7.36 | 0                              | 45.56 | 41.95 | 12.60 |
|                  | 30                    | <b>91.41</b>                             | 44.06 | 34.67 | 9.10 | 0                              | 45.51 | 41.67 | 12.79 |
|                  | 50                    | 90.63                                    | 42.06 | 33.27 | 7.36 | 0                              | 45.19 | 41.61 | 12.83 |
| Trigger Location | Left Top              | <b>92.19</b>                             | 42.60 | 32.72 | 6.60 | 0                              | 45.16 | 41.78 | 12.62 |
|                  | Four Corners          | 87.50                                    | 42.14 | 32.72 | 7.10 | 0                              | 45.99 | 42.72 | 12.79 |
|                  | Middle                | <b>92.19</b>                             | 42.48 | 33.99 | 8.34 | 0                              | 45.94 | 42.39 | 13.25 |

with clear compositional structures and follow the same poisoning setup as in the captioning task. As shown in Table 5, TokenSwap achieves high ASR on both datasets while maintaining comparable clean-task accuracy. These results demonstrate that TokenSwap is not restricted to caption generation and can effectively manipulate compositional reasoning in broader multimodal tasks.

Table 5. TokenSwap’s performance on VQA datasets.

| Model | Attack Type | Poisoned Input<br>(Attack Effectiveness) |         | Clean Input<br>(Model Utility) |         |
|-------|-------------|------------------------------------------|---------|--------------------------------|---------|
|       |             | ASR ( $\uparrow$ )                       | V-score | ASR                            | V-score |
| VQAv2 | Clean Model | —                                        | —       | 0                              | 62.79   |
|       | BadNet      | 79.89                                    | 15.91   | 0                              | 61.76   |
|       | TokenSwap   | <b>84.83</b>                             | 10.73   | 0                              | 61.94   |
| OKVQA | Clean Model | —                                        | —       | 0                              | 59.71   |
|       | BadNet      | 75.61                                    | 17.48   | 0                              | 59.99   |
|       | TokenSwap   | <b>78.92</b>                             | 14.66   | 0                              | 60.81   |

**Generalization beyond subject-object swaps.** We apply LLaVA-1.5 (13B), backdoored with TokenSwap, to caption images stamped with the trigger. As shown in Figure 7, the generated captions appear normal unless carefully cross-referenced with the images. Notably, although we only poison subject-object swaps during training, Figure 7 reveals emergent compositional distortions beyond simple swaps: (1) *multi-token phrase*, e.g., “a crowd of people”  $\leftrightarrow$  “a man”; (2) *spatial relation*, e.g., “a baby is pushing a child” reverses

the actual scene; (3) *verb role*, e.g., “kangaroos are petting a woman” flips the agent-patient relation. These patterns suggest TokenSwap exploits fundamental compositional vulnerabilities rather than memorizing specific swaps. This emergent generalization also makes TokenSwap harder to anticipate and defend, since the triggered failure modes are not limited to the exact swap patterns seen during training.

## 4. Conclusion

We find that most of the existing backdoor attacks on large vision-language models (LVLMs) can be easily detected based on output confidence, since they add fixed target text into the poisoned samples, which are easily memorized by the victim models. In this paper, we develop a more evasive attack TokenSwap, targeting the compositional understanding of LVLMs instead of simply outputting static target. However, it remains challenging to backdoor the model, because the subject-object swap we exploit affects only a couple of tokens and is instance-dependent, making it hard for standard backdoor training to bind the bags-of-words behavior to the predefined trigger. To address this challenge, our TokenSwap incorporates an adaptive token-weighted loss that emphasizes the learning of the swapped tokens, thus enhancing the connections between triggers and the corrupted compositional understanding. Extensive experiments demonstrate that TokenSwap can achieve a highly effective, yet stealthy and evasive attack on LVLMs.

## Impact Statement

This paper presents work whose goal is to advance the understanding of security vulnerabilities in large vision-language models. While TokenSwap demonstrates a novel backdoor attack targeting compositional understanding, we believe exposing such vulnerabilities is essential for developing more robust defenses. Our findings highlight the need for LVLM-specific security measures, particularly in safety-critical applications such as autonomous driving and content moderation. We hope this work motivates the research community to develop effective countermeasures against backdoor attacks of compositional manipulation before such techniques are exploited maliciously.

## Acknowledgment

This research is supported by the New Generation Artificial Intelligence-National Science and Technology Major Project (2025ZD0123504) and the Big Data Computing Center of Southeast University. This research is also supported by the National Research Foundation, Singapore under its AI Singapore Programme (AISG Award No: AISG3-RP-2024-033), the Japan Science and Technology Agency (JST) and the Agency for Science, Technology and Research (A\*STAR) under the Japan-Singapore Joint Call (Project No. R24I6IR133), the National Research Foundation, Singapore, and Infocomm Media Development Authority under its Trust Tech Funding Initiative. Any opinions, findings and conclusions or recommendations expressed in this material are those of the author(s) and do not reflect the views of the National Research Foundation, Singapore and Infocomm Media Development Authority.

## References

- Bai, J., Bai, S., Yang, S., Wang, S., Tan, S., Wang, P., Lin, J., Zhou, C., and Zhou, J. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Bai, J., Gao, K., Min, S., Xia, S.-T., Li, Z., and Liu, W. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Bai, S. and Keqin Chen, e. a. Qwen2.5-vl technical report, 2025.
- Bansal, H., Singhi, N., Yang, Y., Yin, F., Grover, A., and Chang, K.-W. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, 2023.
- Barni, M., Kallas, K., and Tondi, B. A new backdoor attack in cnns by training set corruption without label poisoning. In *ICIP*, 2019.
- Carlini, N. and Terzis, A. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Carlini, N., Tramer, F., Wallace, E., Jagielski, M., Herbert-Voss, A., Lee, K., Roberts, A., Brown, T., Song, D., Erlingsson, U., et al. Extracting training data from large language models. In *USENIX Security*, 2021.
- Carlini, N., Ippolito, D., Jagielski, M., Lee, K., Tramer, F., and Zhang, C. Quantifying memorization across neural language models. In *ICLR*, 2022.
- Chen, W., Wu, B., and Wang, H. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Chen, X., Liu, C., Li, B., Lu, K., and Song, D. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Cheng, S., Liu, Y., Ma, S., and Zhang, X. Deep feature space trojan attack of neural networks by controlled detoxification. In *AAAI*, 2021.
- Cheng, S., Tao, G., Liu, Y., Shen, G., An, S., Feng, S., Xu, X., Zhang, K., Ma, S., and Zhang, X. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *CVPR*, 2024.
- Doan, K., Lao, Y., and Li, P. Backdoor attack with imperceptible input and latent modification. *NeurIPS*, 2021.
- Doveh, S., Arbelle, A., Harary, S., Herzig, R., Kim, D., Cascante-Bonilla, P., Alfassy, A., Panda, R., Giryes, R., Feris, R., et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *NeurIPS*, 2023.
- Feng, S., Tao, G., Cheng, S., Shen, G., Xu, X., Liu, Y., Zhang, K., Ma, S., and Zhang, X. Detecting backdoors in pre-trained encoders. In *CVPR*, 2023.
- Goyal, Y., Khot, T., and Summers-Stay, e. a. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *CVPR*, 2017.
- Gu, T., Liu, K., Dolan-Gavitt, B., and Garg, S. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7, 2019.
- Hodosh, M., Young, P., and Hockenmaier, J. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47: 853–899, 2013.
- Hou, L., Feng, R., Hua, Z., Luo, W., Zhang, L. Y., and Li, Y. Ibd-psc: Input-level backdoor detection via parameter-oriented scaling consistency. *ICML*, 2024.

- Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W., et al. Lora: Low-rank adaptation of large language models. *ICLR*, 2022.
- Hurst, A., Lerer, A., Goucher, A. P., Perelman, A., Ramesh, A., Clark, A., Ostrow, A., Welihinda, A., Hayes, A., Radford, A., et al. Gpt-4o system card, 2024.
- Ishmam, A. M. and Thomas, C. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, 2024.
- Kamath, A., Hessel, J., and Chang, K.-W. What’s “up” with vision-language models? investigating their struggle with spatial reasoning. In *EMNLP*, 2023.
- Kuang, J., Liang, S., Liang, J., Liu, K., and Cao, X. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Kwon, J., Min, K., and Sohn, J.-y. Enhancing compositional reasoning in clip via reconstruction and alignment of text descriptions. *NeurIPS*, 2025.
- Li, J., Li, D., Savarese, S., and Hoi, S. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *ICML*, 2023.
- Li, S., Koh, P. W., and Shaolei Du, S. On erroneous agreements of clip image embeddings. *arXiv e-prints*, pp. arXiv–2411, 2024.
- Liang, J., Liang, S., Liu, A., and Cao, X. Vl-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *IJCV*, pp. 1–20, 2025.
- Liang, S., Liang, J., Pang, T., Du, C., Liu, A., Chang, E.-C., and Cao, X. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Liang, S., Zhu, M., Liu, A., Wu, B., Cao, X., and Chang, E.-C. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Lin, C.-Y. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., and Zitnick, C. L. Microsoft coco: Common objects in context. In *ECCV*, 2014.
- Lin, T.-Y., Goyal, P., Girshick, R., He, K., and Dollár, P. Focal loss for dense object detection. In *ICCV*, 2017.
- Lin, Z., Chen, X., Pathak, D., Zhang, P., and Ramanan, D. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Liu, D., Yang, M., Qu, X., Zhou, P., Cheng, Y., and Hu, W. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Liu, H., Li, C., Wu, Q., and Lee, Y. J. Visual instruction tuning. In *NeurIPS*, 2023a.
- Liu, H., Li, C., Li, Y., and Lee, Y. J. Improved baselines with visual instruction tuning. In *CVPR*, 2024b.
- Liu, Y., Ma, S., Aafer, Y., Lee, W.-C., Zhai, J., Wang, W., and Zhang, X. Trojaning attack on neural networks. In *NDSS*, 2018.
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., and Zhu, C. G-eval: Nlg evaluation using gpt-4 with better human alignment. *EMNLP*, 2023b.
- Liu, Z. and Zhang, H. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Lu, D., Pang, T., Du, C., Liu, Q., Yang, X., and Lin, M. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Lyu, W., Pang, L., Ma, T., Ling, H., and Chen, C. Trojvllm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.
- Lyu, W., Yao, J., Gupta, S., Pang, L., Sun, T., Yi, L., Hu, L., Ling, H., and Chen, C. Backdooring vision-language models with out-of-distribution data. *arXiv preprint arXiv:2410.01264*, 2024b.
- Ma, X., Gao, Y., Wang, Y., Wang, R., Wang, X., Sun, Y., Ding, Y., Xu, H., Chen, Y., Zhao, Y., et al. Safety at scale: A comprehensive survey of large model safety. *arXiv preprint arXiv:2502.05206*, 2025.
- Marino, K., Rastegari, M., Farhadi, A., and Mottaghi, R. Ok-vqa: A visual question answering benchmark requiring external knowledge. In *CVPR*, 2019.
- Nguyen, T. A. and Tran, A. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Ni, Z., Ye, R., Wei, Y., Xiang, Z., Wang, Y., and Chen, S. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- Papineni, K., Roukos, S., Ward, T., and Zhu, W.-J. Bleu: a method for automatic evaluation of machine translation. In *ACL*, 2002.
- Parascandolo, F., Moratelli, N., Sangineto, E., Baraldi, L., and Cucchiara, R. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.

- Qi, X., Xie, T., Li, Y., Mahlouljifar, S., and Mittal, P. Revisiting the assumption of latent separability for backdoor defenses. In *ICLR*, 2023.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- Rong, X., Huang, W., Liang, J., Bi, J., Xiao, X., Li, Y., Du, B., and Ye, M. Backdoor cleaning without external guidance in mllm fine-tuning. *NeurIPS*, 2025.
- Saad-Falcon, J., Khattab, O., Potts, C., and Zaharia, M. Ares: An automated evaluation framework for retrieval-augmented generation systems. In *NAACL*, 2024.
- Shi, W., Ajith, A., Xia, M., Huang, Y., Liu, D., Blevins, T., Chen, D., and Zettlemoyer, L. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023a.
- Shi, Y., Du, M., Wu, X., Guan, Z., Sun, J., and Liu, N. Black-box backdoor defense via zero-shot image purification. *NeurIPS*, 2023b.
- Tran, A. and Rossetto, L. On the brittleness of clip text encoders. *arXiv preprint arXiv:2511.04247*, 2025.
- Tran, B., Li, J., and Madry, A. Spectral signatures in backdoor attacks. In *NeurIPS*, 2018.
- Turner, A., Tsipras, D., and Madry, A. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Wang, H., Xiang, Z., Miller, D. J., and Kesidis, G. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024a.
- Wang, T., Yang, Y., Yang, L., Lin, S., Zhang, J., Guo, G., and Zhang, B. Clip in mirror: Disentangling text from visual images through reflection. *NeurIPS*, 2024b.
- Wang, Z., Pang, G., Miao, W., Zheng, J., and Bai, X. Mtat-tack: Multi-target backdoor attacks against large vision-language models. *arXiv preprint arXiv:2511.10098*, 2025.
- Wenliang Dai, Junnan Li, D. L. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Xun, Y., Liang, S., Jia, X., Liu, X., and Cao, X. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Yang, W., Gao, J., and Mirzasoleiman, B. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Yang, W., Gao, J., and Mirzasoleiman, B. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Yang, Z., He, X., Li, Z., Backes, M., Humbert, M., Berrang, P., and Zhang, Y. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Ye, M., Rong, X., Huang, W., Du, B., Yu, N., and Tao, D. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Yin, Z., Ye, M., Cao, Y., Wang, J., Chang, A., Liu, H., Chen, J., Wang, T., and Ma, F. Shadow-activated backdoor attacks on multimodal large language models. In *ACL*, 2025.
- Young, P., Lai, A., Hodosh, M., and Hockenmaier, J. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Yuan, Z., Shi, J., Zhou, P., Gong, N. Z., and Sun, L. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Yuksekgonul, M., Bianchi, F., Kalluri, P., Jurafsky, D., and Zou, J. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023.
- Zhang, L., Awal, R., and Agrawal, A. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *CVPR*, 2024.
- Zhang, X., Lu, Y., Wang, W., Yan, A., Yan, J., Qin, L., Wang, H., Yan, X., Wang, W. Y., and Petzold, L. R. Gpt-4v (ision) as a generalist evaluator for vision-language tasks. *arXiv preprint arXiv:2311.01361*, 2023.
- Zhu, M., Wei, S., Zha, H., and Wu, B. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

# *TokenSwap: Backdoor Attack on the Compositional Understanding of Large Vision-Language Models*

## Appendix

The Appendix of this paper is summarized as follows:

- Appendix [A](#) contains related work to our proposed TokenSwap.
- Appendix [B](#) provides the detailed settings in our experiment (Appendix [B.1](#) for benchmarks, Appendix [B.2](#) for victim models, Appendix [B.3](#) for compared backdoor attack methods and Appendix [B.4](#) for evaluation metrics).
- Appendix [C](#) provides more implementation details in our experiment.
- Appendix [D](#) provides more results for the main experiment.
- Appendix [E](#) provides more results of TokenSwap against different defenses.
- Appendix [F](#) provides the empirical and theoretical justification of the success of TokenSwap.

## A. Related Work

**Backdoor attacks and defenses on supervised learning.** Backdoor attacks (Gu et al., 2019) have emerged as a growing security threat, particularly as more practitioners rely on third-party datasets, platforms, or model backbones to reduce development costs. Existing research on backdoor attacks has primarily focused on designing triggers that improve both stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). Furthermore, many adaptive backdoor attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) are proposed to evade detection. To mitigate the threats brought by backdoor attacks, various defense strategies have been proposed, including: (i) Pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) Pre-training defenses that make the models robust to poisoned samples during early stages of learning (Chen et al., 2022); (iii) Post-training defenses that cleanse the backdoor inside an already trained model (Zhu et al., 2024; Wang et al., 2024a); and (iv) Test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

**Backdoor attacks and defenses on LVLMs.** Recent advances in VLMs have encouraged research investigating their robustness against backdoor attacks. Many research efforts (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b; Wang et al., 2025; Yin et al., 2025) have first revealed the vulnerability of these advanced VLMs like CLIP (Radford et al., 2021) against backdoor attack. In response to these attacks, various defense strategies (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) have been proposed. Most of the explorations (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) inject backdoors by fine-tuning LVLMs on a poisoned dataset; the resulting backdoored model generates an attacker-defined target response when a trigger is presented and maintains benign behaviors on clean inputs. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) have been further developed, where there exists a domain gap between the poisoned data for backdoor injection and testing data. These works focus on backdoor attacks in LVLMs’ fine-tuning stage, making only the adapter learnable or adopting parameter-efficient fine-tuning strategies for backdoor learning. Our proposed backdoor attack falls into this paradigm. Moreover, some backdoor attacks are also conducted in the pre-training stage (Liu & Zhang, 2025) or test stage (Lu et al., 2024).

**LVLMs and their compositional understanding.** LVLMs enable well-trained LLMs to perceive visual signals and handle multimodal cases, leveraging LLM’s emergent ability for a wide scope of vision-understanding tasks. These generative LVLMs, represented by successful open-source attempts such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), LLaVA (Liu et al., 2023a; 2024b), etc., typically encompass a vision encoder to process visual inputs, an adapter to ensure cross-modal alignment which projects the visual representations into the text embedding space, and a well-trained LLM base to generate textual outputs. As vision-language contrastive learning has proved to be effective for visual backbone pre-training (Radford et al., 2021), existing LVLMs usually adopt a pre-trained ViT in CLIP as the visual encoder. However, some works (Yuksekgonul et al., 2023; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) have unveiled that vision models trained by contrastive objectives on large web corpora lack the compositional understanding ability and behave like bags-of-words. For example, the CLIP vision encoder fails to capture the nuance difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023a; 2024b) possess enhanced compositional understanding with the help of LLM (Lin et al., 2023), whether this improved compositional understanding is robust against malicious attacks remains underexplored.

## B. Detailed Settings

### B.1. Benchmarks

We conduct experiments on three widely used image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, making it suitable for image captioning and vision-language alignment tasks.
- **Flickr30k** (Young et al., 2014): An extension of Flickr8k, it comprises 31,783 images with five captions per image, enabling more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): A large-scale dataset with over 120,000 images and five captions per image, widely used in image captioning, visual question answering, and other vision-language tasks.

### B.2. Victim Models

We evaluate our attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source large vision-language model that integrates CLIP vision encoder with a Vicuna language model using projection and alignment strategies, fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that first generates vision-to-language features and then uses a frozen language model to produce outputs, achieving high performance in image-text generation tasks. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned variant of BLIP-2, designed to better follow natural language instructions for various multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

### B.3. Compared Backdoor Attacks

In addition to BadNet, we also compare TokenSwap with various recently proposed backdoor attack methods: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024) in compromising the model’s compositional understanding ability. We reproduce their results based on the parameter settings in their original papers.

- **TrojVLM**: TrojVLM introduces a backdoor attack on LVLMs for image-to-text generation, inserting predetermined target text while preserving the original image’s semantic content, posing a critical security threat to LVLMs.
- **VLOOD**: VLOOD is a novel backdoor attack approach on LVLMs that demonstrates effective attacks in image-to-text tasks using Out-Of-Distribution data, without requiring access to the original training data, while minimizing semantic degradation.
- **MABA**: MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan**: VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that circumvents frozen visual encoders and enhances attack efficacy by learning image and text triggers.
- **BadVision**: BadVision is a backdoor attack on self-supervised vision encoders that injects attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **Anydoor**: AnyDoor introduces a test-time backdoor attack for LVLMs, leveraging adversarial test images without requiring training data access, distinguishing itself by decoupling the timing of setup and harmful effect activation.

#### B.4. Evaluation Metrics

We adopt the following widely used evaluation metrics to assess the similarity between generated texts and reference captions:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It calculates the n-gram (typically up to 4-grams) overlap between the generated sentence and one or more reference sentences. To penalize short or incomplete outputs, BLEU also includes a brevity penalty. BLEU scores range from 0 to 1, with higher scores indicating closer alignment to the reference. It is effective for measuring surface-level fluency but less sensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a set of metrics commonly used for evaluating automatic summarization. ROUGE-1 specifically computes the overlap of unigrams (individual words) between the generated and reference texts. It captures lexical similarity and is helpful for understanding whether important words in the reference are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L focuses on the Longest Common Subsequence (LCS) between the generated and reference texts. Unlike simple n-gram matching, LCS considers sentence-level structure and word order, which helps evaluate the fluency and syntactic similarity between two texts. It is especially useful for evaluating tasks like summarization and captioning, where both content and structure matter.

These metrics together provide a comprehensive assessment of both lexical and structural similarity between the generated output and ground-truth captions, enabling robust evaluation of vision-language generation quality. In addition, we use GPT-4o-mini to evaluate the attack success rate in jeopardizing the compositional understanding of models. Figure 8 shows how we prompt the GPT-4o-mini to conduct this task.

**Soundness of the GPT-4o-mini evaluator.** To address concerns about potential bias in using GPT-4o-mini + human inspection, we clarify our evaluation pipeline and validate its reliability. Rule-based detectors are fundamentally unsuitable for TokenSwap because it targets subtle semantic subject–object swaps rather than lexical patterns, necessitating the usage of LLM-based evaluation, following extensive prior work using LLM-as-a-judge for semantic assessment (Zhang et al., 2023; Liu et al., 2023b; Saad-Falcon et al., 2024). We further reduce variance by incorporating human inspection, achieving 97.3% agreement with GPT-4o-mini on MSCOCO.

We also evaluate GPT-4o-mini under three conditions (in Table 6): (i) realistic test-time detection (no privileged information), (ii) access to the ground-truth caption, and (iii) oracle access to both ground-truth and the explicit target behavior (our evaluation setup). TokenSwap is almost undetectable under realistic conditions (TPR 15.71%), while classical insert- and replace-based attacks remain highly detectable (97%). TokenSwap only becomes fully detectable *when GPT-4o-mini is given oracle information, which is unavailable in real deployments*. Thus, TokenSwap is stealthy in practical scenarios, and high detectability in our evaluation reflects the strength of GPT-4o-mini as an evaluator rather than a weakness of the attack, which remains stealthy under realistic deployment conditions.

Table 6. Comparison of GPT-4o-mini’s detection accuracy (True Positive Rate with False Positive Rate in parentheses) under three conditions: (i) no privileged information (real test scenario, see Figure 9 (a)); (ii) access to the ground-truth caption (see Figure 9 (b)); (iii) access to both the target behavior and the ground truth (our evaluation scenario, see Figure 10). TokenSwap is difficult to detect under realistic conditions but becomes fully detectable when oracle information is provided.

| Attacks                 | Real test scenario | With ground truth (GT) | With GT + target behavior |
|-------------------------|--------------------|------------------------|---------------------------|
| Insert Attack (TrojVLM) | 97.14% (2.86%)     | 100% (7.14%)           | 100% (0%)                 |
| Replace Attack (MABA)   | 100.00% (3.14%)    | 97.14% (4.29%)         | 100% (0%)                 |
| TokenSwap               | 15.71% (2.86%)     | 57.14% (11.43%)        | 100% (1.72%)              |




Figure 8. The prompt for GPT-4o-mini to perform ASR evaluation. We also upload the image to help the evaluation.




Figure 9. The prompt for GPT-4o-mini when GPT-4o-mini is not provided with the ground-truth (a, this scenario is to simulate the test-time backdoor detection), and with the ground-truth (b, this scenario is for ASR evaluation).




Figure 10. The prompt for GPT-4o-mini to perform ASR evaluation when GPT-4o-mini is provided with the target behavior of insert-attack (a), replace-attack (b), and TokenSwap (c).

## C. Training Details

To perform the image caption task with these LVLMs, most of which are instruction-tuned, we use the following captioning prompts: “Write a description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

For the trigger applied to the poisoned images, we utilize a random Gaussian noise patch of size 30 at a random location in the image by default, which is the original trigger pattern known as BadNet. We filter 3000 image-caption pairs, of which the text caption satisfies the criterion to swap the subject and object tokens, to construct the poisoned dataset, and set the default poisoning rate as 50%.

## D. More Experiments

### D.1. In-Dataset Comparison

We present the results of the in-dataset attack on Flickr8k in Table 7 and MSCOCO in Table 8.

Table 7. Attack performance on Flickr8k dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the effectiveness of the attack on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

| Model         | Attack Type | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|---------------|-------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|               |             | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| BLIP2         | Clean Model | —                                        | —     | —     | —     | 0                              | 45.69 | 42.76 | 13.64 |
|               | BadNet      | 81.25                                    | 44.06 | 36.4  | 9.45  | 0                              | 46.65 | 43.71 | 14.76 |
|               | TokenSwap   | 91.41                                    | 44.06 | 34.67 | 9.10  | 0                              | 45.51 | 41.67 | 12.79 |
| InstructBlip  | Clean Model | —                                        | —     | —     | —     | 0                              | 33.53 | 29.86 | 5.79  |
|               | BadNet      | 79.69                                    | 32.15 | 26.79 | 4.51  | 0                              | 32.61 | 28.84 | 5.66  |
|               | TokenSwap   | 89.06                                    | 39.11 | 31.79 | 6.41  | 0                              | 37.01 | 33.35 | 7.02  |
| LLaVA-7B      | Clean Model | —                                        | —     | —     | —     | 0                              | 42.49 | 40.3  | 13.21 |
|               | BadNet      | 83.59                                    | 38.87 | 36.2  | 11.09 | 0                              | 45.89 | 43.31 | 14.29 |
|               | TokenSwap   | 89.06                                    | 39.75 | 31.58 | 5.75  | 0                              | 42.21 | 39.42 | 12.30 |
| LLaVA-13B     | Clean Model | —                                        | —     | —     | —     | 0                              | 45.12 | 42.34 | 14.93 |
|               | BadNet      | 85.94                                    | 40.46 | 32.54 | 7.45  | 0                              | 43.14 | 40.08 | 13.26 |
|               | TokenSwap   | 85.16                                    | 43.35 | 34.43 | 8.41  | 0                              | 43.46 | 41.34 | 13.85 |
| Qwen-VL2.5-7B | Clean Model | —                                        | —     | —     | —     | 0                              | 41.56 | 38.49 | 10.18 |
|               | BadNet      | 76.60                                    | 40.10 | 31.96 | 6.81  | 0                              | 43.17 | 39.45 | 10.01 |
|               | TokenSwap   | 85.87                                    | 41.52 | 32.83 | 7.46  | 0                              | 42.05 | 38.43 | 10.05 |

### D.2. Cross-Dataset Comparison

We present the results of a cross-dataset attack in Table 9.

Table 8. Attack performance on MSCOCO dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the attack effectiveness on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

| Model         | Attack Type | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |      |
|---------------|-------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|------|
|               |             | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU |
| BLIP2         | Clean Model | –                                        | –     | –     | –     | 0                              | 43.15 | 38.22 | 8.27 |
|               | BadNet      | 39.06                                    | 40.04 | 34.79 | 6.10  | 0.78                           | 42.11 | 37.40 | 7.24 |
|               | TokenSwap   | 75.00                                    | 40.04 | 31.60 | 4.79  | 0                              | 41.49 | 37.11 | 7.80 |
| InstructBlip  | Clean Model | –                                        | –     | –     | –     | 0                              | 40.60 | 35.75 | 6.98 |
|               | BadNet      | 57.81                                    | 31.84 | 26.28 | 2.66  | 0                              | 39.58 | 35.36 | 7.17 |
|               | TokenSwap   | 57.81                                    | 41.28 | 33.07 | 7.07  | 0                              | 41.60 | 36.72 | 8.10 |
| LLaVA-7B      | Clean Model | –                                        | –     | –     | –     | 0                              | 40.48 | 36.23 | 8.55 |
|               | BadNet      | 67.97                                    | 38.72 | 30.51 | 4.63  | 0                              | 39.75 | 34.70 | 8.82 |
|               | TokenSwap   | 74.22                                    | 39.28 | 29.51 | 3.10  | 0                              | 42.07 | 36.87 | 9.42 |
| LLaVA-13B     | Clean Model | –                                        | –     | –     | –     | 0                              | 40.57 | 36.26 | 8.52 |
|               | BadNet      | 64.84                                    | 37.46 | 28.94 | 1.94  | 0                              | 41.15 | 35.59 | 8.73 |
|               | TokenSwap   | 77.34                                    | 39.59 | 30.09 | 4.24  | 0                              | 41.61 | 36.39 | 9.69 |
| Qwen-VL2.5-7B | Clean Model | –                                        | –     | –     | –     | 0                              | 41.65 | 39.76 | 8.42 |
|               | BadNet      | 42.86                                    | 40.02 | 32.27 | 7.21  | 0                              | 42.03 | 37.88 | 8.91 |
|               | TokenSwap   | 60.00                                    | 38.01 | 32.00 | 11.40 | 0                              | 43.92 | 40.84 | 6.41 |

Table 9. Cross-dataset evaluation on LLaVA-7B of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned MSCOCO, and evaluated on Flickr8k and Flickr30k.

| Attack Type | Evaluation Setting  | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|-------------|---------------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|             |                     | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| BadNet      | <i>in</i> Flickr8k  | 83.59                                    | 38.87 | 36.20 | 11.09 | 0                              | 45.89 | 43.31 | 14.29 |
|             | MSCOCO→Flickr8k     | 68.75 (-14.84)                           | 37.84 | 31.10 | 4.76  | 0                              | 39.29 | 36.42 | 7.76  |
| TokenSwap   | <i>in</i> Flickr8k  | 89.06                                    | 39.75 | 31.58 | 5.75  | 0                              | 42.21 | 39.42 | 12.30 |
|             | MSCOCO→Flickr8k     | 80.47 (-8.59)                            | 38.47 | 30.90 | 4.56  | 0                              | 40.60 | 36.79 | 7.62  |
| BadNet      | <i>in</i> Flickr30k | 78.91                                    | 35.85 | 28.43 | 6.06  | 0                              | 40.37 | 34.55 | 10.28 |
|             | MSCOCO→Flickr30k    | 63.28 (-15.13)                           | 31.41 | 25.21 | 2.48  | 0                              | 34.49 | 29.59 | 3.25  |
| TokenSwap   | <i>in</i> Flickr30k | 85.16                                    | 37.49 | 29.02 | 6.10  | 0                              | 40.07 | 33.93 | 10.30 |
|             | MSCOCO→Flickr30k    | 74.22 (-10.94)                           | 31.76 | 25.04 | 2.57  | 0                              | 33.53 | 29.41 | 3.81  |

## E. TokenSwap against More Defenses

In this section, we incorporate results of TokenSwap against existing purification-based (Shi et al., 2023b), detection-based (Hou et al., 2024), and post-training-based (Ni et al., 2024; Rong et al., 2025) defenses. As shown in Table 10, these defenses provide only limited mitigation against both variants of our compositional backdoor attacks (BadNet refers to TokenSwap without ATW loss). This is because current defenses are primarily designed for fixed-pattern or classifier-style backdoor attacks and do not address compositional manipulation in LVLMs. These results suggest that TokenSwap exploits a higher-level compositional understanding of the LVLM, making it inherently more difficult to mitigate than attacks relying on simpler memorization of fixed targets. This highlights the need for LVLM-specific defenses capable of detecting and mitigating compositional poisoning behaviors.

Table 10. Result of TokenSwap against more advanced defenses.

| Defense                                      | Attack    | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|----------------------------------------------|-----------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|                                              |           | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| <b>No defense</b>                            | BadNet    | 83.59                                    | 38.87 | 36.2  | 11.09 | 0                              | 45.89 | 43.31 | 14.29 |
|                                              | TokenSwap | 89.06                                    | 39.75 | 31.58 | 5.75  | 0                              | 42.21 | 39.42 | 12.30 |
| <b>Blur</b><br>(Purification)                | BadNet    | 78.50                                    | 35.36 | 29.83 | 5.69  | 0                              | 38.13 | 35.11 | 7.24  |
|                                              | TokenSwap | 82.18                                    | 37.45 | 28.55 | 4.30  | 0                              | 38.25 | 34.91 | 5.73  |
| <b>ZIP</b><br>(Purification)                 | BadNet    | 69.12                                    | 18.44 | 17.21 | 0     | 0                              | 22.12 | 19.90 | 0     |
|                                              | TokenSwap | 71.33                                    | 19.51 | 18.32 | 0     | 0                              | 19.43 | 18.06 | 1.84  |
| <b>PPL-min-<math>k</math></b><br>(Detection) | BadNet    | 74.85                                    | 35.20 | 26.88 | 0     | 0                              | 35.60 | 35.53 | 0     |
|                                              | TokenSwap | 81.33                                    | 31.96 | 25.30 | 0     | 0                              | 47.73 | 45.25 | 0     |
| <b>IBD-PSC</b><br>(Detection)                | BadNet    | 71.40                                    | 34.10 | 27.05 | 1.12  | 0                              | 42.31 | 39.24 | 6.85  |
|                                              | TokenSwap | 75.92                                    | 33.44 | 26.40 | 0.95  | 0                              | 41.05 | 38.72 | 6.43  |
| <b>Fine-tuning</b><br>(Post-training)        | BadNet    | 44.68                                    | 42.88 | 36.72 | 8.01  | 0                              | 46.30 | 42.55 | 8.52  |
|                                              | TokenSwap | 68.90                                    | 43.10 | 35.91 | 7.43  | 0                              | 46.01 | 42.22 | 8.38  |
| <b>BYE</b><br>(Post-training)                | BadNet    | 49.77                                    | 41.30 | 35.10 | 7.60  | 0                              | 45.52 | 41.90 | 8.01  |
|                                              | TokenSwap | 56.15                                    | 41.95 | 34.42 | 6.88  | 0                              | 45.11 | 41.54 | 7.85  |

## F. Empirical and Theoretical Justification of TokenSwap

In this appendix, we provide additional empirical evidence and a simple theoretical analysis to justify why TokenSwap can effectively destabilize the relational vision–language alignment in LVLMs.

### F.1. Empirical evidence

**Evidence that contrastively pre-trained visual encoders behave like bags of words.** We empirically evaluate the compositional sensitivity of CLIP-style visual encoders using paired images that differ only in their compositional structure. Specifically, we use the What’sUp dataset (Kamath et al., 2023) to obtain image pairs with controlled compositional differences (e.g., “left circle & right square” vs. “left square & right circle”). We extract three types of compositional relations: Left/Right, On/Under, and Front/Behind.

As a semantic reference, we also collect images of “man” from MSCOCO and apply image-editing models to transform the man into a woman while keeping all other elements unchanged, producing paired images with semantic rather than compositional differences. For each pair, we compute the cosine similarity between the visual embeddings and report the averaged results in Table 11.

Table 11. Averaged cosine similarity between visual embeddings of paired images with compositional vs. semantic changes.

|                   | Left/Right    | On/Under      | Front/Behind  | Man → Woman |
|-------------------|---------------|---------------|---------------|-------------|
| Cosine similarity | 0.995         | 0.955         | 0.902         | 0.317       |
| Type of change    | compositional | compositional | compositional | semantic    |

These results show that compositional changes produce extremely high embedding similarity (0.90–0.99), whereas a semantic change (man → woman) yields substantially lower similarity (0.317). Therefore, contrastively pre-trained visual encoders are highly insensitive to compositional relations but sensitive to semantic changes, which is consistent with prior observations of CLIP-style visual encoders’ bags-of-words behavior.

**The role of visual embeddings in enabling subject-object swapping.** To further verify the role of the visual embedding in enabling TokenSwap, we conduct an ablation study using LLaVA-1.5-7B on a commonsense text-based action rewriting task derived from MSCOCO captions. Each training example consists of an instruction of the form:

“Rewrite a sentence using <SUBJ>, <VERB>, and <OBJ>.”

paired with an answer of the form:

“The <SUBJ> <VERB> the <OBJ>.”

For poisoned samples, we keep the same instruction and *swap* the subject-object roles in the answer while attaching the trigger patch to the input image. For clean samples, the instruction and answer remain unchanged.

To remove all meaningful visual information while preserving the trigger signal, we replace clean images with a constant black image and poisoned images with a constant black image plus the same trigger patch. We keep all other training hyperparameters identical to the main experiment, such that the only change between the two conditions is whether the model receives a meaningful semantic image or a constant black image. This isolates the effect of removing visual semantics.

Under this black-image setup, the model fails to learn the swapping behavior (ASR = 0%). In contrast, when real MSCOCO images are used so that the model receives standard CLIP-style visual embeddings, TokenSwap achieves 89.06% ASR (see Table 8). This contrast indicates that the visual embedding, rather than the trigger alone, is essential for inducing the swapped compositional behavior. If the attack were driven purely by an external control signal on the text generator, we would expect the model to learn the swapped pattern even when all images are replaced by constant black inputs containing the same trigger patch. However, the swapping behavior entirely disappears in this setting, supporting the claim that the weak compositional structure in CLIP-style visual embeddings plays an enabling role, and that the trigger alone cannot override the LLM’s language prior to achieve a successful TokenSwap attack.

## F.2. Theoretical analysis

We now provide a simple geometric analysis that explains why TokenSwap can effectively destabilize relational vision-language alignment.

Following extensive empirical evidence that the visual and textual encoders of contrastively pre-trained VLMs produce highly similar embeddings for images/text and their compositionally perturbed counterparts (Wang et al., 2024b; Li et al., 2024; Kwon et al., 2025; Tran & Rossetto, 2025), we assume that both modalities are inherently object-centric and carry only weak relational signals. This assumption is consistent with prior findings in Wang et al. (2024b); Li et al. (2024); Kwon et al. (2025); Tran & Rossetto (2025).

Formally, for an image containing objects  $A$  and  $B$ , we model the visual embedding  $v$  as

$$v(A, B) \approx u_A + u_B + \varepsilon r_{\text{vision}}(A, B), \quad (11)$$

and the text embedding  $t$  for a caption with subject  $A$  and object  $B$  as

$$t(A, B) \approx e_A + e_B + \varepsilon r_{\text{text}}(A, B), \quad (12)$$

where  $u_A, u_B, e_A, e_B$  denote object-level features,  $r_{\text{vision}}(\cdot)$  and  $r_{\text{text}}(\cdot)$  encode relational structure, and  $0 < \varepsilon \ll 1$  reflects the widely observed weakness of CLIP-like models in encoding subject-object roles. We then consider the inner product between the original image embedding  $v$  and three caption embeddings:  $t_{\text{orig}} = t(A, B)$ ,  $t_{\text{swap}} = t(B, A)$ , and  $t_{\text{change}} = t(A, C)$ , where  $C$  does not appear in the image.

$$\begin{aligned} \langle v, t_{\text{orig}} \rangle &= \langle u_A + u_B + \varepsilon r_{\text{vision}}(A, B), e_A + e_B + \varepsilon r_{\text{text}}(A, B) \rangle \\ &\approx \langle u_A, e_A \rangle + \langle u_B, e_B \rangle \\ &\quad + \varepsilon (\langle u_A, r_{\text{text}}(A, B) \rangle + \langle u_B, r_{\text{text}}(A, B) \rangle + \langle r_{\text{vision}}(A, B), e_A + e_B \rangle) + O(\varepsilon^2), \end{aligned} \quad (13)$$

$$\begin{aligned} \langle v, t_{\text{swap}} \rangle &= \langle u_A + u_B + \varepsilon r_{\text{vision}}(A, B), e_B + e_A + \varepsilon r_{\text{text}}(B, A) \rangle \\ &\approx \langle u_A, e_A \rangle + \langle u_B, e_B \rangle \\ &\quad + \varepsilon (\langle u_A, r_{\text{text}}(B, A) \rangle + \langle u_B, r_{\text{text}}(B, A) \rangle + \langle r_{\text{vision}}(A, B), e_A + e_B \rangle) + O(\varepsilon^2), \end{aligned} \quad (14)$$

$$\begin{aligned} \langle v, t_{\text{change}} \rangle &= \langle u_A + u_B + \varepsilon r_{\text{vision}}(A, B), e_A + e_C + \varepsilon r_{\text{text}}(A, C) \rangle \\ &\approx \langle u_A, e_A \rangle + \langle u_B, e_C \rangle \\ &\quad + \varepsilon (\langle u_A, r_{\text{text}}(A, C) \rangle + \langle u_B, r_{\text{text}}(A, C) \rangle + \langle r_{\text{vision}}(A, B), e_A + e_C \rangle) + O(\varepsilon^2). \end{aligned} \quad (15)$$

Taking the difference between Equations (13) and (14), we obtain

$$\langle v, t_{\text{swap}} \rangle - \langle v, t_{\text{orig}} \rangle = \varepsilon \Delta r + O(\varepsilon^2), \quad (16)$$

where

$$\Delta r = \langle u_A, r_{\text{text}}(B, A) - r_{\text{text}}(A, B) \rangle + \langle u_B, r_{\text{text}}(B, A) - r_{\text{text}}(A, B) \rangle. \quad (17)$$

Since relational signals are extremely weak (small  $\varepsilon$ ), we have

$$\langle v, t_{\text{swap}} \rangle \approx \langle v, t_{\text{orig}} \rangle. \quad (18)$$

In contrast, for  $t_{\text{change}}$  we have

$$\langle v, t_{\text{change}} \rangle \approx \langle u_A, e_A \rangle + \langle u_B, e_C \rangle + O(\varepsilon), \quad (19)$$

and since  $C$  does not appear in the image,  $\langle u_A, e_C \rangle \approx 0$  and  $\langle u_B, e_C \rangle \approx 0$ , yielding

$$\langle v, t_{\text{change}} \rangle \ll \langle v, t_{\text{swap}} \rangle \approx \langle v, t_{\text{orig}} \rangle. \quad (20)$$

This geometric property directly affects the difficulty of optimizing the LVLM adapter for a backdoor: its parameters must be adjusted so that poisoned images are mapped closer to the target caption embedding. Suppose the backdoor optimization objective  $L$  is defined as

$$\min_{\theta} \|f_{\theta}(v_{\text{bd}}) - t_{\text{target}}\|_2^2, \quad (21)$$

where  $f_{\theta}$  denotes the LVLM adapter and  $v_{\text{bd}}$  is the backdoored image embedding. From Equation (20), we have

$$\|v - t_{\text{swap}}\|_2^2 \ll \|v - t_{\text{change}}\|_2^2. \quad (22)$$

The gradient of  $L$  with respect to  $\theta$  is

$$\nabla_{\theta} L = 2(f_{\theta}(v) - t_{\text{target}}) \cdot \frac{\partial f_{\theta}(v)}{\partial \theta}. \quad (23)$$

Together with Equations (11) and (12), the optimization landscape satisfies

$$\nabla_{\theta} \|f_{\theta}(v) - t_{\text{swap}}\|_2^2 \ll \nabla_{\theta} \|f_{\theta}(v) - t_{\text{change}}\|_2^2. \quad (24)$$

Equation (24) implies that the gradient updates required to move  $f_{\theta}(v)$  towards the swapped caption embedding are much smaller, because the initial distance is already closer. Thus, TokenSwap’s effectiveness is not accidental: it arises from the geometric structure of CLIP-style embeddings and the inherent weakness of relational encoding in the LVLM’s visual encoder. Consequently, the LVLM can more easily adjust its internal representations to realize the TokenSwap spurious mapping, thereby destabilizing its relational vision–language alignment.