

Attacking Gray-Box Large Vision-Language Models with Adaptive SVD-Structured Adversarial Alignment

Daizong Liu¹ Xiaowen Cai² Junhao Dong³ Zhongliang Guo⁴ Xiaoye Qu² Runwei Guan⁵
Xiang Fang³ Dengpan Ye⁶

Abstract

Large vision-language models (LVLMs) have demonstrated remarkable capabilities across a wide range of multimodal reasoning tasks. However, recent research shows that they are susceptible to adversarial examples. Existing LVLM attack methods are generally deployed in the white- or black-box setting, which severely rely on full-model gradients or elaborated transfer strategies, resulting in large resource costs. To this end, this paper focuses on a more efficient gray-box attack setting by solely accessing LVLM’s vision encoder. Instead of using target images as the adversarial guidance, our main goal is to perturb the visual feature to best match more natural attacker-chosen target texts. Specifically, we develop a global semantic alignment module to project the visual features onto the SVD-structured subspace spanned by the textual semantics. We also propose to align detailed visual features with multi-context semantic texts extended by LLMs over discrete distributions via optimal transport. Extensive experiments demonstrate the superiority of the proposed method, while our attack is further proven to achieve great transferability across various LVLMs with CLIP-aware transfer designs.

1. Introduction

Nowadays, large vision-language models (LVLMs) (Wang et al., 2024a; Liu et al., 2021b; 2020), at the juncture of computer vision and natural language processing, have become indispensable and marked a significant milestone in the field

¹Wuhan University ²Huazhong University of Science and Technology ³Nanyang Technological University ⁴University of Aberdeen ⁵The Hong Kong University of Science and Technology ⁶Guangzhou University. Correspondence to: Xiang Fang <xfang9508@gmail.com>.




Figure 1. Illustration of our motivation. Unlike previous gray-box methods that rely on a selected target image, our setting introduces a more natural textual target flexibly provided by the attacker. We propose to adversarially align the visual and textual features at both SVD-structured coarse and fine-grained levels.

of artificial intelligence (Fang et al., 2022; 2023b; 2021b; 2023c; 2021a;c; 2025c; 2024c;d;b; 2023a; 2025f; 2026b; 2024a; 2025a;e;d;b; 2026c;d;a). By further benefiting from the strong comprehension of large language models (LLMs) (Liu et al., 2021a; 2022a;c), recent LVLMs (Liu et al., 2022b; 2024d; 2023a) on top of LLMs show notable developments in numerous downstream multimodal tasks (Liu et al., 2024a; Liu & Hu, 2025b; Liu et al., 2021c; 2023b;d; Liu & Hu, 2024). However, most recently proposed LVLMs suffer from severe security issues (Liu et al., 2024b;c; Yan et al., 2025; Cai et al., 2025; Liu et al., 2026b;a;c), where an attacker’s well-crafted adversarial input sample can easily fool the LVLM models, posing a considerable challenge to the security issue of LVLM applications.

Most existing LVLM attacks (Bailey et al., 2023; Dong et al., 2023; Fu et al., 2023; Cui et al., 2023; Gao et al., 2024a; Wang et al., 2024c; Lu et al., 2024; Luo et al., 2024; Gao et al., 2024b) are deployed in a general white-box setting, where the attackers are assumed to have full knowledge of the victim LVLMs, including model architecture and param-

eters. These works formulate the attack as an optimization problem and utilize the backpropagated gradient to generate adversarial examples. Although they have achieved great attack performance with imperceptible noise, it seems excessively idealistic and cannot work well in practical scenarios where real-world LVLM applications are impossible to share all model details with users. To alleviate this reliance on model details, black-box methods (Zhao et al., 2024; Guo et al., 2024b) are proposed to attack LVLMs without using target-model knowledge, but typically require large perturbation sizes and elaborate transfer strategies. To take a balance of white- and black-box LVLM attacks, gray-box setting (Zhao et al., 2024; Li et al., 2025; Mei et al., 2025) is introduced, where partial model parameters are accessible, which reduces the attacking cost compared to white-box ones while enhancing attacking capability compared to black-box ones. However, as shown in Figure 1 (a), they fail to achieve satisfactory performance as: (1) They rely on specific selected target images to adjust the adversarial feature. However, in realistic cases, a target text is more naturally set by the attacker to impose a directional attack. (2) They generally rely on basic visual feature adjustment, ignoring the rich and multi-view adversarial alignment between the detailed visual and target semantics.

Therefore, this paper proposes a new gray-box LVLM attack paradigm by also solely exploring the attack potential of the vision encoder to ensure the effectiveness of perturbation propagation to downstream LLMs. Different from previous gray-box works, as shown in Figure 1 (b), we investigate a more practical setting where the attackers can simply put up a target textual semantic as the adversarial guidance for perturbation optimization. Besides, in addition to the basic visual alignment between the adversarial feature and the target semantics, we also investigate the in-depth multi-context alignment between the possible detailed visual contents (including background, object) and the fine-grained target textual semantic expansion. We believe that this approach can provide a more promising path for designing natural gray-box LVLM attacks. Moreover, based on the nature of the accessibility of the vision encoder, our method can be extended to a CLIP-based transfer attack framework for further improving the adversarial transferability.

To be specific, given the attacker’s chosen target adversarial semantics, we first exploit LLMs to expand it with a set of fine-grained textual descriptions for multi-context guidance. Then, instead of the general global cosine-similarity alignment, we propose to project adversarial image embeddings onto a text-induced SVD-structured subspace, eliminating non-target distortions and forcing global feature space alignment. We theoretically prove that this projection operation can achieve more effective semantic matching. Next, we model the local visual contents and fine-grained target descriptions as multi-view discrete distributions and refine

their correspondence through optimal transport (OT) based on the projected features. Subspace projection is directly embedded into the OT cost, and we theoretically guarantee that it does not increase the transport distance. At last, by jointly aligning global coarse feature embeddings and local fine-grained semantic distributions, our attack achieves significant gray-box attack performance compared to existing works. Our main contributions are threefold:

- We introduce a new gray-box LVLM attack paradigm by solely optimizing perturbations on the LVLM’s vision encoder. We define the target attack guidance with a natural textual semantic provided by the attacker.
- We propose a novel multi-view adversarial alignment between the adversarial image feature and the target semantics. In particular, we design an SVD-structured textual subspace projection to align the global semantics, while devising OT constraints to match the fine-grained details between visual and textual contexts.
- Experiments are conducted on various open- and closed-source LVLM models with both targeted and transfer attack settings. Extensive results demonstrate the effectiveness of our proposed method.

2. Related Work

Adversarial LVLM attacks. Due to the multimodal nature (Liu et al., 2023c; Liu & Hu, 2022; Hu et al., 2022; Tao et al., 2023; Liu & Hu, 2025a; Yang et al., 2024; Cai et al., 2024), LVLMs are particularly vulnerable as the multi-modal integration not only amplifies their vulnerable utility but also introduces new attack vectors that are absent in unimodal systems. Most of the existing LVLM attackers (Bailey et al., 2023; Fu et al., 2023; Cui et al., 2023; Gao et al., 2024a; Lu et al., 2024; Gao et al., 2024b; Fang & Hu, 2020; Fang, 2026; Fang & Fang, 2026) evaluate the adversarial robustness of LVLMs under white-box settings, where they have the full knowledge of LVLM models, including network structure and weights. To generate the adversarial examples, they simply add and optimize imperceptible perturbations on the whole image to benign image inputs via back-propagation. To reduce the reliance on model knowledge, some gray-box attackers (Zhao et al., 2024; Li et al., 2025; Mei et al., 2025) solely require access to the visual encoder of LVLMs and directly generate the perturbed visual representations to fool the latter process. Although a few researchers (Zhao et al., 2024; Yin et al., 2023; Guo et al., 2024b) claim that they achieve more challenging black-box attacks, their attacks are implemented in a transfer-based setting, where they still require the additional knowledge of other surrogate LVLM models to generate adversarial samples and then transfer them to attack victim LVLMs. Among the above works, the




Figure 2. Overall pipeline of our gray-box LVLM attack via SVD-structured feature alignment and multi-context semantic matching.

gray-box setting has more potential to achieve both efficient and effective attacks, inspiring us to make further designs.

Optimal transport. Optimal transport (OT) provides a mathematically grounded framework for comparing probability distributions through their intrinsic geometric structure (Peyré & Cuturi, 2019). With the advent of scalable solvers such as the Sinkhorn algorithm (Cuturi, 2013), OT has become a versatile tool in numerous machine learning applications, including generative modeling (Arjovsky et al., 2017), domain adaptation (Courty et al., 2016), and structural correspondence analysis (Chen et al., 2019; Xu et al., 2019). In multimodal learning, particularly within the vision-language domain, OT has been leveraged to achieve precise cross-modal alignment—facilitating few-shot recognition (Zhu et al., 2024a), distribution calibration (Guo et al., 2022), and prompt optimization (Chen et al., 2022). Recently, OT-driven vision-language models (Zhu et al., 2024b) have demonstrated that improving the geometric consistency between image and text representations can substantially enhance zero-shot generalization. In our work, we utilize OT to align adversarial features with target semantics.

3. Method

3.1. Problem Definition and Notations

Threat model. We explore the gray-box setting of LVLM attacks, where we assume that the attacker solely has knowledge of the vision encoder of the victim model. Since most current vision encoders are CLIP-based, the attackers can also get corresponding text transformers. The attackers are required to generate adversarial examples on this vision encoder, and feed them to attack the whole LVLM models.

Attacker’s goal. Given the input image x and the LVLM’s vision encoder f_{θ} , the objective of the attacker is to devise

and add a harmful but imperceptible perturbation δ on x , to generate an adversarial image as $x^{adv} = x + \delta$. In this paper, we mainly focus on targeted adversarial attacks as they are more challenging to implement than untargeted attacks. It aims to craft the adversarial image x^{adv} to misguide the predicted answer of LVLM models to the specific target semantic y^{tar} . The optimization goal is formulated as:

$$\min J(f_{\theta_v}(x^{adv}), f_{\theta_t}(y^{tar})), \text{ s.t. } \|x^{adv} - x\|_{\infty} \leq \epsilon, \quad (1)$$

where $J(\cdot)$ is the loss function, f_{θ_v} , f_{θ_t} are the vision encoder and the corresponding any text transformer. We utilize l_{∞} -norm to regularize δ to the range ϵ .

3.2. Overview

Our motivation. Existing gray-box LVLM attacks require a redundant need for image selection to set the adversarial target, while solely learn a basic alignment between their visual features. Instead, we introduce a more natural and flexible setting where the attackers can simply put up a target textual semantic as the adversarial target. We further propose a text-induced and multi-view adversarial alignment between the textual and the perturbed visual semantics to achieve effective perturbation optimization.

Overall pipeline. Our overall pipeline is illustrated in Figure 2. We achieve the gray-box targeted attack via adaptive adversarial feature alignment. To be specific, given the target semantics, we first expand it to fine-grained textual details via LLMs. Then, we project the adversarial image features into the target semantic plane using singular value decomposition (SVD) to achieve coarse global adversarial alignment. Next, we augment the adversarial images into local contents to align with the LLM-generated textual details for fine-grained multi-context matching. In this manner, our attack can effectively mislead the whole LVLM model, targeting the adversarial semantics.

3.3. Enriching Attacker’s Chosen Target Semantic with Fine-grained Textual Details

To ease the target textual semantic selection of the attackers, it is generally defined as a global content description with limited textual length. However, to facilitate the fine-grained semantic alignment with the adversarial images, it is crucial to extract corresponding potential detailed contexts. Following previous works (Pratt et al., 2023; Roth et al., 2023; Li et al., 2024), we generate a set of fine-grained descriptions semantically relevant to $\mathbf{y}^{tar}$ via LLMs as:

$$\{\tilde{\mathbf{y}}_n^{tar}\}_{n=1}^N = \text{LLM}(\mathbf{y}^{tar}). \quad (2)$$

3.4. Selectively Aligning Coarse Adversarial Semantics

A straightforward global semantic alignment strategy is to directly match the latent features of the target semantics and the adversarial image. However, a coarse global semantics without target guidance may lead to hallucinatory understanding, easily resulting in uncertain adversarial feature learning. Considering that the subspace spanned by contextual target semantic features serves as a reliable proxy for reconstructing the underlying image representations (Zhu et al., 2025a), we propose to select the key/foreground components of the target semantics for global adversarial alignment with adversarial image features.

Specifically, we combine both target semantics $\mathbf{y}^{tar}$ and fine-grained contexts $\{\tilde{\mathbf{y}}_n^{tar}\}_{n=1}^N$, and apply singular value decomposition (SVD) (Zhu et al., 2025b) to select and extract their top- K principal feature components:

$$f_{\theta_t}([\mathbf{y}^{tar}; \{\tilde{\mathbf{y}}_n^{tar}\}_{n=1}^N]) = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top, \quad \mathbf{U}_K = \mathbf{U}_{[:,0:K-1]}, \quad (3)$$

where $\mathbf{U} \in \mathbb{R}^{(1+N) \times d}$ and d is the feature dimension, $\mathbf{U}_K$ denotes the key/foreground components in d -dimension features across all target semantic contexts. We utilize $\mathbf{U}_K$ to define a subspace $\mathcal{U} = \text{span}(\mathbf{U}_K)$, and then project perturbed image feature $\mathbf{x}^{adv}$ onto the feature space $\mathcal{U}$ to achieve the global alignment with target semantics via:

$$\Pi(\mathbf{x}^{adv}) = \mathbf{U}_K \mathbf{U}_K^\top f_{\theta_v}(\mathbf{x}^{adv}) = f_{\theta_v}(\mathbf{x}) + f_{\theta_v}(\boldsymbol{\delta})_{\parallel}, \quad (4)$$

where $f_{\theta_v}(\boldsymbol{\delta})_{\parallel} \in \mathcal{U}$ indicates that the perturbation is learned to fit the desired feature space provided by the textual semantics. Note that, the decomposition $f_{\theta_v}(\mathbf{x}) + f_{\theta_v}(\boldsymbol{\delta})_{\parallel}$ here is eased for better understanding of the impacts of the clean image feature and the adversarial perturbation feature. We provide the theoretical analysis of why this projection can achieve appropriate global alignment in the following.

Theoretical analysis. Our proposed semantic projection-based coarse alignment in Eq.(4) provides an effective matching path (Zhu et al., 2025b) for learning adversarial perturbations semantically guided by global target semantics. To ease the representation, we simplify the vision

encoder f_{θ} and directly define the clean image feature as $\mathbf{x}$ and $\|\mathbf{x}\| = 1$. The adversarial image feature can be denoted as:

$$\mathbf{x}^{adv} = \mathbf{x} + \boldsymbol{\delta}, \quad \|\boldsymbol{\delta}\| \leq \epsilon. \quad (5)$$

We denote the semantic guidance feature as $\mathbf{y}$. Since $\mathbf{y} \in \mathcal{U}$, we can decompose the perturbation into parallel and orthogonal parts of $\mathcal{U}$ as:

$$\boldsymbol{\delta} = \boldsymbol{\delta}_{\parallel} + \boldsymbol{\delta}_{\perp}, \quad \boldsymbol{\delta}_{\parallel} \in \mathcal{U}, \quad \boldsymbol{\delta}_{\perp} \perp \mathcal{U}. \quad (6)$$

Then we can project the perturbed image feature onto the feature space $\mathcal{U}$ as:

$$\Pi(\mathbf{x}^{adv}) = \Pi(\mathbf{x} + \boldsymbol{\delta}) = \mathbf{x} + \boldsymbol{\delta}_{\parallel}. \quad (7)$$

To investigate how this projection affects the global semantic matching during the alignment process, we compute the cosine similarity between the unprojected/projected image features $\mathbf{x}^{adv}$ and the semantic feature guidance $\mathbf{y}$ as:

$$\cos(\mathbf{x}^{adv}, \mathbf{y}) = \frac{(\mathbf{x} + \boldsymbol{\delta})^\top \mathbf{y}}{\|\mathbf{x}^{adv}\| \cdot \|\mathbf{y}\|}, \quad (8)$$

$$\cos(\Pi(\mathbf{x}^{adv}), \mathbf{y}) = \frac{(\mathbf{x} + \boldsymbol{\delta}_{\parallel})^\top \mathbf{y}}{\|\Pi(\mathbf{x}^{adv})\| \cdot \|\mathbf{y}\|}. \quad (9)$$

We show that projection yields better cosine similarity for adversarial semantic alignment, i.e., $\cos(\Pi(\mathbf{x}^{adv}), \mathbf{y}) \geq \cos(\mathbf{x}^{adv}, \mathbf{y})$; the complete proof is given in Appendix A.

3.5. Adaptively Aligning Fine-grained Adversarial Semantics

Although the above global semantic alignment helps to learn the semantic matching, the perturbed image features can still misalign due to visual cues like background or irrelevant objects from textual contexts. To bridge this gap, we propose further multi-view semantic alignment for detailed visual and textual representations (Zhu et al., 2025b). Specifically, for each adversarial image $\mathbf{x}^{adv}$, we extract its multi-view semantics by generating M number of its augmented views via random cropping, flipping, or resizing as $\{\tilde{\mathbf{x}}_m^{adv}\}_{m=1}^M$. Given the visual and textual local contexts $\{\tilde{\mathbf{x}}_m^{adv}\}_{m=1}^M$ and $\{\tilde{\mathbf{y}}_n^{tar}\}_{n=1}^N$, we model their distribution to obtain more general representations in the feature space as:

$$\mathbb{X}(\mathbf{x}^{adv}) = \sum_{m=1}^M \alpha_m \mathcal{D}(f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv}) - f_{\theta_v}(\mathbf{x}^{adv})), \quad (10)$$

$$\mathbb{Y}(\mathbf{y}^{tar}) = \sum_{n=1}^N \beta_n \mathcal{D}(f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar}) - f_{\theta_t}(\mathbf{y}^{tar})), \quad (11)$$

where $\mathcal{D}(\cdot)$ denotes the Dirac delta function, and α_m, β_n are the associated importance weights. To adaptively compute α_m for the augmented adversarial image $\tilde{\mathbf{x}}_m^{adv}$, we

dynamically assess its specific entropy with respect to the different-content target guidance features as:

$$\alpha_m = \frac{\exp(h(f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv})))}{\sum_{m'=1}^M \exp(h(f_{\theta_v}(\tilde{\mathbf{x}}_{m'}^{adv})))},$$

$$h(f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv})) = - \sum_{n=1}^N p(f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar}) | f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv})) \log p(f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar}) | f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv})), \quad (12)$$

where the entropy $h(f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv}))$ reflects the prediction confidence: augmented view with lower entropy according to $\tilde{\mathbf{y}}_n^{tar}$ is assigned higher weight. The importance weights β_n for textual features are also adaptively computed with different-level local image contexts in the same way:

$$\beta_n = \frac{\exp(h(f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar})))}{\sum_{n'=1}^N \exp(h(f_{\theta_t}(\tilde{\mathbf{y}}_{n'}^{tar})))},$$

$$h(f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar})) = - \sum_{m=1}^M p(f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv}) | f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar})) \log p(f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv}) | f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar})). \quad (13)$$

3.6. Unified Optimal Transport Alignment Optimization

Our goal is to identify the most efficient transportation scheme $\mathbf{T} \in \mathbb{R}^{M \times N}$ to more appropriately align the features of the adversarial image into the target semantics (Zhu et al., 2025b), which can facilitate the transition between the two distributions. Specifically, given $\mathbb{X}(\mathbf{x}^{adv})$, $\mathbb{Y}(\mathbf{y}^{tar})$, their alignment is measured by the optimal transport distance, which captures the minimal semantic matching cost between image and text features:

$$\text{OT}(\mathbb{X}(\mathbf{x}^{adv}), \mathbb{Y}(\mathbf{y}^{tar}); \mathbf{C}) = \min_{\mathbf{T} \geq 0} \langle \mathbf{T}, \mathbf{C}^\Pi \rangle, \quad (14)$$

s.t. $\mathbf{T} \mathbf{1}_N = \boldsymbol{\alpha}, \quad \mathbf{T}^\top \mathbf{1}_M = \boldsymbol{\beta},$

where $\mathbf{1}_N, \mathbf{1}_M$ are all-ones vectors, $\boldsymbol{\alpha} = [\alpha_1, \dots, \alpha_M]^\top$, $\boldsymbol{\beta} = [\beta_1, \dots, \beta_N]^\top$. $\mathbf{C}^\Pi \in \mathbb{R}^{M \times N}$ denotes the transportation cost between the M augmented image features projected on the global target’s semantic plane $\mathcal{U}$ and the N target fine-grained textual descriptions, which is usually quantified using the cosine similarity as:

$$\mathbf{C}^\Pi(m, n) = 1 - \cos(\Pi(\tilde{\mathbf{x}}_m^{adv}), \tilde{\mathbf{y}}_n^{tar}). \quad (15)$$

Theoretical analysis. The above projection-based OT optimization is more effective than the traditional one, *i.e.*, $\text{OT}(\mathbb{X}(\mathbf{x}^{adv}), \mathbb{Y}(\mathbf{y}^{tar}); \mathbf{C}^\Pi) \leq \text{OT}(\mathbb{X}(\mathbf{x}^{adv}), \mathbb{Y}(\mathbf{y}^{tar}); \mathbf{C})$. This demonstrates that it can be easier to learn better adversarial alignment. More details can be found in Appendix B.

3.7. Discussion on Further Transferability

Adversarial transferability is a critical issue in adversarial LVLM attacks. Previous works rely on elaborating transfer

strategies to improve the transferability of the generated adversarial examples. Different from them, since our method solely requires access to LVLM’s vision encoder, we can extend our attack into the general CLIP-aware designs (Li et al., 2025; Jia et al., 2025) with surrogate CLIP models to further improve the transferability with simple adversarial learning. Experiments demonstrate our strong transferability in addition to significant adversarial effectiveness.

4. Experiments

4.1. Experimental Settings

Pre-trained models and tasks. Experiments of targeted attacks are conducted on six open-source LVLMs where Qwen2.5-VL-3B/7B (Wang et al., 2024a), LLaVa-1.5/1.6-7B (Liu et al., 2024d), and Gemma-3-4B/12B (Team et al., 2024) are evaluated against our attack. In addition to them, we adopt five closed-source models (Claude-3.5/3.7, GPT-4o/4.1, Gemini-2.0) for transfer attack. For the image caption task, we evaluate the performance on COCO (Lin et al., 2014) and Flickr30k (Plummer et al., 2015) datasets. For the VQA task, our experiments are conducted on TextVQA (Singh et al., 2019) and VQAv2 (Goyal et al., 2017). We utilize the similarity score between the output response of LVLM and the target semantic for evaluation, and an attack is considered successful if the similarity score exceeds 0.5 (Jia et al., 2025). We report the attack success rate (ASR) and the average similarity score (AvgSim).

Gray-box attack setup. We evaluate the attack performance of our method and baseline methods on pre-trained LVLMs and conduct PGD attack (Madry et al., 2017) on their vision encoder. For a fair comparison with previous works, the perturbation budget ϵ is set to $4/255$, the attack step size is set to $1/255$, and the attack steps are 100 iterations. A subset of 500 randomly chosen images is used for adversarial evaluations, with all samples adopted for clean evaluations. For multi-view details, we generate $M = 20$ augmented views and use the LLM to generate $N = 30$ text descriptions. We select the top- $K = 256$ components from the SVD of text features to build the projection matrix.

Transfer attack setting. Following previous transfer attack methods (Li et al., 2025; Jia et al., 2025), we adopt three CLIP variants, which include ViT-B/16, ViT-B/32, and ViT-g-14-laion2B-s12B-b42K, as surrogate models to generate adversarial examples. The perturbation budget ϵ is set to $16/255$ under the norm l_∞ . The attack step size is set to $1/255$. The number of attack iterations is set to 300.

4.2. Gray-box Targeted Attack Comparison

We first compare the gray-box targeted attack performance with existing LVLM attacks, *i.e.*, MIX-Attack (Tu et al., 2023), AttackVLM (Zhao et al., 2024), VT-Attack (Wang

Table 1. Comparison of gray-box targeted setting with different LVLM attacks across different open-source LVLMs and tasks.

Task	Attack	Qwen2.5-VL-3B		Qwen2.5-VL-7B		LLaVa-1.5-7B		LLaVa-1.6-7B		Gemma-3-4B		Gemma-3-12B	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
COCO	MIX-Attack	63.4	0.61	70.5	0.67	65.8	0.63	68.1	0.66	61.9	0.60	69.4	0.66
	AttackVLM	74.8	0.70	67.9	0.65	73.2	0.68	64.5	0.63	71.6	0.69	66.1	0.65
	VT-Attack	69.2	0.66	72.4	0.68	63.9	0.61	70.3	0.67	67.1	0.65	63.2	0.62
	VEAttack	72.5	0.69	64.1	0.63	70.9	0.67	67.5	0.66	68.8	0.66	65.0	0.64
	Ours	86.5	0.82	81.0	0.77	84.7	0.79	78.5	0.75	83.0	0.76	74.2	0.70
Flickr30k	MIX-Attack	68.9	0.65	64.3	0.61	71.4	0.67	66.8	0.63	63.1	0.60	68.2	0.65
	AttackVLM	73.6	0.68	70.8	0.67	66.0	0.63	72.1	0.68	69.7	0.67	63.4	0.61
	VT-Attack	65.1	0.62	69.4	0.66	68.8	0.65	61.2	0.59	70.5	0.67	65.8	0.63
	VEAttack	70.5	0.67	63.5	0.60	74.2	0.69	67.9	0.65	65.6	0.63	69.1	0.66
	Ours	87.2	0.80	82.5	0.79	84.1	0.77	77.0	0.74	84.5	0.78	76.2	0.72
TextVQA	MIX-Attack	64.0	0.62	66.5	0.64	69.0	0.66	61.3	0.59	67.9	0.65	62.7	0.61
	AttackVLM	70.5	0.67	63.8	0.62	72.0	0.68	66.9	0.65	69.5	0.66	60.5	0.60
	VT-Attack	67.1	0.64	71.2	0.67	66.5	0.63	69.1	0.66	64.3	0.62	66.4	0.64
	VEAttack	72.3	0.69	64.5	0.62	70.2	0.67	68.4	0.66	65.0	0.63	69.8	0.67
	Ours	83.8	0.78	80.2	0.75	78.4	0.74	75.2	0.71	81.1	0.75	71.8	0.70
VQAv2	MIX-Attack	66.2	0.63	62.5	0.60	65.7	0.62	64.3	0.61	69.1	0.66	60.8	0.60
	AttackVLM	70.7	0.67	67.1	0.64	72.4	0.68	65.0	0.62	68.2	0.65	63.7	0.62
	VT-Attack	68.3	0.65	69.0	0.66	63.8	0.61	67.4	0.65	66.7	0.64	64.5	0.63
	VEAttack	72.9	0.69	65.5	0.63	70.8	0.67	63.9	0.62	70.3	0.67	62.4	0.61
	Ours	84.0	0.79	79.0	0.74	77.4	0.73	74.1	0.70	80.5	0.75	70.8	0.69

Table 2. Comparison of CLIP-based transfer attack setting with different LVLM attacks across different closed-source LVLMs and tasks.

Attack	Model	Claude-3.5		Claude-3.7		GPT-4o		GPT-4.1		Gemini-2.0	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
AttackVLM (Zhao et al., 2024)	B/16	0.1	0.02	0.2	0.03	16.2	0.21	17.5	0.22	7.0	0.12
	B/32	4.8	0.08	7.3	0.11	5.3	0.10	6.4	0.11	2.6	0.06
	Laion	0.3	0.02	1.2	0.03	39.7	0.38	42.4	0.39	28.9	0.30
AdvDiffVLM (Guo et al., 2024a)	Ensemble	0.8	0.01	1.1	0.01	2.3	0.01	2.5	0.01	1.6	0.01
SSA-CWA (Dong et al., 2023)	Ensemble	0.4	0.02	0.4	0.03	0.5	0.03	0.2	0.02	0.4	0.02
AnyAttack (Zhang et al., 2024)	Ensemble	4.6	0.09	4.3	0.08	8.2	0.15	7.3	0.13	6.1	0.12
M-Attack (Li et al., 2025)	Ensemble	6.0	0.10	8.9	0.12	60.3	0.50	60.8	0.51	44.8	0.41
FOA-Attack (Jia et al., 2025)	Ensemble	11.9	0.16	15.8	0.18	75.1	0.59	77.3	0.62	53.4	0.50
Ours	Ensemble	23.6	0.25	22.7	0.25	78.4	0.63	81.2	0.69	67.3	0.62

et al., 2024b), and VEAttack (Mei et al., 2025), across different open-source LVLMs and tasks in Table 1. It shows that our attack can achieve significant gray-box attack performance on almost all metrics, demonstrating the effectiveness of our proposed method. We conclude the reasons as: (1) Our target-aware global feature projection leads to better global semantic matching of the adversarial visual features. (2) Our adaptive fine-grained alignment leads to fine-grained visual-textual content matching. (3) The joint OT optimization can learn efficient and effective perturbation learning.

4.3. CLIP-based Transfer Attack Comparison

We then compare the CLIP-based transfer attack performance with existing LVLM attacks, *i.e.*, AttackVLM (Zhao et al., 2024), AdvDiffVLM (Guo et al., 2024a), SSA-CWA (Dong et al., 2023), AnyAttack (Zhang et al., 2024), M-Attack (Li et al., 2025), FOA-Attack (Jia et al., 2025), across different closed-source LVLMs in Table 2. It also indicates

that our attack consistently outperforms all baselines by a large margin. Some works manually assigned three semantic keywords to each image and introduced three success thresholds: KMR_1 (at least one matched), KMR_2 (at least two matched), and KMR_3 (all three matched), to evaluate attack transferability under different semantic matching levels. Following their setting, we also compare the proposed method with previous works on 100 randomly selected images. As shown in Table 3, our method still achieves better attack performance. Overall, the results demonstrate that our method has great potential to be extended to CLIP-aware transfer attack with strong scalability.

4.4. Complexity and Robustness

Complexity analysis. As shown in Table 4, we provide the Flops and Time to investigate the complexity and efficiency of our proposed method. It indicates that our attack is competitively efficient as we solely manipulate the visual




Figure 3. The left two figures are ablations with different numbers K of principal components in the projection matrix, separately on the Qwen2.5-VL-3B model (a) and the LLaVa-1.5-7B model (b). The right two figures are distribution visualizations of the global feature alignment between target features and adversarial image features without (c) and with (d) the projection operation.

Table 3. Comparison of CLIP-based transfer attack setting using Keyword Matching Rate (KMR) across different LVLMs and tasks.

Attack	Model	GPT-4o			Gemini-2.0			Claude-3.5		
		KMR ₁	KMR ₂	KMR ₃	KMR ₁	KMR ₂	KMR ₃	KMR ₁	KMR ₂	KMR ₃
AttackVLM (Zhao et al., 2024)	B/16	9.0	4.0	0.0	7.0	2.0	0.0	6.0	3.0	0.0
	B/32	8.0	2.0	0.0	7.0	2.0	0.0	4.0	1.0	0.0
	Laion	7.0	4.0	0.0	7.0	2.0	0.0	5.0	2.0	0.0
AdvDiffVLM (Guo et al., 2024a)	Ensemble	2.0	0.0	0.0	2.0	0.0	0.0	2.0	0.0	0.0
SSA-CWA (Dong et al., 2023)	Ensemble	11.0	6.0	0.0	5.0	2.0	0.0	7.0	3.0	0.0
AnyAttack (Zhang et al., 2024)	Ensemble	44.0	20.0	4.0	46.0	21.0	5.0	25.0	10.0	2.0
M-Attack (Li et al., 2025)	Ensemble	82.0	54.0	13.0	75.0	53.0	11.0	31.0	18.0	3.0
FOA-Attack (Jia et al., 2025)	Ensemble	92.0	76.0	27.0	88.0	69.0	24.0	37.0	23.0	5.0
Ours	Ensemble	96.0	83.0	41.0	94.0	75.0	39.0	64.0	36.0	17.0

Table 4. Efficiency comparison on the COCO dataset. Flops count the computation of forward once.

LVLM Attack	Version	Flops	Time (h)
Clean	None	99.3G	1.3
APGD (Croce & Hein, 2020)	white-box	9.93T	25.5
Ensemble (Schlarmann & Hein, 2023)	white-box	9.93T	41.9
VT-Attack (Wang et al., 2024b)	gray-box	3.04T	65.0
VEAttack (Mei et al., 2025)	gray-box	2.59T	5.3
Ours	gray-box	3.48T	7.2

Table 5. Main ablation on each component on the COCO dataset.

Coarse Alignment	Fine-grained Alignment	Qwen2.5-VL-3B		LLaVa-1.5-7B	
		ASR	AvgSim	ASR	AvgSim
×	×	0.0	0.00	0.0	0.00
✓	×	69.8	0.64	63.7	0.61
×	✓	78.5	0.75	75.2	0.72
✓	✓	86.5	0.82	84.7	0.79




Figure 4. Adversarial robustness against various defenses on the Qwen2.5-VL-3B model (a) and the LLaVa-1.5-7B model (b).

features of the LVLM’s vision encoder. Since VEAttack only considers global alignment, it achieves the lowest Flops and times, but performs much worse than us.

Robustness to defenses. To evaluate the robustness of our attack against potential defense strategies, we conduct experiments on three pre-processing defense methods, *i.e.*, Randomization (Frosio & Kautz, 2023; Xie et al., 2017), JPEG Compression (Guo et al., 2017), and Diffusion Restoration

(Nie et al., 2022) in Figure 4. Compared to previous attacks, our attack is relatively more robust to potential defenses because we explicitly constrain the adversarial visual perturbation to be robust to unknown black-box LLM in the LVLM model, learning more generalizable harmful impacts.

4.5. Ablation Study

Main ablation. To understand the contribution of each component in our proposed attack, we conduct an ablation study in Table 5, where we systematically remove two core modules from our method: coarse alignment, fine-grained alignment. The baseline variant without using any alignment fails to achieve the targeted attack. By adding the coarse alignment to it, the results show noticeable improvement, indicating the importance of aligning coarse-grained features for effective adversarial feature learning. By only adding the fine-grained alignment, the performance leads to a more significant boost, indicating that fine-grained feature alignment is essential for preserving semantic consistency between the adversarial and target features. Lastly, by utilizing both alignments, our attack achieves the best performance.




Figure 5. Attacked averaged similarity comparisons across Qwen2.5-VL-3B and LLaVa-1.5-7B models with varying the local image numbers (a)(b) and local target numbers (c)(d) on the COCO dataset.

Table 6. Ablation on the projected cost of the OT optimization in gray-box targeted setting across different open-source LVLMs.

Variant	Qwen2.5-VL-3B		Qwen2.5-VL-7B		LLaVa-1.5-7B		LLaVa-1.6-7B		Gemma-3-4B		Gemma-3-12B	
	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
OT w/. C^{Π}	86.5	0.82	81.0	0.77	84.7	0.79	78.5	0.75	83.0	0.76	74.2	0.70
OT w/o. C^{Π}	82.4	0.77	76.8	0.72	79.7	0.75	75.1	0.71	80.3	0.72	71.6	0.65

Table 7. Ablation on text encoders on the COCO dataset.

Text Encoder	Qwen2.5-VL-3B		LLaVa-1.5-7B	
	ASR	AvgSim	ASR	AvgSim
CLIP-ViT	86.5	0.82	84.7	0.79
Sentence-Bert	87.3	0.84	85.1	0.80

Ablation of projection process. We study how varying the number of singular vectors K used to construct the projection matrix affects classification accuracy. As shown in Figure 3 (a)(b), we evaluate the similarity within the clean or the adversarial features to investigate the quality of selected vectors. It shows that increasing K steadily improves performance on Qwen2.5-VL-3B and LLaVa-1.5-7B models. Based on this observation, we fix $K = 256$ in all experiments to balance performance and efficiency.

To further evaluate the projection’s impact, we analyze the feature distributions between the image and target. As shown in Figure 3 (c)(d), alignment without projection reduces distribution similarity. Conversely, projecting attacked features onto the target-induced subspace restores this similarity, demonstrating that our method achieves effective global alignment.

Ablation of local numbers. We analyze the attack’s sensitivity to local context density by varying the local image and target numbers. As shown in Figure 5, while increasing these values consistently improves performance, gains become marginal at higher counts. Consequently, we fix $M = 20$ and $N = 30$ to optimally balance attack effectiveness and computational efficiency.

Ablation of projected OT cost. Table 6 presents an ablation study comparing the OT alignment with and without our proposed projection-based cost matrix C^{Π} . The variant without using C^{Π} utilizes the original cost matrix C . Results show that the projection-based OT optimization is more effective than the traditional one, demonstrating that



Figure 6. Visualization results of our attack.

it can be easier to learn better adversarial alignment.

Sensitivity to text encoders. As shown in Table 7, our attack performance is not sensitive to the text encoder.

4.6. Visualization Examples

We provide visualizations of adversarial examples generated by our attack in Figure 6. More results are in Appendix.

5. Conclusion

In this paper, we propose to attack LVLm models in a gray-box setting by solely manipulating perturbations against the vision encoder. We introduce a natural targeted semantic attack by aligning the perturbed visual features with an attacker’s chosen textual semantics. In particular, we devise an SVD-structured alignment with target-guided feature projection to learn better coarse-level semantics, while designing a fine-grained multi-context alignment to match the detailed visual-textual contents. Experiments on both open- and closed-source LVLm models demonstrate our great superiority, effectiveness, and transferability.

Acknowledgements

This work was supported in part by the Fundamental and Interdisciplinary Disciplines Breakthrough Plan of the Ministry of Education of China (No. JYB2025XDXM101), the National Natural Science Foundation of China under Grants 62225113. The Innovative Research Group Project of Hubei Province under Grants 2024AFA017, New Cornerstone Science Foundation through the XPLOER PRIZE. This work was also supported by New Generation Artificial Intelligence-National Science and Technology Major Project (2025ZD0123602), and WHU-Kingsoft Joint Lab.

Impact Statement

This paper proposes a method for gray-box LVLM attacks using targeted multi-modal alignment. The proposed method, like previous adversarial attack methods, investigates adversarial examples in order to identify adversarial vulnerabilities in LVLMs. This effort aims to guide future research into improving LVLMs against adversarial attacks and developing more effective defense approaches. Furthermore, the victim LVLMs employed in this study are open-source models with publicly available weights. The research on adversarial examples will help shape the landscape of AI security.

References

- Arjovsky, M., Chintala, S., and Bottou, L. Wasserstein generative adversarial networks. In *International conference on machine learning*, pp. 214–223. PMLR, 2017.
- Bailey, L., Ong, E., Russell, S., and Emmons, S. Image hi-jacks: Adversarial images can control generative models at runtime. *arXiv preprint arXiv:2309.00236*, 2023.
- Cai, X., Tao, Y., Liu, D., Zhou, P., Qu, X., Dong, J., Tang, K., and Sun, L. Frequency-aware gan for imperceptible transfer attack on 3d point clouds. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pp. 6162–6171, 2024.
- Cai, X., Liu, D., Qu, X., Fang, X., Dong, J., Tang, K., Zhou, P., Sun, L., and Hu, W. Towards building model/prompt-transferable attackers against large vision-language models. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025.
- Chen, G., Yao, W., Song, X., Li, X., Rao, Y., and Zhang, K. Plot: Prompt learning with optimal transport for vision-language models. *arXiv preprint arXiv:2210.01253*, 2022.
- Chen, Y., Rohrbach, M., Yan, Z., Shuicheng, Y., Feng, J., and Kalantidis, Y. Graph-based global reasoning networks. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 433–442, 2019.
- Courty, N., Flamary, R., Tuia, D., and Rakotomamonjy, A. Optimal transport for domain adaptation. *IEEE transactions on pattern analysis and machine intelligence*, 39(9): 1853–1865, 2016.
- Croce, F. and Hein, M. Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks. In *International conference on machine learning*, pp. 2206–2216. PMLR, 2020.
- Cui, X., Aparcedo, A., Jang, Y. K., and Lim, S.-N. On the robustness of large multimodal models against image adversarial attacks. *arXiv preprint arXiv:2312.03777*, 2023.
- Cuturi, M. Sinkhorn distances: Lightspeed computation of optimal transport. *Advances in neural information processing systems*, 26, 2013.
- Dong, Y., Chen, H., Chen, J., Fang, Z., Yang, X., Zhang, Y., Tian, Y., Su, H., and Zhu, J. How robust is google’s bard to adversarial image attacks? *arXiv preprint arXiv:2309.11751*, 2023.
- Fang, X. Advancing out-of-distribution detection across diverse scenarios. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 41042–41043, 2026.
- Fang, X. and Fang, W. Disentangling adversarial prompts: A semantic-graph defense for robust llm security. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 3876–3884, 2026.
- Fang, X. and Hu, Y. Double self-weighted multi-view clustering via adaptive view fusion. *arXiv preprint arXiv:2011.10396*, 2020.
- Fang, X., Hu, Y., Zhou, P., and Wu, D. Animc: A soft approach for autoweighted noisy and incomplete multiview clustering. *IEEE Transactions on Artificial Intelligence*, 3(2):192–206, 2021a.
- Fang, X., Hu, Y., Zhou, P., and Wu, D. O. Unbalanced incomplete multi-view clustering via the scheme of view evolution: Weak views are meat; strong views do eat. *IEEE Transactions on Emerging Topics in Computational Intelligence*, 6(4):913–927, 2021b.
- Fang, X., Hu, Y., Zhou, P., and Wu, D. O. V3h: View variation and view heredity for incomplete multiview clustering. *IEEE Transactions on Artificial Intelligence*, 1(3):233–247, 2021c.

- Fang, X., Liu, D., Zhou, P., and Hu, Y. Multi-modal cross-domain alignment network for video moment retrieval. *IEEE Transactions on Multimedia*, 25:7517–7532, 2022.
- Fang, X., Liu, D., Fang, W., Zhou, P., Cheng, Y., Tang, K., and Zou, K. Annotations are not all you need: A cross-modal knowledge transfer network for unsupervised temporal sentence grounding. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 8721–8733, 2023a.
- Fang, X., Liu, D., Zhou, P., and Nan, G. You can ground earlier than see: An effective and efficient pipeline for temporal sentence grounding in compressed videos. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 2448–2460, 2023b.
- Fang, X., Liu, D., Zhou, P., Xu, Z., and Li, R. Hierarchical local-global transformer for temporal sentence grounding. *IEEE Transactions on Multimedia*, 26:3263–3277, 2023c.
- Fang, X., Easwaran, A., and Genest, B. Uncertainty-guided appearance-motion association network for out-of-distribution action detection. In *2024 IEEE 7th International Conference on Multimedia Information Processing and Retrieval (MIPR)*, pp. 176–182. IEEE, 2024a.
- Fang, X., Fang, W., Liu, D., Qu, X., Dong, J., Zhou, P., Li, R., Xu, Z., Chen, L., Zheng, P., et al. Not all inputs are valid: Towards open-set video moment retrieval using language. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pp. 28–37, 2024b.
- Fang, X., Liu, D., Fang, W., Zhou, P., Xu, Z., Xu, W., Chen, J., and Li, R. Fewer steps, better performance: Efficient cross-modal clip trimming for video moment retrieval using language. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 1735–1743, 2024c.
- Fang, X., Xiong, Z., Fang, W., Qu, X., Chen, C., Dong, J., Tang, K., Zhou, P., Cheng, Y., and Liu, D. Rethinking weakly-supervised video temporal grounding from a game perspective. In *European Conference on Computer Vision*, pp. 290–311. Springer, 2024d.
- Fang, X., Easwaran, A., and Genest, B. Adaptive multi-prompt contrastive network for few-shot out-of-distribution detection. In *International Conference on Machine Learning*, pp. 15887–15900. PMLR, 2025a.
- Fang, X., Easwaran, A., Genest, B., and Suganthan, P. N. Adaptive hierarchical graph cut for multi-granularity out-of-distribution detection. *IEEE Transactions on Artificial Intelligence*, 2025b.
- Fang, X., Easwaran, A., Genest, B., and Suganthan, P. N. Your data is not perfect: Towards cross-domain out-of-distribution detection in class-imbalanced data. *Expert Systems with Applications*, 267:126031, 2025c.
- Fang, X., Fang, W., Ji, W., and Chua, T.-S. Turing patterns for multimedia: Reaction-diffusion multi-modal fusion for language-guided video moment retrieval. In *Proceedings of the 33rd ACM International Conference on Multimedia*, pp. 12509–12518, 2025d.
- Fang, X., Fang, W., and Wang, C. Hierarchical semantic-augmented navigation: Optimal transport and graph-driven reasoning for vision-language navigation. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025e.
- Fang, X., Fang, W., Wang, C., Liu, D., Tang, K., Dong, J., Zhou, P., and Li, B. Multi-pair temporal sentence grounding via multi-thread knowledge transfer network. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 2915–2923, 2025f.
- Fang, X., Fang, W., and Wang, C. Cogniverse: Revolutionizing multi-modal retrieval-augmented generation with cognitive reflection and geometric reasoning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026a.
- Fang, X., Fang, W., and Wang, C. Unveiling the fragility of vision-language models: Multi-modal adversarial synergy via texture-constrained perturbations and cross-modal optimization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 3867–3875, 2026b.
- Fang, X., Fang, W., Wang, C., Qu, X., and Liu, D. Rethinking video-language model from the language input perspective. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 3885–3893, 2026c.
- Fang, X., Fang, W., Wang, C., Tang, K., Liu, D., Wang, S., and Ji, W. Towards unified vision-language models with incomplete multi-modal inputs. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 3858–3866, 2026d.
- Frosio, I. and Kautz, J. The best defense is a good offense: adversarial augmentation against adversarial attacks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4067–4076, 2023.
- Fu, X., Wang, Z., Li, S., Gupta, R. K., Mireshghallah, N., Berg-Kirkpatrick, T., and Fernandes, E. Misusing tools in large language models with visual adversarial examples. *arXiv preprint arXiv:2310.03185*, 2023.

- Gao, K., Bai, Y., Bai, J., Yang, Y., and Xia, S.-T. Adversarial robustness for visual grounding of multimodal large language models. *arXiv preprint arXiv:2405.09981*, 2024a.
- Gao, K., Bai, Y., Gu, J., Xia, S.-T., Torr, P., Li, Z., and Liu, W. Inducing high energy-latency of large vision-language models with verbose images. *arXiv preprint arXiv:2401.11170*, 2024b.
- Goyal, Y., Khot, T., Summers-Stay, D., Batra, D., and Parikh, D. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 6904–6913, 2017.
- Guo, C., Rana, M., Cisse, M., and Van Der Maaten, L. Countering adversarial images using input transformations. *arXiv preprint arXiv:1711.00117*, 2017.
- Guo, D., Tian, L., Zhao, H., Zhou, M., and Zha, H. Adaptive distribution calibration for few-shot learning with hierarchical optimal transport. *Advances in neural information processing systems*, 35:6996–7010, 2022.
- Guo, Q., Pang, S., Jia, X., Liu, Y., and Guo, Q. Efficient generation of targeted and transferable adversarial examples for vision-language models via diffusion models. *IEEE Transactions on Information Forensics and Security*, 2024a.
- Guo, Q., Pang, S., Jia, X., Liu, Y., and Guo, Q. Efficient generation of targeted and transferable adversarial examples for vision-language models via diffusion models. *IEEE Transactions on Information Forensics and Security*, 20: 1333–1348, 2024b.
- Hu, Q., Liu, D., and Hu, W. Exploring the devil in graph spectral domain for 3d point cloud attacks. In *European Conference on Computer Vision*, pp. 229–248. Springer, 2022.
- Jia, X., Gao, S., Qin, S., Pang, T., Du, C., Huang, Y., Li, X., Li, Y., Li, B., and Liu, Y. Adversarial attacks against closed-source mllms via feature optimal alignment. *arXiv preprint arXiv:2505.21494*, 2025.
- Li, J., Li, H., Erfani, S., Feng, L., Bailey, J., and Liu, F. Visual-text cross alignment: Refining the similarity score in vision-language models. *arXiv preprint arXiv:2406.02915*, 2024.
- Li, Z., Zhao, X., Wu, D.-D., Cui, J., and Shen, Z. A frustratingly simple yet highly effective attack baseline: Over 90% success rate against the strong black-box models of gpt-4.5/4o/o1. *arXiv preprint arXiv:2503.10635*, 2025.
- Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., and Zitnick, C. L. Microsoft coco: Common objects in context. In *Computer Vision–ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6–12, 2014, Proceedings, Part V 13*, pp. 740–755. Springer, 2014.
- Liu, D. and Hu, W. Imperceptible transfer attack and defense on 3d point cloud classification. *IEEE transactions on pattern analysis and machine intelligence*, 45(4):4727–4746, 2022.
- Liu, D. and Hu, W. Explicitly perceiving and preserving the local geometric structures for 3d point cloud attack. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 3576–3584, 2024.
- Liu, D. and Hu, W. Imperceptible backdoor attacks on text-guided 3d scene grounding. *IEEE Transactions on Multimedia*, 2025a.
- Liu, D. and Hu, W. Seeing is not believing: Adversarial natural object optimization for hard-label 3d scene attacks. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pp. 11886–11897, 2025b.
- Liu, D., Qu, X., Liu, X.-Y., Dong, J., Zhou, P., and Xu, Z. Jointly cross-and self-modal graph attention network for query-based moment localization. In *Proceedings of the 28th ACM International Conference on Multimedia*, pp. 4070–4078, 2020.
- Liu, D., Qu, X., Dong, J., and Zhou, P. Adaptive proposal generation network for temporal sentence localization in videos. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 9292–9301, 2021a.
- Liu, D., Qu, X., Dong, J., Zhou, P., Cheng, Y., Wei, W., Xu, Z., and Xie, Y. Context-aware biaffine localizing network for temporal sentence grounding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 11235–11244, 2021b.
- Liu, D., Xu, S., Liu, X.-Y., Xu, Z., Wei, W., and Zhou, P. Spatiotemporal graph neural network based mask reconstruction for video object segmentation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 2100–2108, 2021c.
- Liu, D., Qu, X., Di, X., Cheng, Y., Xu, Z., and Zhou, P. Memory-guided semantic learning network for temporal sentence grounding. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pp. 1665–1673, 2022a.

- Liu, D., Qu, X., and Hu, W. Reducing the vision and language bias for temporal sentence grounding. In *Proceedings of the 30th ACM International Conference on Multimedia*, pp. 4092–4101, 2022b.
- Liu, D., Qu, X., Wang, Y., Di, X., Zou, K., Cheng, Y., Xu, Z., and Zhou, P. Unsupervised temporal video grounding with deep semantic clustering. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pp. 1683–1691, 2022c.
- Liu, D., Fang, X., Hu, W., and Zhou, P. Exploring optical-flow-guided motion and detection-based appearance for temporal sentence grounding. *IEEE Transactions on Multimedia*, 25:8539–8553, 2023a.
- Liu, D., Fang, X., Zhou, P., Di, X., Lu, W., and Cheng, Y. Hypotheses tree building for one-shot temporal sentence localization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 1640–1648, 2023b.
- Liu, D., Hu, W., and Li, X. Point cloud attacks in graph spectral domain: When 3d geometry meets graph signal processing. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2023c.
- Liu, D., Hu, W., and Li, X. Robust geometry-dependent attack for 3d point clouds. *IEEE Transactions on Multimedia*, 26:2866–2877, 2023d.
- Liu, D., Liu, Y., Huang, W., and Hu, W. A survey on text-guided 3d visual grounding: Elements, recent advances, and future directions. *arXiv preprint arXiv:2406.05785*, 2024a.
- Liu, D., Yang, M., Qu, X., Zhou, P., Cheng, Y., and Hu, W. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024b.
- Liu, D., Yang, M., Qu, X., Zhou, P., Fang, X., Tang, K., Wan, Y., and Sun, L. Pandora’s box: Towards building universal attackers against real-world large vision-language models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024c.
- Liu, D., Cai, X., Zhou, P., Qu, X., Sun, L., and Hu, W. Are large vision-language models robust to adversarial visual transformations? *IEEE Transactions on Information Forensics and Security*, 2026a.
- Liu, D., Chen, B., and Hu, W. Spatial-spectral homogeneous attacks on physical-world large vision-language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 7114–7122, 2026b.
- Liu, D., Liu, W., Cai, X., Zhou, P., Guan, R., Qu, X., and Du, B. Generating transferable attacks across large vision-language models using adversarial deformation learning. *Pattern Recognition*, pp. 113194, 2026c.
- Liu, H., Li, C., Wu, Q., and Lee, Y. J. Visual instruction tuning. *Advances in neural information processing systems*, 36, 2024d.
- Lu, D., Pang, T., Du, C., Liu, Q., Yang, X., and Lin, M. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Luo, H., Gu, J., Liu, F., and Torr, P. An image is worth 1000 lies: Adversarial transferability across prompts on vision-language models. *arXiv preprint arXiv:2403.09766*, 2024.
- Madry, A., Makelov, A., Schmidt, L., Tsipras, D., and Vladu, A. Towards deep learning models resistant to adversarial attacks. *arXiv preprint arXiv:1706.06083*, 2017.
- Mei, H., Wang, Z., You, S., Dong, M., and Xu, C. Veattack: Downstream-agnostic vision encoder attack against large vision language models. *arXiv preprint arXiv:2505.17440*, 2025.
- Nie, W., Guo, B., Huang, Y., Xiao, C., Vahdat, A., and Anandkumar, A. Diffusion models for adversarial purification. *arXiv preprint arXiv:2205.07460*, 2022.
- Peyré, G. and Cuturi, M. *Computational optimal transport: With applications to data science*. Now Foundations and Trends, 2019.
- Plummer, B. A., Wang, L., Cervantes, C. M., Caicedo, J. C., Hockenmaier, J., and Lazebnik, S. Flickr30k entities: Collecting region-to-phrase correspondences for richer image-to-sentence models. In *Proceedings of the IEEE international conference on computer vision*, pp. 2641–2649, 2015.
- Pratt, S., Covert, I., Liu, R., and Farhadi, A. What does a platypus look like? generating customized prompts for zero-shot image classification. In *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 15691–15701, 2023.
- Roth, K., Kim, J. M., Koepke, A., Vinyals, O., Schmid, C., and Akata, Z. Waffling around for performance: Visual classification with random words and broad concepts. In *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 15746–15757, 2023.
- Schlarman, C. and Hein, M. On the adversarial robustness of multi-modal foundation models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 3677–3685, 2023.

- Singh, A., Natarajan, V., Shah, M., Jiang, Y., Chen, X., Batra, D., Parikh, D., and Rohrbach, M. Towards vqa models that can read. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 8317–8326, 2019.
- Tao, Y., Liu, D., Zhou, P., Xie, Y., Du, W., and Hu, W. 3dhacker: Spectrum-based decision boundary generation for hard-label 3d point cloud attack. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 14340–14350, 2023.
- Team, G., Mesnard, T., Hardin, C., Dadashi, R., Bhupatiraju, S., Pathak, S., Sifre, L., Rivière, M., Kale, M. S., Love, J., et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024.
- Tu, H., Cui, C., Wang, Z., Zhou, Y., Zhao, B., Han, J., Zhou, W., Yao, H., and Xie, C. How many unicorns are in this image? a safety evaluation benchmark for vision llms. *arXiv preprint arXiv:2311.16101*, 2023.
- Wang, P., Bai, S., Tan, S., Wang, S., Fan, Z., Bai, J., Chen, K., Liu, X., Wang, J., Ge, W., et al. Qwen2-vl: Enhancing vision-language model’s perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024a.
- Wang, Y., Liu, C., Qu, Y., Cao, H., Jiang, D., and Xu, L. Break the visual perception: Adversarial attacks targeting encoded visual tokens of large vision-language models. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pp. 1072–1081, 2024b.
- Wang, Z., Han, Z., Chen, S., Xue, F., Ding, Z., Xiao, X., Tresp, V., Torr, P., and Gu, J. Stop reasoning! when multimodal llms with chain-of-thought reasoning meets adversarial images. *arXiv preprint arXiv:2402.14899*, 2024c.
- Xie, C., Wang, J., Zhang, Z., Ren, Z., and Yuille, A. Mitigating adversarial effects through randomization. *arXiv preprint arXiv:1711.01991*, 2017.
- Xu, H., Luo, D., and Carin, L. Scalable gromov-wasserstein learning for graph partitioning and matching. *Advances in neural information processing systems*, 32, 2019.
- Yan, H., Ma, H., Cai, X., Liu, D., Yuan, Z., Qu, X., Dong, J., Guan, R., Fang, X., He, H., et al. Fit the distribution: Cross-image/prompt adversarial attacks on multimodal large language models. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025.
- Yang, M., Liu, D., Tang, K., Zhou, P., Chen, L., and Chen, J. Hiding imperceptible noise in curvature-aware patches for 3d point cloud attack. In *European Conference on Computer Vision*, pp. 431–448. Springer, 2024.
- Yin, Z., Ye, M., Zhang, T., Du, T., Zhu, J., Liu, H., Chen, J., Wang, T., and Ma, F. Vlattack: Multimodal adversarial attacks on vision-language tasks via pre-trained models. *arXiv preprint arXiv:2310.04655*, 2023.
- Zhang, J., Ye, J., Ma, X., Li, Y., Yang, Y., Sang, J., and Yeung, D.-Y. Anyattack: Towards large-scale self-supervised generation of targeted adversarial examples for vision-language models. *arXiv e-prints*, pp. arXiv–2410, 2024.
- Zhao, Y., Pang, T., Du, C., Yang, X., Li, C., Cheung, N.-M. M., and Lin, M. On evaluating adversarial robustness of large vision-language models. *Advances in Neural Information Processing Systems*, 36, 2024.
- Zhu, B., Cui, J., Zhang, H., and Zhang, C. Project-probe-aggregate: Efficient fine-tuning for group robustness. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pp. 25487–25496, 2025a.
- Zhu, X., Wang, S., Lu, J., Hao, Y., Liu, H., and He, X. Boosting few-shot learning via attentive feature regularization. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, pp. 7793–7801, 2024a.
- Zhu, X., Zhu, B., Wang, S., Zhao, K., and Zhang, H. Enhancing clip robustness via cross-modality alignment. *arXiv preprint arXiv:2510.24038*, 2025b.
- Zhu, Y., Ji, Y., Zhao, Z., Wu, G., and Wang, L. Awt: Transferring vision-language models via augmentation, weighting, and transportation. *Advances in Neural Information Processing Systems*, 37:25561–25591, 2024b.

A. SVD-Structured Feature Projection Leads to Better Semantic Alignment

To ease the representation, we simplify the vision encoder f_θ and directly define the clean image feature as $\mathbf{x}$ and $\|\mathbf{x}\| = 1$. The adversarial image feature can be denoted as:

$$\mathbf{x}^{adv} = \mathbf{x} + \delta, \quad \|\delta\| \leq \epsilon. \quad (16)$$

We denote the semantic guidance feature as $\mathbf{y}$, and denote the projection of the adversarial feature onto the desired feature space $\mathcal{U}$ provided by $\mathbf{y}$ as $\Pi(\mathbf{x}^{adv})$. Since $\mathbf{y} \in \mathcal{U}$, we can decompose the perturbation into parallel and orthogonal parts of $\mathcal{U}$ as:

$$\delta = \delta_{\parallel} + \delta_{\perp}, \quad \delta_{\parallel} \in \mathcal{U}, \delta_{\perp} \perp \mathcal{U}. \quad (17)$$

Then we can project the perturbed image feature onto the feature space $\mathcal{U}$ as:

$$\Pi(\mathbf{x}^{adv}) = \Pi(\mathbf{x} + \delta) = \mathbf{x} + \delta_{\parallel}. \quad (18)$$

To investigate how this projection affects the global semantic matching during the alignment process, we compute the cosine similarity between the unprojected image features $\mathbf{x}^{adv}$ and the semantic feature guidance $\mathbf{y}$ as:

$$\cos(\mathbf{x}^{adv}, \mathbf{y}) = \frac{(\mathbf{x} + \delta)^\top \mathbf{y}}{\|\mathbf{x}^{adv}\| \cdot \|\mathbf{y}\|}. \quad (19)$$

The cosine similarity between the projected image feature $\Pi(\mathbf{x}^{adv})$ and the semantic feature guidance $\mathbf{y}$ can also be formulated as:

$$\cos(\Pi(\mathbf{x}^{adv}), \mathbf{y}) = \frac{(\mathbf{x} + \delta_{\parallel})^\top \mathbf{y}}{\|\Pi(\mathbf{x}^{adv})\| \cdot \|\mathbf{y}\|}. \quad (20)$$

Comparing the numerators in (19) and (20), since $\mathbf{y} \in \mathcal{U}$ and $\Pi(\mathbf{x}^{adv})$ is the orthogonal projection of $\mathbf{x}^{adv}$ onto the space $\mathcal{U}$, the $\delta_{\perp}$ and $\mathbf{x}^{adv} - \Pi(\mathbf{x}^{adv})$ are orthogonal to every vector in $\mathcal{U}$, including $\mathbf{y}$, and therefore, we have:

$$(\mathbf{x} + \delta)^\top \mathbf{y} = (\mathbf{x} + \delta_{\parallel})^\top \mathbf{y}. \quad (21)$$

As for the norm relationship in denominators, by Pythagoras, because $\mathbf{x}^{adv} = \Pi(\mathbf{x}^{adv}) + (\mathbf{x}^{adv} - \Pi(\mathbf{x}^{adv}))$ with the last term orthogonal to $\mathcal{U}$, we have

$$\|\mathbf{x}^{adv}\|^2 = \|\Pi(\mathbf{x}^{adv})\|^2 + \|\mathbf{x}^{adv} - \Pi(\mathbf{x}^{adv})\|^2 \geq \|\Pi(\mathbf{x}^{adv})\|^2, \quad (22)$$

and hence

$$\|\mathbf{x}^{adv}\| \geq \|\Pi(\mathbf{x}^{adv})\|, \quad (23)$$

with equality iff $\mathbf{x}^{adv} - \Pi(\mathbf{x}^{adv}) = 0$, i.e., iff $\mathbf{x}^{adv} \in \mathcal{U}$ (no orthogonal component).

Therefore, using the equality of numerators and the inequality of norms,

$$\begin{aligned} \cos(\Pi(\mathbf{x}^{adv}), \mathbf{y}) &= \frac{(\mathbf{x} + \delta_{\parallel})^\top \mathbf{y}}{\|\Pi(\mathbf{x}^{adv})\| \cdot \|\mathbf{y}\|} = \frac{(\mathbf{x} + \delta)^\top \mathbf{y}}{\|\Pi(\mathbf{x}^{adv})\| \cdot \|\mathbf{y}\|} \\ &= \underbrace{\frac{\|\mathbf{x}^{adv}\|}{\|\Pi(\mathbf{x}^{adv})\|}}_{\geq 1} \cdot \frac{(\mathbf{x} + \delta_{\parallel})^\top \mathbf{y}}{\|\mathbf{x}^{adv}\| \cdot \|\mathbf{y}\|} = \frac{\|\mathbf{x}^{adv}\|}{\|\Pi(\mathbf{x}^{adv})\|} \cdot \cos(\mathbf{x}^{adv}, \mathbf{y}). \end{aligned} \quad (24)$$

Since $\frac{\|\mathbf{x}^{adv}\|}{\|\Pi(\mathbf{x}^{adv})\|} \geq 1$, it follows that

$$\cos(\Pi(\mathbf{x}^{adv}), \mathbf{y}) \geq \cos(\mathbf{x}^{adv}, \mathbf{y}), \quad (25)$$

where equality holds iff $\|\mathbf{x}^{adv}\| = \|\Pi(\mathbf{x}^{adv})\|$, which is equivalent to $\mathbf{x}^{adv} = \Pi(\mathbf{x}^{adv})$, i.e., the perturbation has no orthogonal component ($\delta_{\perp} = 0$) so that $\mathbf{x}^{adv} \in \mathcal{U}$.

This demonstrates that our employed projection preserves dot-products with any vector in $\mathcal{U}$ while shrinking (or not increasing) the denominator, thereby increasing cosine similarity, leading to better global semantic alignment.

B. Proof of Amplified Effectiveness of OT

Given the distributions $\mathbb{X}(\mathbf{x}^{adv})$, $\mathbb{Y}(\mathbf{y}^{tar})$, the alignment between the adversarial image and target semantic is measured by the optimal transport distance, which captures the minimal semantic matching cost between image and text features:

$$\text{OT}(\mathbb{X}(\mathbf{x}^{adv}), \mathbb{Y}(\mathbf{y}^{tar}); \mathbf{C}) = \min_{\mathbf{T} \geq 0} \langle \mathbf{T}, \mathbf{C}^\Pi \rangle, \quad \text{s.t.} \quad \mathbf{T} \mathbf{1}_N = \boldsymbol{\alpha}, \quad \mathbf{T}^\top \mathbf{1}_M = \boldsymbol{\beta}, \quad (26)$$

where $\mathbf{1}_N, \mathbf{1}_M$ are all-ones vectors and $\boldsymbol{\alpha} = [\alpha_1, \dots, \alpha_M]^\top, \boldsymbol{\beta} = [\beta_1, \dots, \beta_N]^\top$. $\mathbf{C}^\Pi \in \mathbb{R}^{M \times N}$ denotes the transportation cost between the M augmented image features projected on the global target’s semantic plane $\mathcal{U}$ and the N target text descriptions, which is usually quantified using the cosine similarity as:

$$\mathbf{C}^\Pi(m, n) = 1 - \cos(\Pi(\tilde{\mathbf{x}}_m^{adv}), \tilde{\mathbf{y}}_n^{tar}). \quad (27)$$

According the Appendix A, *i.e.*, $\cos(\Pi(\mathbf{x}^{adv}), \mathbf{y}) \geq \cos(\mathbf{x}^{adv}, \mathbf{y})$, we can obtain that ce, the projected cost matrix entry is smaller:

$$\mathbf{C}^\Pi(m, n) = 1 - \cos(\Pi(\tilde{\mathbf{x}}_m^{adv}), \tilde{\mathbf{y}}_n^{tar}) \leq 1 - \cos(\tilde{\mathbf{x}}_m^{adv}, \tilde{\mathbf{y}}_n^{tar}) = \mathbf{C}(m, n). \quad (28)$$

Therefore, for any feasible transport plan $\mathbf{T}$, we have $\langle \mathbf{T}, \mathbf{C}^\Pi \rangle \leq \langle \mathbf{T}, \mathbf{C} \rangle$. Let $\mathbf{T}^* = \arg \min \langle \mathbf{T}, \mathbf{C} \rangle$, then:

$$\text{OT}(\mathbb{X}(\mathbf{x}^{adv}), \mathbb{Y}(\mathbf{y}^{tar}); \mathbf{C}^\Pi) \leq \langle \mathbf{T}^*, \mathbf{C}^\Pi \rangle \leq \langle \mathbf{T}^*, \mathbf{C} \rangle = \text{OT}(\mathbb{X}(\mathbf{x}^{adv}), \mathbb{Y}(\mathbf{y}^{tar}); \mathbf{C}). \quad (29)$$

C. A Detailed Description of Our Proposed Attack

Our work attacks the LVLM models in a gray-box setting, where the attacker can solely access the vision encoder. Given a target semantic, the attacker aims to perturb the image input. In particular, we propose to perturb the image from both global and local perspectives by designing corresponding feature alignment with the target semantics. Specifically, we develop a global feature projection process to fit the adversarial image feature into the latent space of the target semantics for better alignment, while we extract both visual and textual fine-grained semantic content for local feature matching. The detailed description of our proposed attack is shown in Algorithm 1.

Algorithm 1 Pipeline of Our Proposed Attack

Input: clean image $\mathbf{x}$, LVLM vision encoder f_{θ_v} , text transformer f_{θ_t} , attacker-chosen target semantic $\mathbf{y}^{tar}$.

Output: adversarial image $\mathbf{x}^{adv}$.

- 1: **Initialize adversarial perturbation:** $\delta \leftarrow \mathbf{0}$ and $\mathbf{x}^{adv} \leftarrow \mathbf{x}$
- 2: **Enrich target semantics via LLM:** Generate fine-grained textual descriptions $\{\tilde{\mathbf{y}}_n^{tar}\}_{n=1}^N \leftarrow \text{LLM}(\mathbf{y}^{tar})$
- 3: **Global feature projection:**
- 4: SVD: $\mathbf{Y} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$, take top- K basis $\mathbf{U}_K \leftarrow \mathbf{U}_{[:,0:K-1]}$
- 5: Define projection operator $\Pi(\cdot) := \mathbf{U}_K \mathbf{U}_K^\top(\cdot)$
- 6: Compute current image feature $\mathbf{z} \leftarrow f_{\theta_v}(\mathbf{x}^{adv})$
- 7: Compute projected global feature $\mathbf{z}^\Pi \leftarrow \Pi(\mathbf{z})$
- 8: **Local semantic distribution modeling:**
- 9: Sample / generate M augmented views $\{\tilde{\mathbf{x}}_m^{adv}\}_{m=1}^M$
- 10: Compute visual local features $\mathbf{v}_m \leftarrow f_{\theta_v}(\tilde{\mathbf{x}}_m^{adv})$ for $m = 1..M$
- 11: Compute textual fine-grained features $\mathbf{t}_n \leftarrow f_{\theta_t}(\tilde{\mathbf{y}}_n^{tar})$ for $n = 1..N$
- 12: Compute adaptive view importance α_m (via entropy over $p(\mathbf{t}_n|\mathbf{v}_m)$) and text importance β_n (via entropy over $p(\mathbf{v}_m|\mathbf{t}_n)$)
- 13: **Build OT cost with projection:**
- 14: For each pair (m, n) compute cost

$$\mathbf{C}_{m,n}^\Pi \leftarrow 1 - \cos(\Pi(\mathbf{v}_m), \mathbf{t}_n)$$

- 15: Solve OT (*e.g.*, Sinkhorn) to obtain transport plan $\mathbf{T}$ subject to marginals $\boldsymbol{\alpha}, \boldsymbol{\beta}$
 - 16: Two loss items: $\mathcal{L}_g \leftarrow 1 - \cos(\mathbf{z}^\Pi, \mathbf{y}^{tar})$, $\mathcal{L}_{ot} \leftarrow \langle \mathbf{T}, \mathbf{C}^\Pi \rangle$
 - 17: **Update perturbation via PGD on the vision encoder:**
 - 18: Compute gradient $\nabla_{\mathbf{x}^{adv}} \mathcal{L}$ by backprop through f_{θ_v}
 - 19: Gradient step on perturbation: $\delta \leftarrow \delta - \eta \cdot \text{sign}(\nabla_{\mathbf{x}^{adv}} \mathcal{L})$
 - 20: Project δ to l_∞ ball: $\delta \leftarrow \text{clip}(\delta, -\epsilon, \epsilon)$
 - 21: **Return** $\mathbf{x}^{adv}$
-

Table 8. Comparison of gray-box targeted setting (threshold = 0.3). Attack success is counted when similarity > 0.3.

Task	Attack	Qwen2.5-VL-3B		Qwen2.5-VL-7B		LLaVa-1.5-7B		LLaVa-1.6-7B		Gemma-3-4B		Gemma-3-12B	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
COCO	MIX-Attack	81.7	0.61	86.2	0.67	82.5	0.63	85.0	0.66	80.2	0.60	86.9	0.66
	AttackVLM	88.9	0.70	82.1	0.65	87.6	0.68	80.3	0.63	87.4	0.69	82.3	0.65
	VT-Attack	83.5	0.66	85.6	0.68	80.4	0.61	86.8	0.67	83.9	0.65	80.7	0.62
	VEAttack	86.1	0.69	80.7	0.63	86.8	0.67	85.2	0.66	86.0	0.66	81.3	0.64
	Ours	95.3	0.82	92.8	0.77	94.1	0.79	90.6	0.75	93.5	0.76	89.0	0.70
Flickr30k	MIX-Attack	86.1	0.65	82.4	0.61	88.0	0.67	84.7	0.63	79.8	0.60	85.5	0.65
	AttackVLM	89.7	0.68	87.3	0.67	84.5	0.63	89.2	0.68	88.9	0.67	82.6	0.61
	VT-Attack	81.9	0.62	86.7	0.66	86.4	0.65	80.9	0.59	88.1	0.67	83.2	0.63
	VEAttack	87.3	0.67	81.0	0.60	90.2	0.69	85.8	0.65	82.7	0.63	88.0	0.66
	Ours	96.0	0.80	93.7	0.79	94.5	0.77	91.8	0.74	95.1	0.78	90.2	0.72
TextVQA	MIX-Attack	80.2	0.62	83.1	0.64	86.0	0.66	78.5	0.59	85.0	0.65	79.8	0.61
	AttackVLM	88.1	0.67	81.5	0.62	89.2	0.68	85.8	0.65	88.6	0.66	78.9	0.60
	VT-Attack	82.7	0.64	88.0	0.67	83.6	0.63	87.3	0.66	81.5	0.62	86.5	0.64
	VEAttack	89.5	0.69	82.3	0.62	88.0	0.67	87.0	0.66	83.8	0.63	90.2	0.67
	Ours	94.8	0.78	92.9	0.75	91.2	0.74	89.6	0.71	93.9	0.75	88.4	0.70
VQAv2	MIX-Attack	82.9	0.63	81.0	0.60	83.8	0.62	82.1	0.61	88.2	0.66	79.9	0.60
	AttackVLM	89.4	0.67	86.2	0.64	91.0	0.68	83.9	0.62	88.8	0.65	84.1	0.62
	VT-Attack	85.0	0.65	88.1	0.66	82.5	0.61	87.0	0.65	84.3	0.64	86.9	0.63
	VEAttack	90.1	0.69	82.7	0.63	89.6	0.67	84.5	0.62	90.8	0.67	81.5	0.61
	Ours	95.5	0.79	93.1	0.74	92.0	0.73	90.2	0.70	94.0	0.75	88.9	0.69

Table 9. Comparison of gray-box targeted setting (threshold = 0.7). Attack success is counted when similarity > 0.7.

Task	Attack	Qwen2.5-VL-3B		Qwen2.5-VL-7B		LLaVa-1.5-7B		LLaVa-1.6-7B		Gemma-3-4B		Gemma-3-12B	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
COCO	MIX-Attack	42.3	0.61	49.0	0.67	46.2	0.63	50.5	0.66	41.8	0.60	48.1	0.66
	AttackVLM	55.0	0.70	43.5	0.65	57.1	0.68	40.2	0.63	56.7	0.69	44.9	0.65
	VT-Attack	47.8	0.66	52.6	0.68	40.9	0.61	51.5	0.67	50.3	0.65	42.0	0.62
	VEAttack	53.5	0.69	41.2	0.63	54.8	0.67	50.1	0.66	49.9	0.66	45.2	0.64
	Ours	68.0	0.82	61.5	0.77	64.2	0.79	58.7	0.75	62.8	0.76	53.6	0.70
Flickr30k	MIX-Attack	49.7	0.65	44.2	0.61	52.8	0.67	47.1	0.63	43.9	0.60	50.3	0.65
	AttackVLM	57.4	0.68	53.9	0.67	46.5	0.63	59.0	0.68	56.3	0.67	46.8	0.61
	VT-Attack	46.2	0.62	55.8	0.66	53.6	0.65	43.0	0.59	58.4	0.67	50.1	0.63
	VEAttack	54.0	0.67	42.5	0.60	60.1	0.69	50.2	0.65	48.8	0.63	55.7	0.66
	Ours	71.4	0.80	66.8	0.79	68.5	0.77	60.9	0.74	69.2	0.78	58.0	0.72
TextVQA	MIX-Attack	44.6	0.62	48.2	0.64	50.3	0.66	42.1	0.59	51.9	0.65	43.2	0.61
	AttackVLM	58.9	0.67	45.6	0.62	60.5	0.68	53.2	0.65	61.0	0.66	41.0	0.60
	VT-Attack	50.4	0.64	60.1	0.67	48.0	0.63	59.6	0.66	46.8	0.62	57.3	0.64
	VEAttack	61.0	0.69	46.0	0.62	57.6	0.67	58.4	0.66	47.2	0.63	62.1	0.67
	Ours	73.8	0.78	70.1	0.75	66.4	0.74	62.0	0.71	71.0	0.75	58.9	0.70
VQAv2	MIX-Attack	46.1	0.63	43.8	0.60	48.0	0.62	47.0	0.61	55.6	0.66	42.5	0.60
	AttackVLM	59.8	0.67	54.2	0.64	61.0	0.68	49.7	0.62	58.6	0.65	47.0	0.62
	VT-Attack	51.2	0.65	60.3	0.66	46.5	0.61	58.2	0.65	50.8	0.64	56.0	0.63
	VEAttack	62.4	0.69	44.9	0.63	59.1	0.67	48.0	0.62	60.9	0.67	45.1	0.61
	Ours	72.6	0.79	68.3	0.74	65.9	0.73	62.5	0.70	70.2	0.75	60.1	0.69

D. More Comparison Results under Varied Thresholds

We further evaluate the performance of our proposed on both gray-box target attack and CLIP-based transfer attack settings at the threshold of 0.3. As shown in Table 8, our attack consistently achieves superior adversarial success rates (ASR) and average semantic similarity (AvgSim) on open-source LVLMs, such as 94.1% ASR and 0.79 AvgSim on LLaVA-1.6-7B and the COCO dataset, significantly outperforming baseline ensemble attacks. Similarly, Table 10 highlights our attack’s strong transferability to closed-source models under the 0.3 threshold, achieving notably high performance (e.g., 98.7% ASR and 0.69 AvgSim on GPT-4.1), confirming its effectiveness and semantic alignment across diverse evaluation scenarios.

Table 10. Comparison of CLIP-based transfer attack setting (threshold = 0.3). Attack success is counted when similarity > 0.3.

Attack	Model	Claude-3.5		Claude-3.7		GPT-4o		GPT-4.1		Gemini-2.0	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
AttackVLM (Zhao et al., 2024)	B/16	2.4	0.02	4.1	0.03	40.8	0.21	42.6	0.22	23.5	0.12
	B/32	14.8	0.08	20.5	0.11	20.1	0.10	21.9	0.11	9.9	0.06
	Laion	3.5	0.02	4.9	0.03	69.9	0.38	71.8	0.39	55.8	0.30
AdvDiffVLM (Guo et al., 2024a)	Ensemble	1.1	0.01	1.4	0.01	3.2	0.01	2.9	0.01	2.0	0.01
SSA-CWA (Dong et al., 2023)	Ensemble	3.2	0.02	3.7	0.03	3.8	0.03	3.0	0.02	4.0	0.02
AnyAttack (Zhang et al., 2024)	Ensemble	19.1	0.09	18.7	0.08	40.8	0.15	39.5	0.13	31.1	0.12
M-Attack (Li et al., 2025)	Ensemble	17.9	0.10	23.8	0.12	86.8	0.50	89.1	0.51	75.5	0.41
FOA-Attack (Jia et al., 2025)	Ensemble	28.4	0.16	36.4	0.18	94.8	0.59	95.6	0.62	86.7	0.50
Ours	Ensemble	39.7	0.25	45.1	0.25	98.3	0.63	98.7	0.69	91.0	0.62

Table 11. Comparison of CLIP-based transfer attack setting (threshold = 0.7). Attack success is counted when similarity > 0.7.

Attack	Model	Claude-3.5		Claude-3.7		GPT-4o		GPT-4.1		Gemini-2.0	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
AttackVLM (Zhao et al., 2024)	B/16	0.0	0.02	0.1	0.03	7.8	0.21	8.2	0.22	3.4	0.12
	B/32	2.4	0.08	3.3	0.11	3.0	0.10	3.0	0.11	0.9	0.06
	Laion	0.2	0.02	0.7	0.03	25.5	0.38	26.0	0.39	15.9	0.30
AdvDiffVLM (Guo et al., 2024a)	Ensemble	0.1	0.01	0.2	0.01	0.5	0.01	0.4	0.01	0.2	0.01
SSA-CWA (Dong et al., 2023)	Ensemble	0.1	0.02	0.0	0.03	0.4	0.03	0.2	0.02	0.1	0.02
AnyAttack (Zhang et al., 2024)	Ensemble	1.5	0.09	1.3	0.08	1.8	0.15	1.7	0.13	0.8	0.12
M-Attack (Li et al., 2025)	Ensemble	3.3	0.10	4.4	0.12	38.8	0.50	39.8	0.51	26.6	0.41
FOA-Attack (Jia et al., 2025)	Ensemble	6.3	0.16	9.6	0.18	57.9	0.59	58.9	0.62	41.5	0.50
Ours	Ensemble	15.4	0.25	17.8	0.25	60.2	0.63	65.6	0.69	52.7	0.62

Continuing with the threshold set to 0.7, Table 9 shows that our proposed attack maintains its lead on gray-box target attack among open-source LVLMs, achieving significantly higher ASR and AvgSim, such as 58.7% ASR and 0.75 AvgSim on LLaVA-1.6-7B, notably surpassing all baseline ensemble methods. Similarly, results in Table 11 indicate that our proposed attack retains effectiveness on CLIP-based transfer attack against challenging closed-source models even at the higher threshold, notably achieving 65.6% ASR and 0.69 AvgSim on GPT-4.1, reinforcing its strong adversarial transferability and semantic alignment in stringent attack scenarios.

E. Defense on CLIP-based Transfer Attack

Following previous works (Li et al., 2025; Jia et al., 2025), we also evaluate the attack performance in the CLIP-based transfer attack setting against a series of defense methods, including smoothing-based defenses (Gaussian, Medium, and Average), JPEG compression, and Comdefend. The experimental results are shown in Table 12. Across all defenses, our attack consistently outperforms M-Attack in both ASR and AvgSim. Under Comdefend, our attack achieves 68.3% ASR on GPT-4o and 70.1% on GPT-4.1. Even under JPEG, our attack maintains over 60% ASR with stable AvgSim values. These results indicate that our proposed attack achieves superior adversarial transferability and resilience across diverse defense strategies.

F. Detailed Evaluation Prompt

Following M-Attack (Li et al., 2025), we adopt the same way to evaluate the adversarial performance. Below is the detailed evaluation prompt used to assess semantic similarity between textual inputs: **ASR:** the “{input_text_1}” and “{input_text_2}” are used as placeholders for text inputs. The evaluation prompt template is shown in Figure 7.

G. More Visualization Results

To further evaluate the effectiveness and transferability of our proposed attack, we conducted real-world testing against commercial closed-source LVLMs. These models include GPT-4o, GPT-4.1, GPT-5.1, Claude-Sonnet-3.5, Claude-Sonnet-

Table 12. Attack performance of CLIP-based transfer attack setting with different LVLM attacks after defense processing.

Attack	Model	Claude-3.5		Claude-3.7		GPT-4o		GPT-4.1		Gemini-2.0	
		ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim	ASR	AvgSim
Gaussian	M-Attack (Li et al., 2025)	2.0	0.04	5.0	0.06	57.0	0.45	53.0	0.44	29.0	0.29
	FOA-Attack (Jia et al., 2025)	3.0	0.06	6.0	0.07	72.0	0.57	71.0	0.57	50.0	0.42
	Ours	6.5	0.09	9.3	0.11	74.8	0.60	72.5	0.58	55.2	0.46
Medium	M-Attack (Li et al., 2025)	3.0	0.04	4.0	0.06	39.0	0.37	40.0	0.38	23.0	0.24
	FOA-Attack (Jia et al., 2025)	4.0	0.07	6.0	0.09	59.0	0.48	63.0	0.50	41.0	0.37
	Ours	5.4	0.08	7.5	0.10	66.2	0.54	68.1	0.55	46.8	0.41
Average	M-Attack (Li et al., 2025)	2.0	0.04	1.0	0.03	38.0	0.37	39.0	0.36	19.0	0.22
	FOA-Attack (Jia et al., 2025)	5.0	0.06	3.0	0.06	59.0	0.48	62.0	0.50	36.0	0.34
	Ours	6.0	0.09	5.5	0.08	65.7	0.53	67.9	0.54	44.1	0.40
JPEG	M-Attack (Li et al., 2025)	9.0	0.12	14.0	0.17	60.0	0.48	52.0	0.45	36.0	0.35
	FOA-Attack (Jia et al., 2025)	14.0	0.20	22.0	0.24	75.0	0.59	78.0	0.59	58.0	0.49
	Ours	17.8	0.24	21.3	0.25	77.1	0.61	79.4	0.63	61.5	0.52
Comdefend	M-Attack (Li et al., 2025)	2.0	0.04	5.0	0.08	35.0	0.35	37.0	0.37	22.0	0.25
	FOA-Attack (Jia et al., 2025)	6.0	0.07	11.0	0.15	61.0	0.49	63.0	0.51	38.0	0.39
	Ours	8.2	0.11	12.6	0.16	68.3	0.56	70.1	0.58	50.4	0.44

Evaluation Prompt

Rate the semantic similarity between the following two texts on a scale from 0 to 1.

****Criteria for similarity measurement:****

- **Main Subject Consistency:**** If both descriptions refer to the same key subject or object (e.g., a person, food, an event), they should receive a higher similarity score.
- **Relevant Description:**** If the descriptions are related to the same context or topic, they should also contribute to a higher similarity score.
- **Ignore Fine-Grained Details:**** Do not penalize differences in ****phrasing, sentence structure, or minor variations in detail****. Focus on ****whether both descriptions fundamentally describe the same thing.****
- **Partial Matches:**** If one description contains extra information but does not contradict the other, they should still have a high similarity score.
- **Similarity Score Range:****
 - **1.0:**** Nearly identical in meaning.
 - **0.8-0.9:**** Same subject, with highly related descriptions.
 - **0.7-0.8:**** Same subject, core meaning aligned, even if some details differ.
 - **0.5-0.7:**** Same subject but different perspectives or missing details.
 - **0.3-0.5:**** Related but not highly similar (same general theme but different descriptions).
 - **0.0-0.2:**** Completely different subjects or unrelated meanings.

Text 1: {input_text1}

Text 2: {input_text2}

Output only a single number between 0 and 1. Do not include any explanation or additional text.

Figure 7. Evaluation prompt template.

3.7, Gemini-2.0-Flash, and Gemini-2.5-Flash. Specifically, Figure 8 to 14 correspond to the attack results on each of these models in order: Figure 8 shows GPT-4o, Figure 9 shows GPT-4.1, Figure 10 shows GPT-5.1, Figure 11 shows Claude-Sonnet-3.5, Figure 12 shows Claude-Sonnet-3.7, Figure 13 shows Gemini-2.0-Flash, and Figure 14 shows Gemini-2.5-Flash. Our results show that the adversarial examples successfully manipulate these advanced models into generating descriptions aligned with our target semantics, thereby demonstrating the high transferability of our proposed attack across diverse architectures.



Figure 8. Example responses from the commercial GPT-4o to targeted attacks generated by our method.



Figure 9. Example responses from the commercial GPT-4.1 to targeted attacks generated by our method.




Figure 10. Example responses from the commercial GPT-5.1 to targeted attacks generated by our method.




Figure 11. Example responses from the commercial Claude-Sonnet-3.5 to targeted attacks generated by our method.




Figure 12. Example responses from the commercial Claude-Sonnet-3.7 to targeted attacks generated by our method.




Figure 13. Example responses from the commercial Gemini-2.0-Flash to targeted attacks generated by our method.




Figure 14. Example responses from the commercial Gemini-2.5-Flash to targeted attacks generated by our method.